

# Modern Blockchains through the Lens of Network Security

Alberto Sonnino

# Byzantine Fault Tolerance



# Byzantine Fault Tolerance



$> 2/3$





1. make transaction



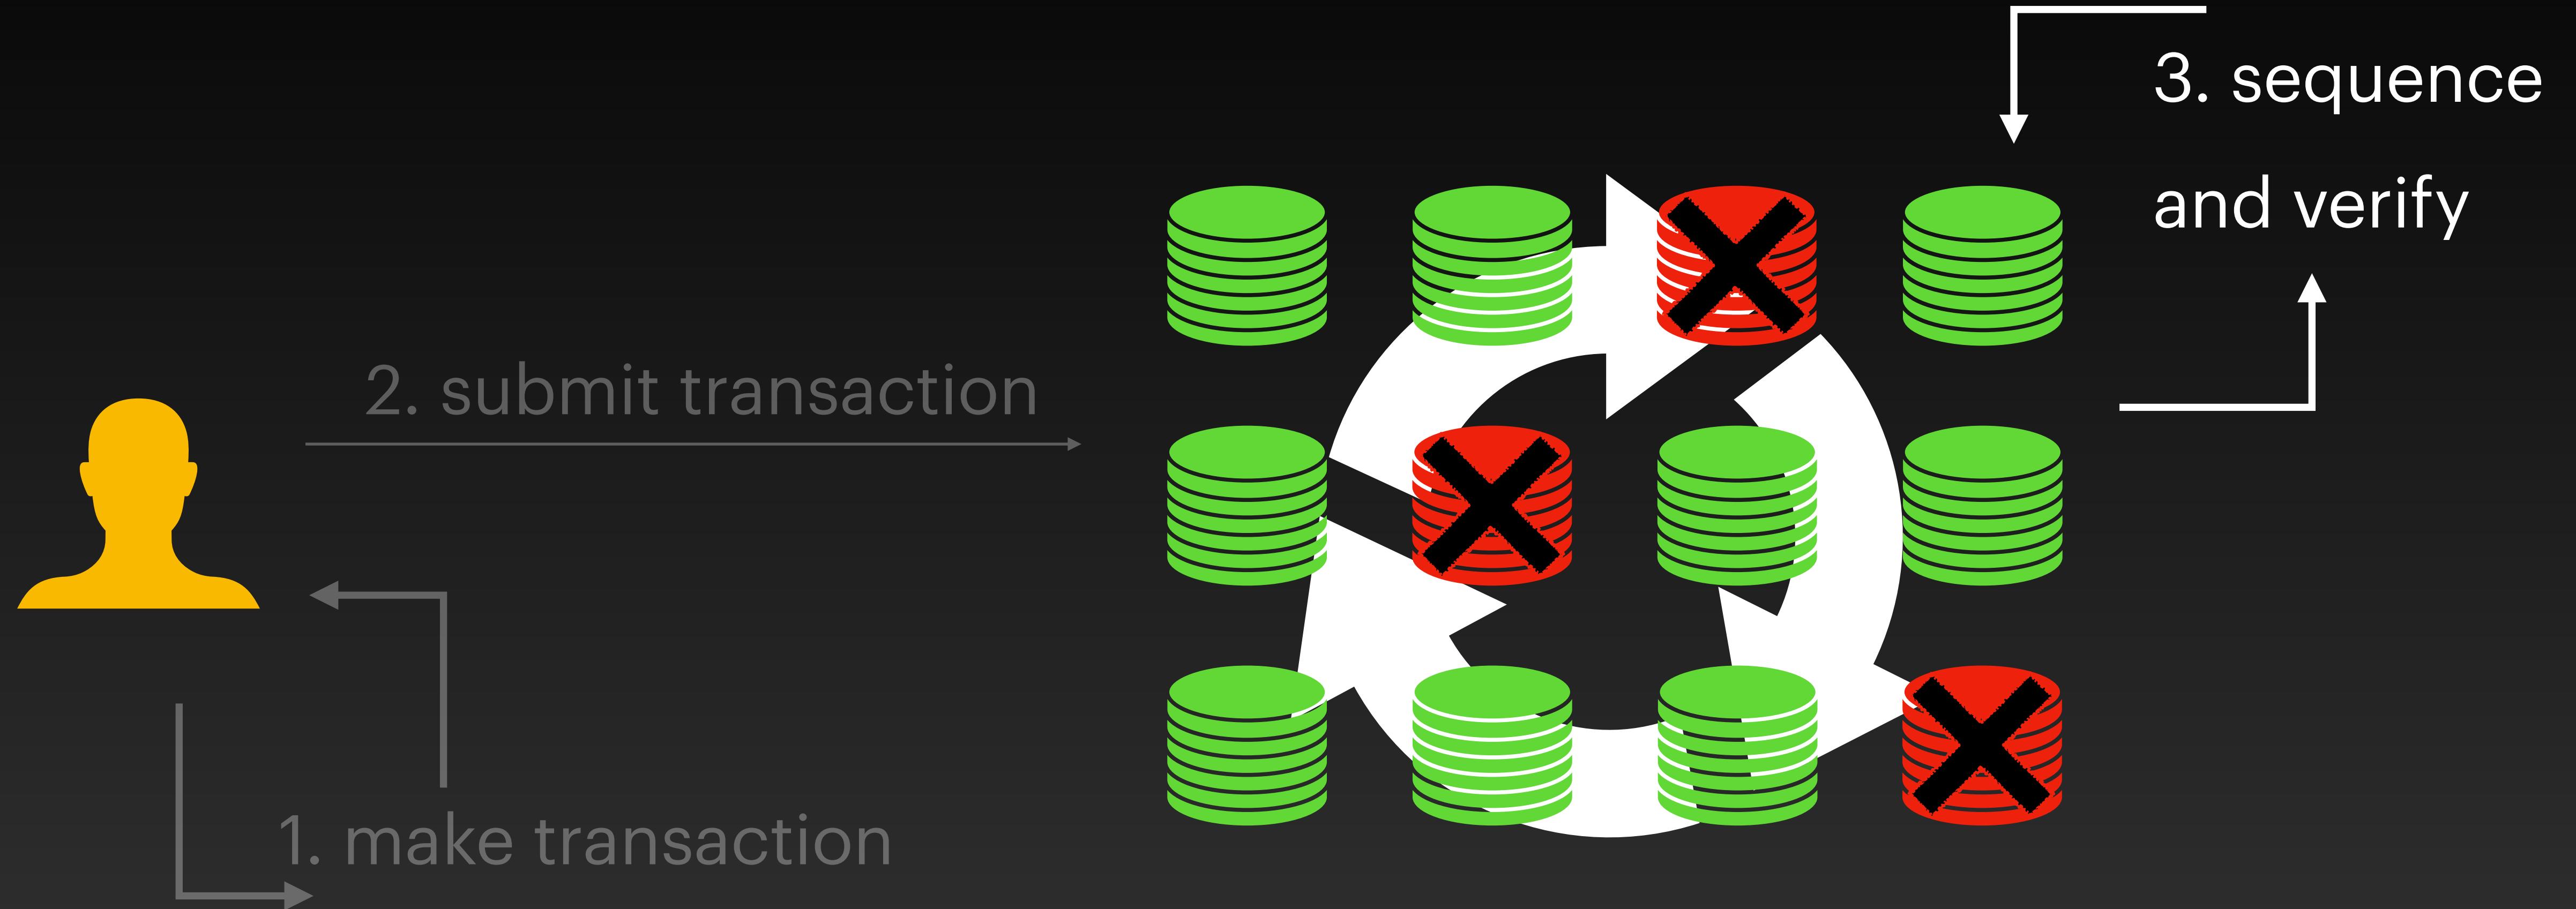


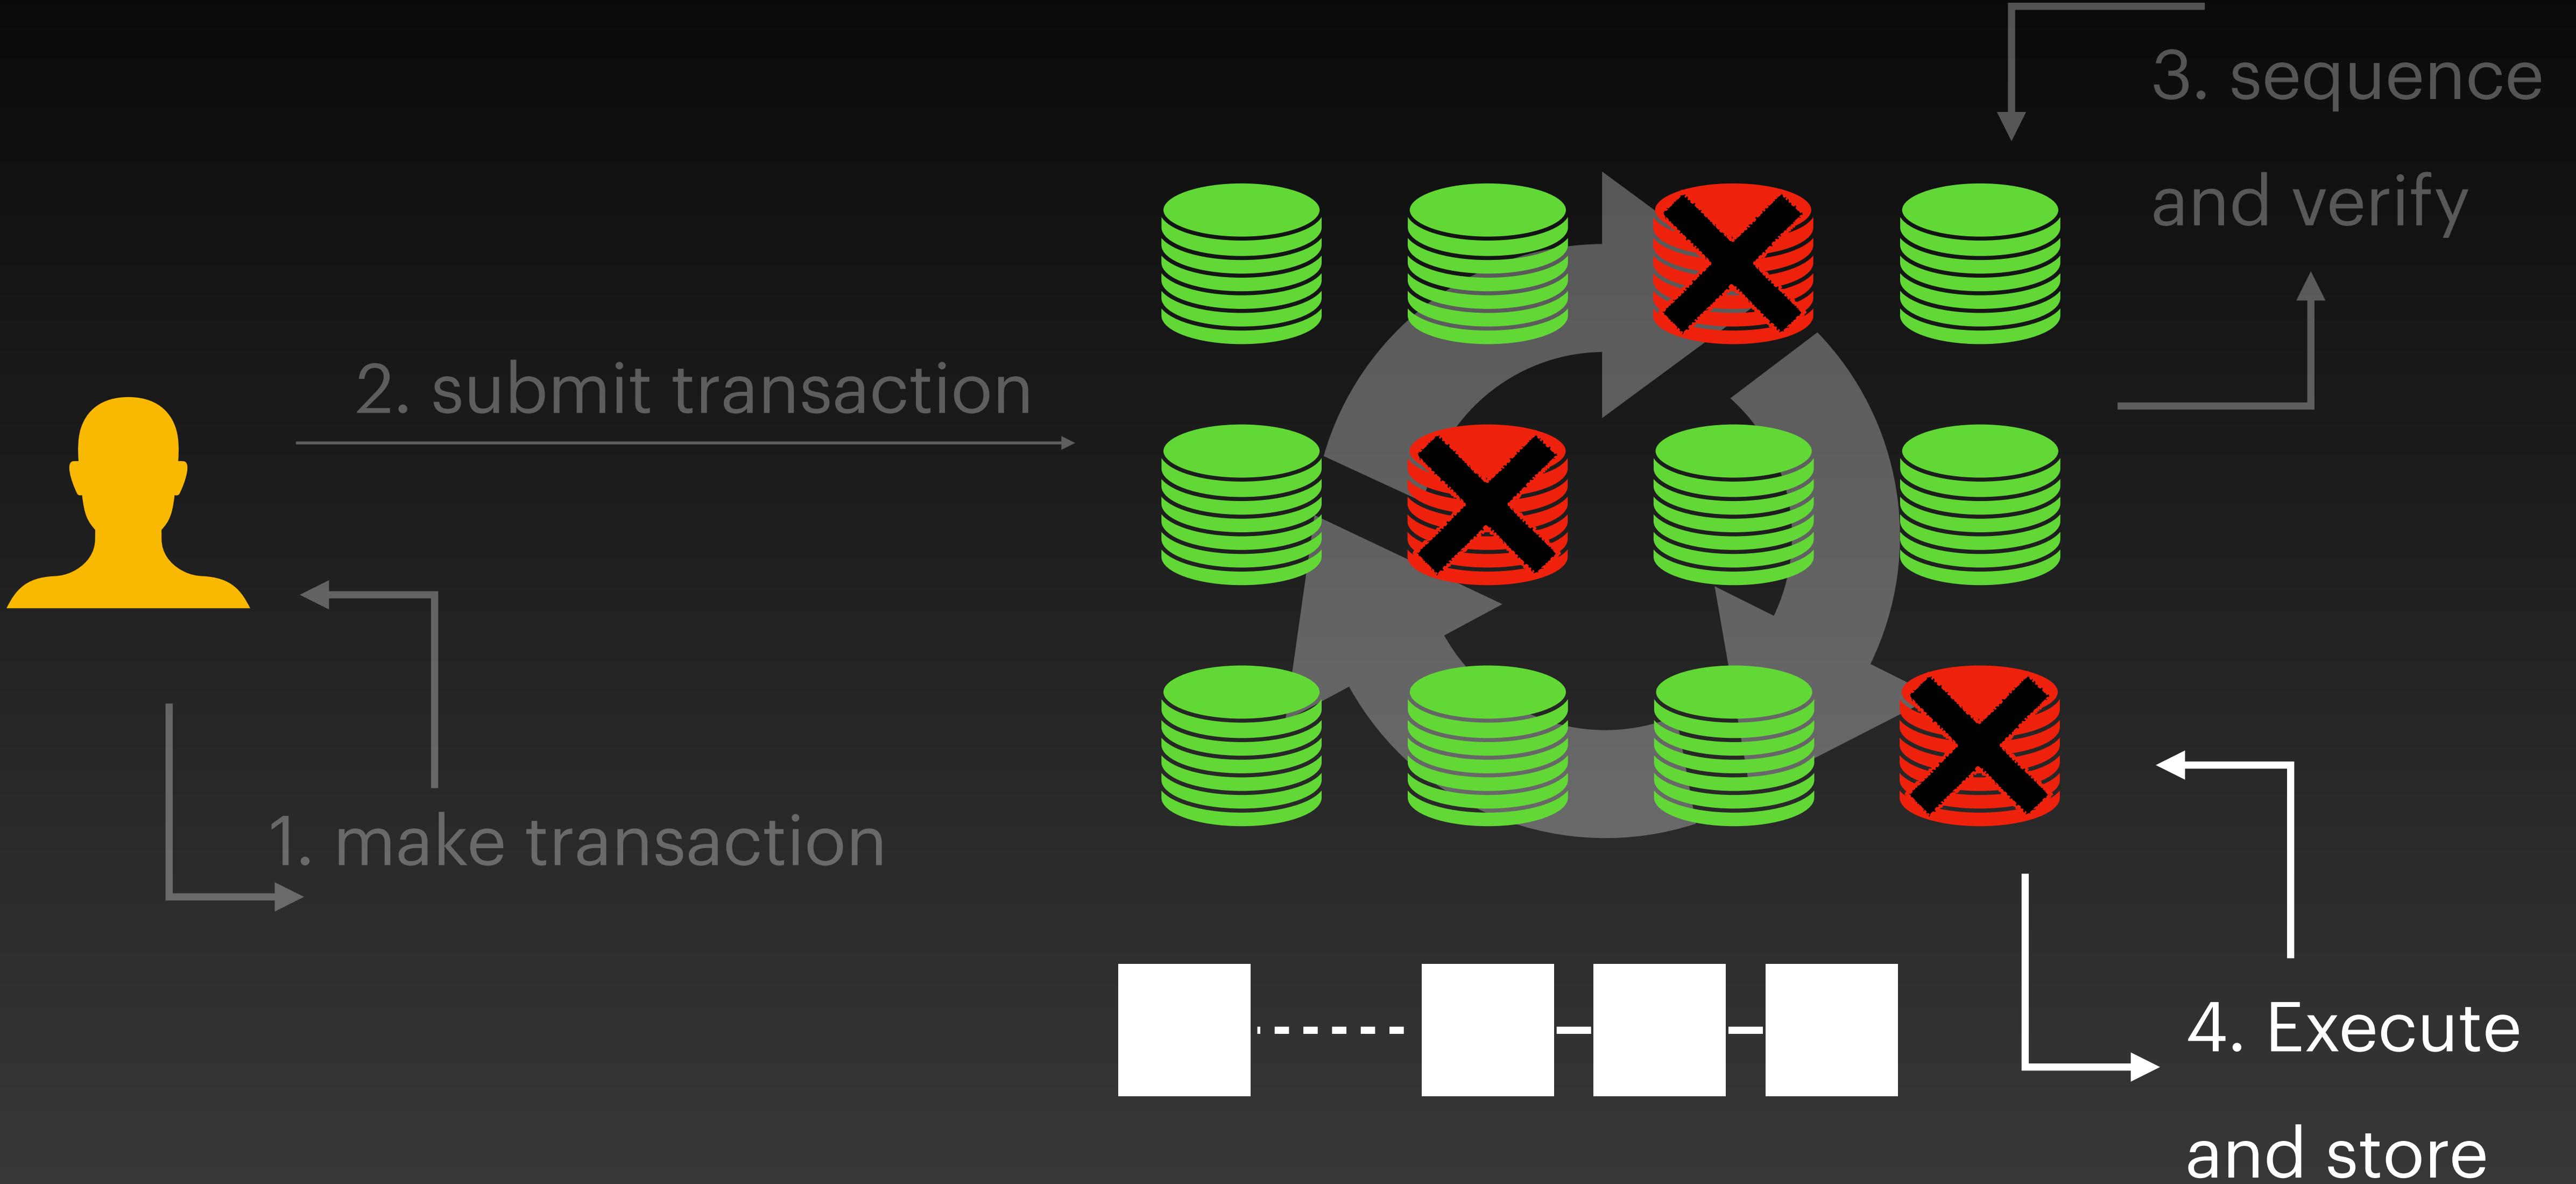
2. submit transaction



1. make transaction







- **Distributed Systems**
  - But not like a DB running in my datacenter
  - Adversarial network and Byzantine adversaries
- **Systems Security**
  - Both network and systems security
  - Interaction between networked components
- **Programming Languages**
  - Execute the smart contract & ensure determinism
  - Solidity, Move
- **Cryptography**
  - Validators cannot use secrets to execute smart contracts
  - Anonymous credentials, ZK-proofs

- **Distributed Systems**
  - But not like a DB running in my datacenter
  - Adversarial network and Byzantine adversaries
- **Systems Security**
  - **Both network and systems security**
  - **Interaction between networked components**
- **Programming Languages**
  - Execute the smart contract & ensure determinism
  - Solidity, Move
- **Cryptography**
  - Validators cannot use secrets to execute smart contracts
  - Anonymous credentials, ZK-proofs

# Security Properties

## Safety

**Undesirable things never happen**

## Liveness

**Desirable things eventually happen**

# Adversary

## #1 The Network: Worst possible schedule

### Properties

- **Synchronous:** A message sent will be delivered before a maximum (known) delay.
- **Asynchronous:** A message sent will eventually be delivered at an arbitrary time before a maximum (unknown) delay.
- **Partial Synchronous:** the network is asynchronous but after some time it enters a period of synchrony.
- 

### Challenges

- Theoretical models: Need careful implementation to ensure we approximate them, e.g., retransmissions.
- Memory: Naive implementations use infinite buffers. Identify conditions after which retransmissions are not necessary and buffers can be freed.
- Asynchrony means the protocol should maintain properties for any re-ordering of message deliveries.
- Unknown delay means delay should be adaptive to ensure robustness.

# Adversary

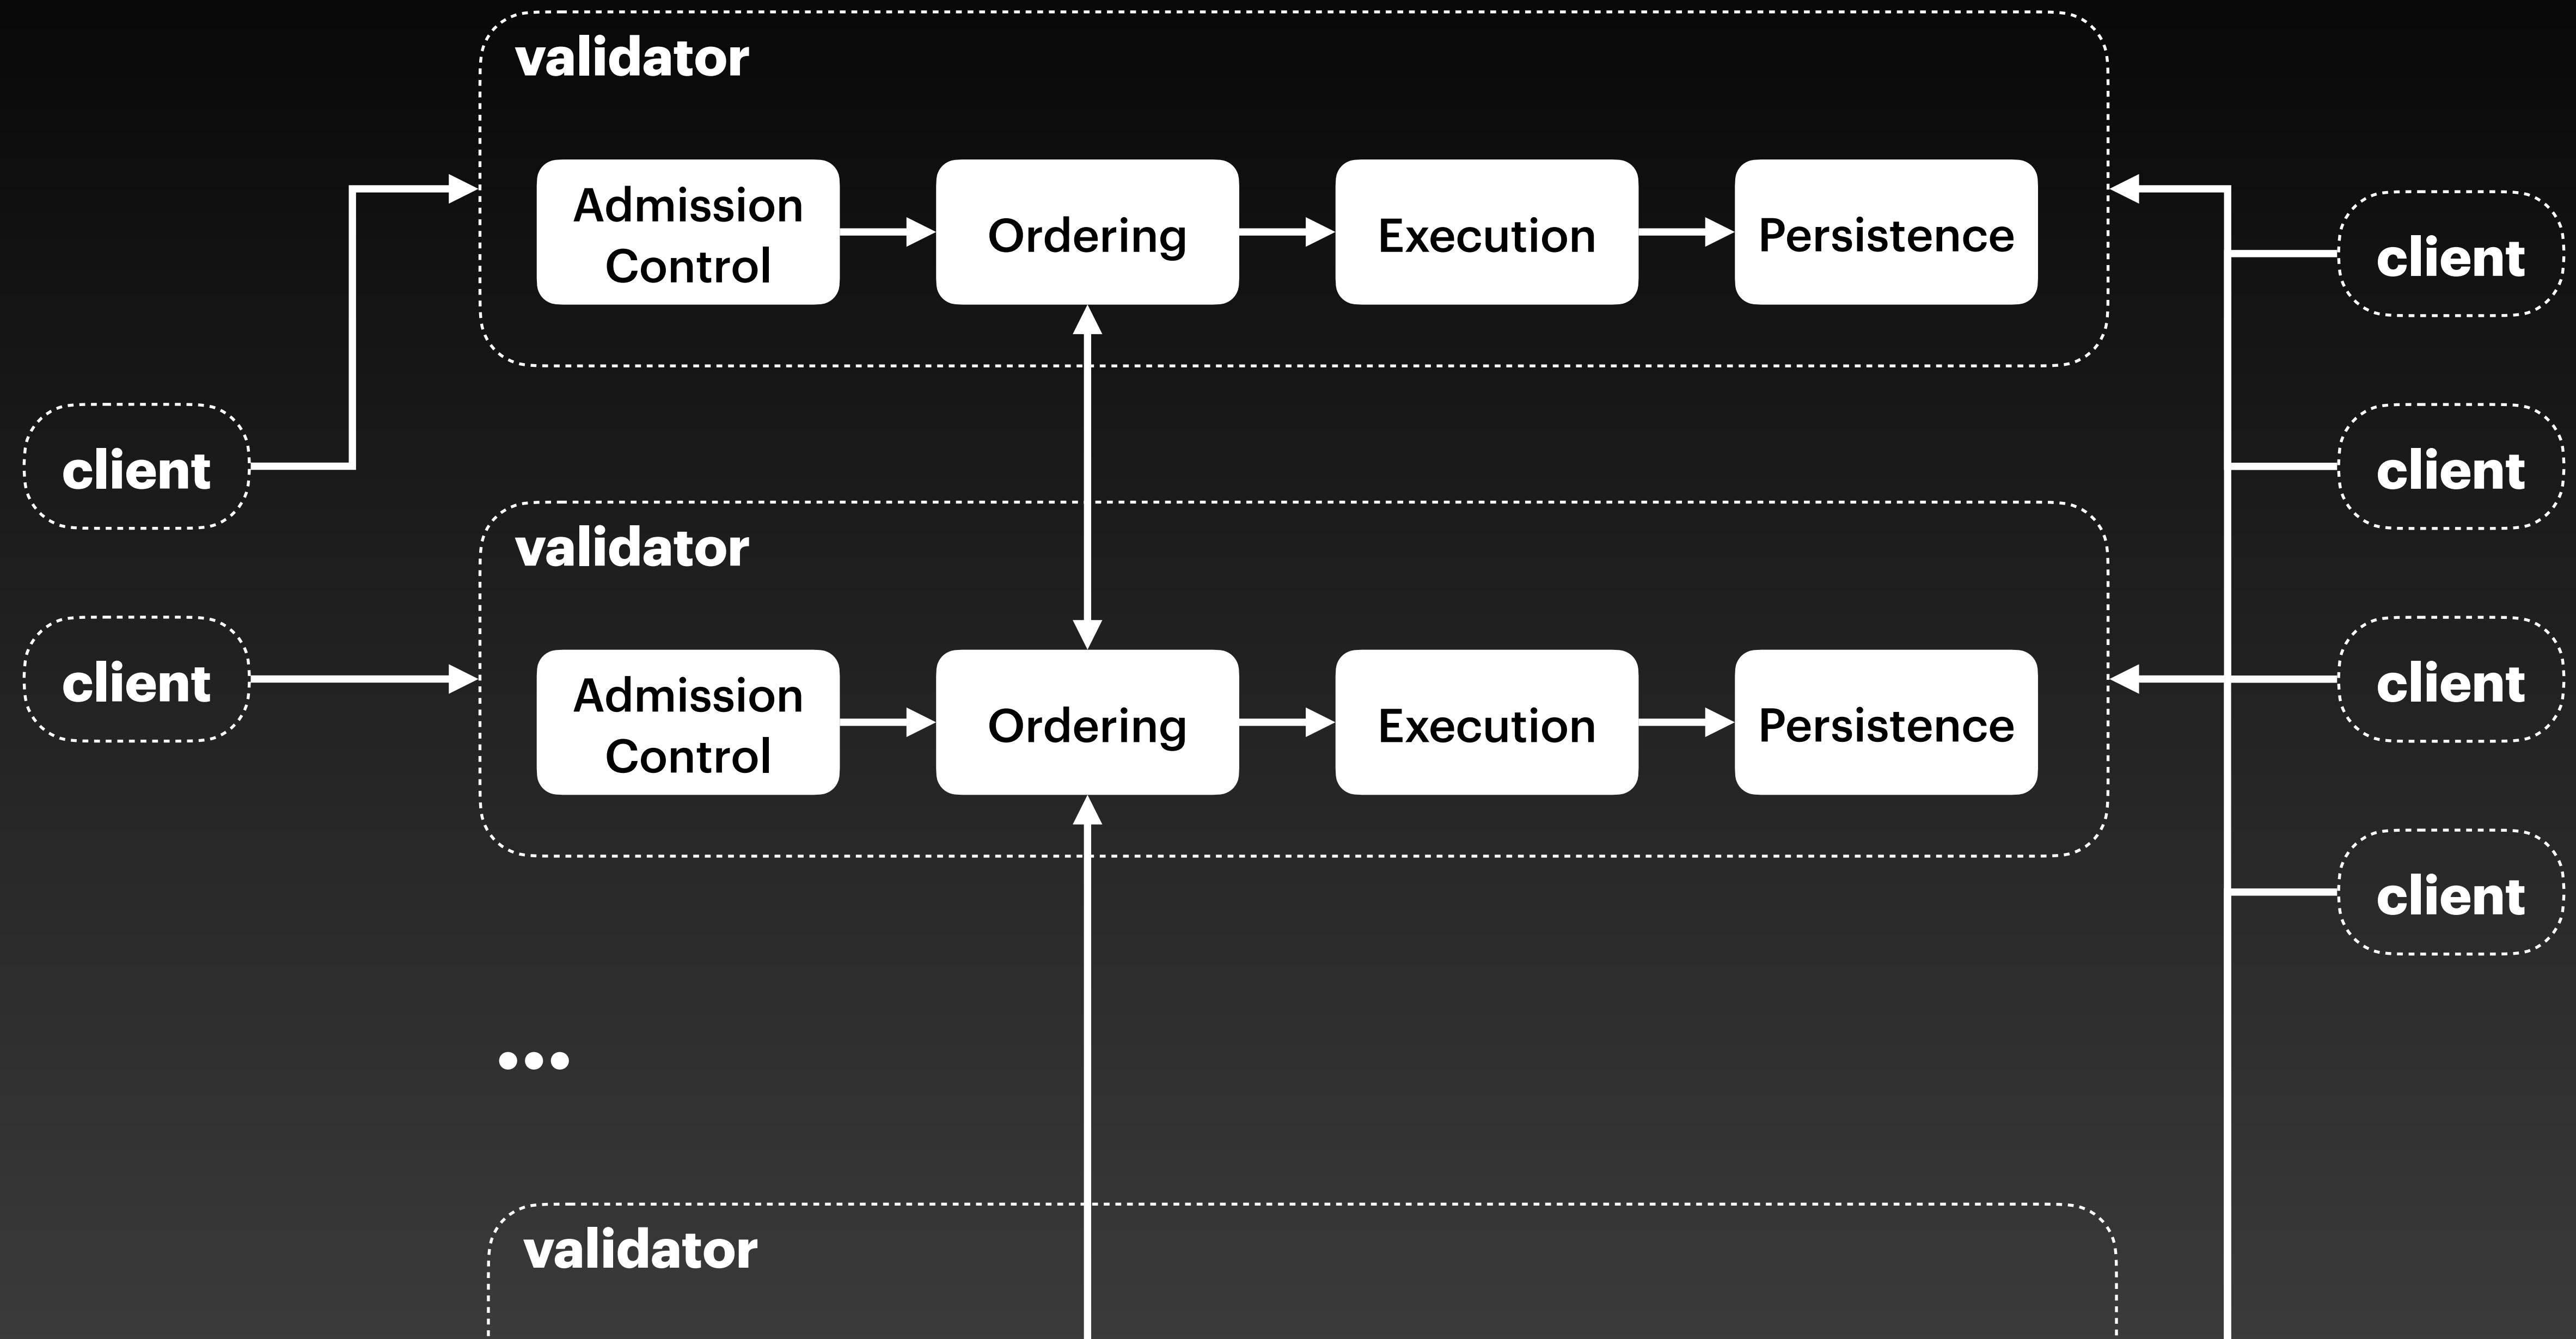
## #2 Bad Nodes: Arbitrary behaviour

### Properties

- **Correct / honest / good:** Will remain live and follow the protocol as specified by the designers of the system.
- **Byzantine:** will deviate arbitrarily from the protocol. May respond incorrectly or not at all.

### Challenges

- **Crash & recover:** still a correct validators with very high latency. Need persistence to ensure this
- **Rational:** honest validators may have some discretion. They may use it to maximise profit



# Network Security

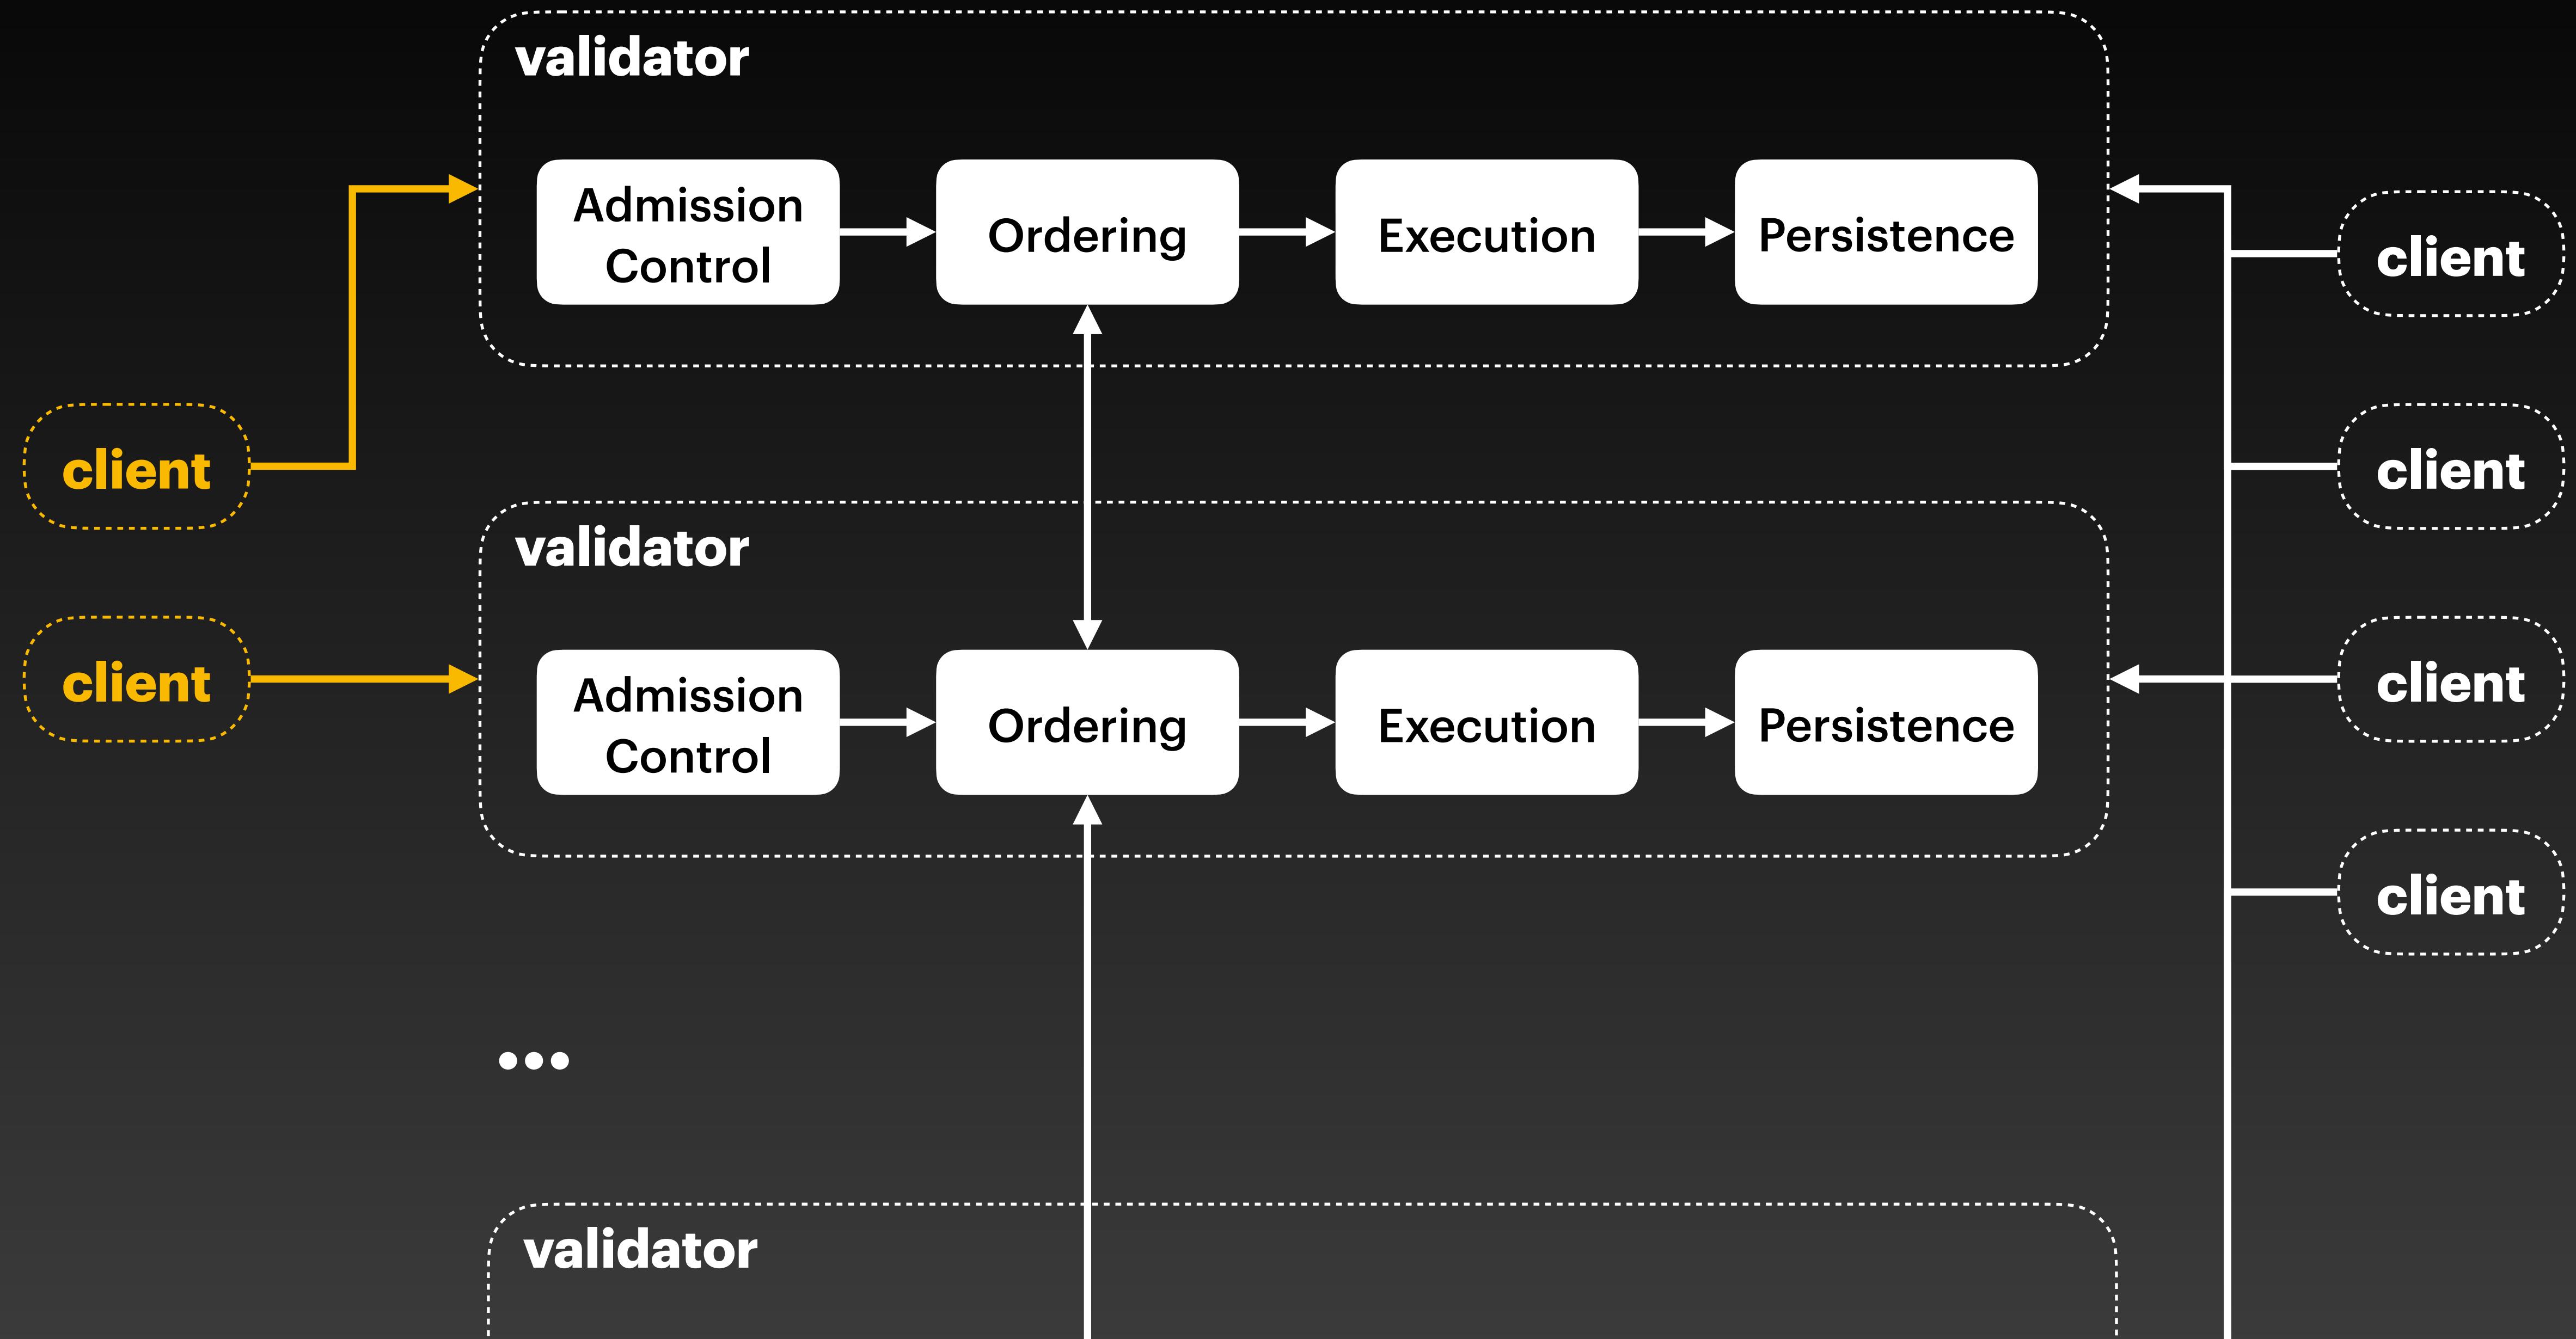
## Challenge #1: Nodes

- **Validators are exposed (not in datacenter no on beefy machines)**

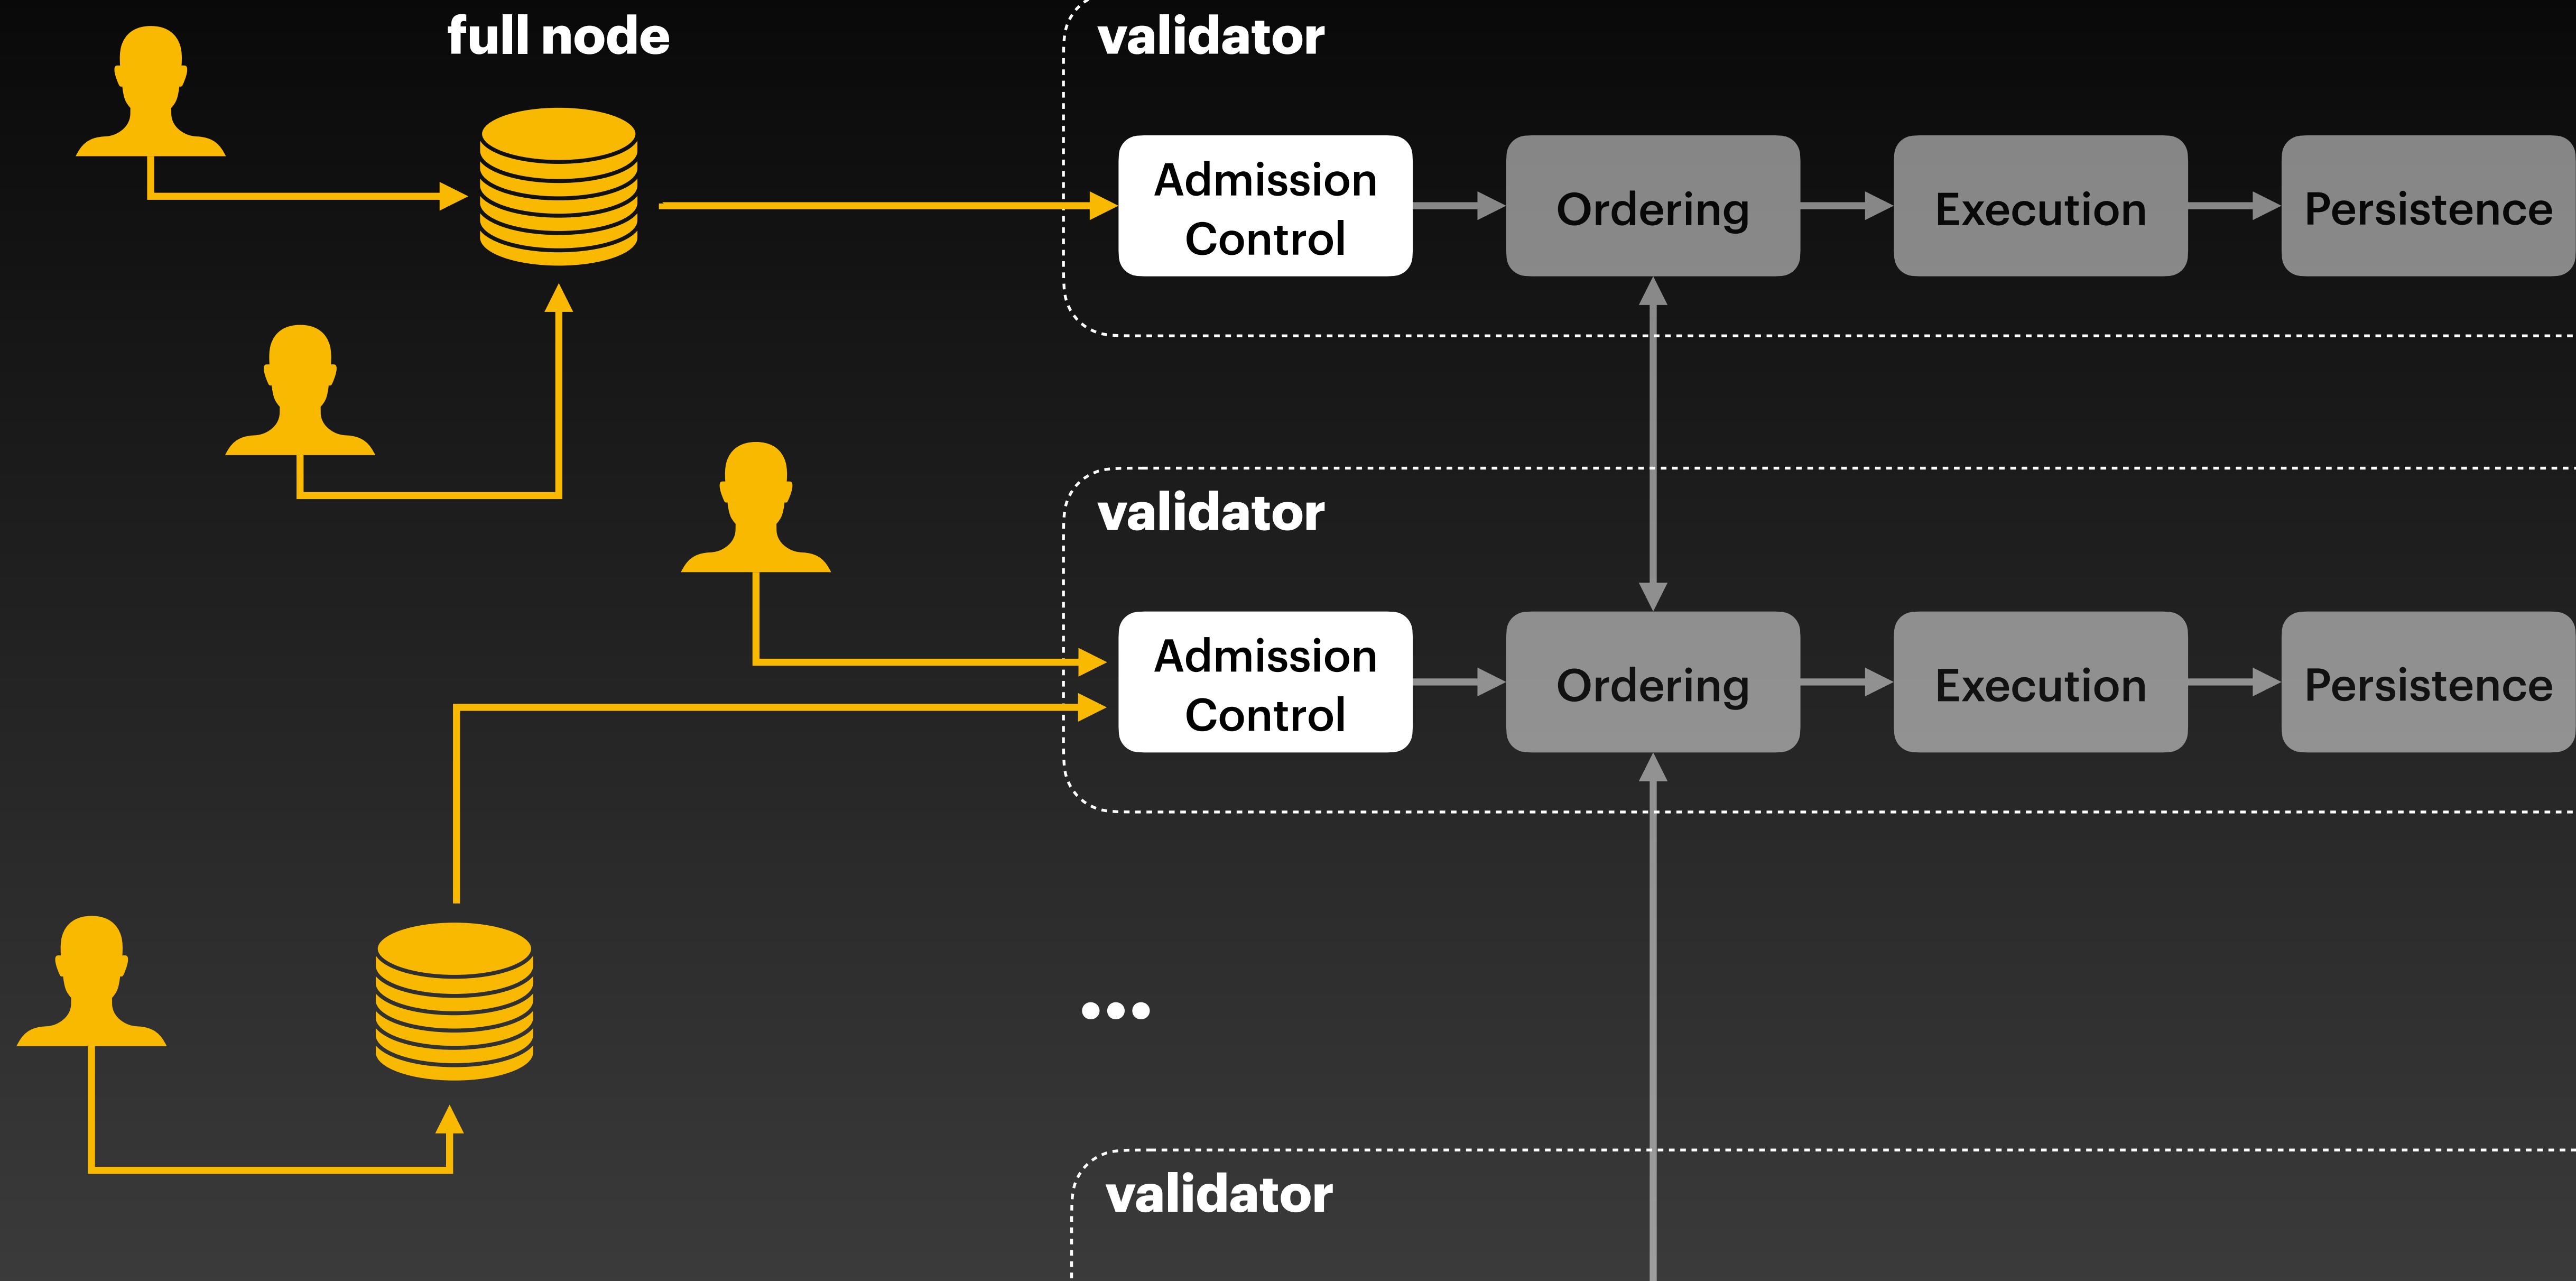
# Network Security

## Challenge #1: Nodes

- Validators are exposed (not in datacenter no on beefy machines)
- **Highly dynamic set of nodes**



**light client**



# Network Security

## Challenge #2: Clients

- **Different types of target links: clients-validator and validator-validator**

# Network Security

## Challenge #2: Clients

- Different types of target links: clients-validator and validator-validator
- **Highly dynamic clients**

# Network Security

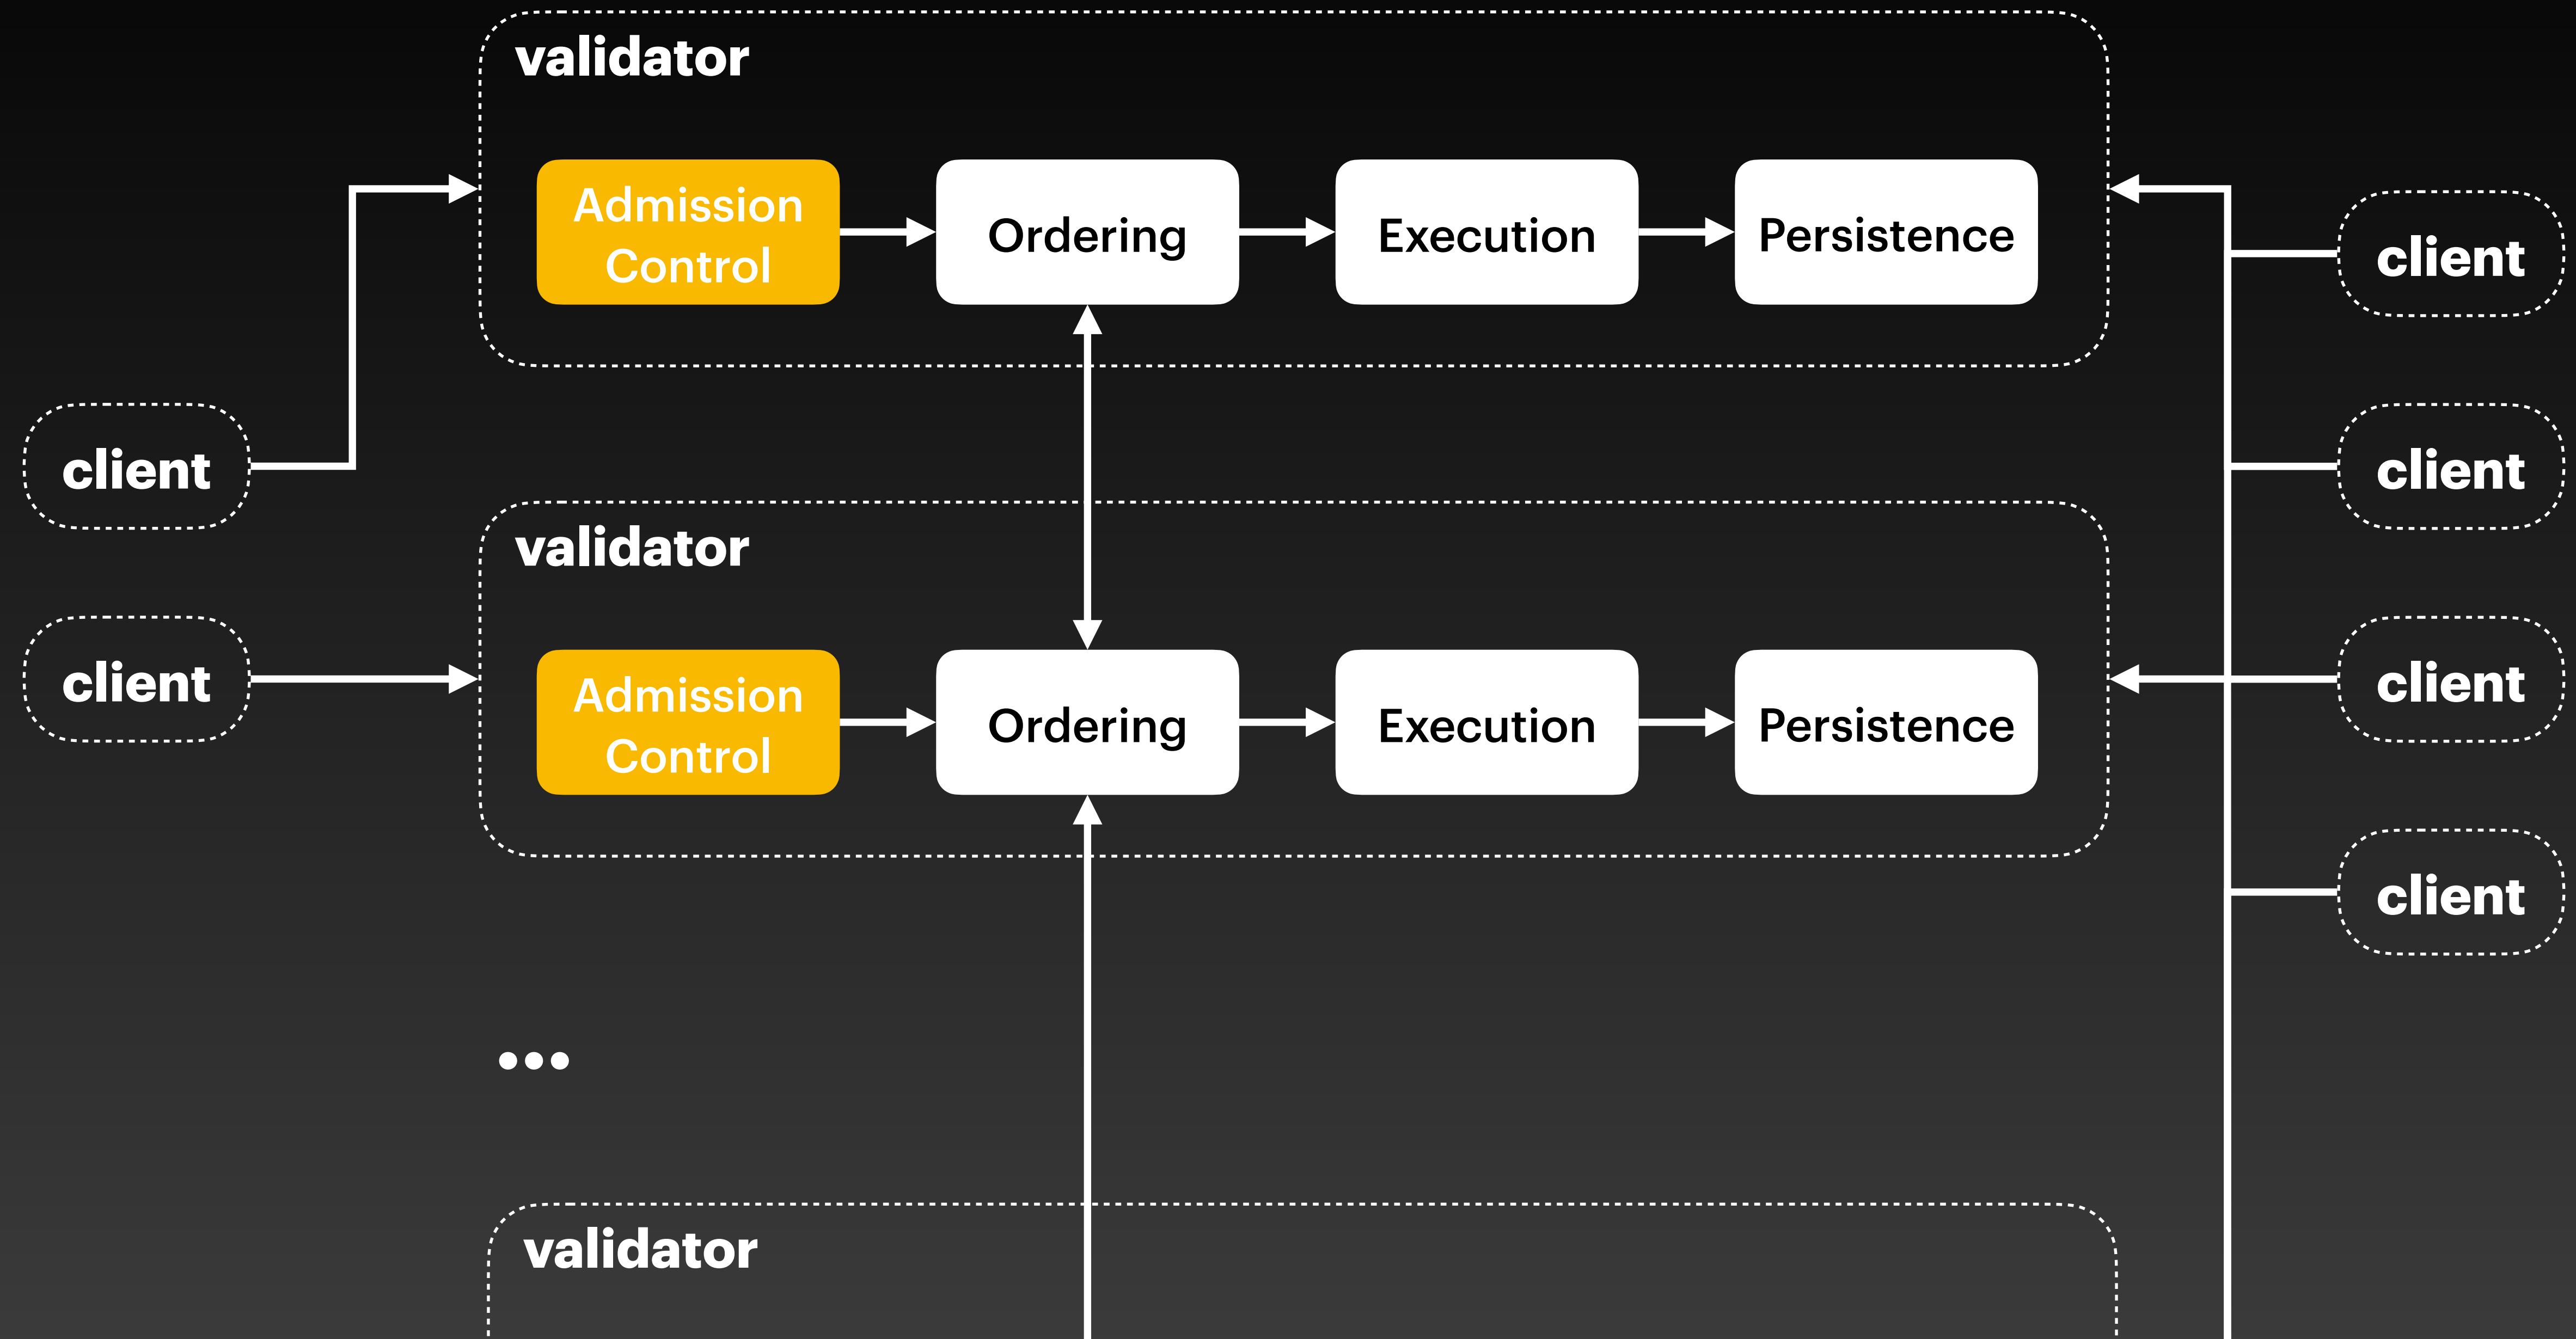
## Challenge #2: Clients

- Different types of target links: clients-validator and validator-validator
- Highly dynamic clients
- **Clients have no fixed identity**

# Network Security

## Challenge #2: Clients

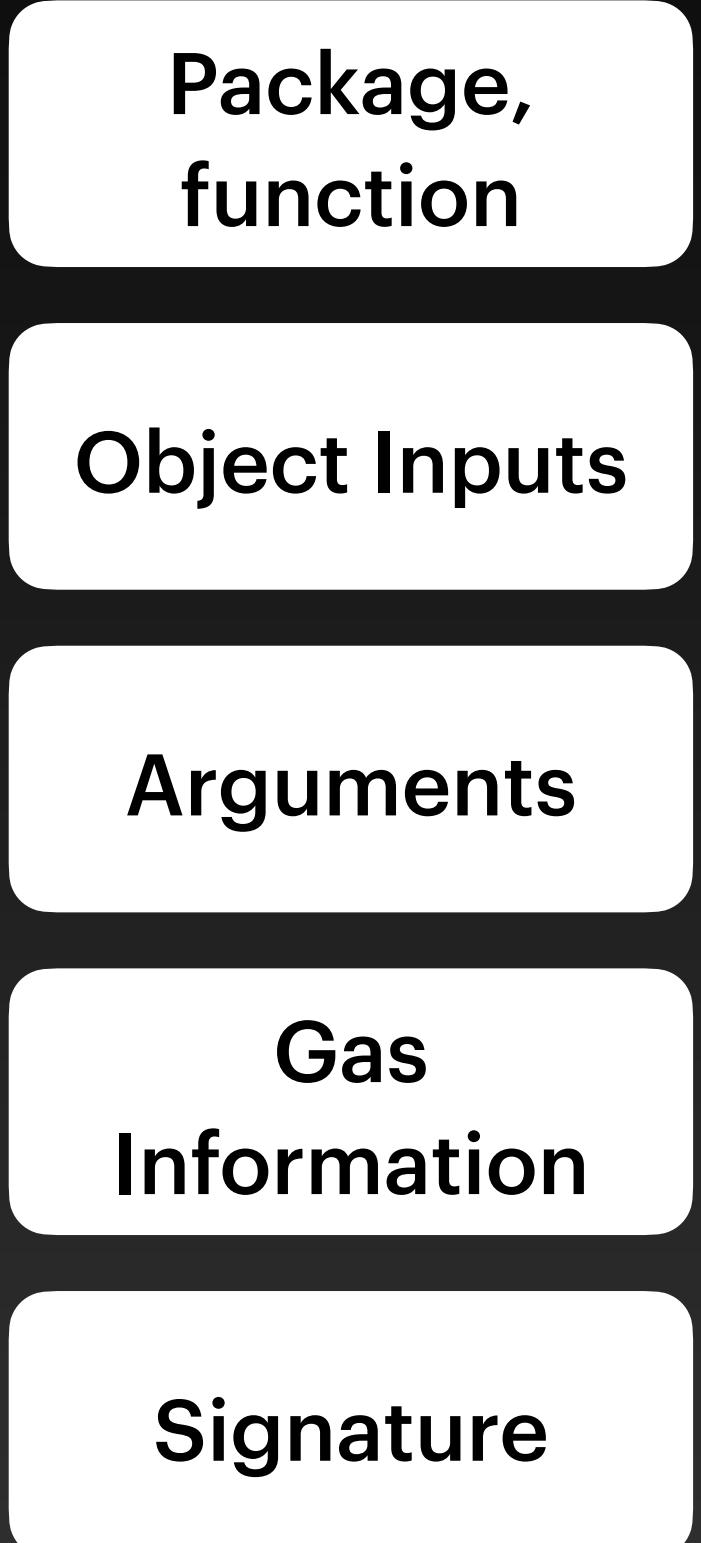
- Different types of target links: clients-validator and validator-validator
- Highly dynamic clients
- Clients have no fixed identity
- **Unclear validator selection algorithm**



## Objects:

- Unique ID
- Version number
- Ownership Information
- Type

Transaction's  
content



Coin::Send

Alice's account

Bob's account,  
Balance=5

0.001, max=0.005

## Example Transaction

**T1**

**Inputs:** O1, O2, O3

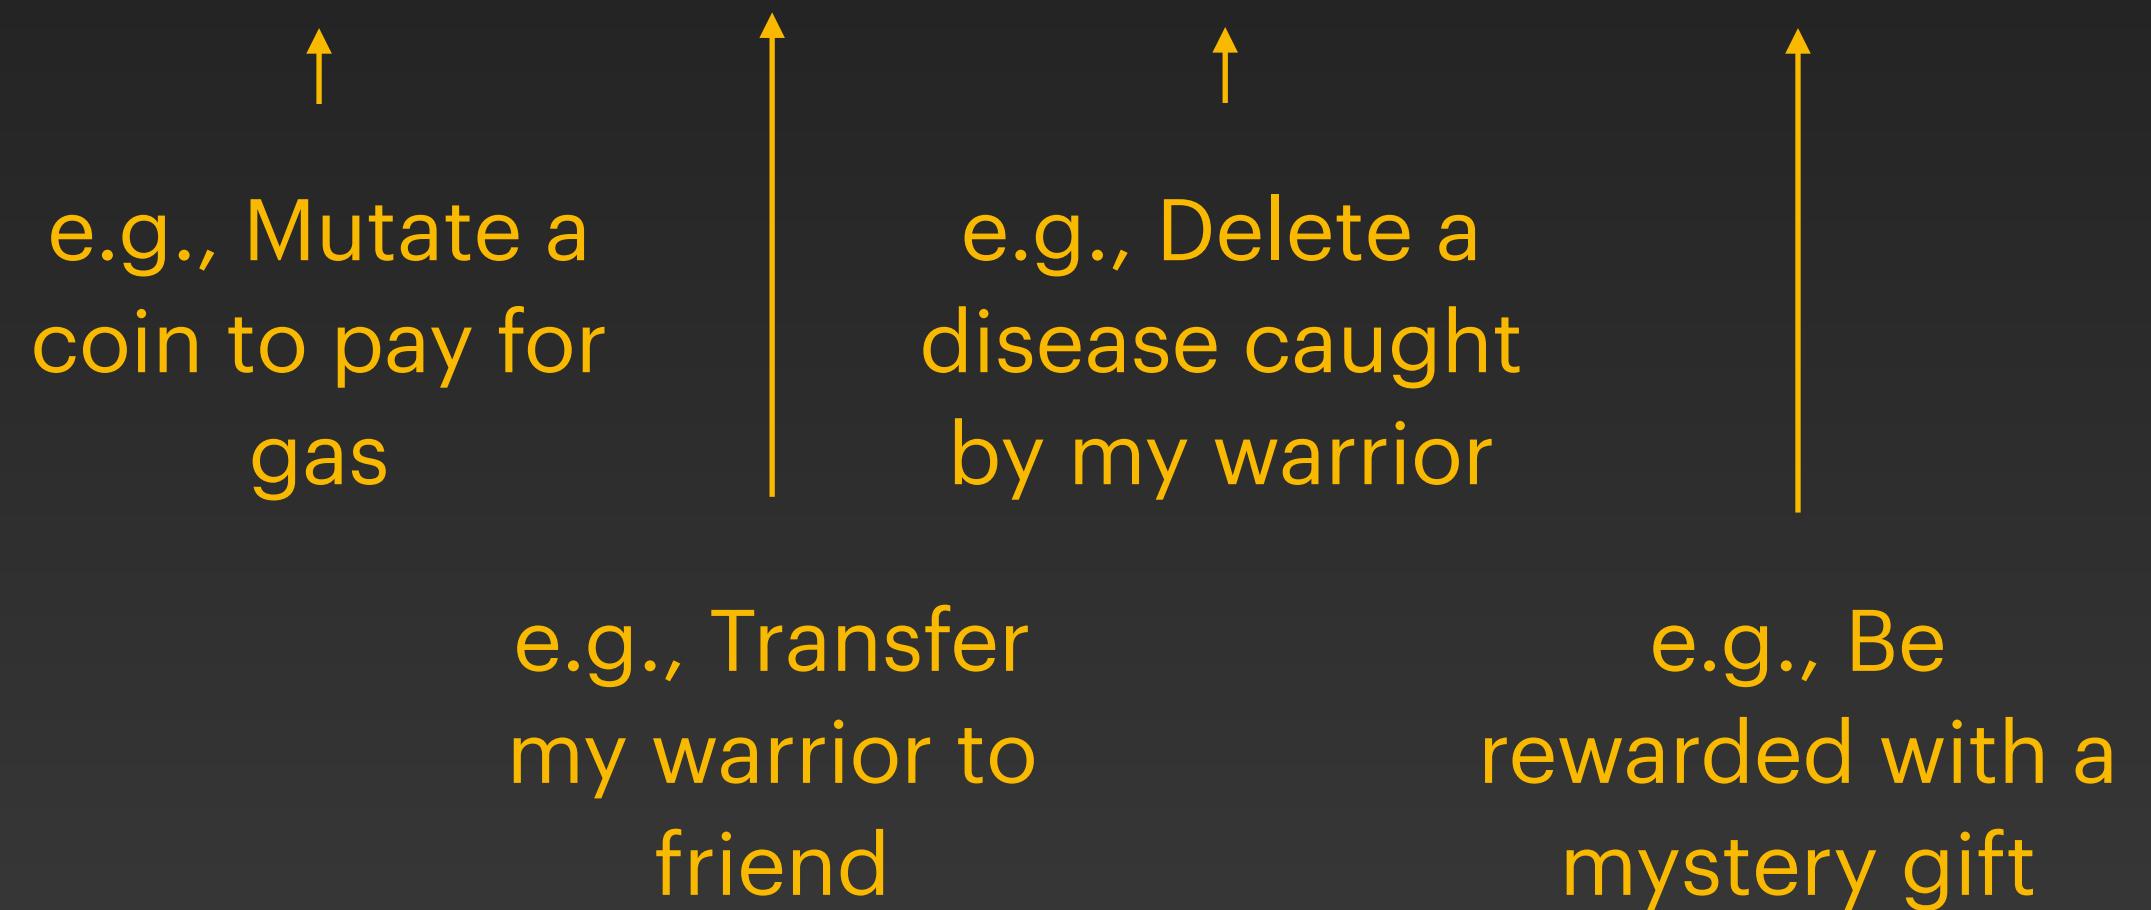
**Output:** Mutate O1, Transfer O2, Delete O3, Create O4

## Example Transaction

T1

**Inputs:** O1, O2, O3

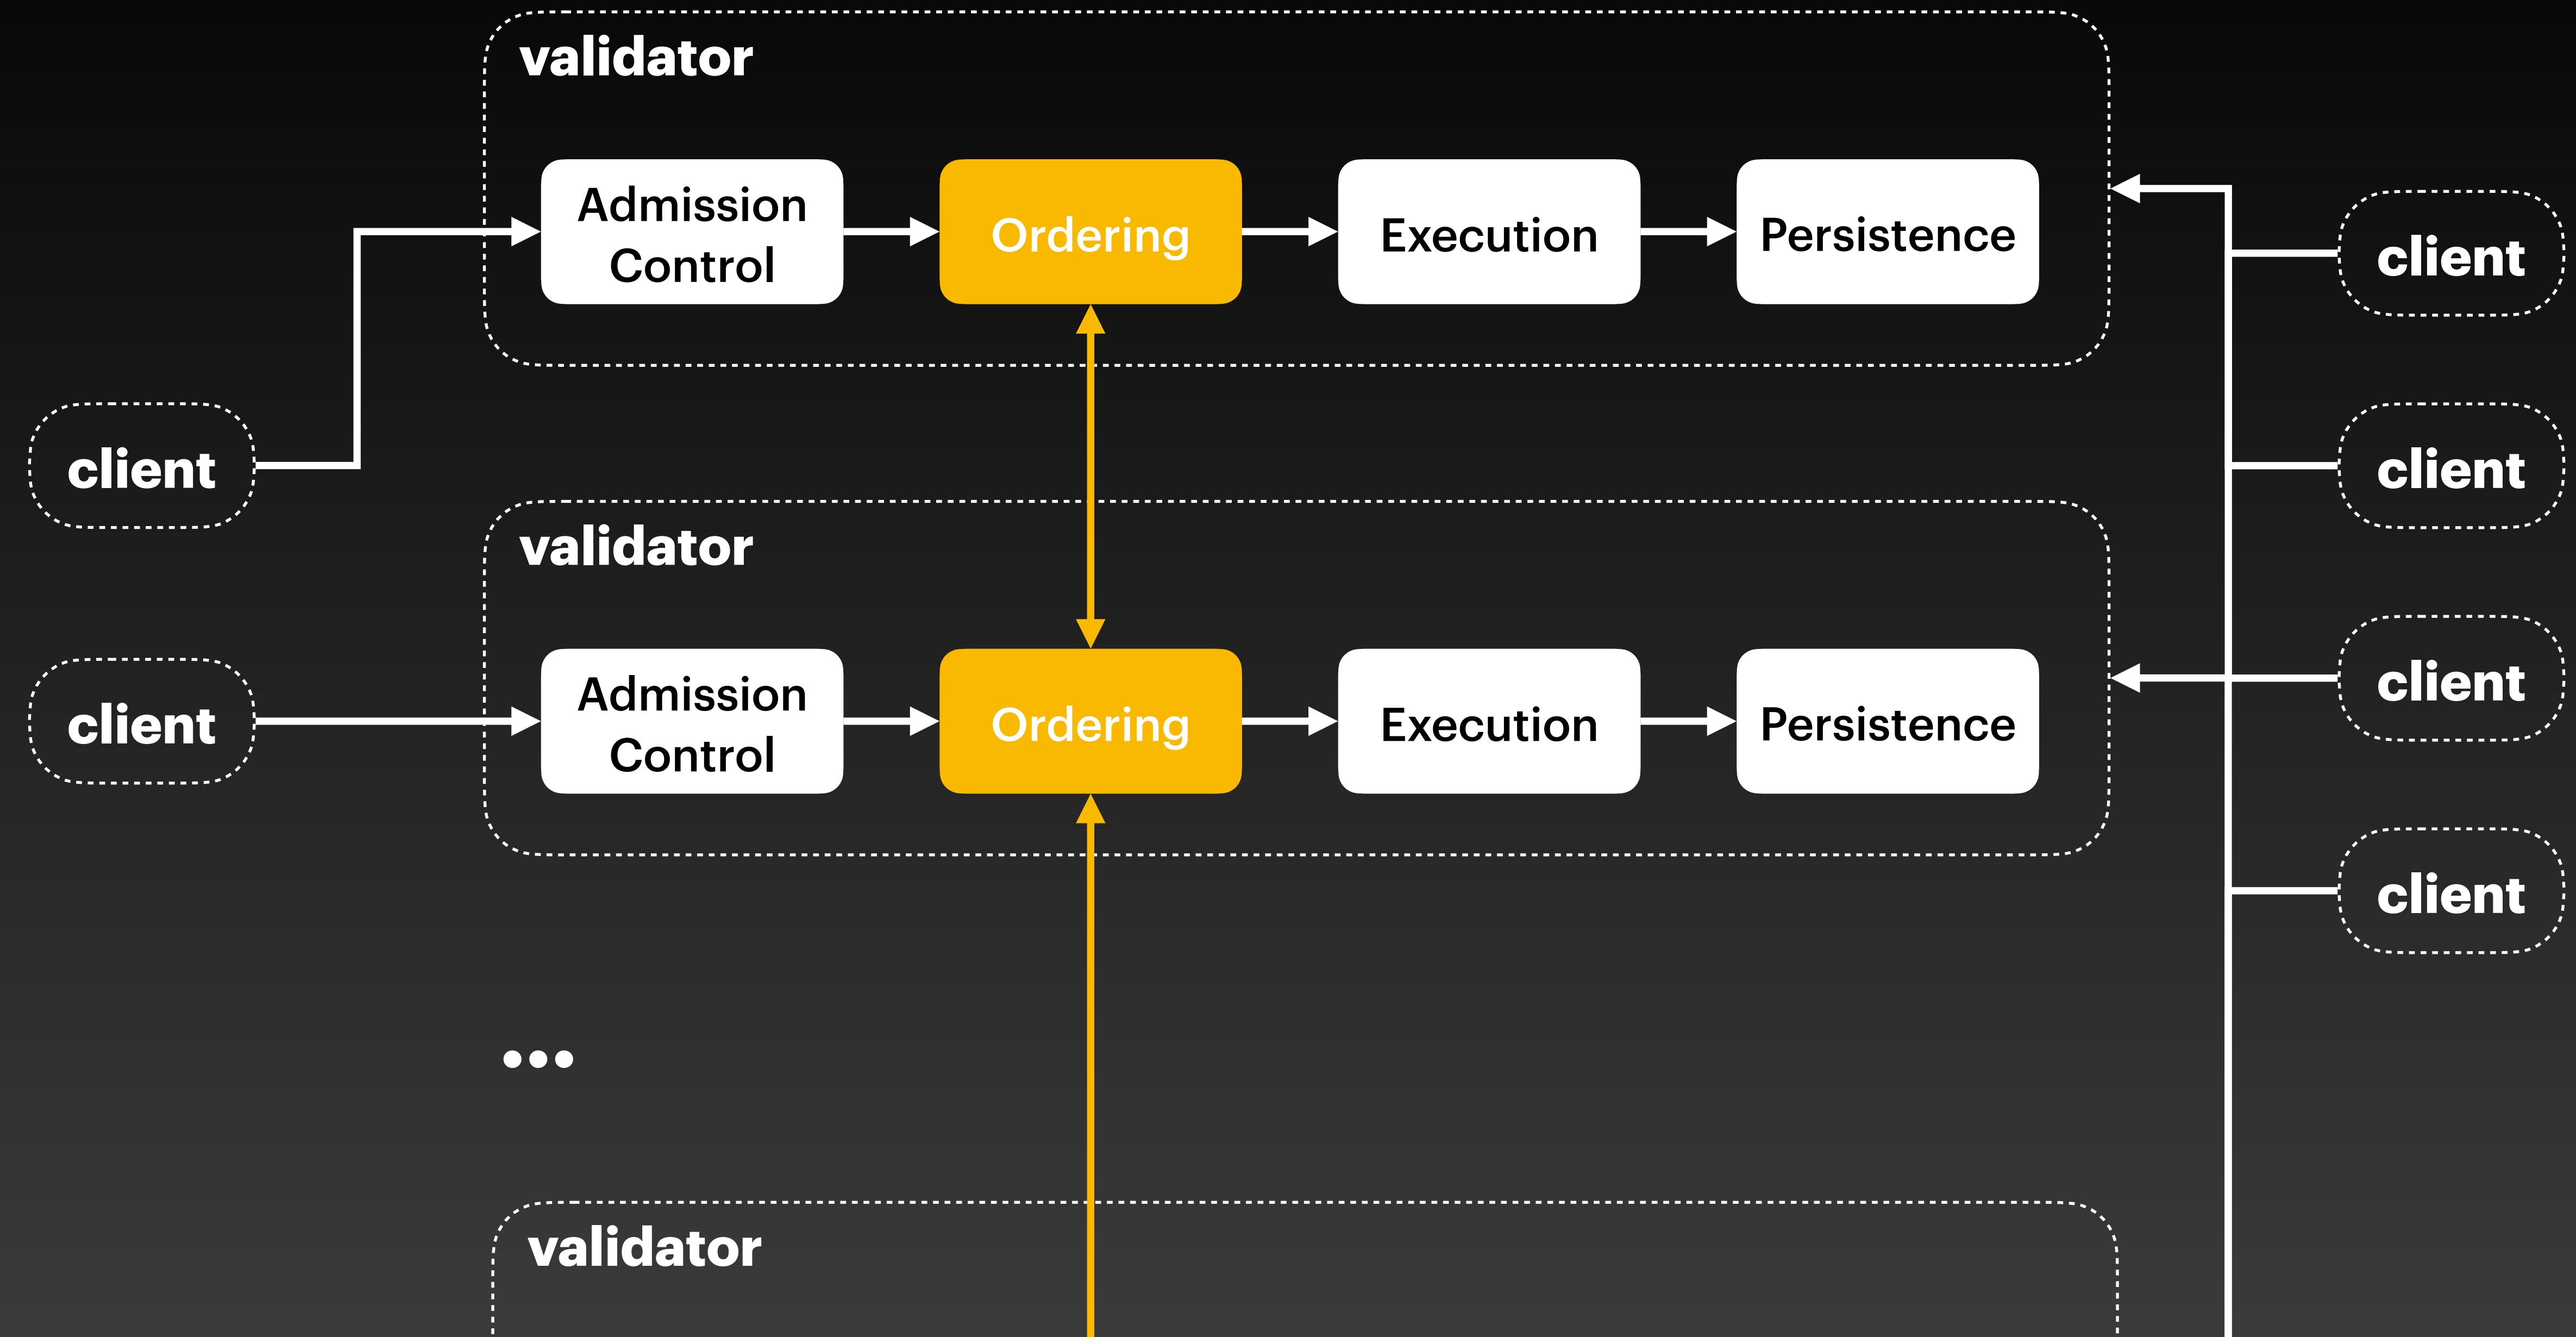
**Output:** Mutate O1, Transfer O2, Delete O3, Create O4

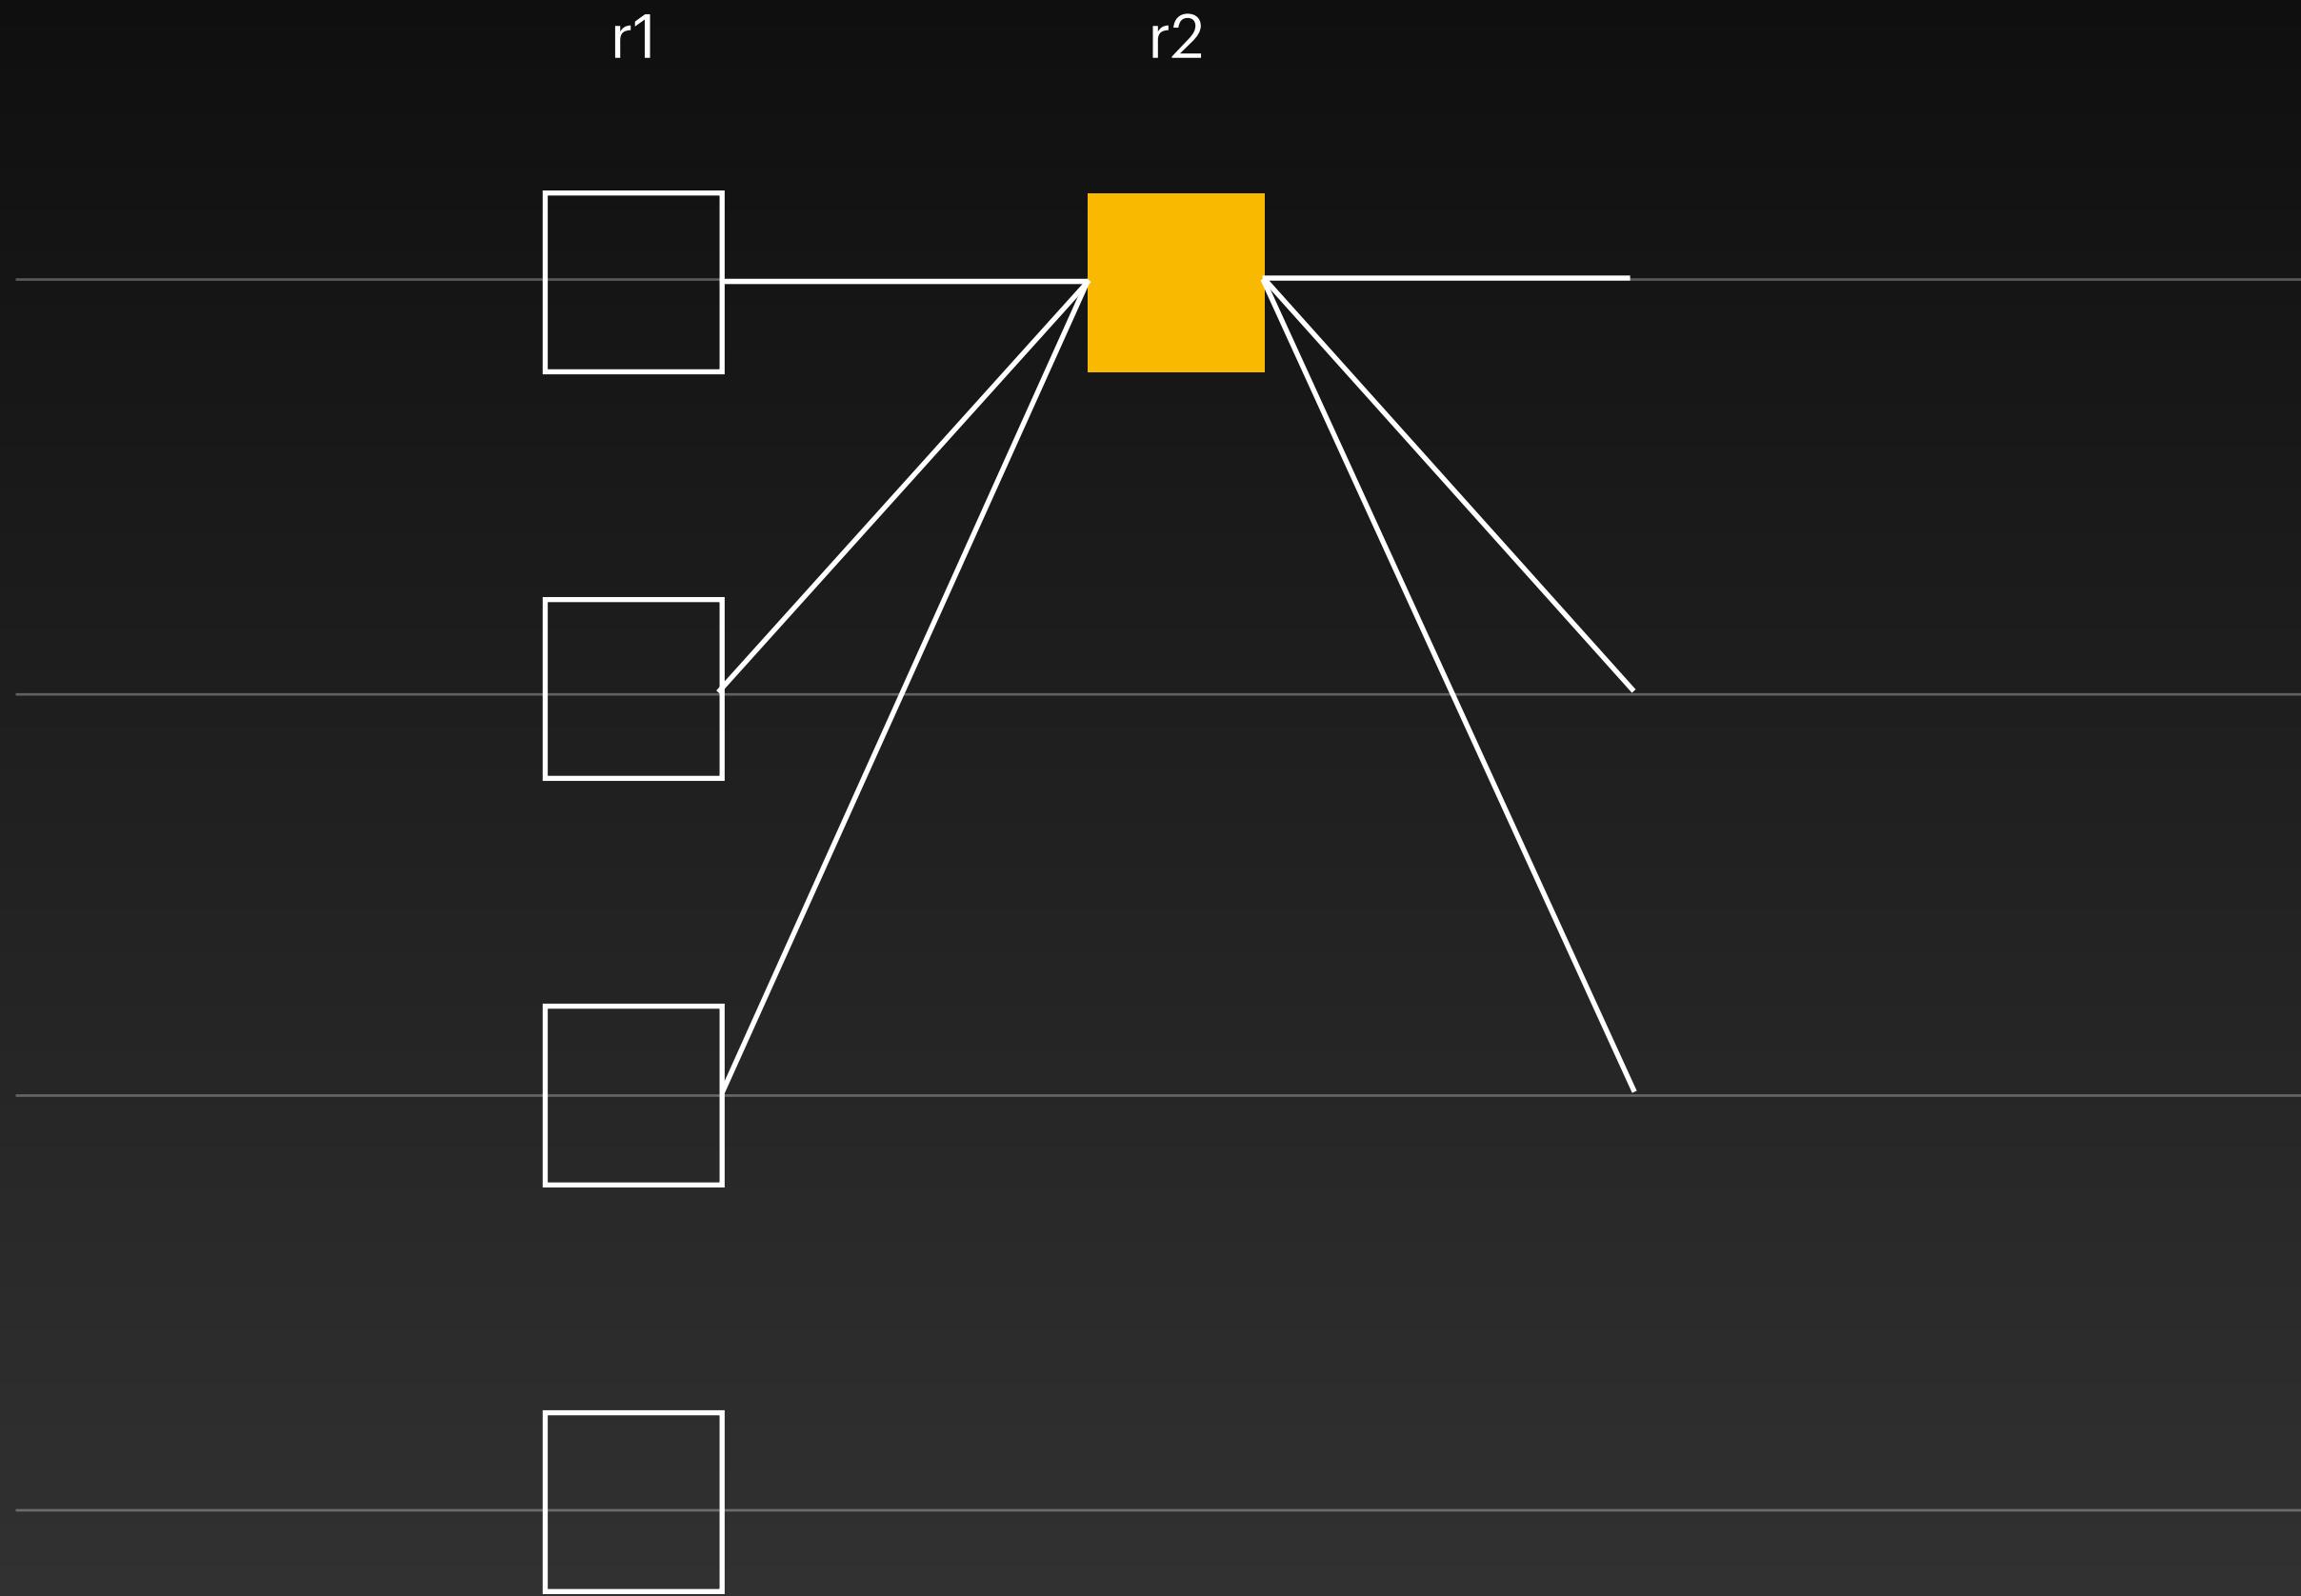


# Network Security

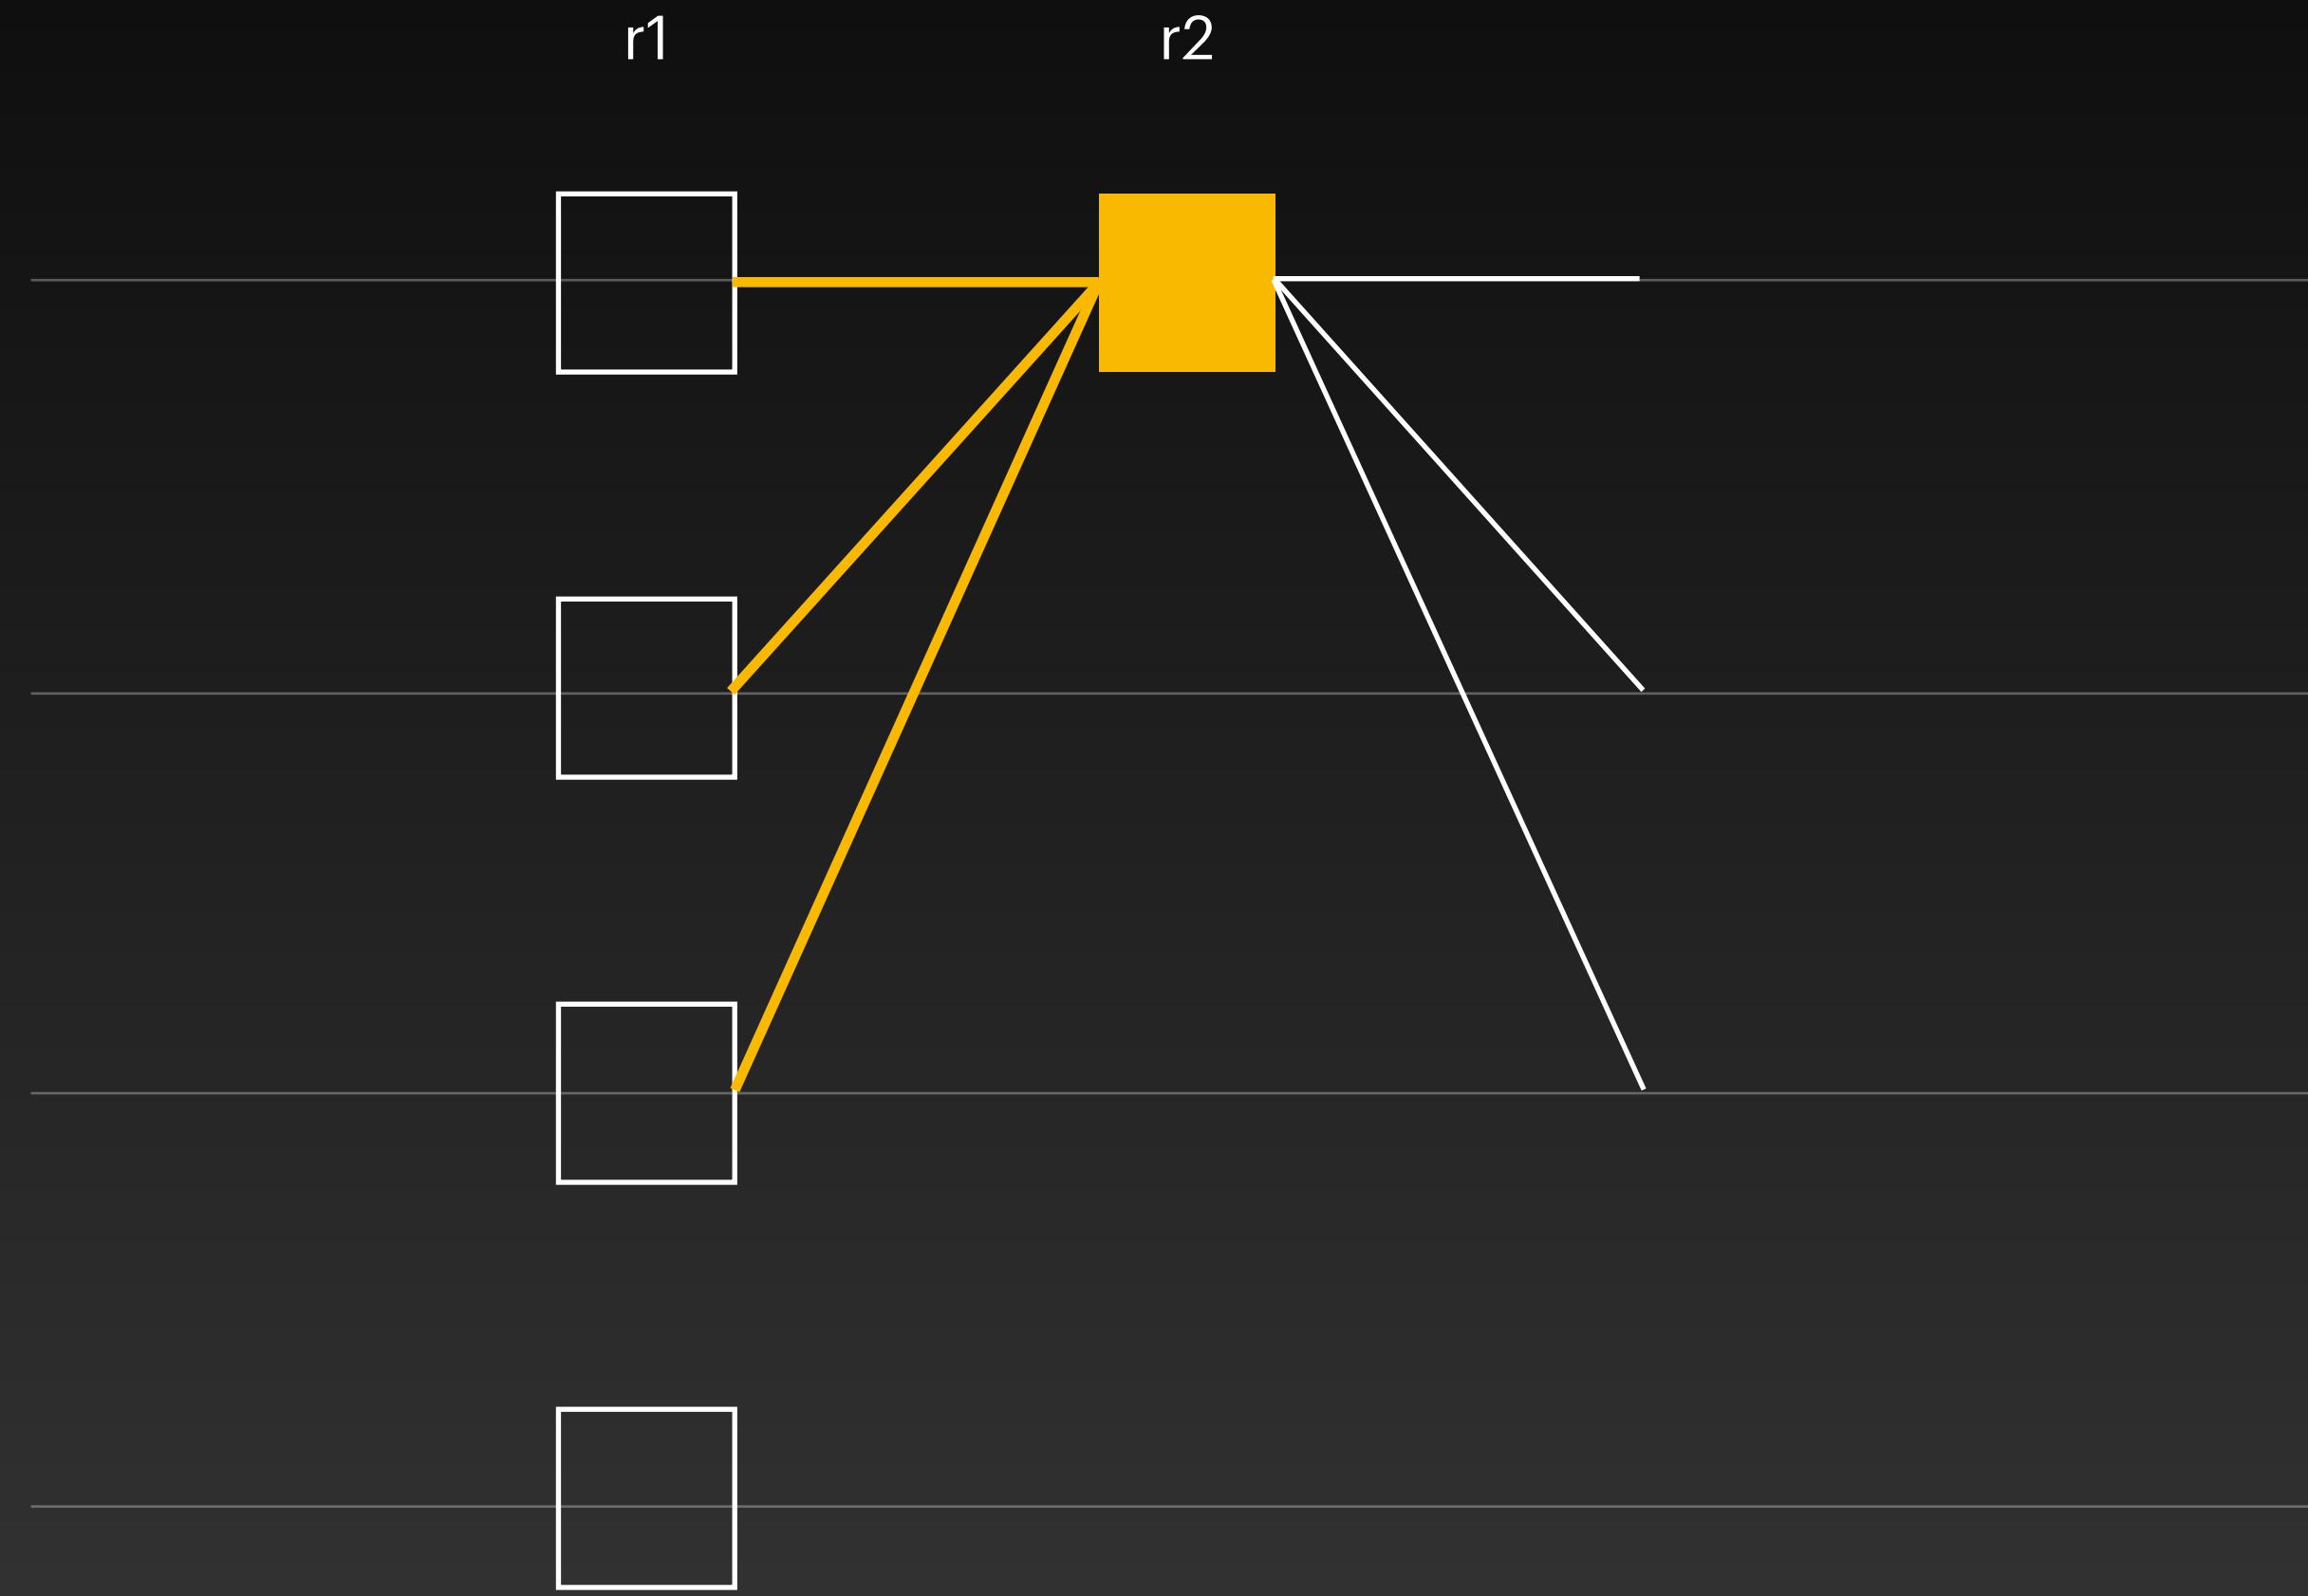
## Challenge #3: Admission Control

- **No established way to run pre-checks on input transactions**

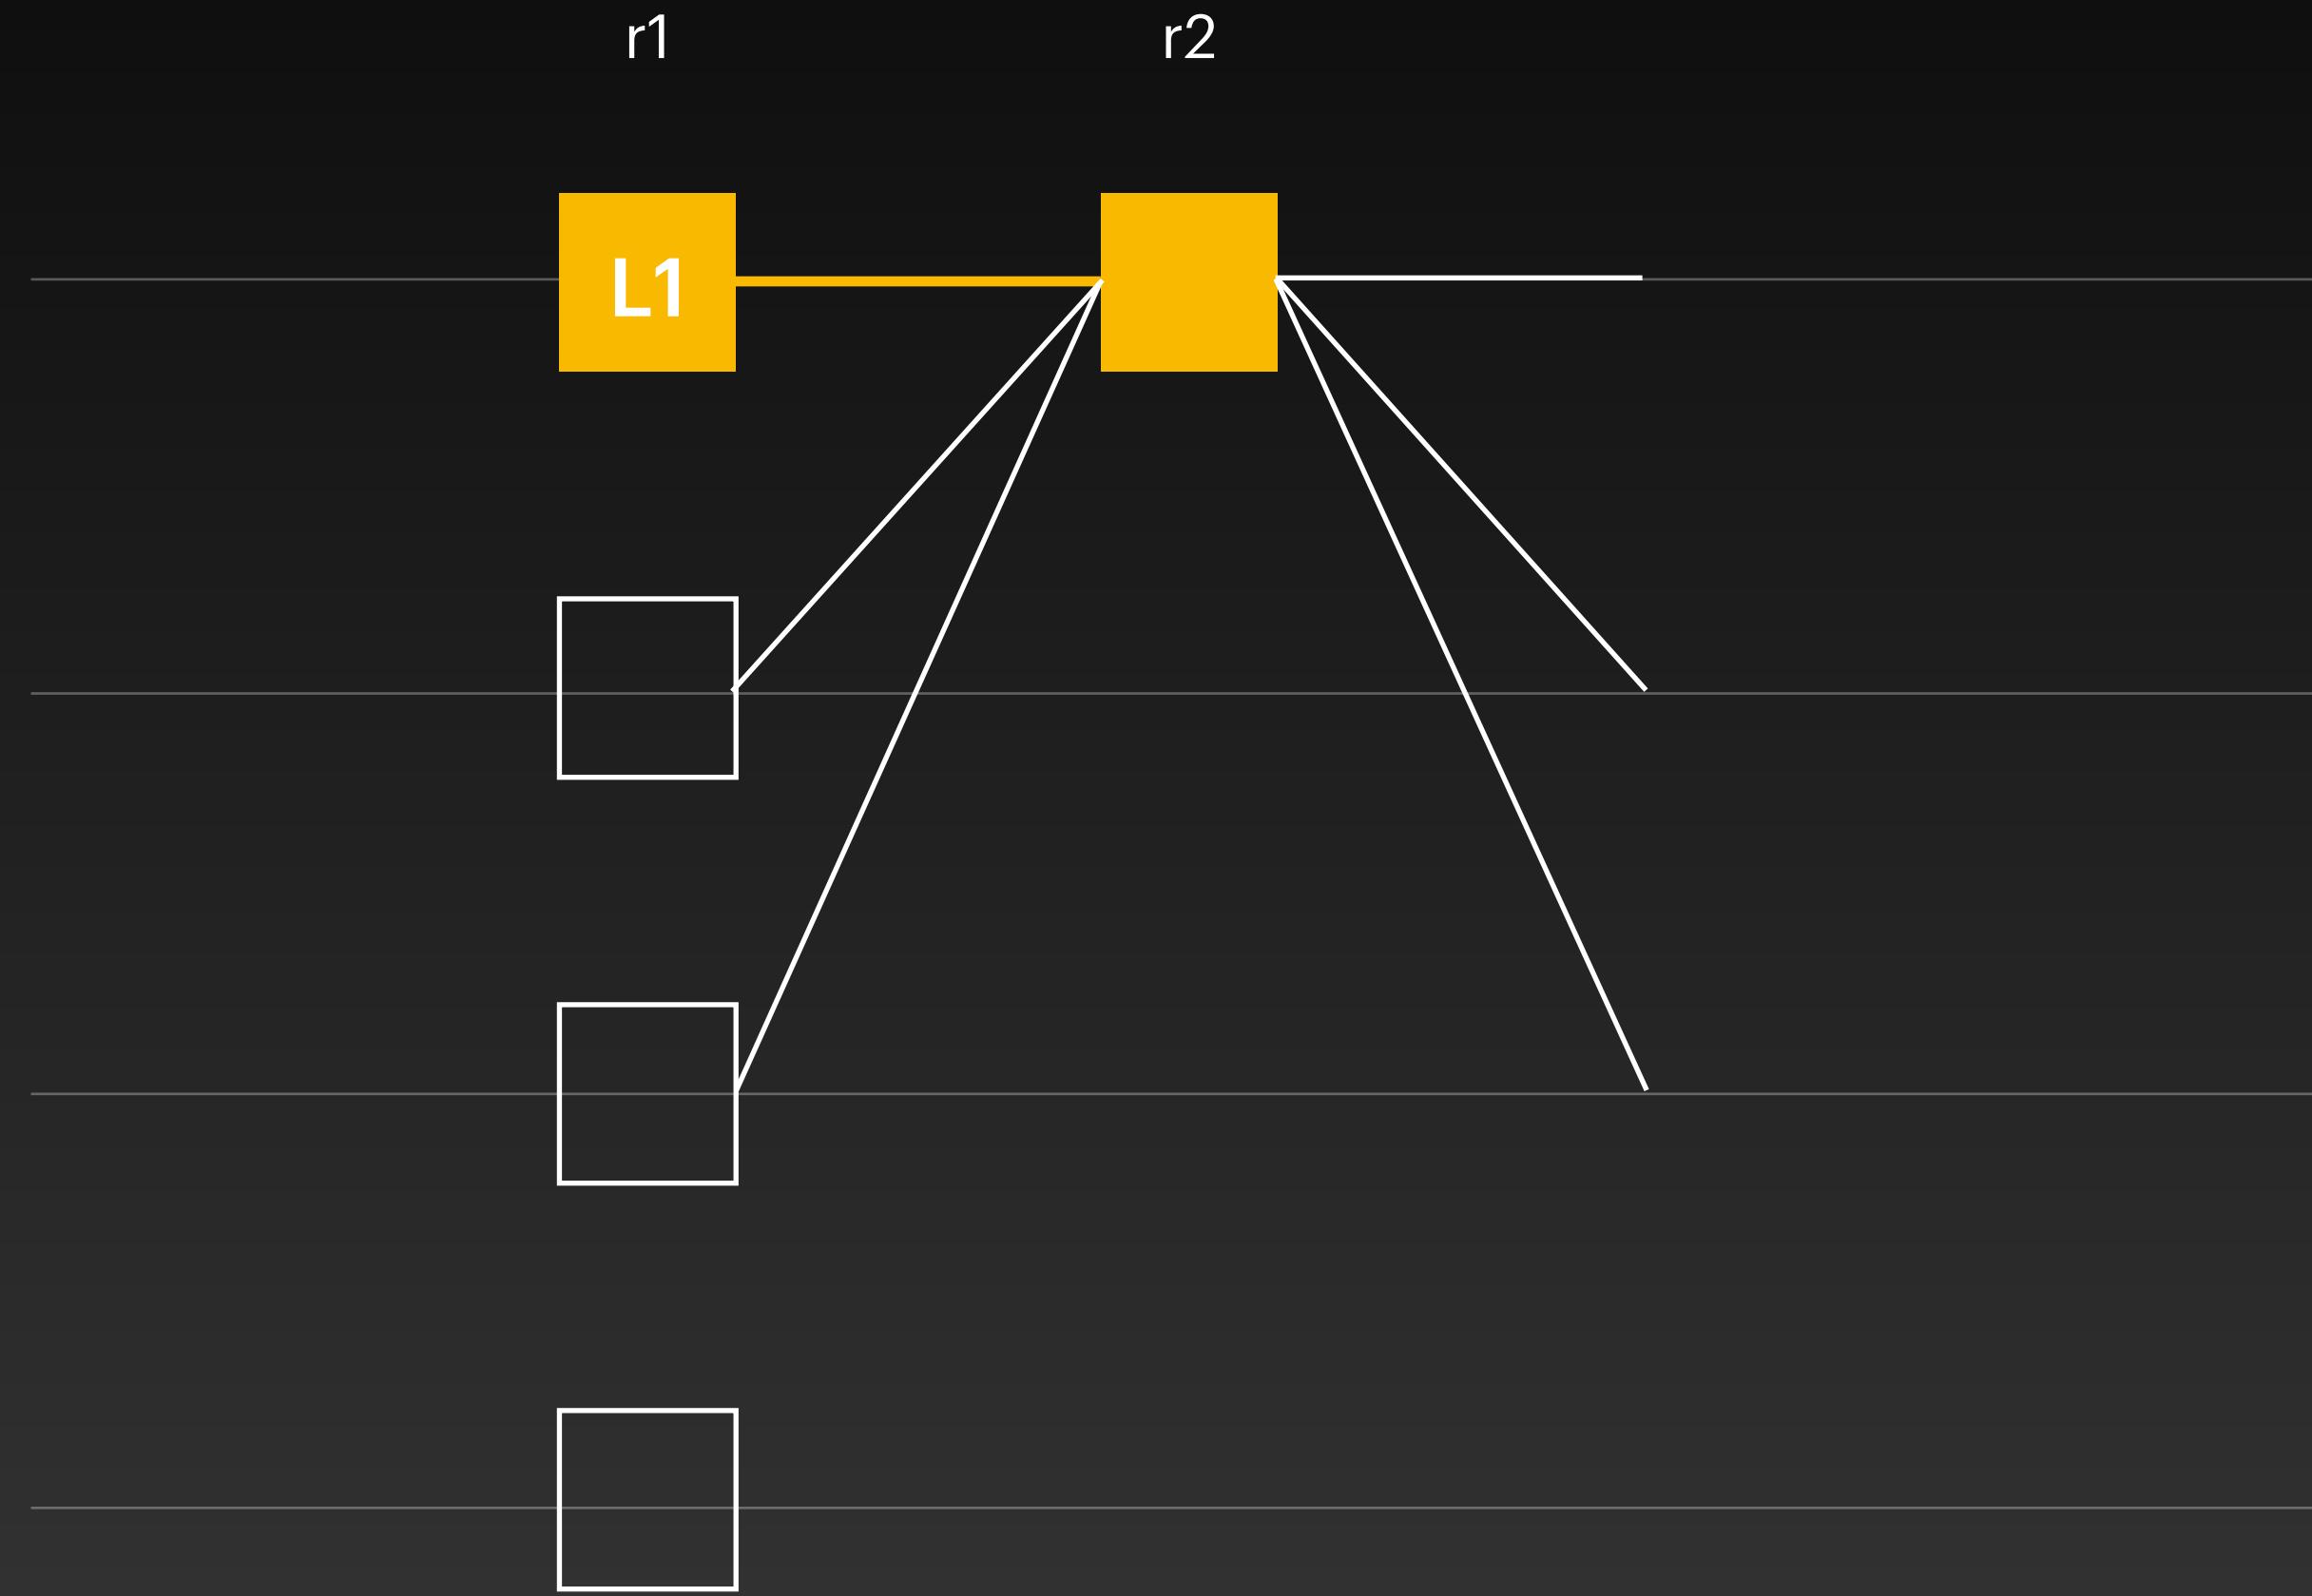




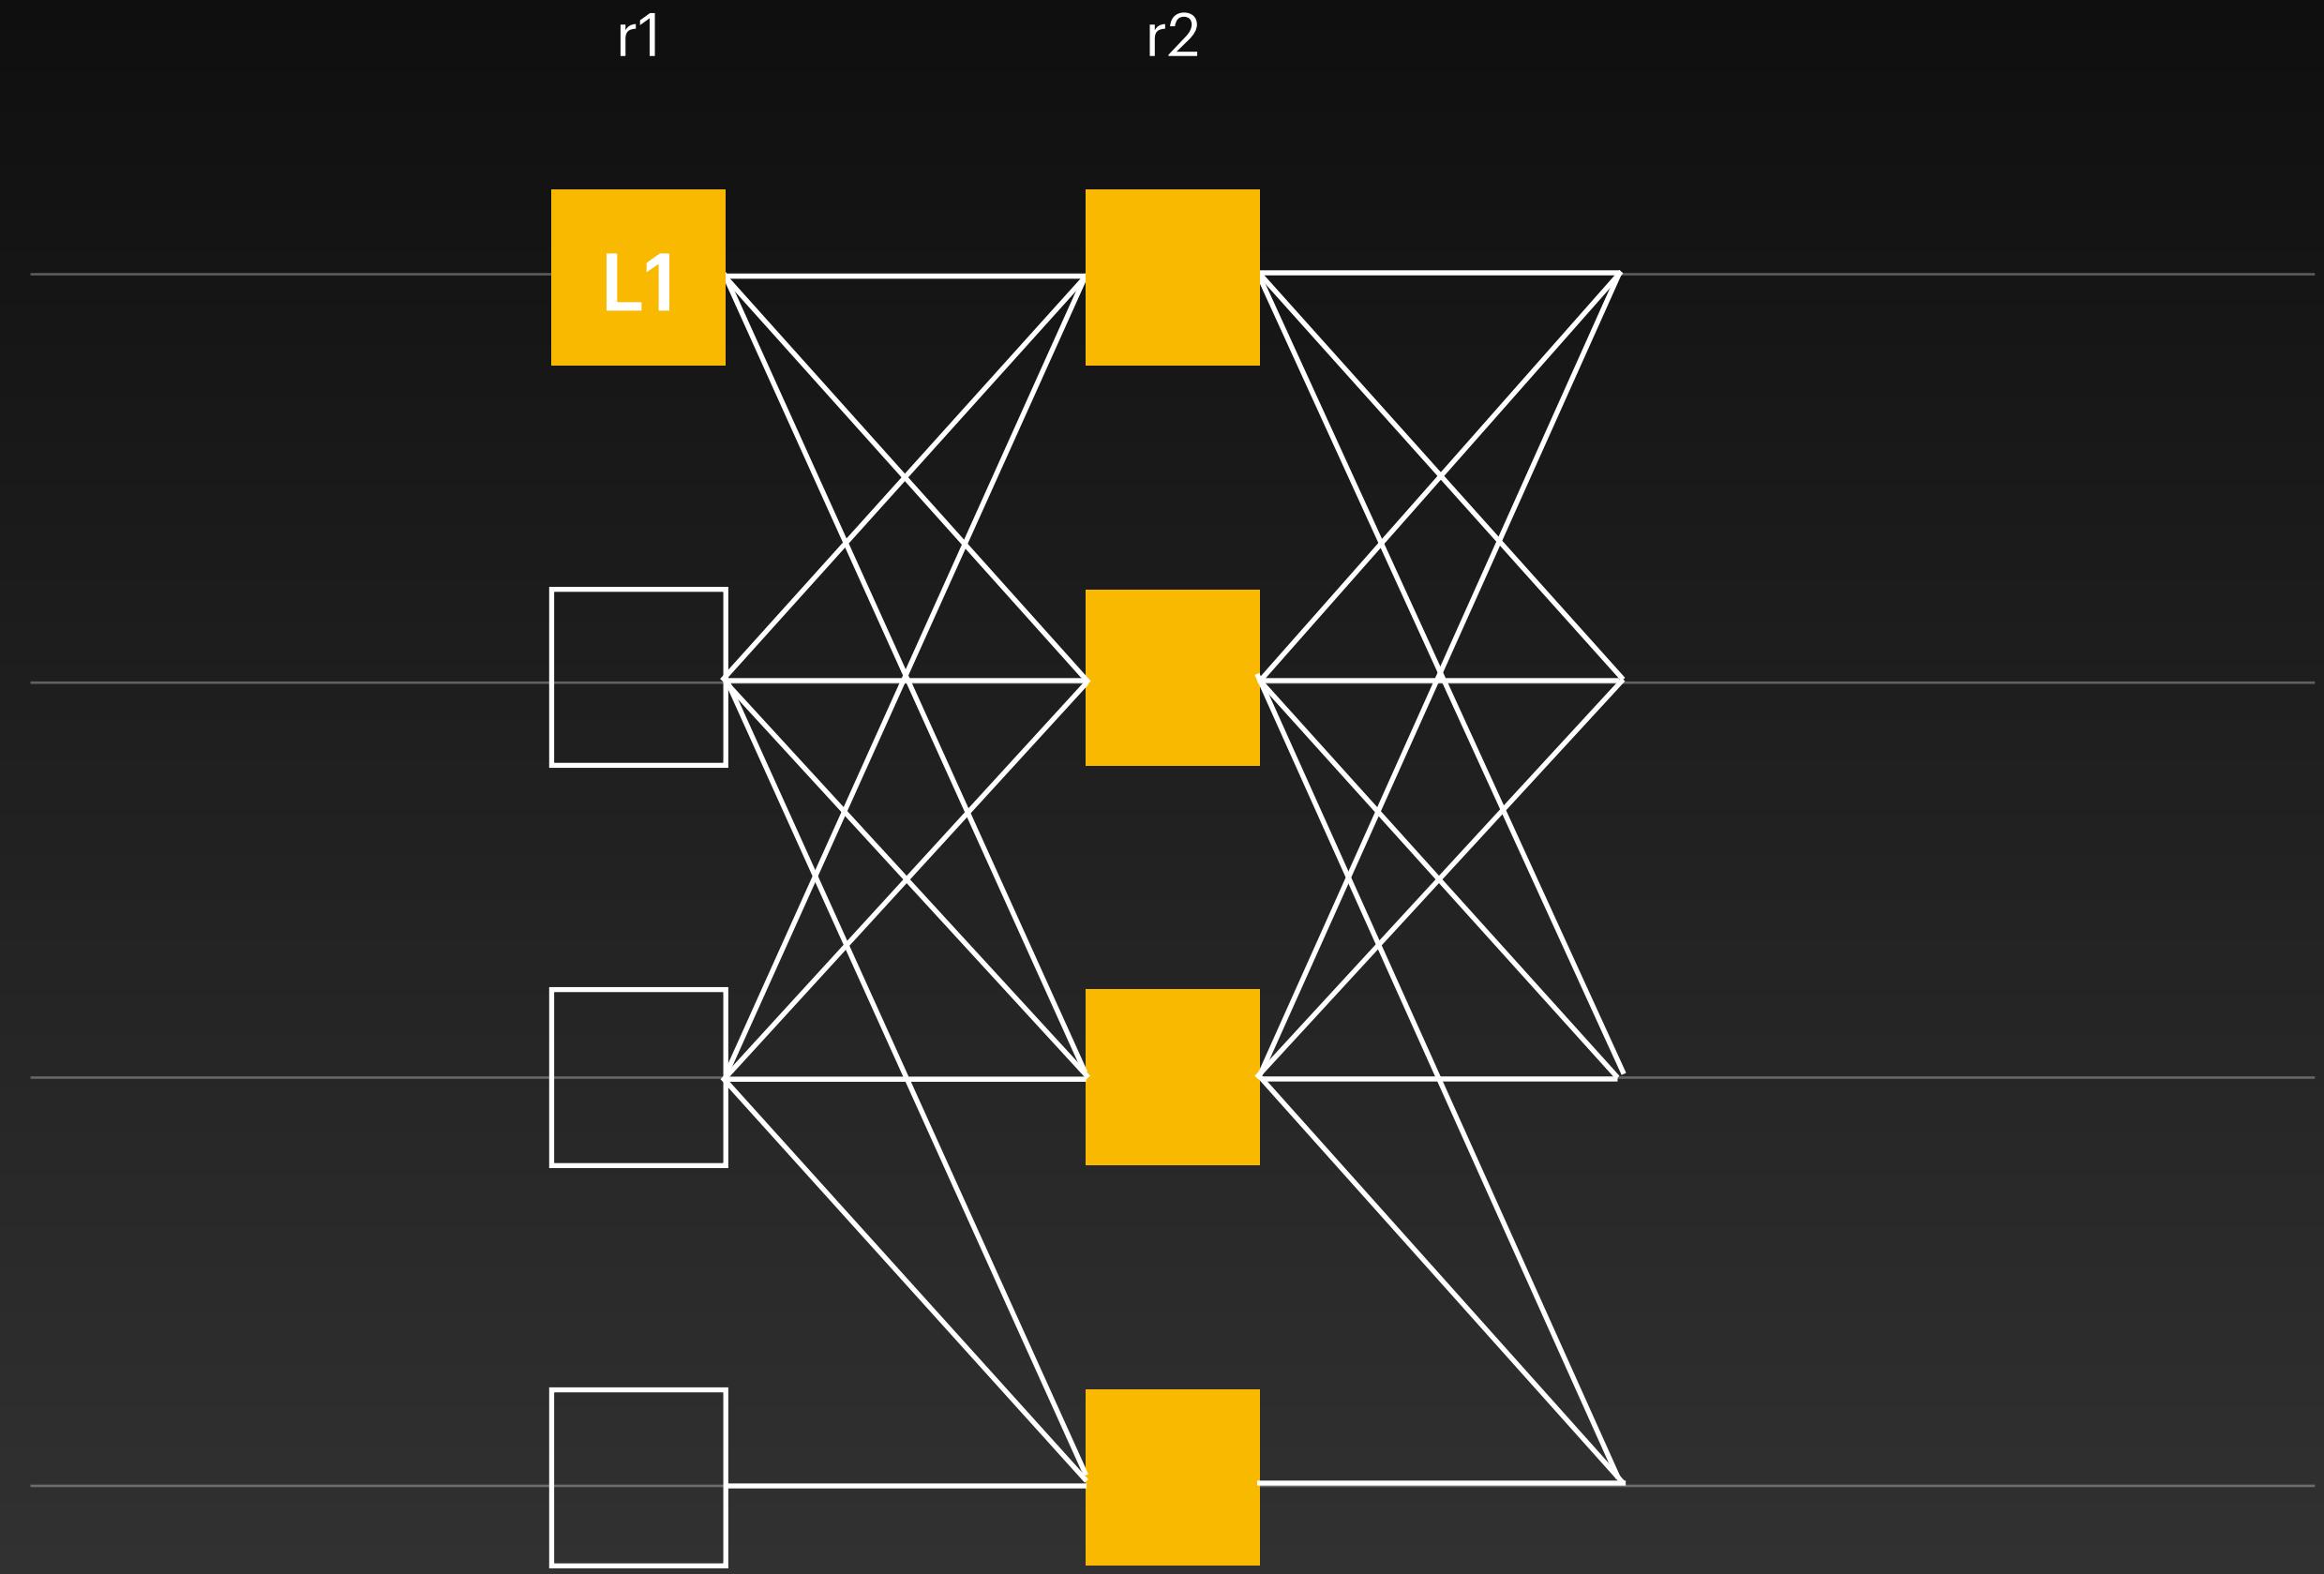
- Round number
- Author
- Payload (transactions)
- Signature



- Link to previous blocks



- Wait for the leader

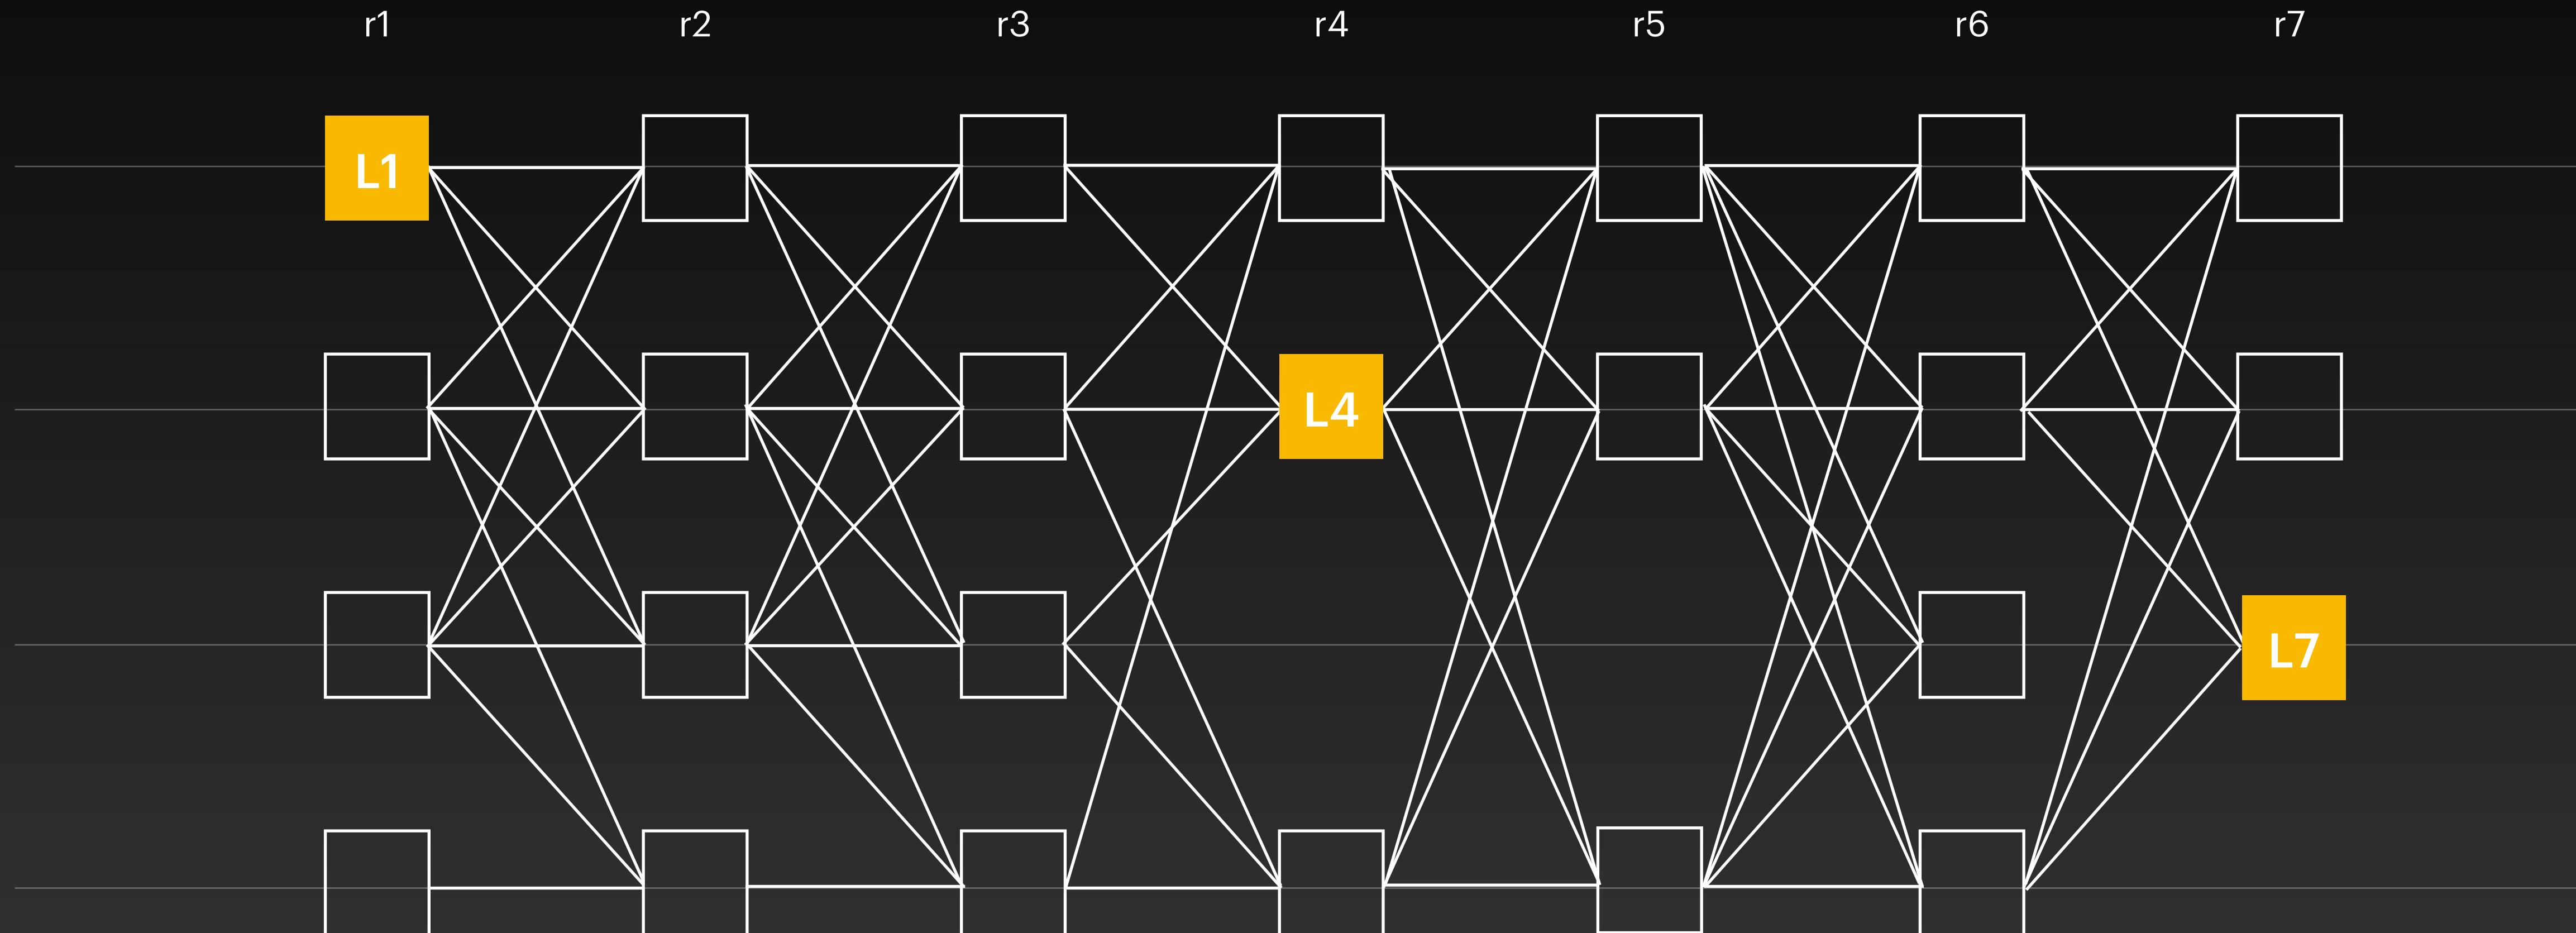


- All validators run in parallel

# Network Security

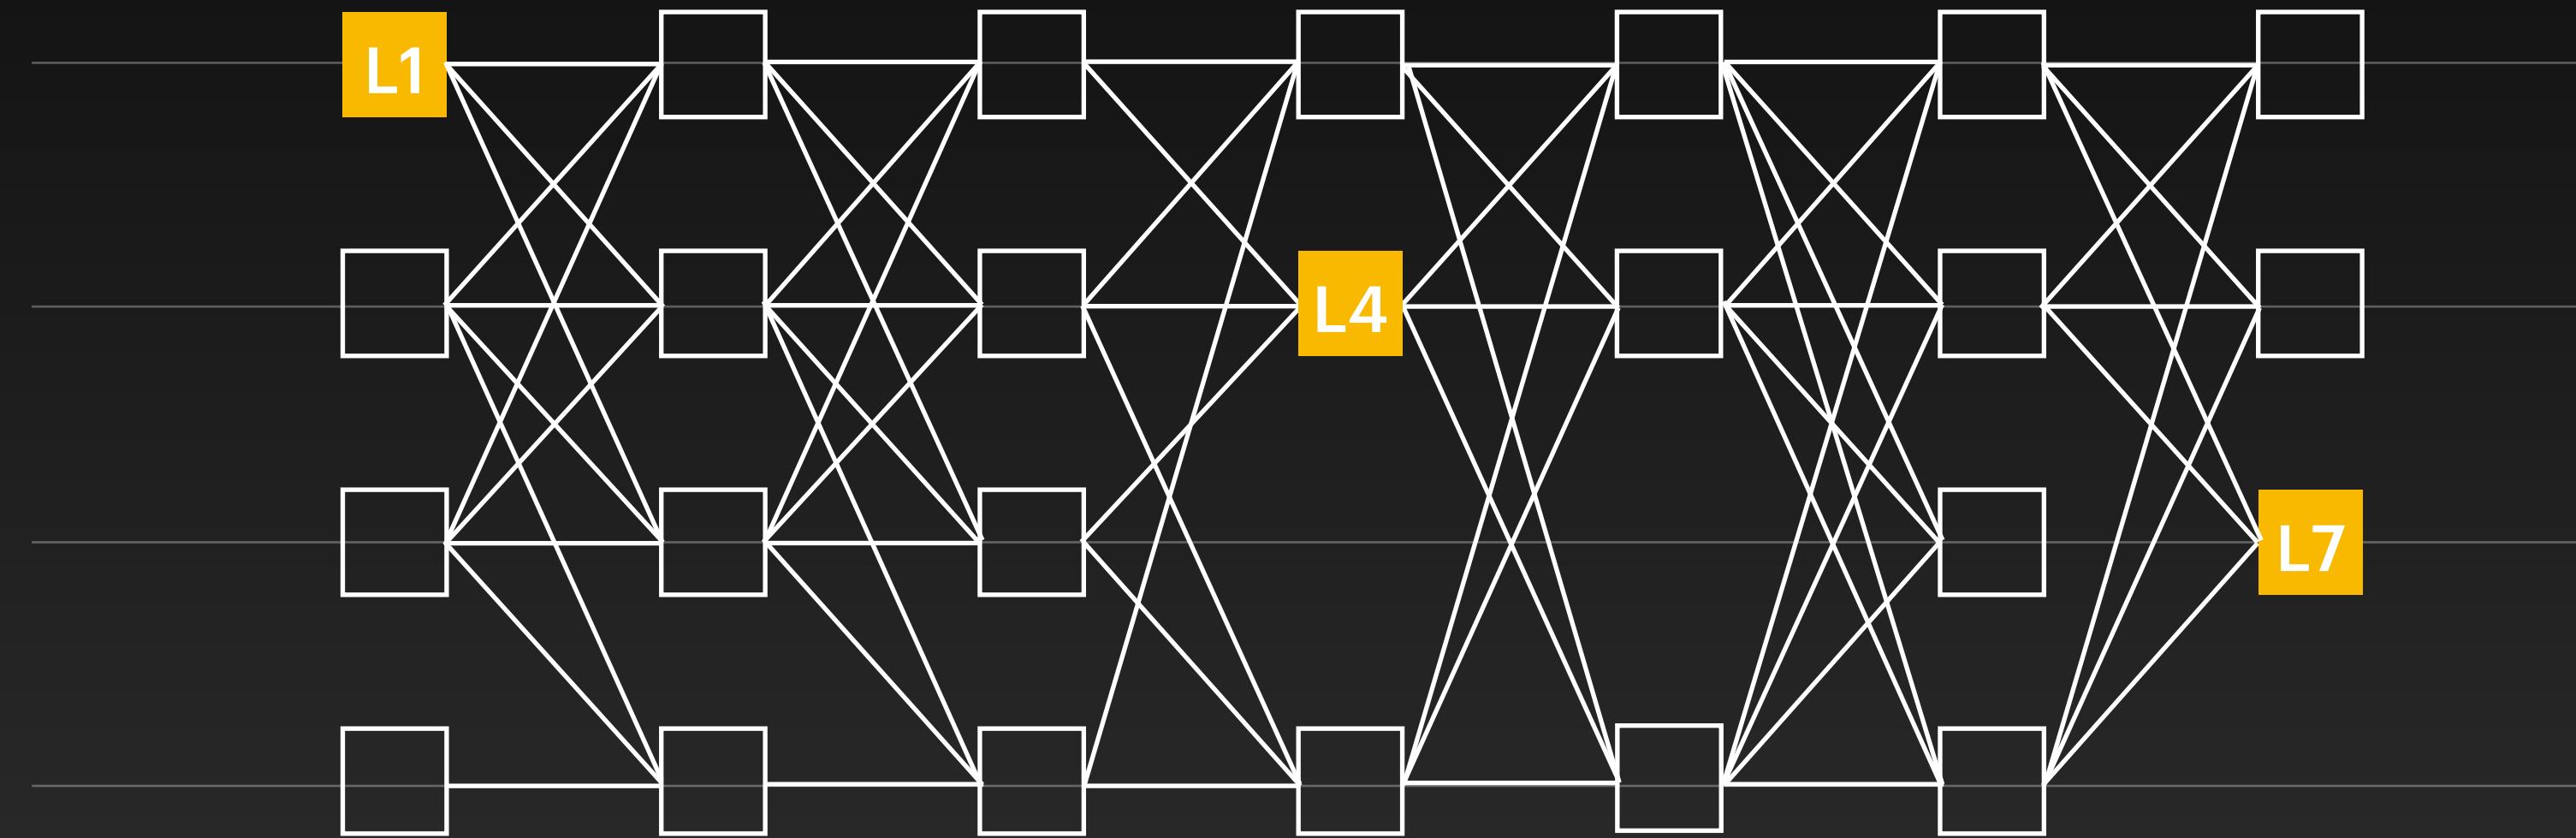
## Challenge #4: Ordering

- **How to find the best path to send the block to another node?**



# End Goal

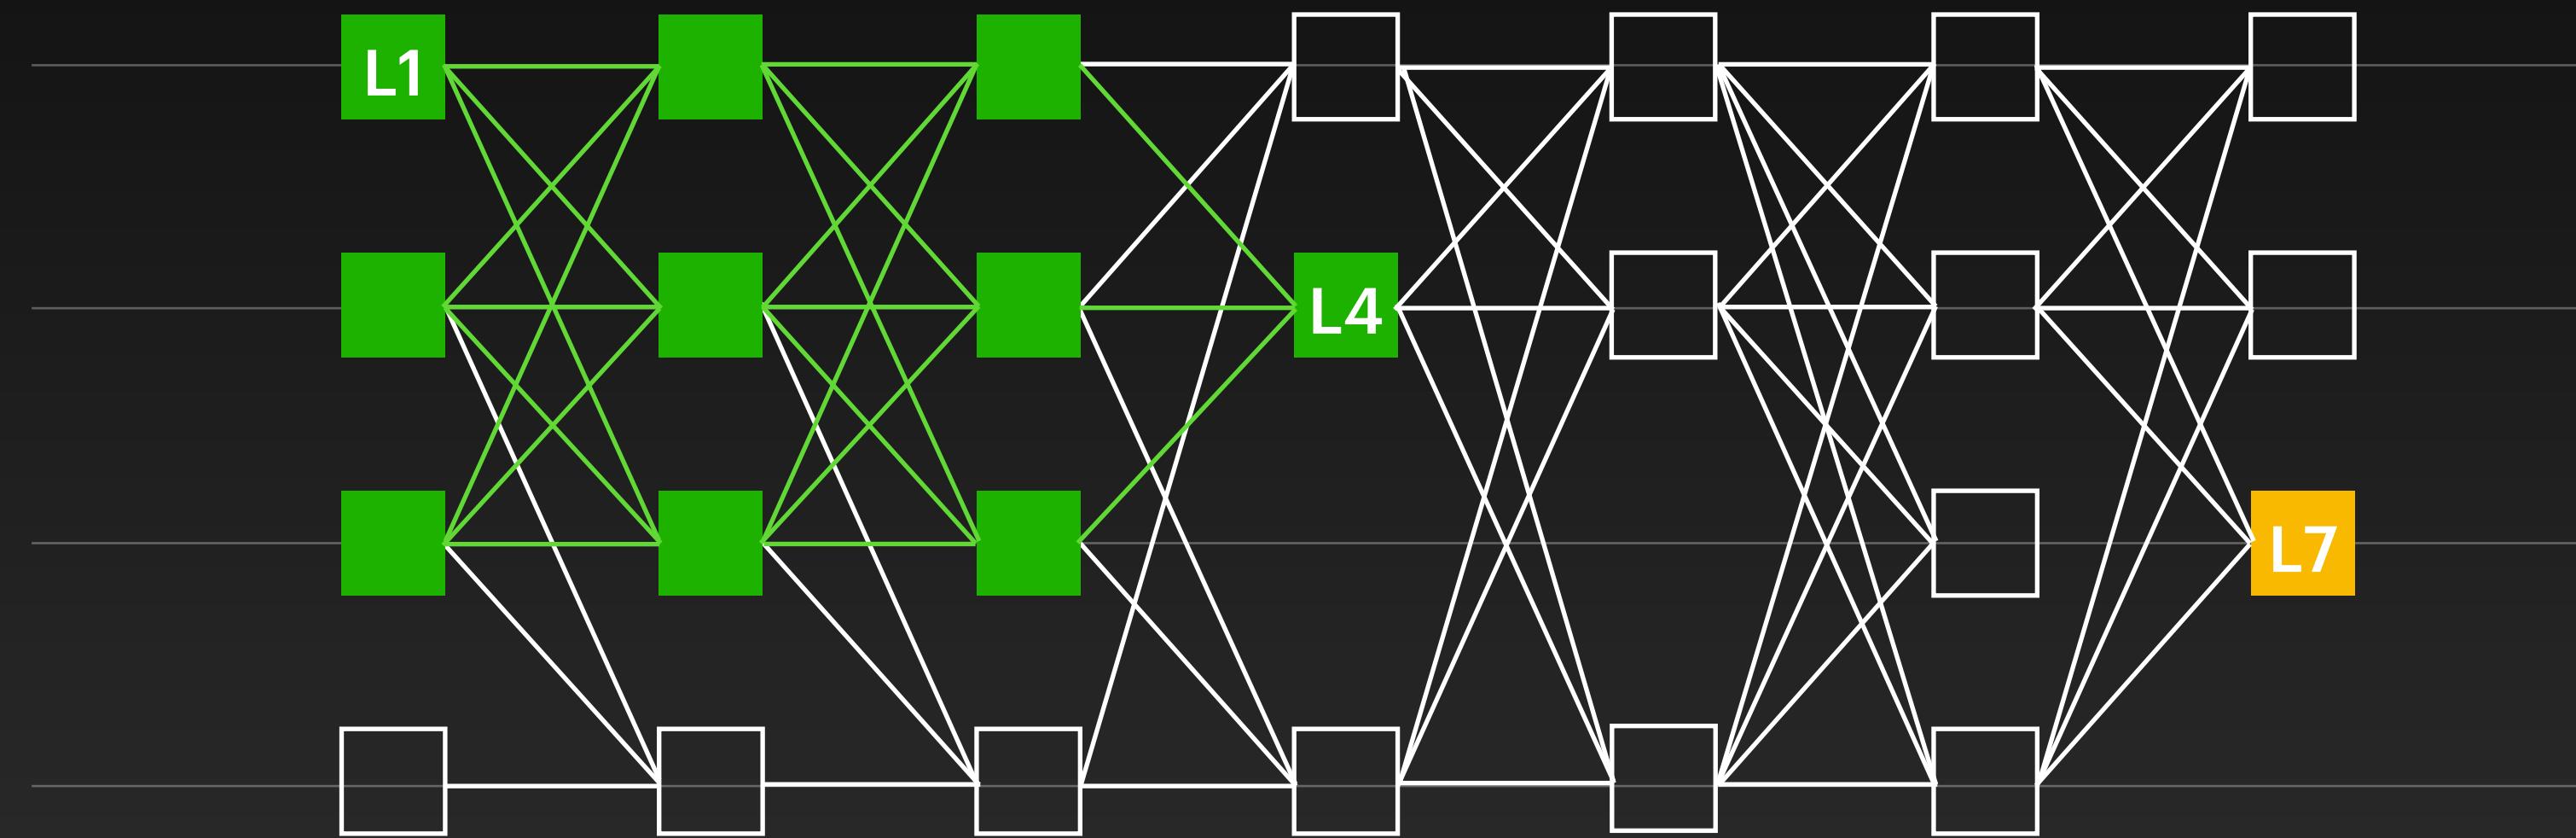
## Ordering leaders



- We focus on ordering leaders: L1 L4 L7

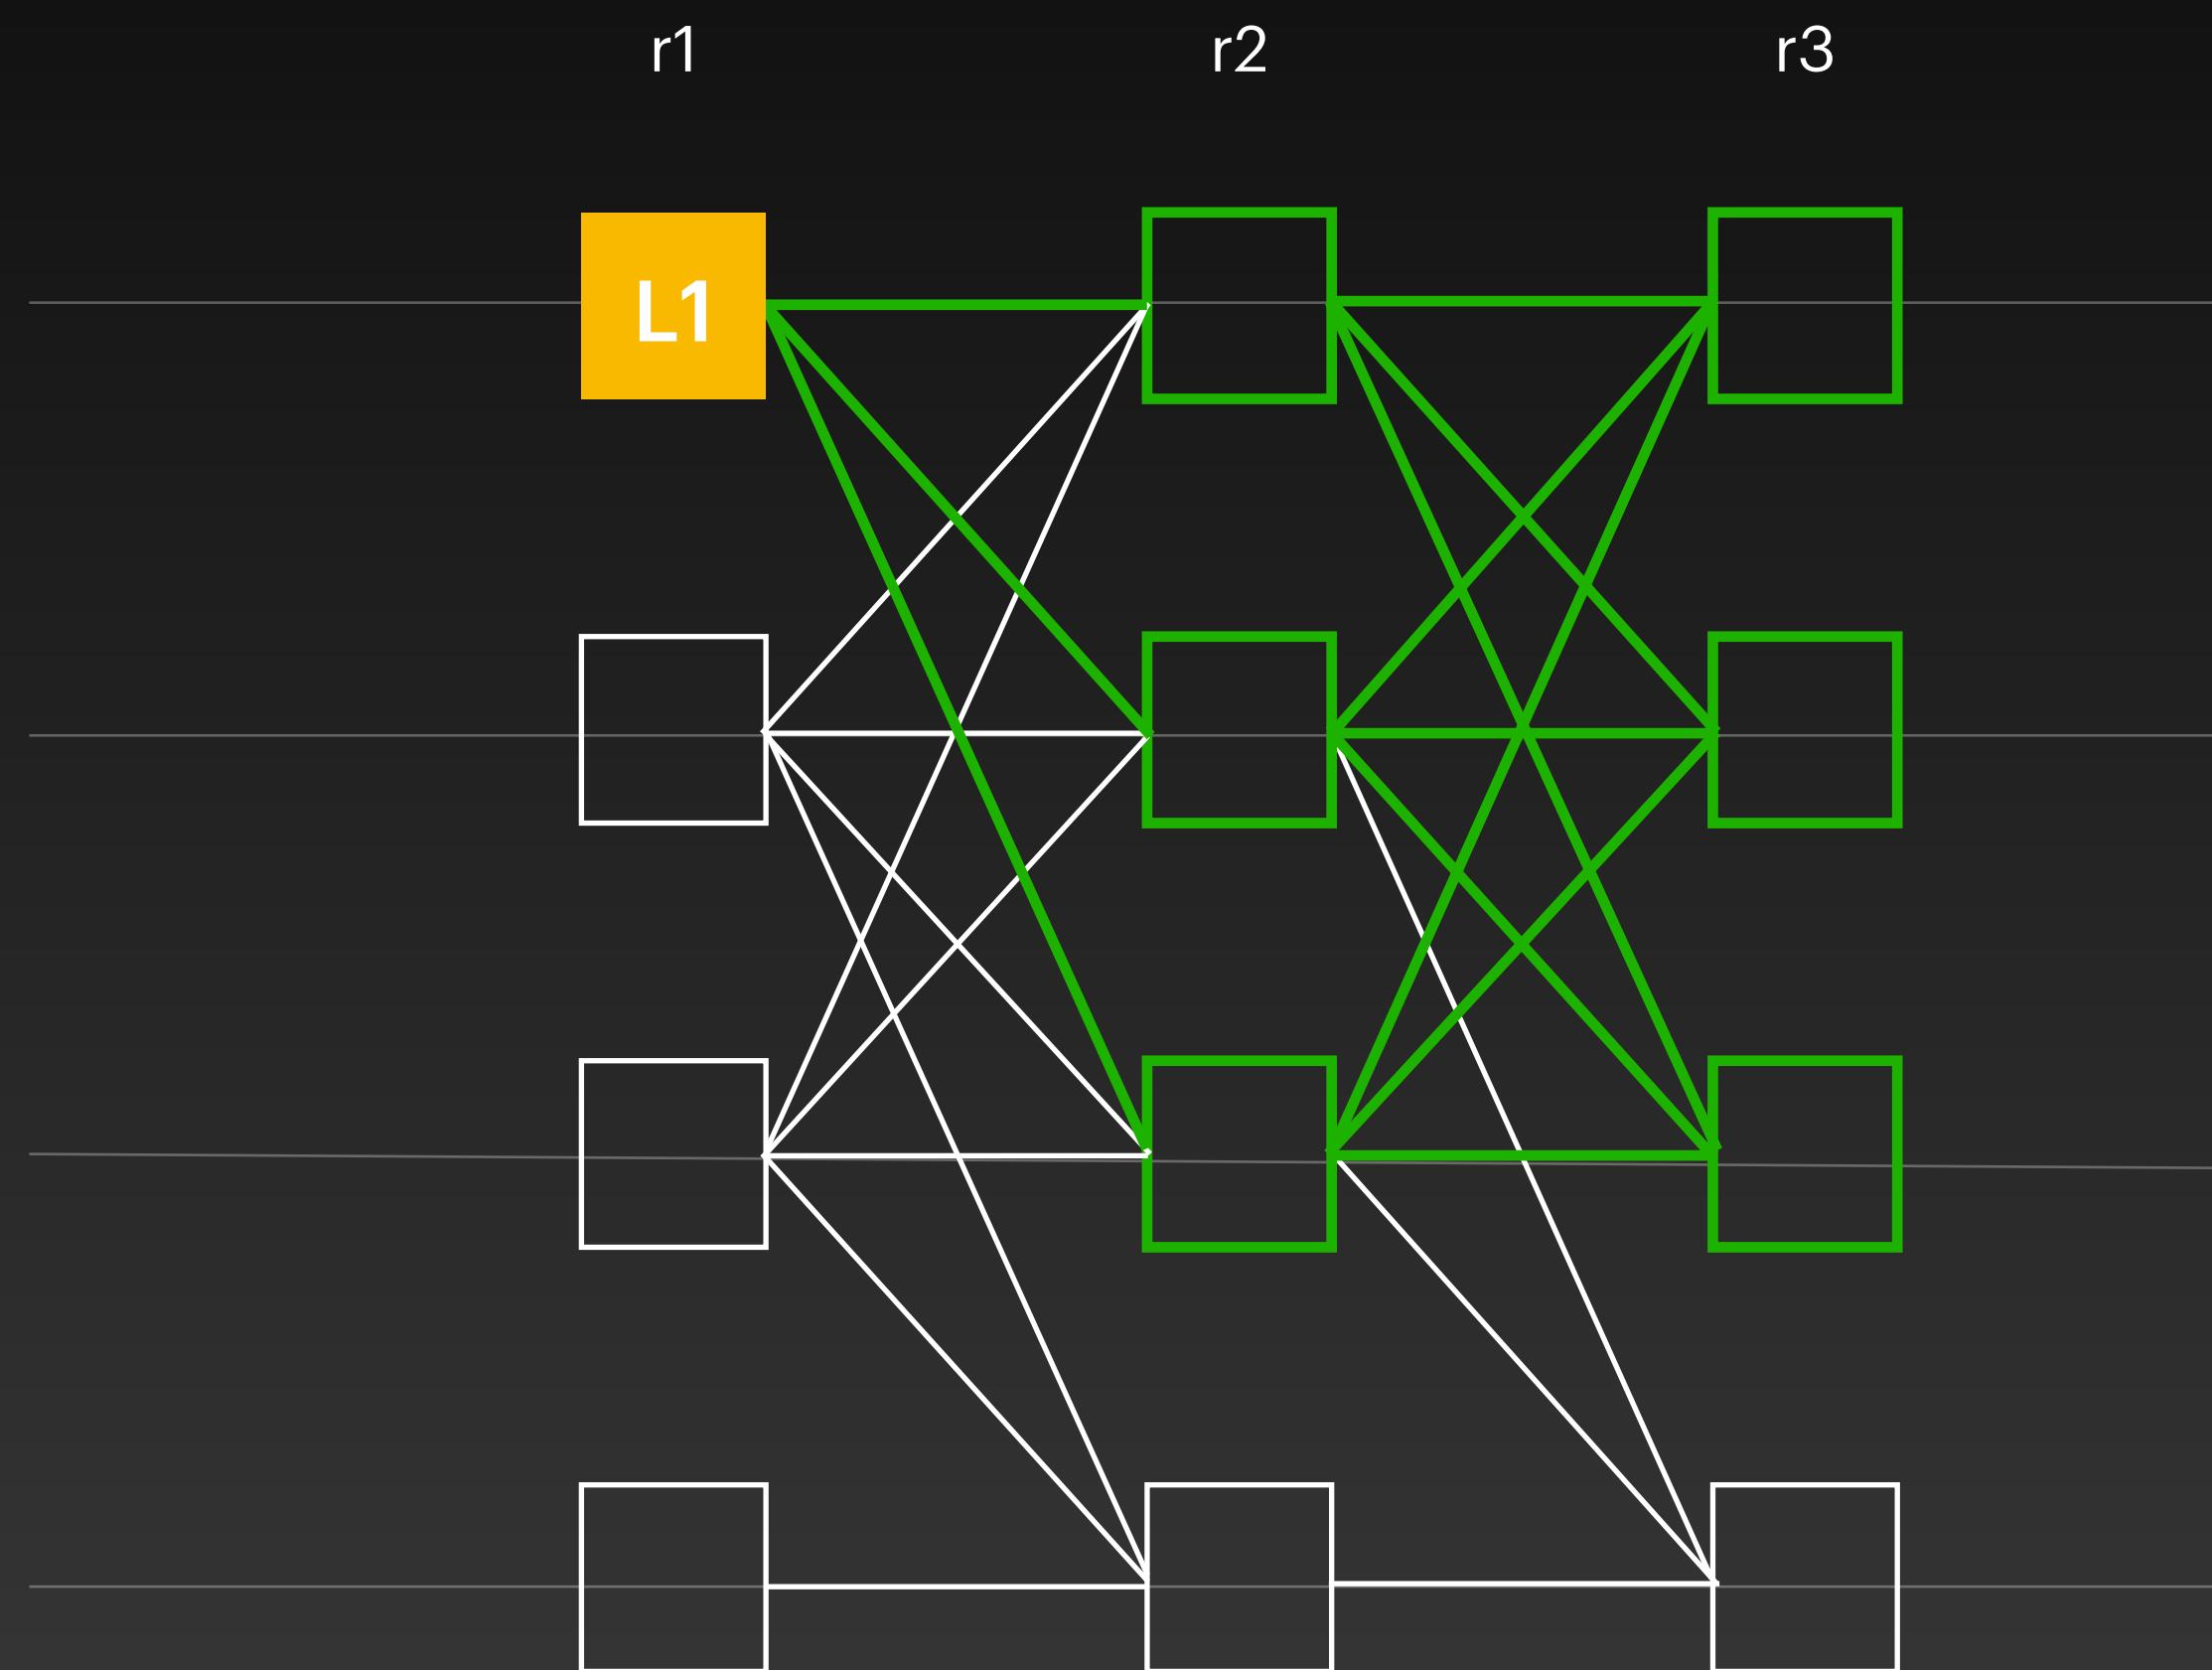
# End Goal

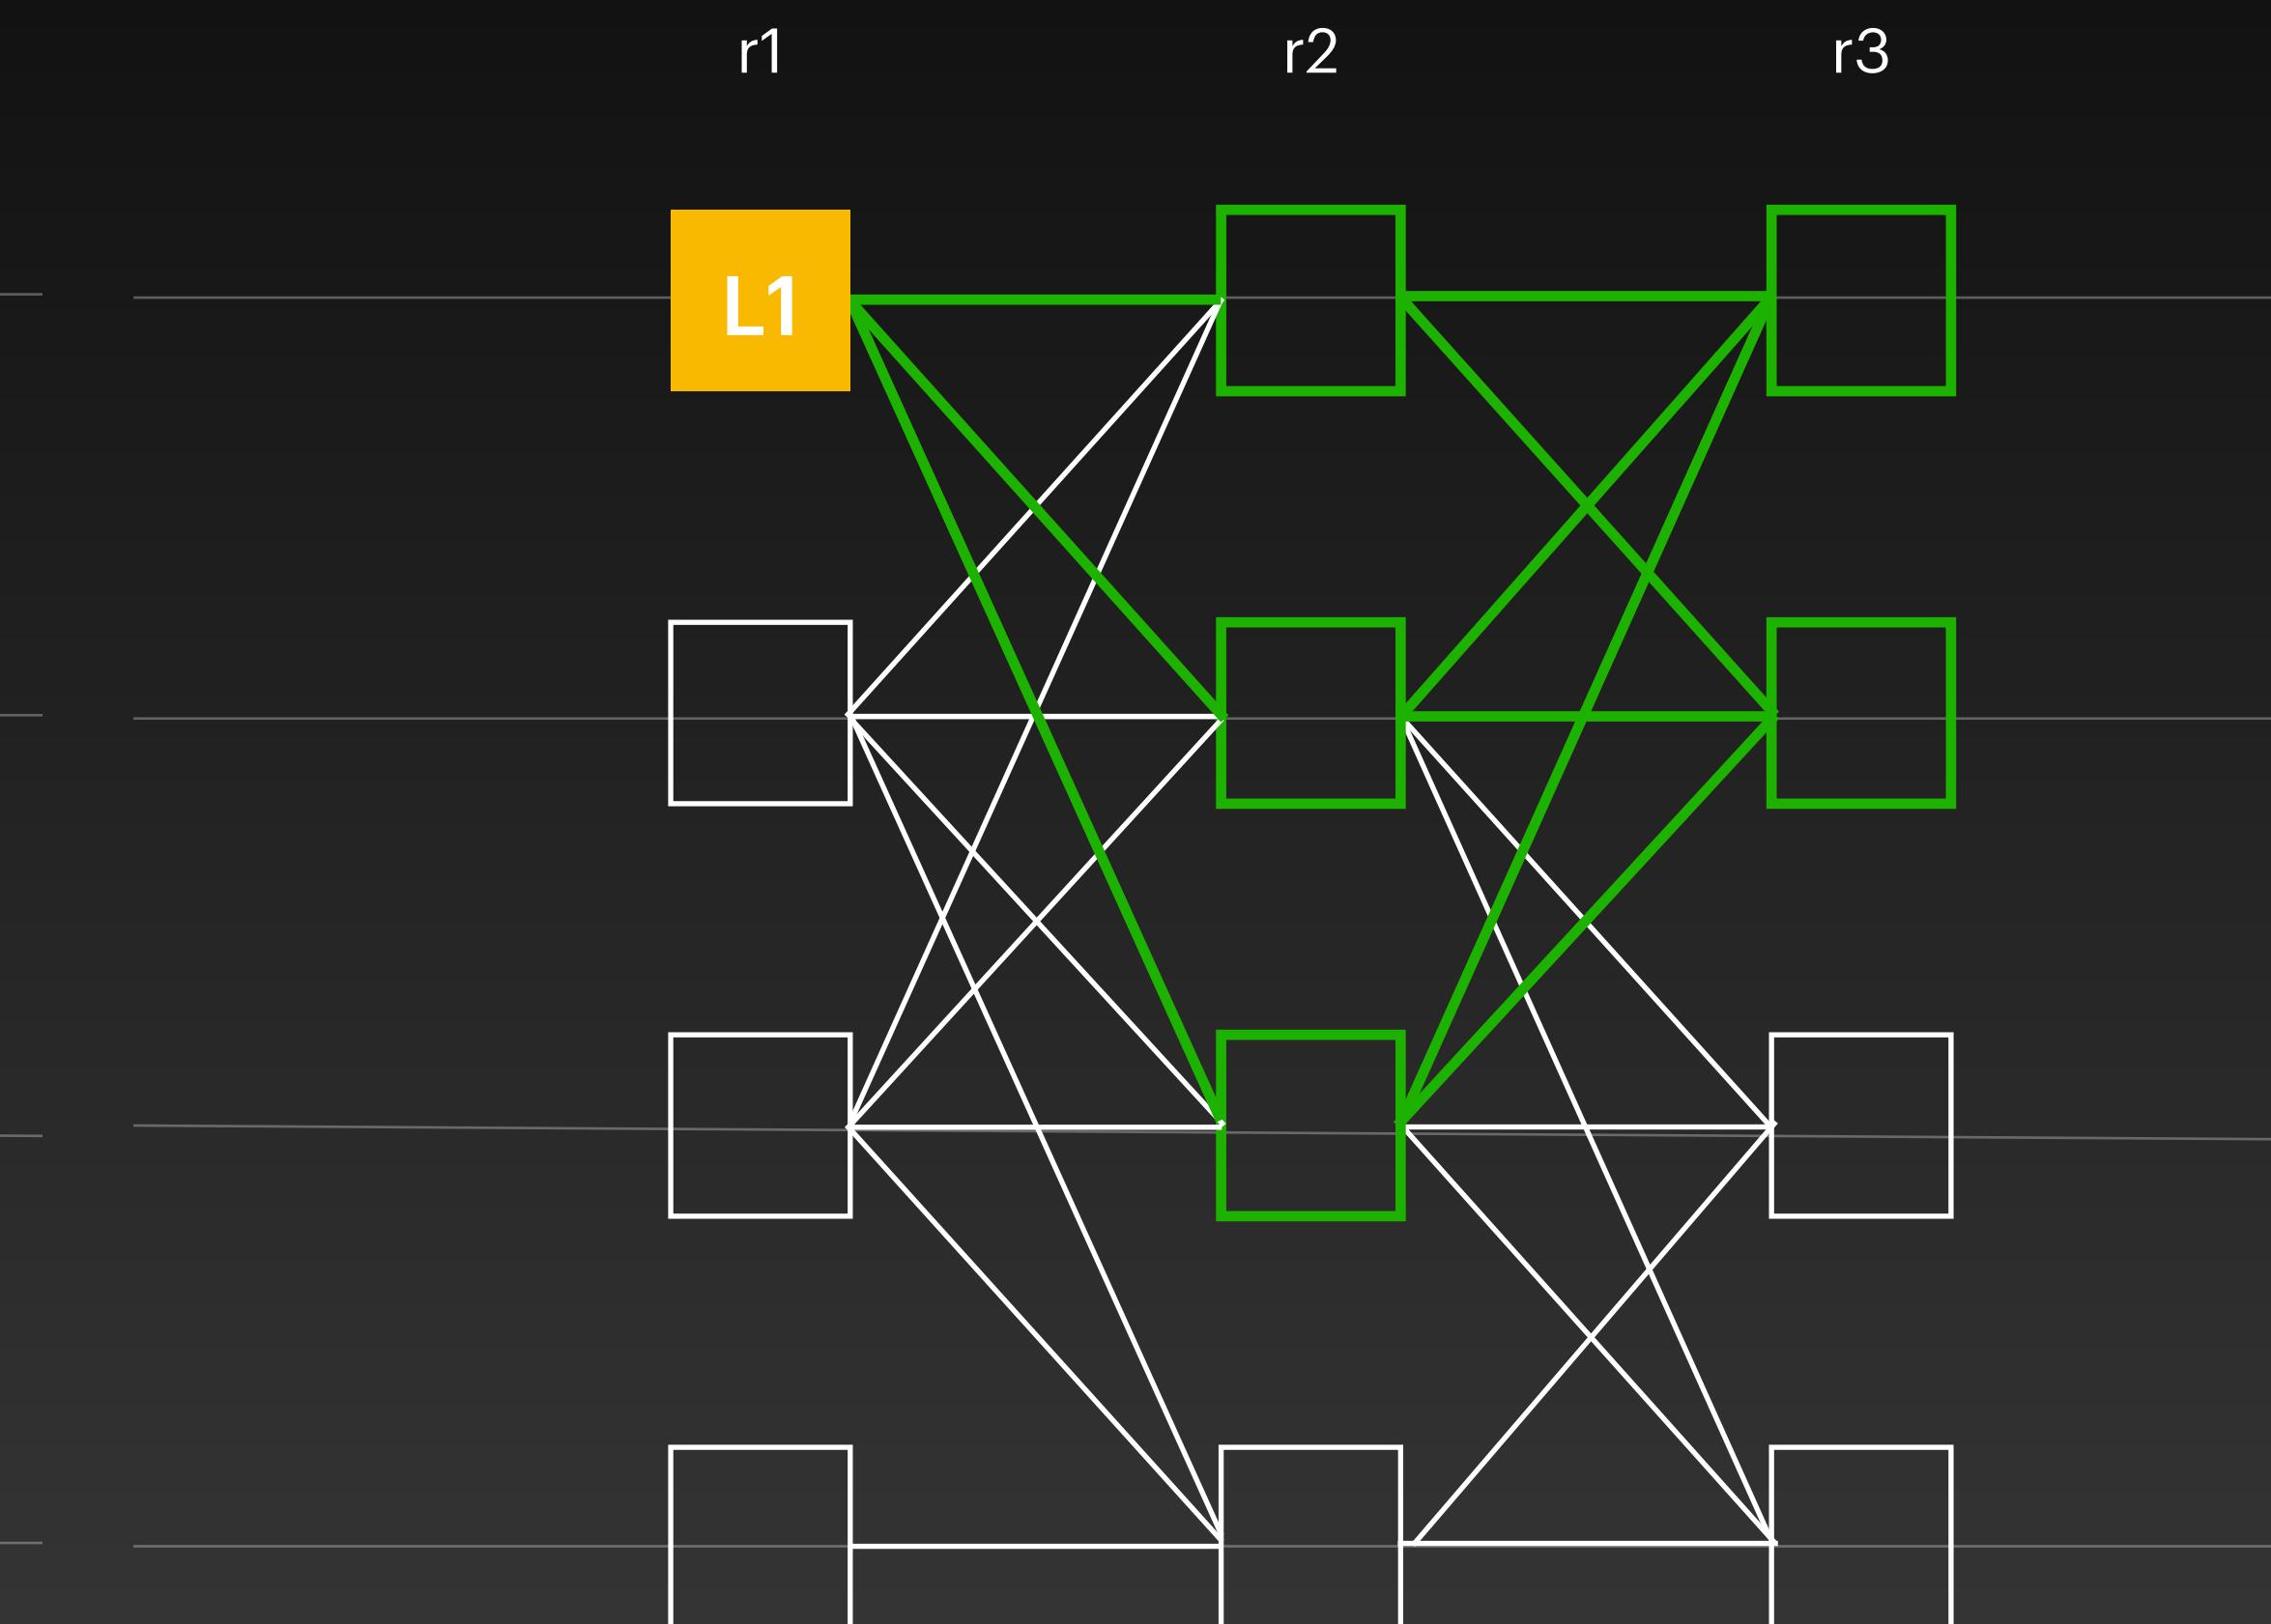
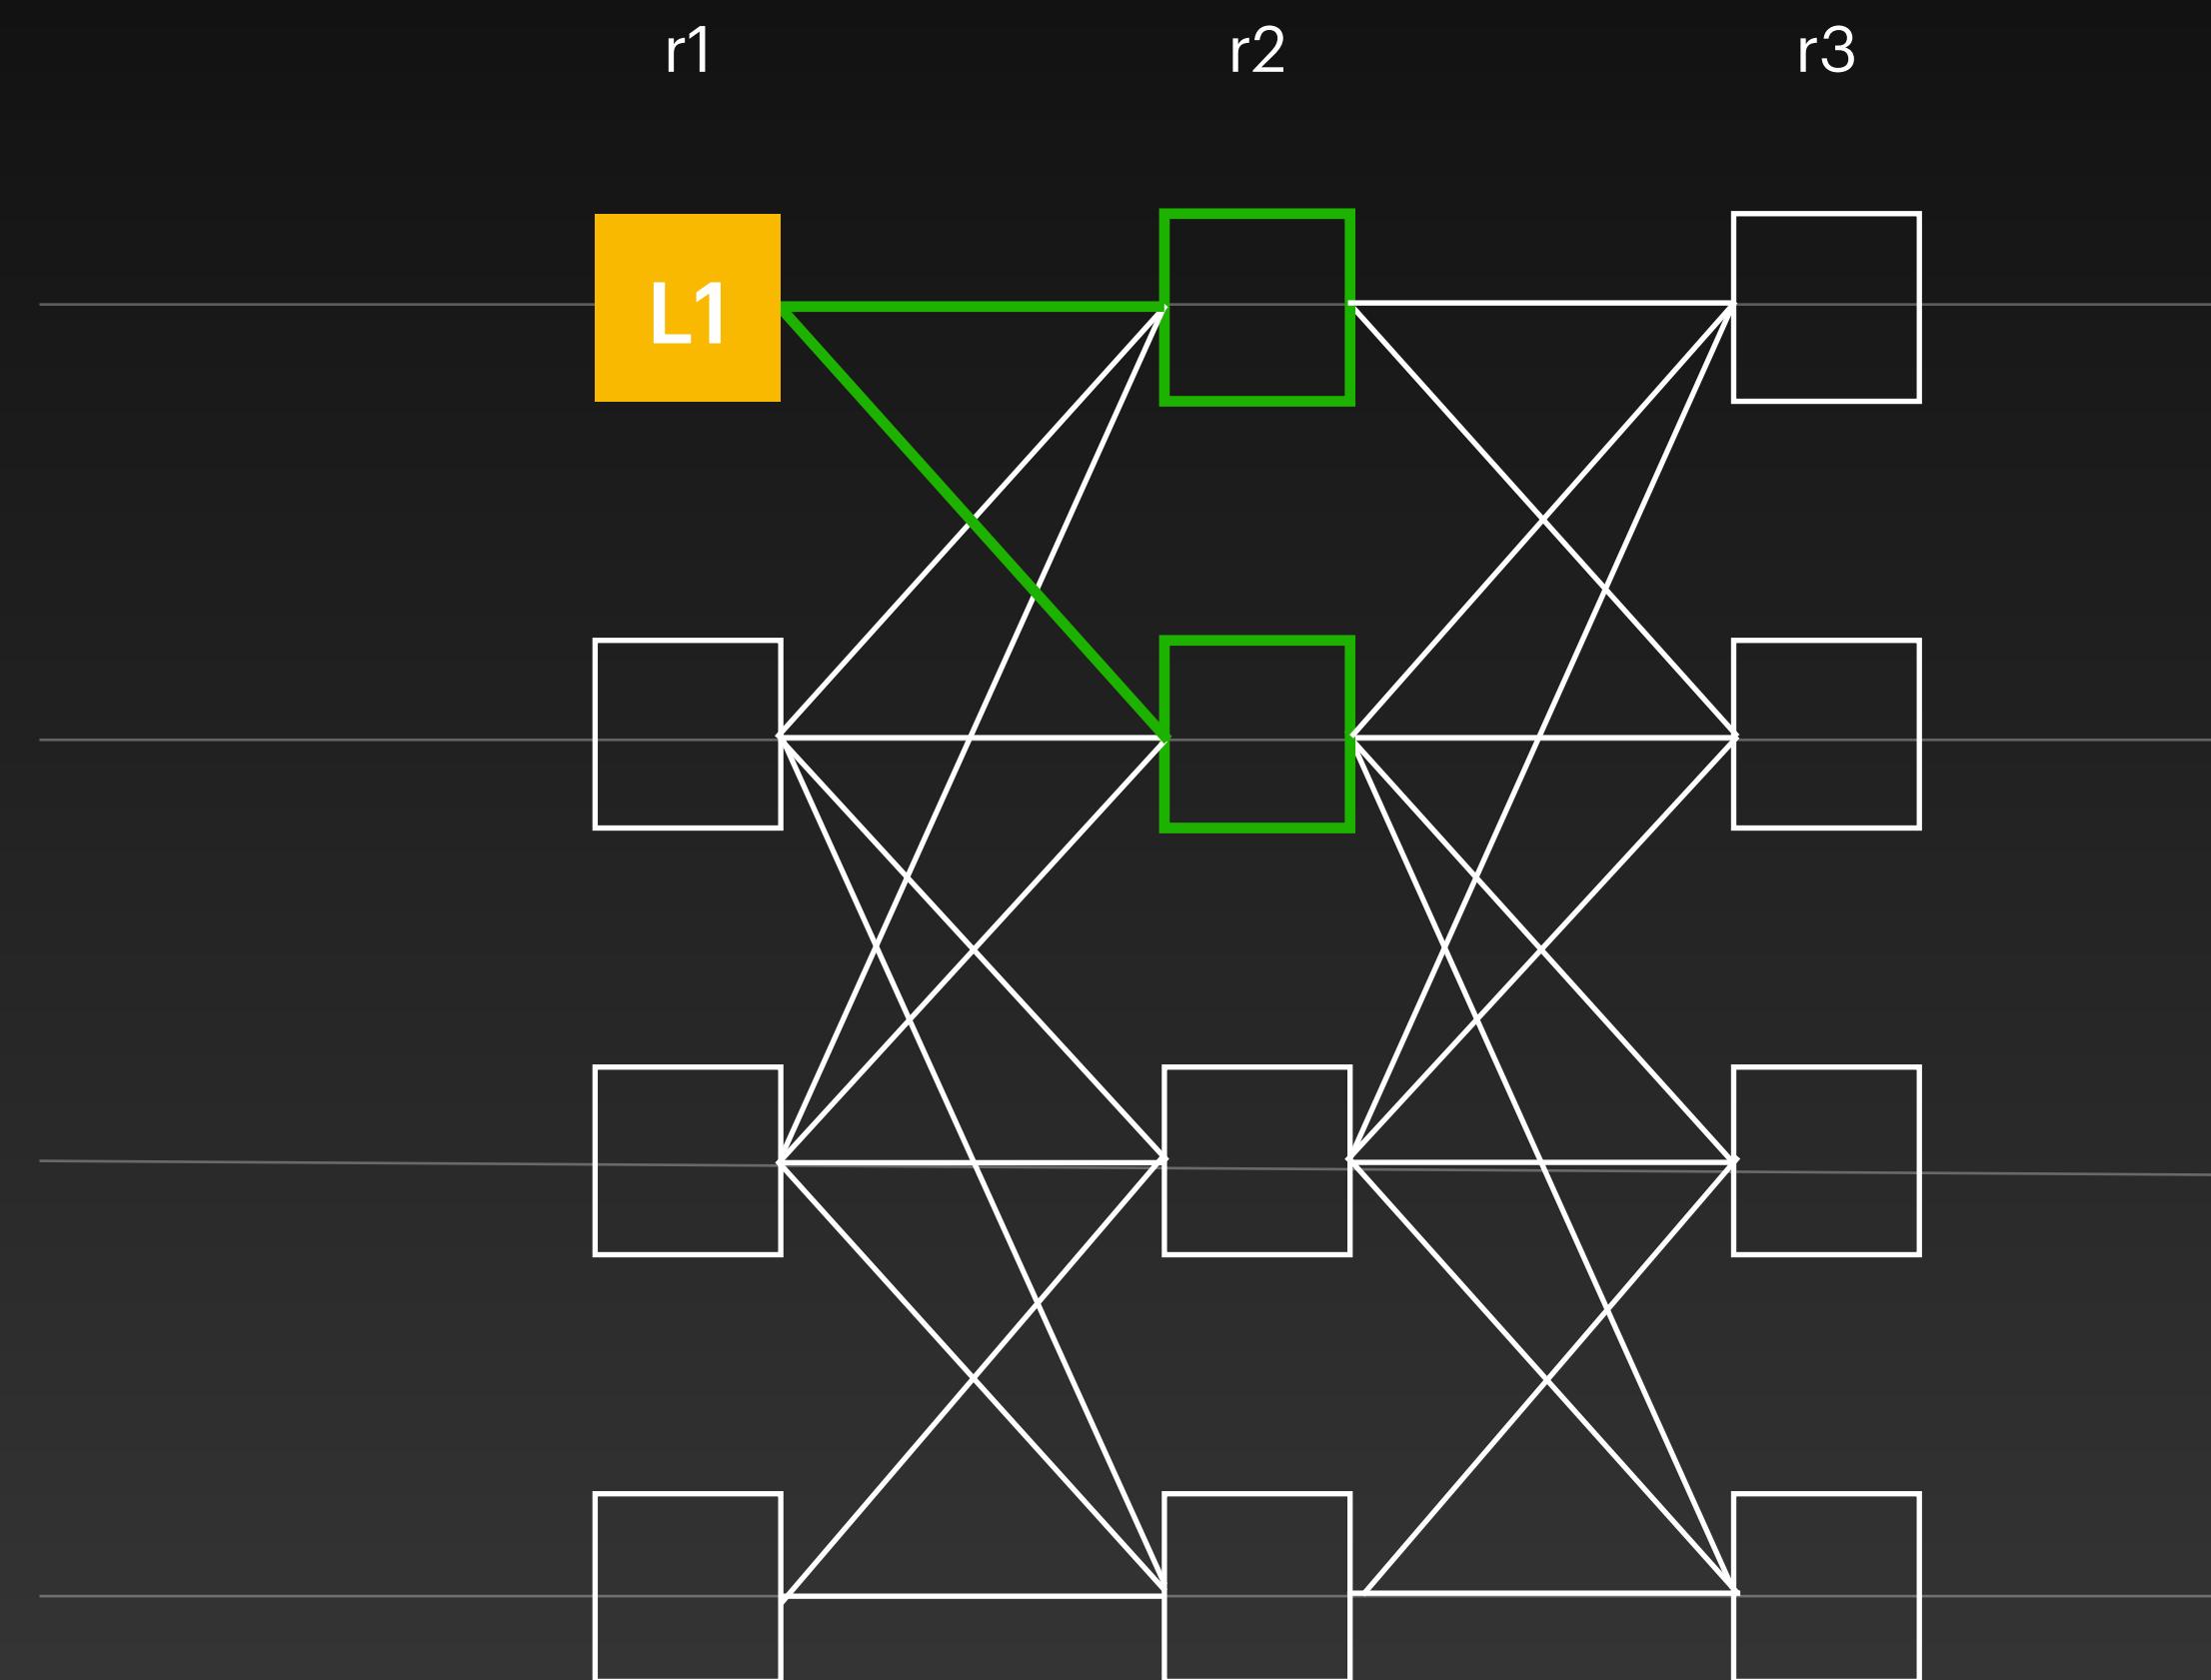
## Ordering leaders



- We focus on ordering leaders: **L1** **L4** **L7**
- Linearising the sub-DAG is simple

# How is it done?



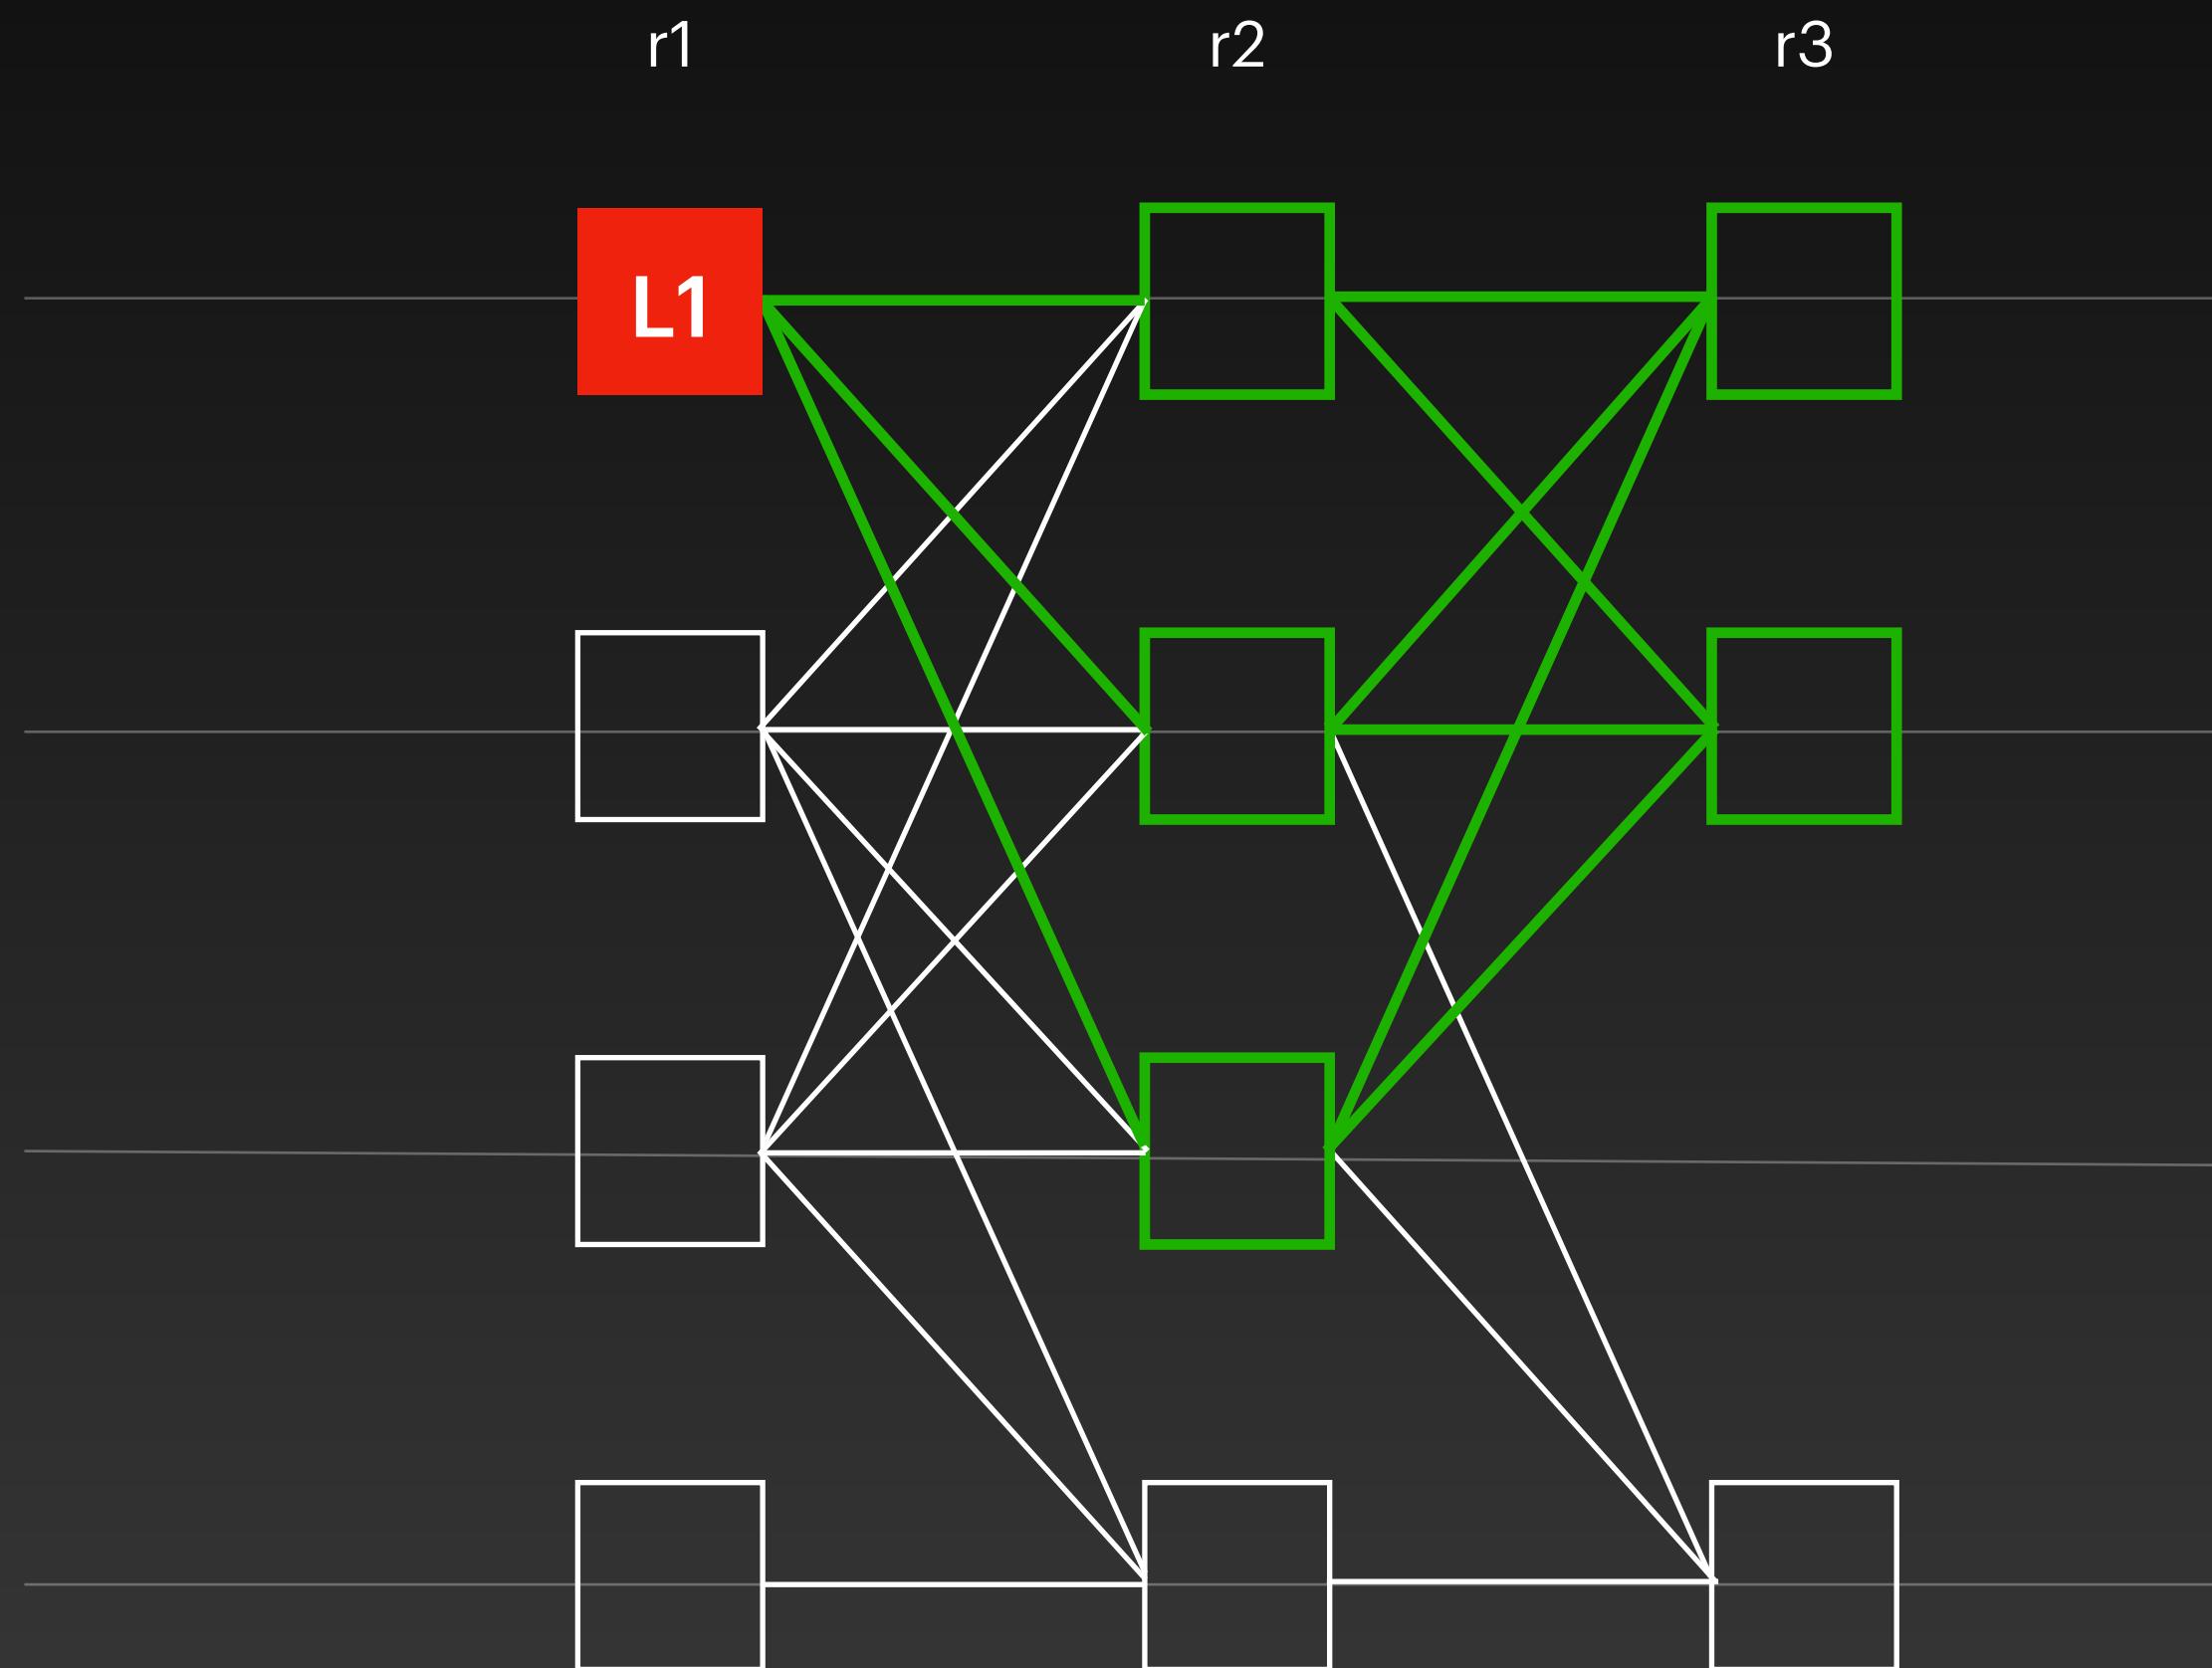


# Network Security

## Challenge #4: Ordering

- How to find the best path to send the block to another node?
- **DoS against the leader are particularly effective**

# Message not received in order?



- Bad leader?
- Or bad network?

# Network Security

## Challenge #4: Ordering

- How to find the best path to send the block to another node?
- DoS against the leader are particularly effective
- **Reordering messages causes massive slowdowns**

# Network Security

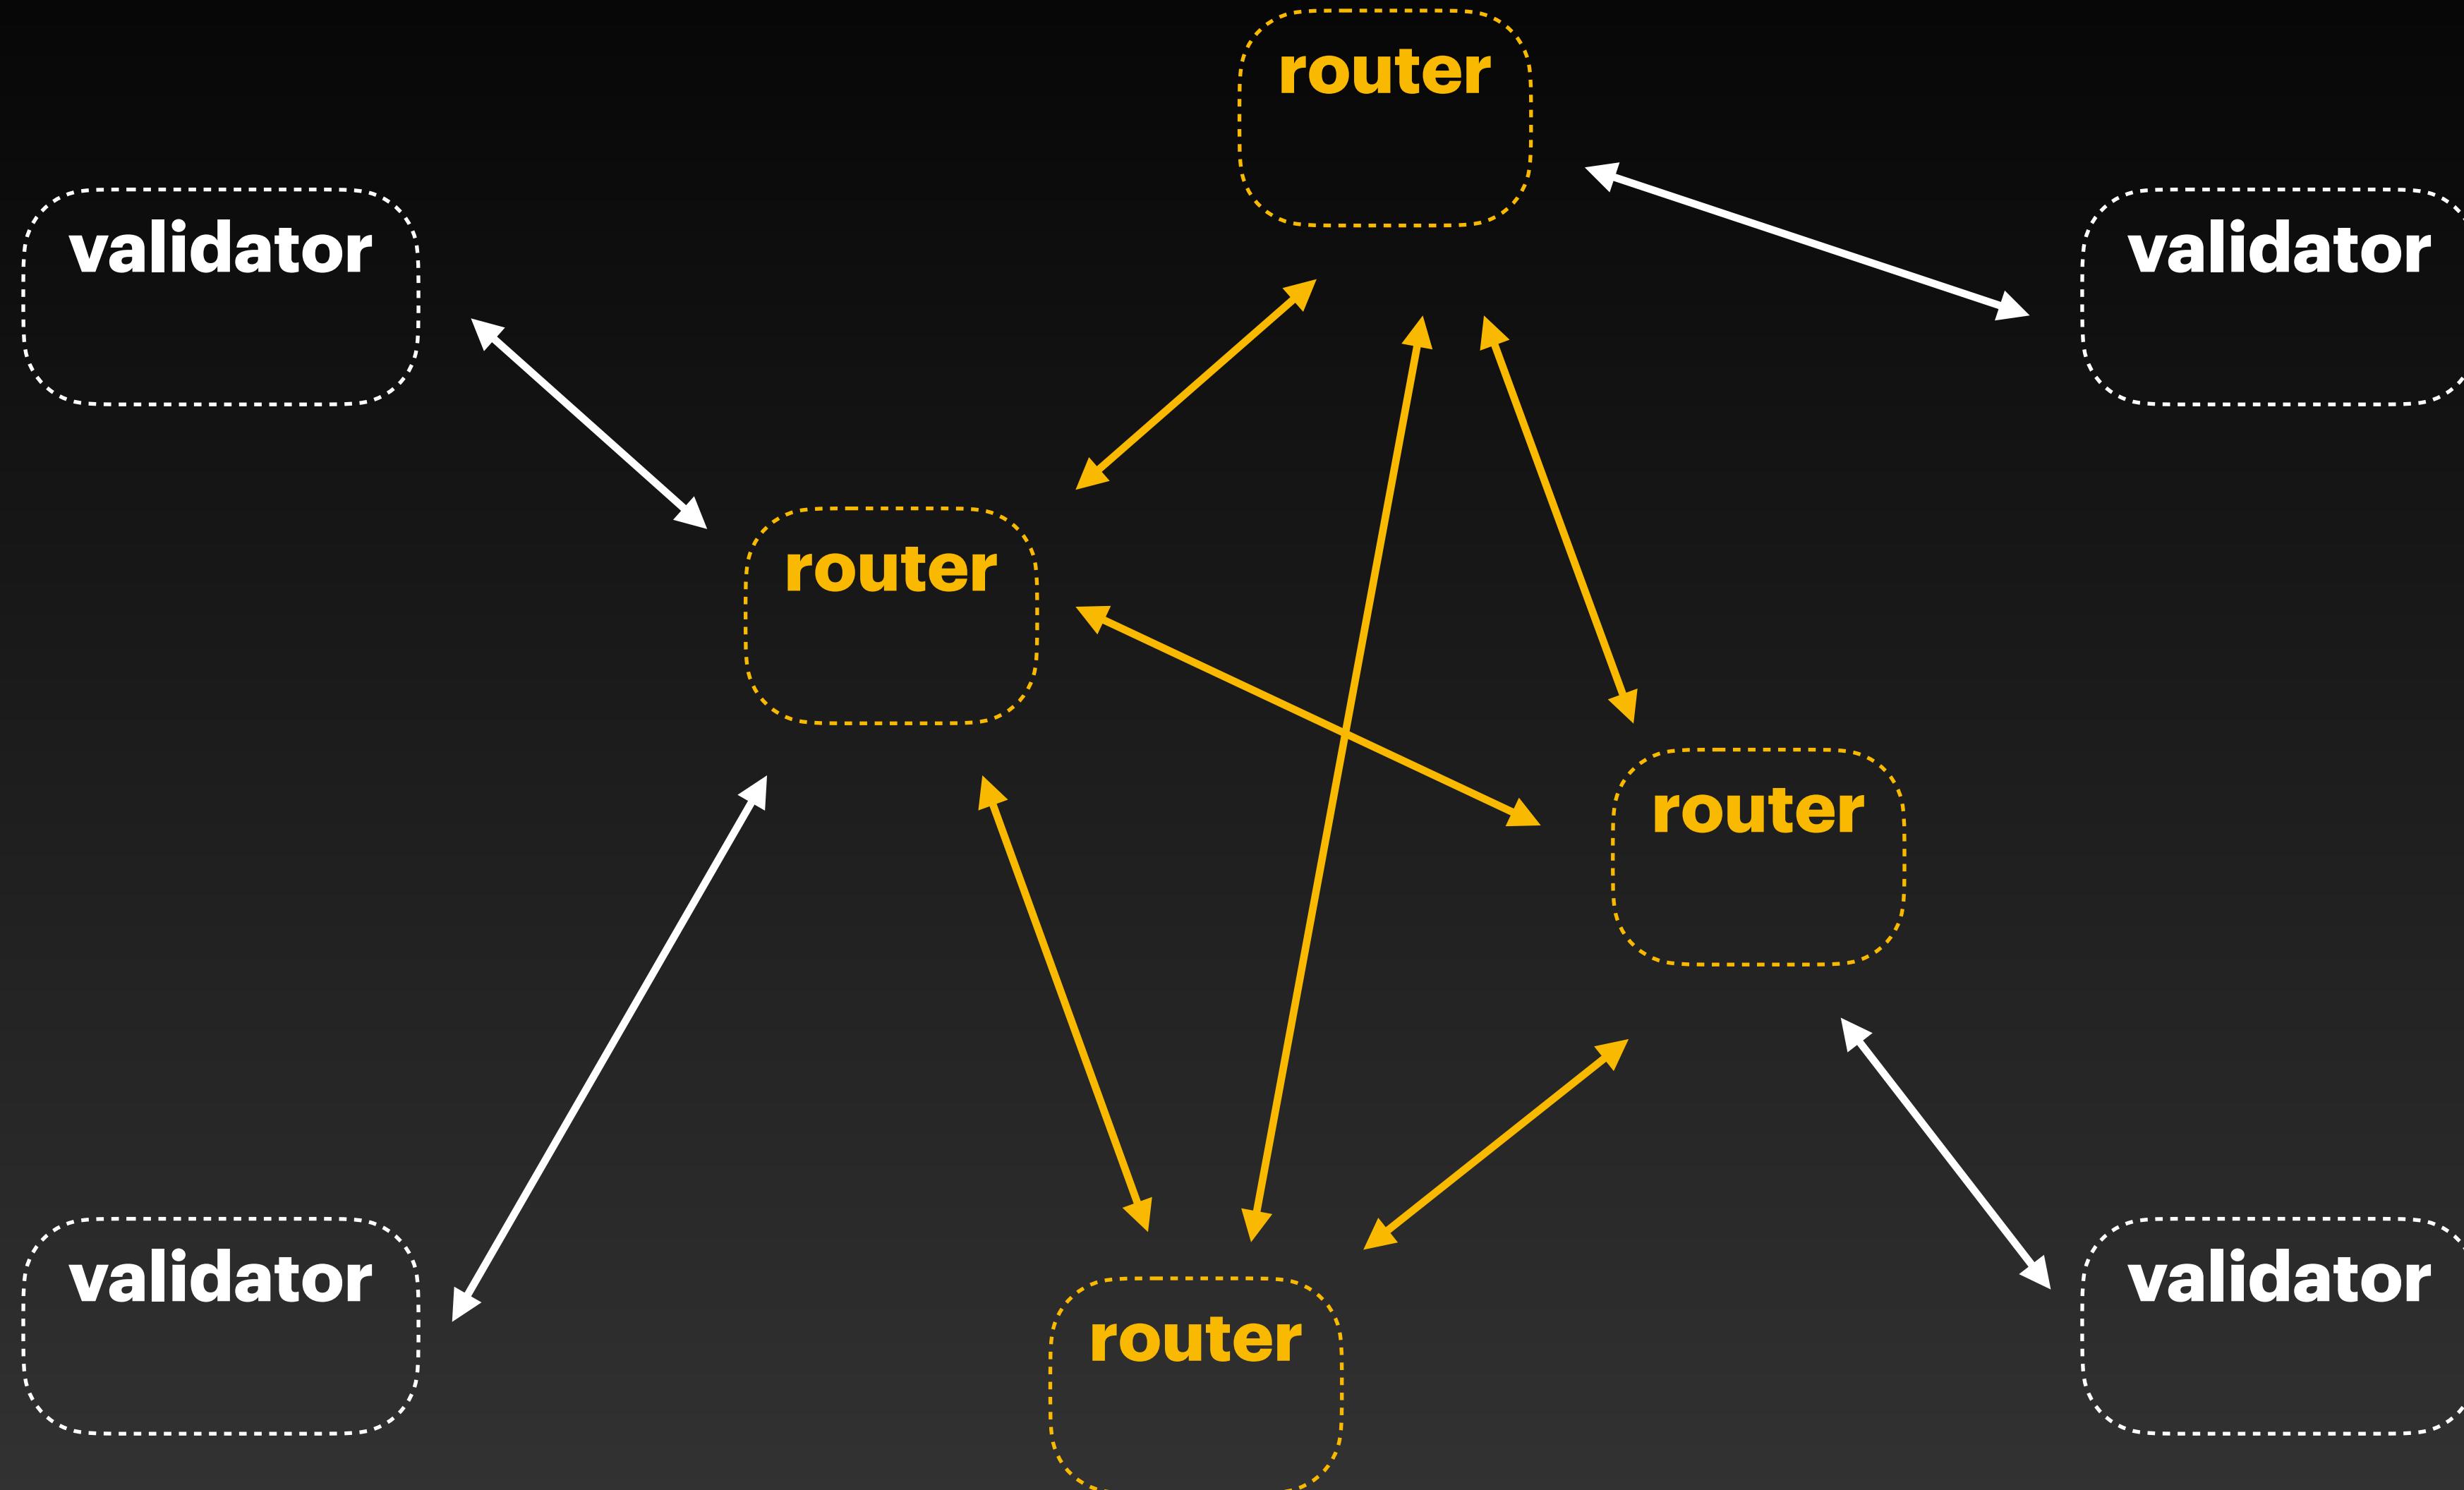
## Challenge #4: Ordering

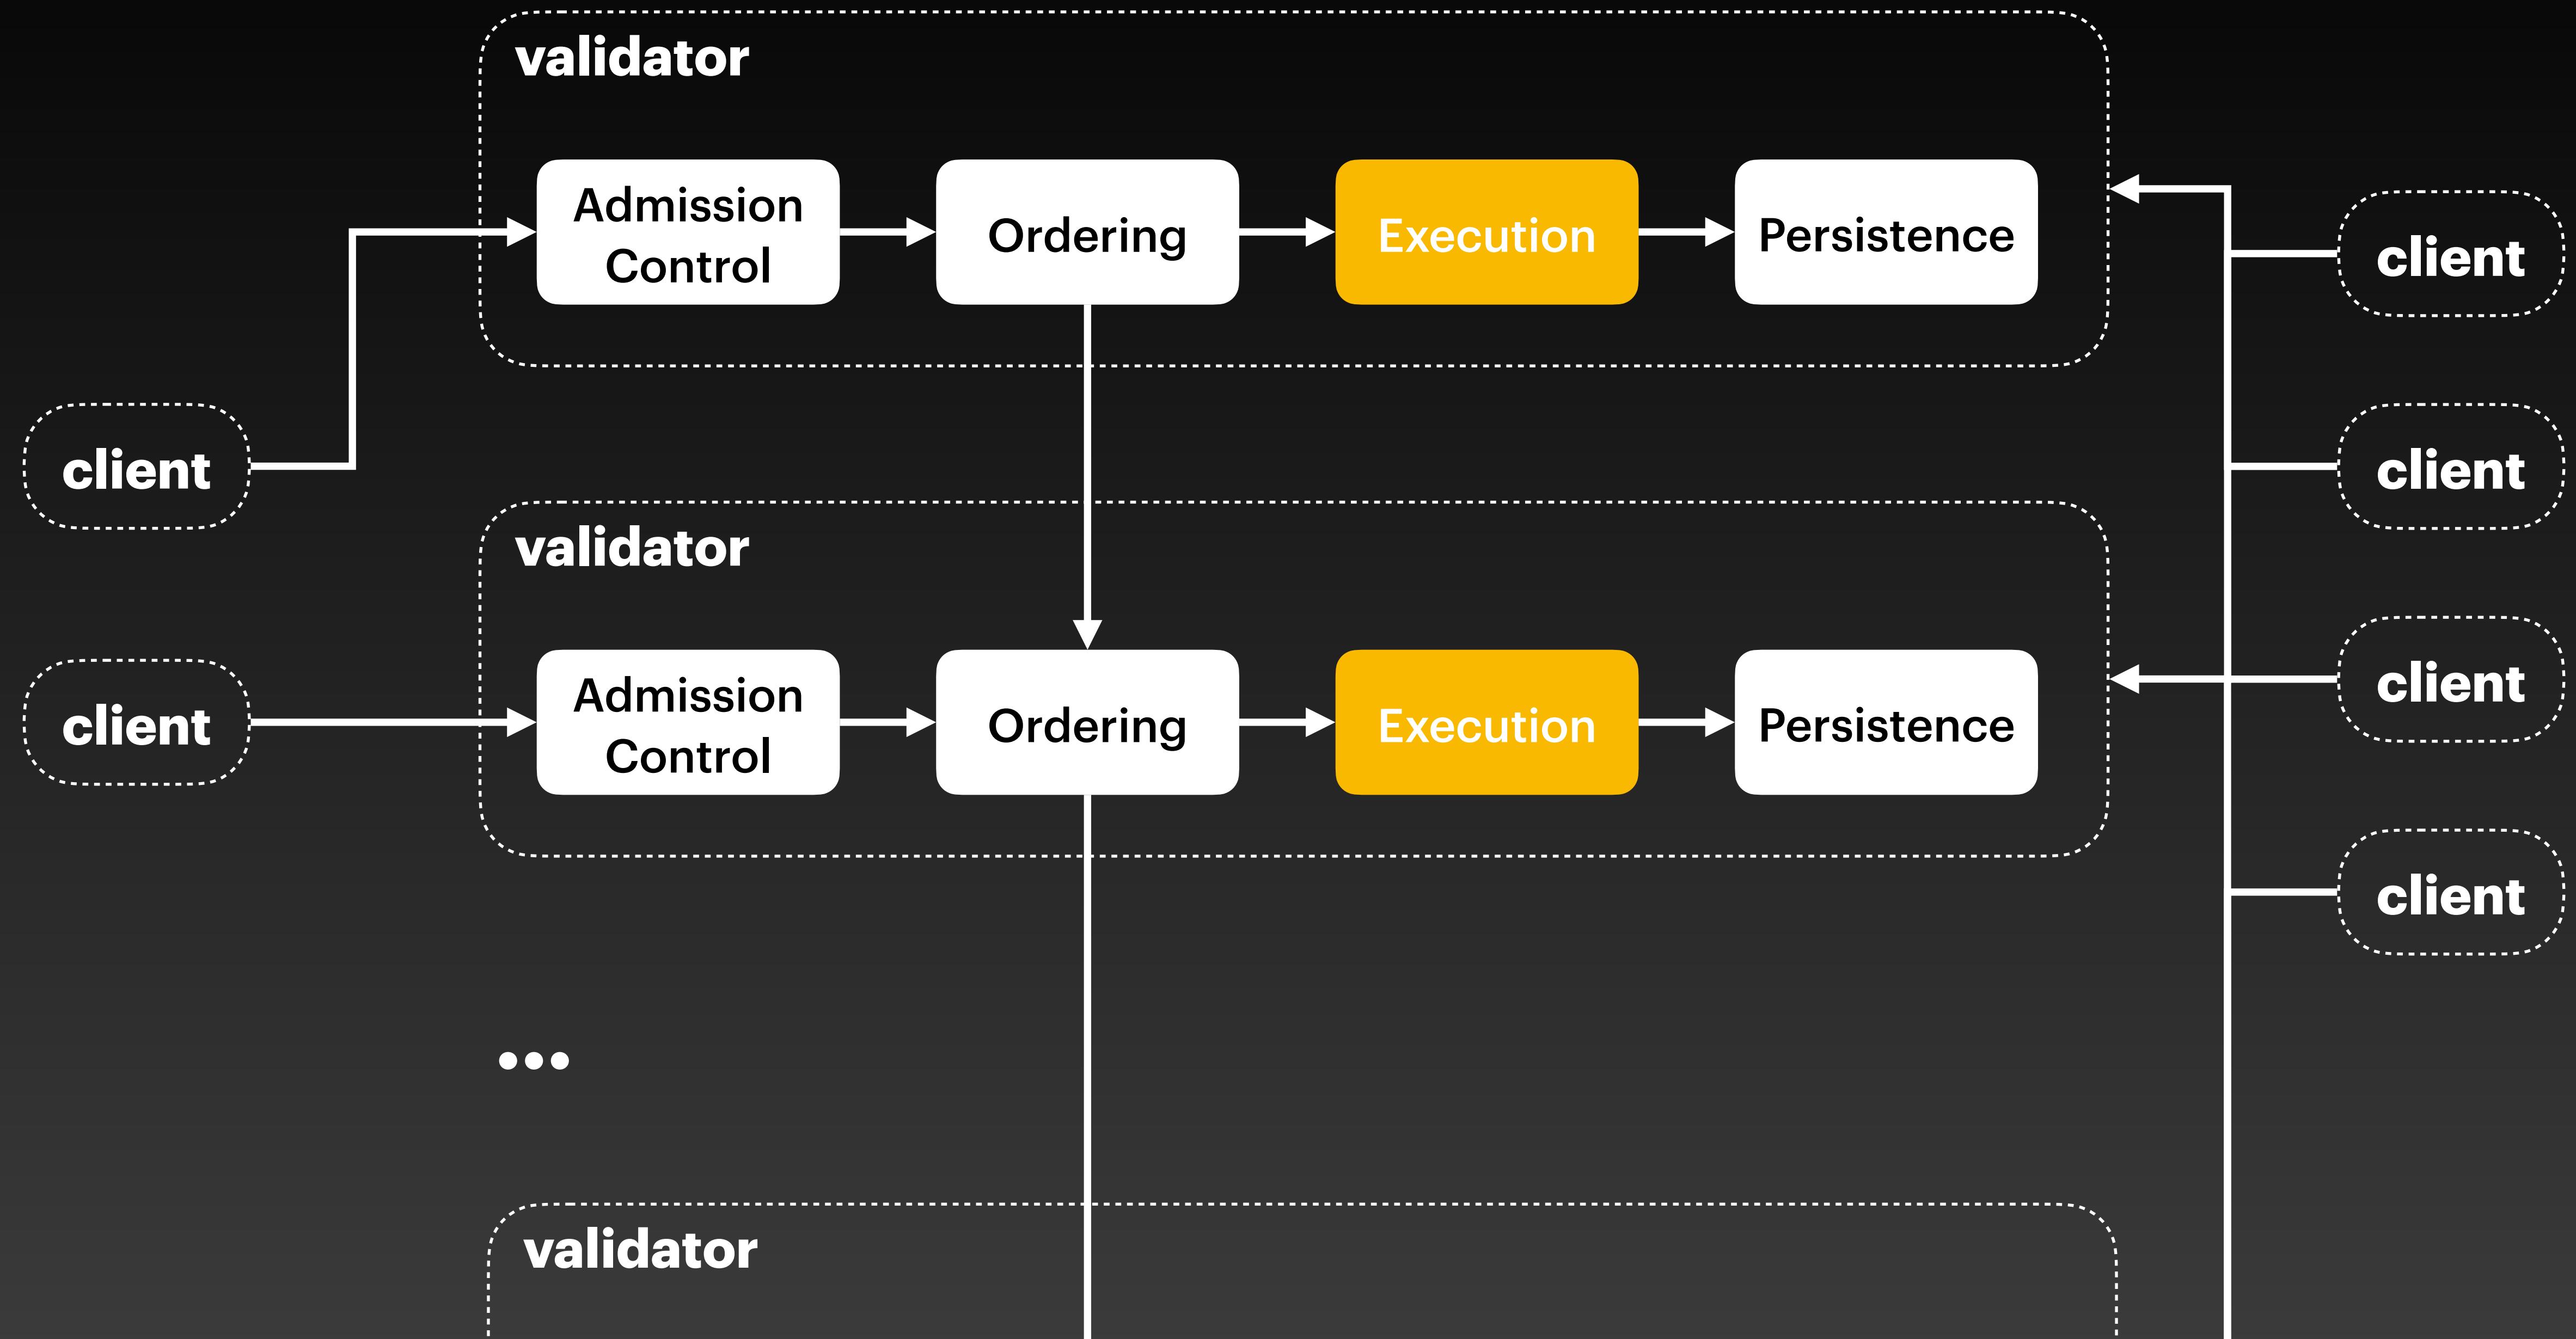
- How to find the best path to send the block to another node?
- DoS against the leader are particularly effective
- Reordering messages causes massive slowdowns
- **Nodes don't know whether they are connected to a malicious node**

# Network Security

## Challenge #4: Ordering

- How to find the best path to send the block to another node?
- DoS against the leader are particularly effective
- Reordering messages causes massive slowdowns
- Nodes don't know whether they are connected to a malicious node
- **Bad nodes have access to insider information (committee addresses)**





## Example Transaction

**T1**

**Inputs:** O1, O2, O3

**Output:** Mutate O1, Transfer O2, Delete O3, Create O4

## Check transaction, assign locks

O1

Version = 10

Owner = Alice

O2

Version = 27

Owner = Alice

O3

Version = 1001

Owner = Alice

### Checks

Input objects exist

Function call details

Signature of Alice

## Execute in parallel

O1

Version = 11

Owner = X

O2

Version = 28

Owner = Bob

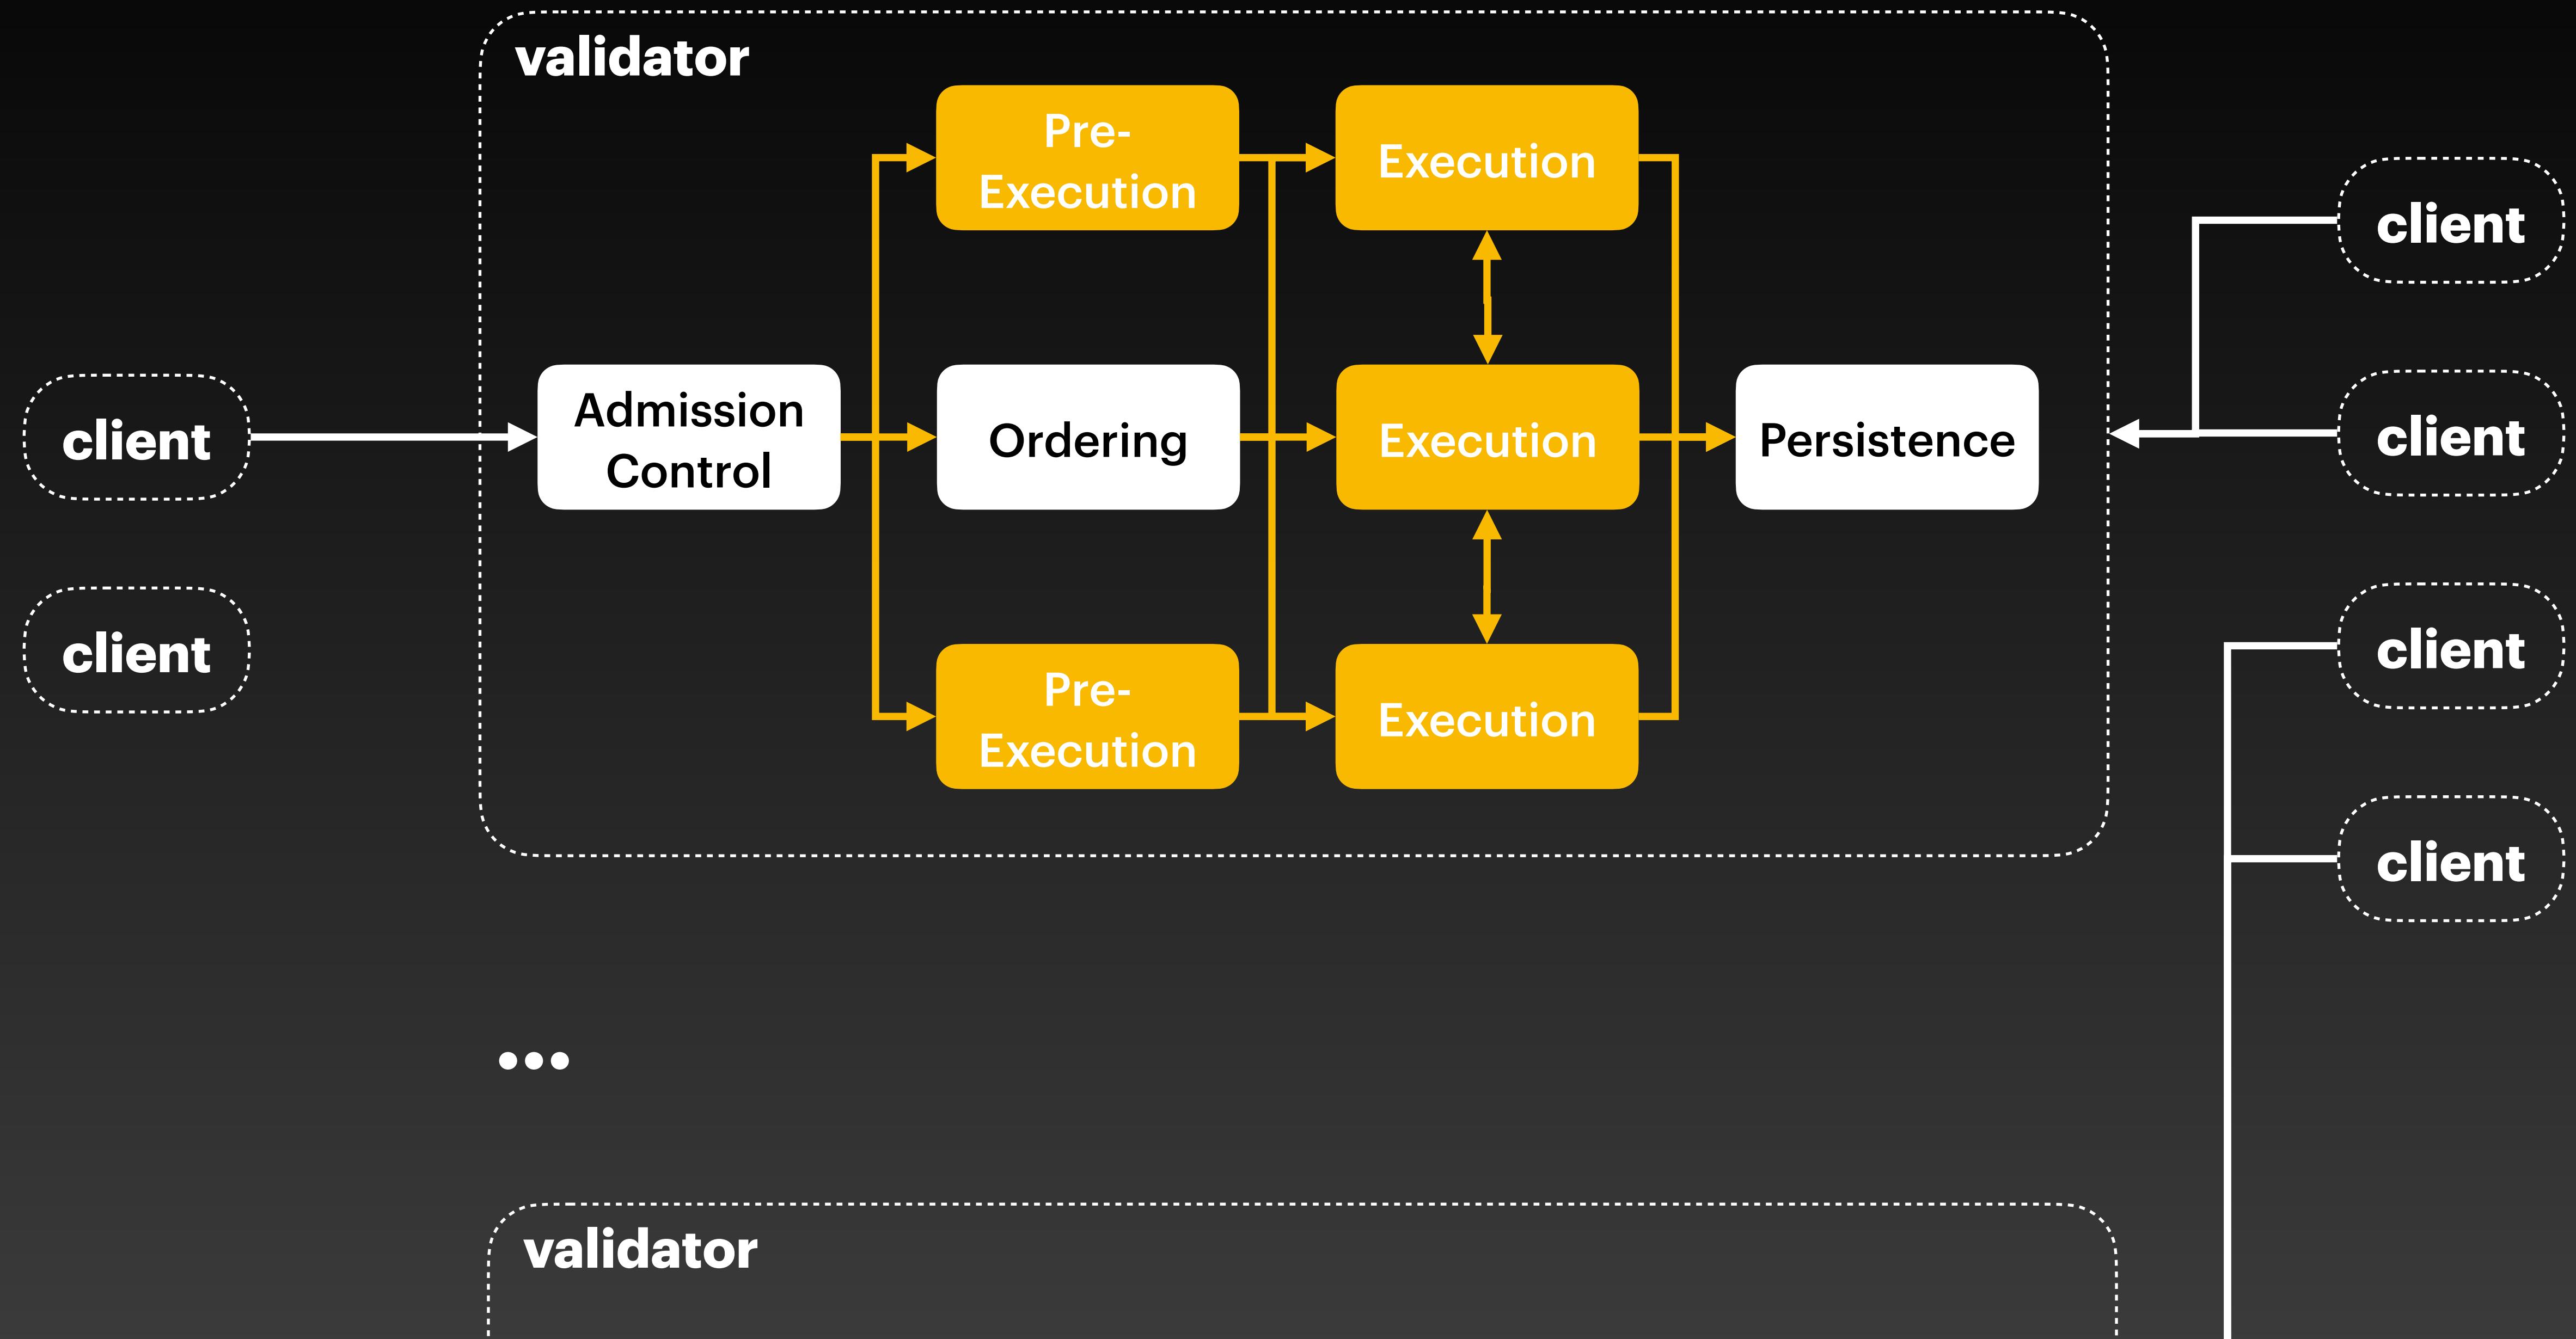
O4

Version = 1

Owner = Alice

### Execute T1

- O1 mutated
- O2 transferred
- O3 deleted
- O4 created



# Network Security

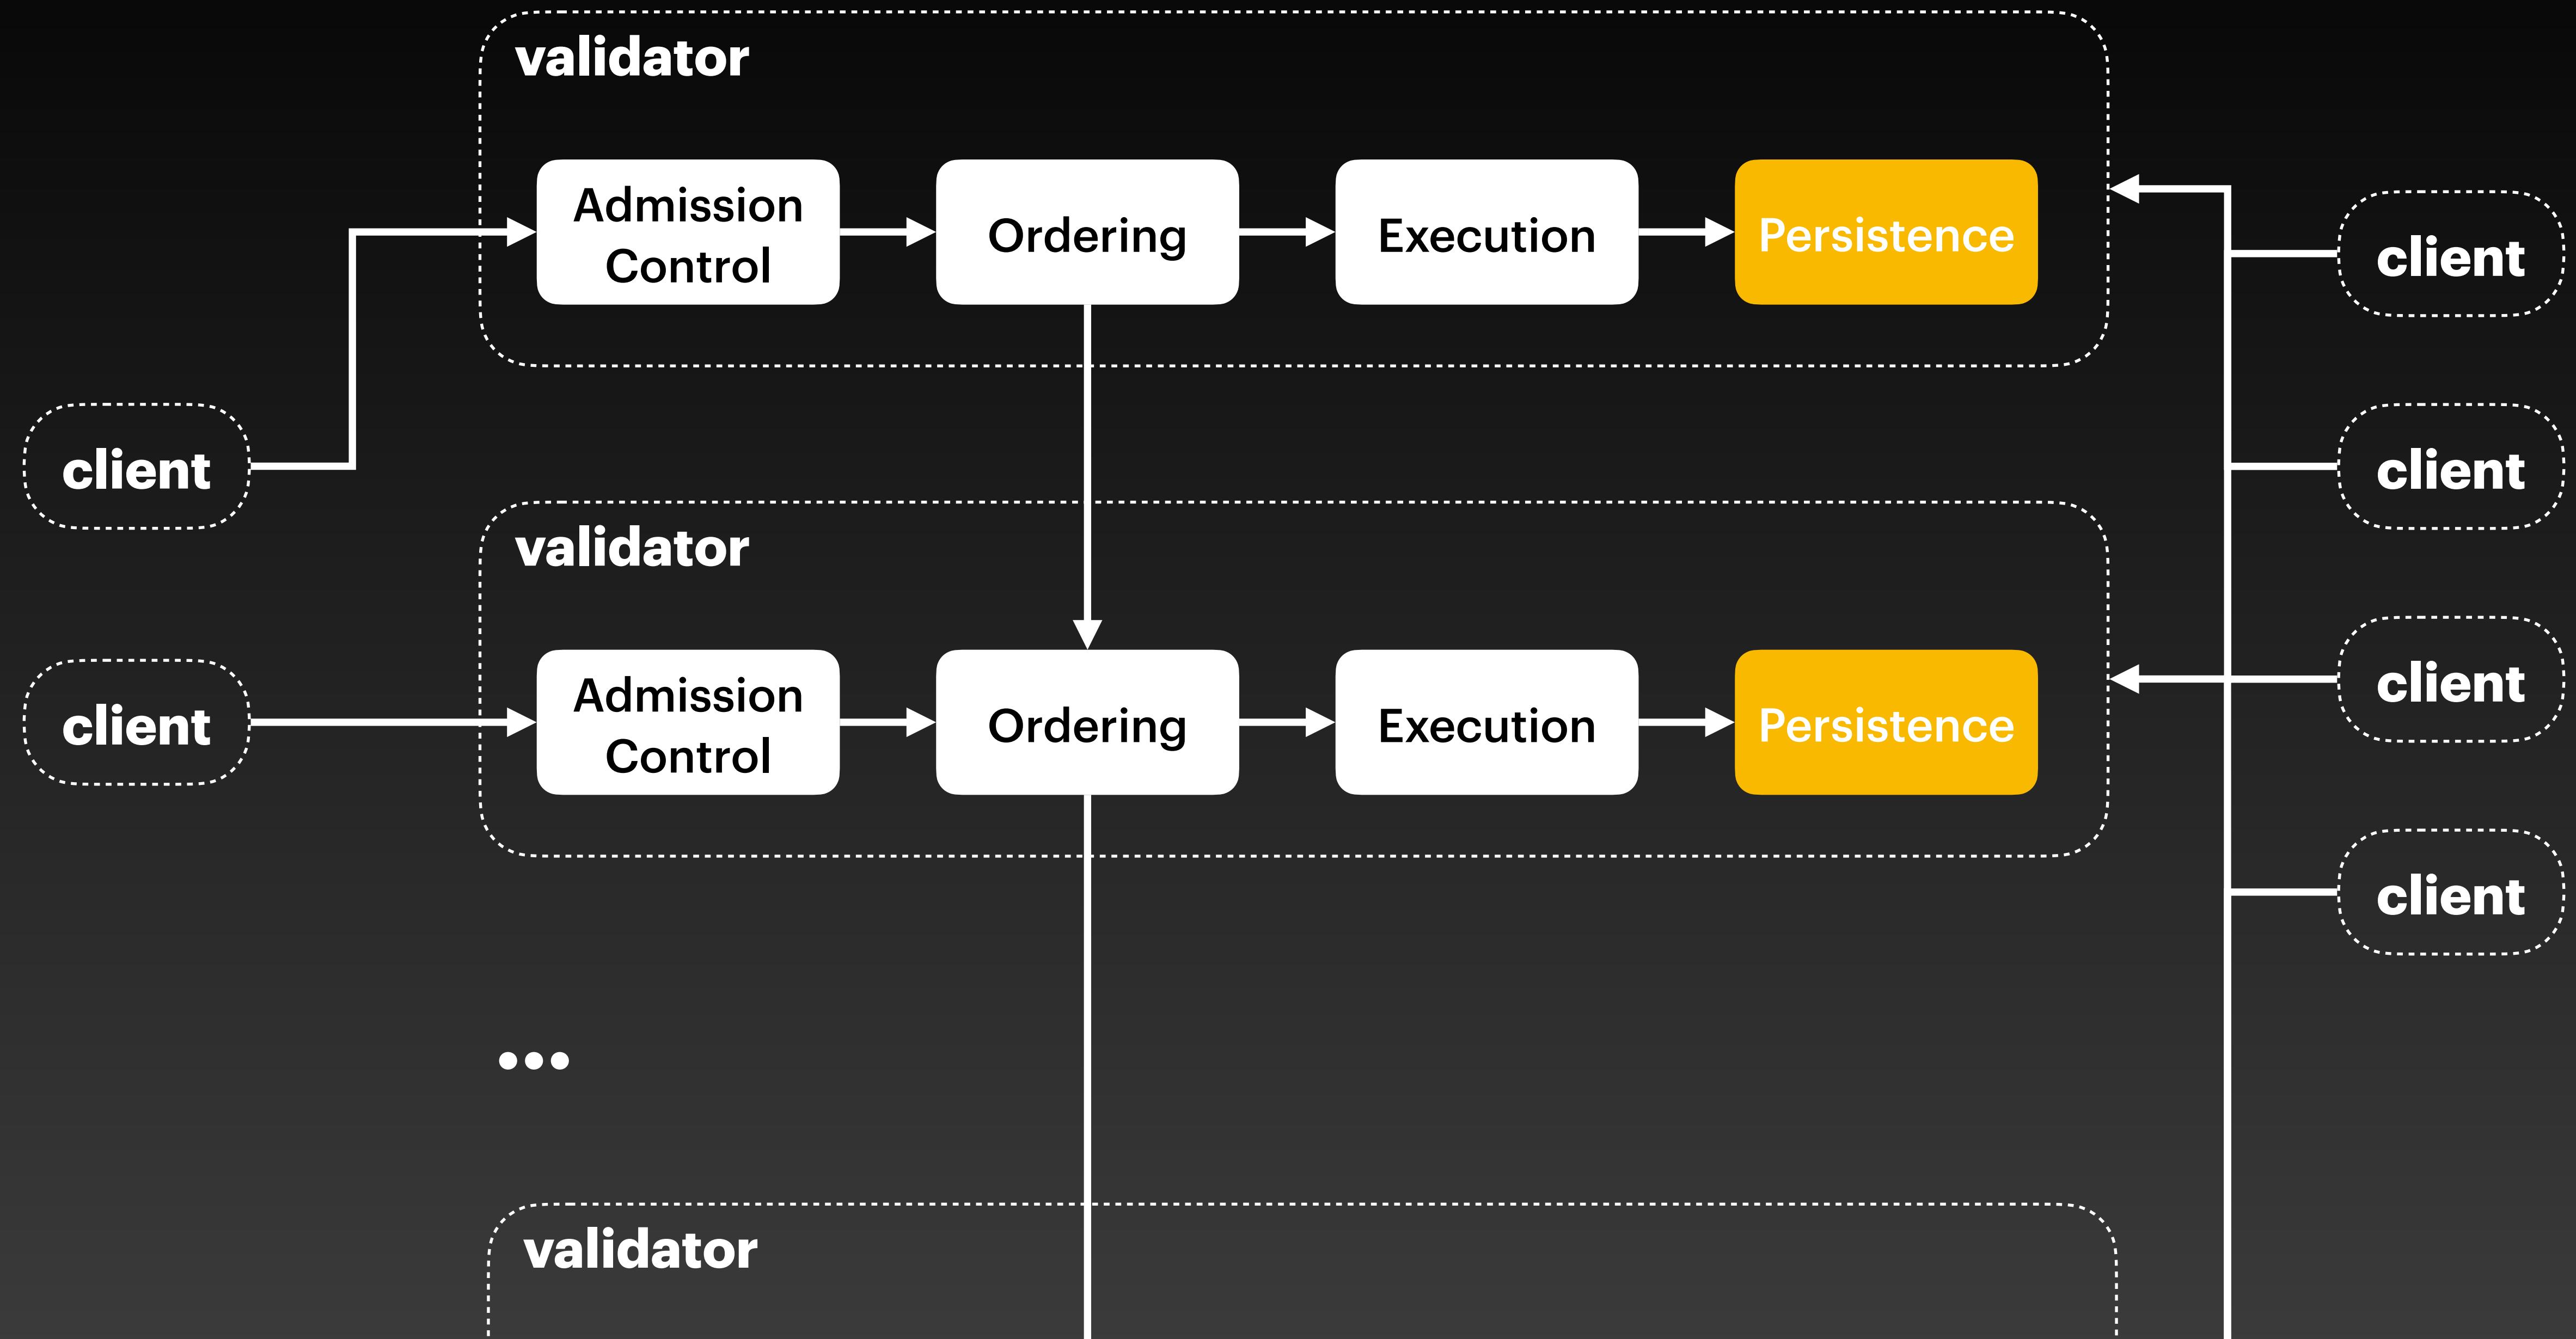
## Challenge #5: Execution

- **Intra-datacenter connections but on low power machines**

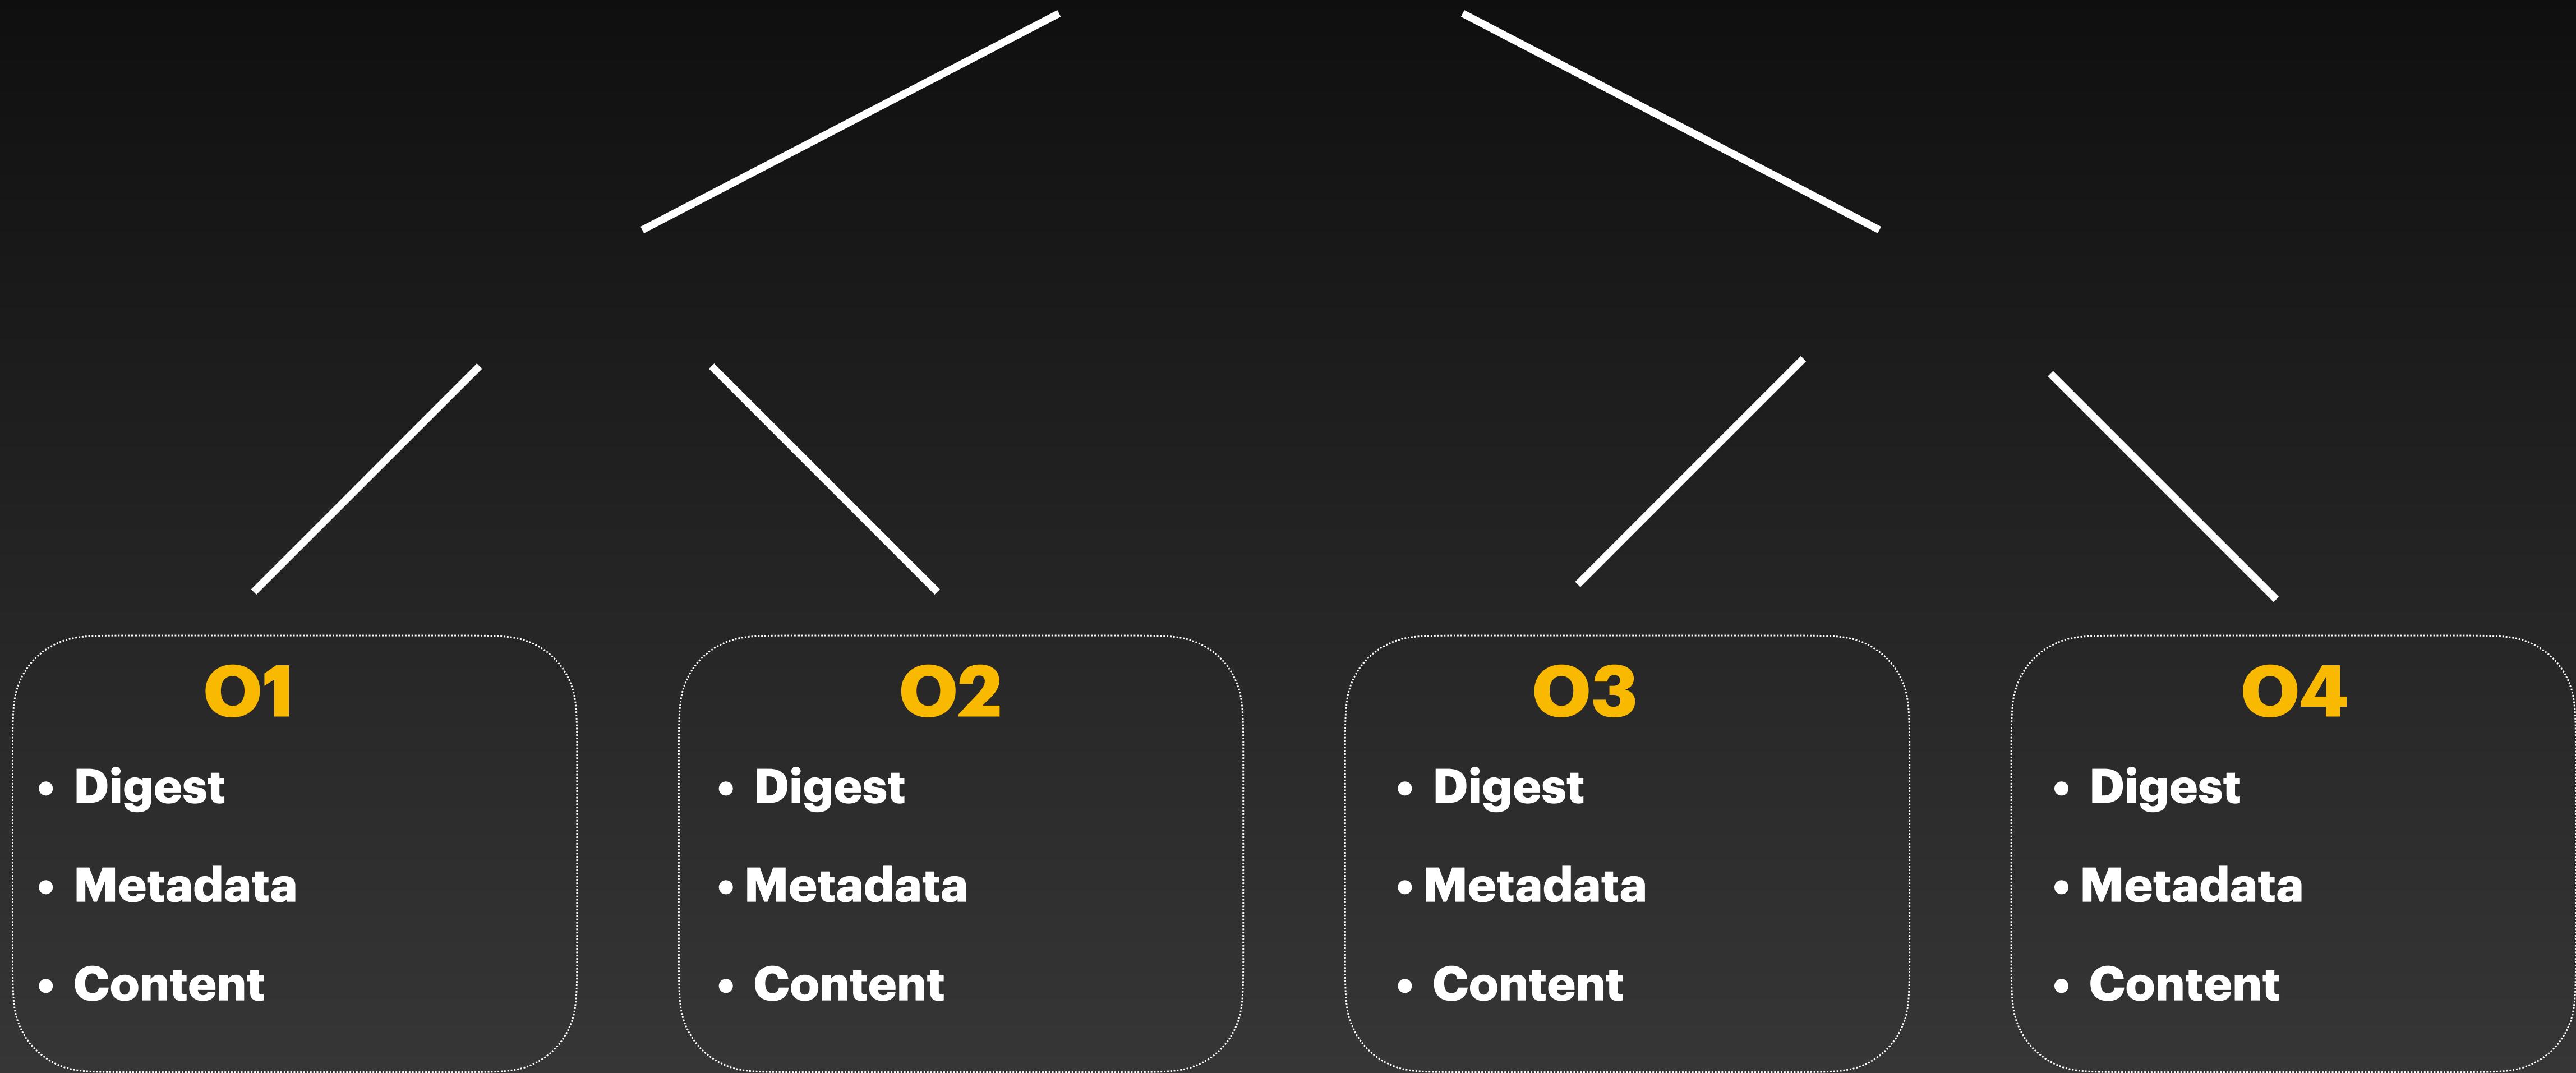
# Network Security

## Challenge #5: Execution

- Intra-datacenter connections but on low power machines
- **Load drastically varies: need elasticity**



# Root



**Root**

**H(O1,O2)**

**H(O3,O4)**

**O1**

- Digest
- Metadata
- Content

**O2**

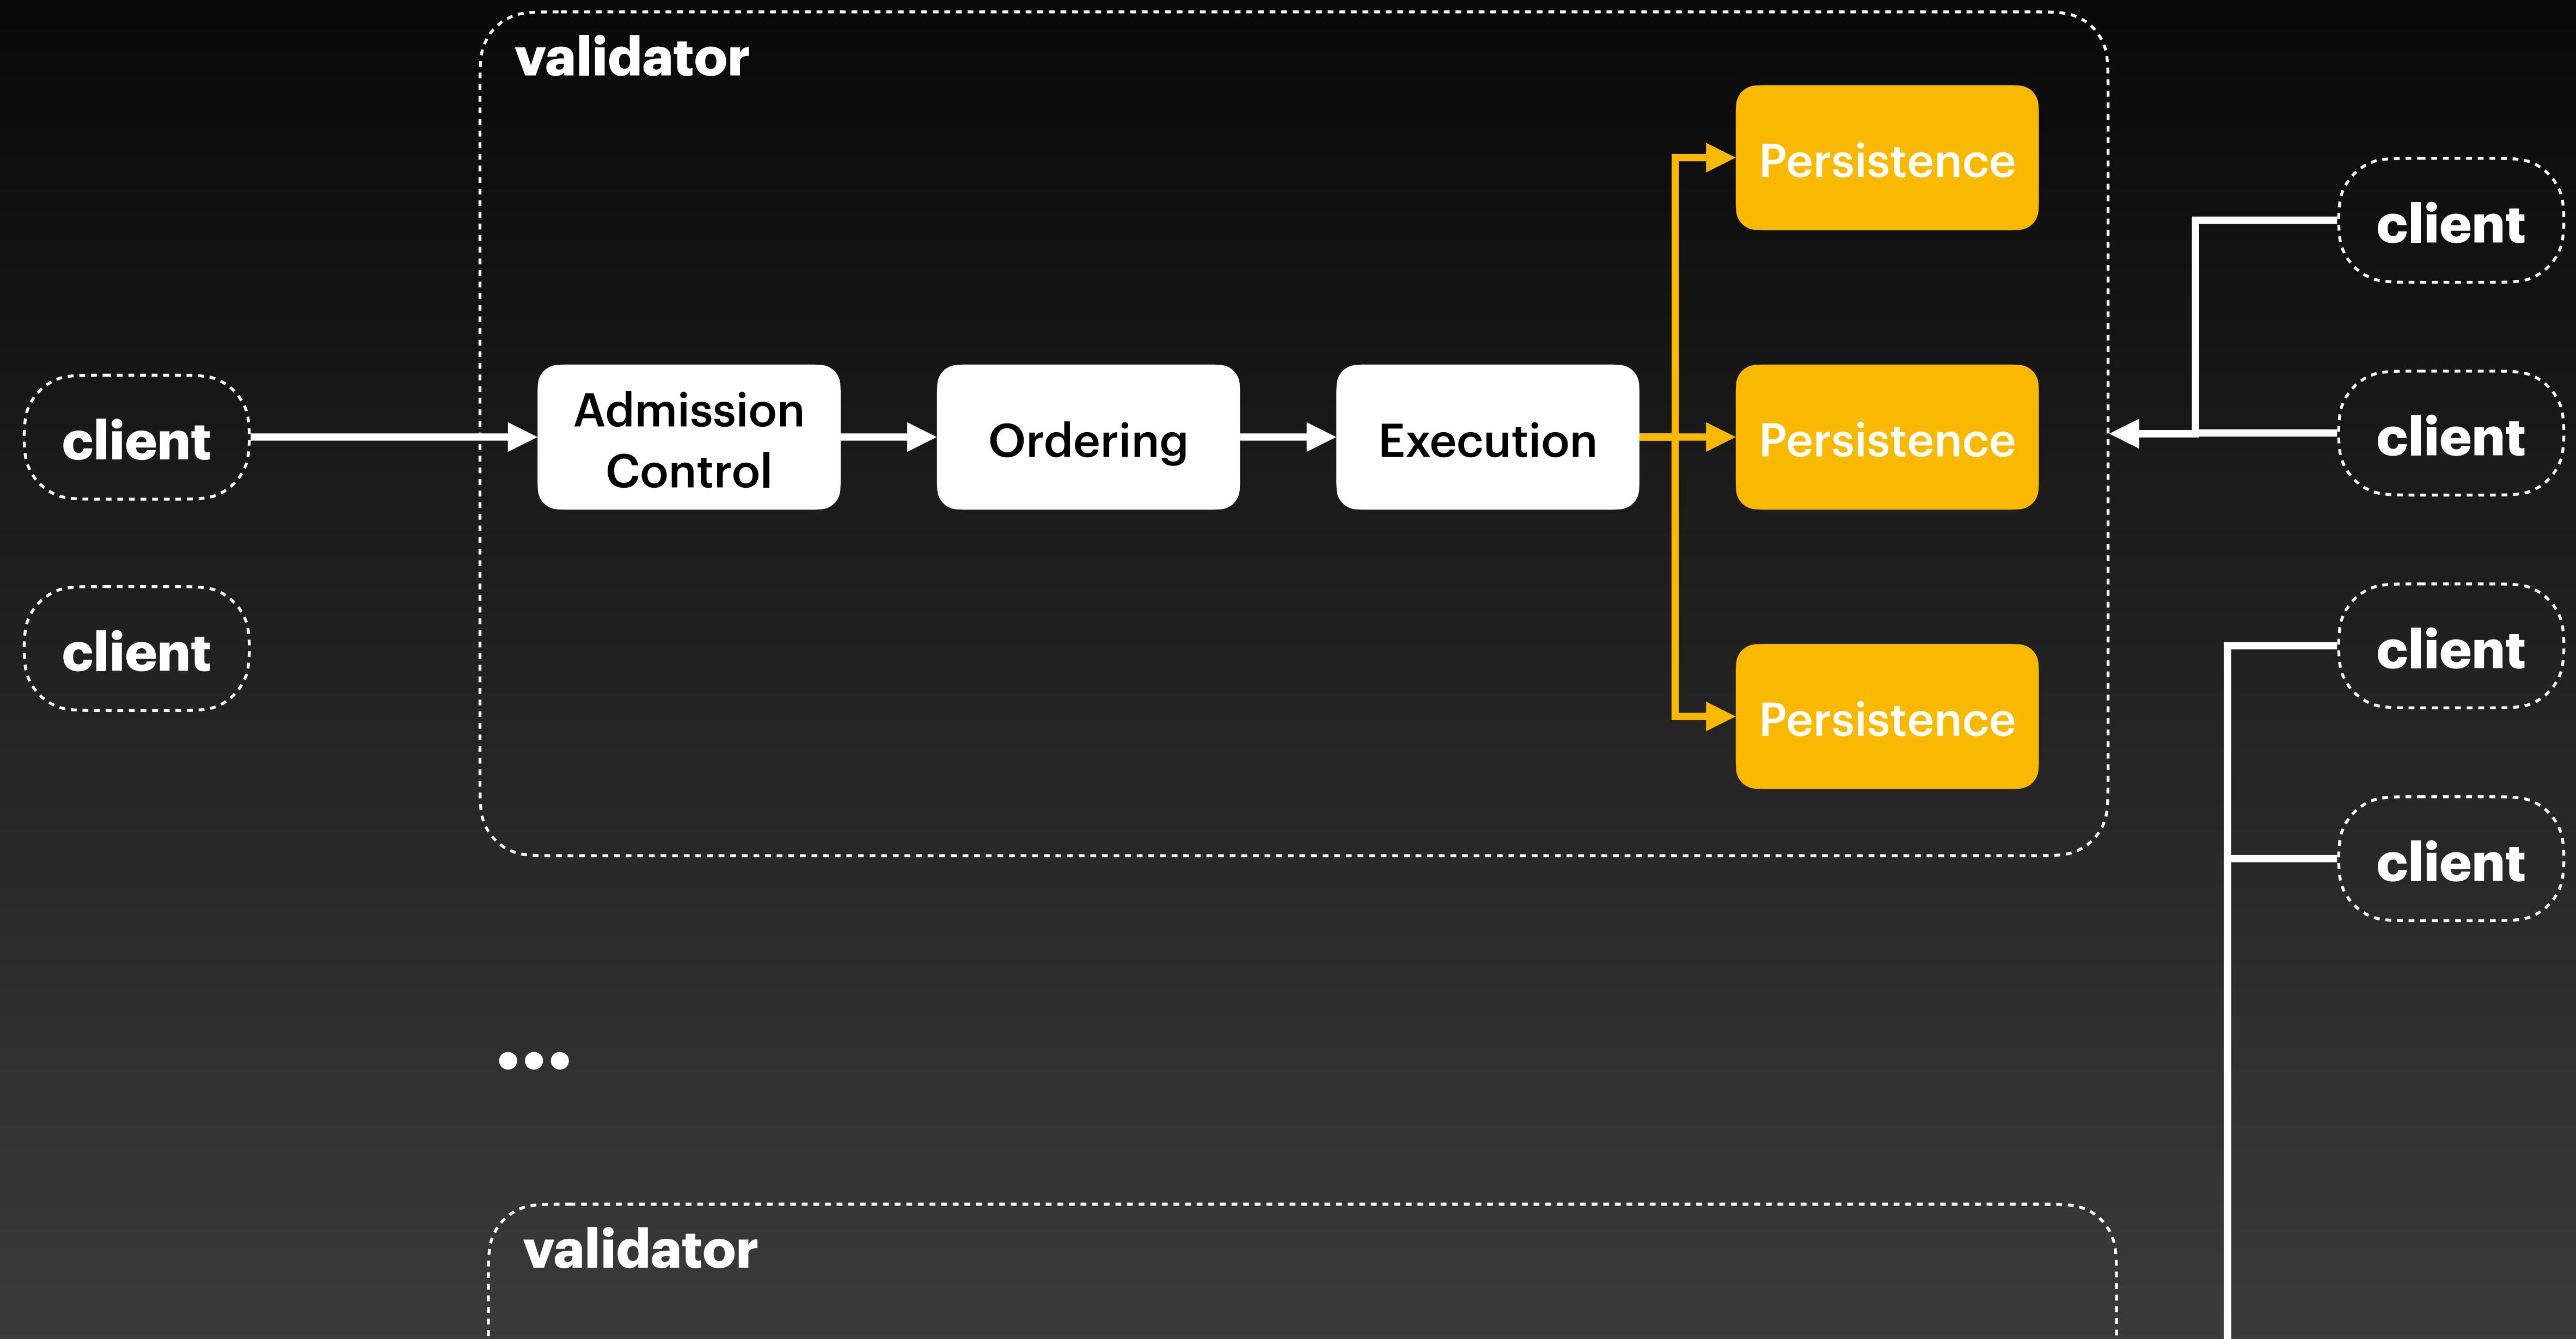
- Digest
- Metadata
- Content

**O3**

- Digest
- Metadata
- Content

**O4**

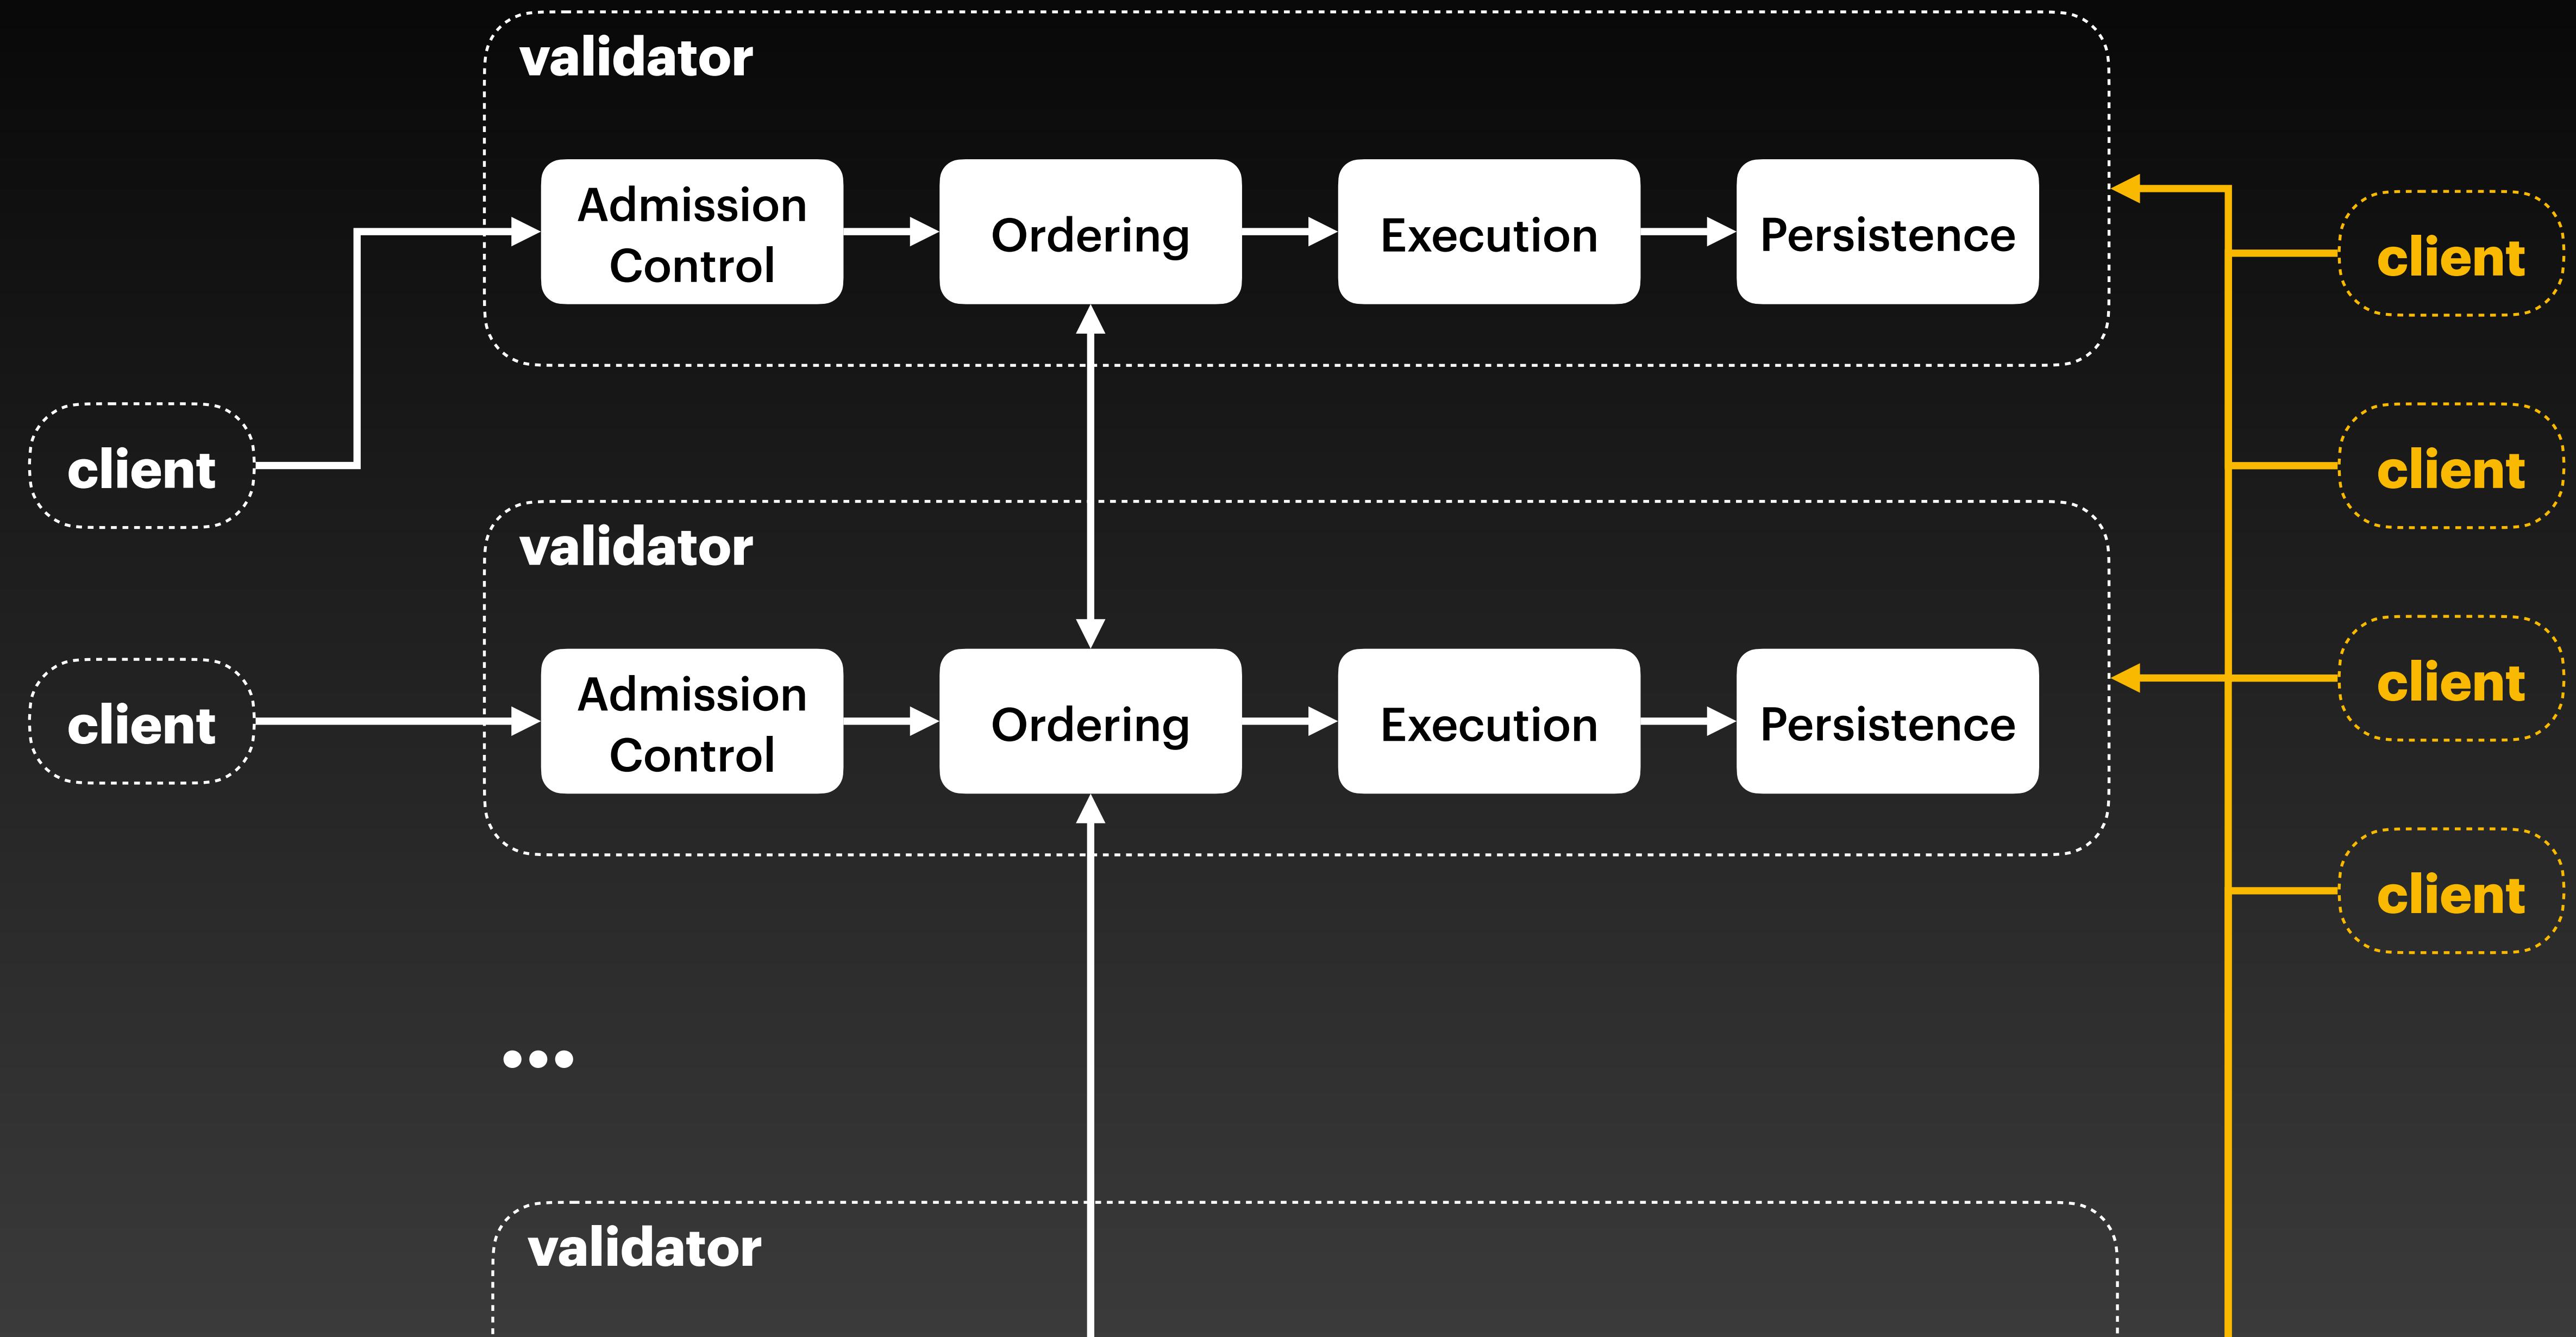
- Digest
- Metadata
- Content

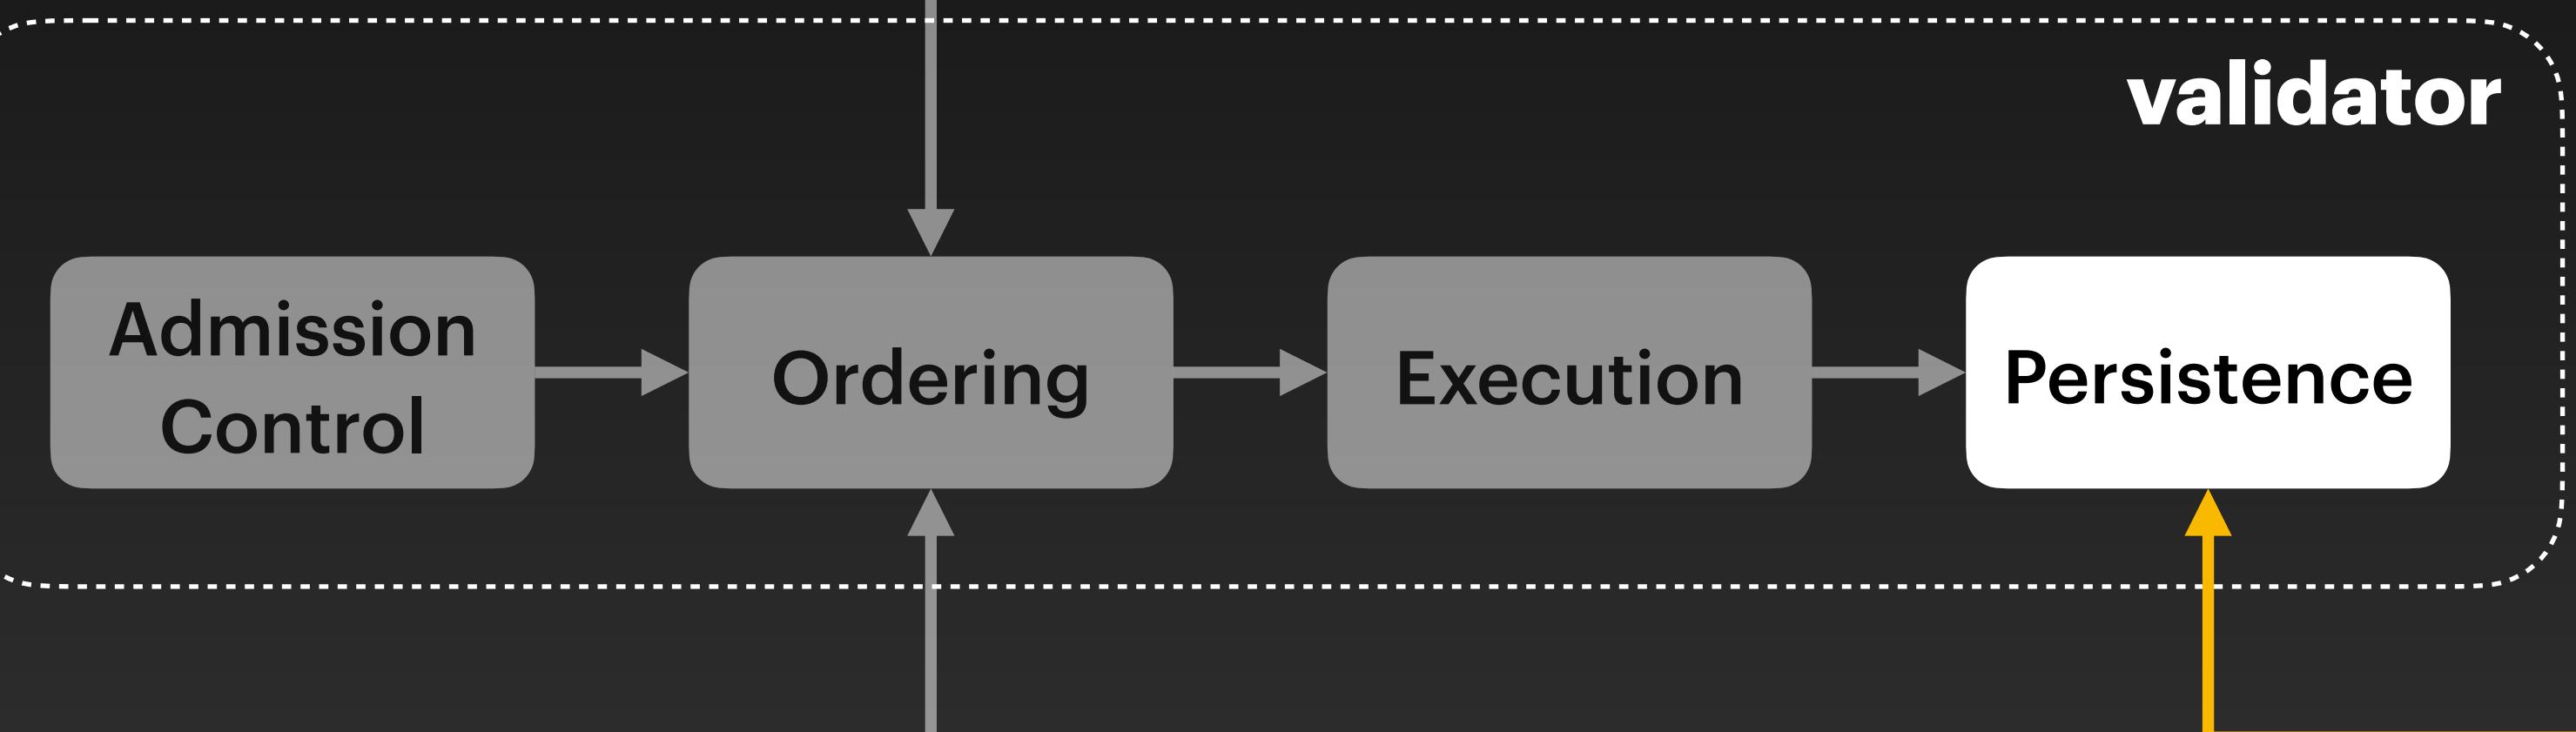
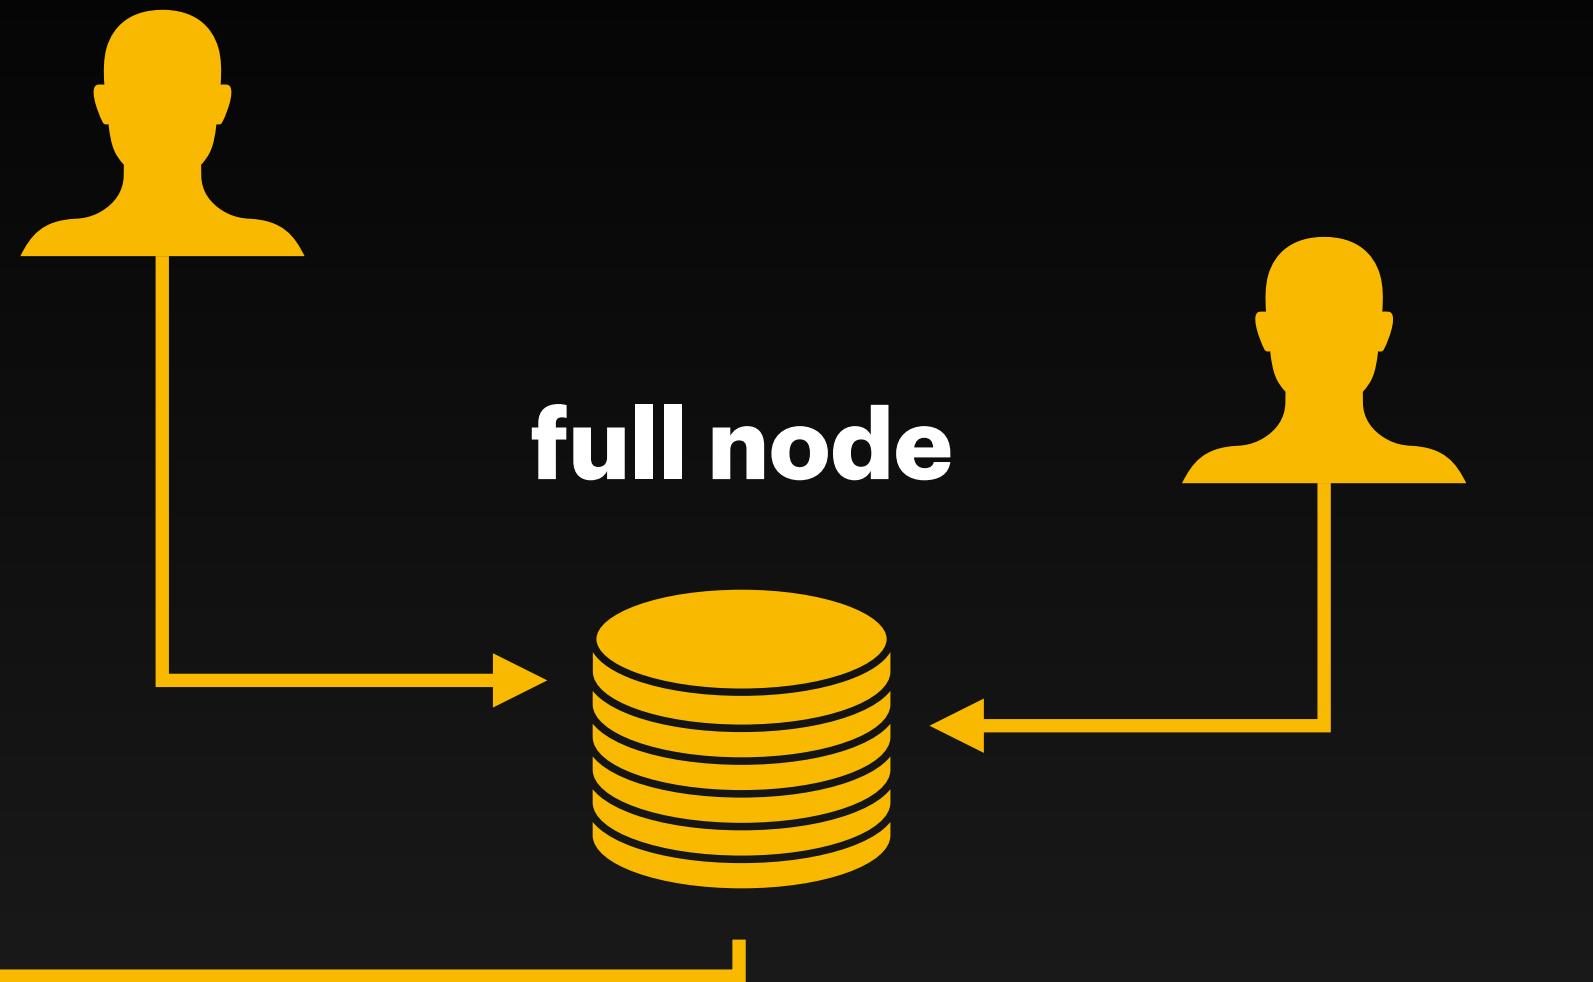
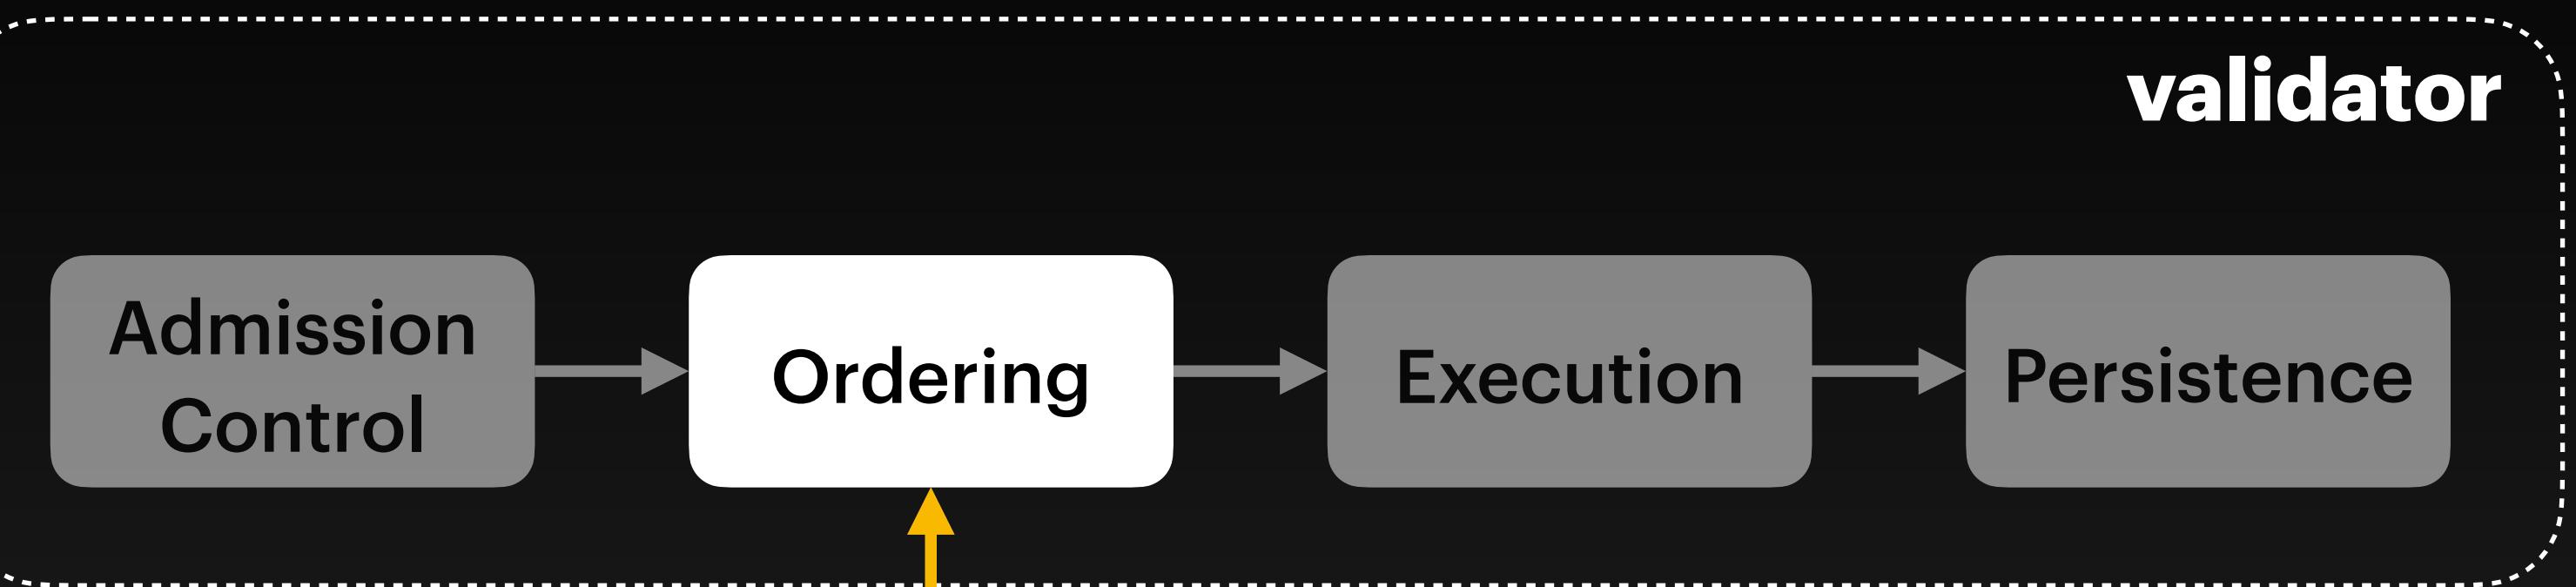


# Network Security

## Challenge #6: Persistence

- Need low-latency networking to distribute the tree creation





...

**validator**



# Network Security

## Challenge #7: Reads

- **Potentially very large number of readers (>400)**

# Network Security

## Challenge #7: Reads

- Potentially very large number of readers (>400)
- **Unpredictable, may read arbitrary data**

# Network Security

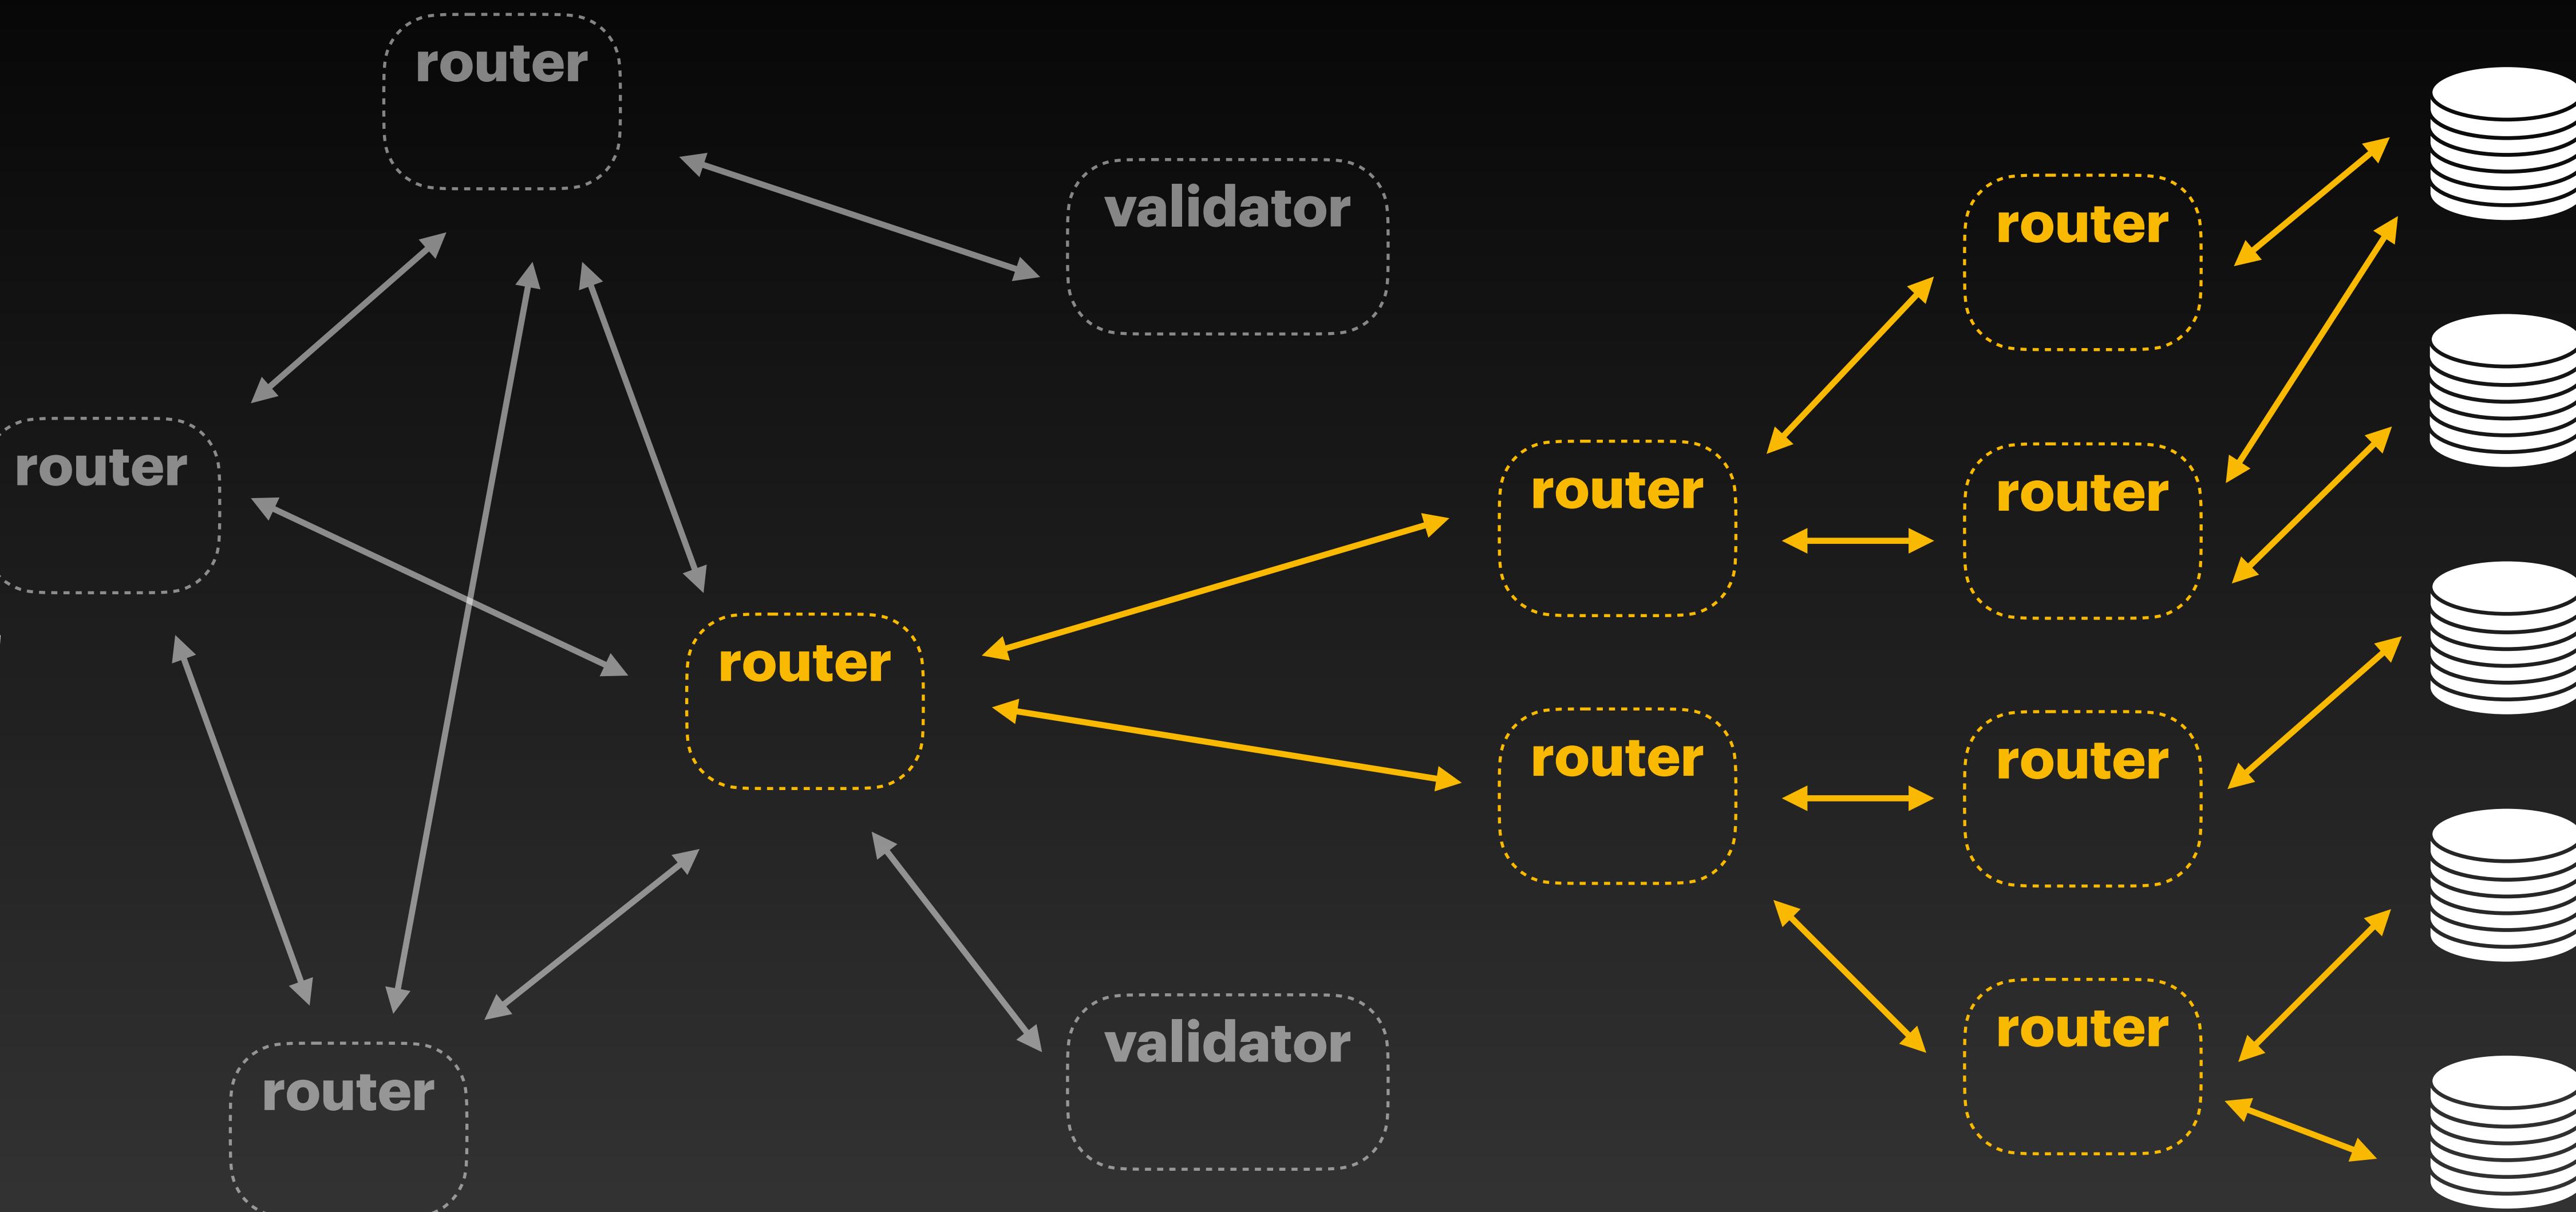
## Challenge #7: Reads

- Potentially very large number of readers (>400)
- Unpredictable, may read arbitrary data
- **Sometimes require extreme performance**

# Network Security

## Challenge #7: Reads

- Potentially very large number of readers (>400)
- Unpredictable, may read arbitrary data
- Sometimes require extreme performance
- **Most reads must be free**



alberto@mystenlabs.com