

Mahi-Mahi

Low-Latency Asynchronous BFT DAG-based Consensus

Byzantine Fault Tolerance



$> 2/3$



Byzantine Fault Tolerance

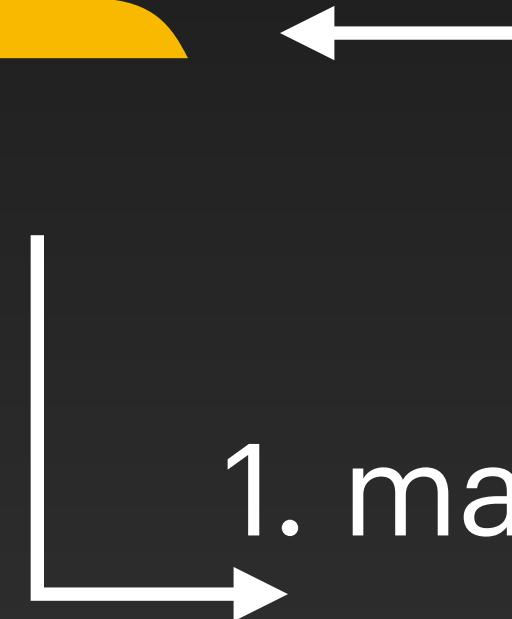


$\geq 2f+1$



$3f+1$

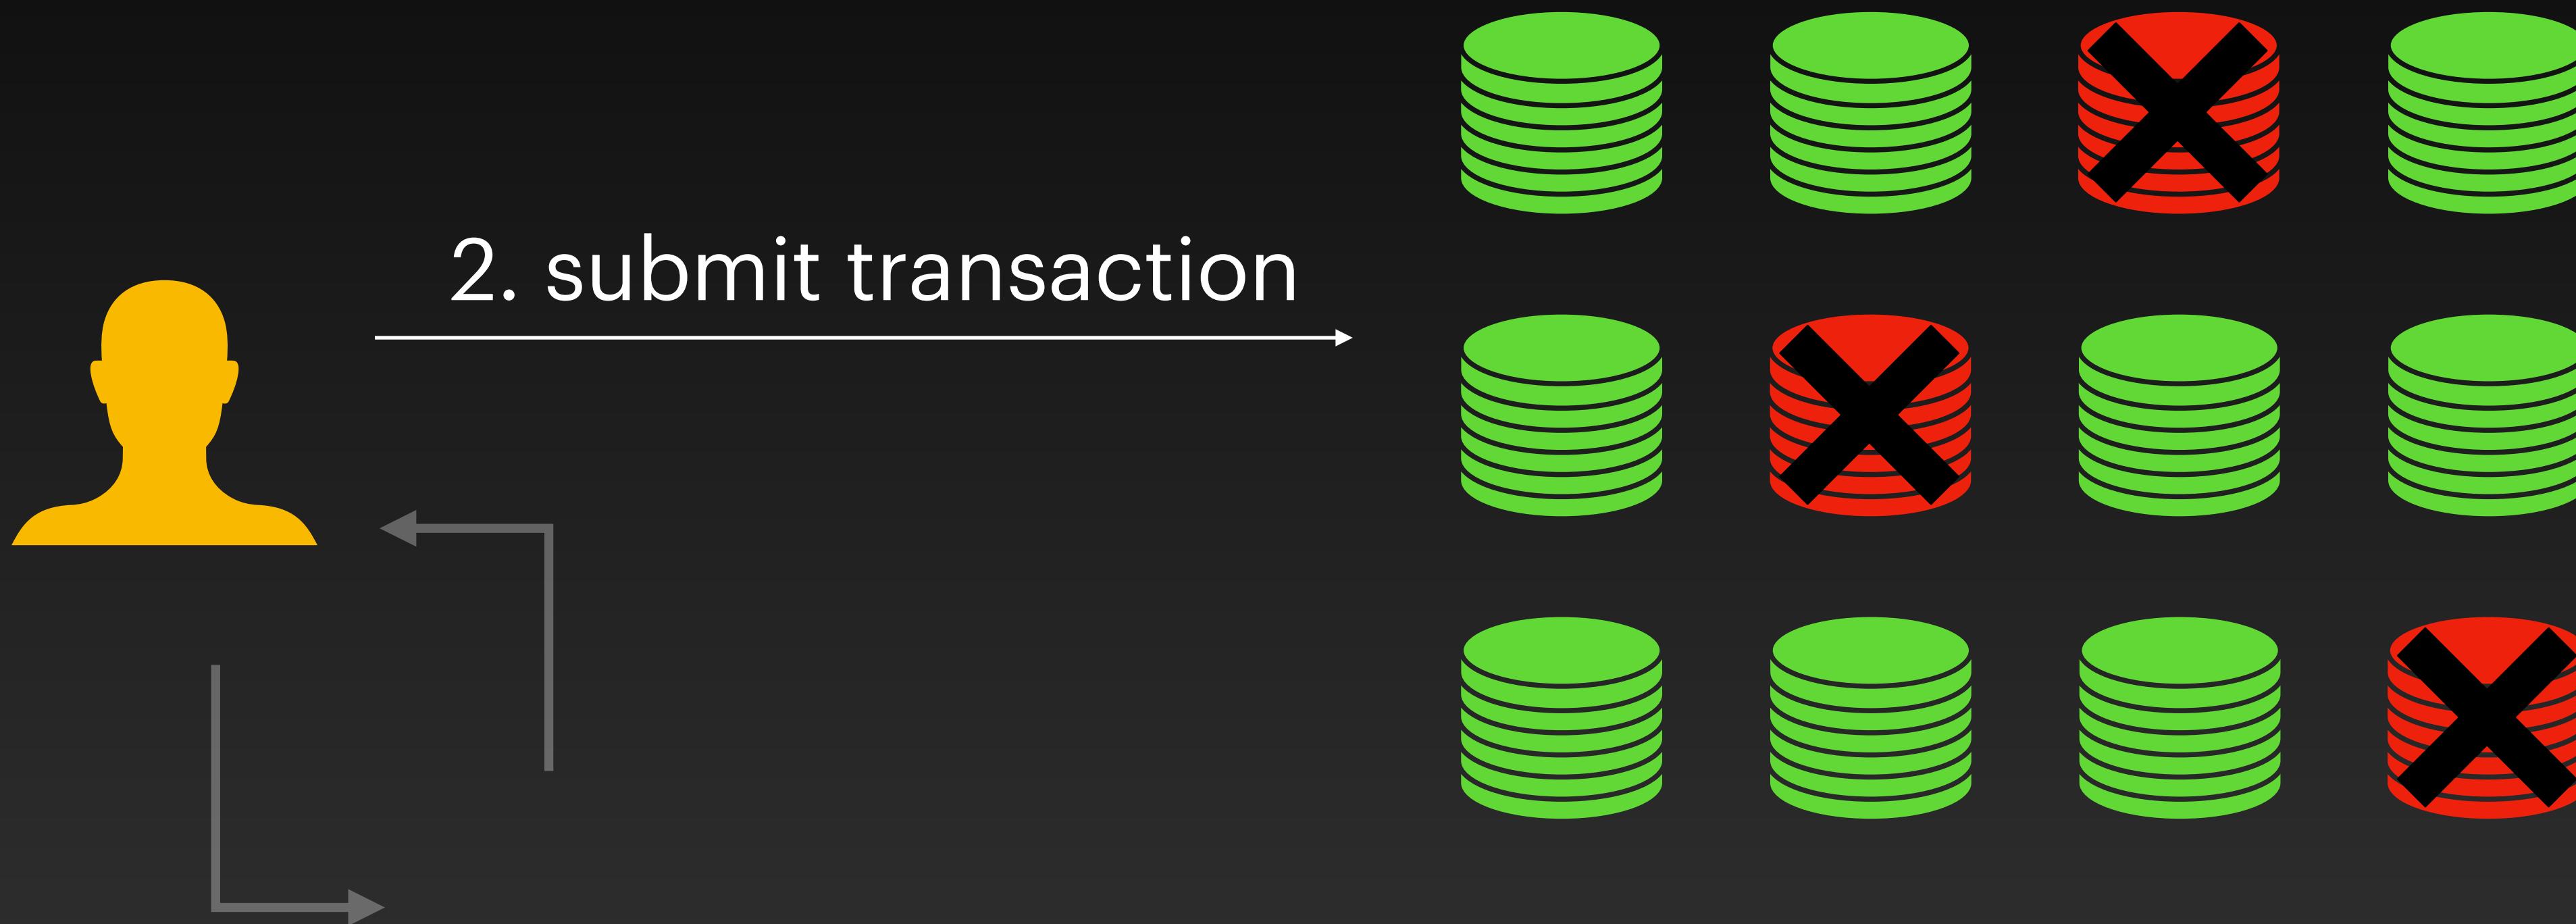
Blockchains



1. make transaction



Blockchains



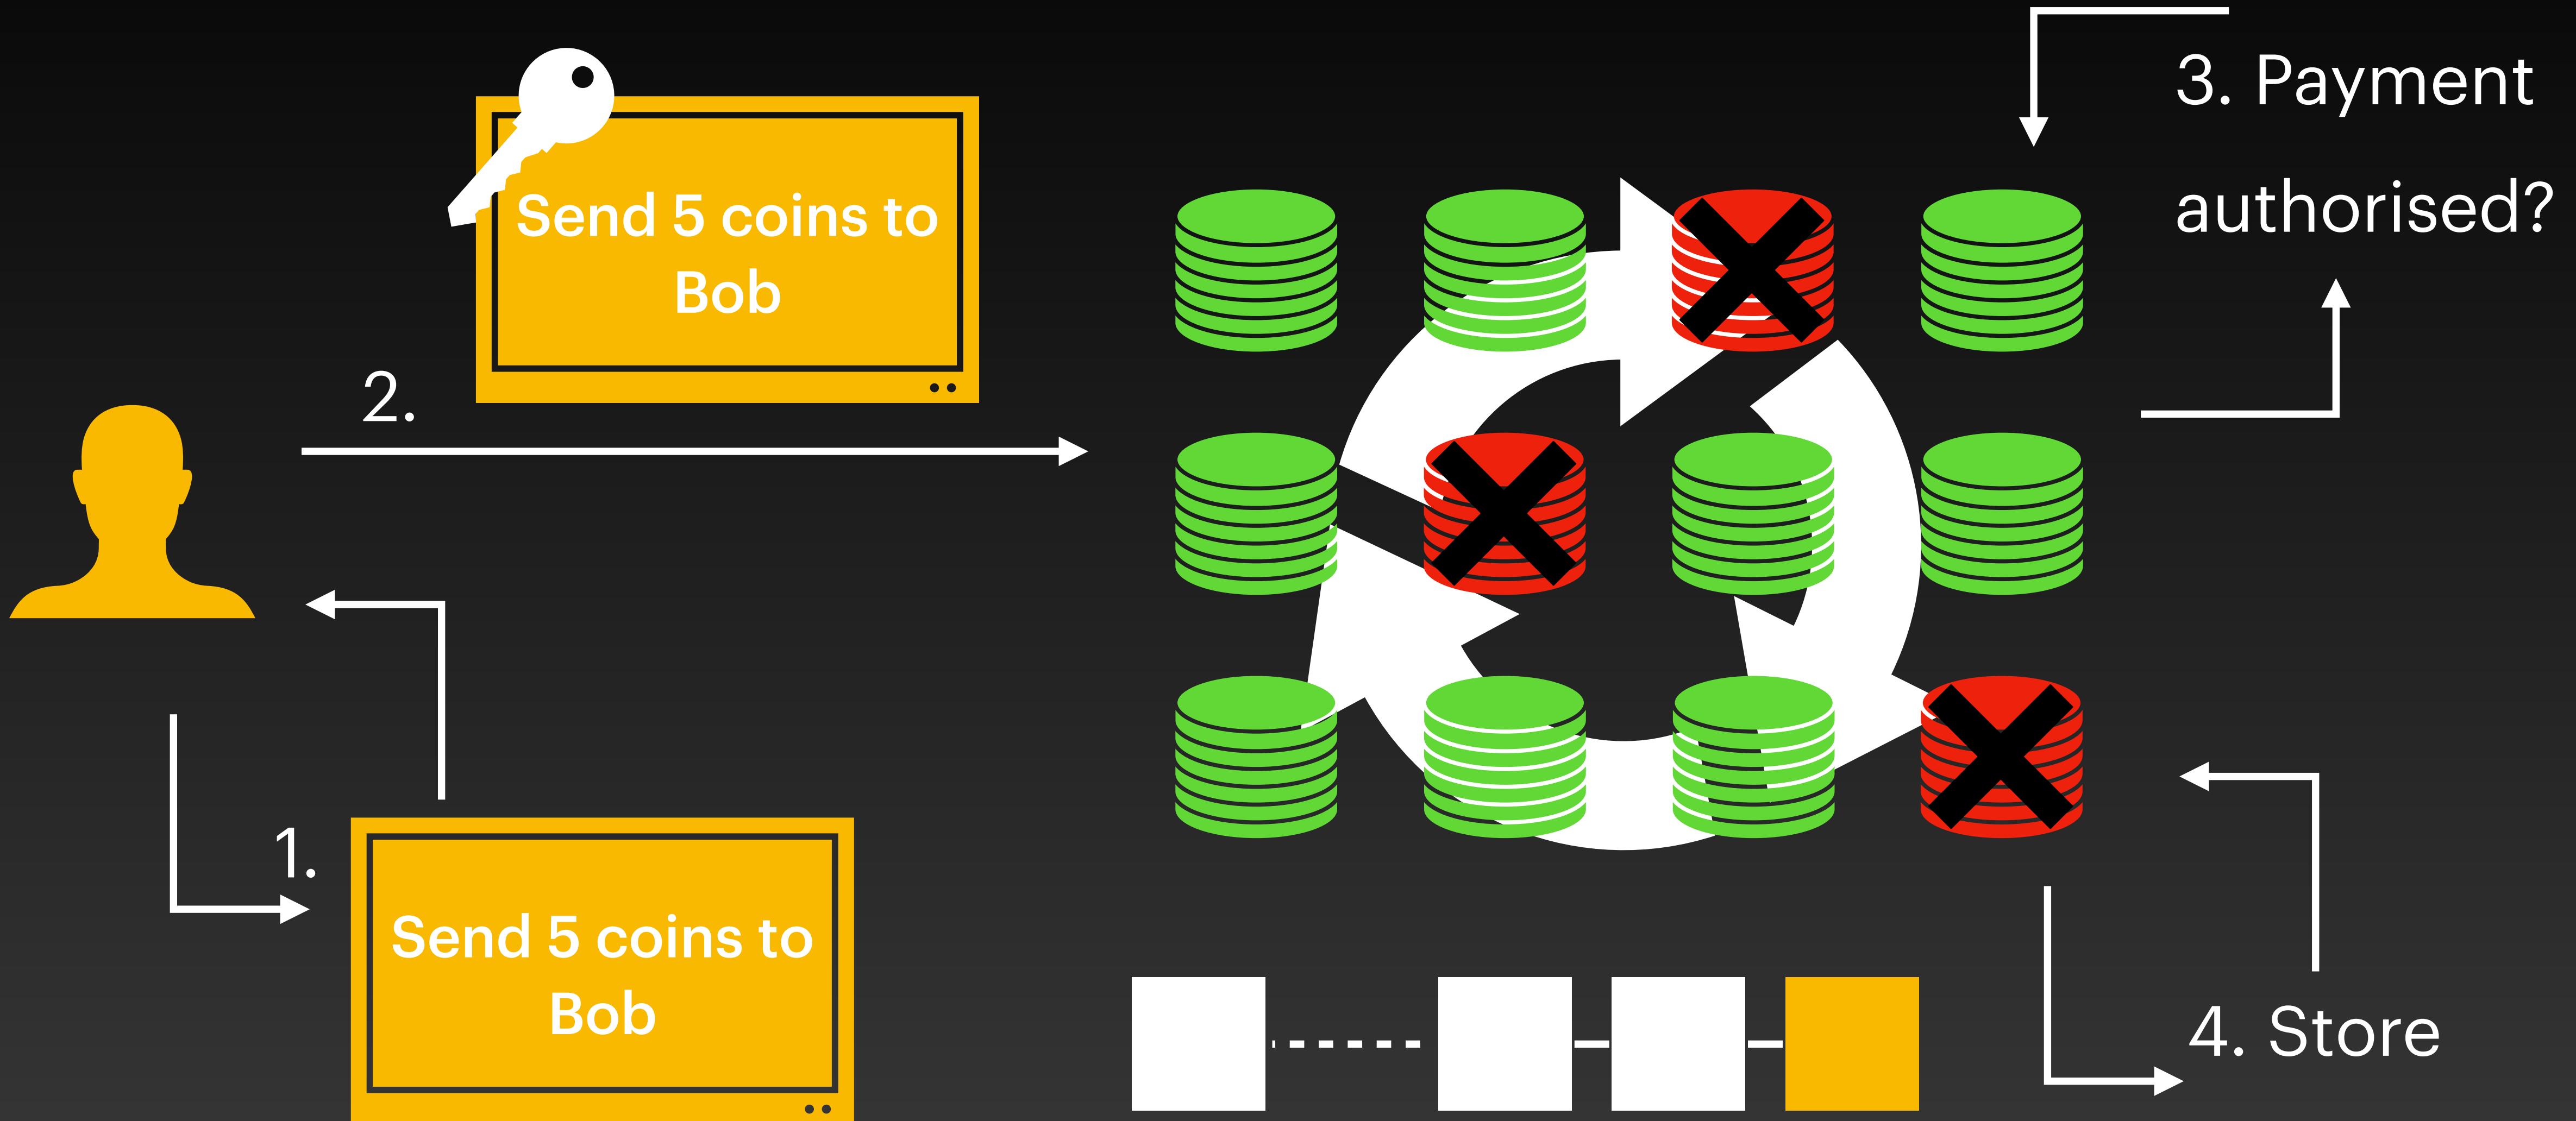
Blockchains



Blockchains



The Best Example



Keeping the Talk Short

In scope

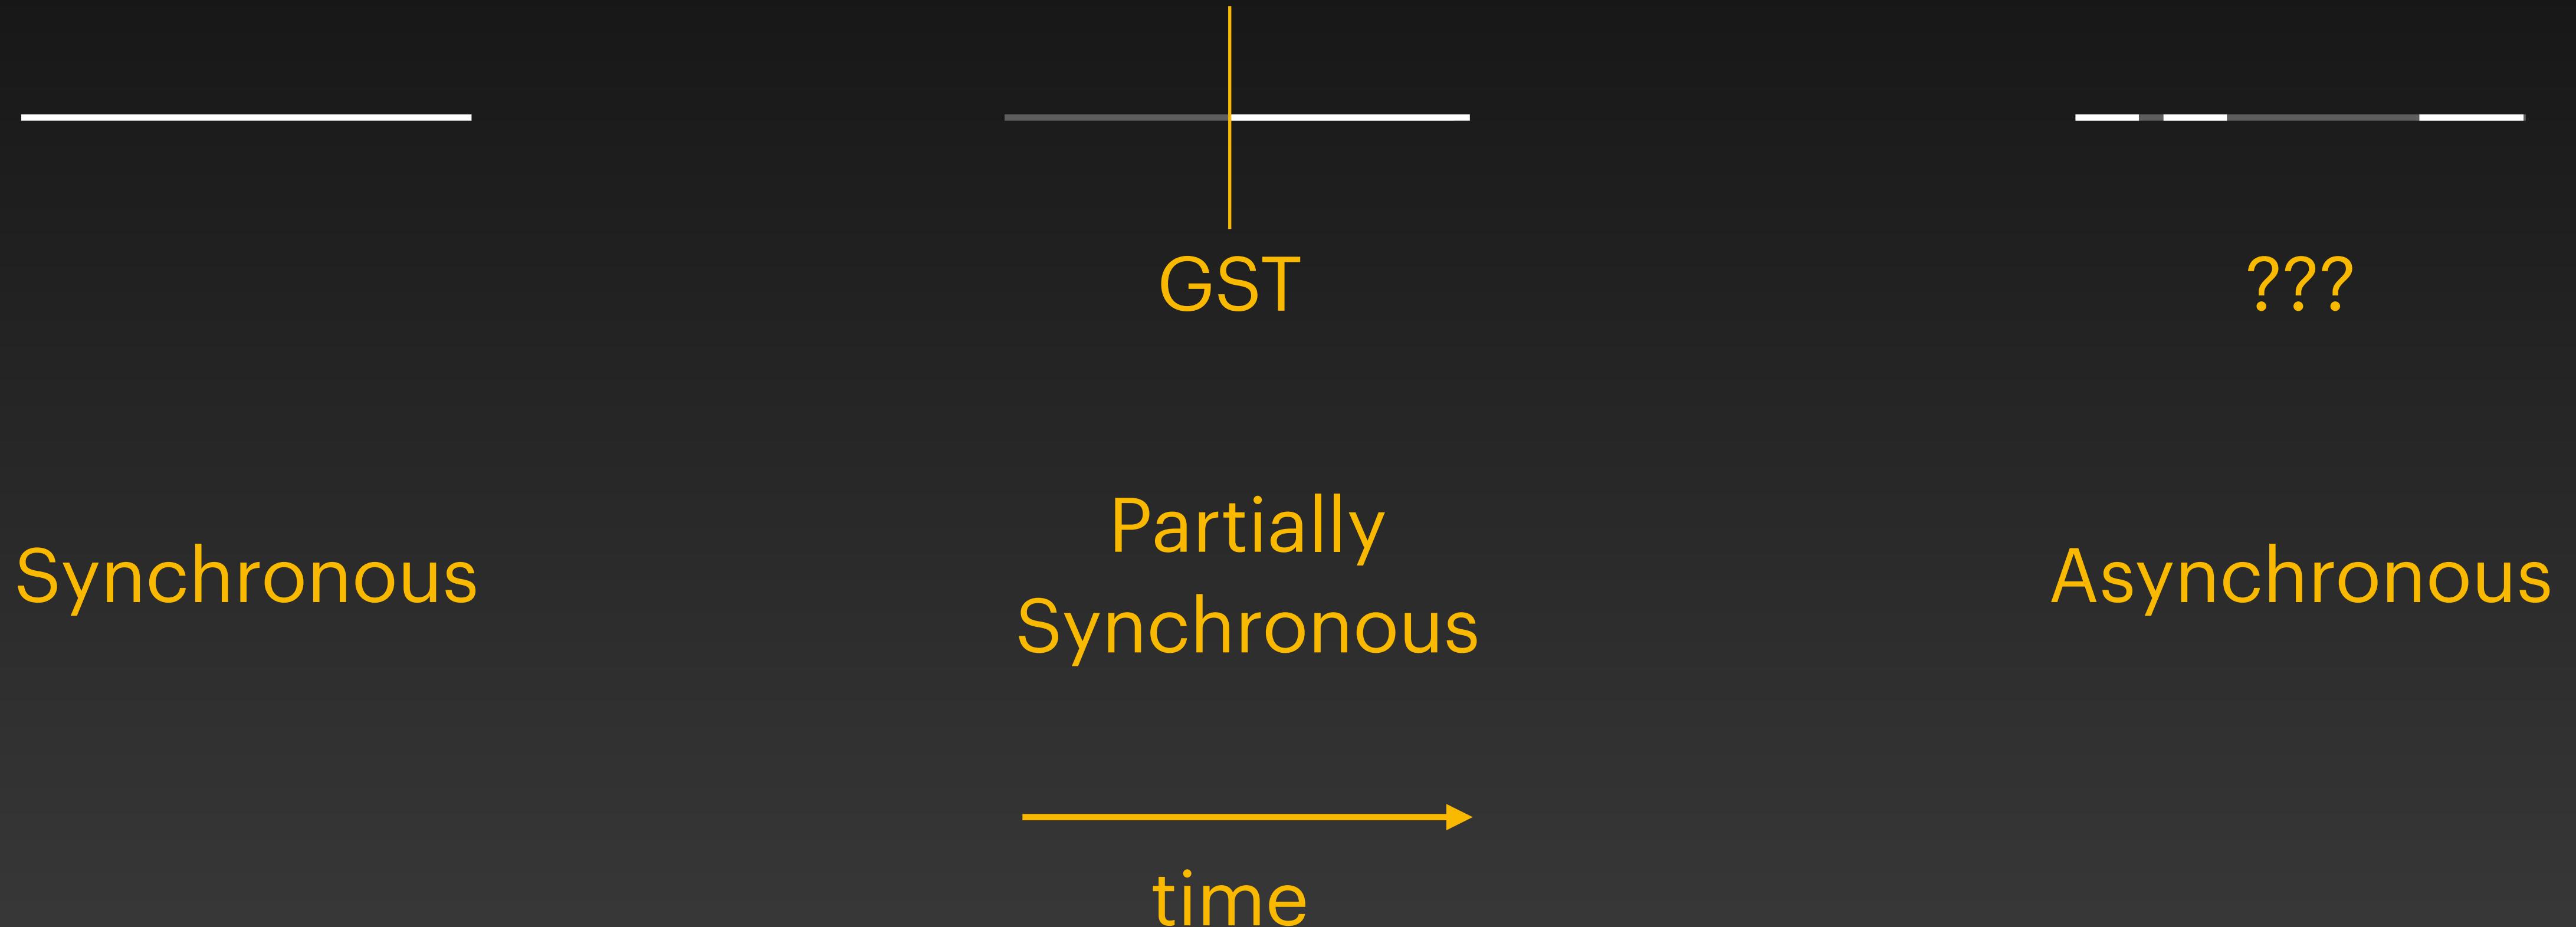
- Ordering (quorum-based)



Not in scope

- Nodes selection?
- Committee reconfiguration?
- Transactions execution?
- Transactions language?
- Financial incentives?
- etc

3 Timing Models



What do we want in a consensus protocol?

Simplicity

Dual mode is complicated

Performance

Can we improve throughput?

Resilience

Can we resist DDoS?

Can we have all three?

Recent Related Consensus Protocols

DAG-Rider: Asynchronous protocol needing **12 messages** to commit

Tusk: Asynchronous protocol but needs **9 messages** to commit

Mysticeti: Fast and simple protocol but **only partially synchronous**

FIN: Asynchronous protocol but **low throughput**

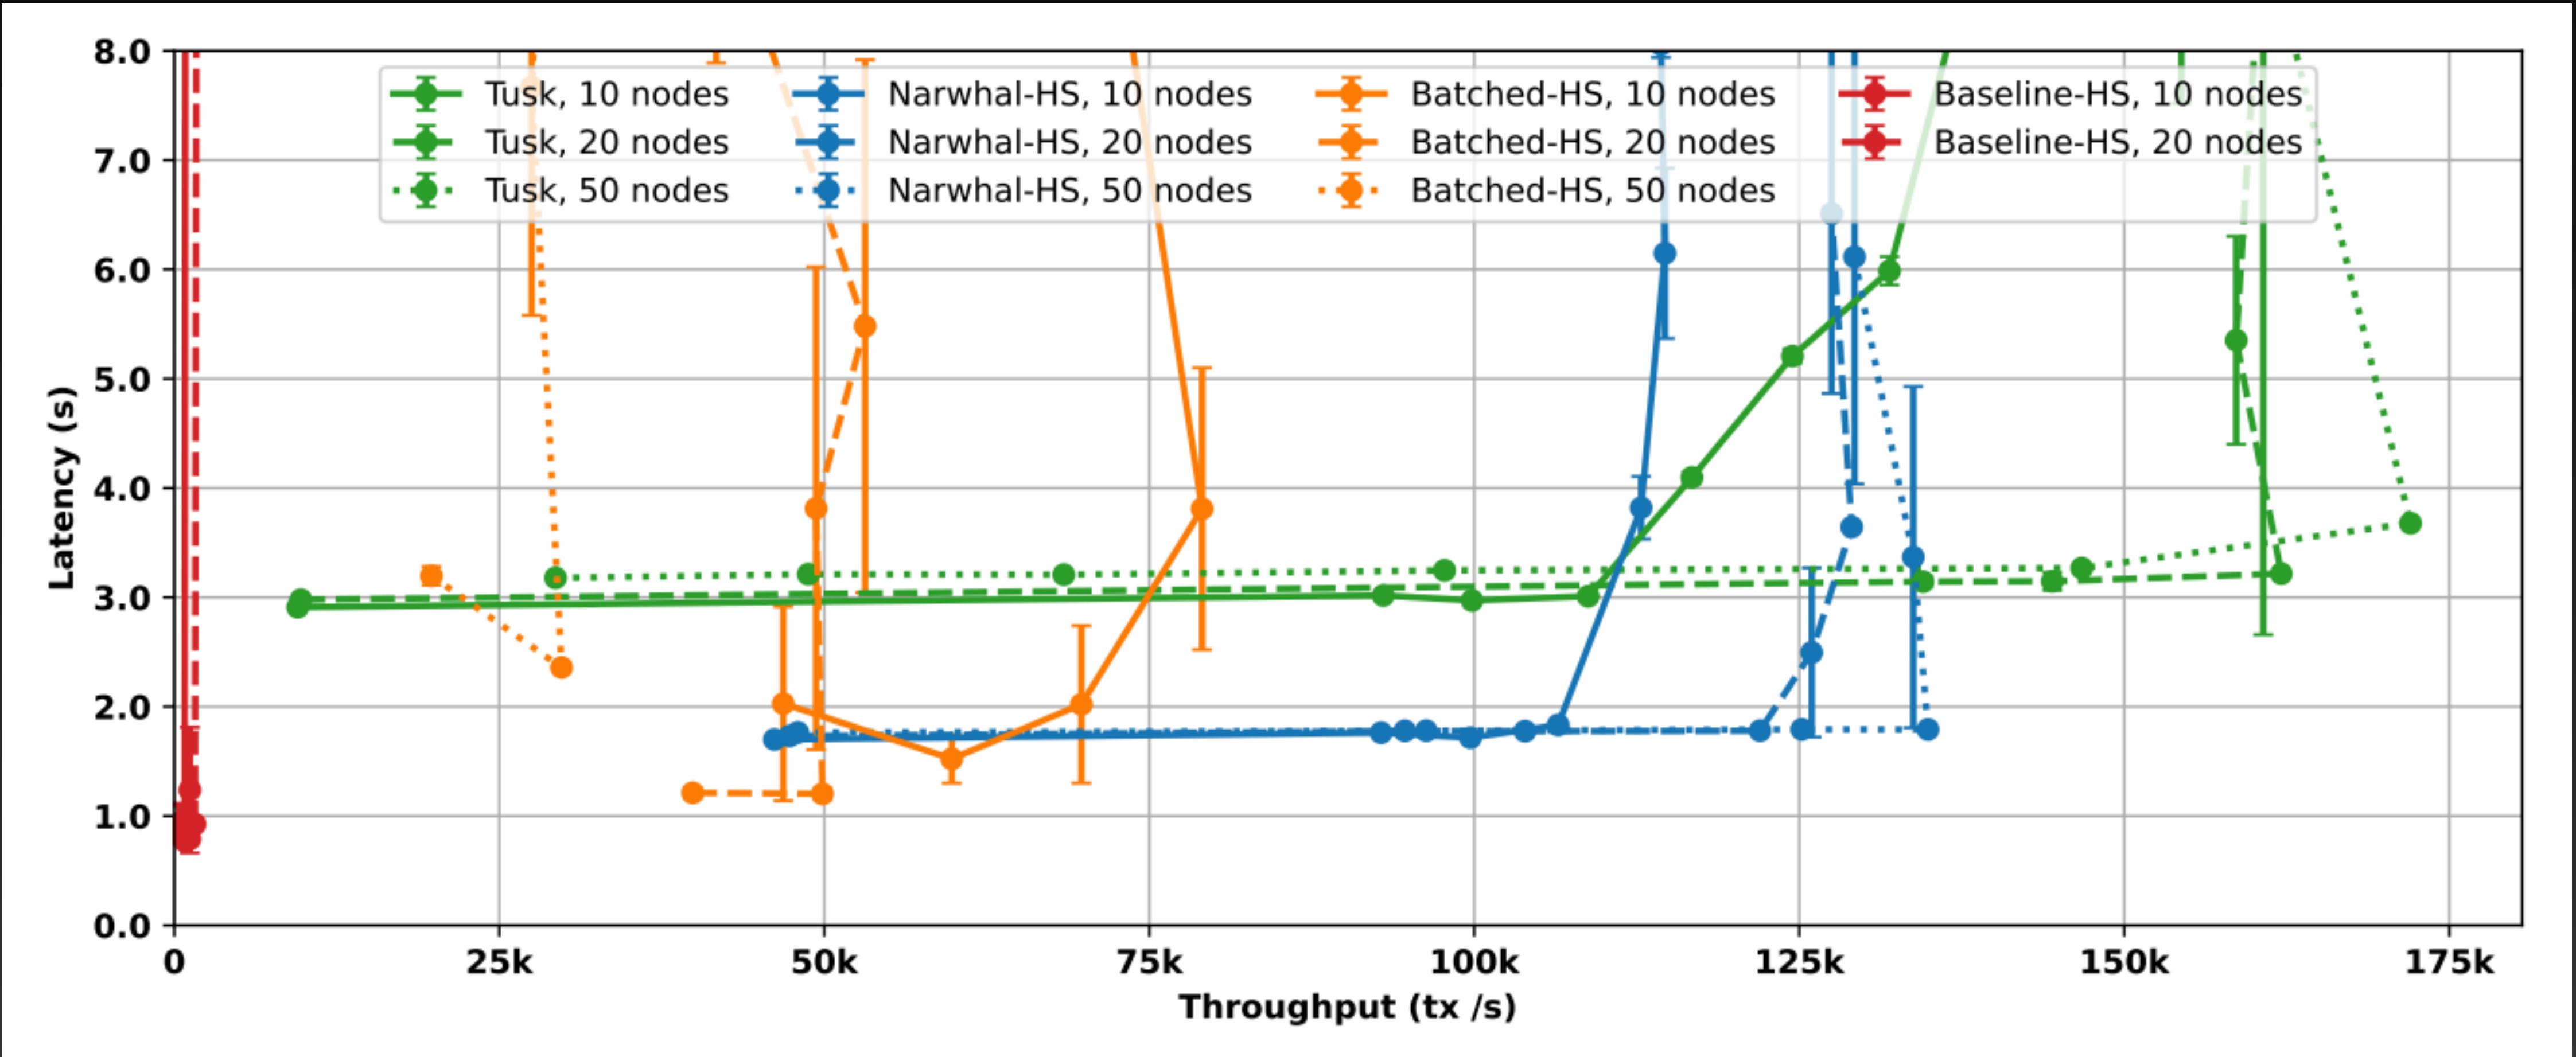
Cordial Miners: Asynchronous protocol but **higher latency**

Why Mahi-Mahi?

Fast, Simple, Resilient

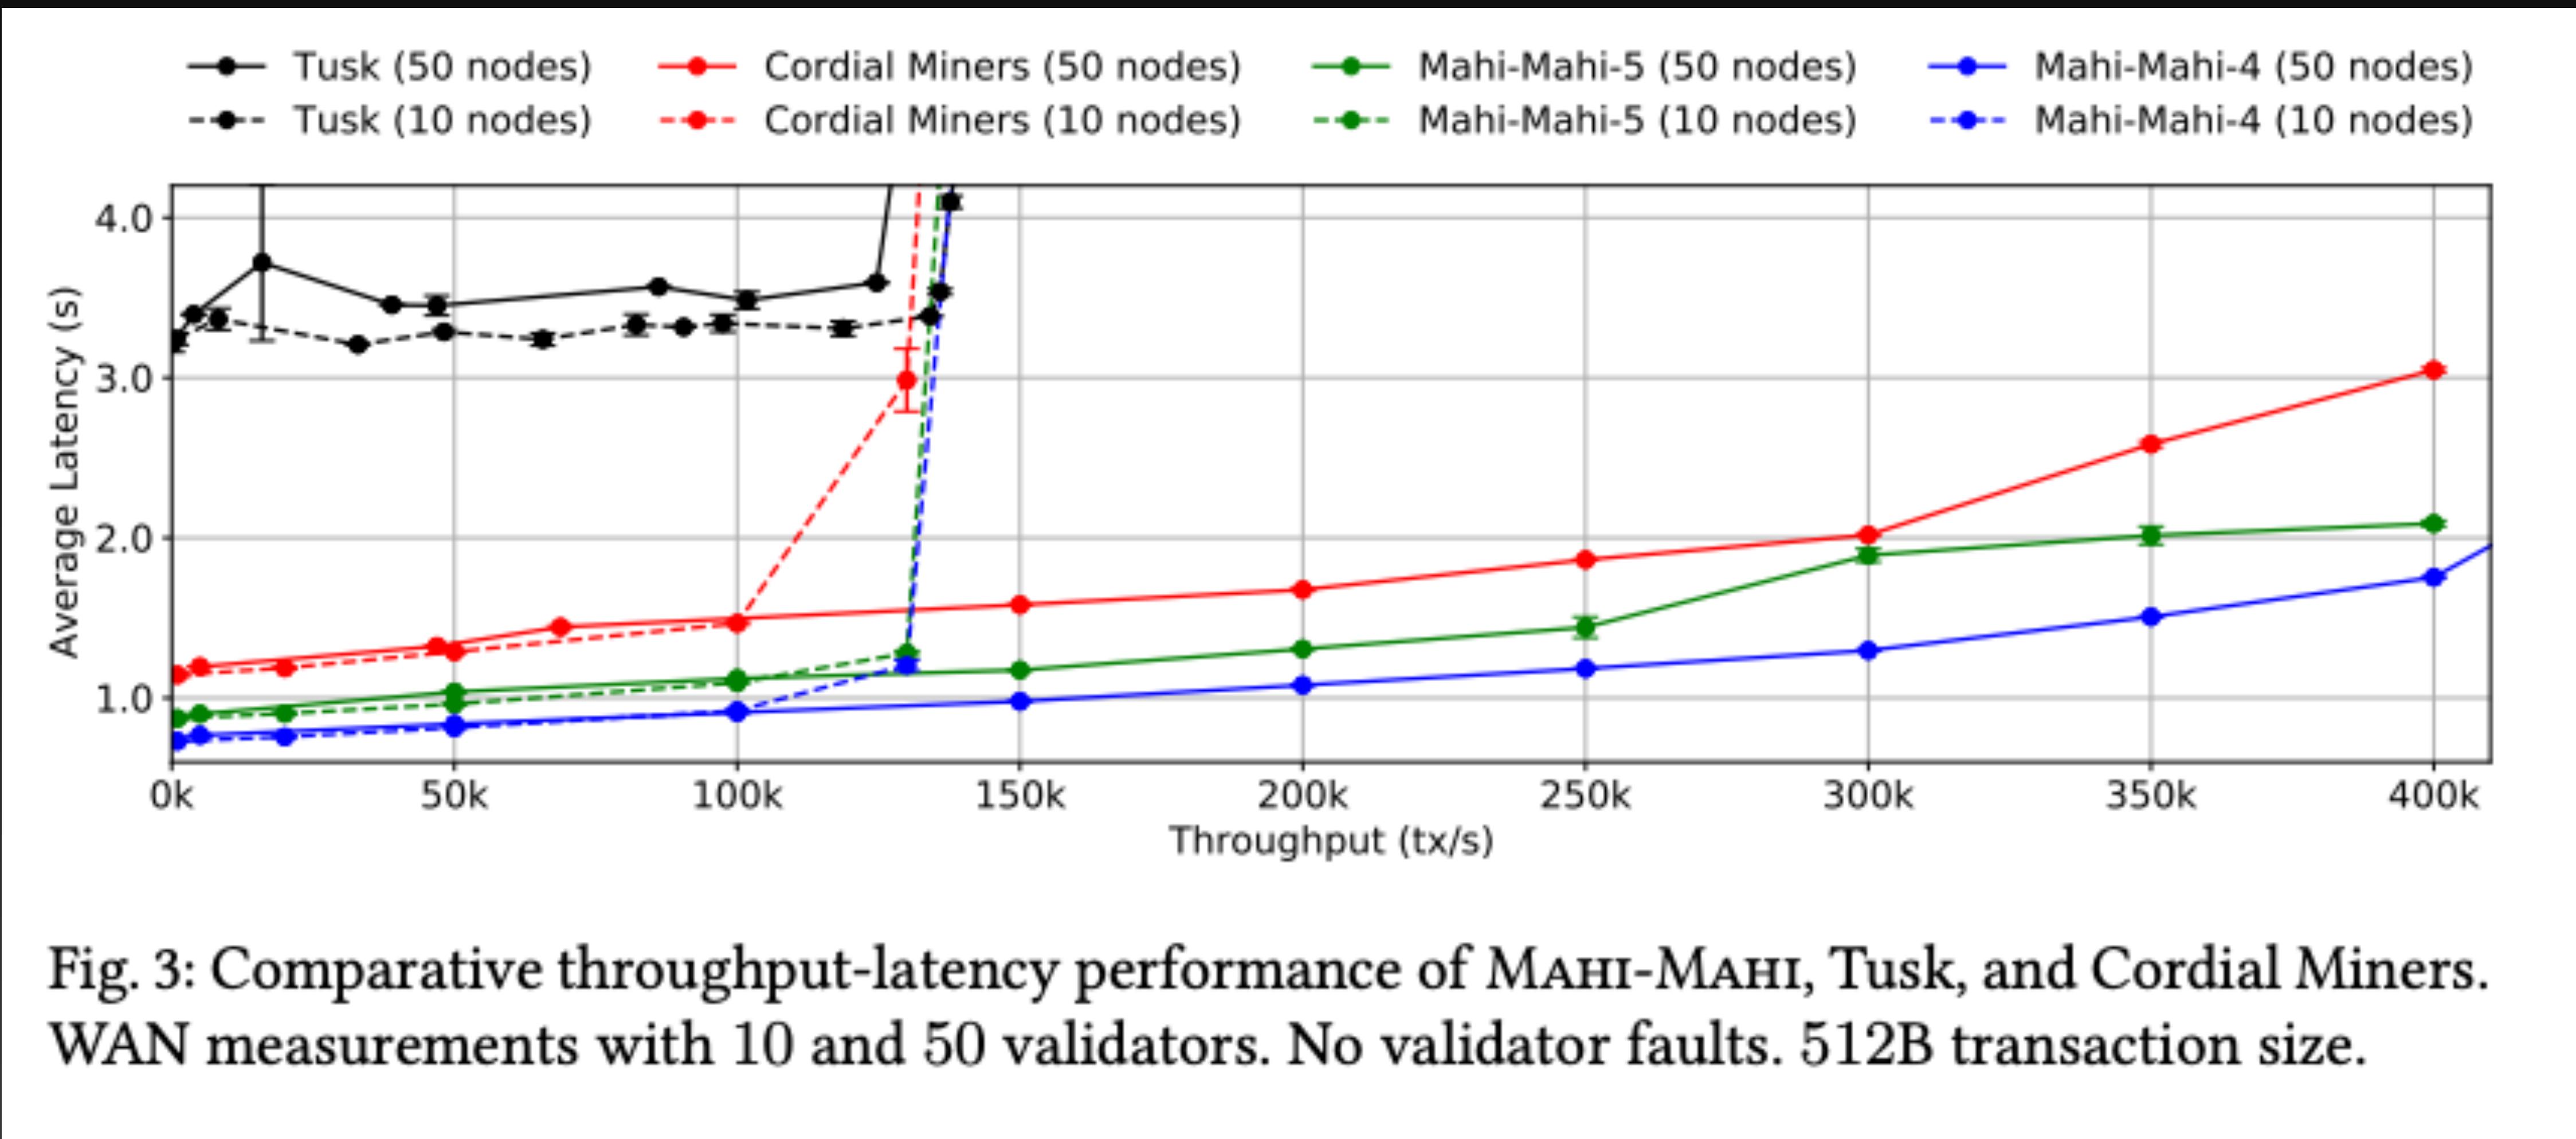
- Is live and can progress under asynchronous conditions
- Uncertified DAG to reduce latency
- Sub-second latency with >100k throughput per second
- Simple to implement
- Can be parameterized

How did previous protocols perform?



How Fast is Mahi-Mahi?

Faster than any prior protocols!

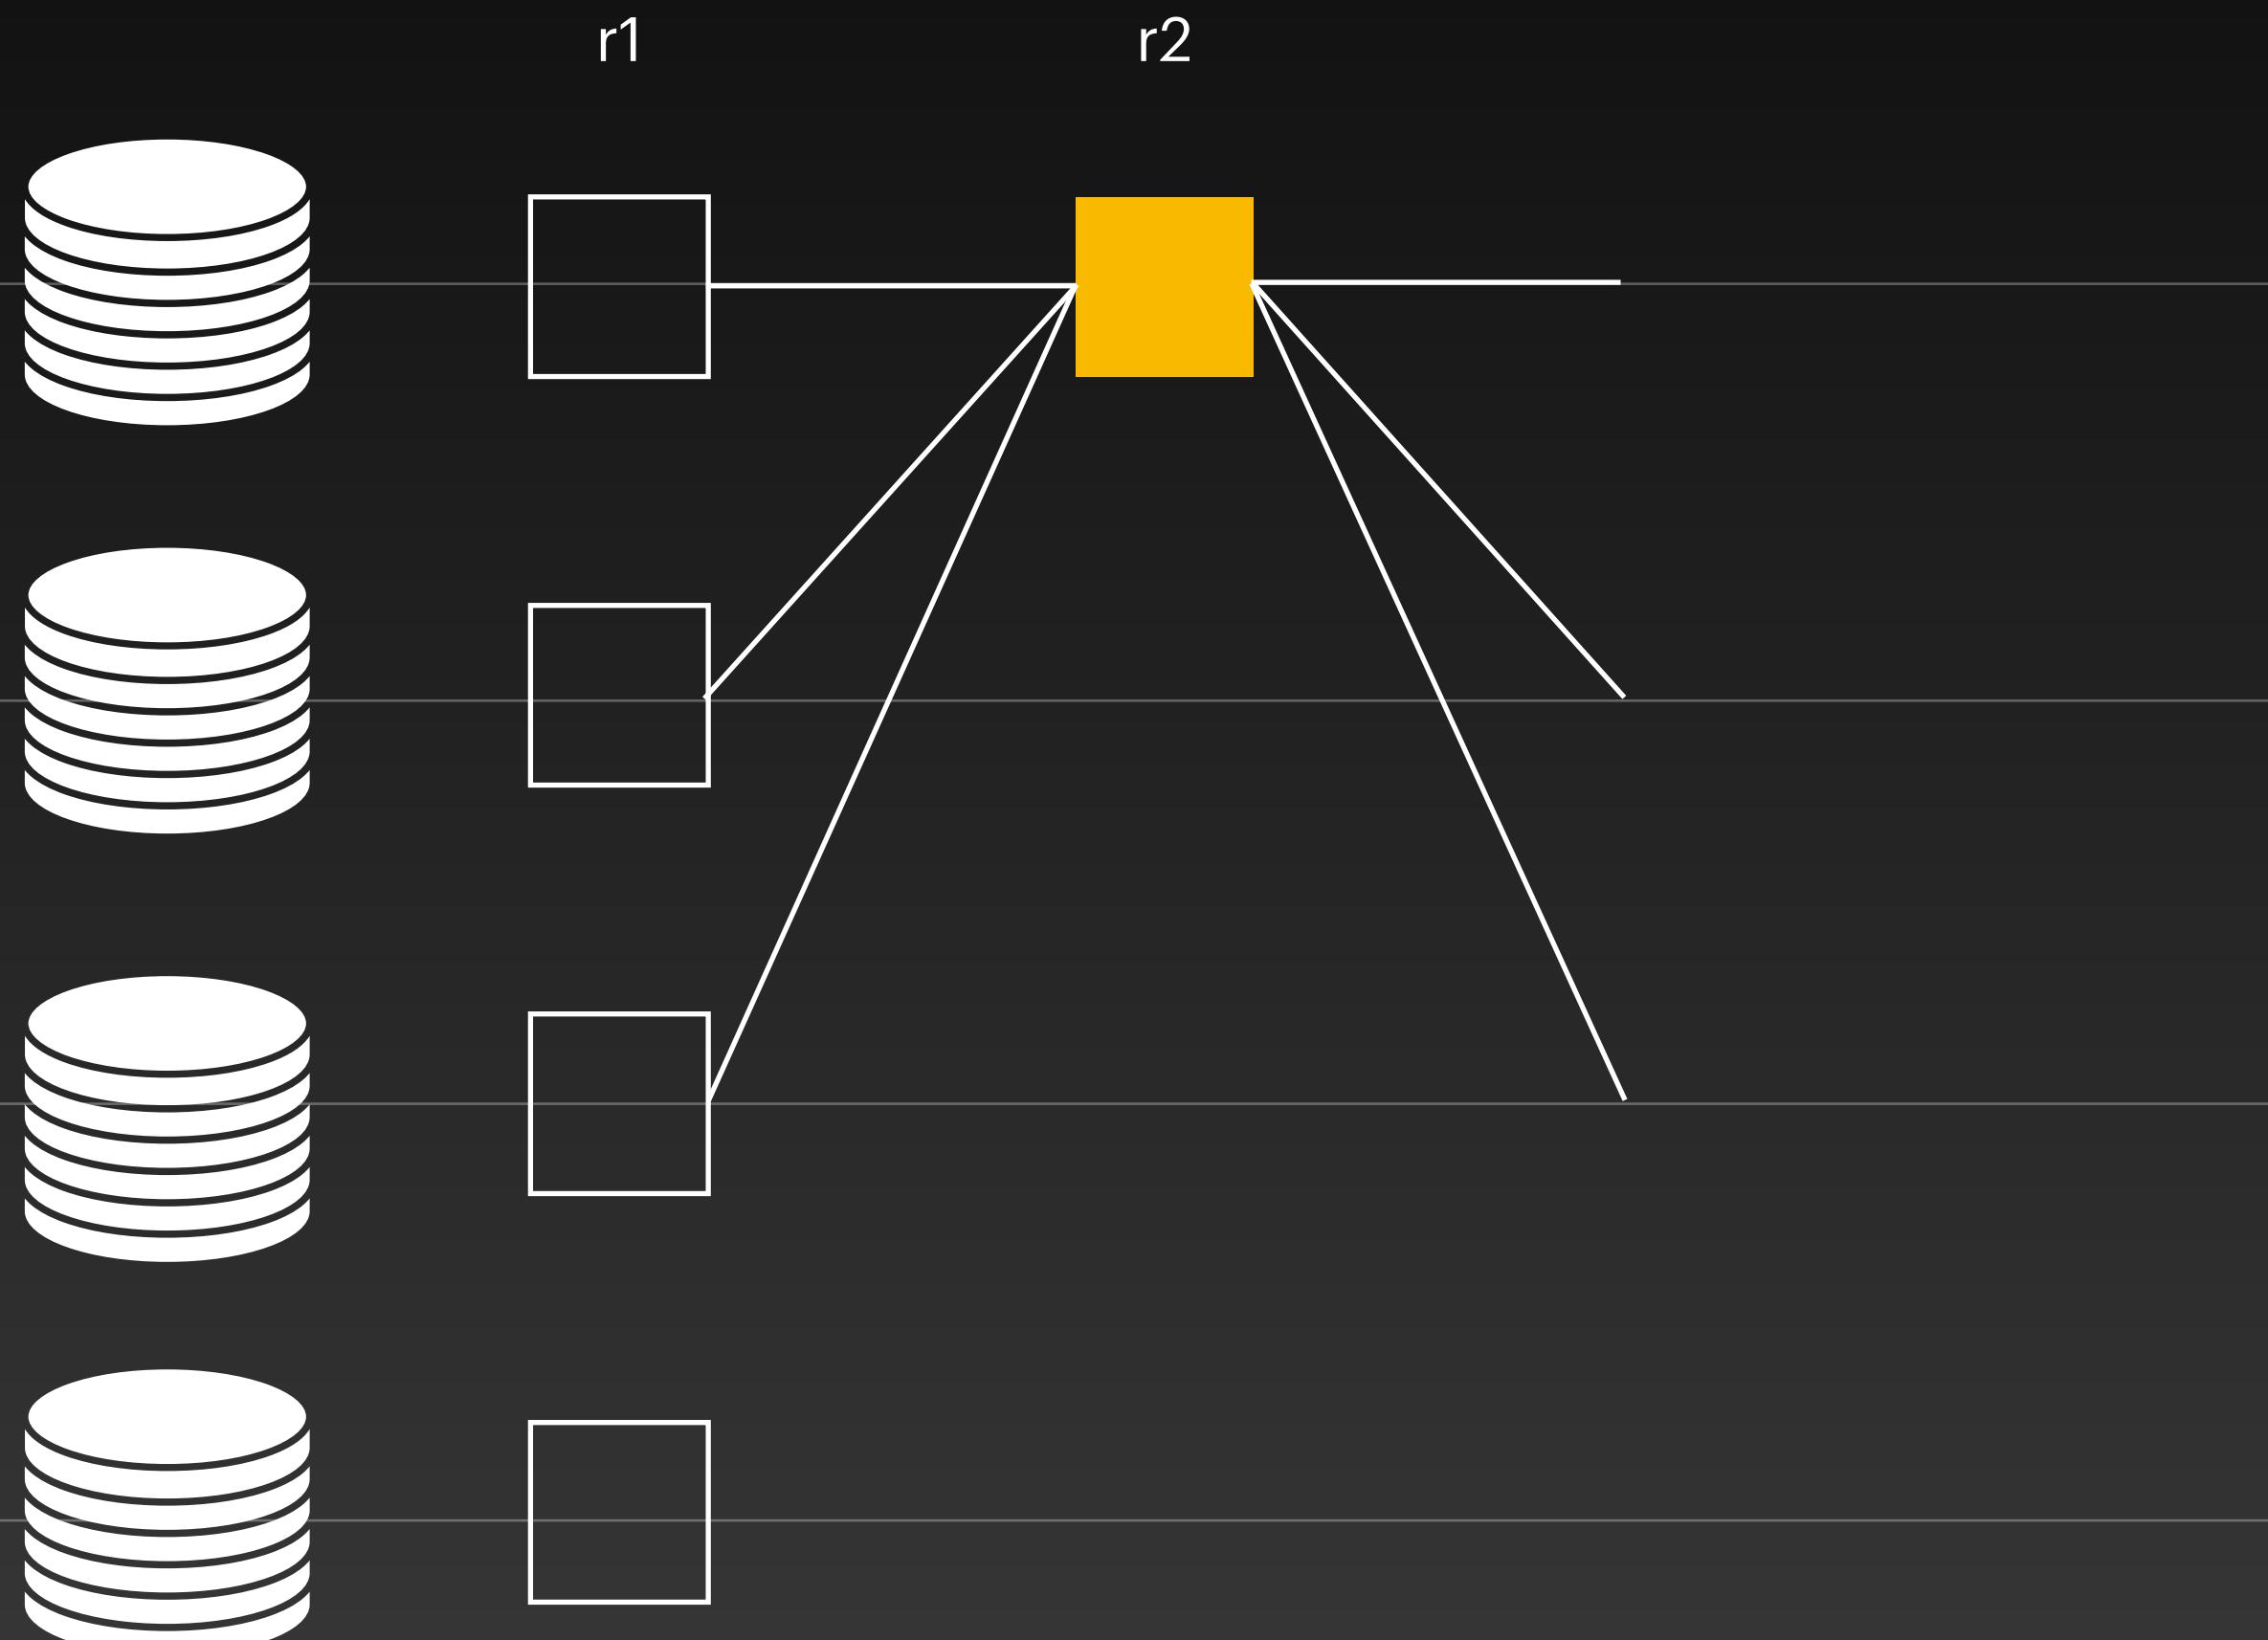


The Mahi-Mahi DAG

Implicit Certification Consensus

Uncertified DAGs

Block Creation

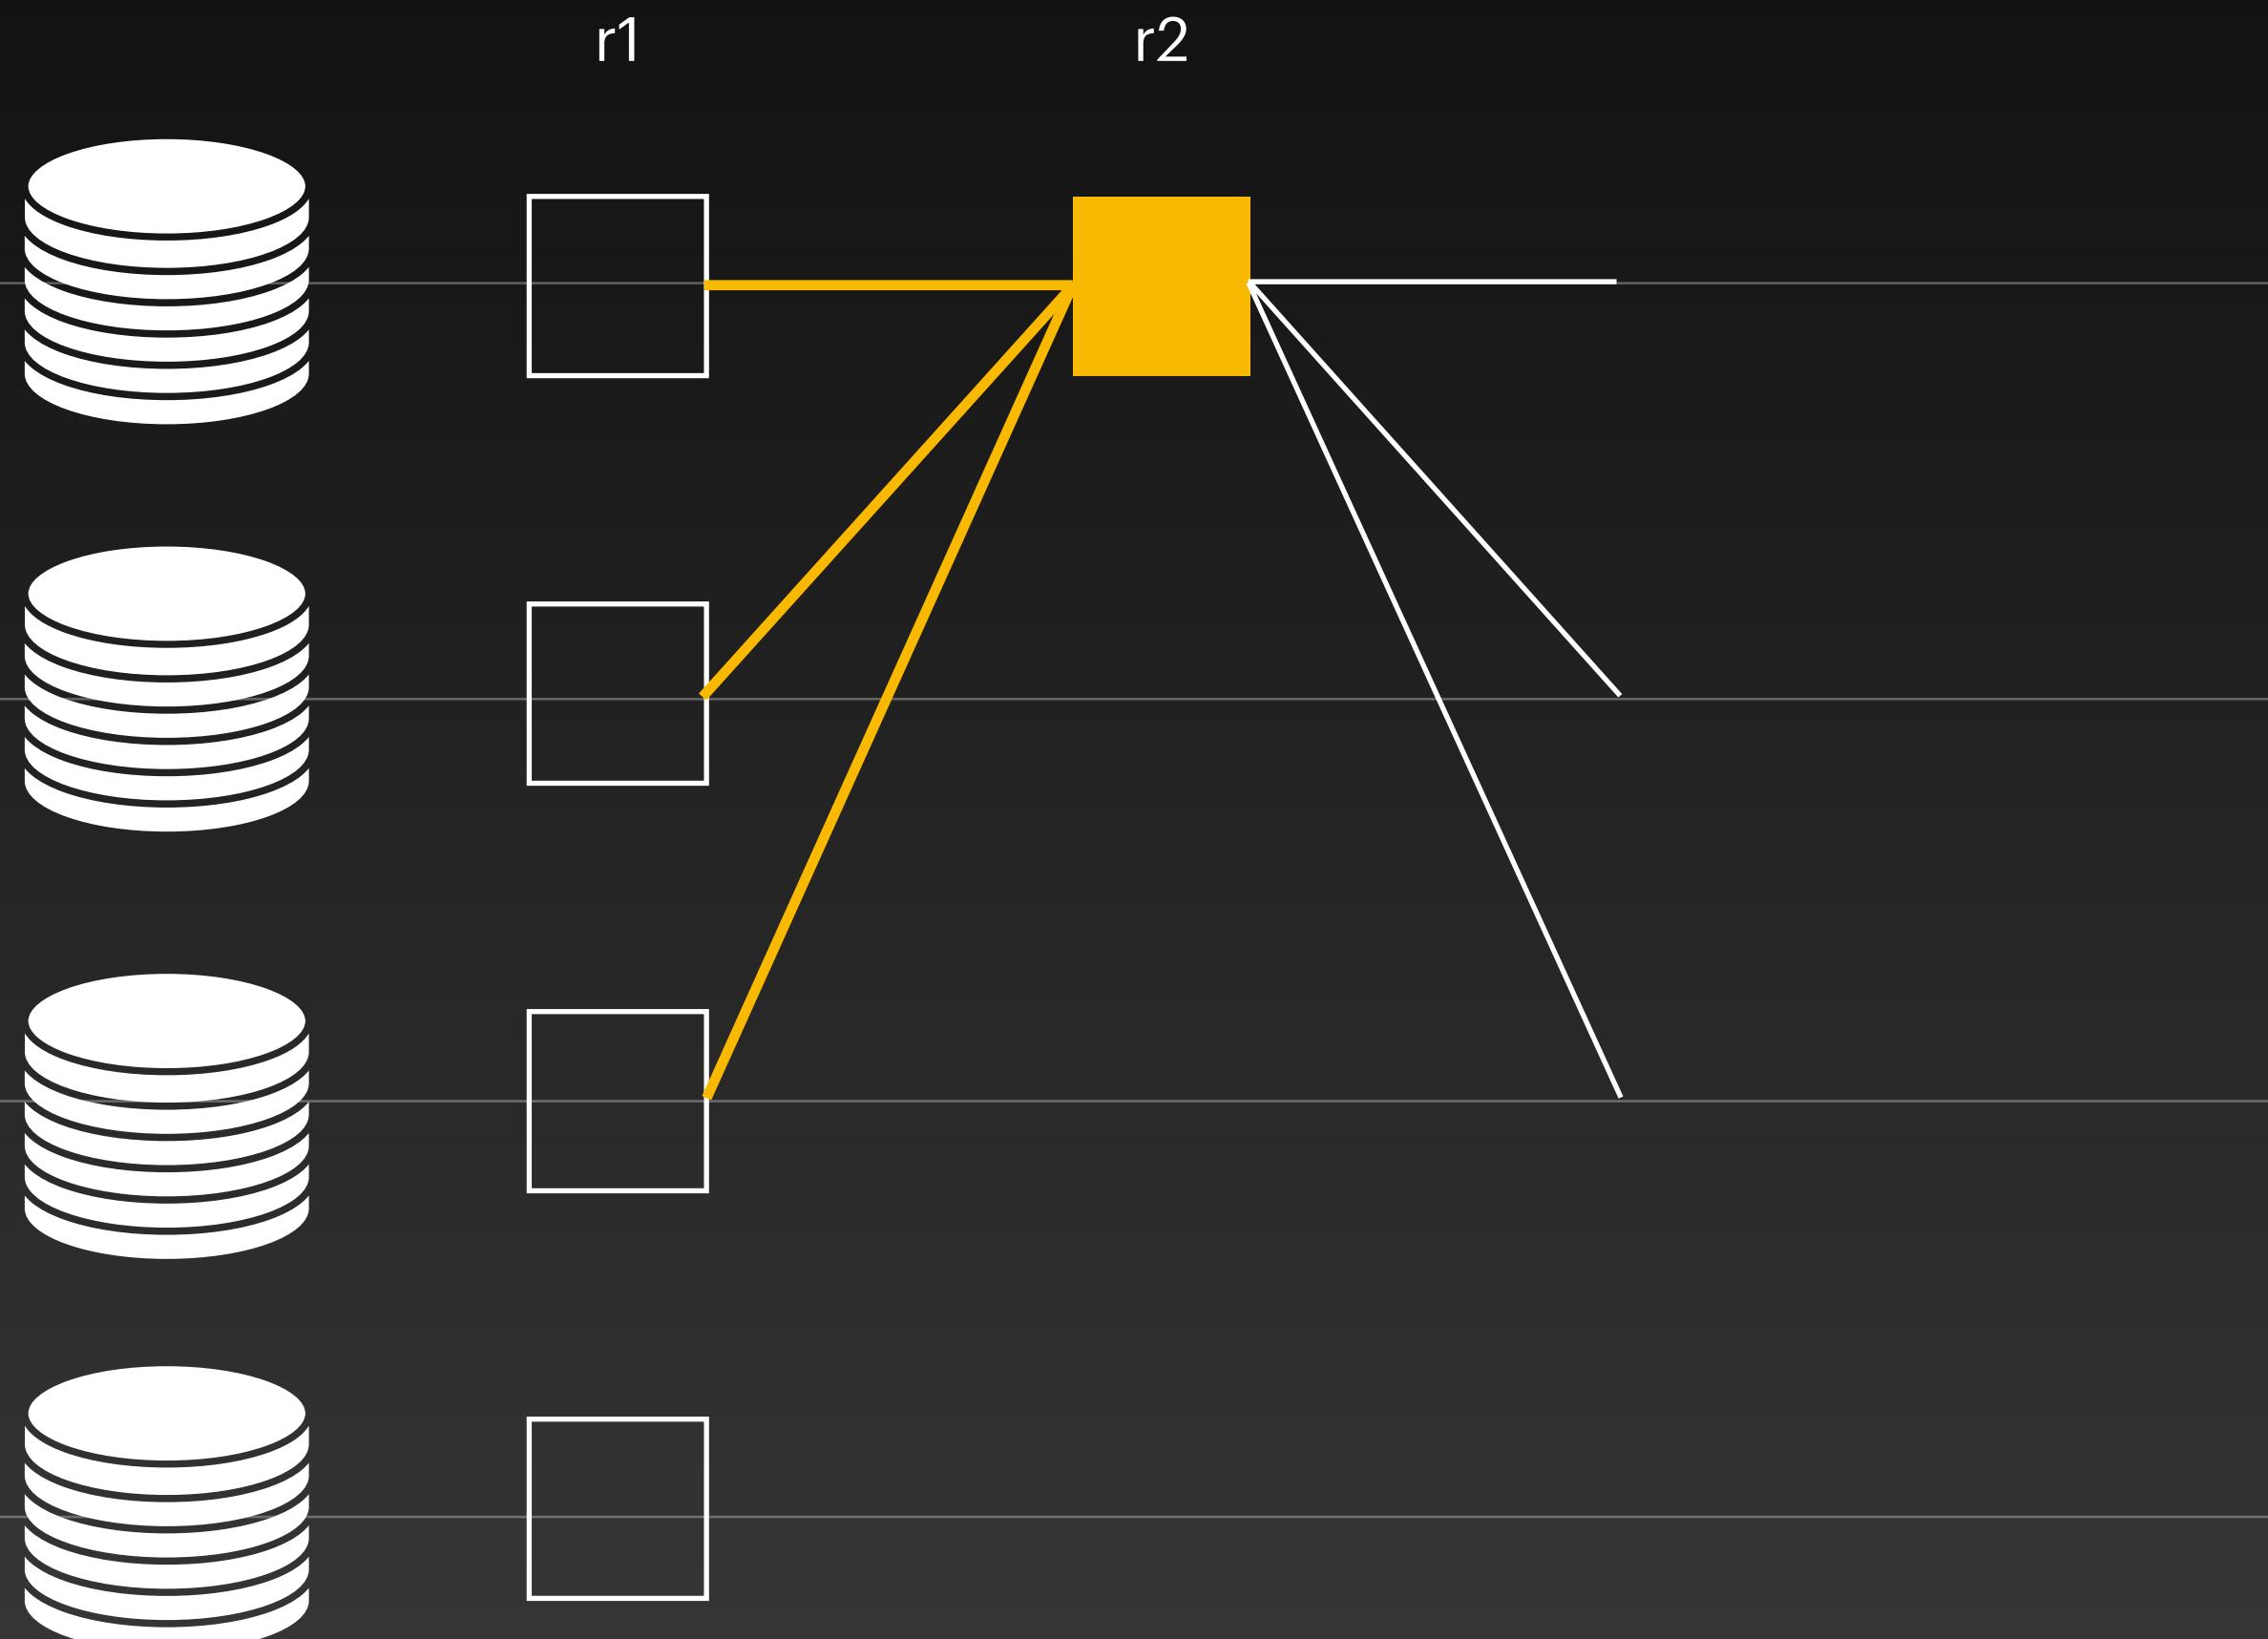


What's in a block:

- Round number
- Author
- Payload (transactions)
- Signature

Uncertified DAGs

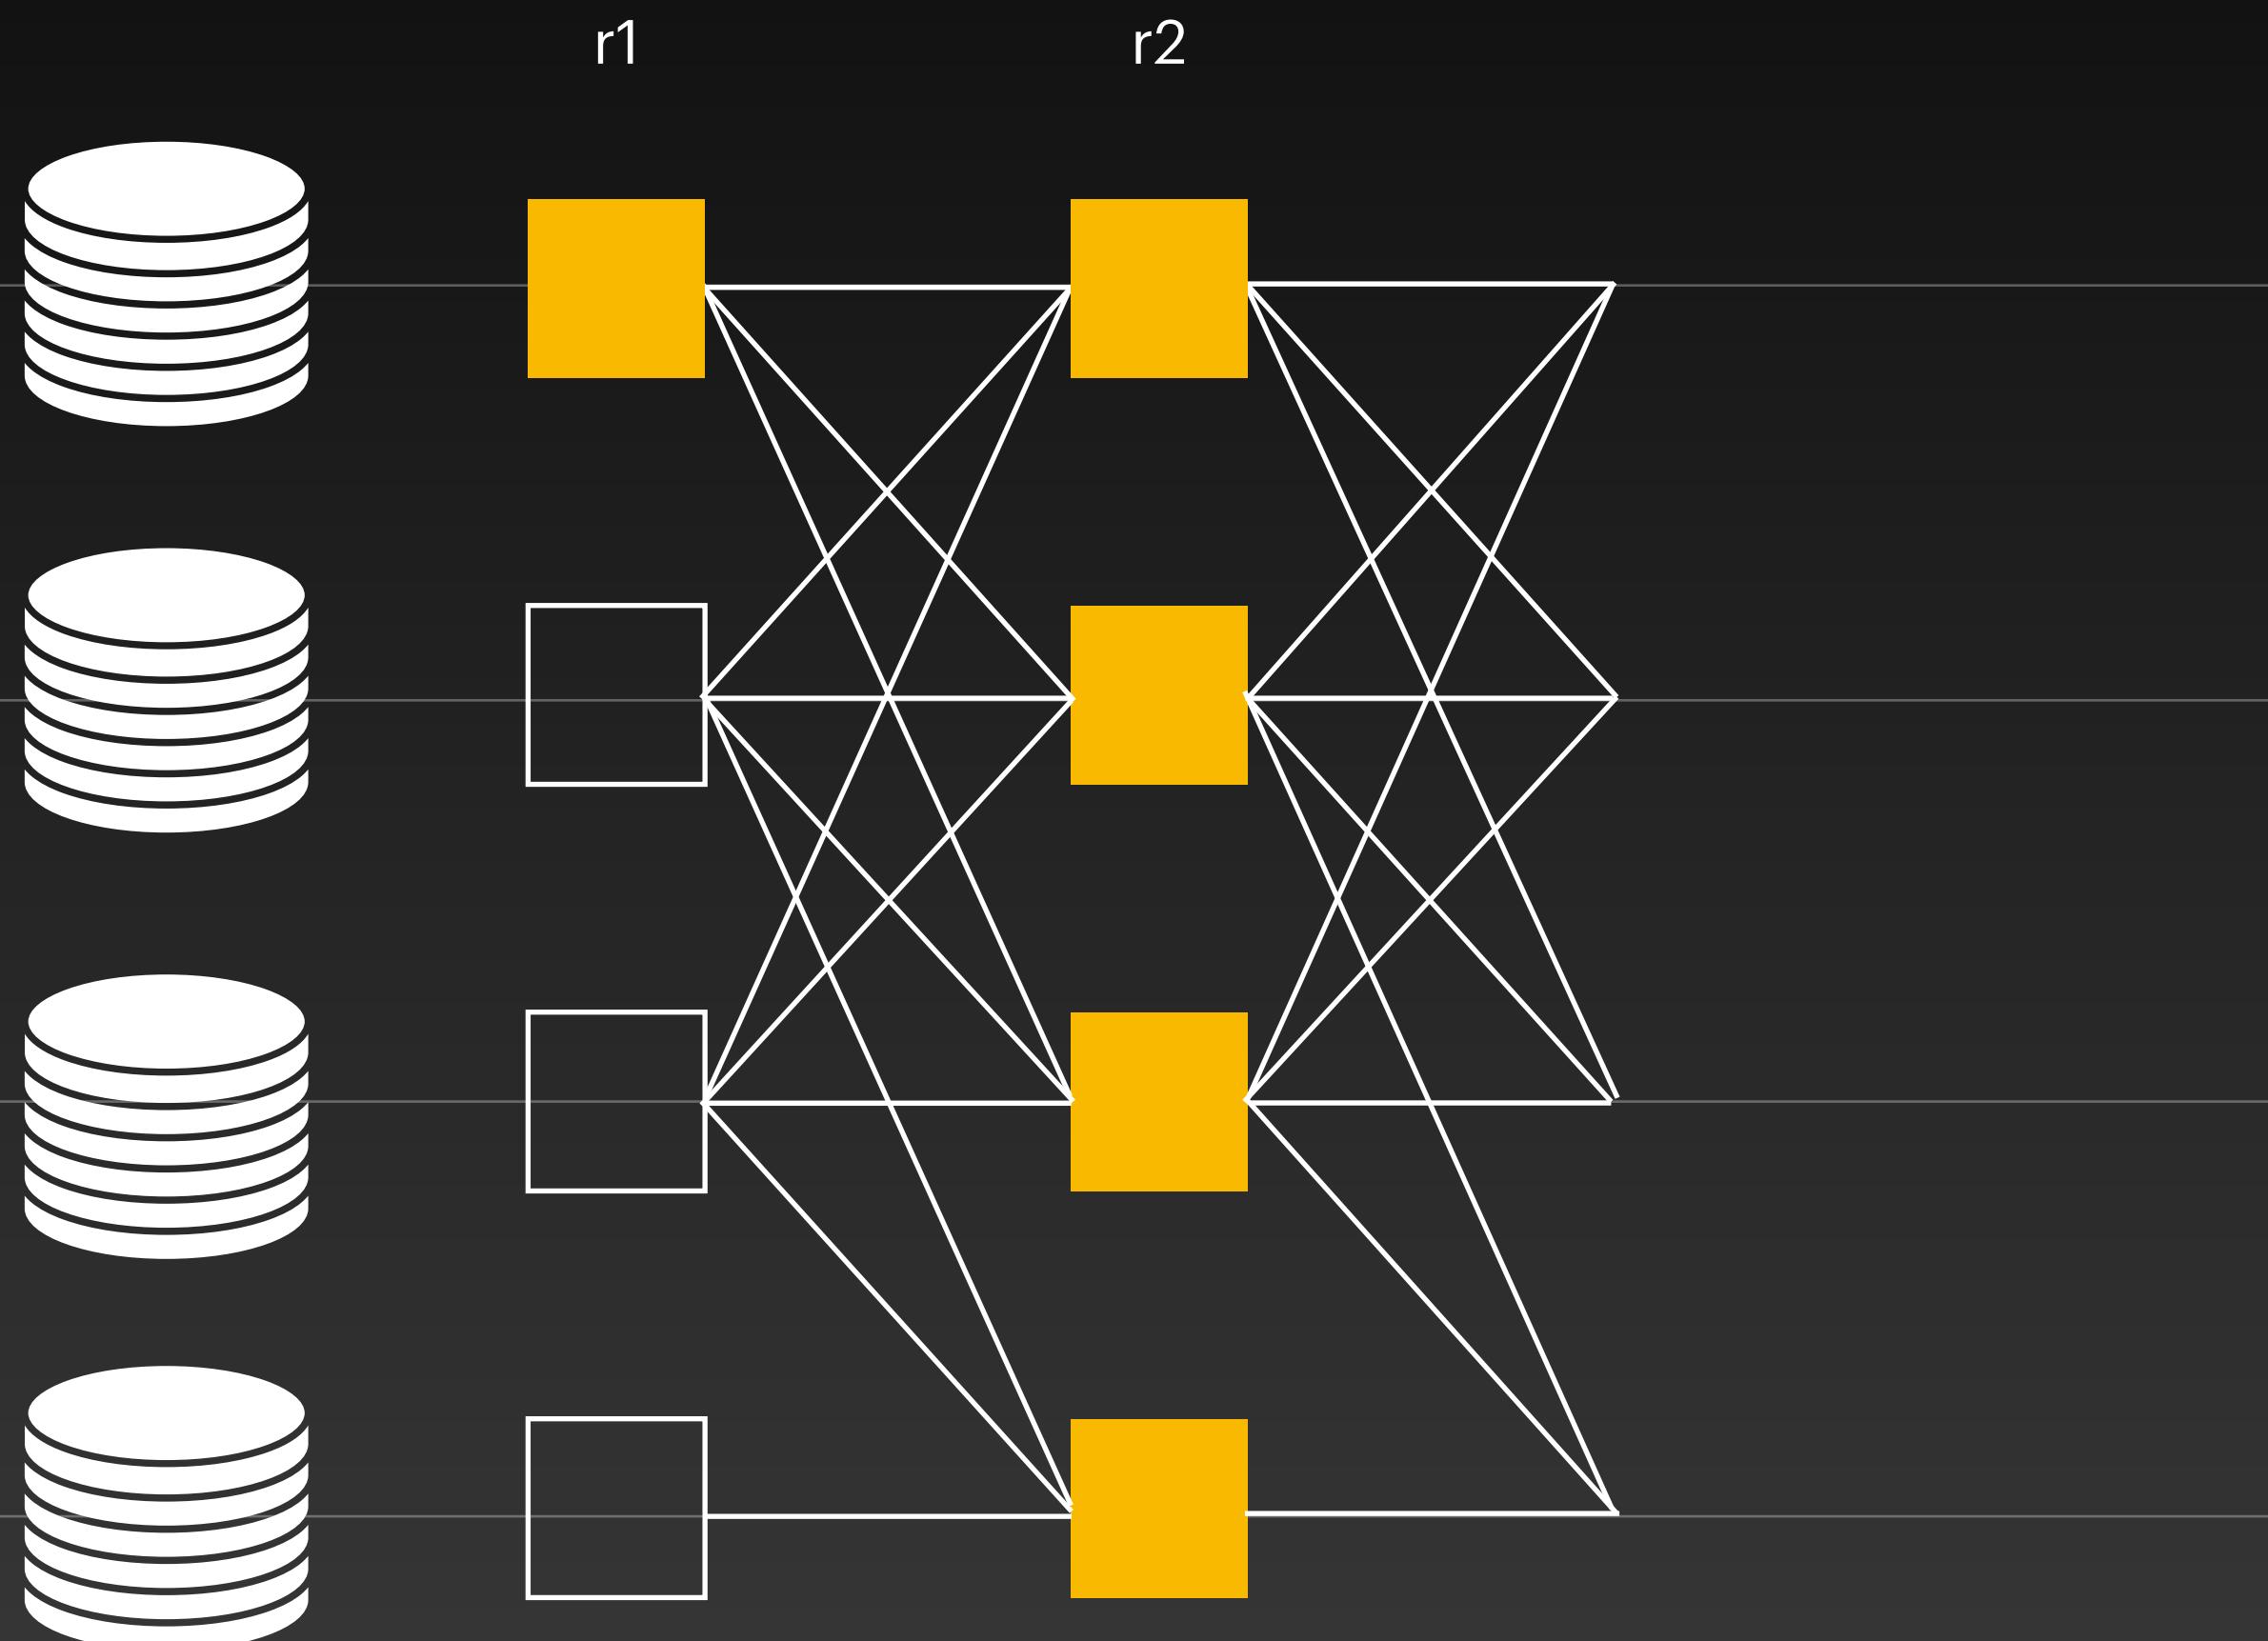
Rule 1: Link to $2f+1$ parents



- Total nodes: **$3f+1 = 4$**
- Quorum: **$2f+1 = 3$**
- Generate randomness to retrospectively elect leaders with common coin

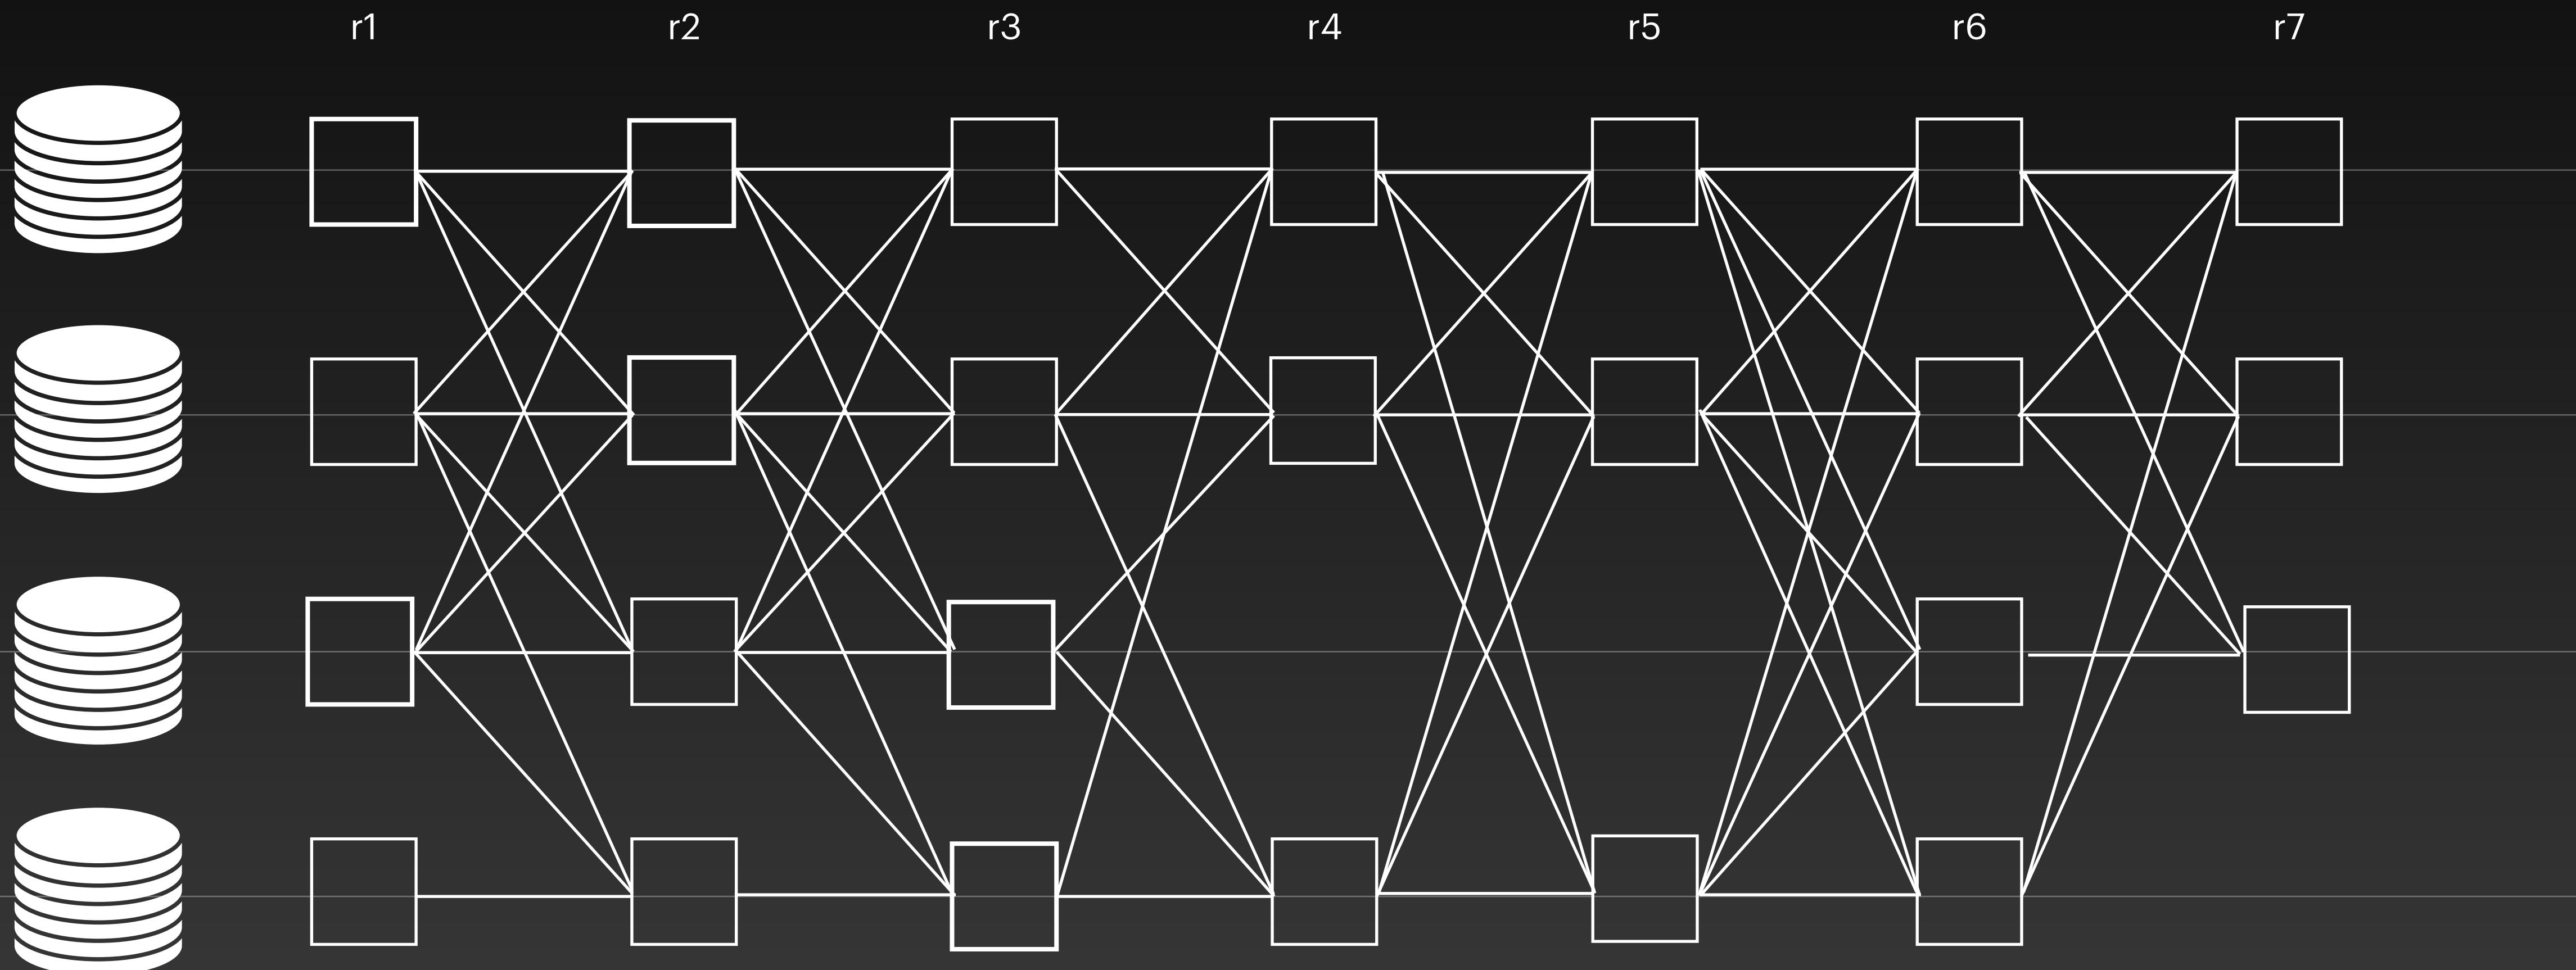
Uncertified DAGs

Rule 2: All nodes run in parallel



Uncertified DAGs

An Example



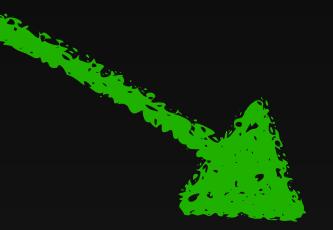
Main Ingredient:

All messages embedded in the DAG

- Fewer signatures
- Isolated engineering component
- Define interpretable patterns on the DAG
- Run multiple protocols on the same DAG

Certified vs Uncertified DAG

Mahi-Mahi is an Uncertified DAG



Certified DAG

- Every round is a reliable broadcast
- $2f+1$ quorum

Uncertified DAG

- Every round is a best effort dissemination
- $2f+1$ quorum

How to we elect a Leader?

Blocks should be committed



Need Randomness

- Cannot reach consensus without randomness (FLP Theorem)
- Randomness retrospectively determines a leader

Use Common Coin

- Generates a random number for honest participants
- Require $2f+1$ shares to execute coin toss.
- Can be implemented with threshold signatures

What do we want in a Consensus Protocol?

Validity

If an honest participant outputs a message, it is eventually delivered to every honest participant.

Total Order

All honest participants will output a consistent sequence of messages.

Agreement

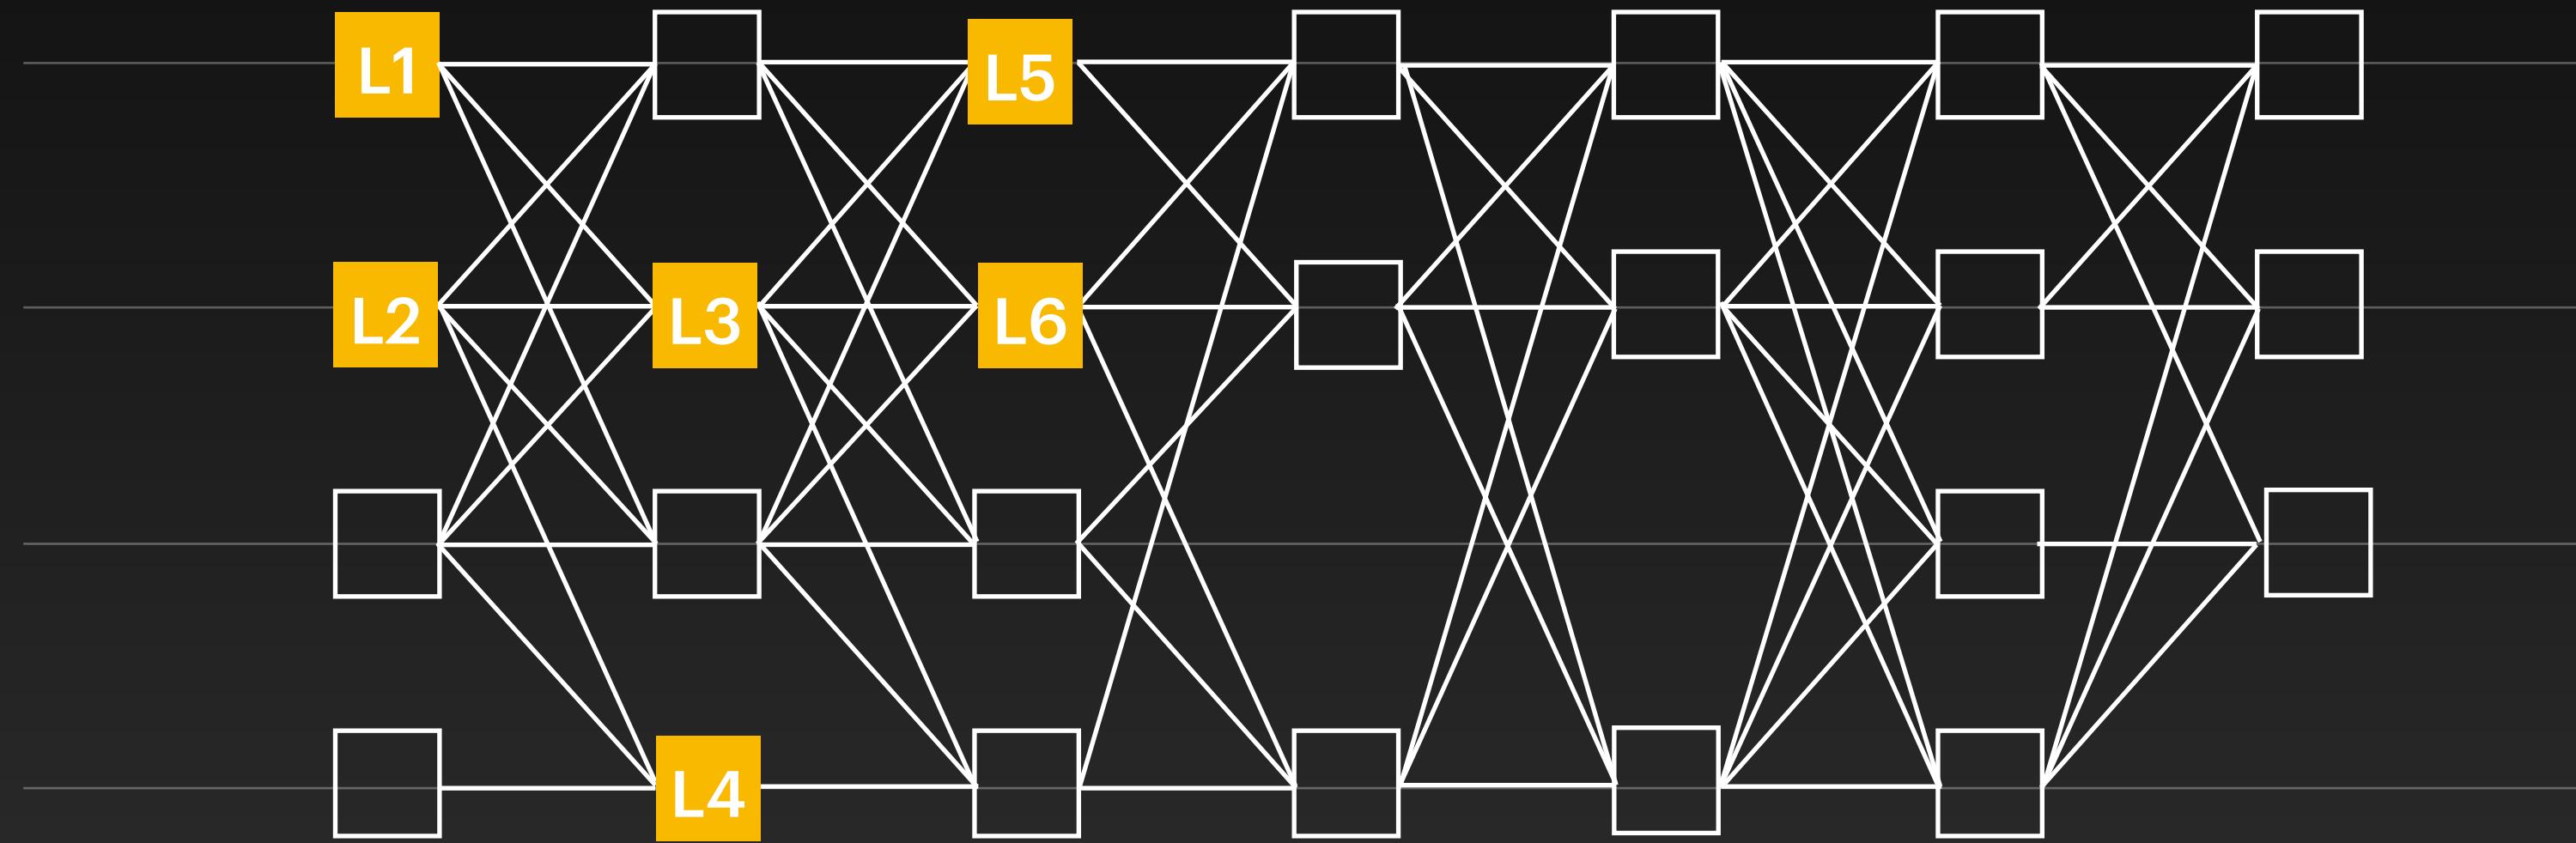
If an honest participant outputs a message, eventually every honest participant will output this message.

Integrity

Spurious messages cannot be attributed to honest participants.

End Goal

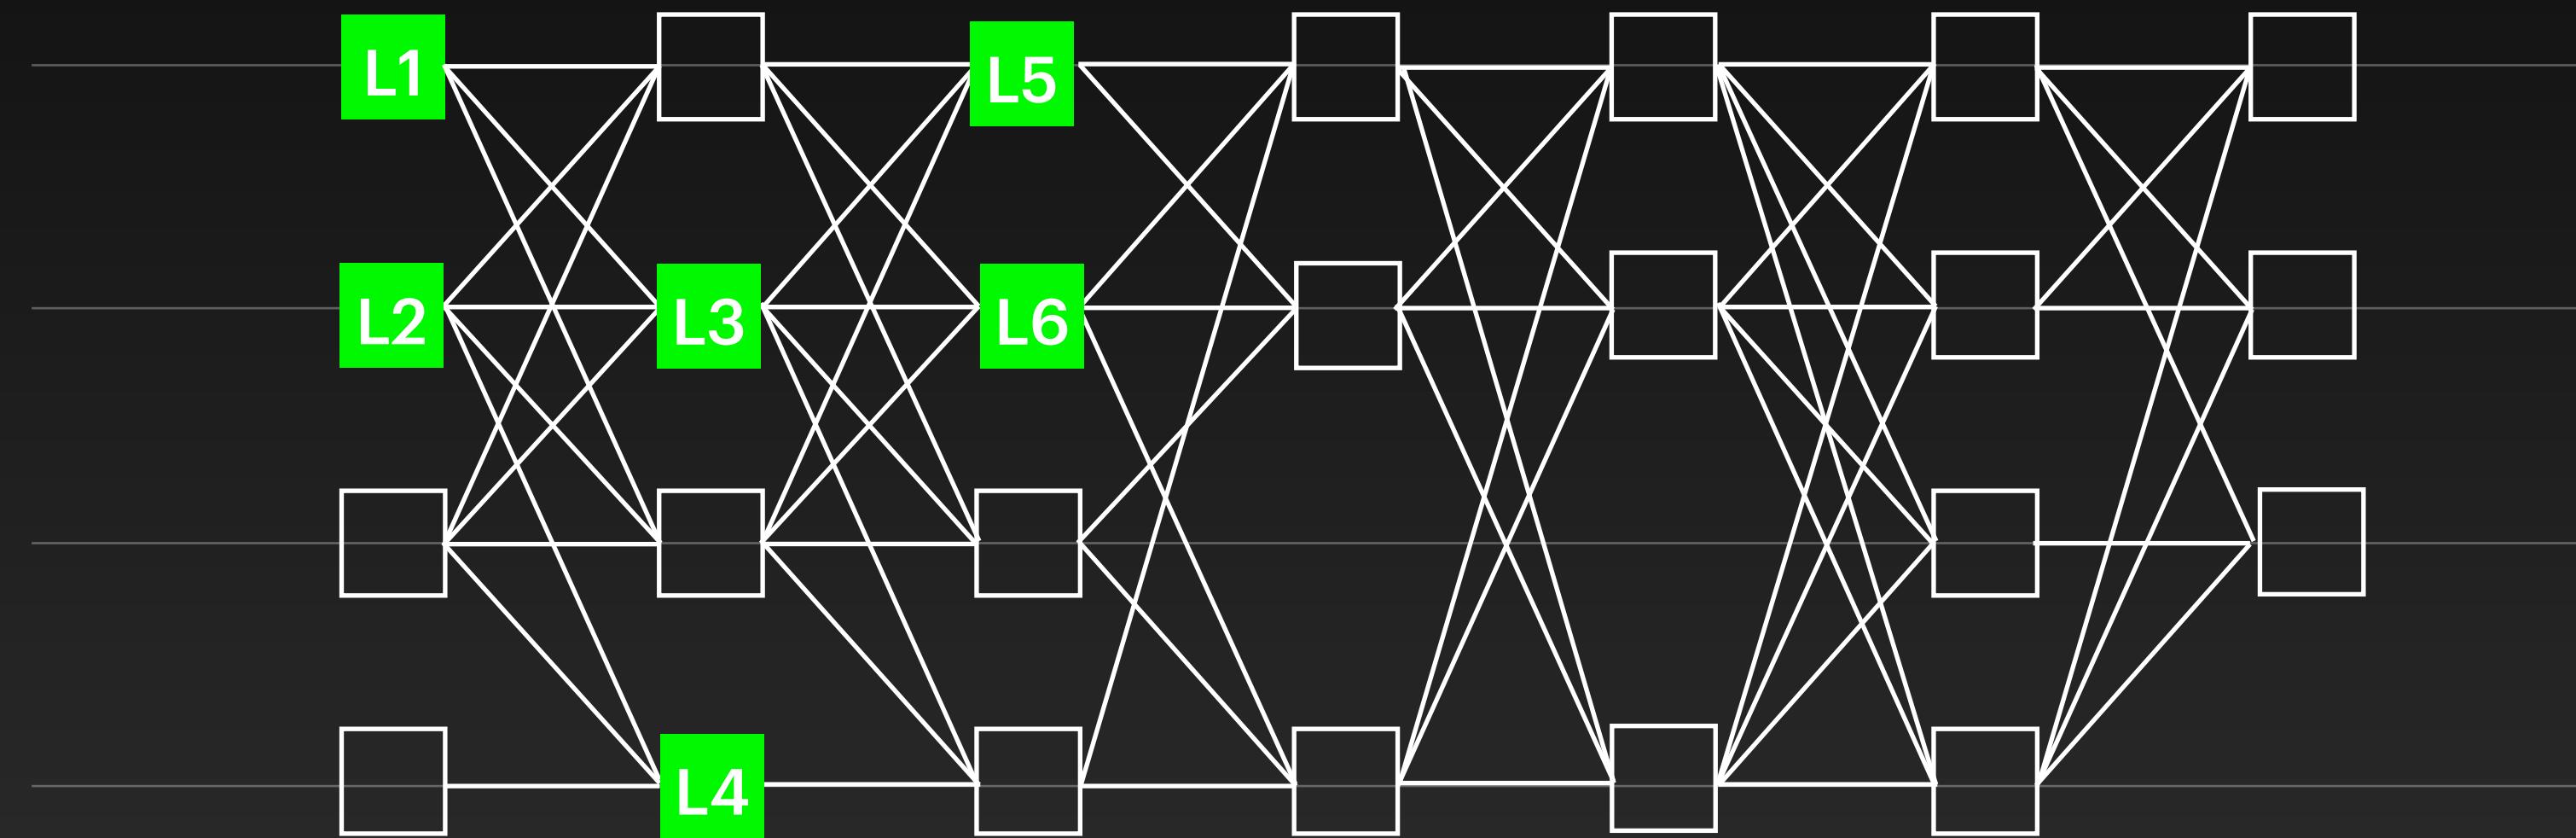
Ordering leaders



- We focus on ordering leaders: L1 L2 L3 L4 L5 L6

End Goal

Ordering leaders



- We focus on ordering leaders: L1 L2 L3 L4 L5 L6
- Linearising the sub-DAG is simple

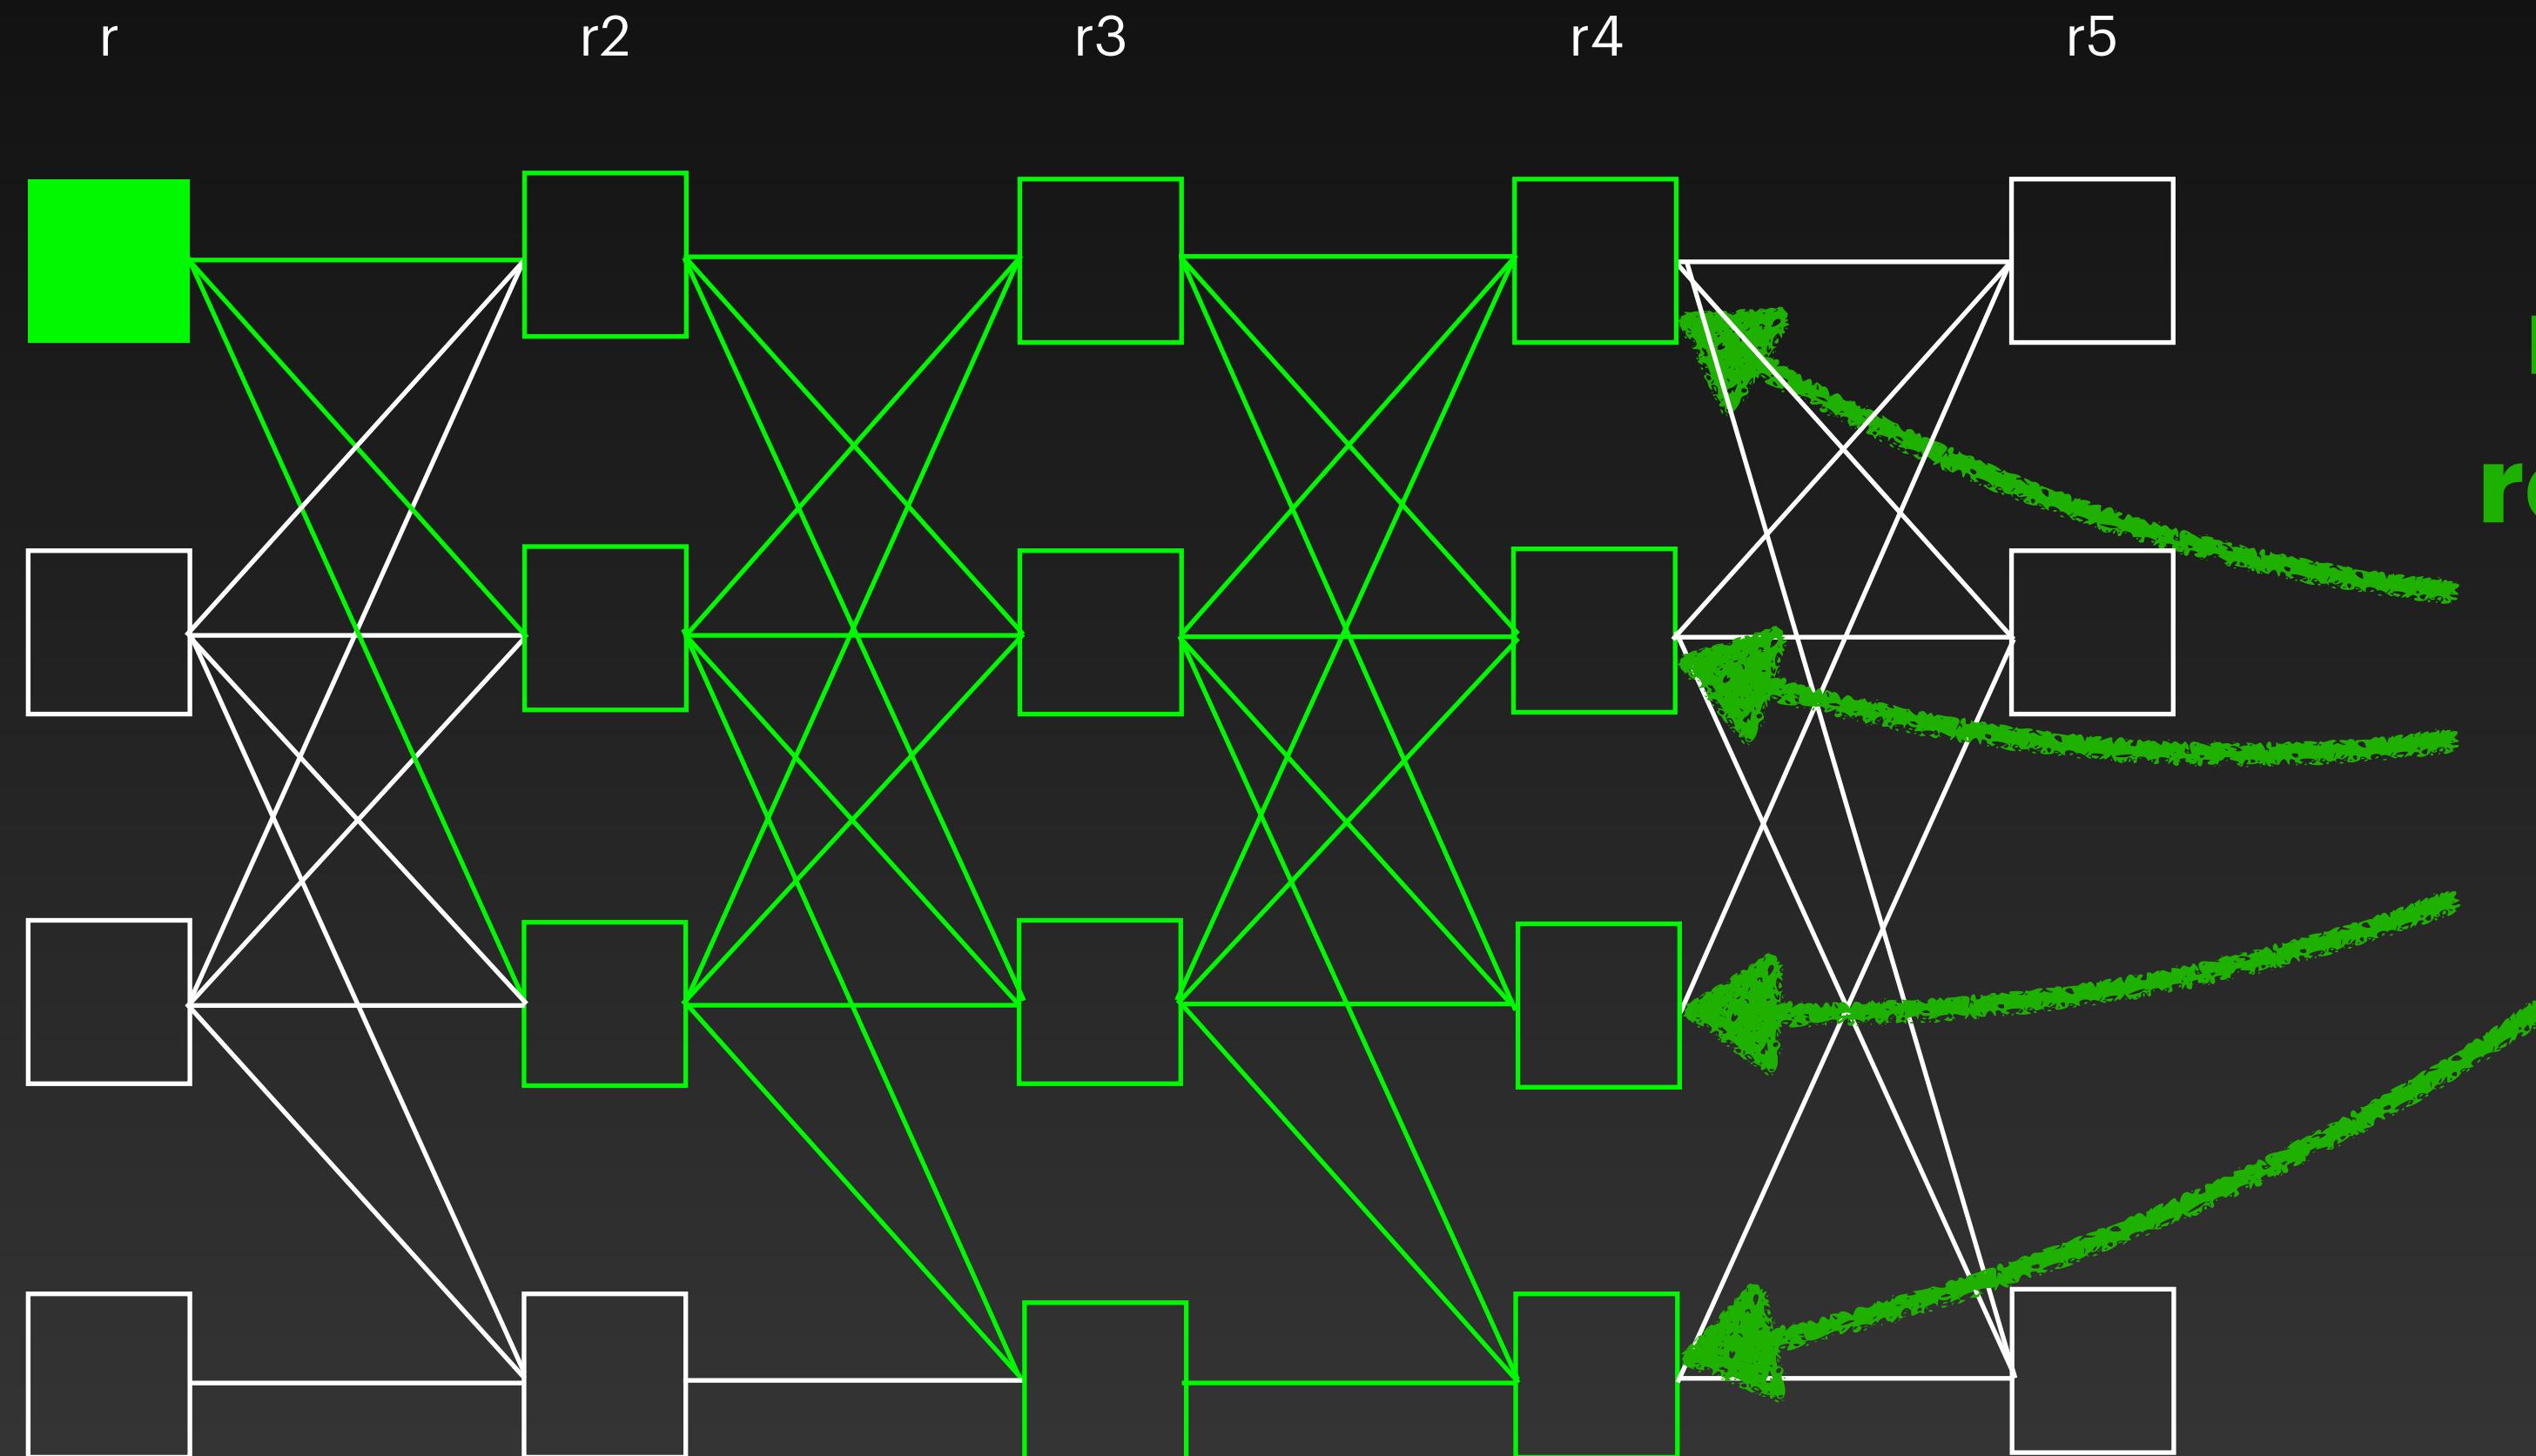
Components for Consensus

How to commit blocks?

Blocks can do the following:

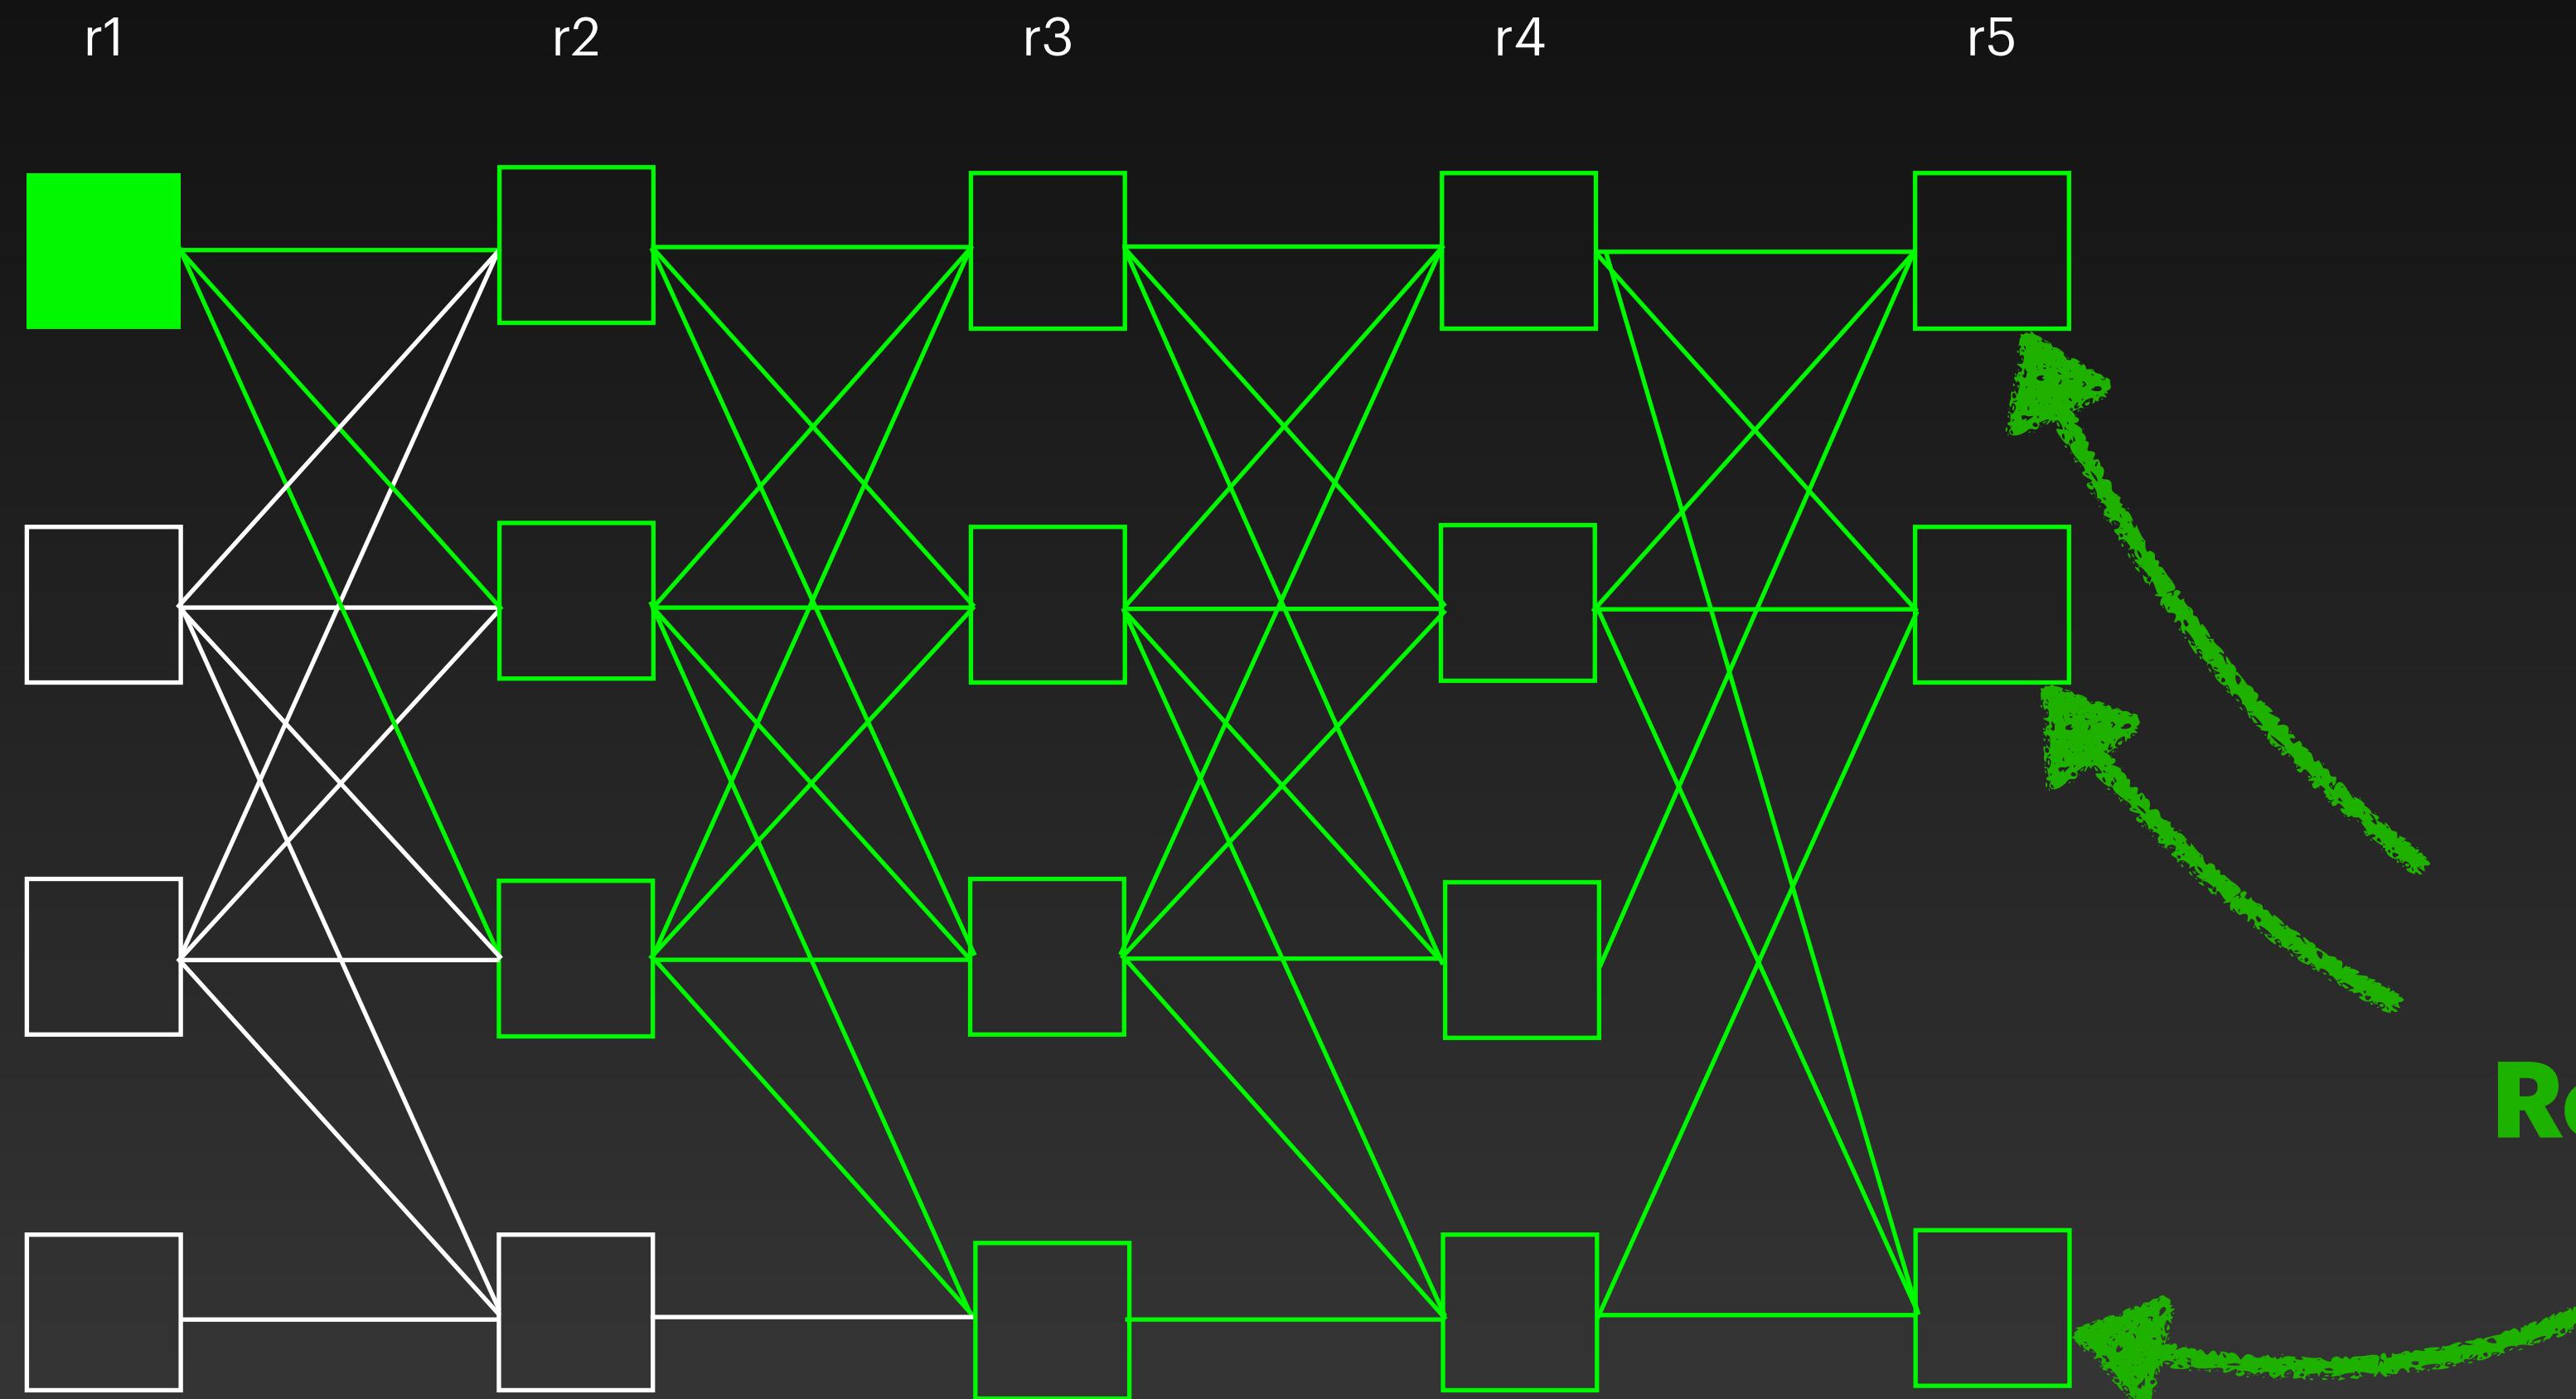
Vote	Certificate	Blame
<u>Path from $r+3$ to r</u>	<u>Reference $2f+1$ votes</u> for a block	<u>No path in $r+4$ to r</u>

Voting for a block



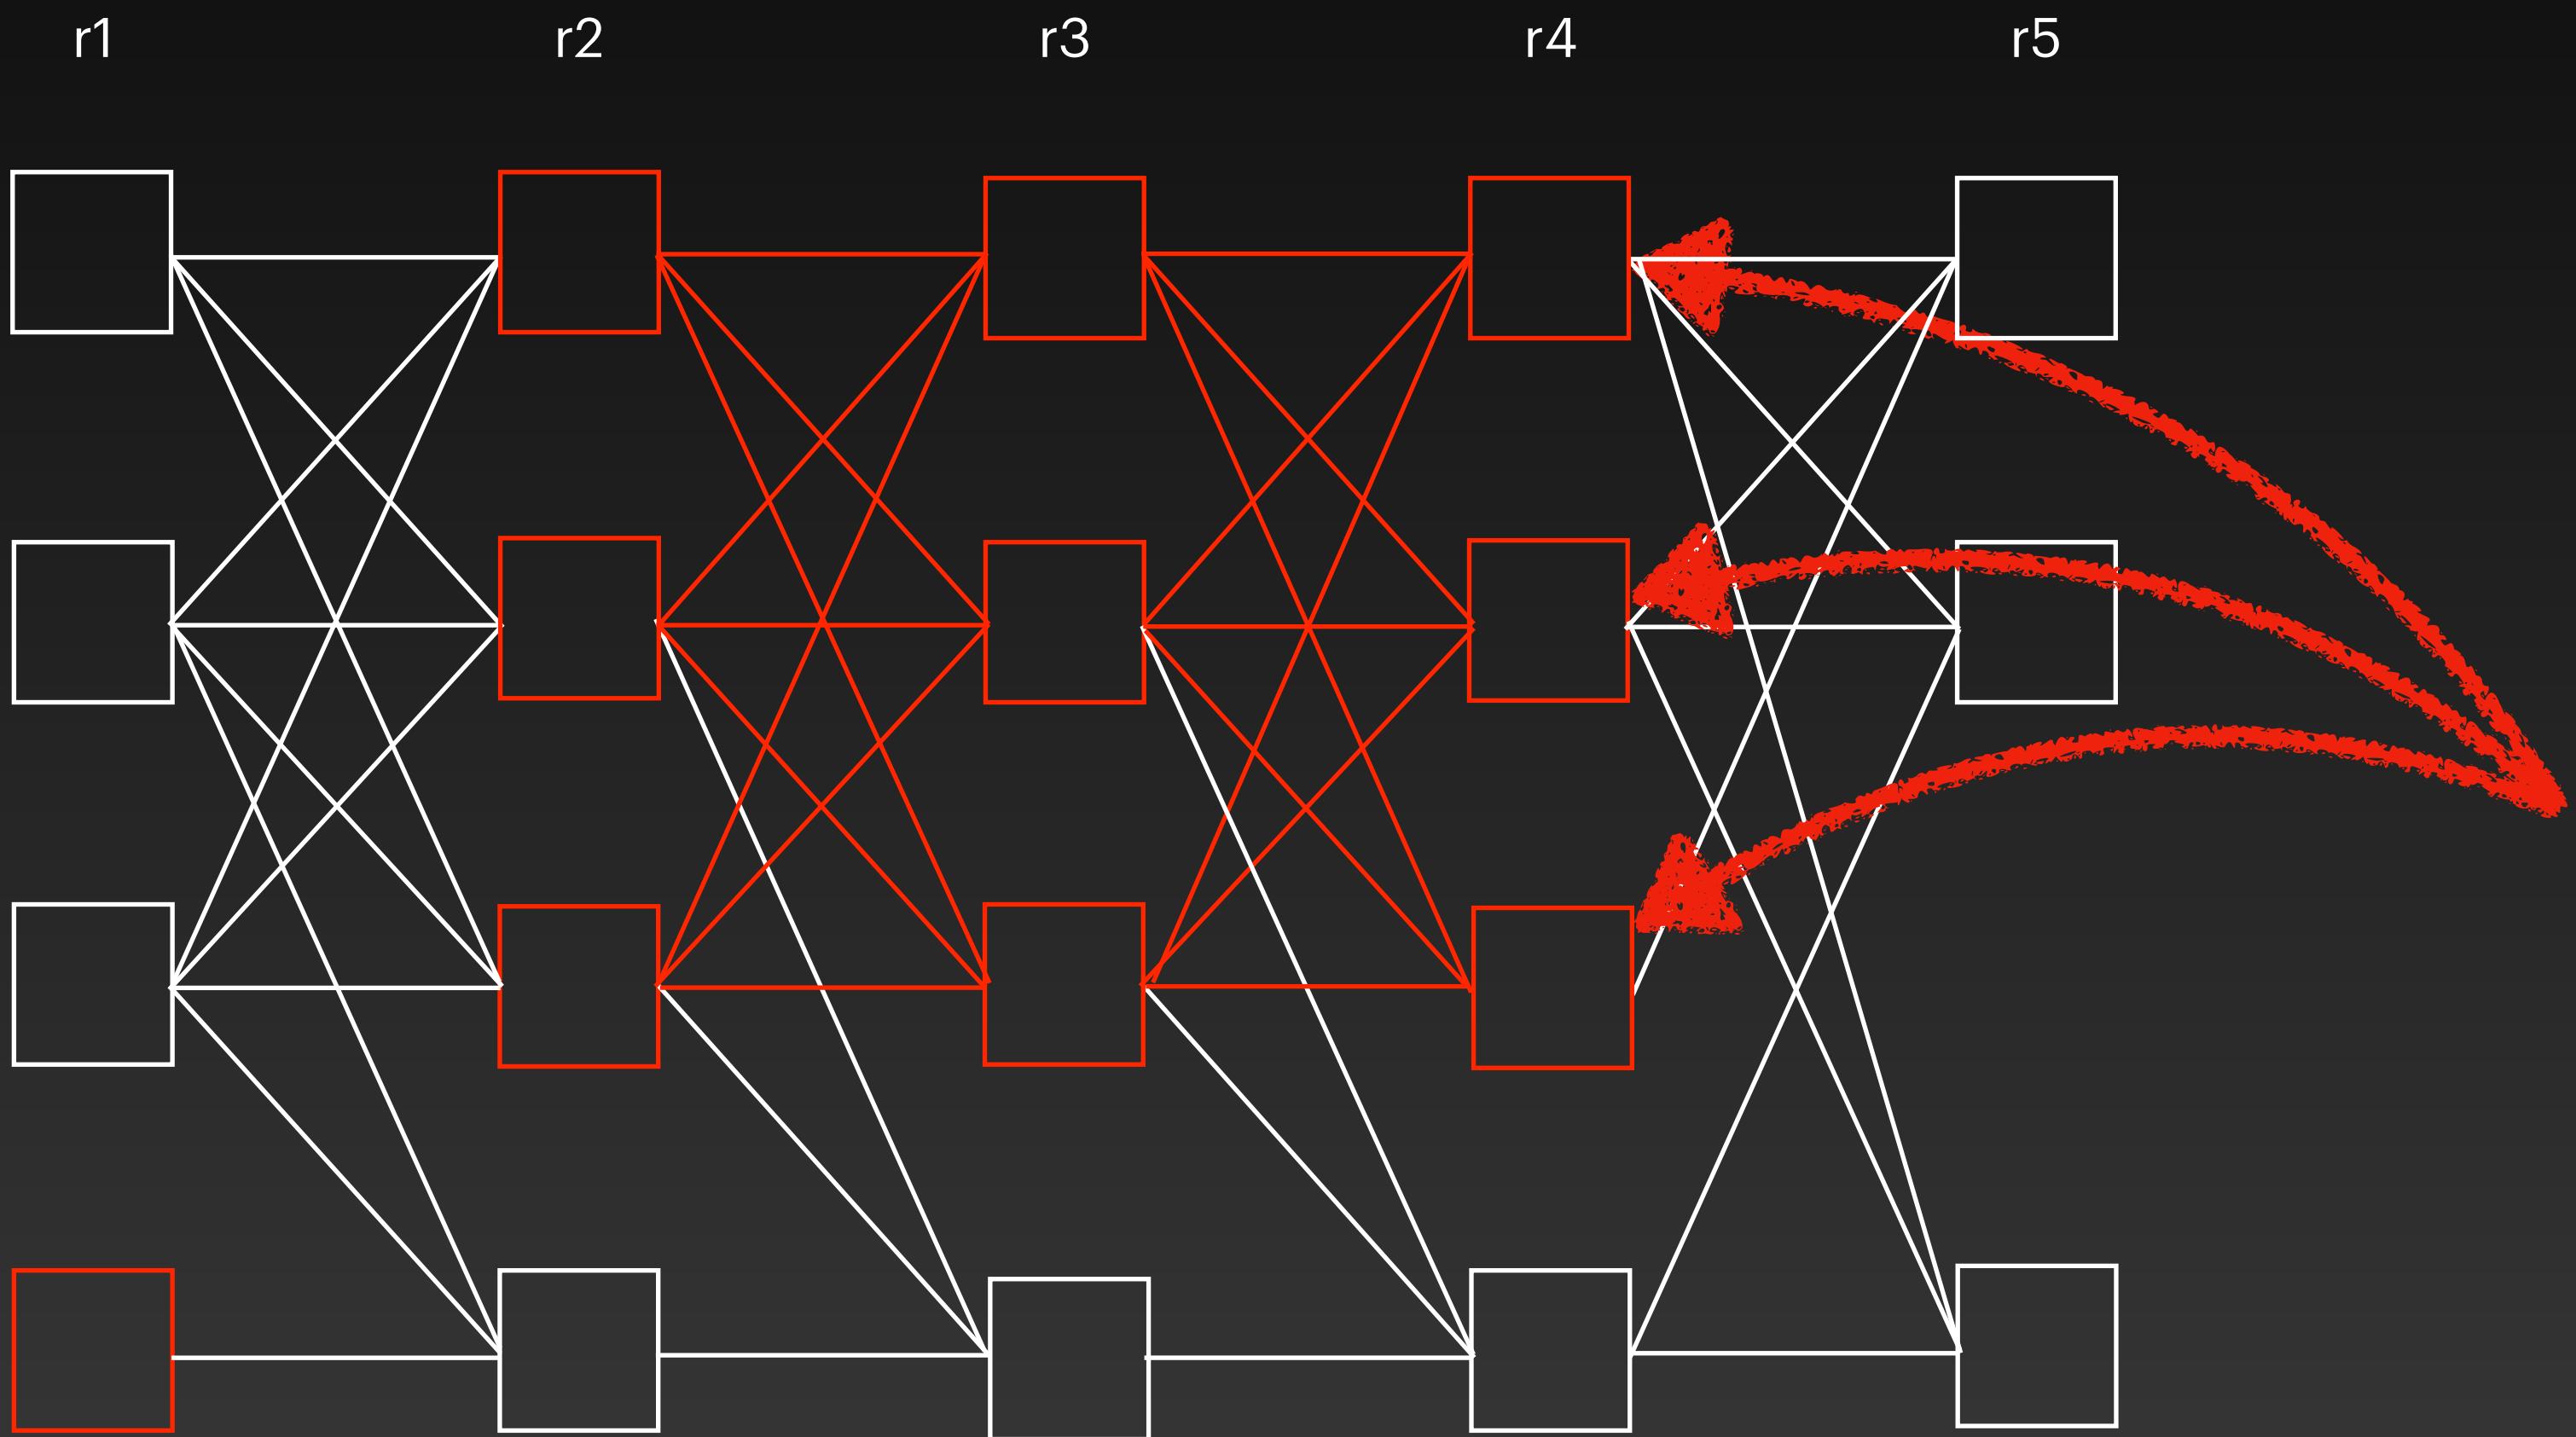
**Vote: A path from
round $r+3$ block to
round r block *make
consistent with
round notation in
dag**

Certificate for a block



Certificate:
Reference 2f+1
votes

Blames for a block

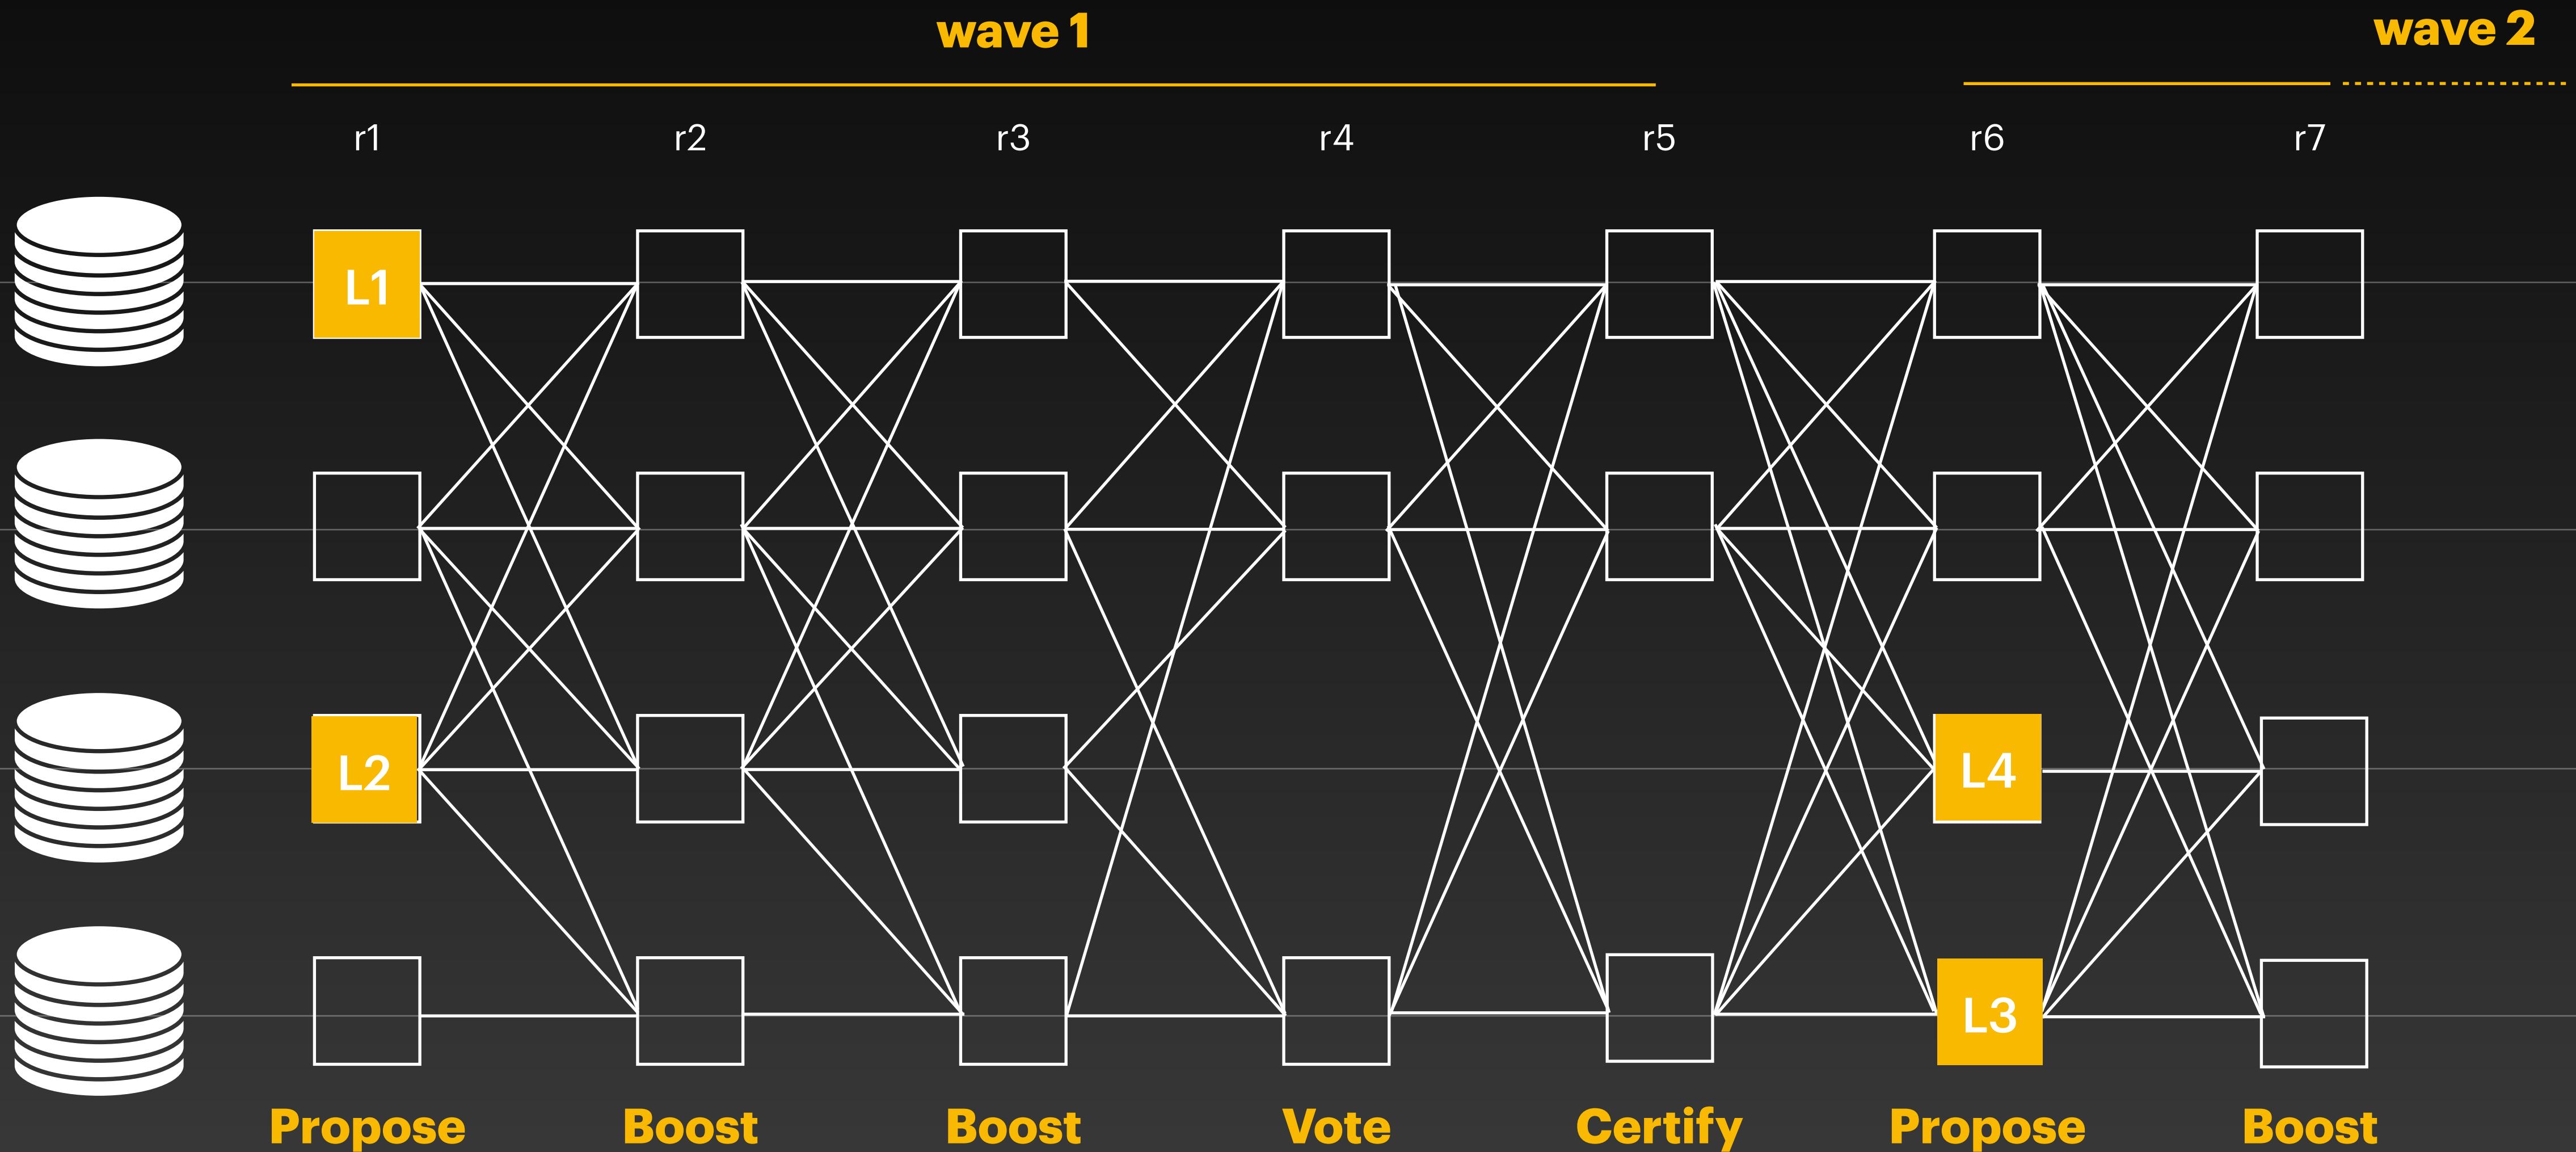


**Blame: No path
from $r+3$ block
to r block**

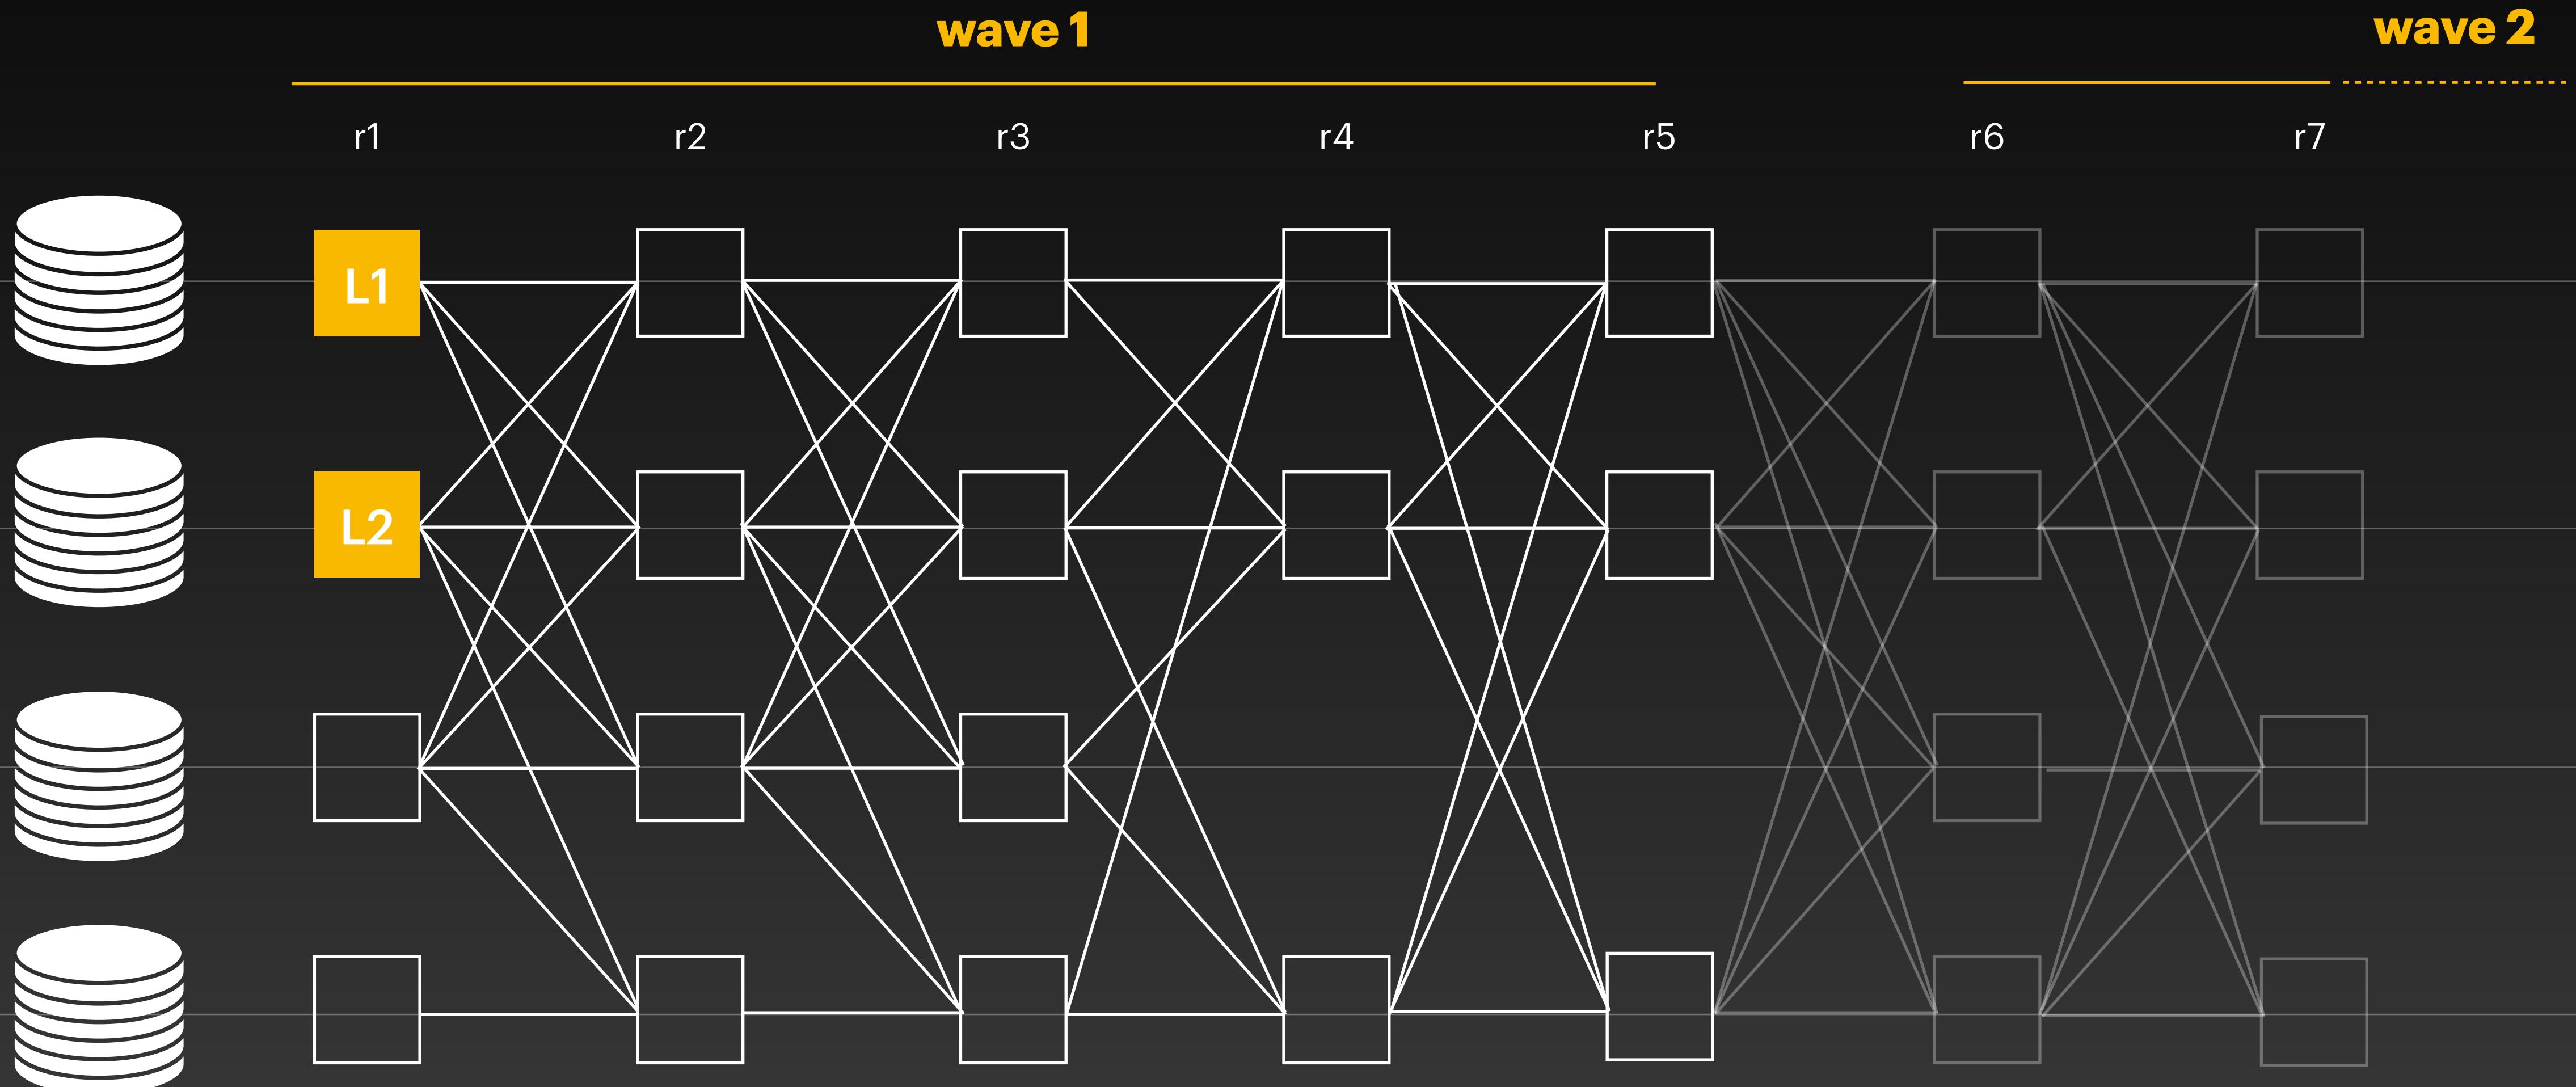
Mahi-Mahi Commit rule

How to reach consensus

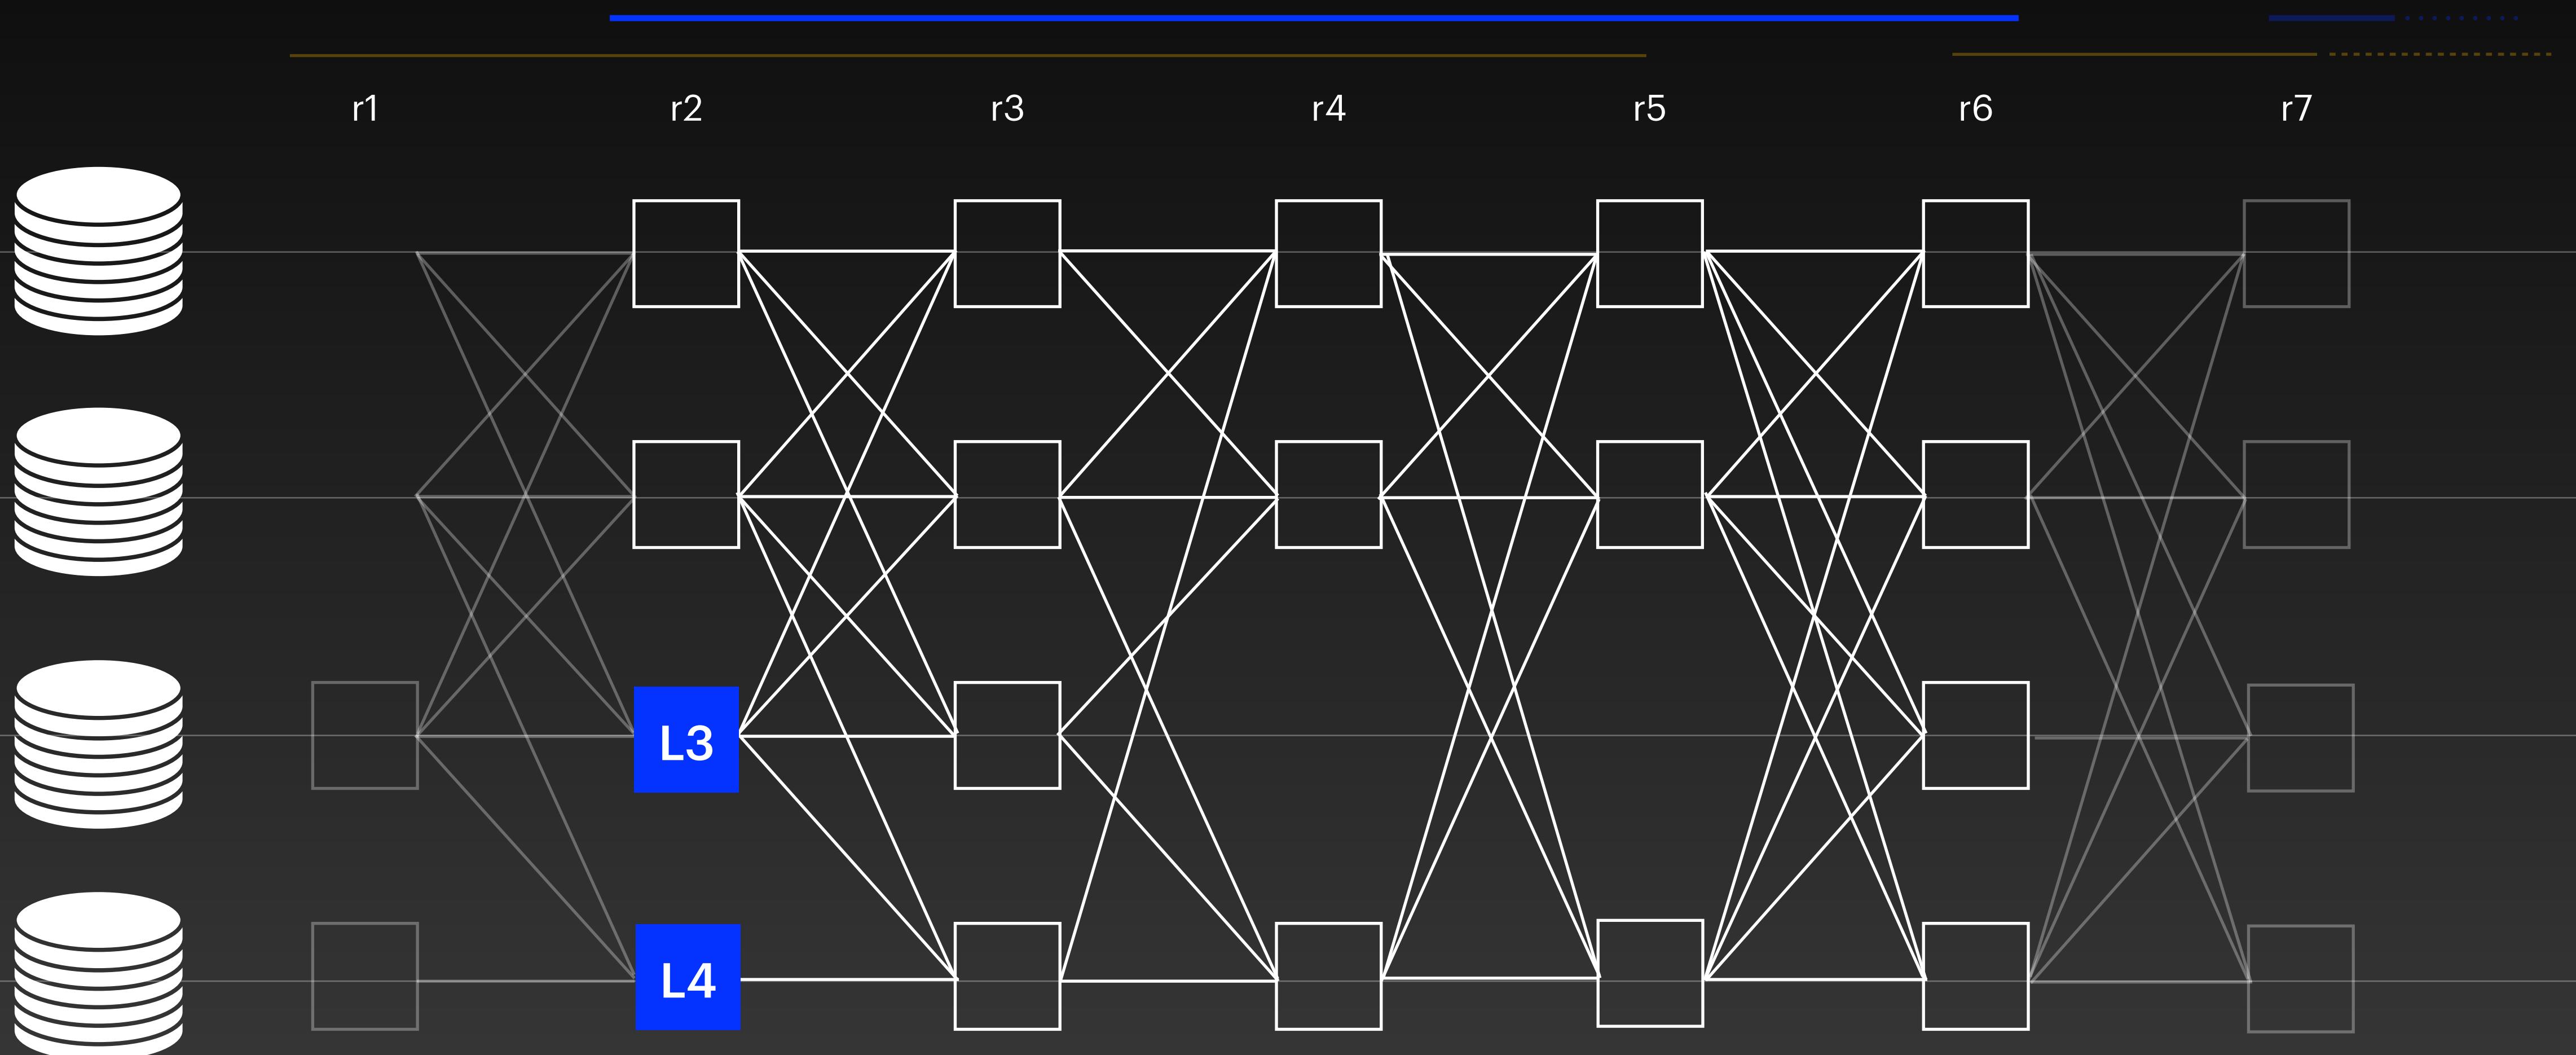
Mahi-Mahi: Outline of a wave



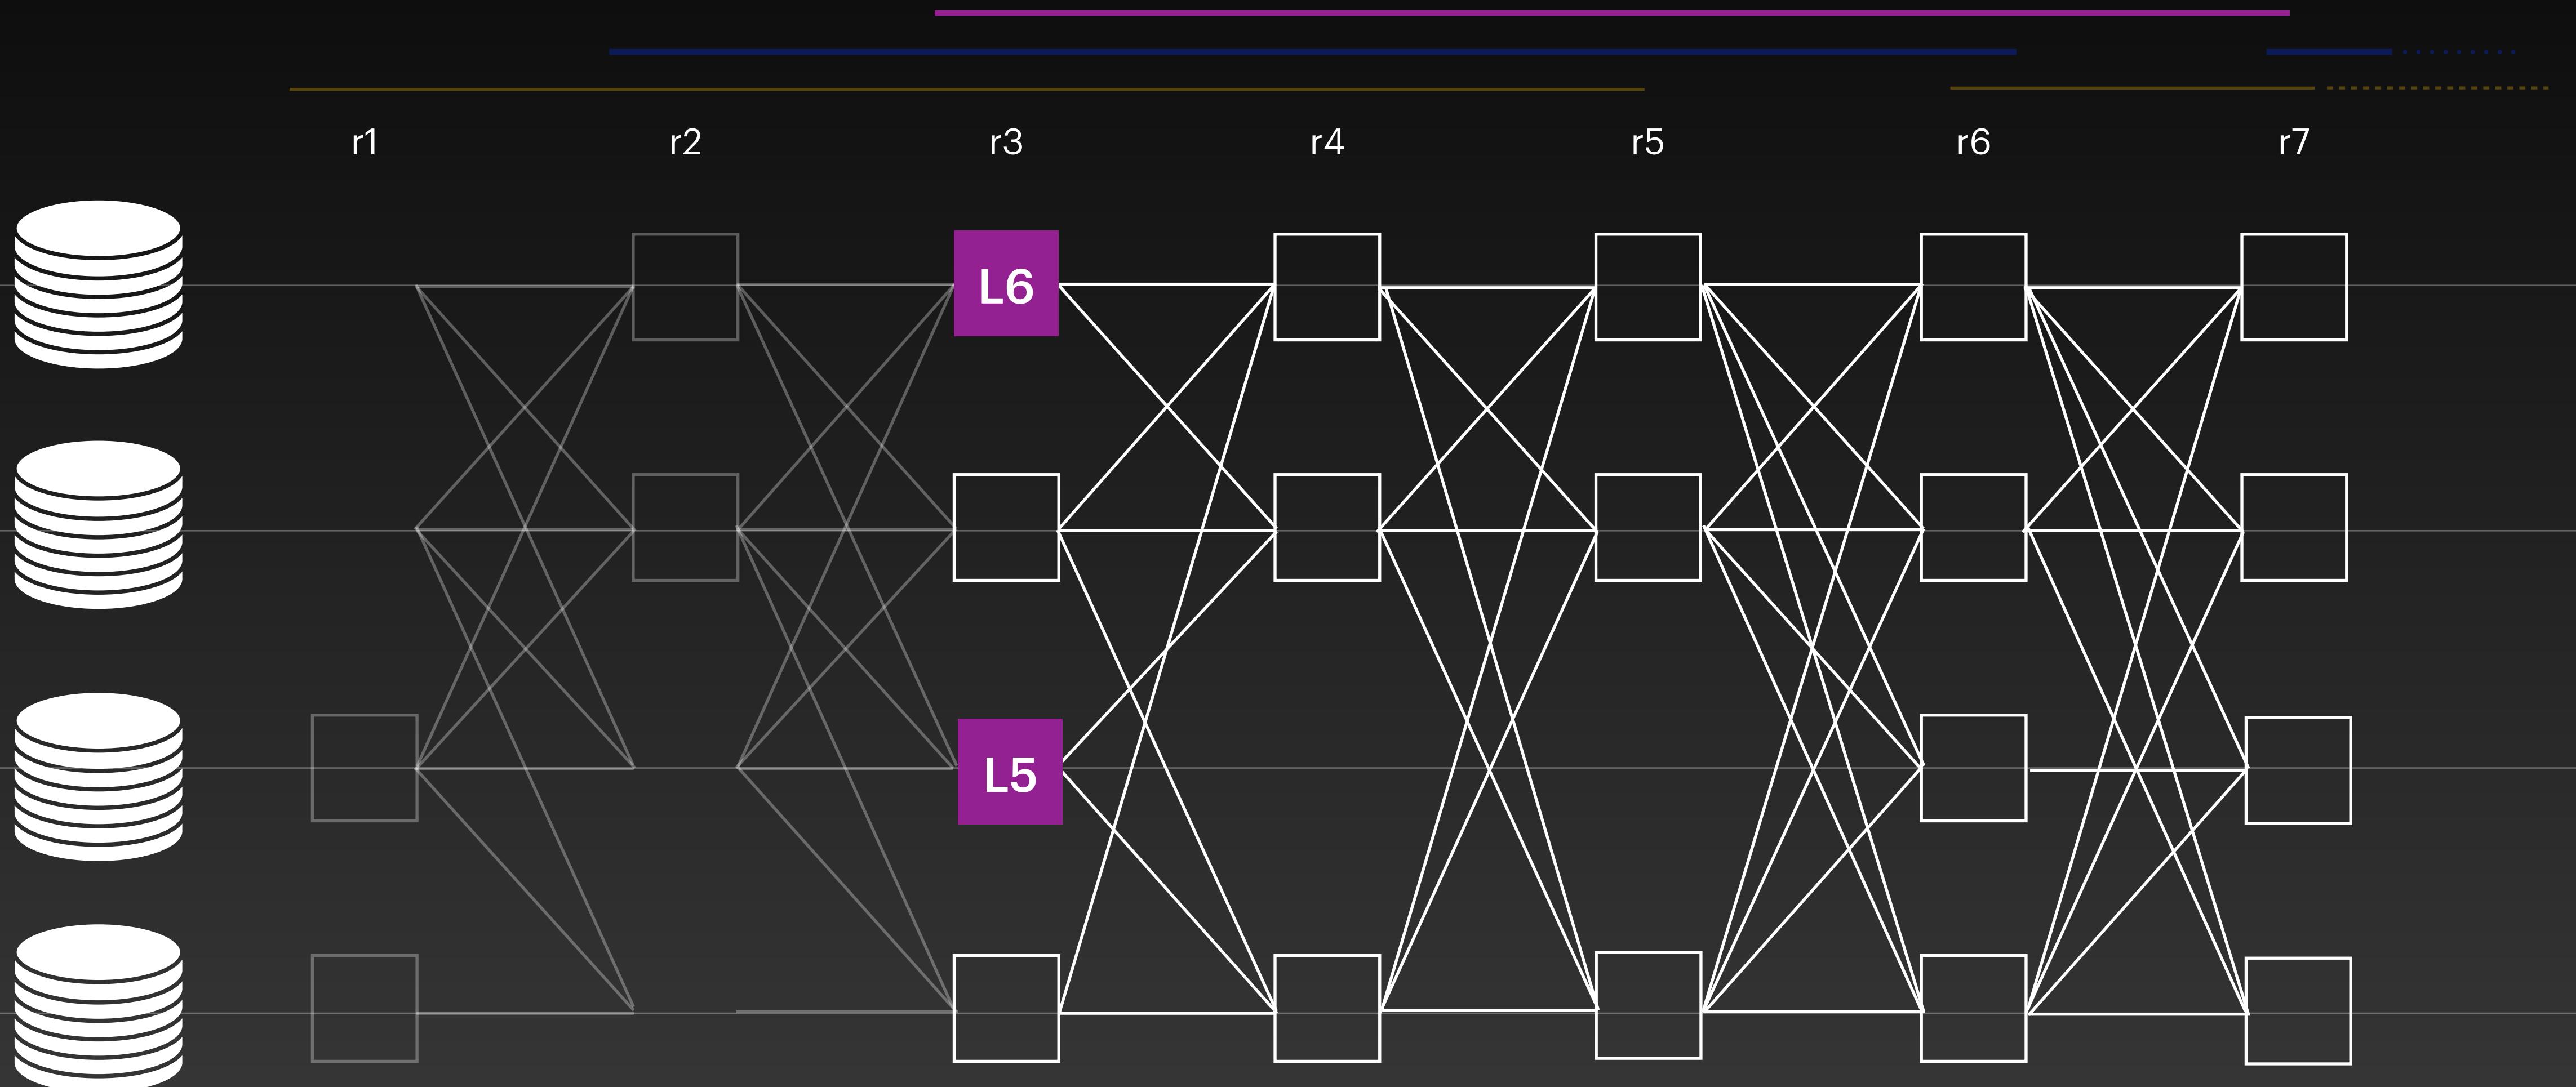
Pipelining in Mahi-Mahi



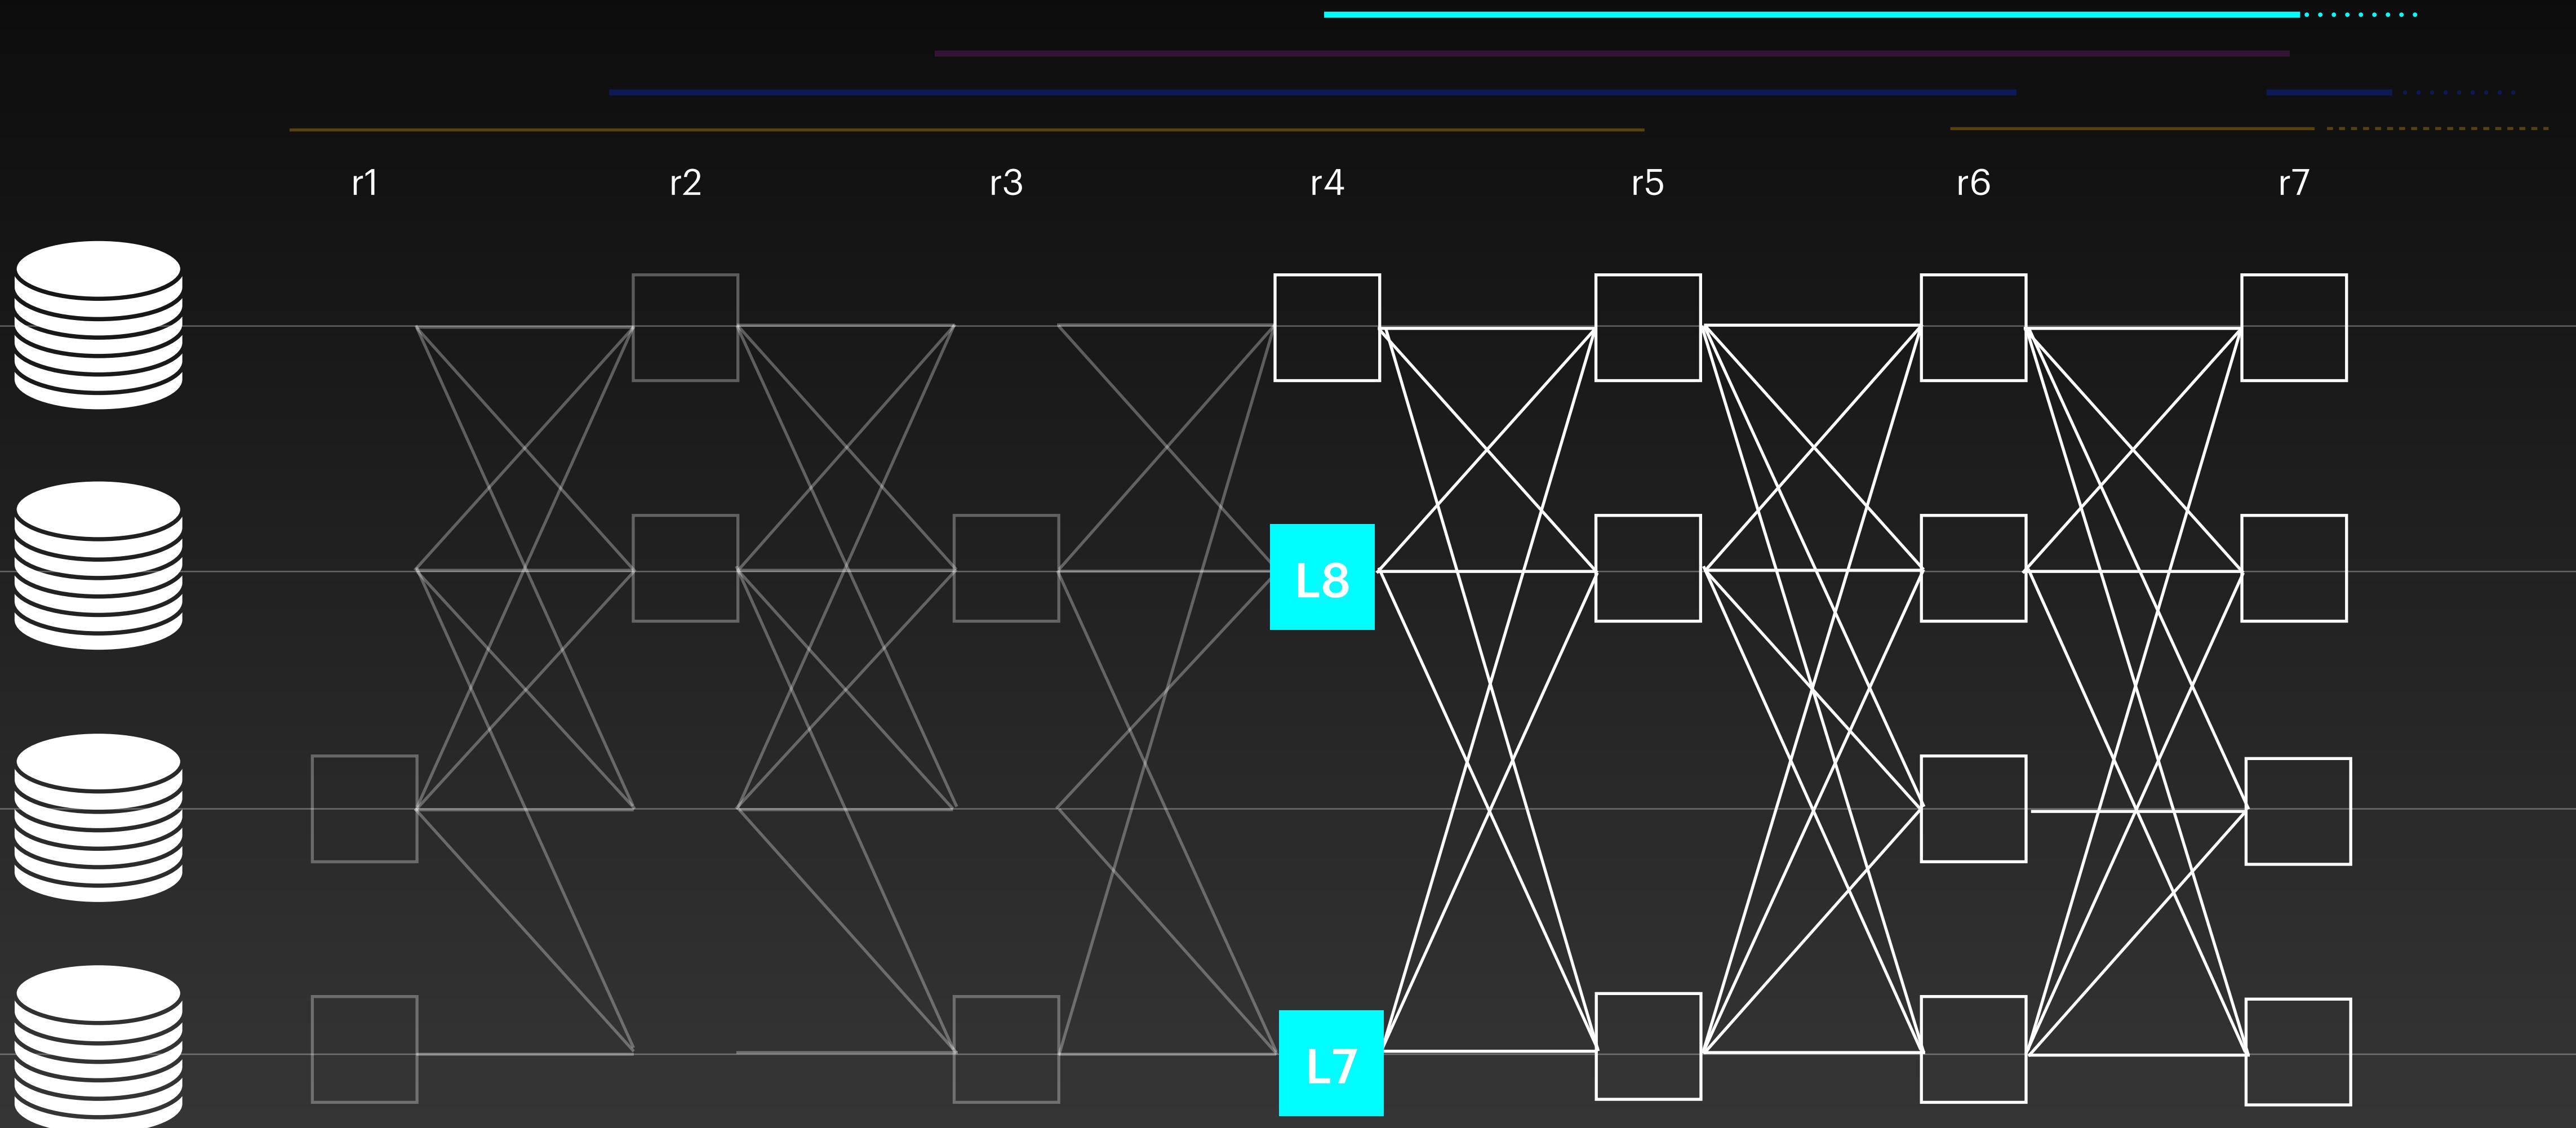
Pipelining in Mahi-Mahi



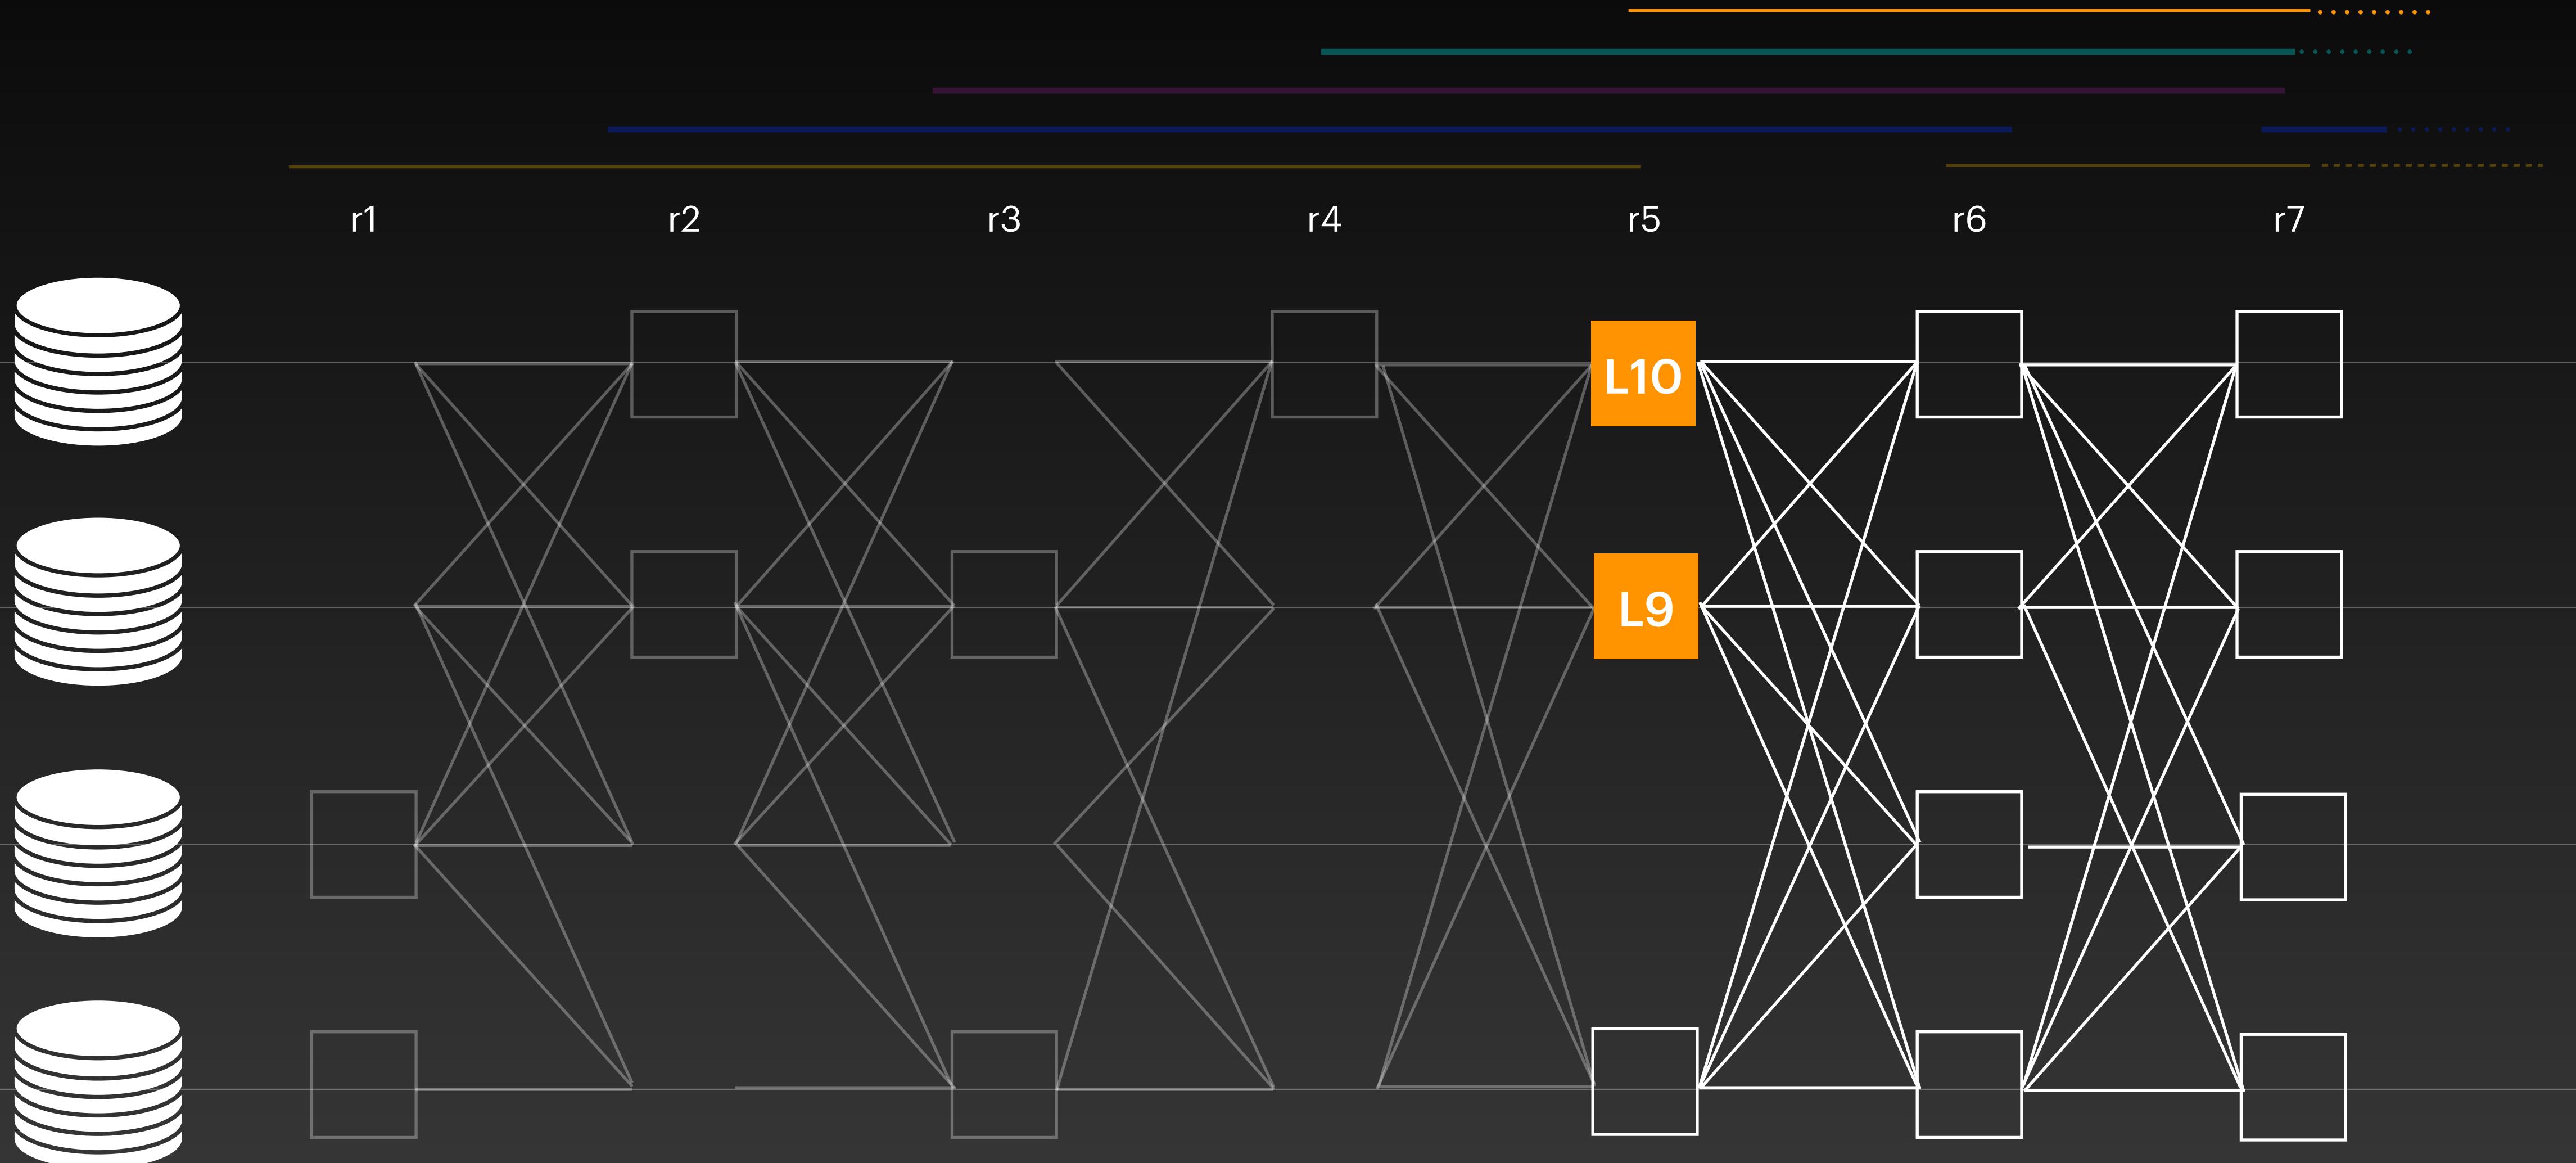
Pipelining in Mahi-Mahi



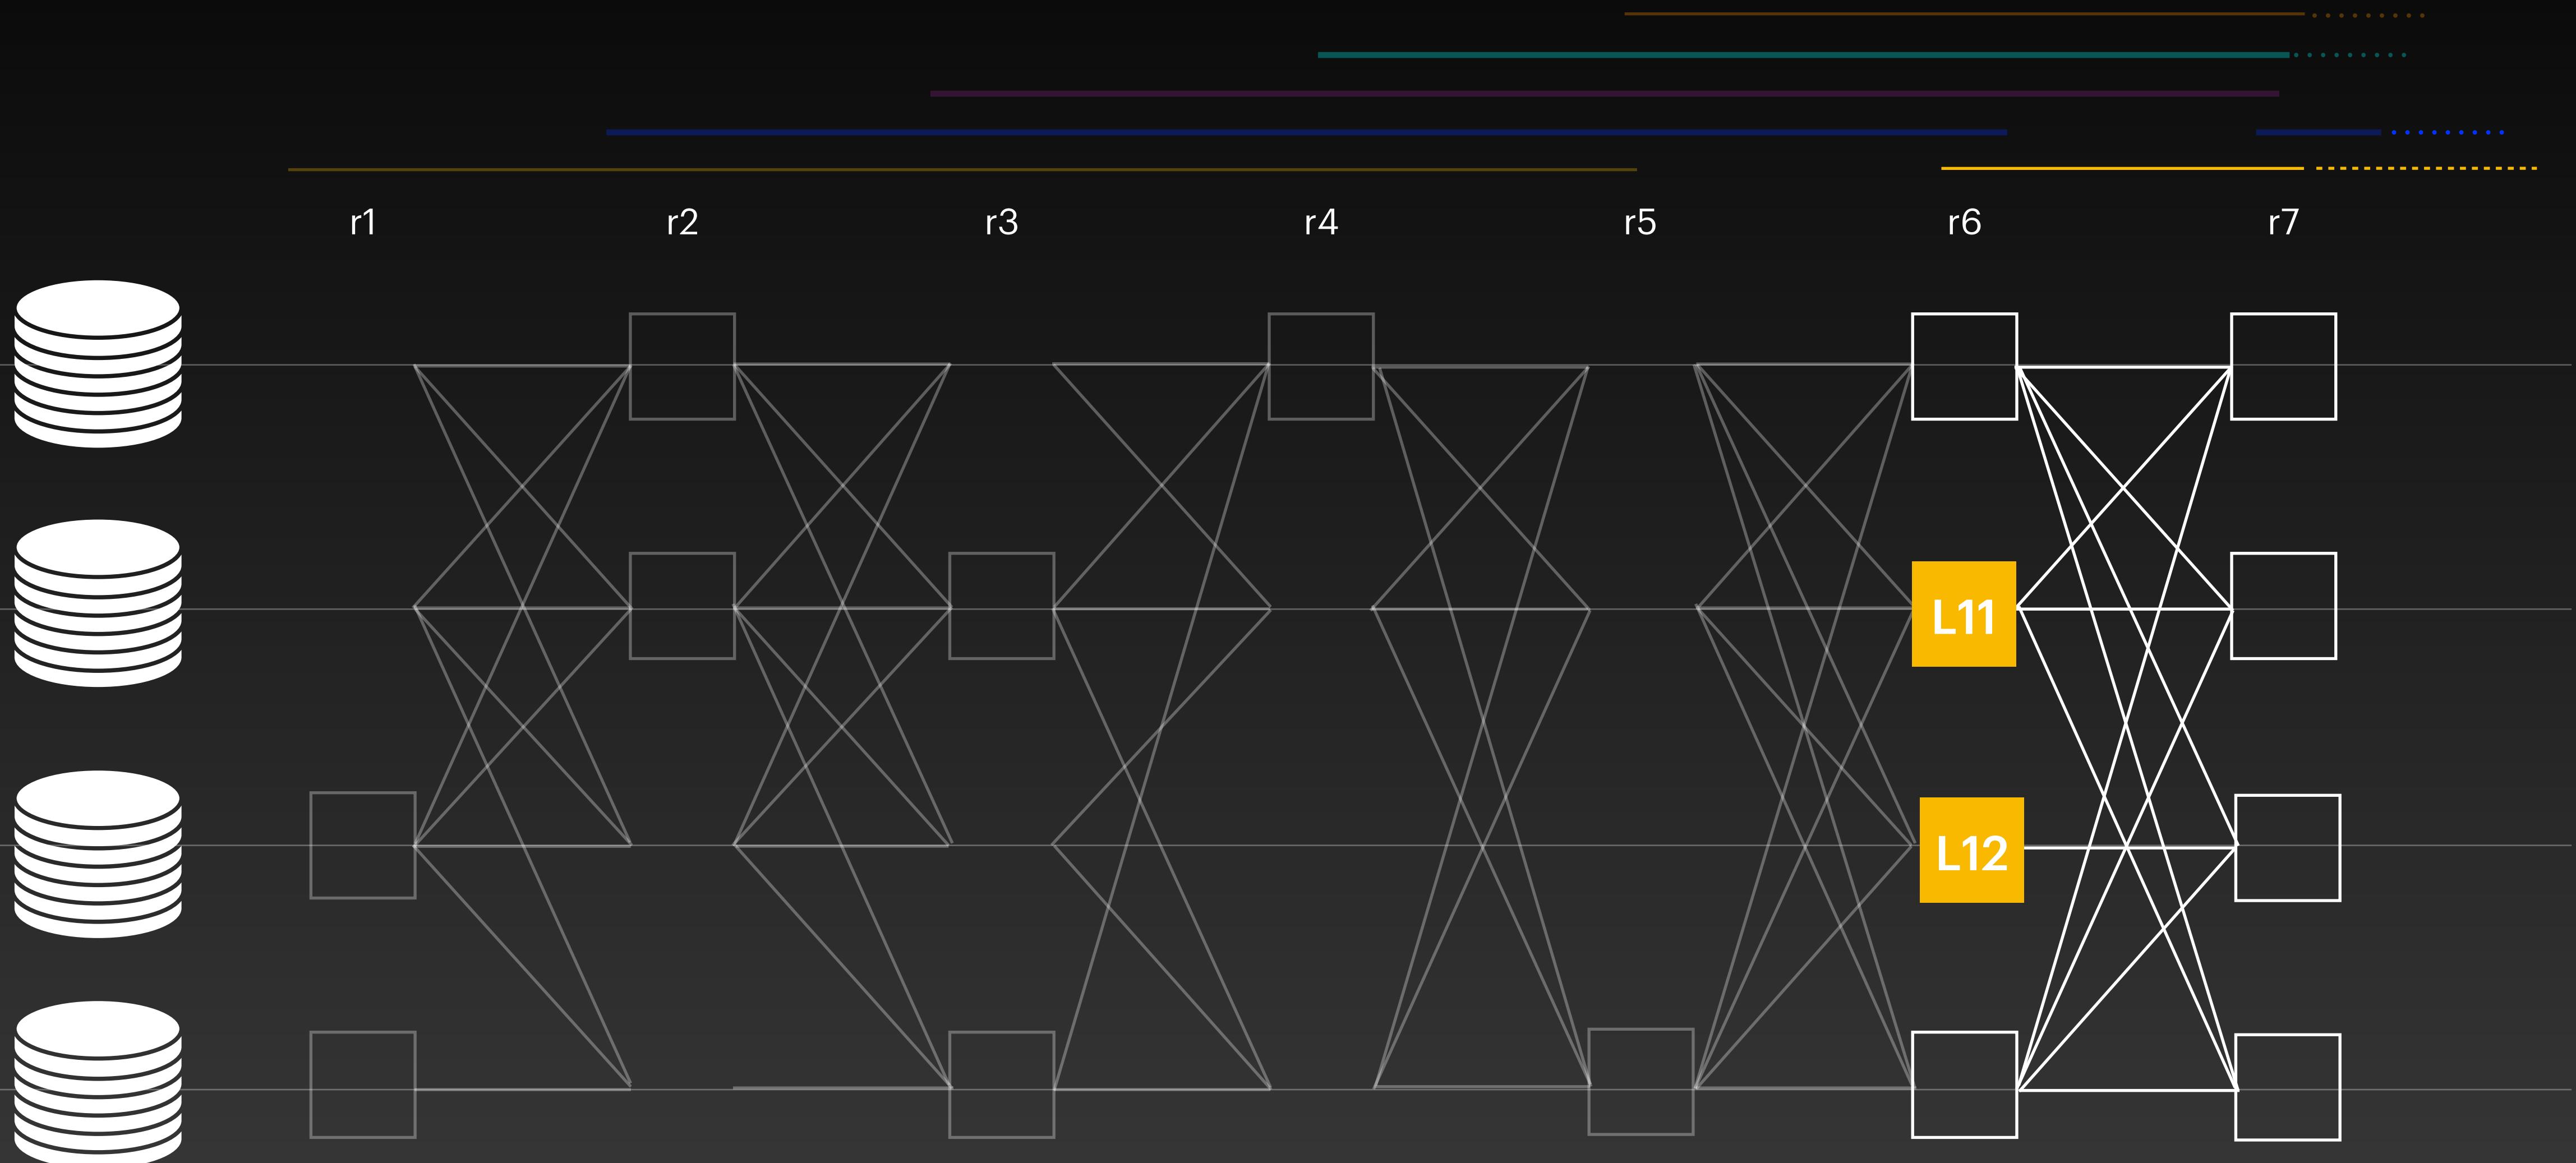
Pipelining in Mahi-Mahi



Pipelining in Mahi-Mahi



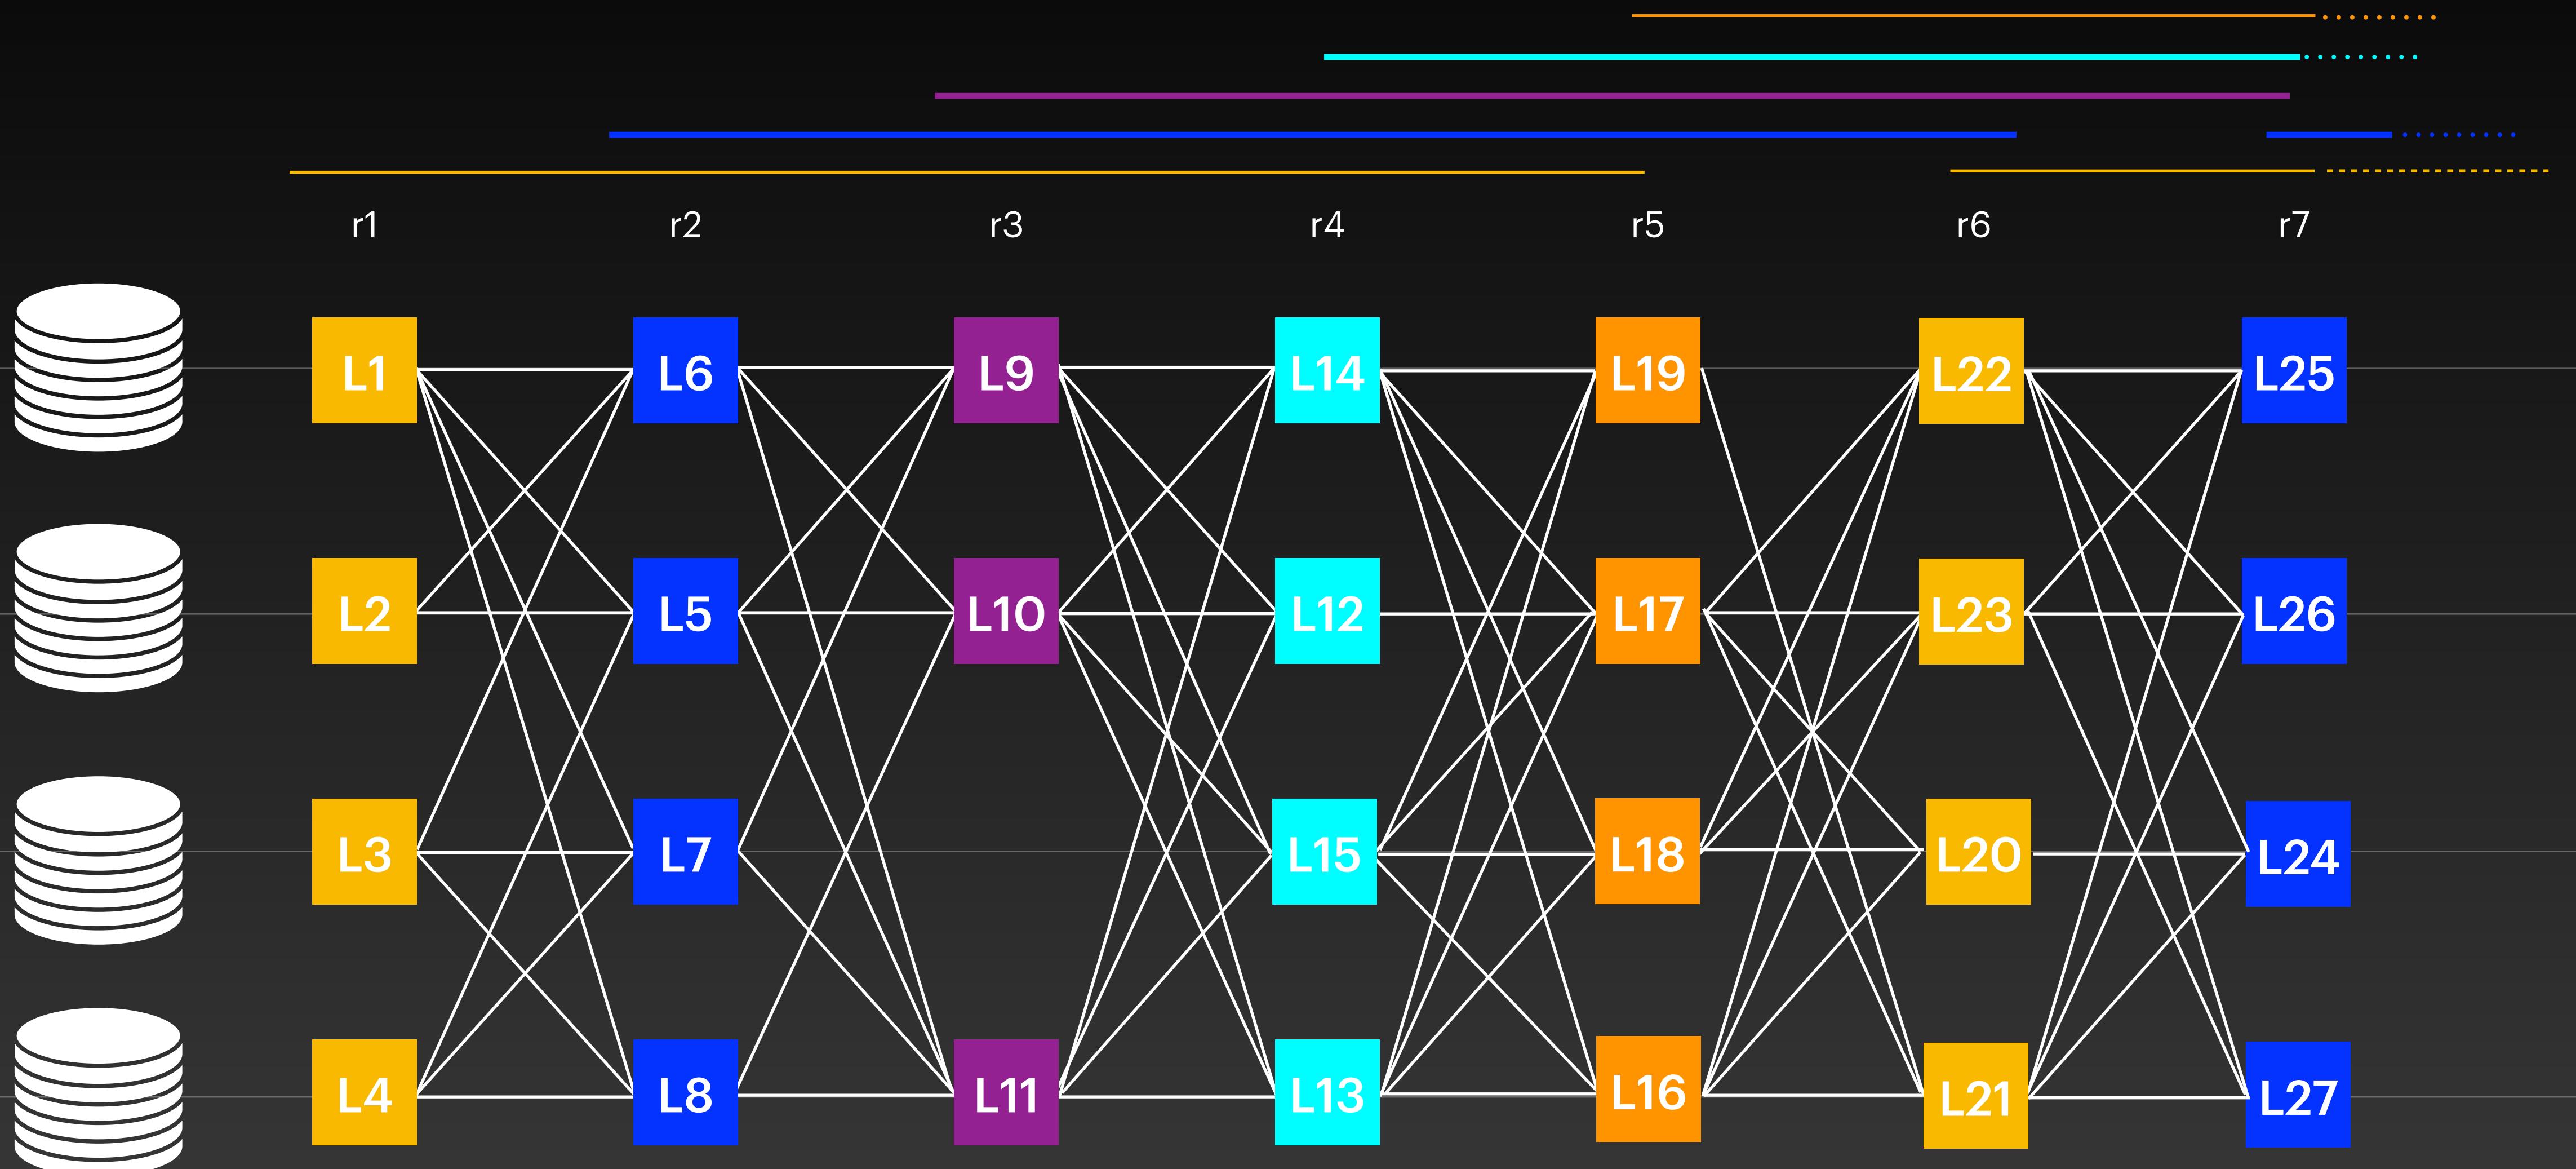
Pipelining in Mahi-Mahi



Pipelining in Mahi-Mahi

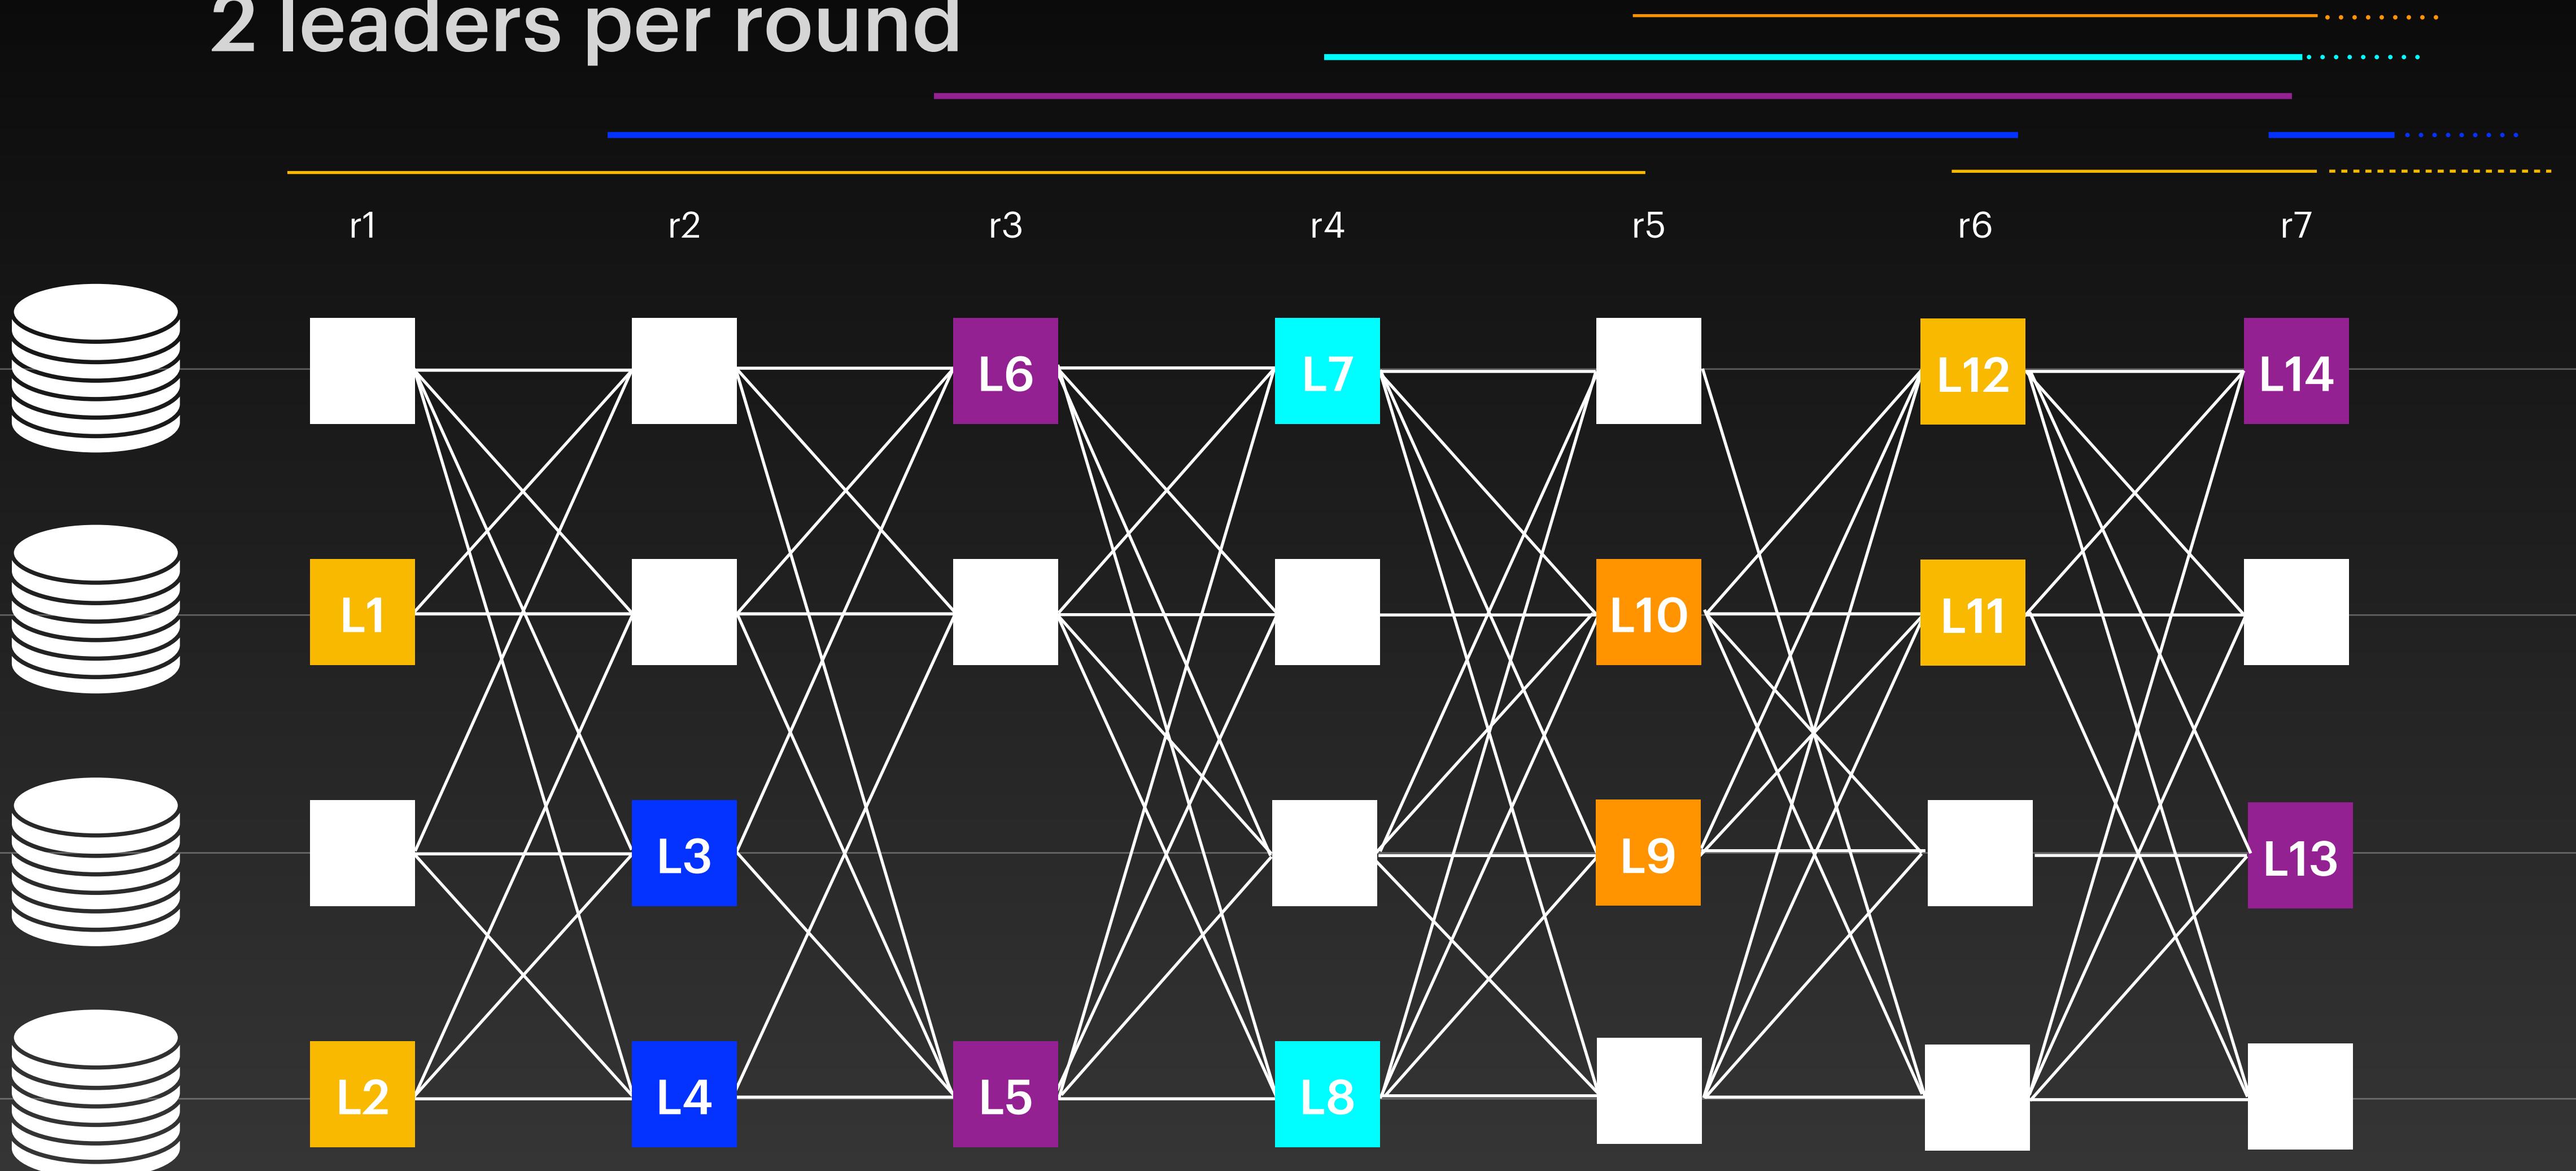


Theoretically, every validator can be a leader



Practical Implementation of Mahi-Mahi

2 leaders per round



Direct Decision Rule

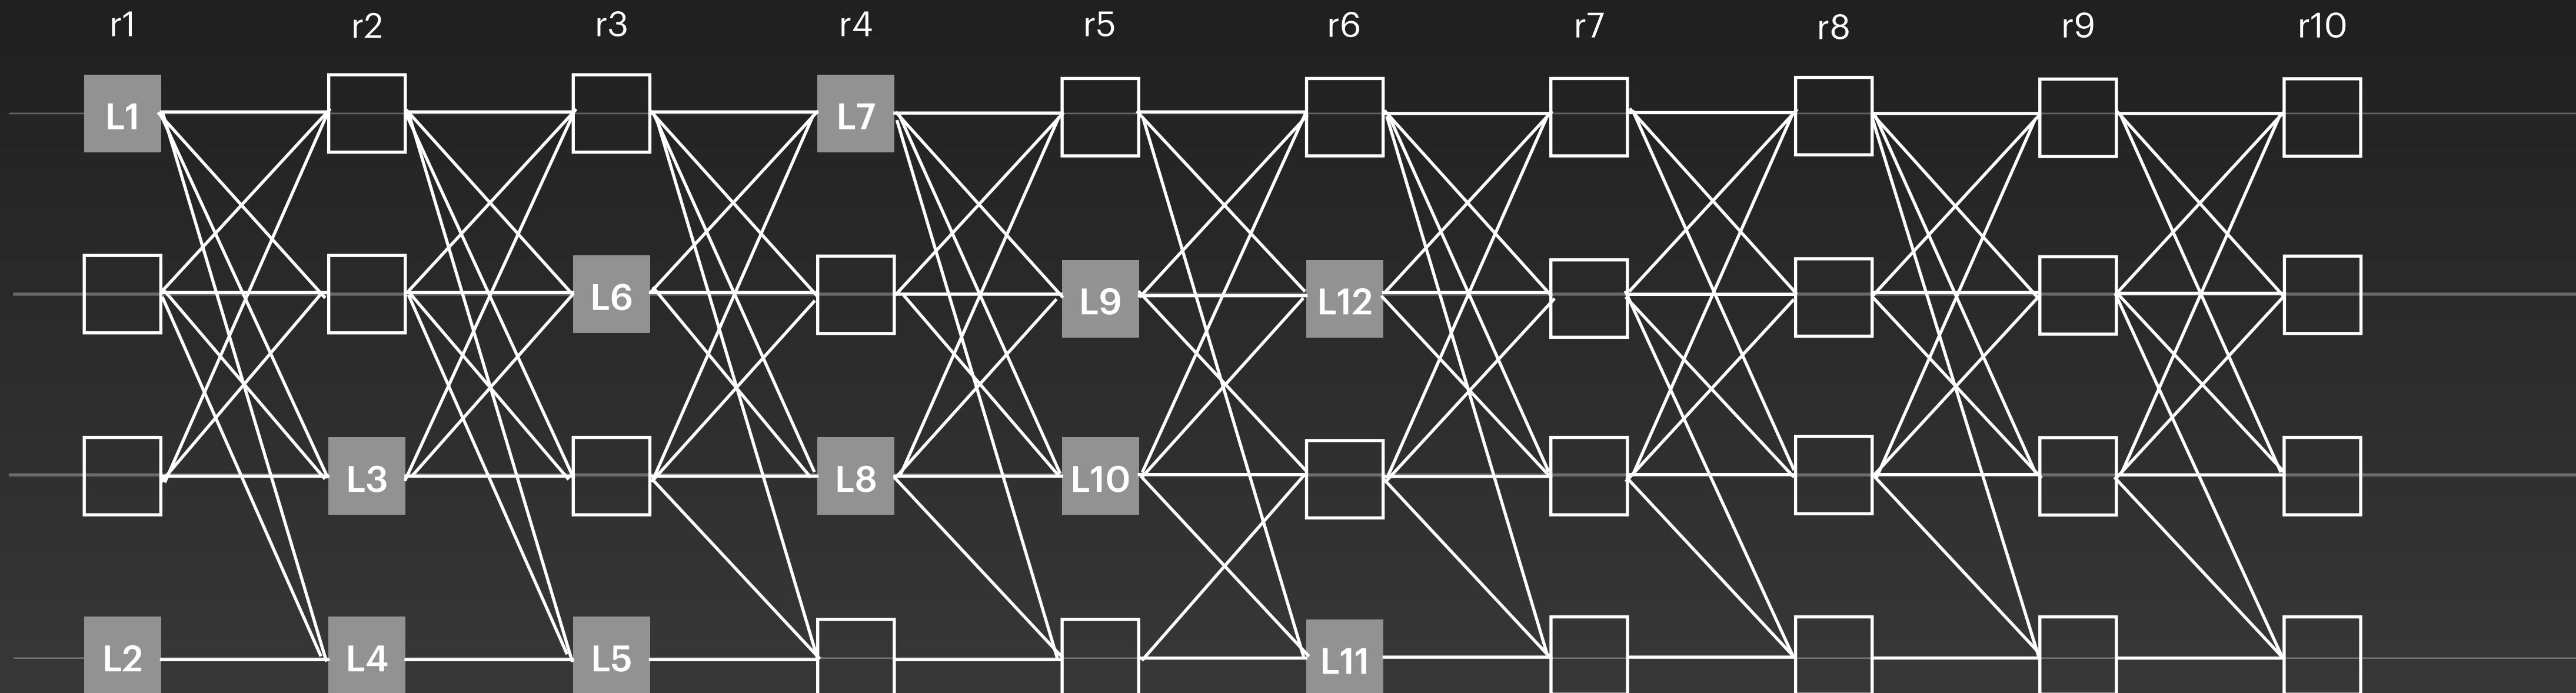
On each leader starting from highest round:

- **Skip** if $2f+1$ blames
- **Commit** if $2f+1$ certificates
- **Undecided** otherwise

Direct Decision Rule

On each leader starting from highest round:

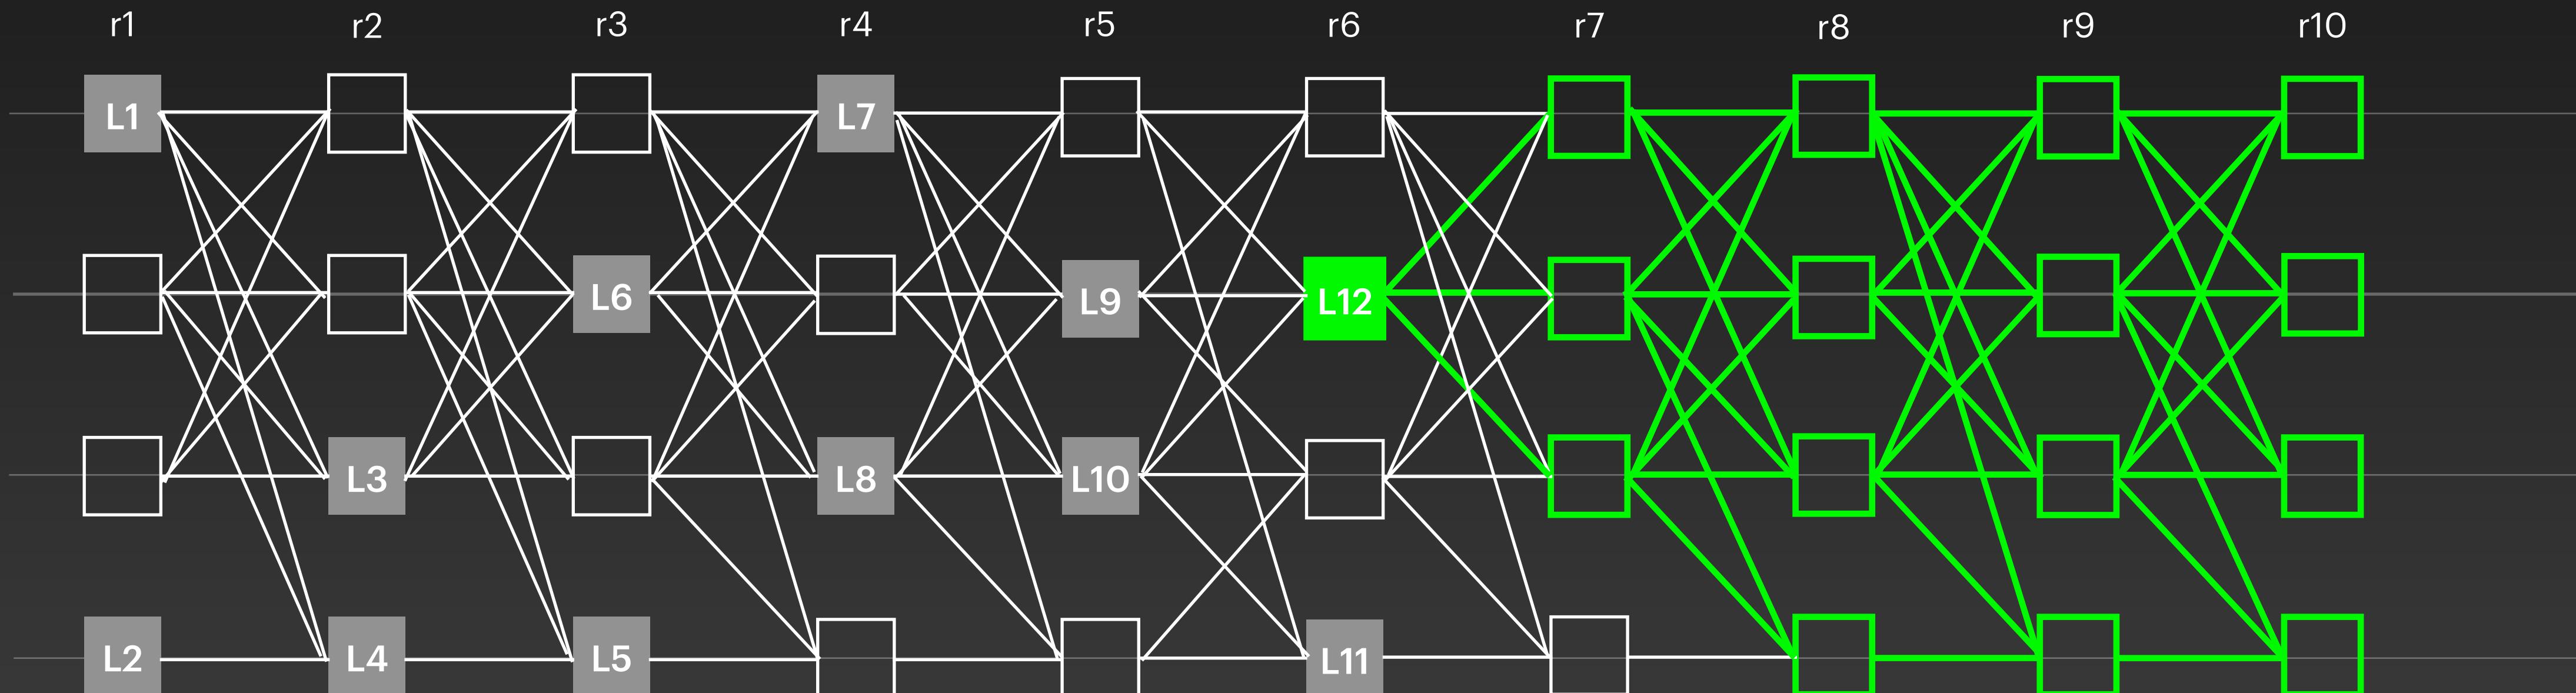
- **Skip** if $2f+1$ blames
- **Commit** if $2f+1$ certificates
- **Undecided** otherwise



Direct Decision Rule

On each leader starting from highest round:

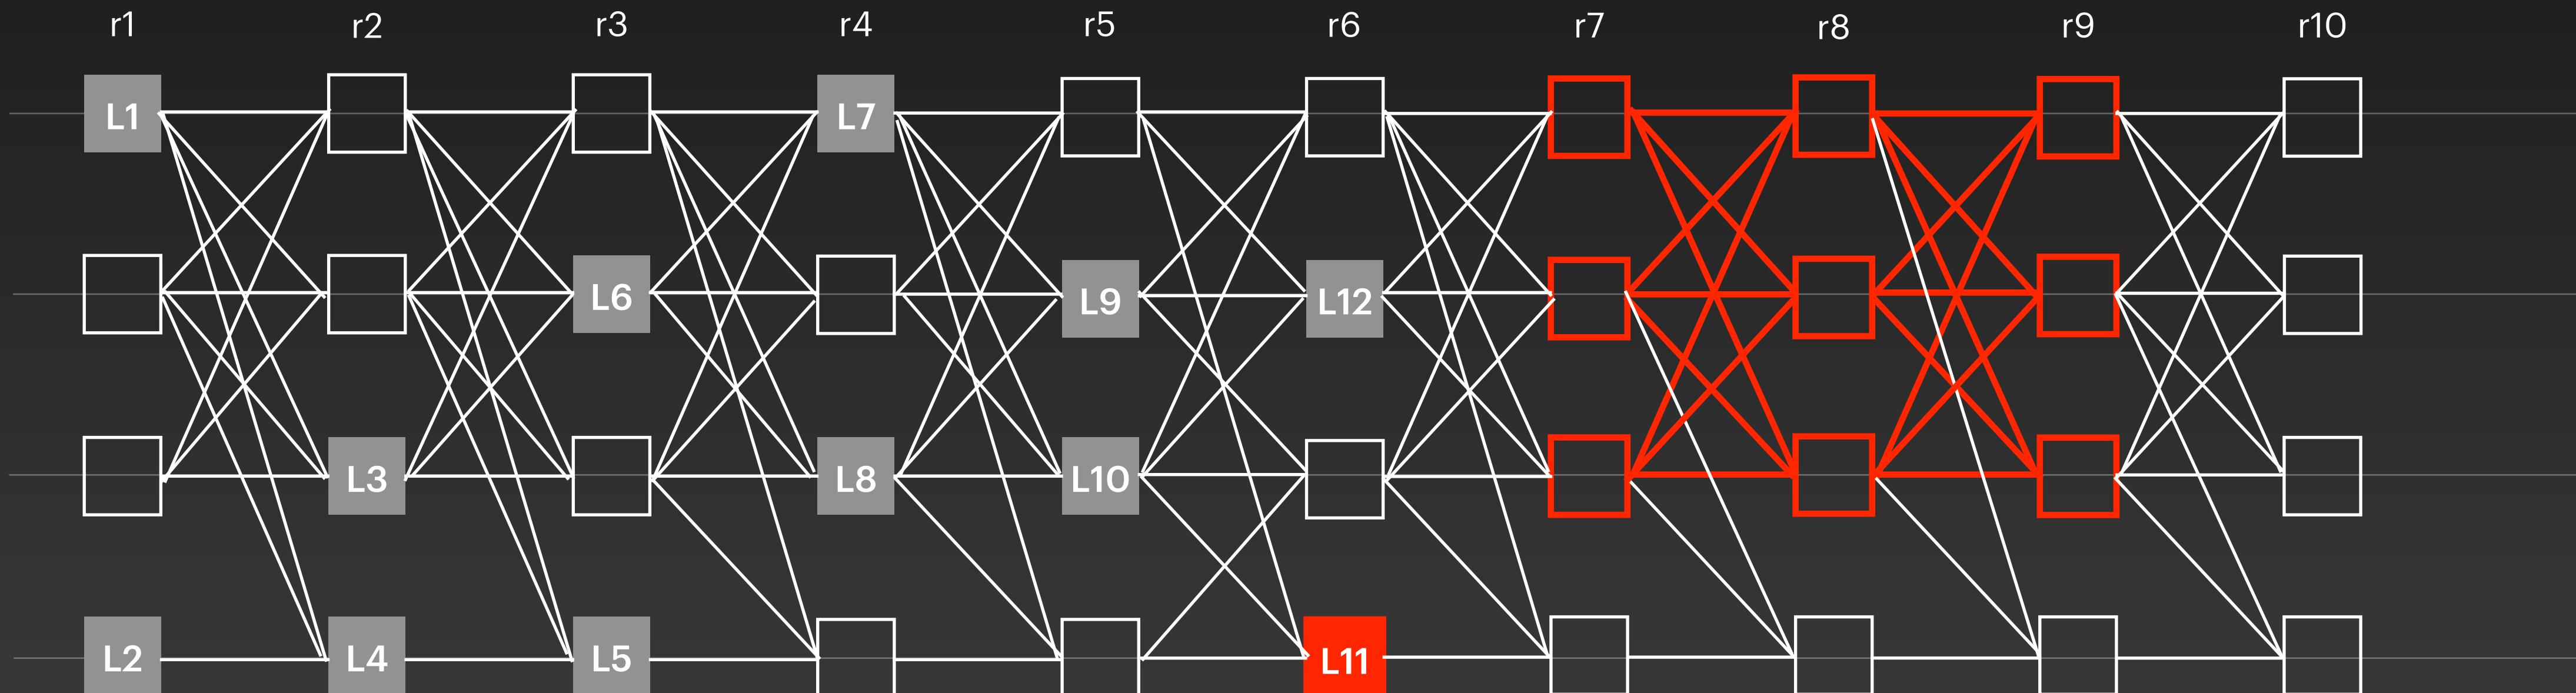
- **Skip** if $2f+1$ blames
- **Commit** if $2f+1$ certificates
- **Undecided** otherwise



Direct Decision Rule

On each leader starting from highest round:

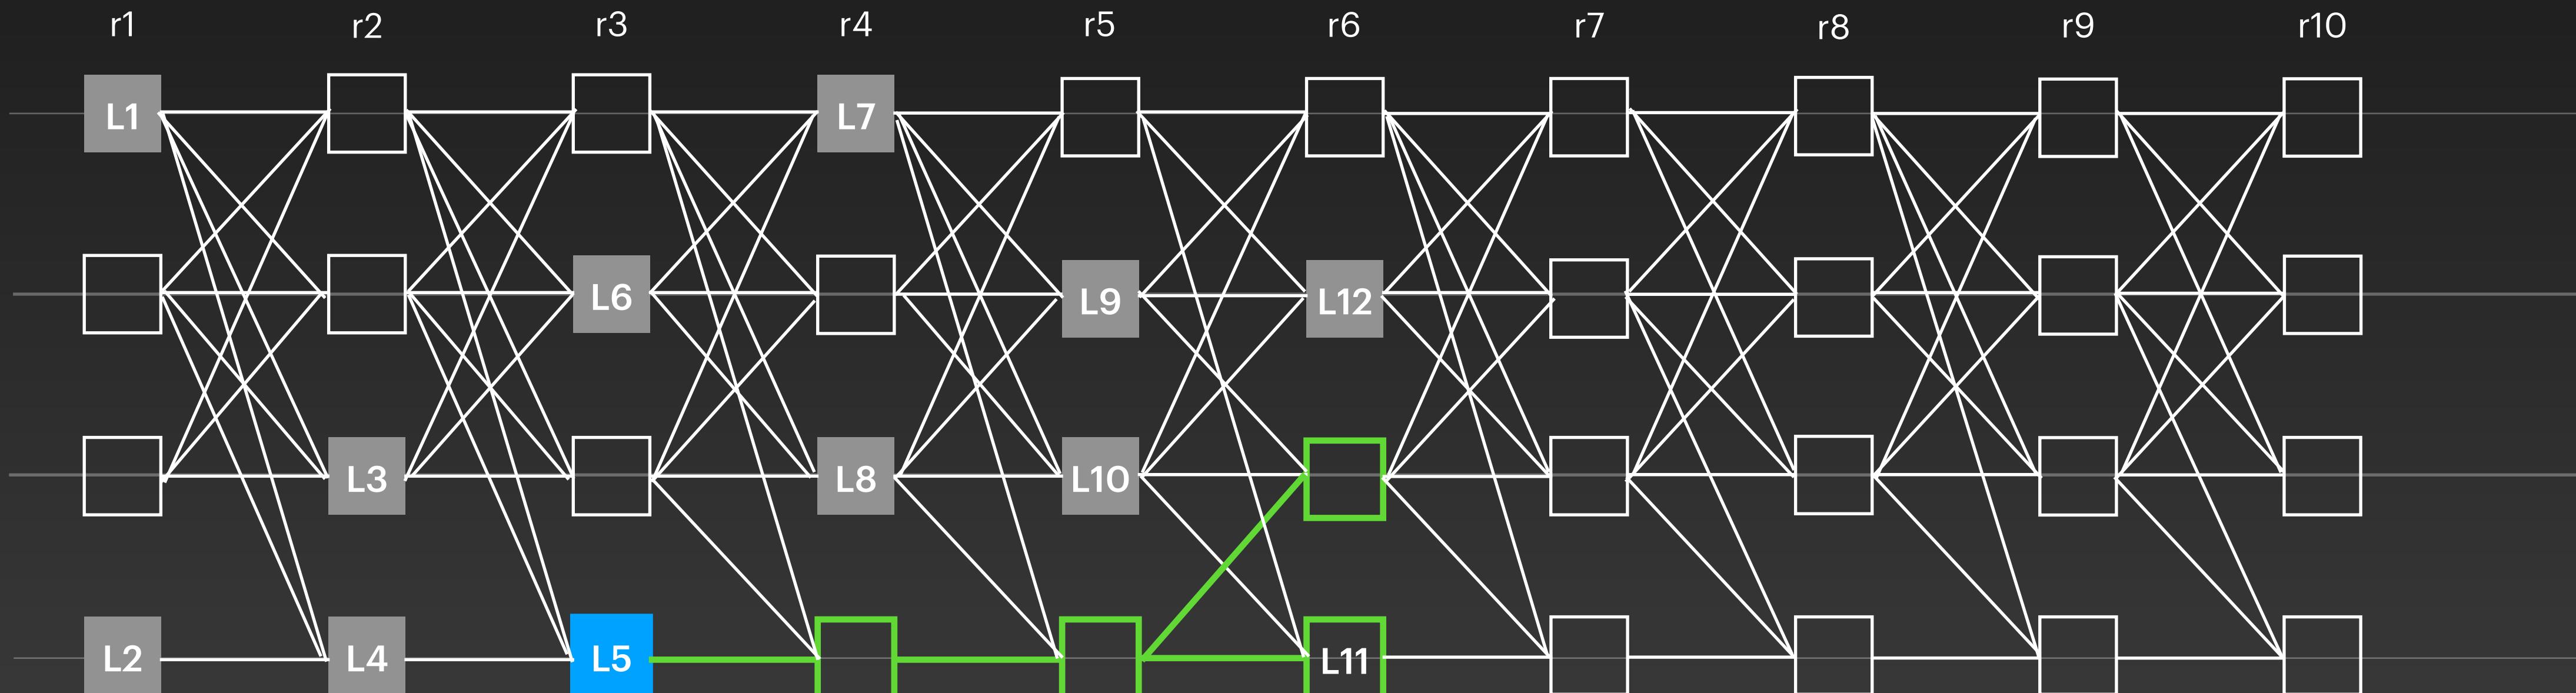
- **Skip** if $2f+1$ blames
- **Commit** if $2f+1$ certificates
- **Undecided** otherwise



Direct Decision Rule

On each leader starting from highest round:

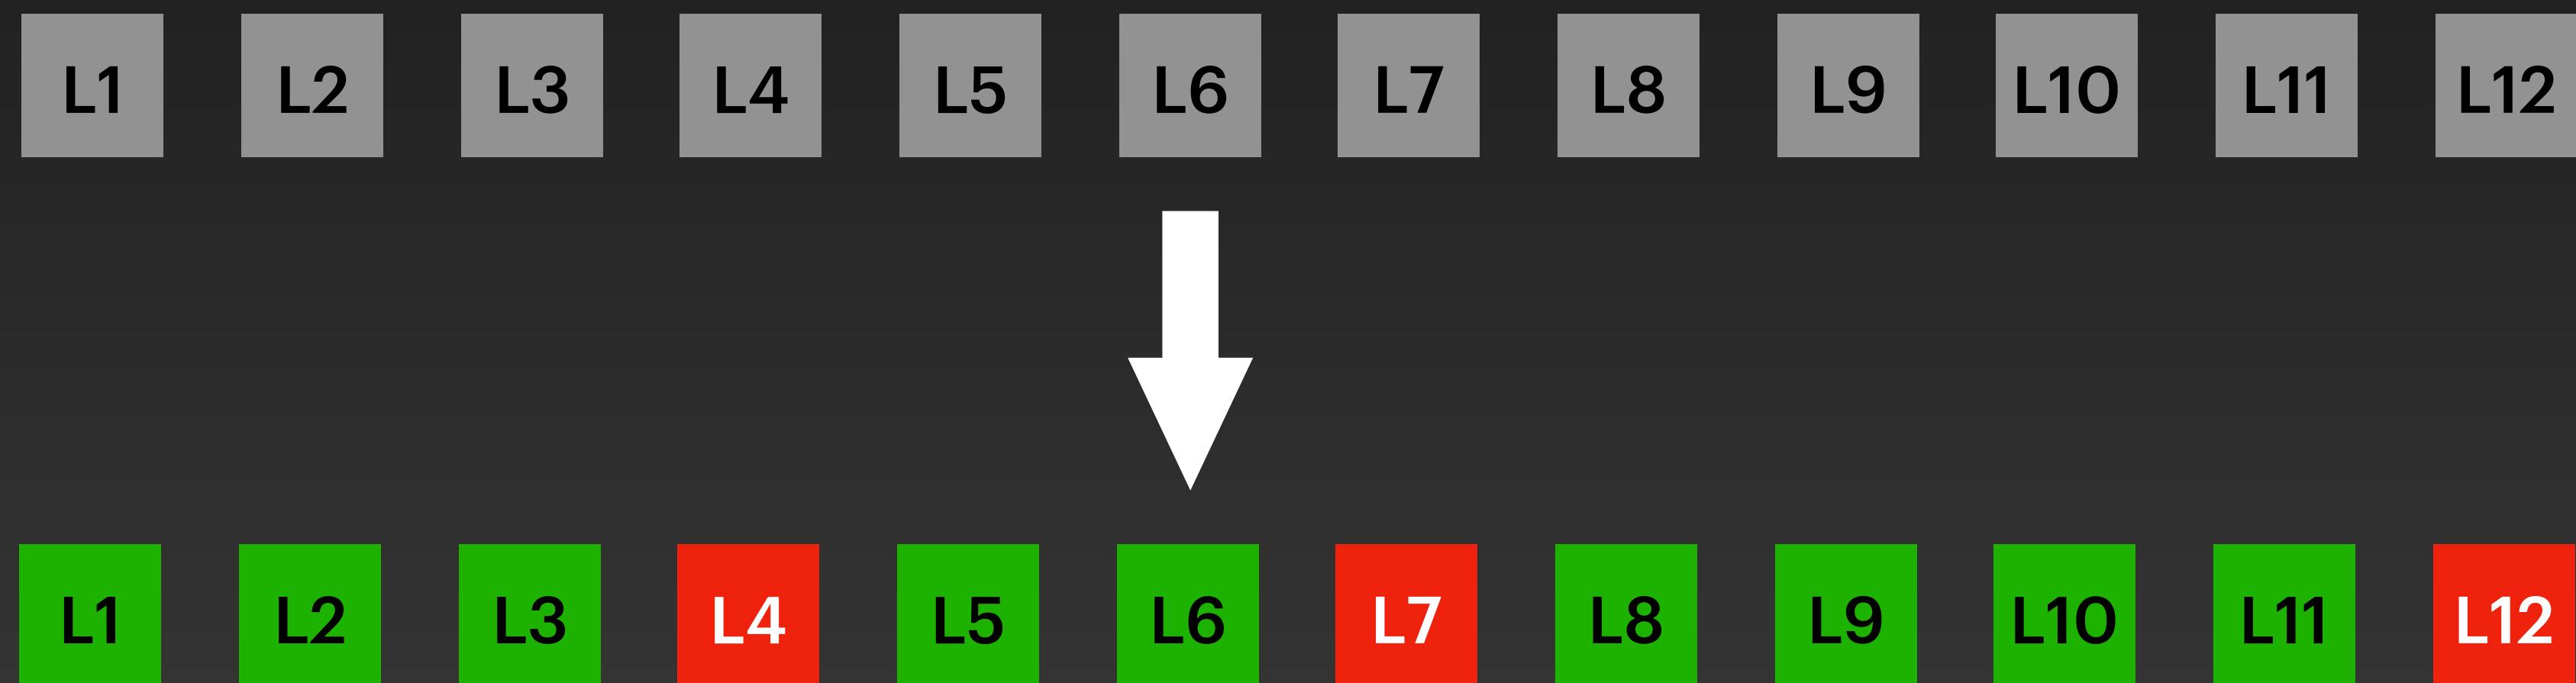
- **Skip** if $2f+1$ blames
- **Commit** if $2f+1$ certificates
- **Undecided** otherwise



End Goal

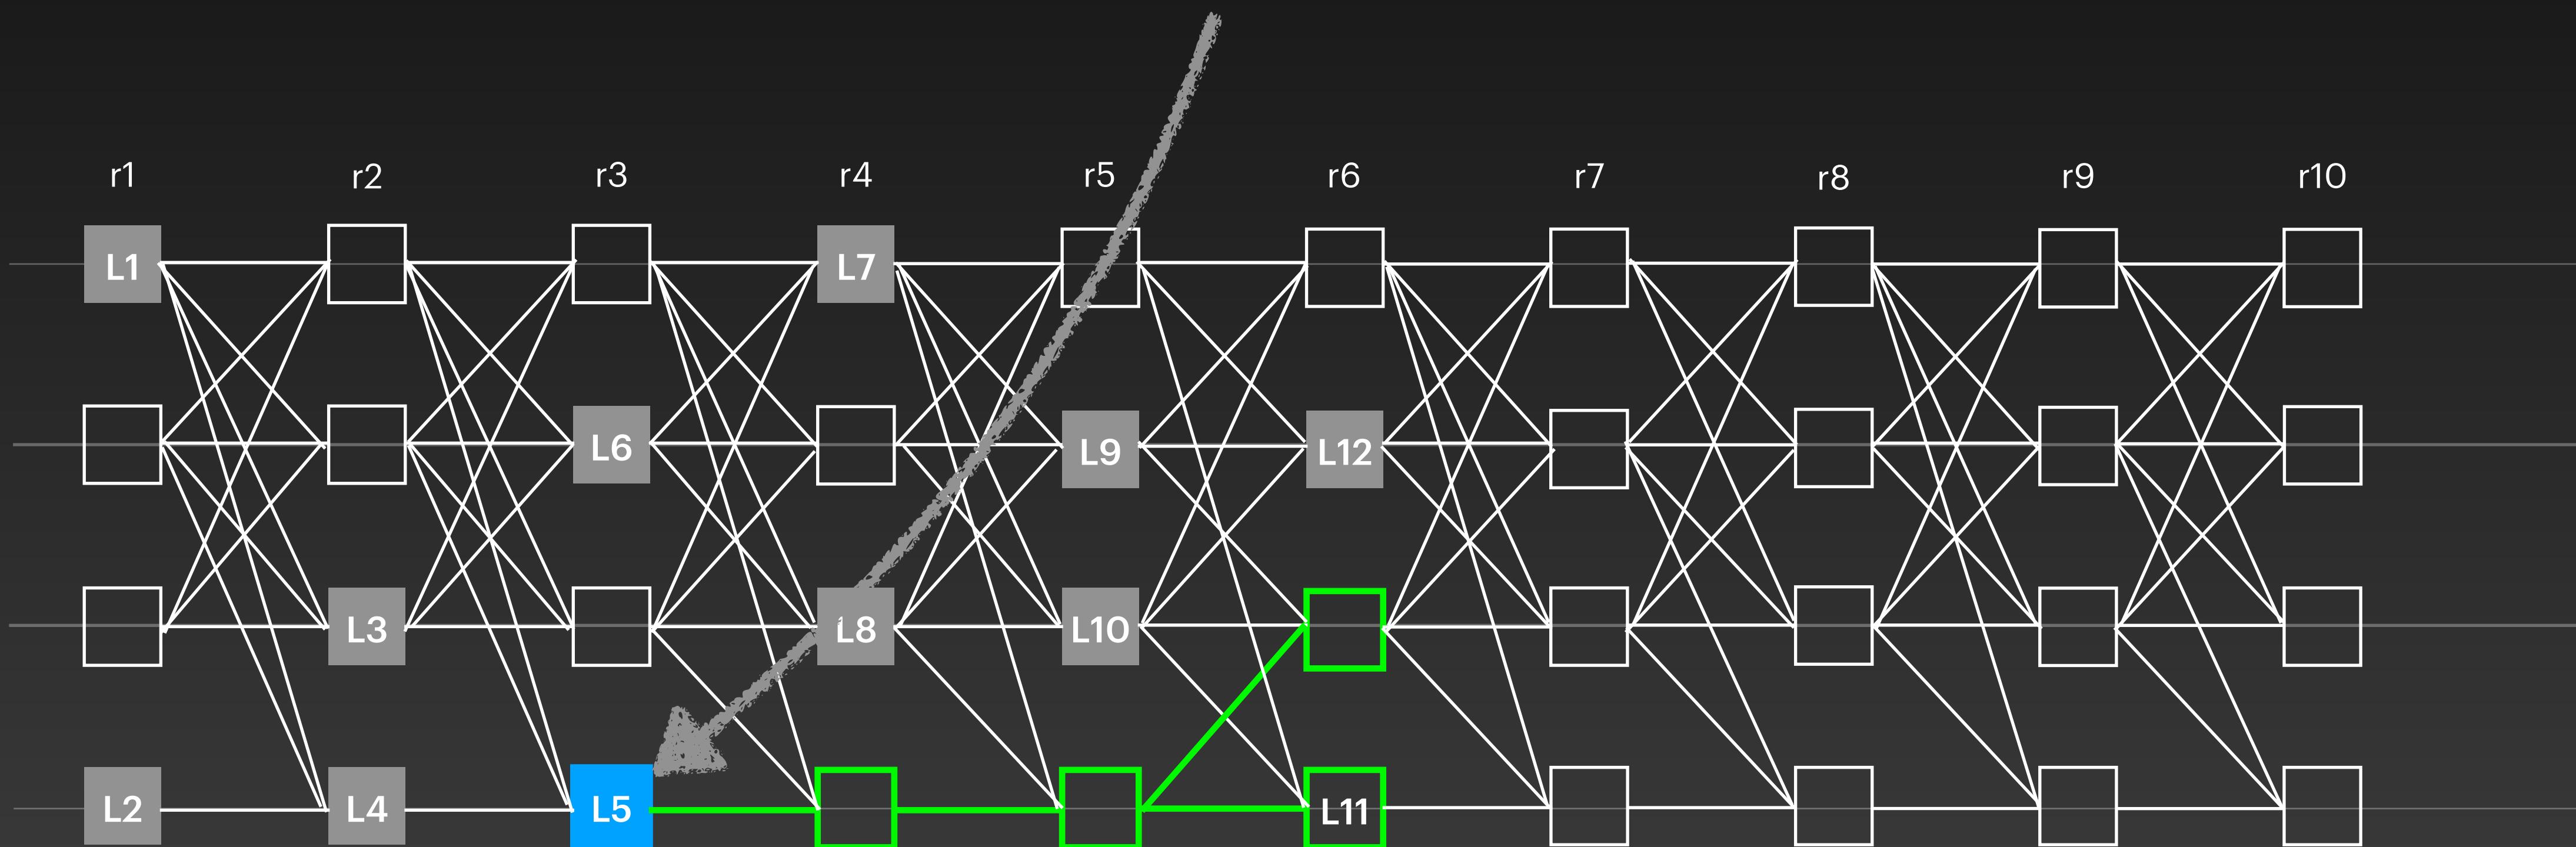
Output a sequence of leaders that are:

- Marked as **Commit** or **Skip** (not **Undecided**)
- Common to all nodes (despite Byzantine faults)



What to do about undecided blocks?

- **Undecided** if there's less than $2f+1$ blames or less than $2f+1$ certificates

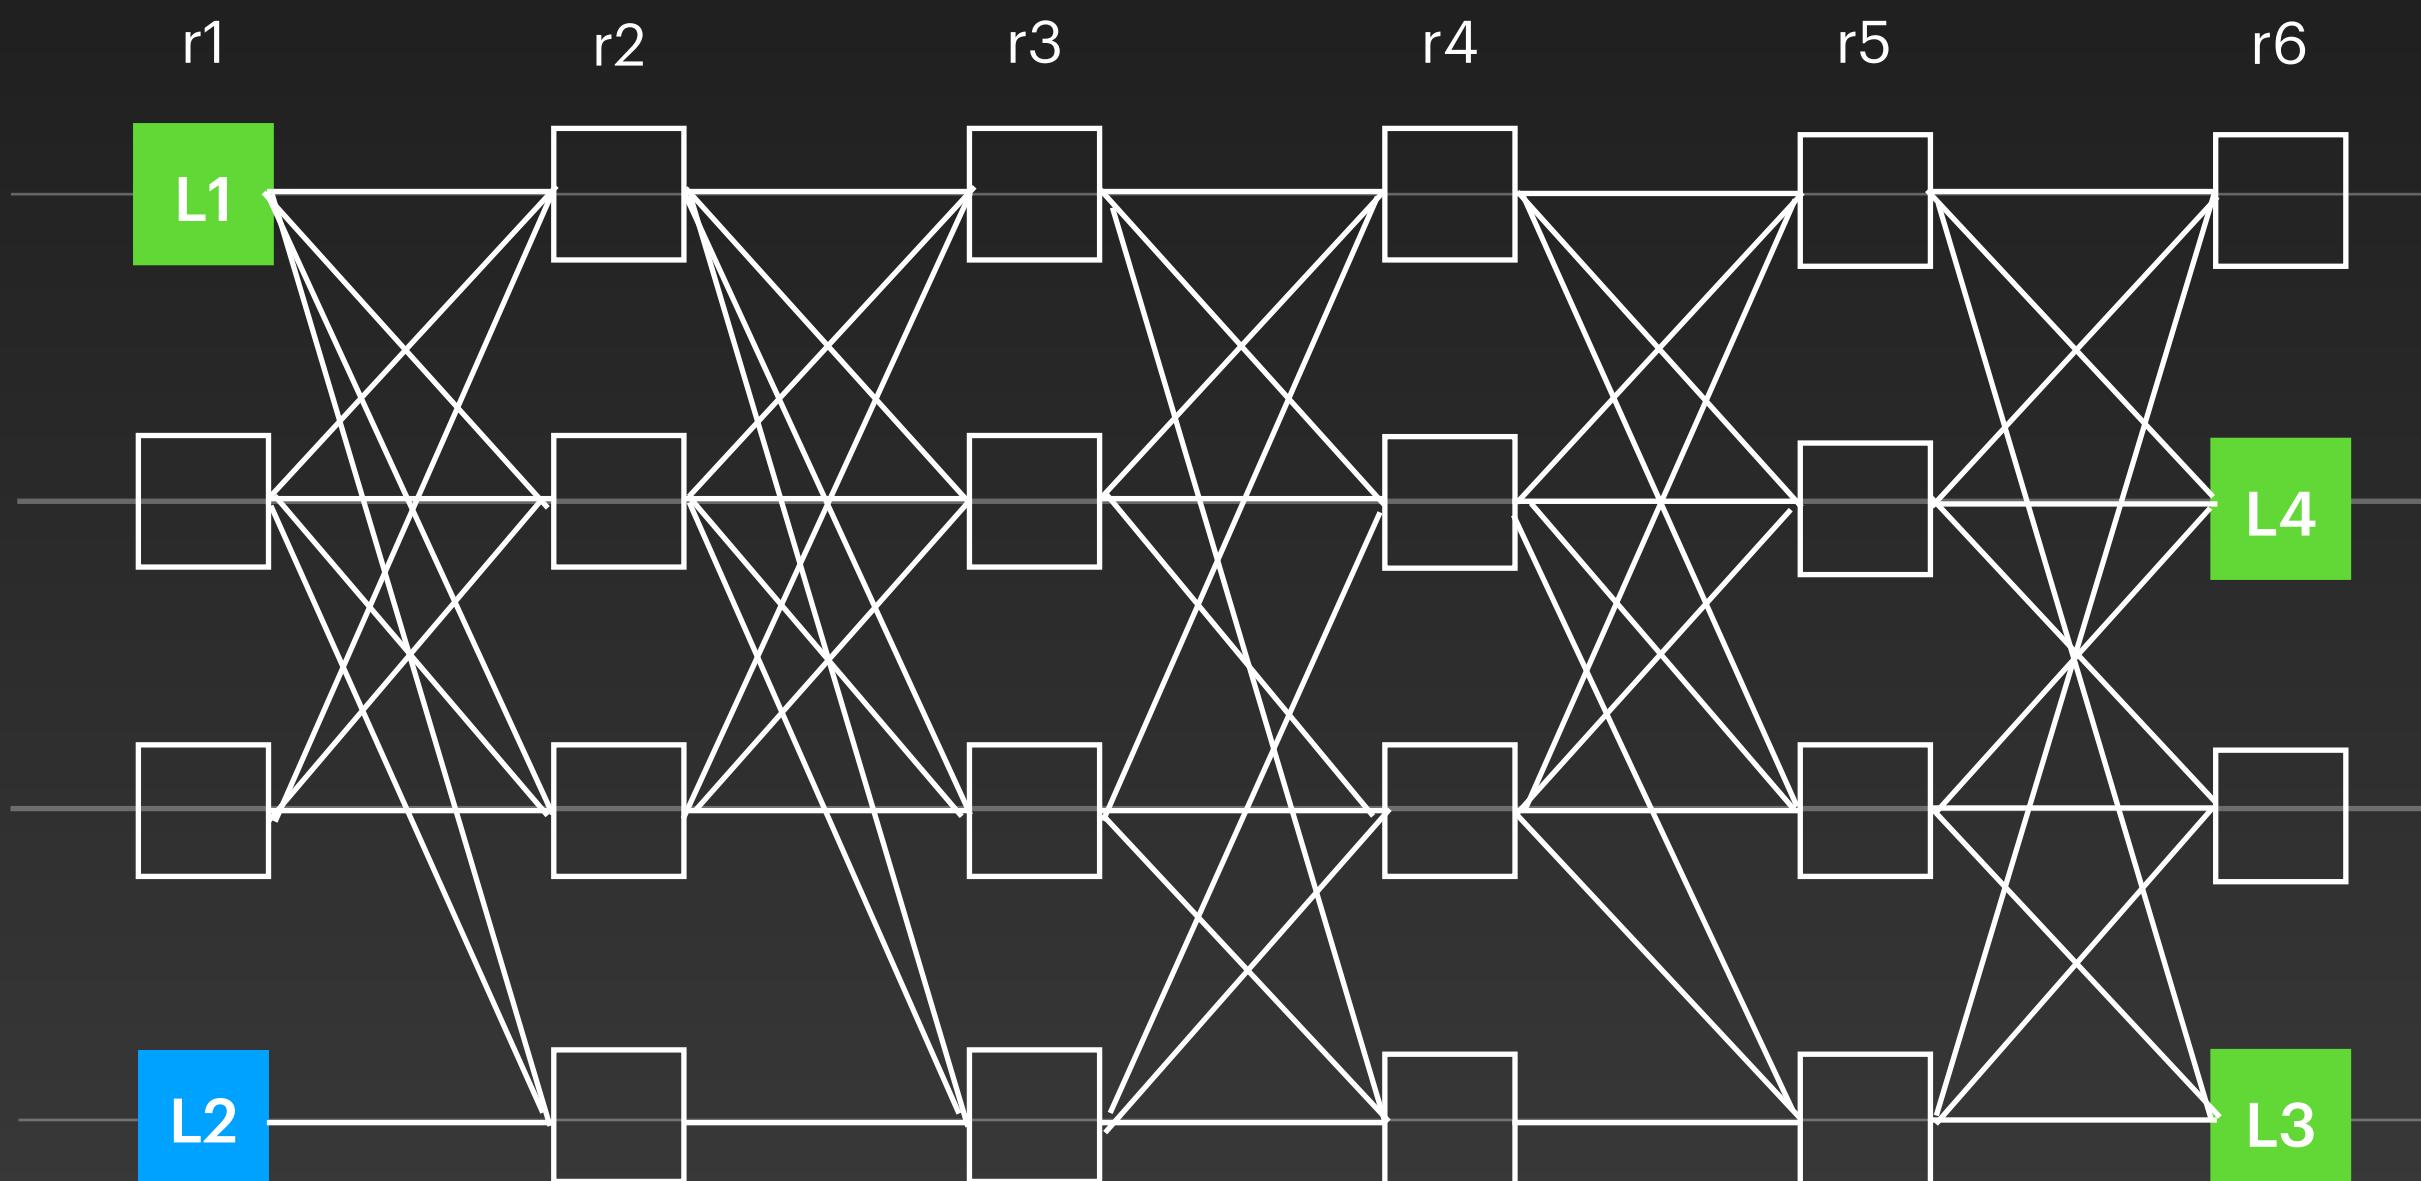


Indirect Decision Rule

Indirect Decision Rule

1. Find Anchor

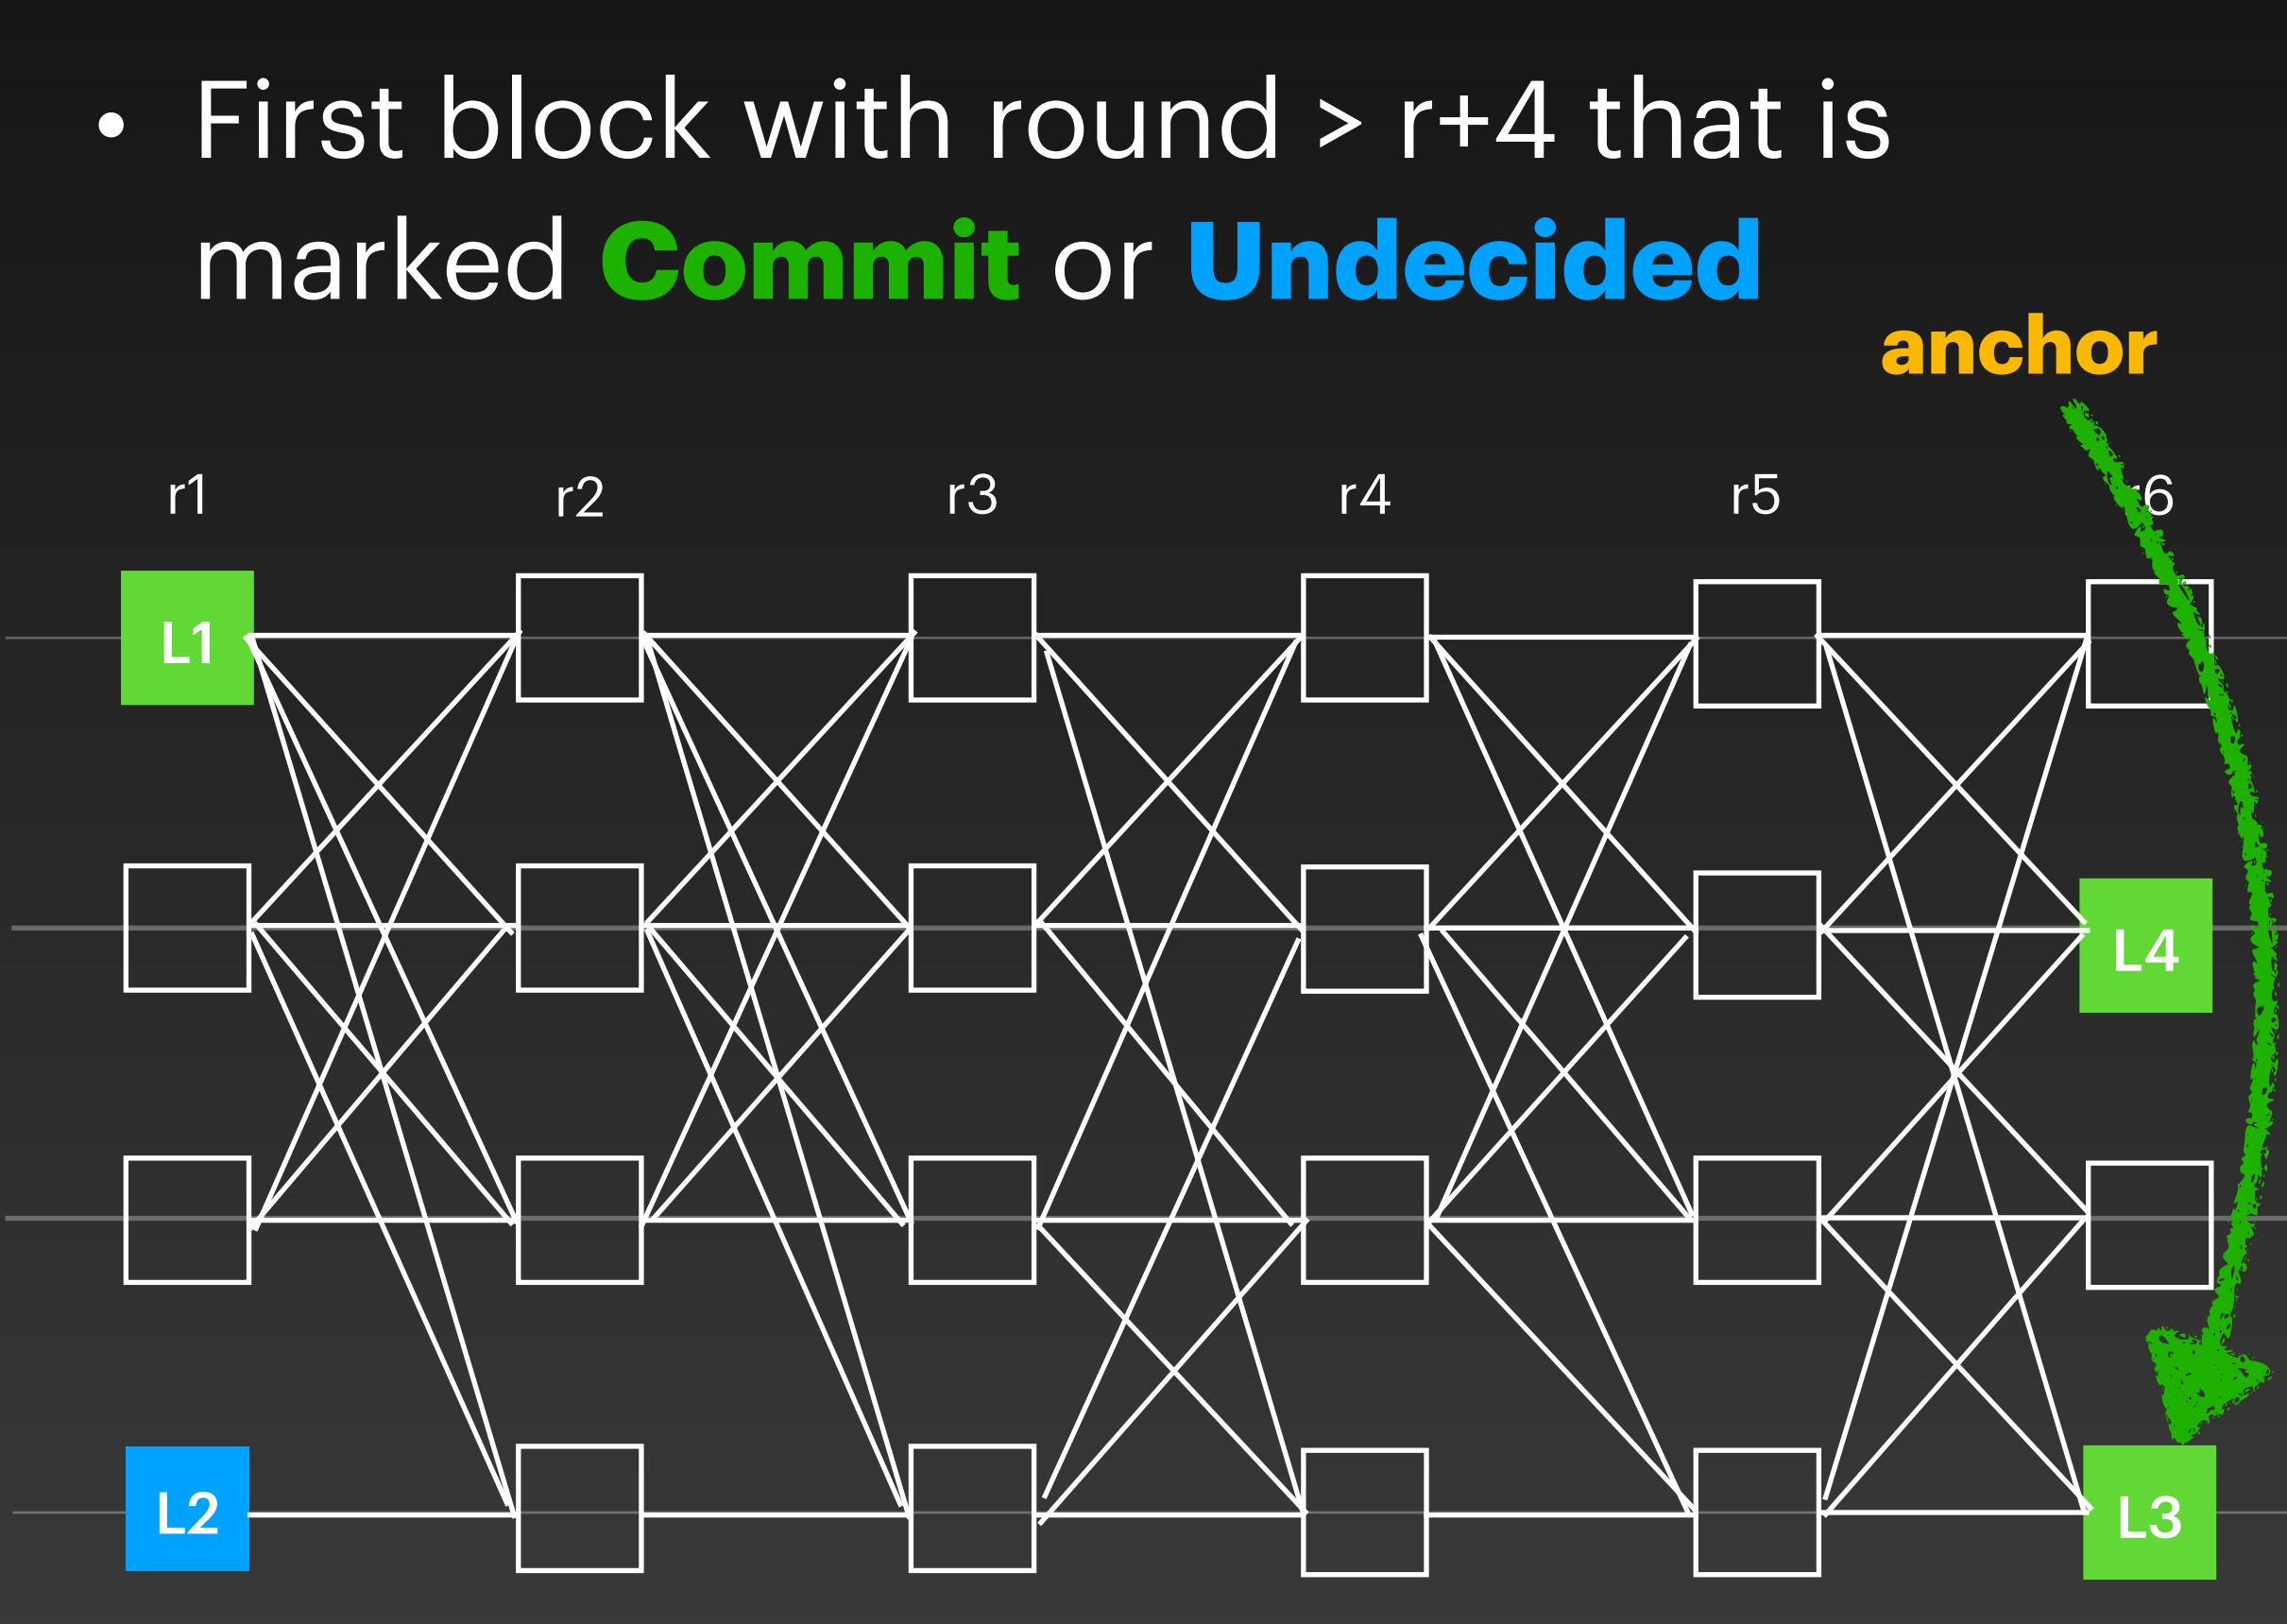
- First block with round $> r+4$ that is
Commit or **Undecided**



Indirect Decision Rule

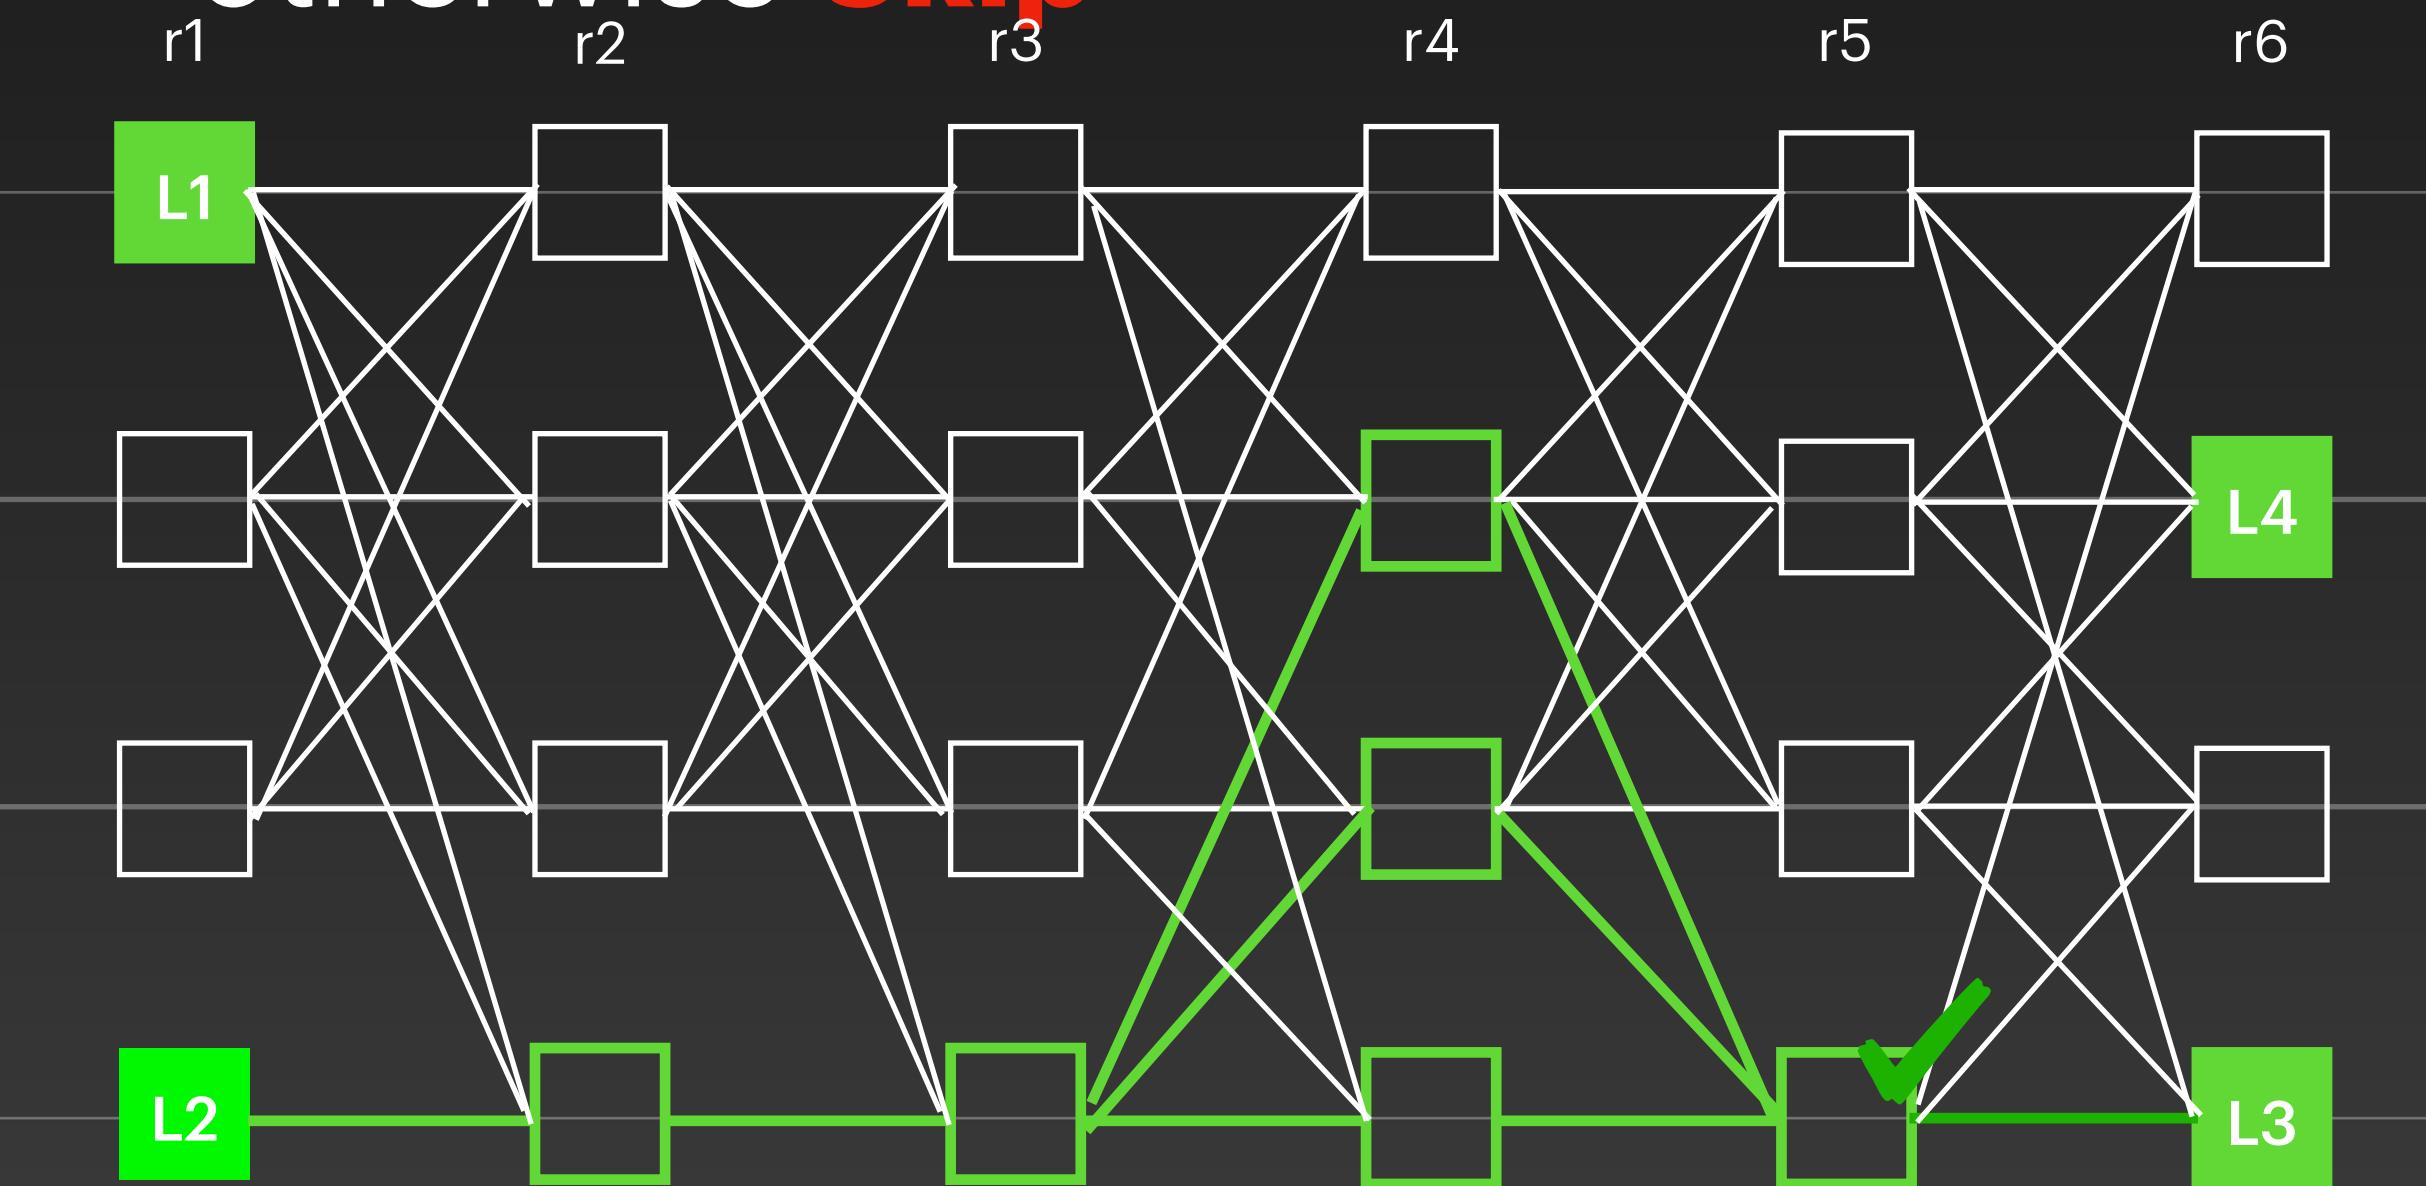
1. Find Anchor

- First block with round $> r+4$ that is marked **Commit** or **Undecided**

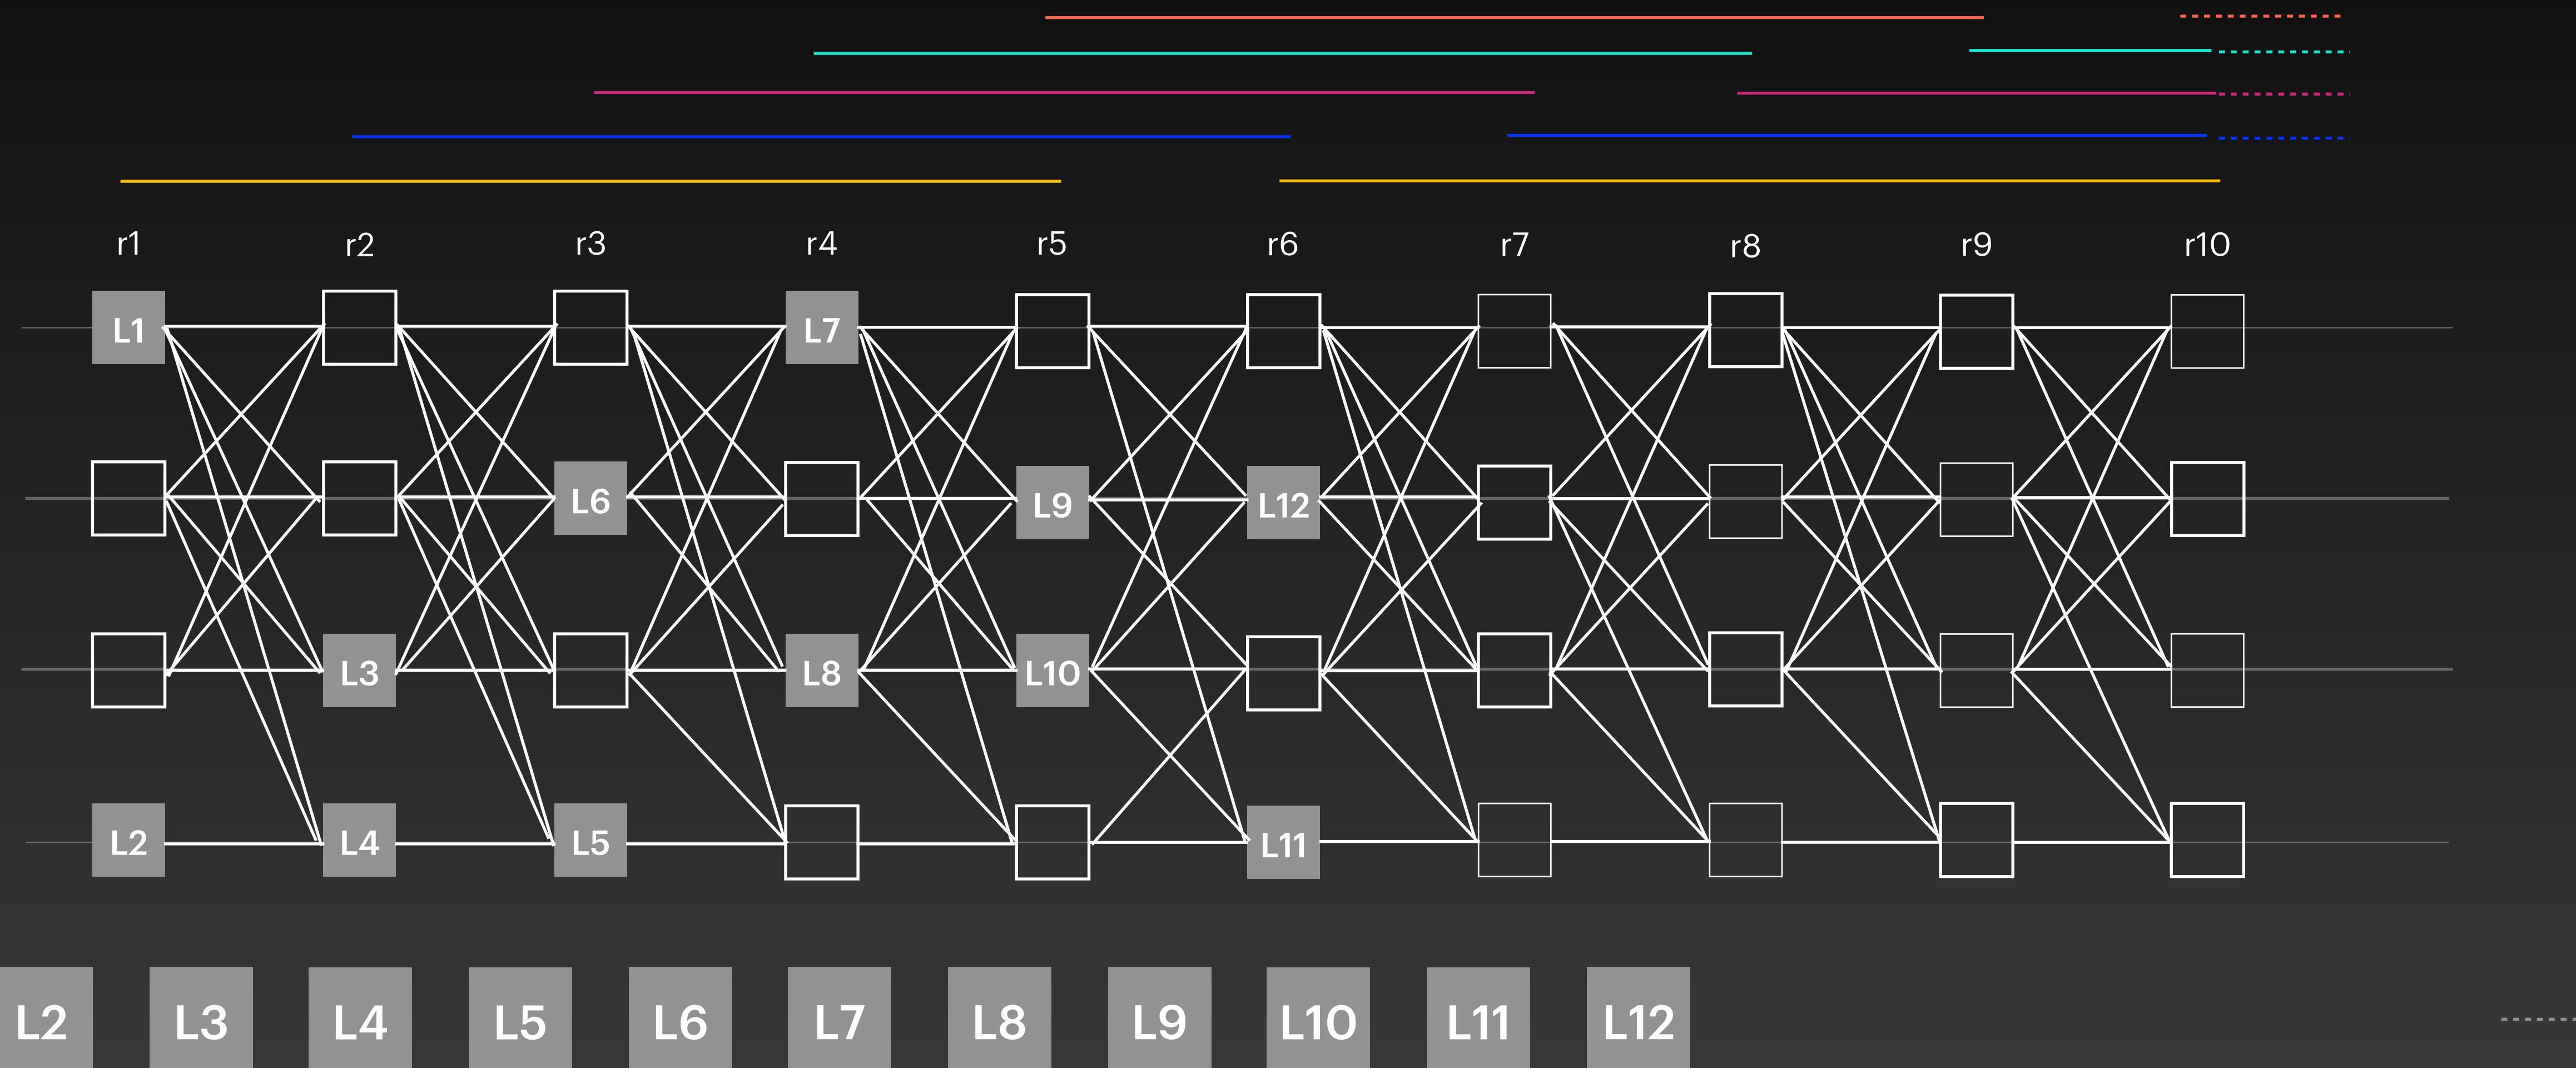


2. Certified link

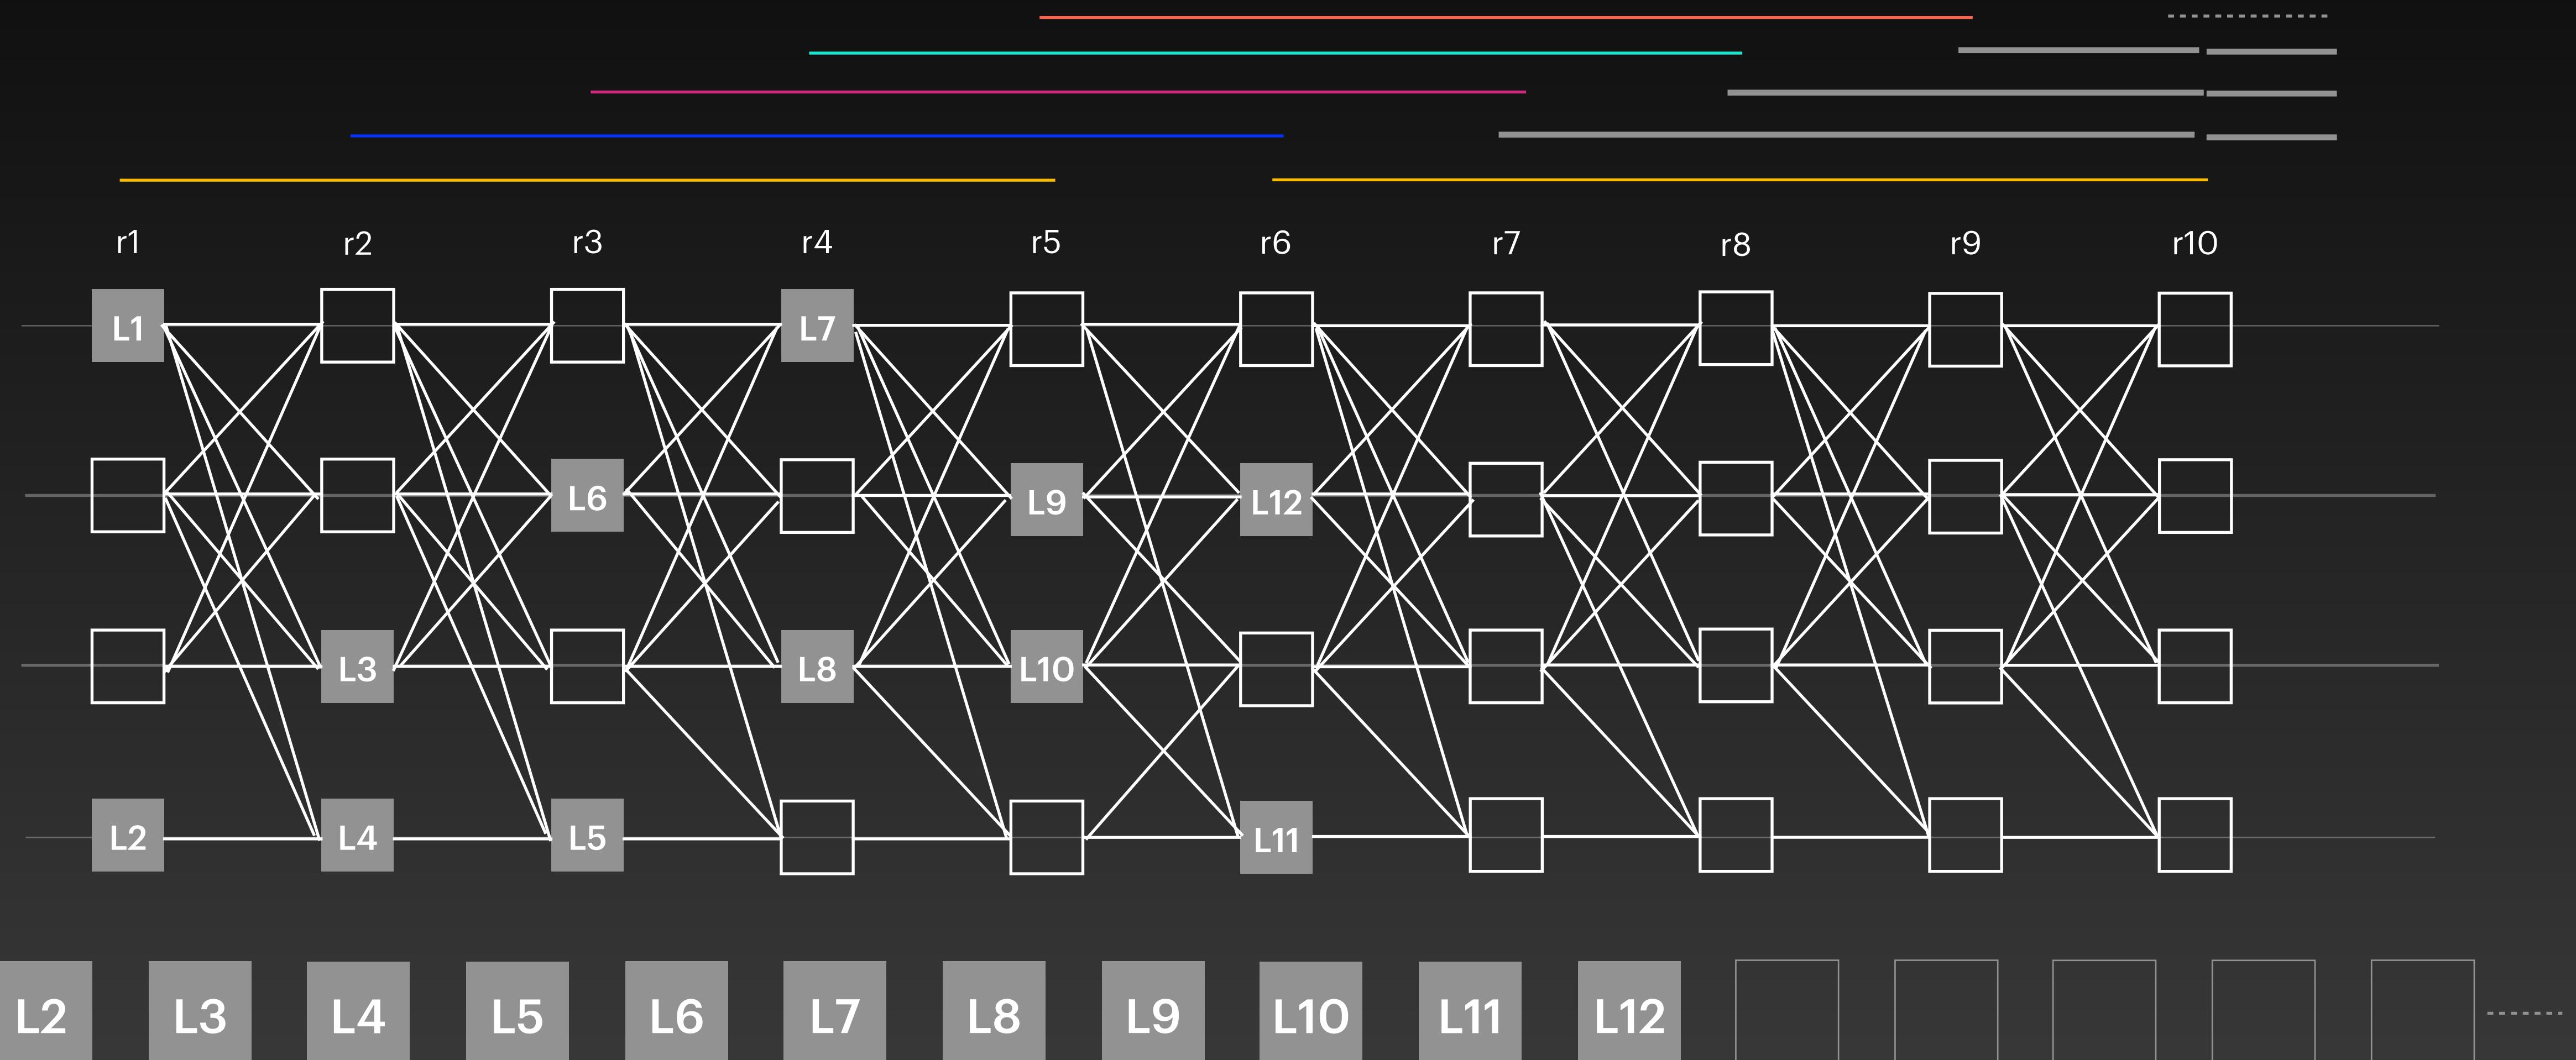
- Commit** if $B <->$ certified link $<-> A$
otherwise **Skip**



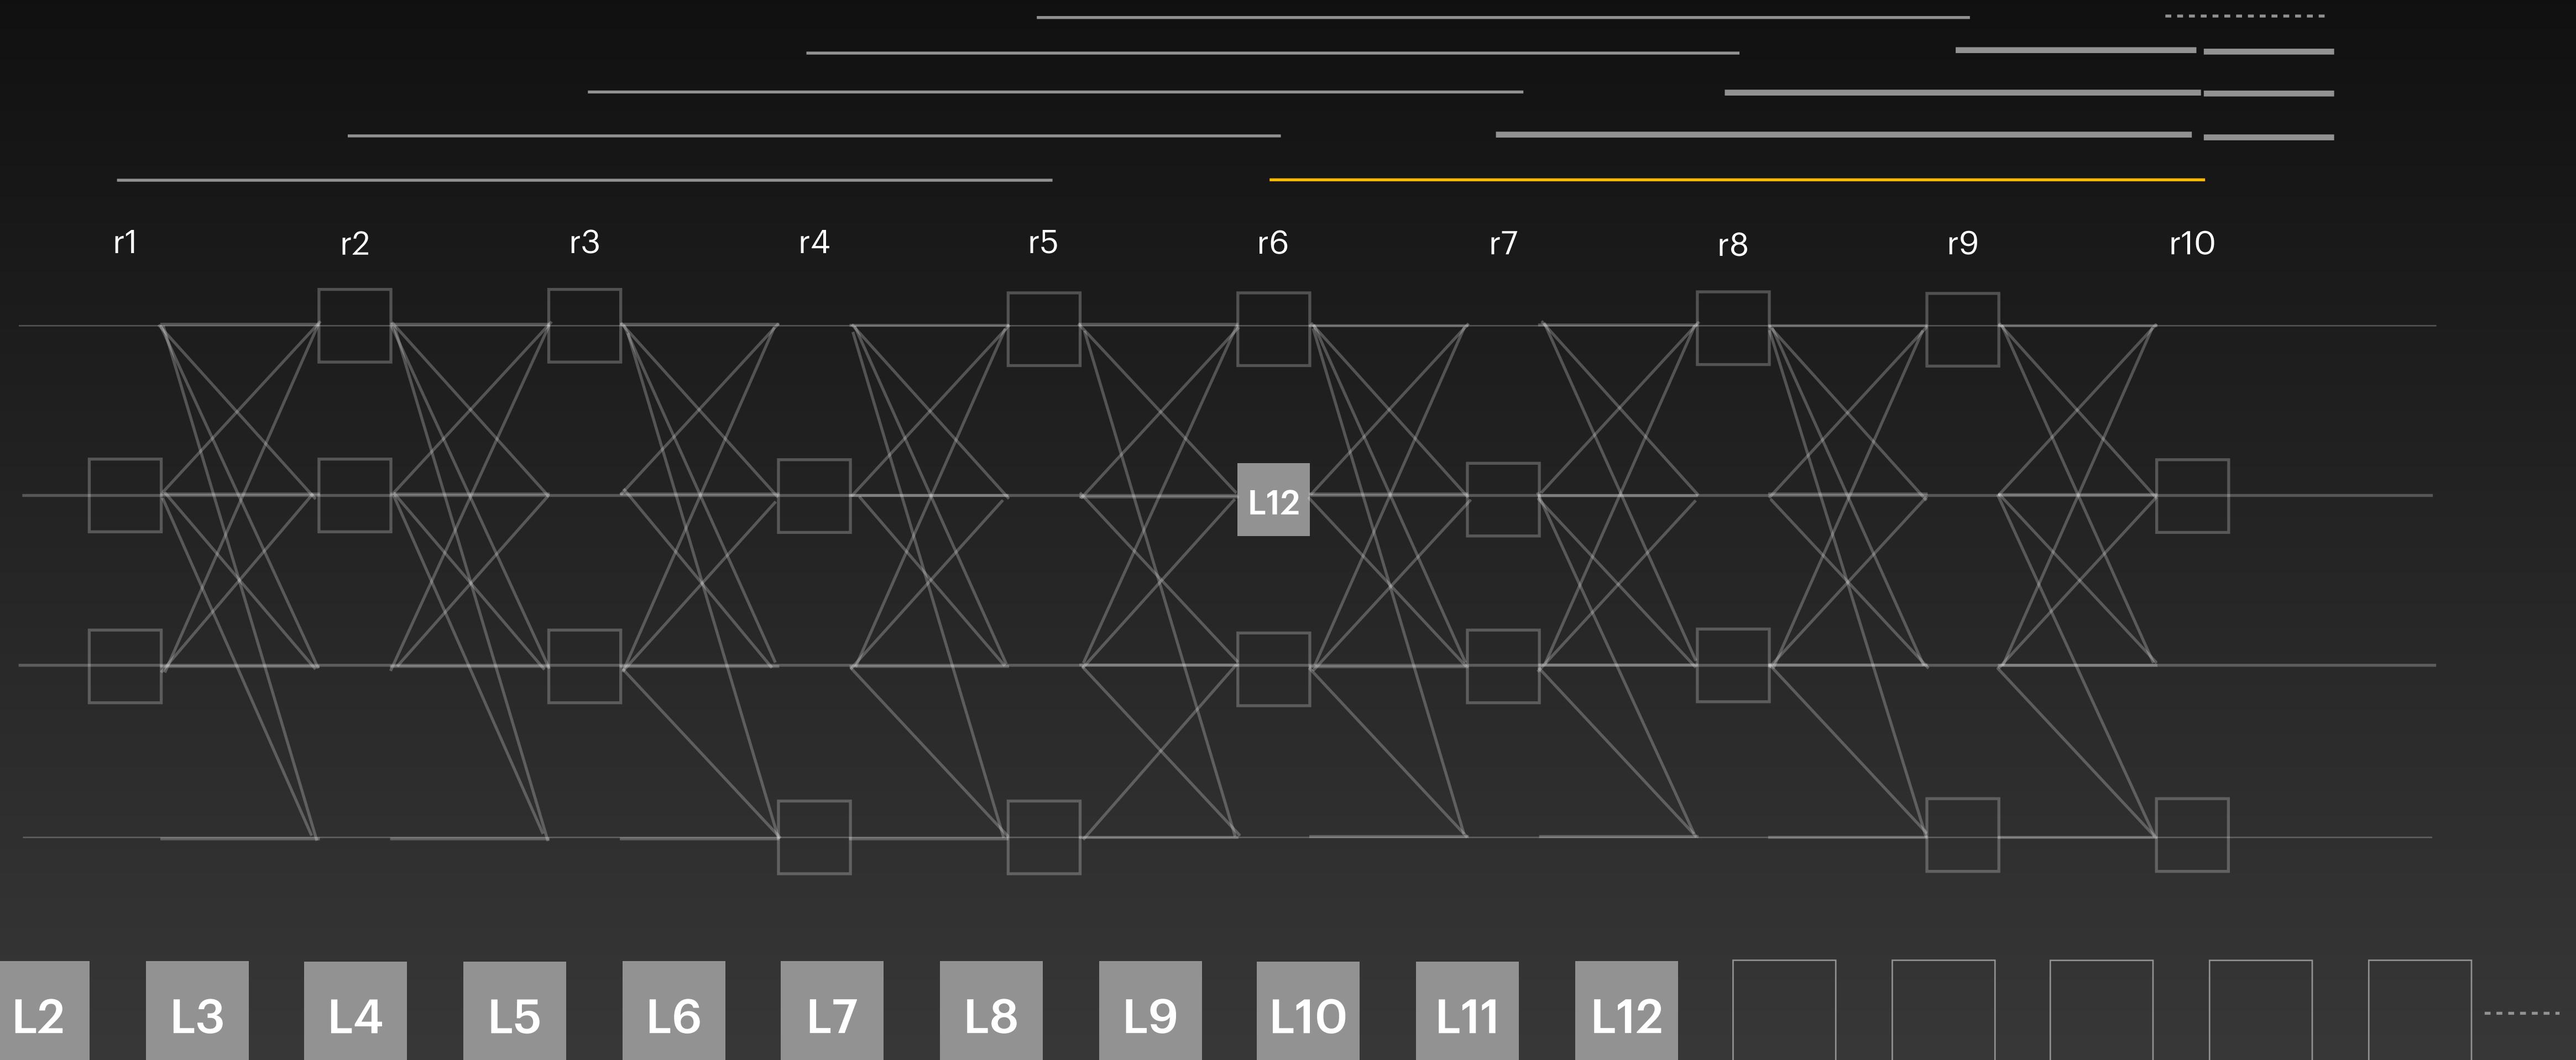
All Blocks Start As Undecided



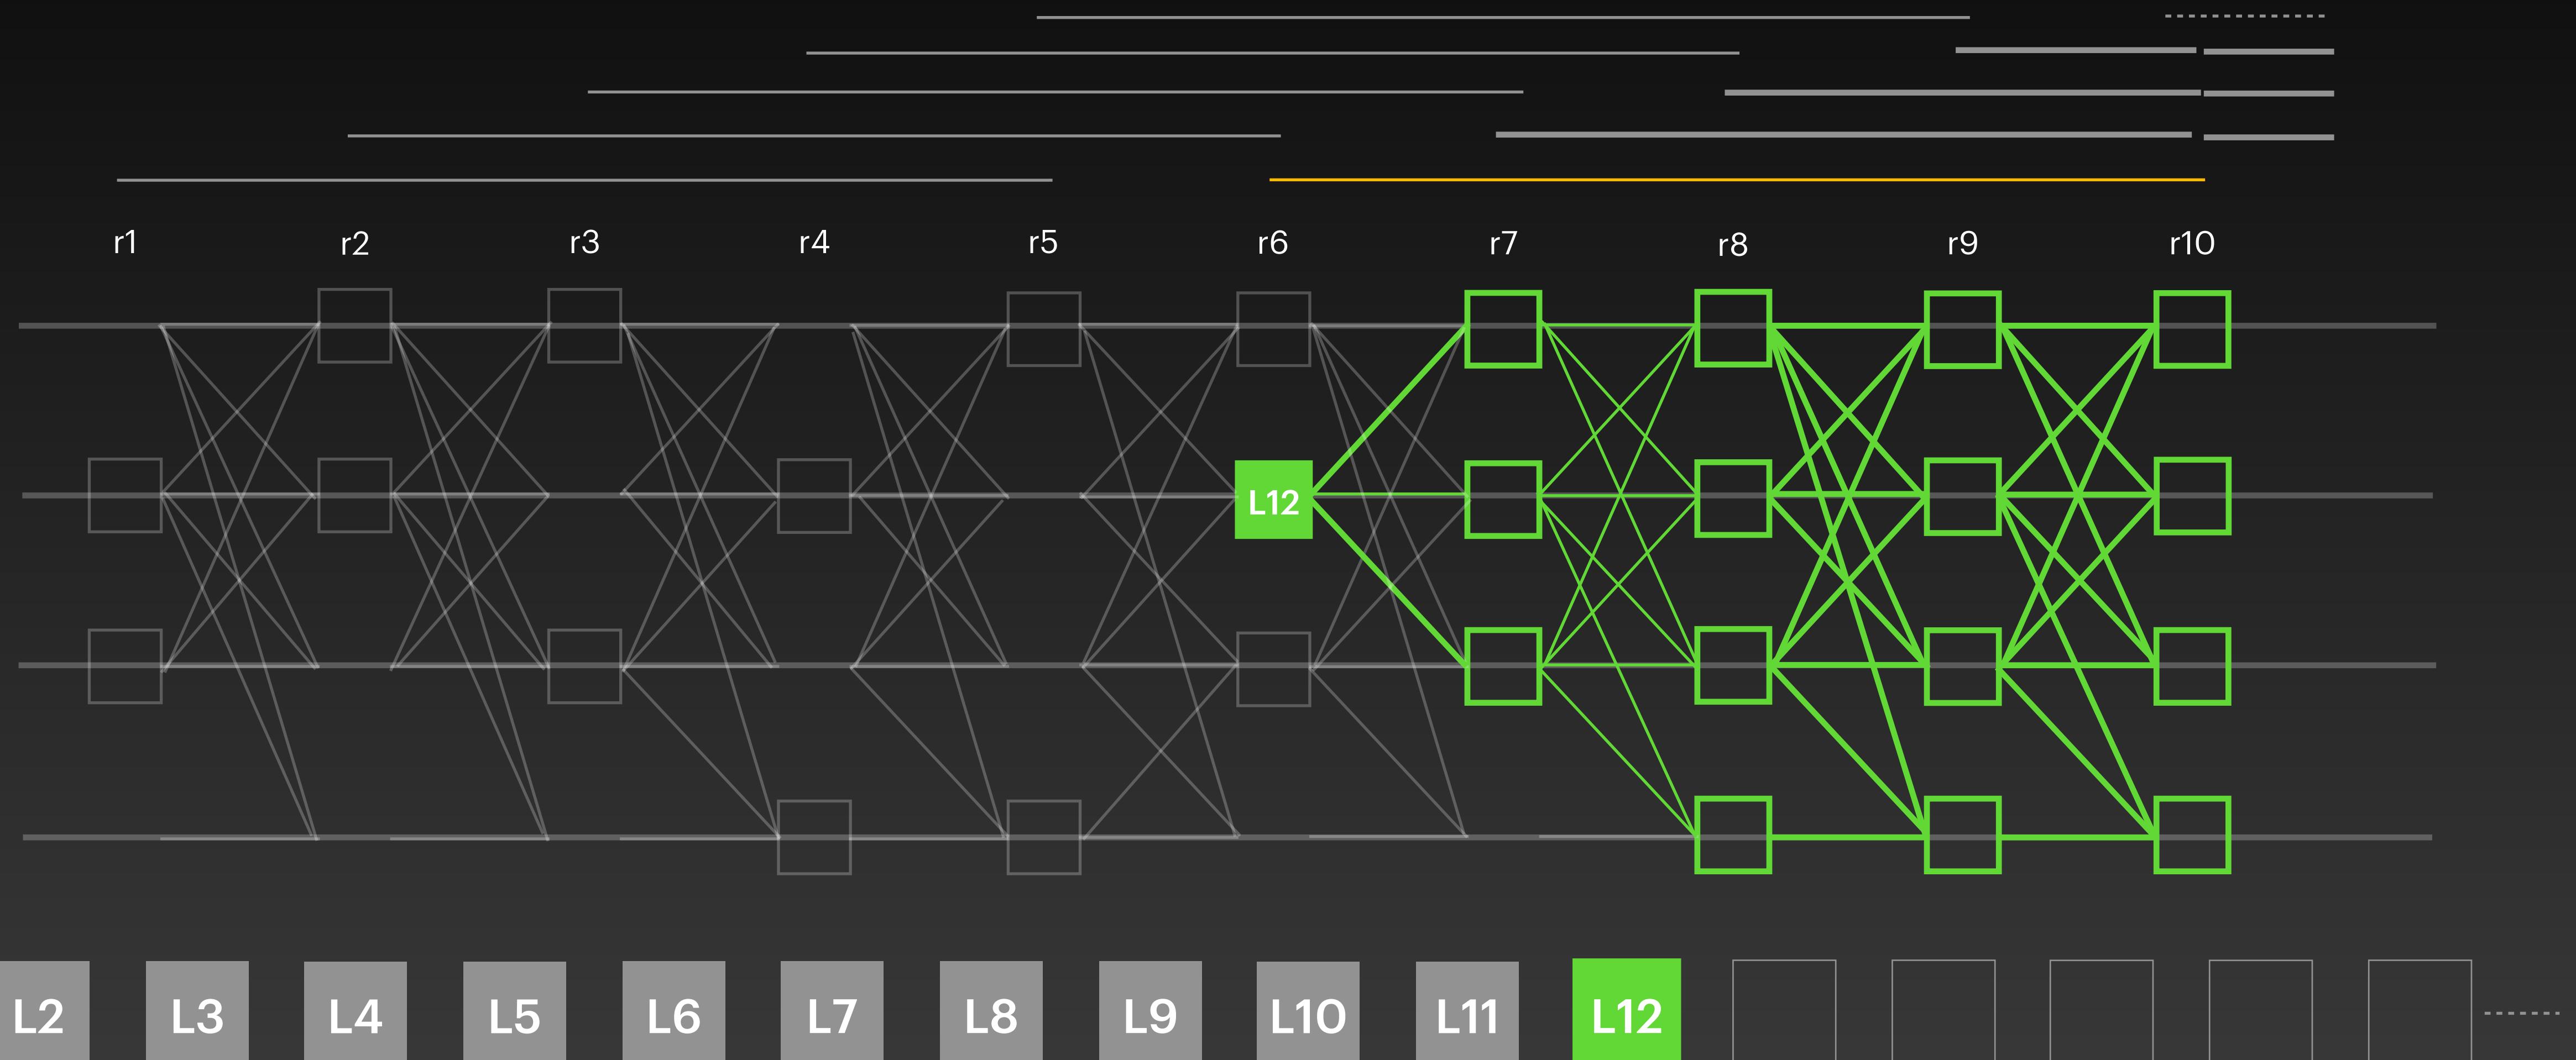
Ignore Incomplete Waves



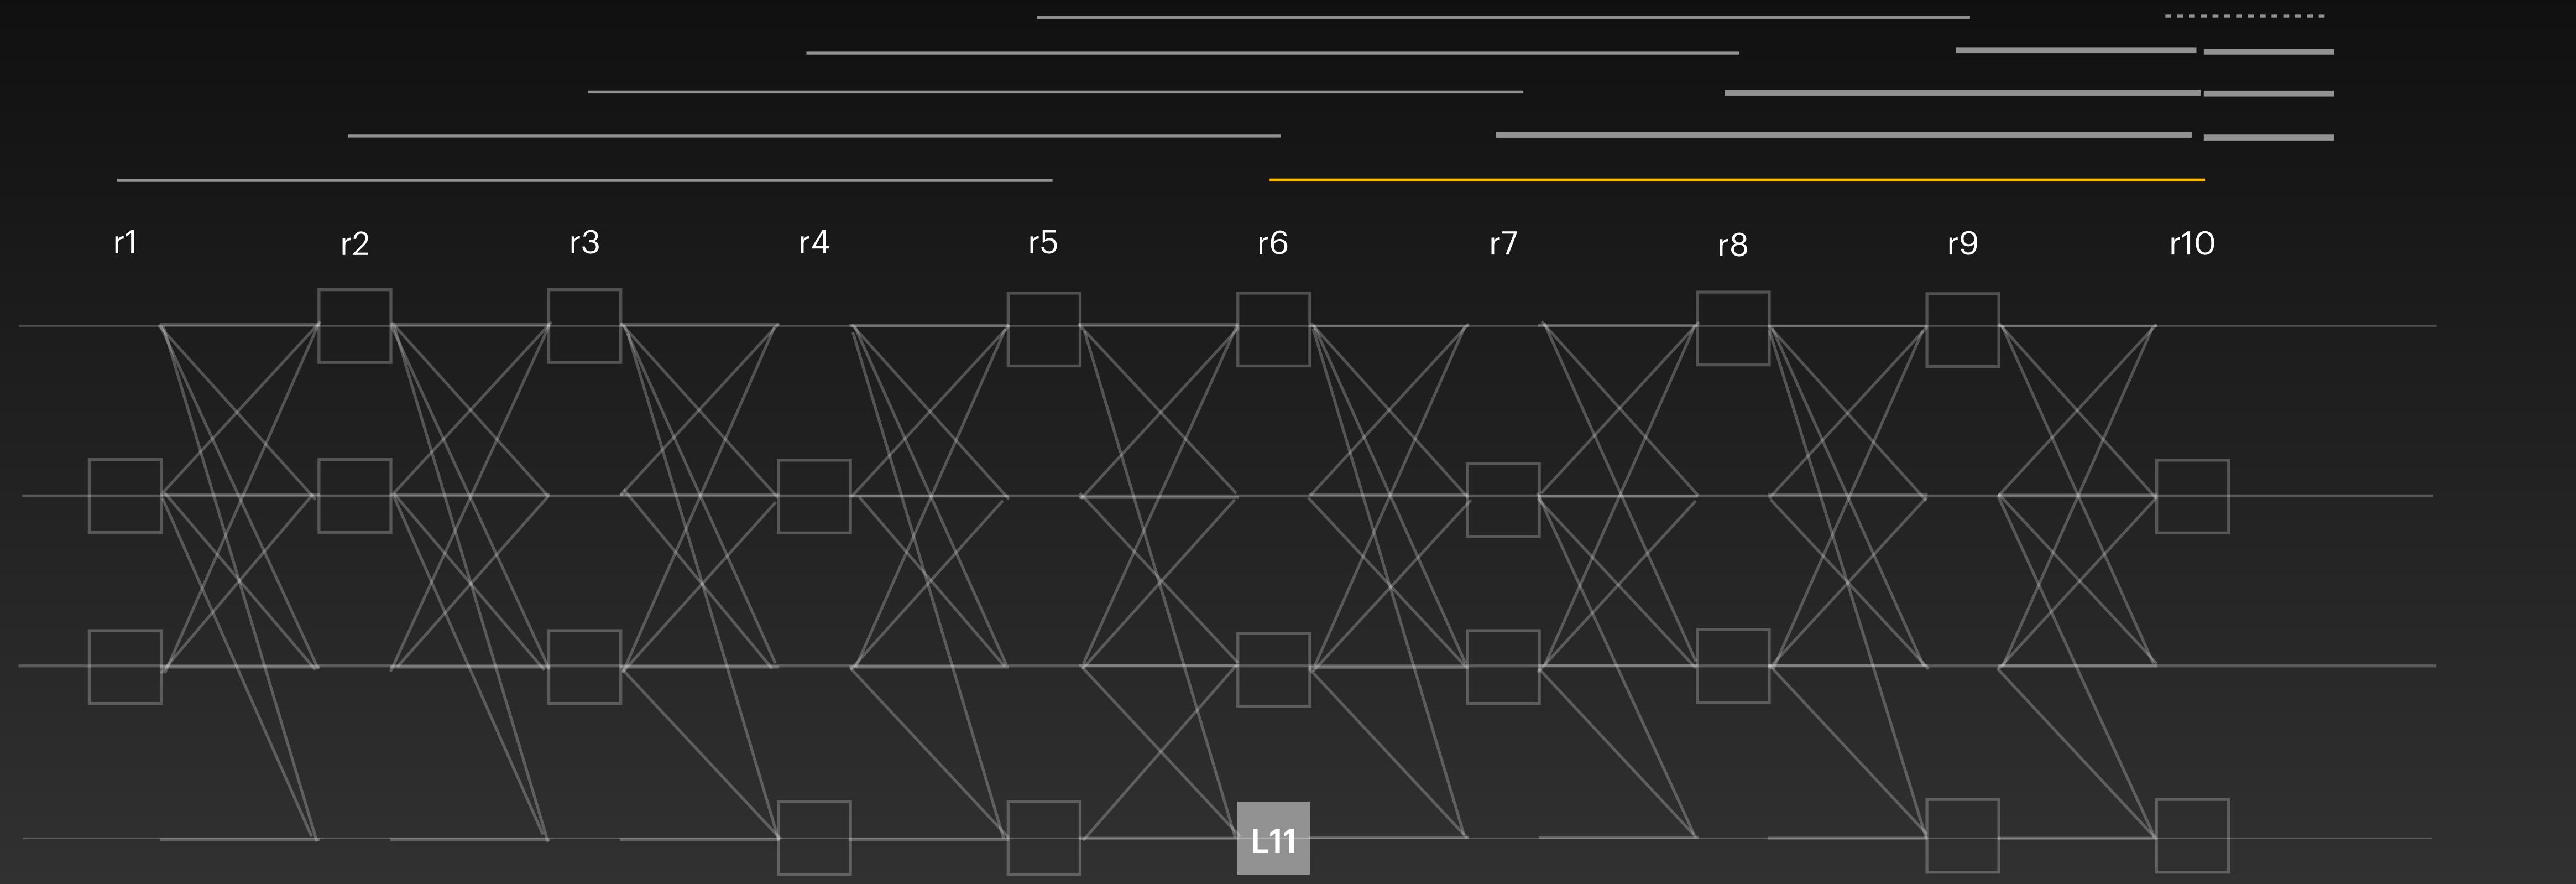
Apply Direct Rule



Apply Direct Rule



Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

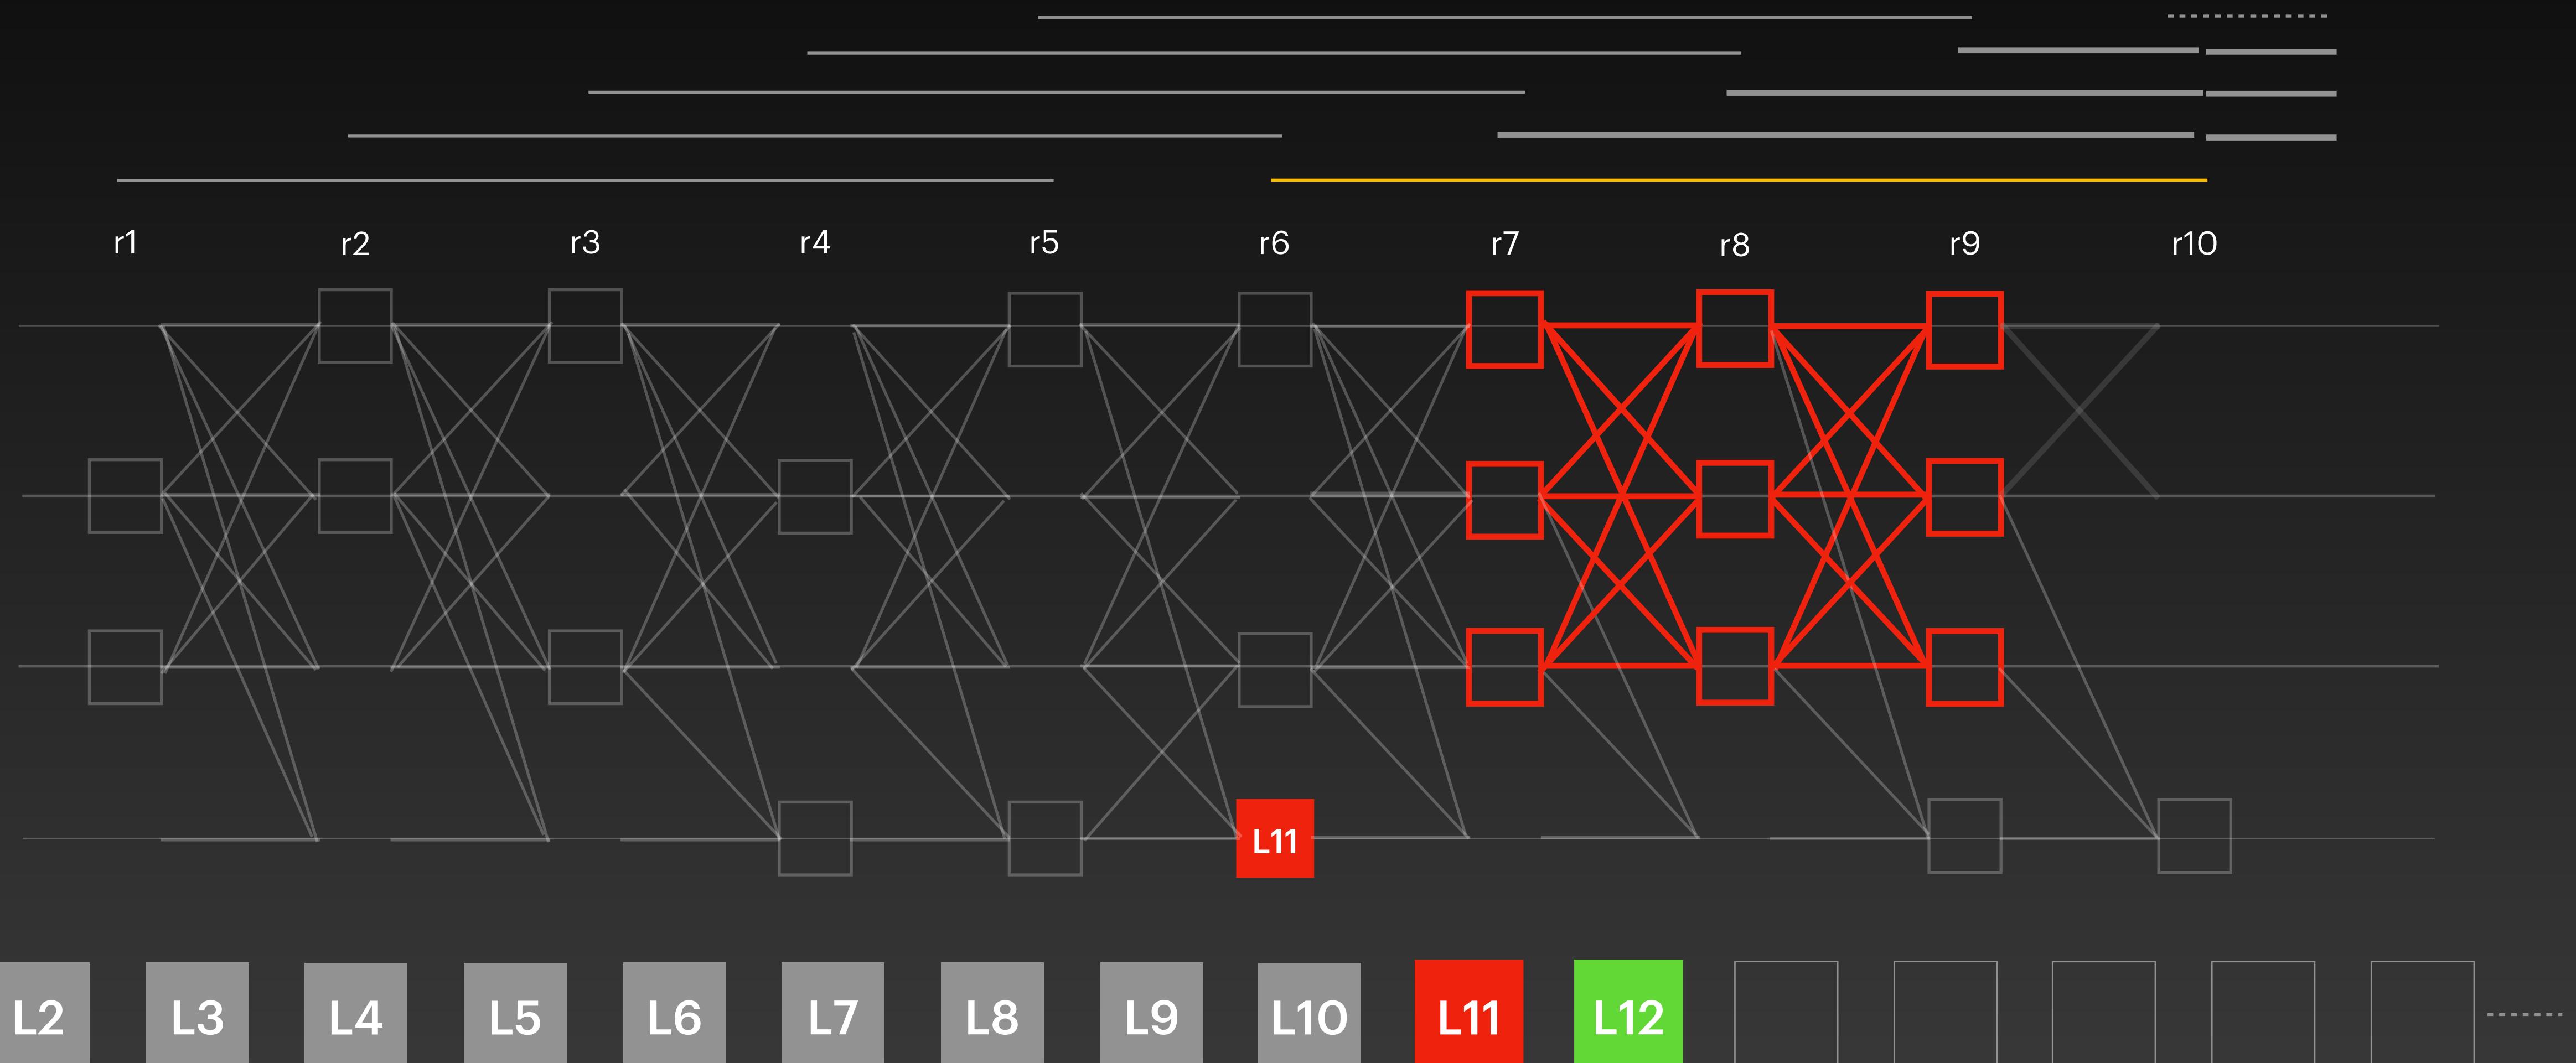
L10

L11

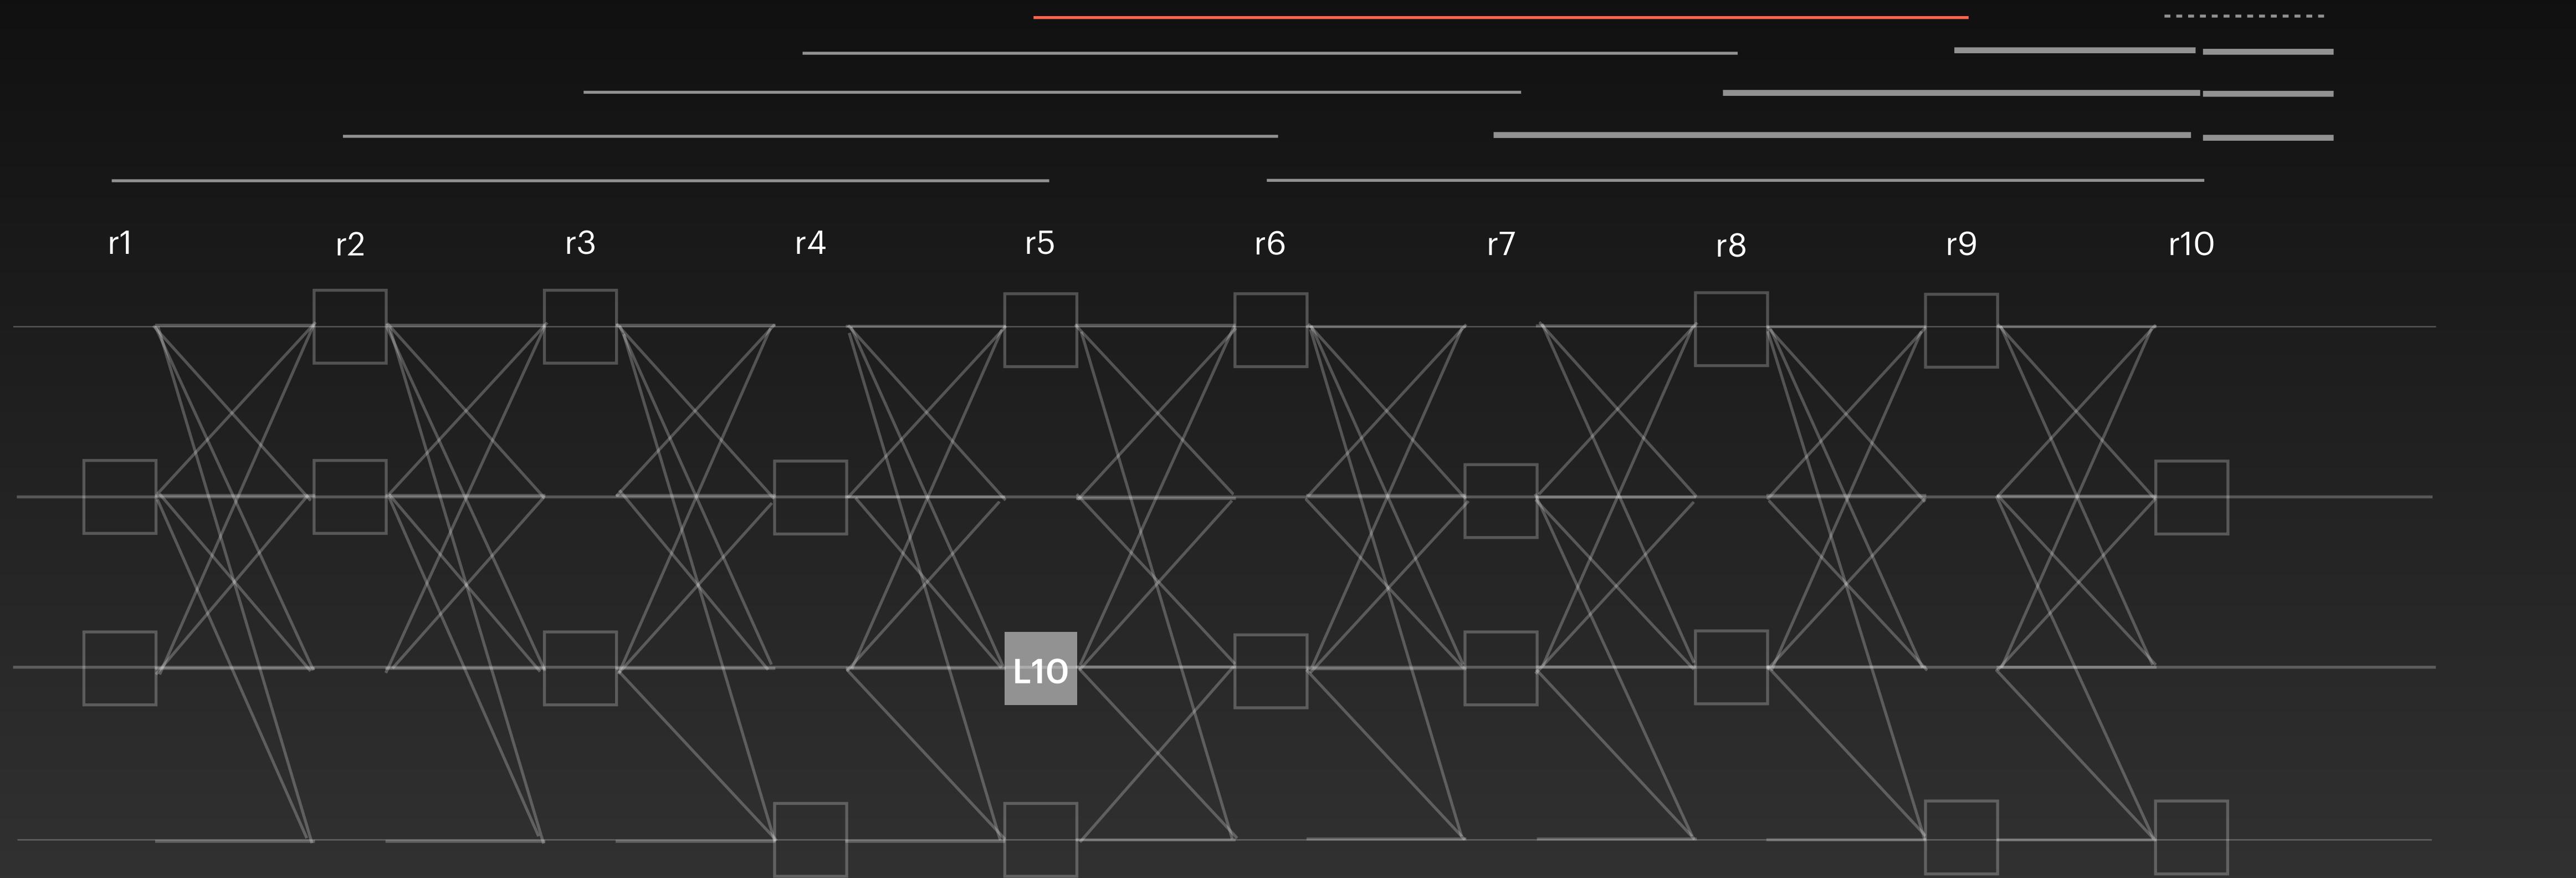
L12

...

Apply Direct Rule



Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

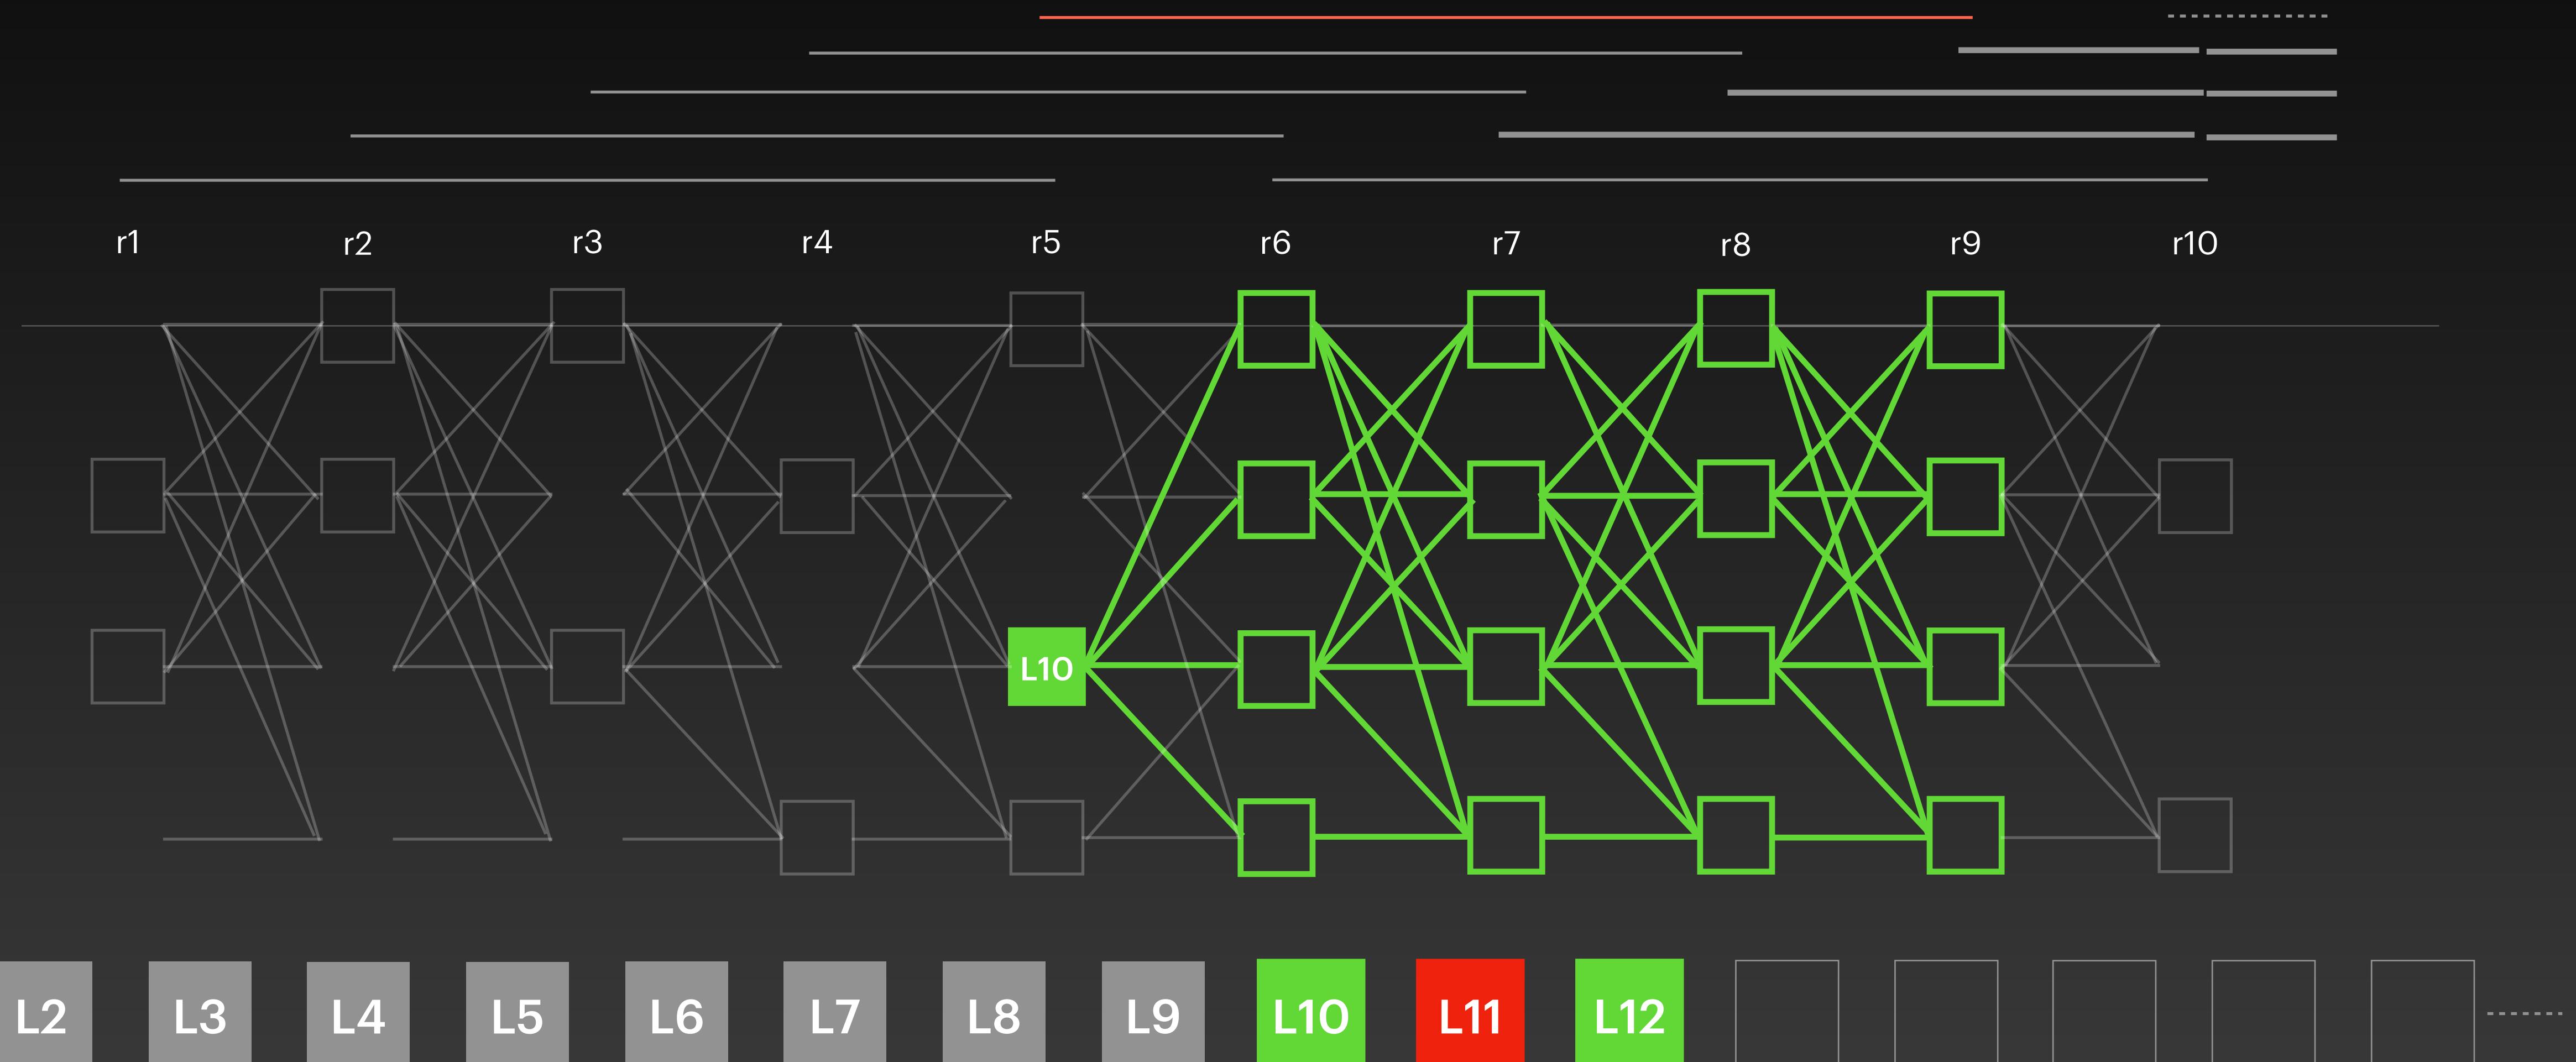
L10

L11

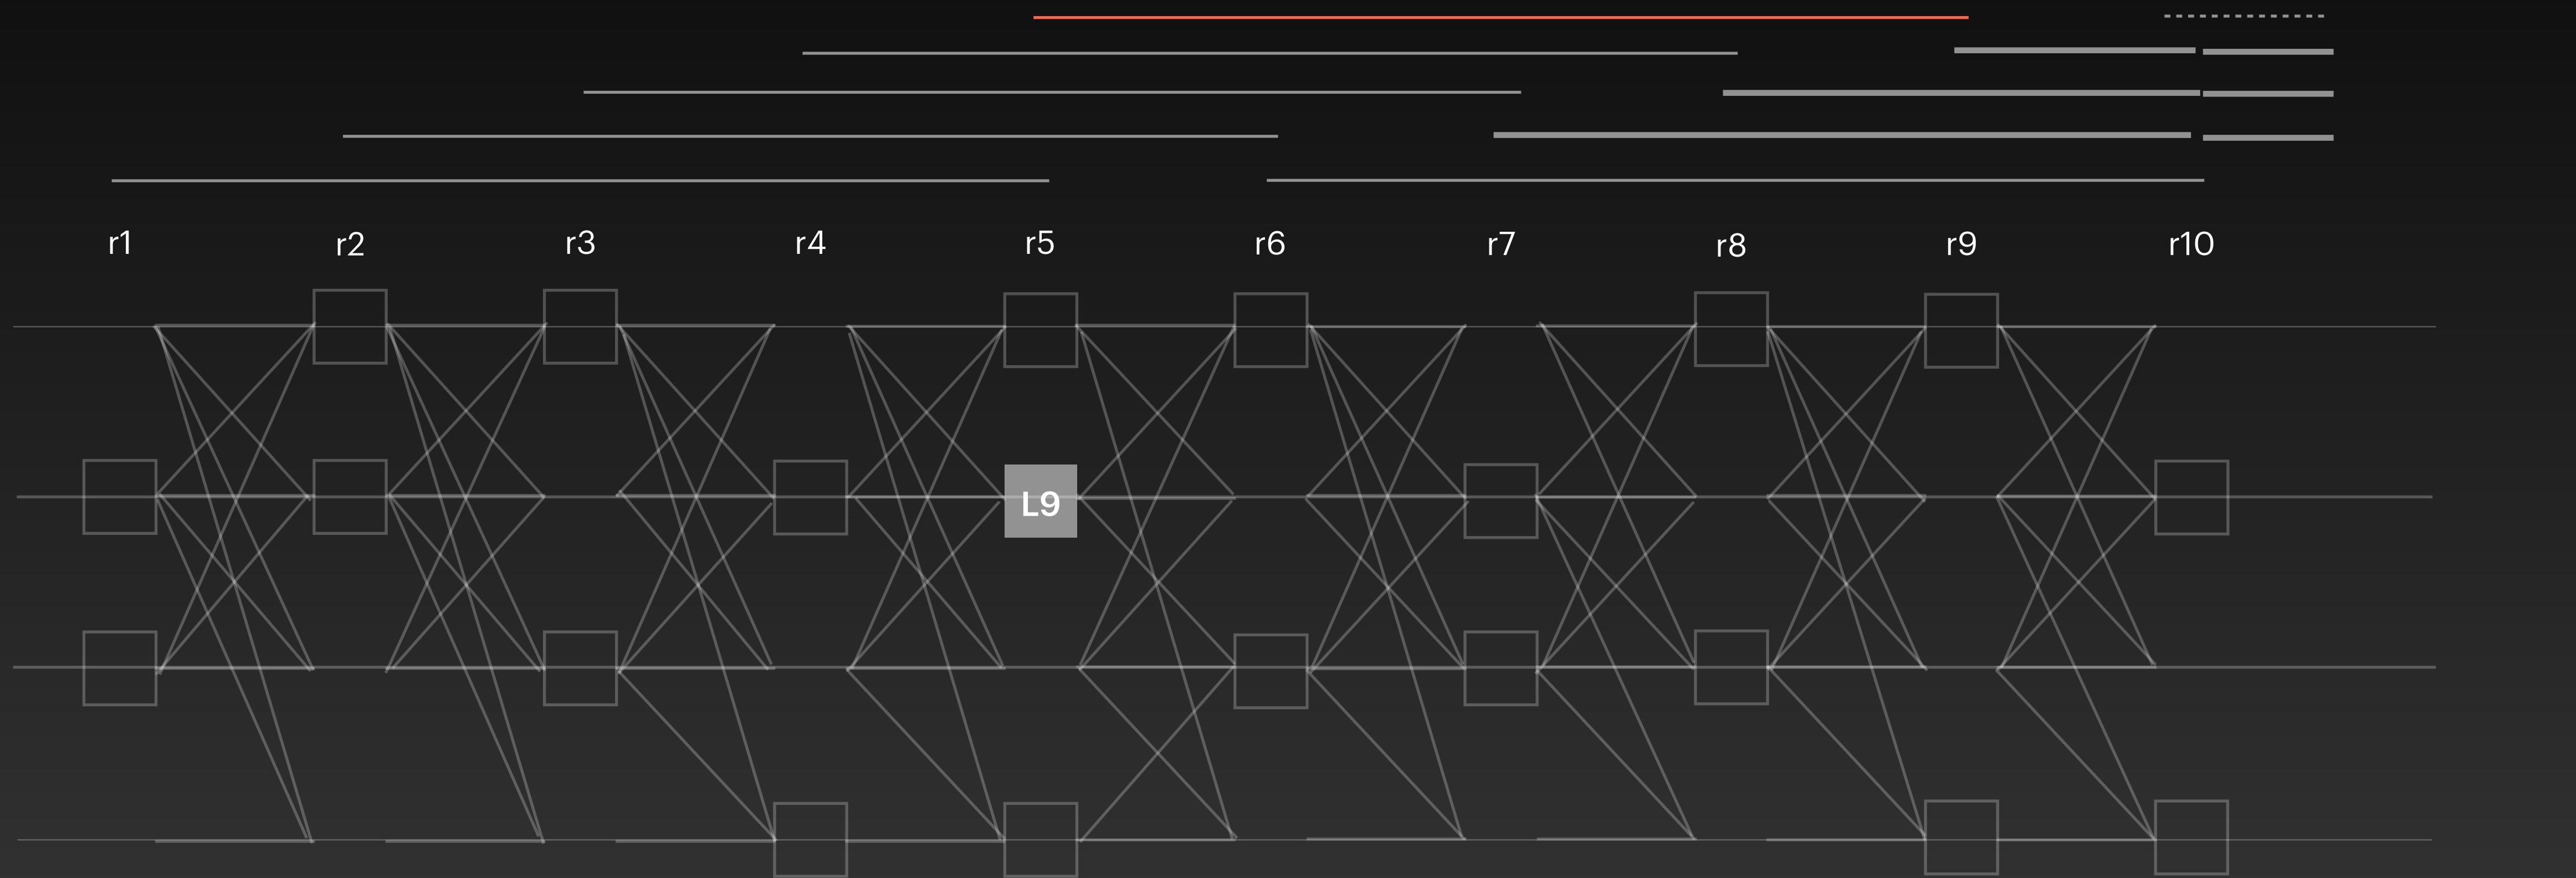
L12

.....

Apply Direct Rule



Apply Direct Rule



L_1

L_2

L_3

L_4

L_5

L_6

L_7

L_8

L_9

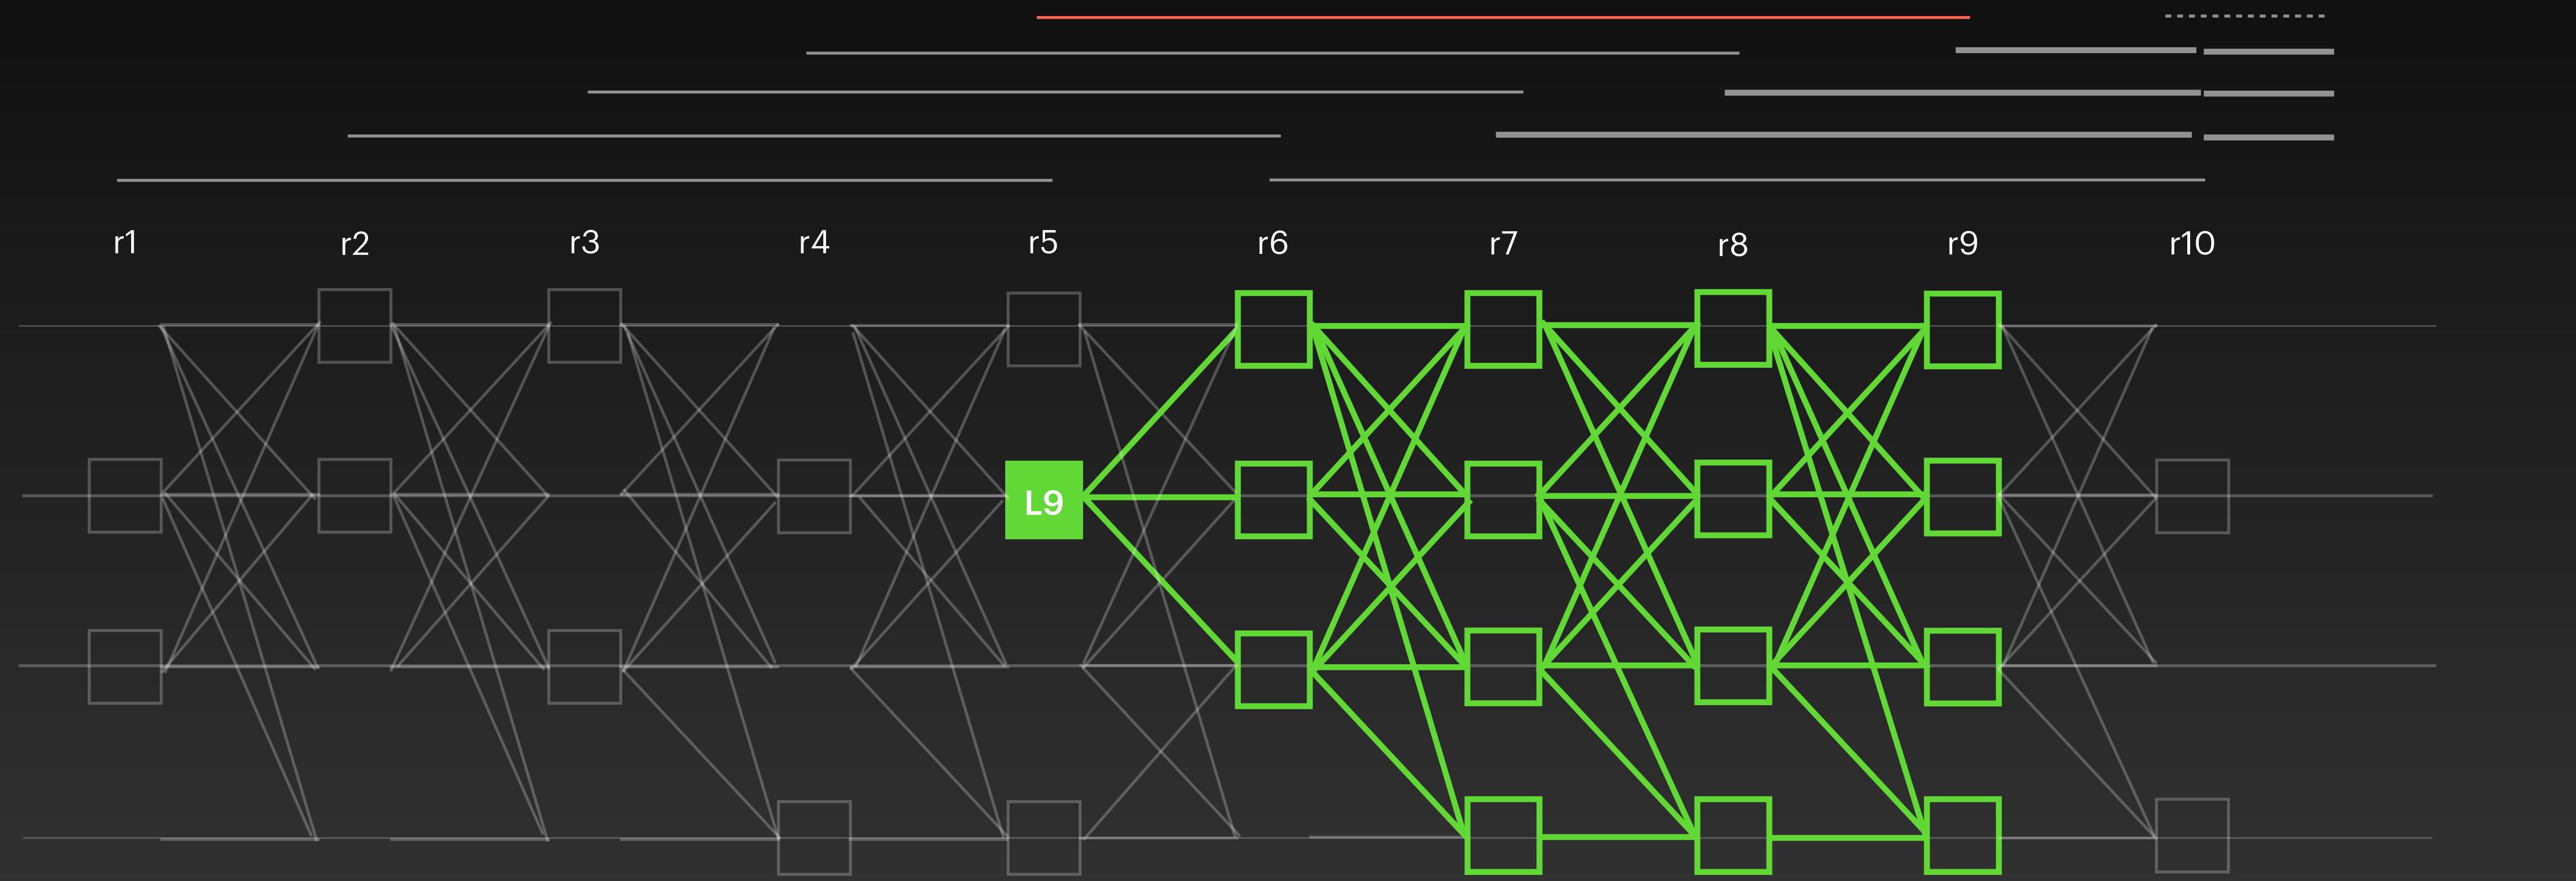
L_{10}

L_{11}

L_{12}

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

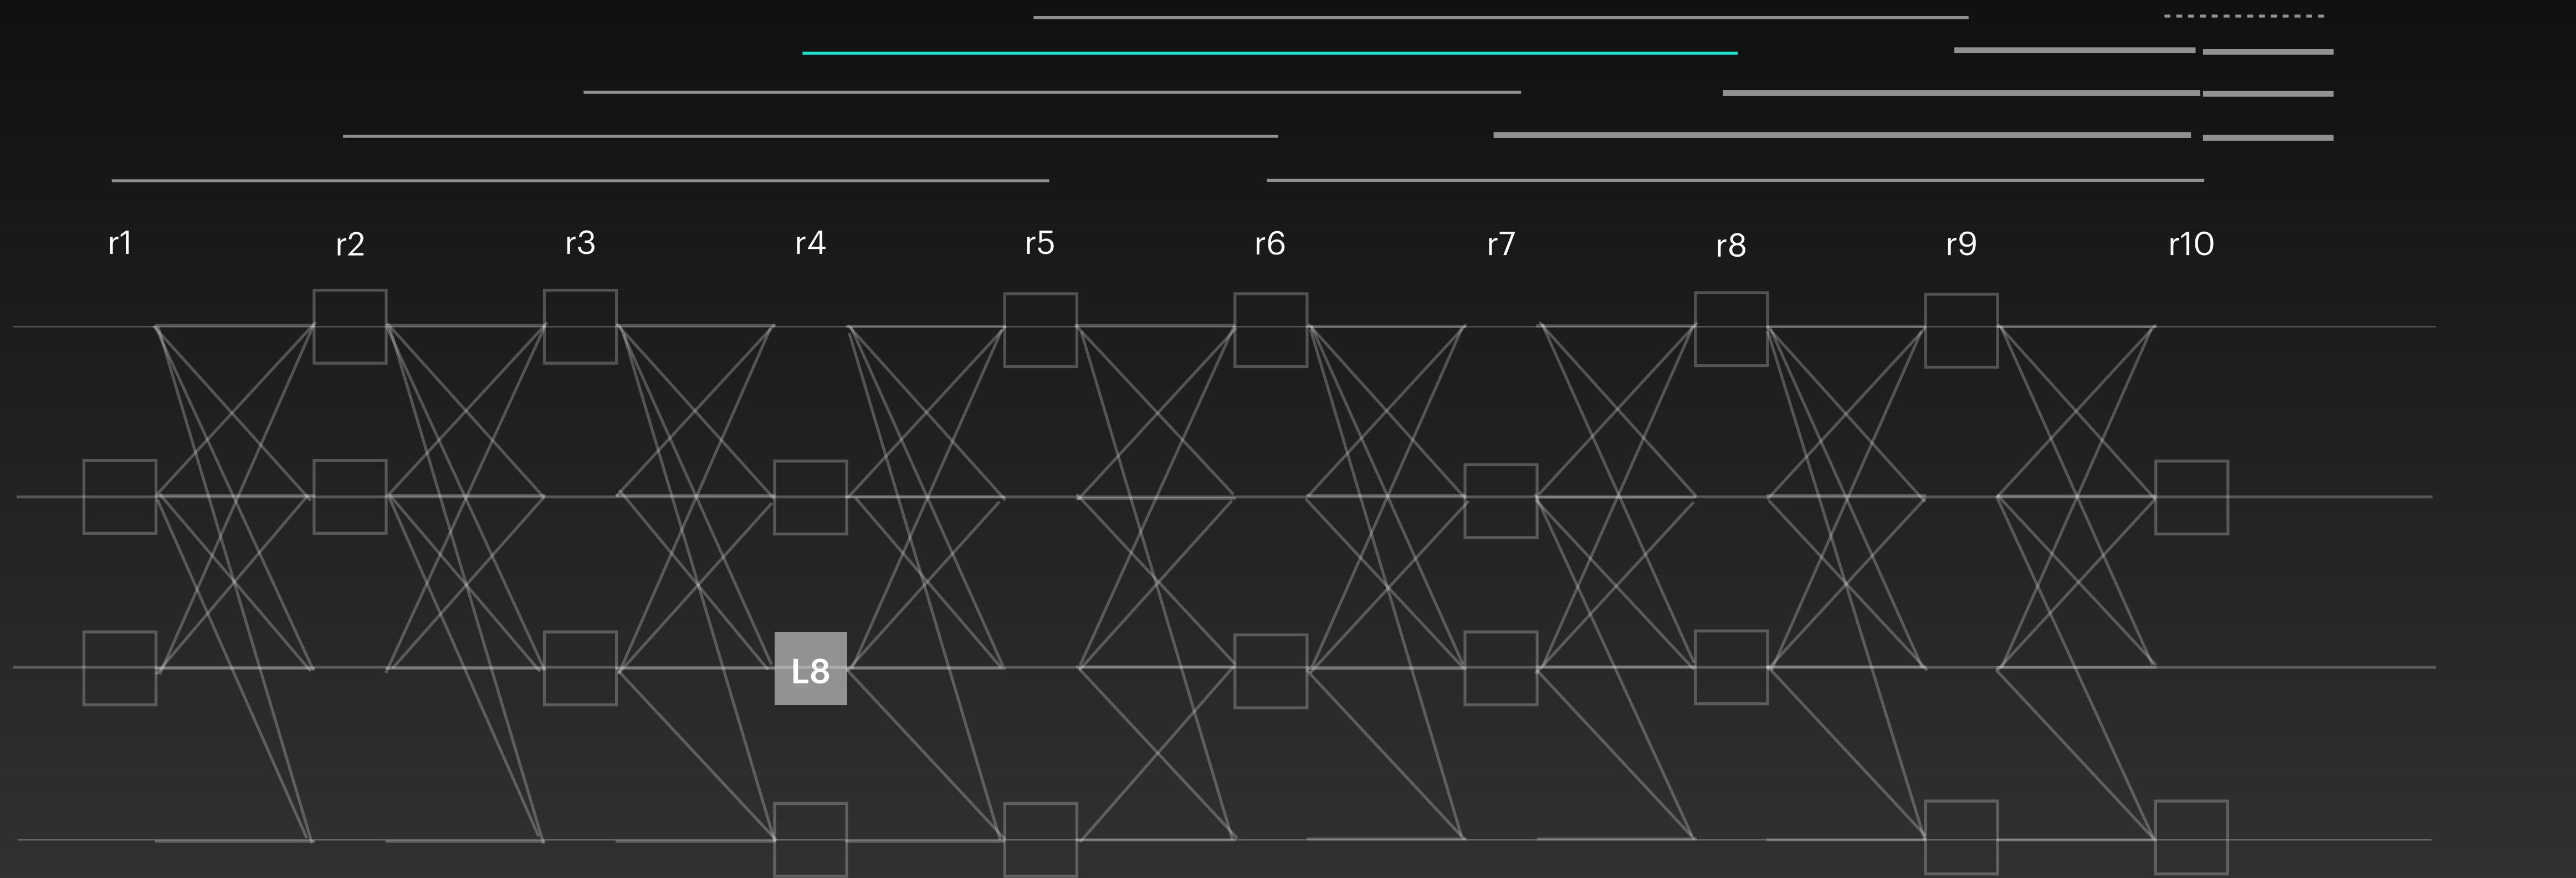
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

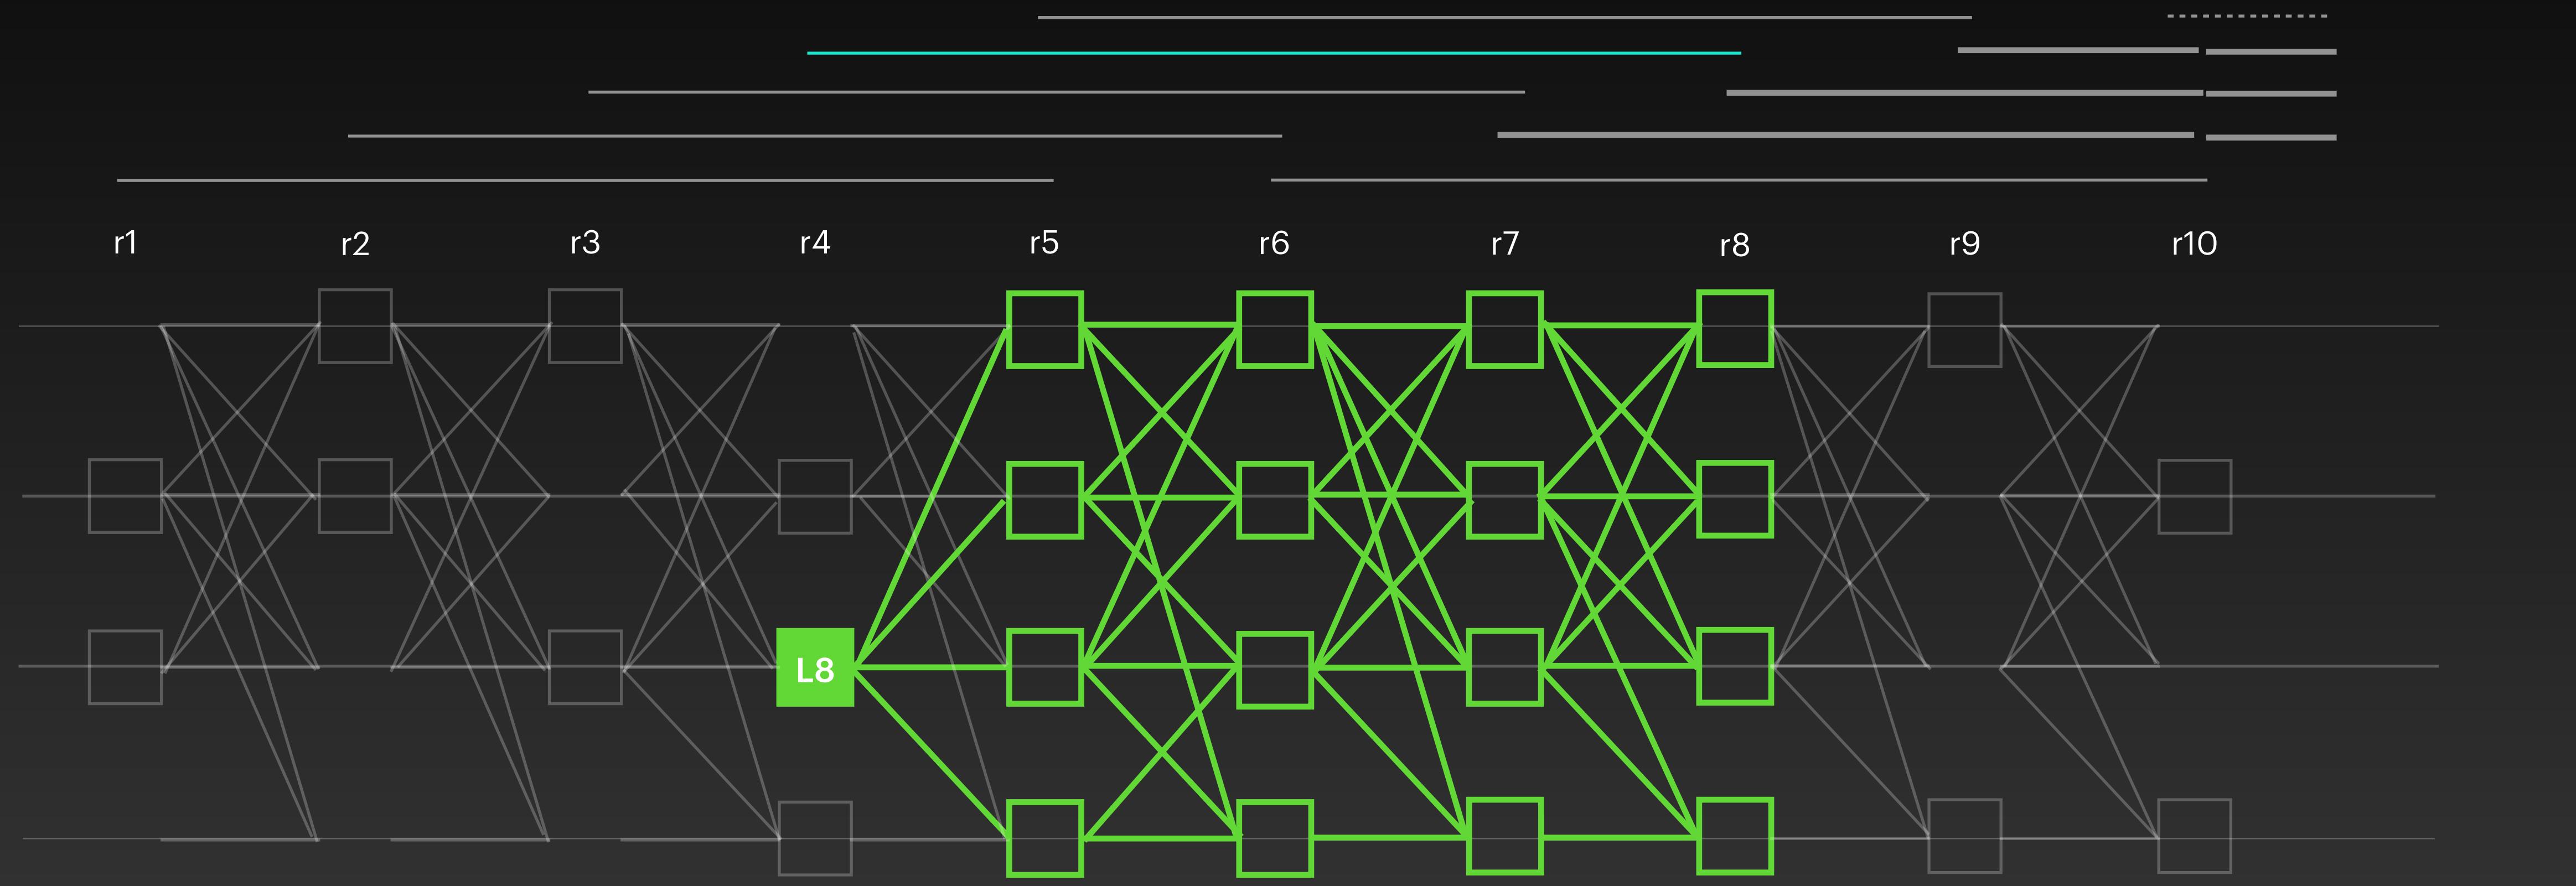
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

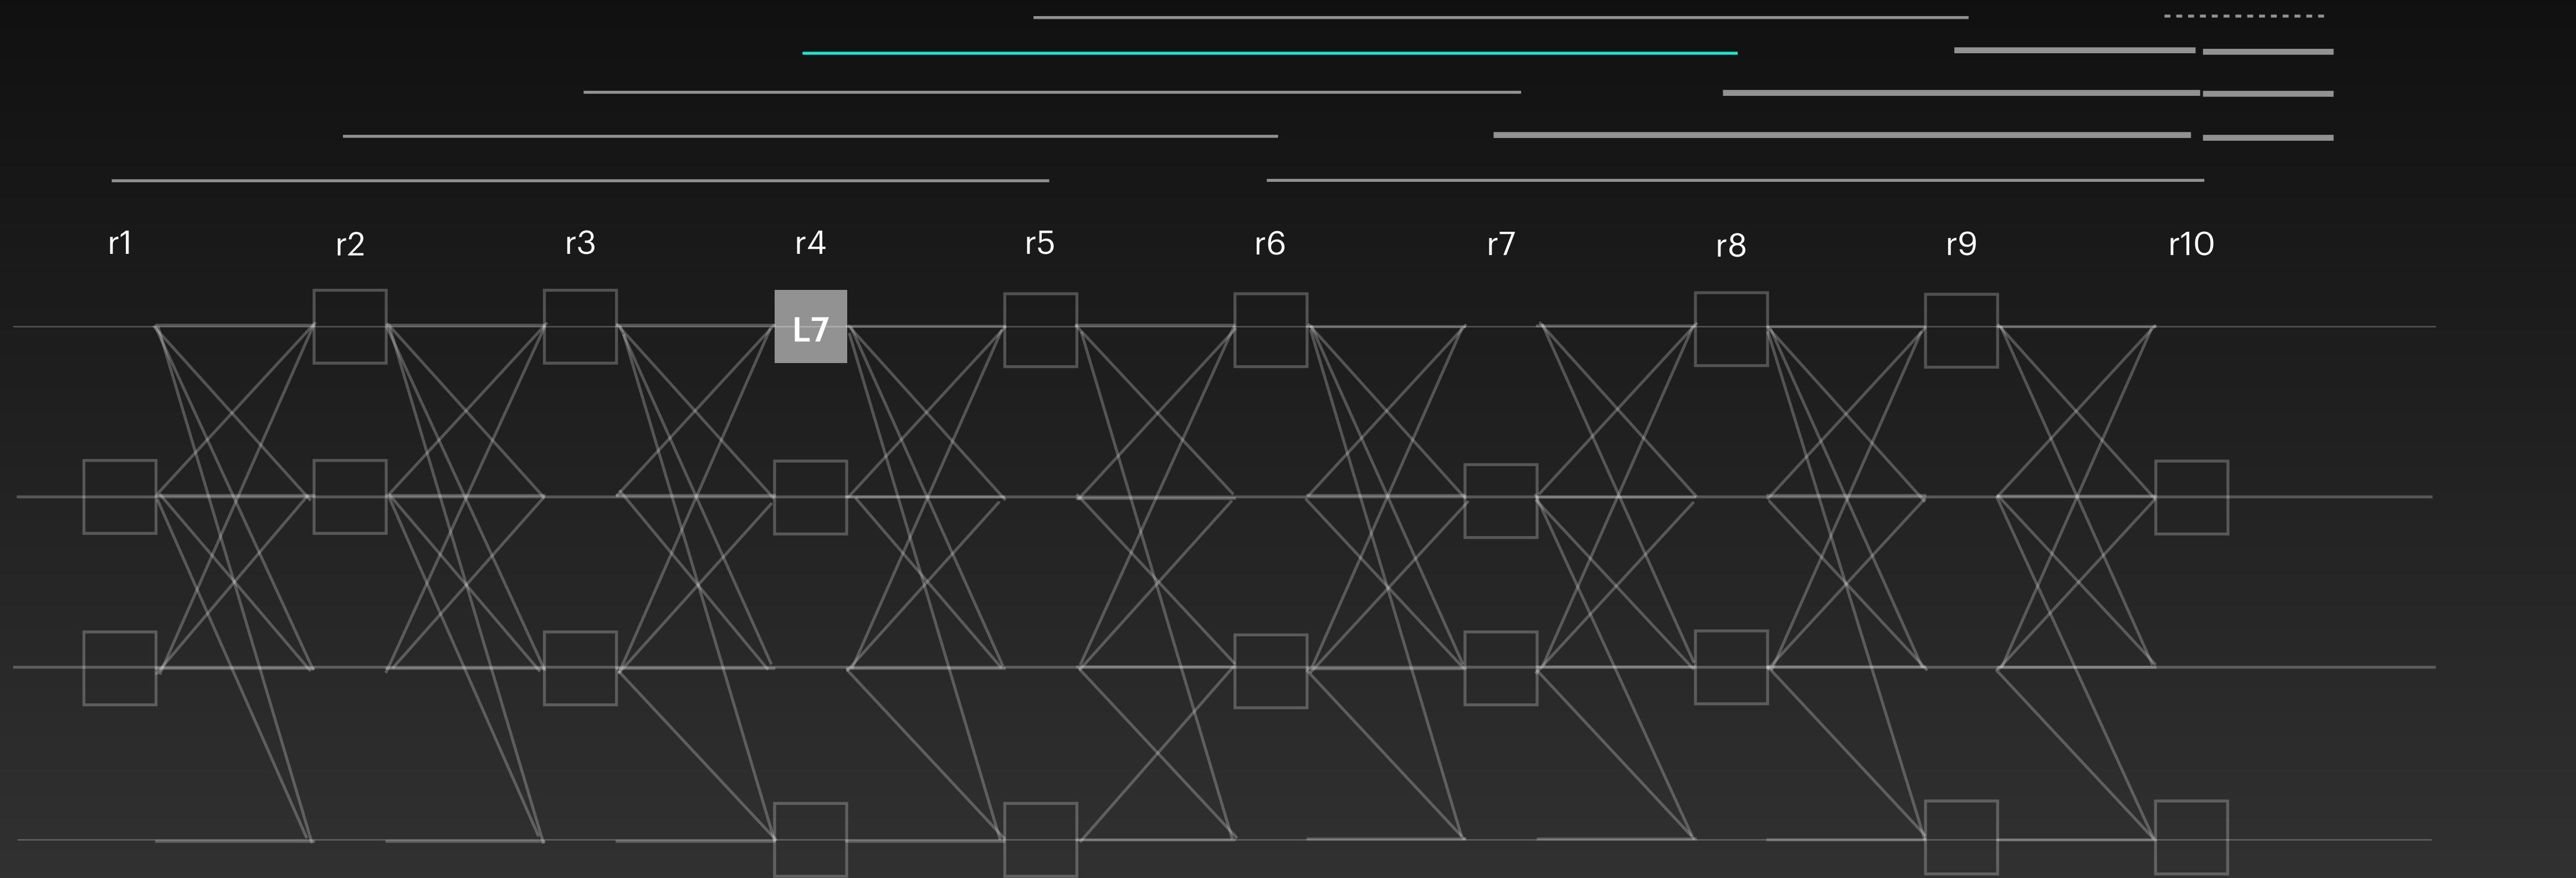
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

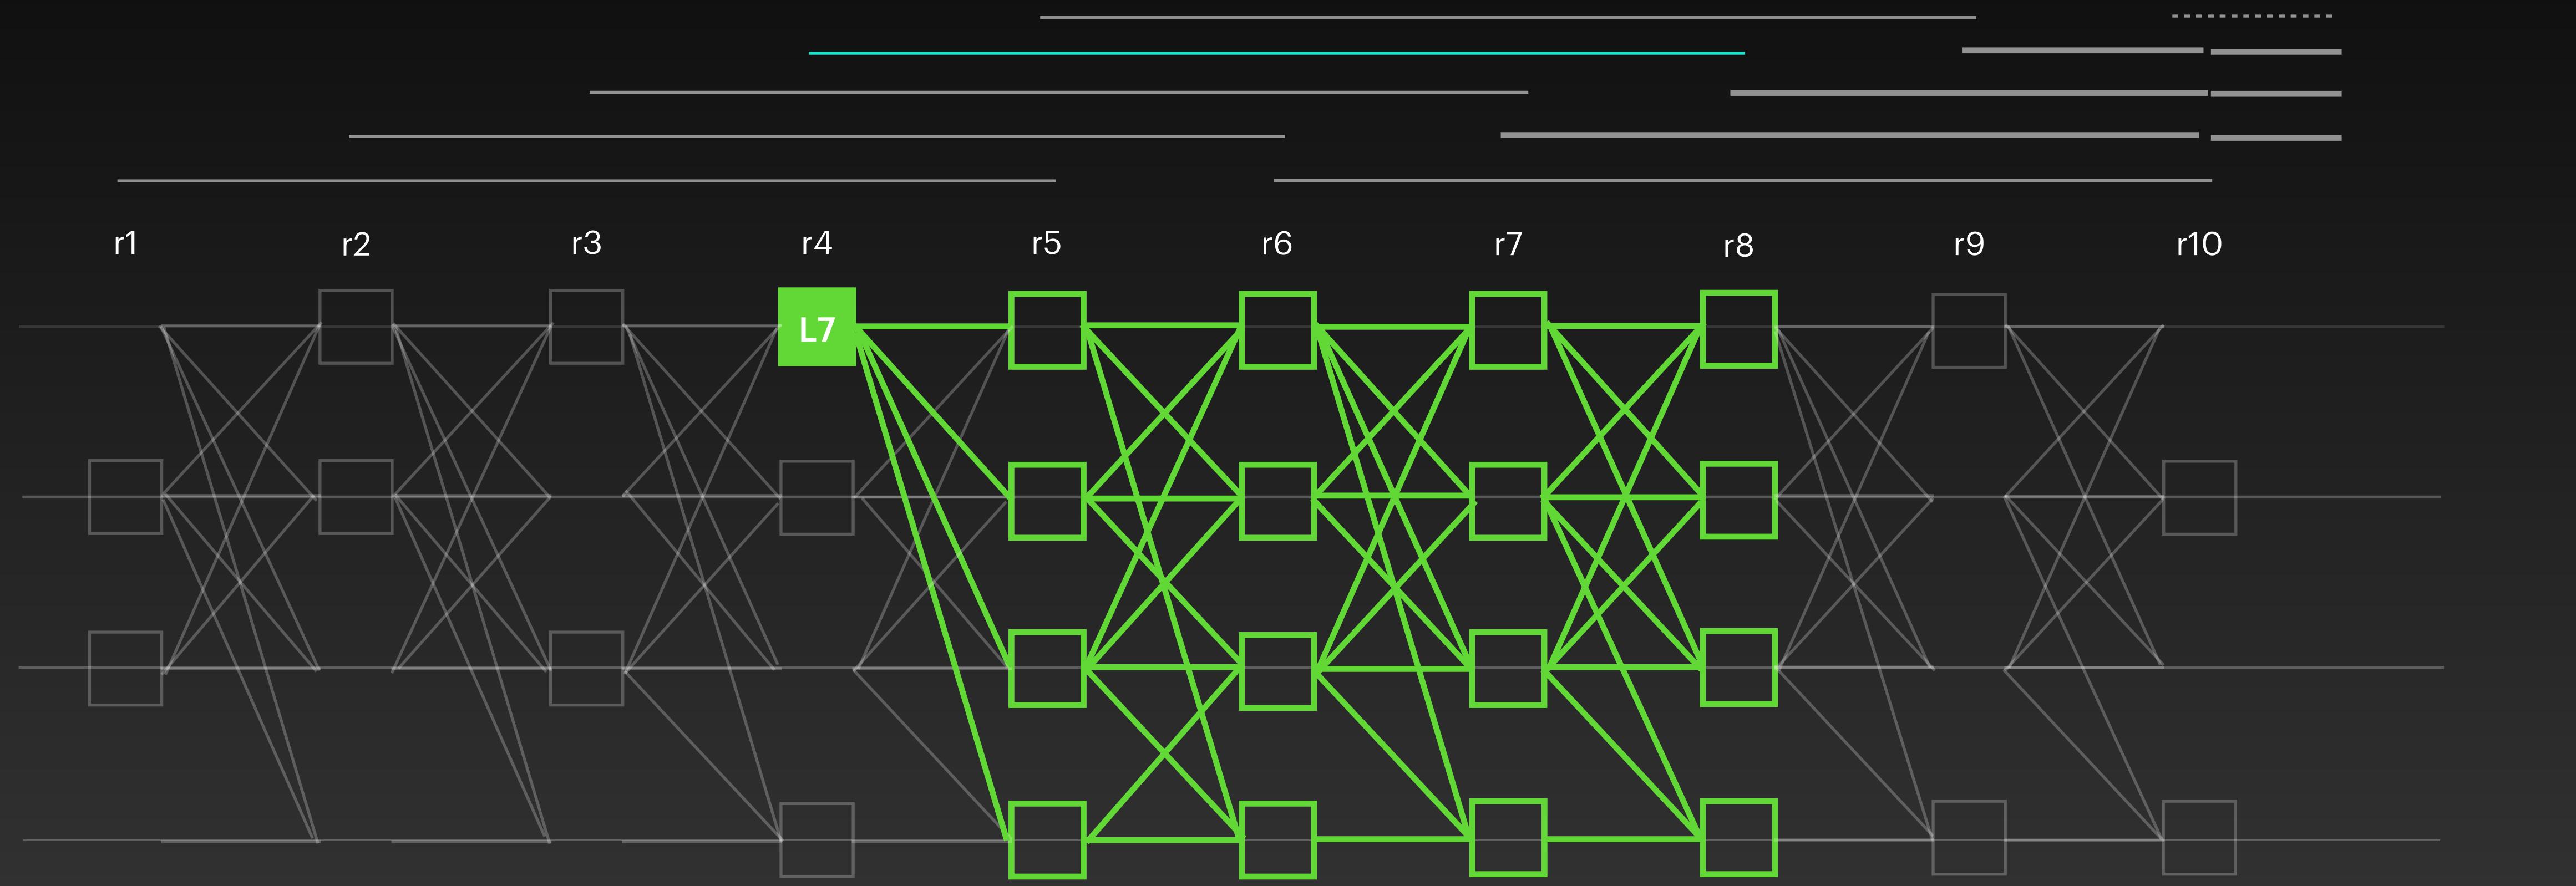
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

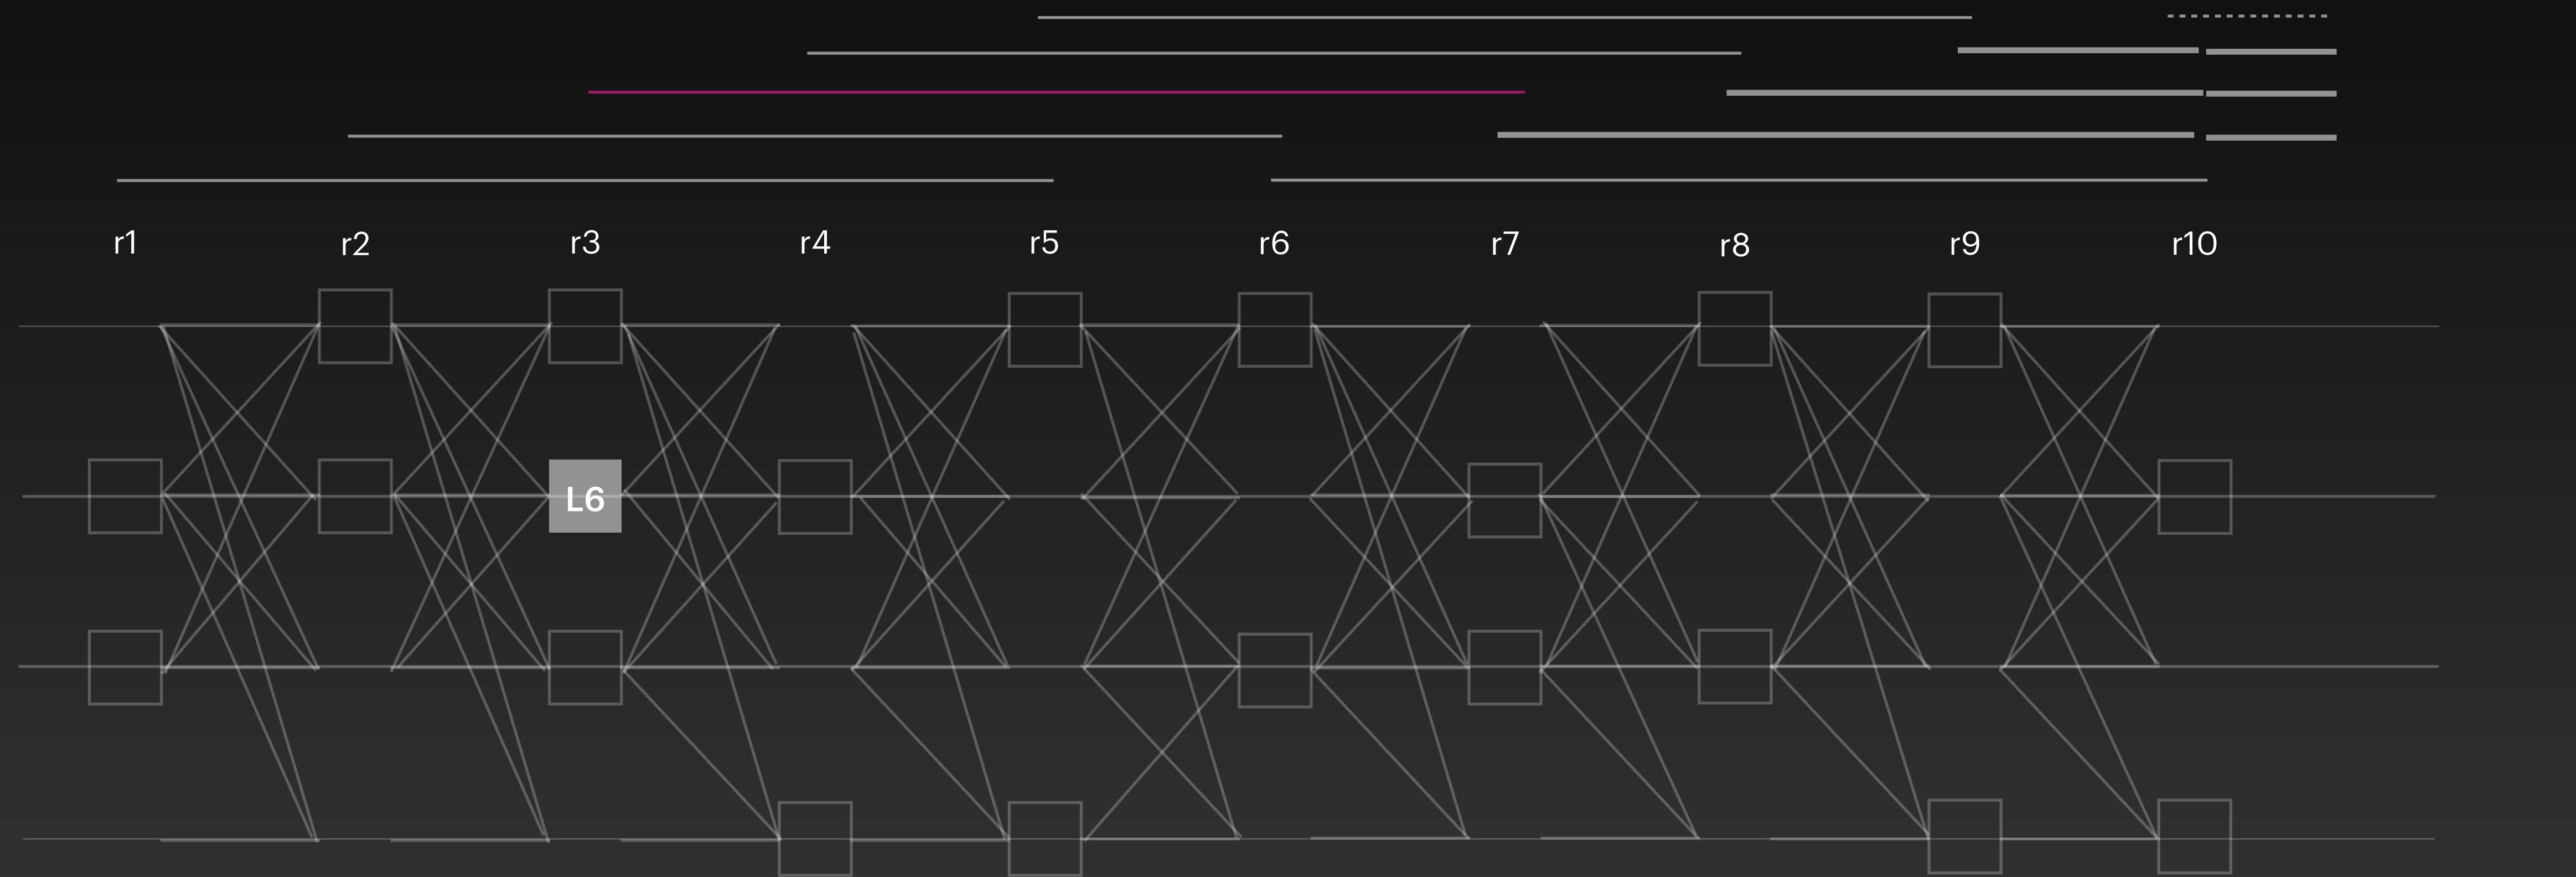
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

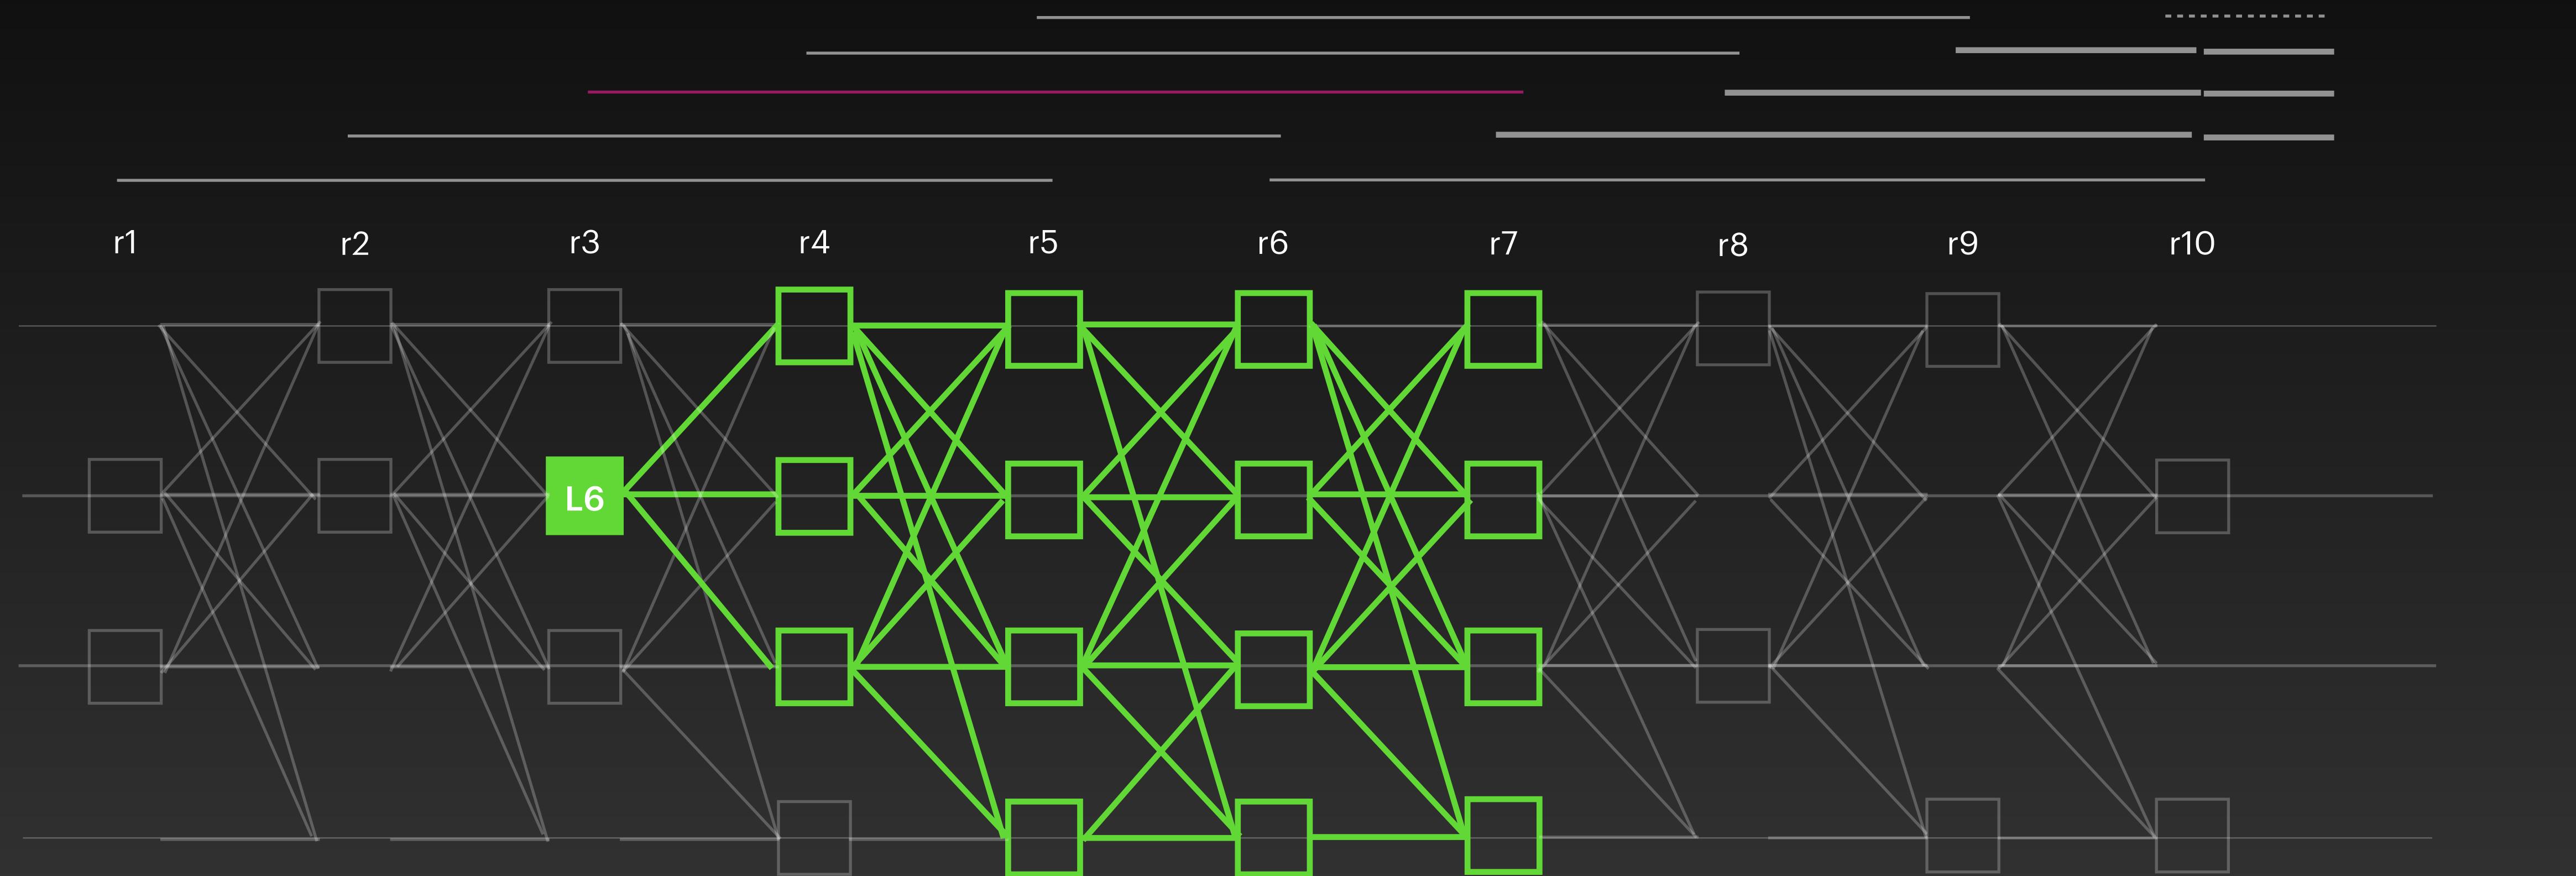
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

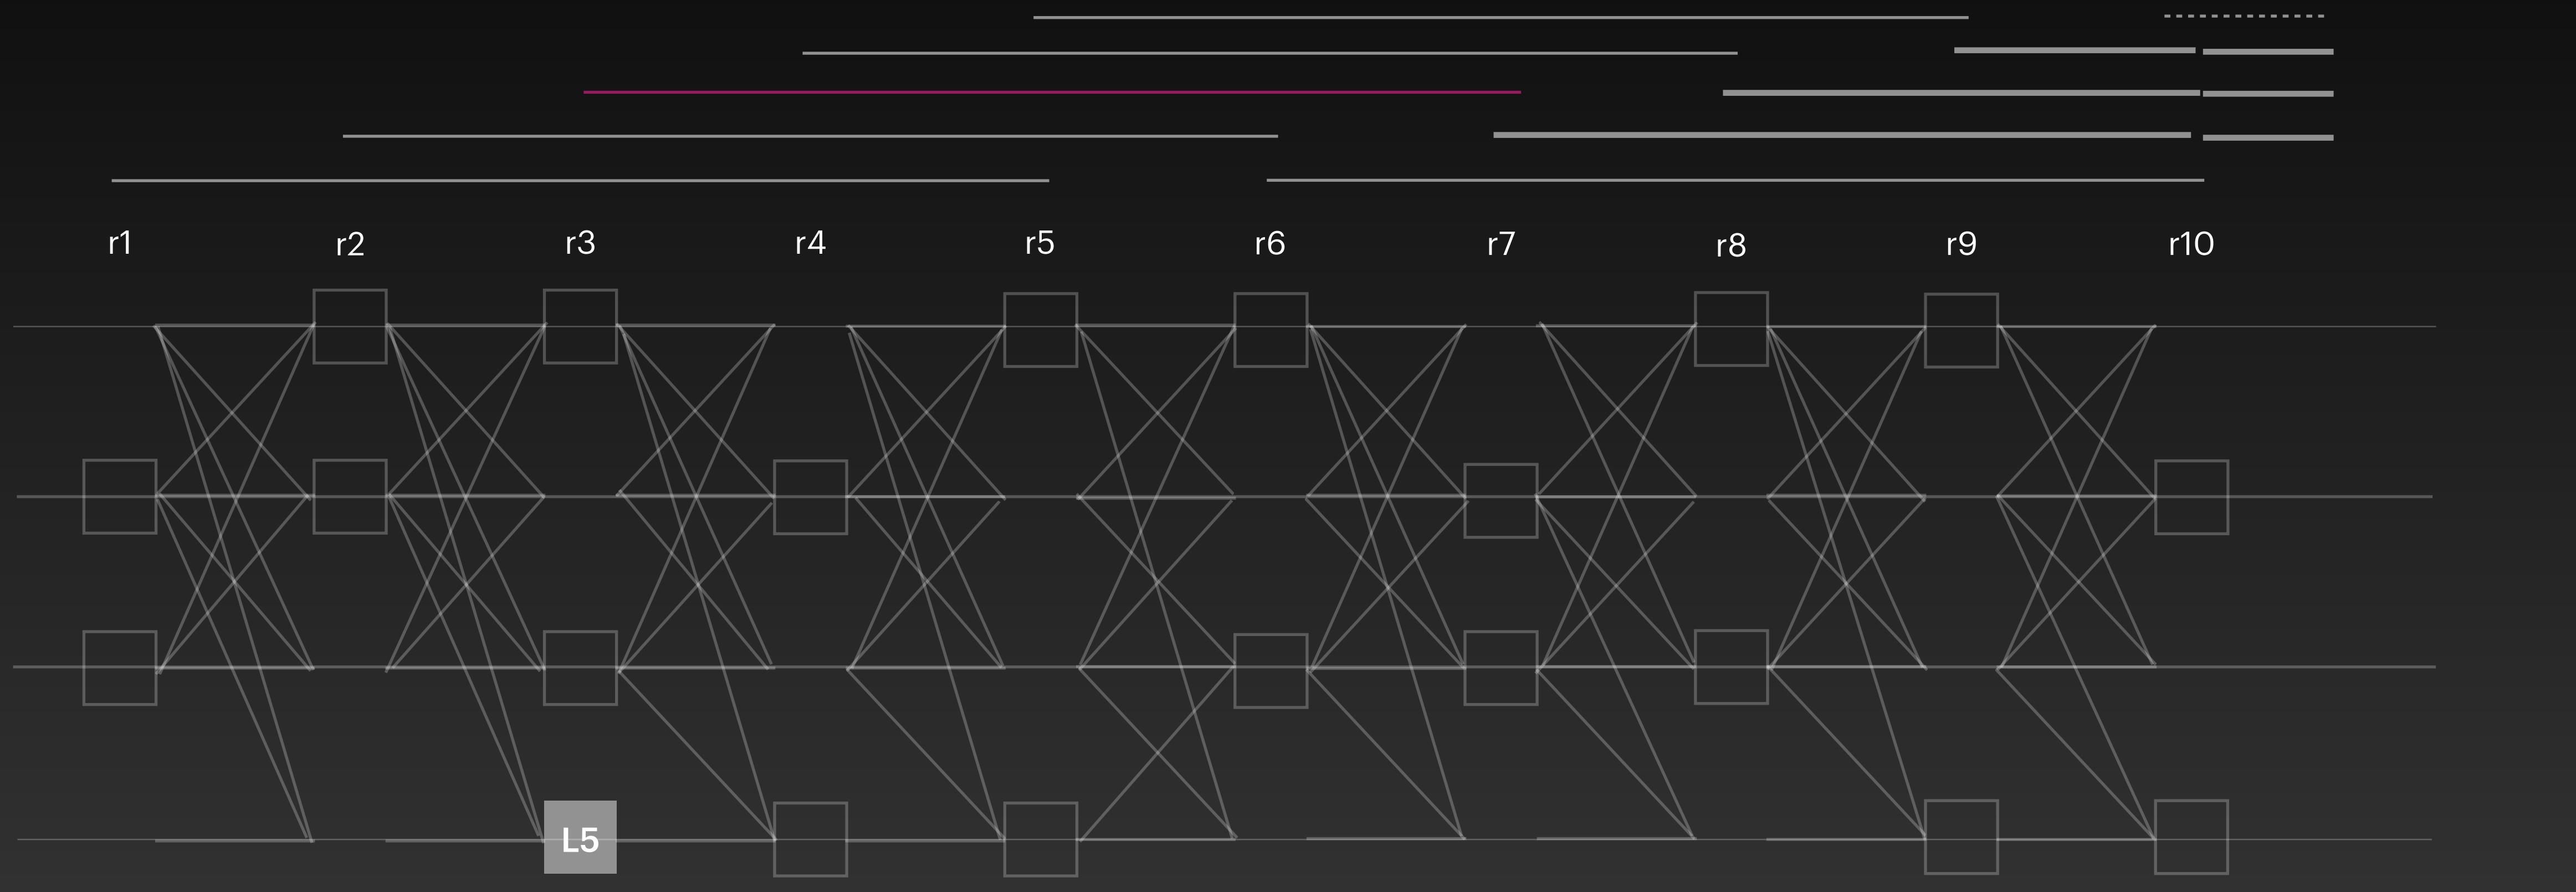
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

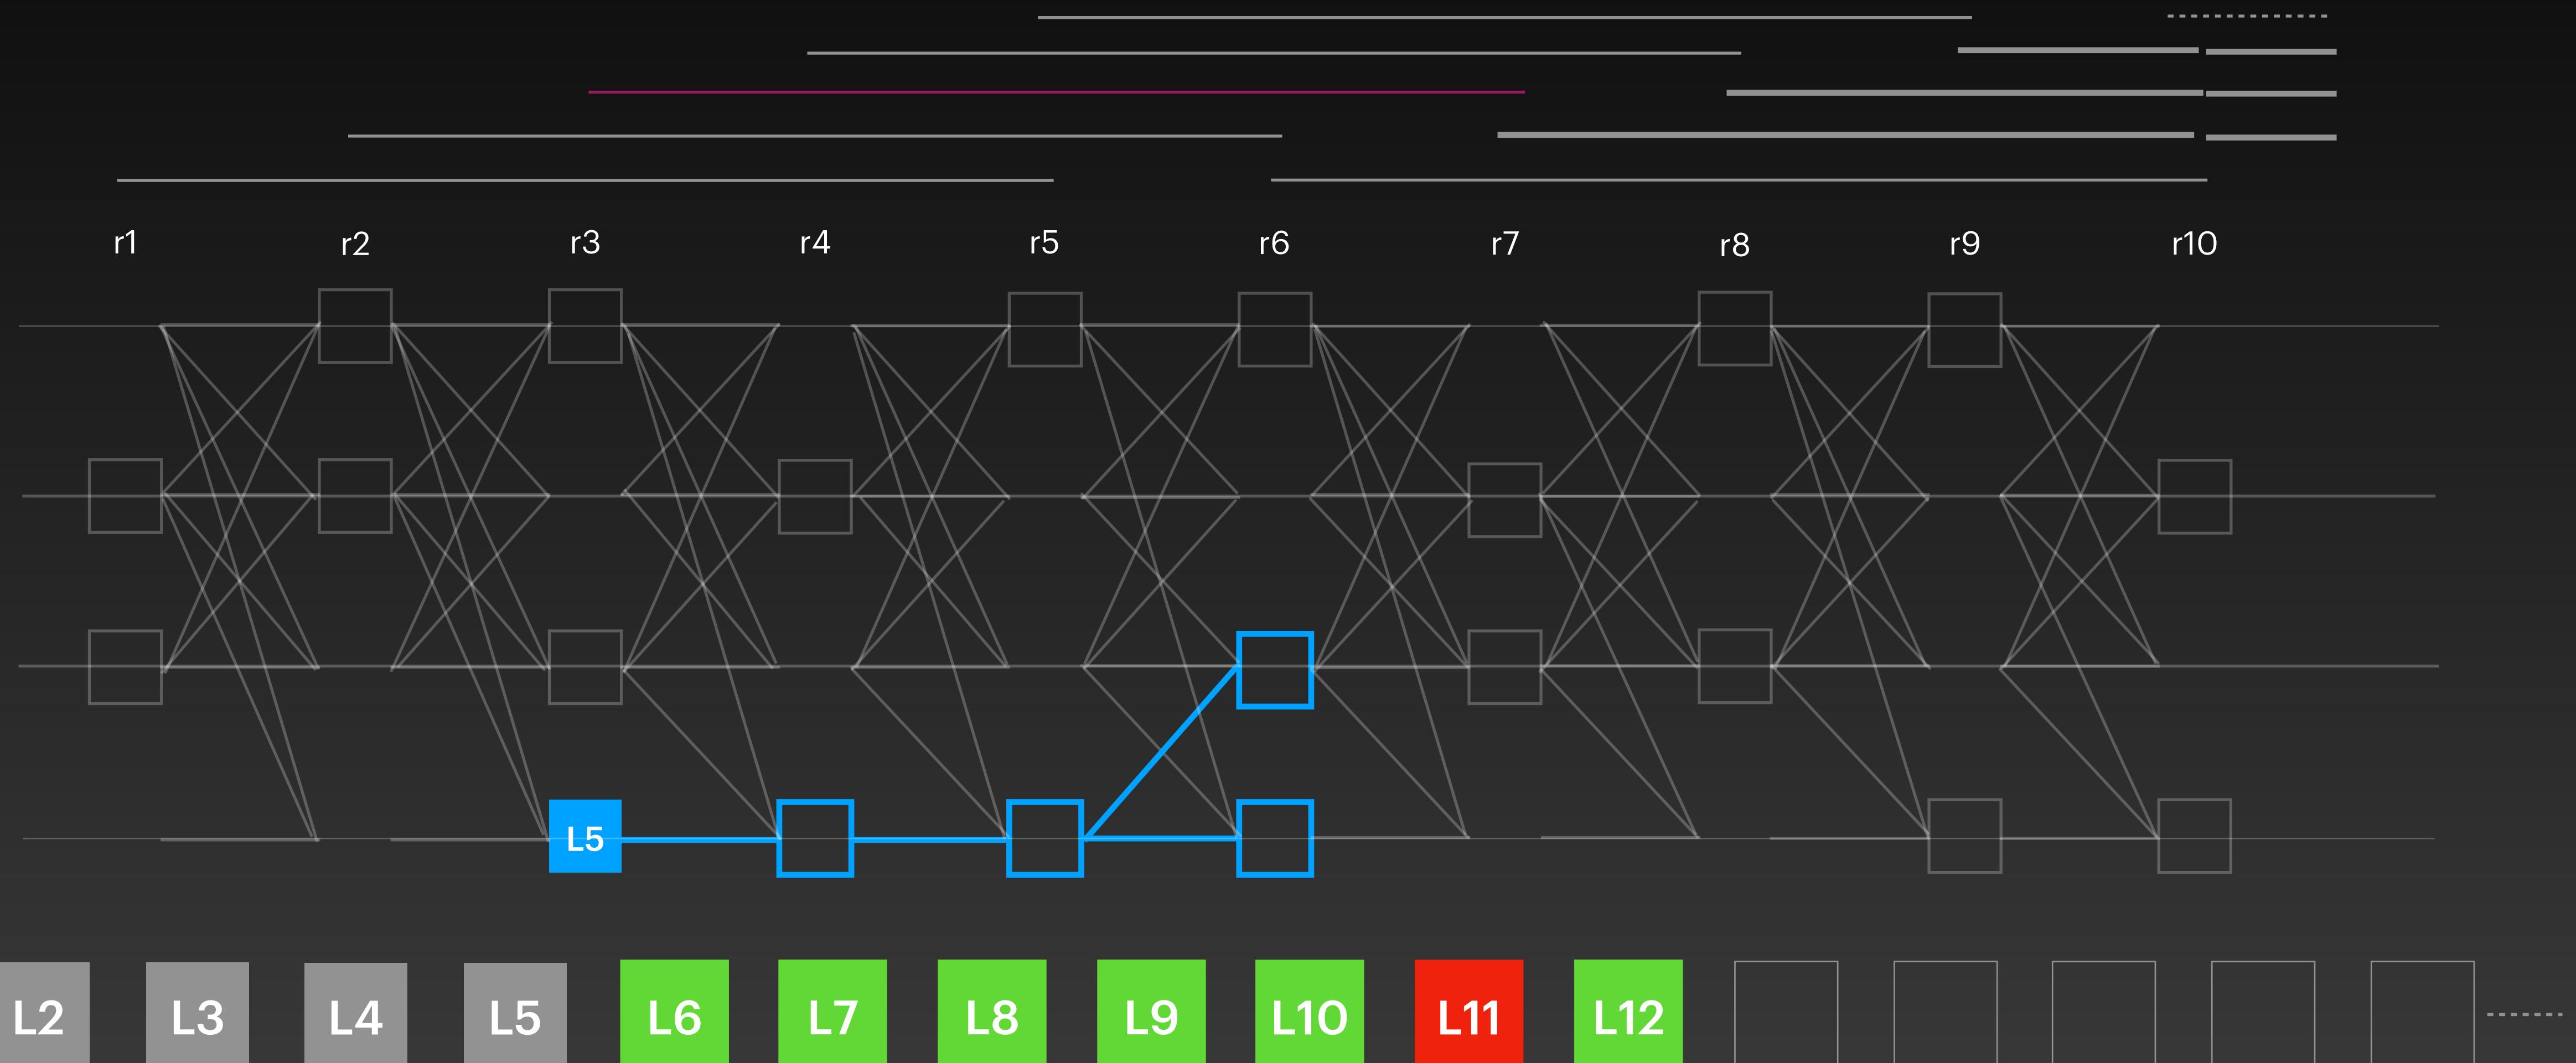
L10

L11

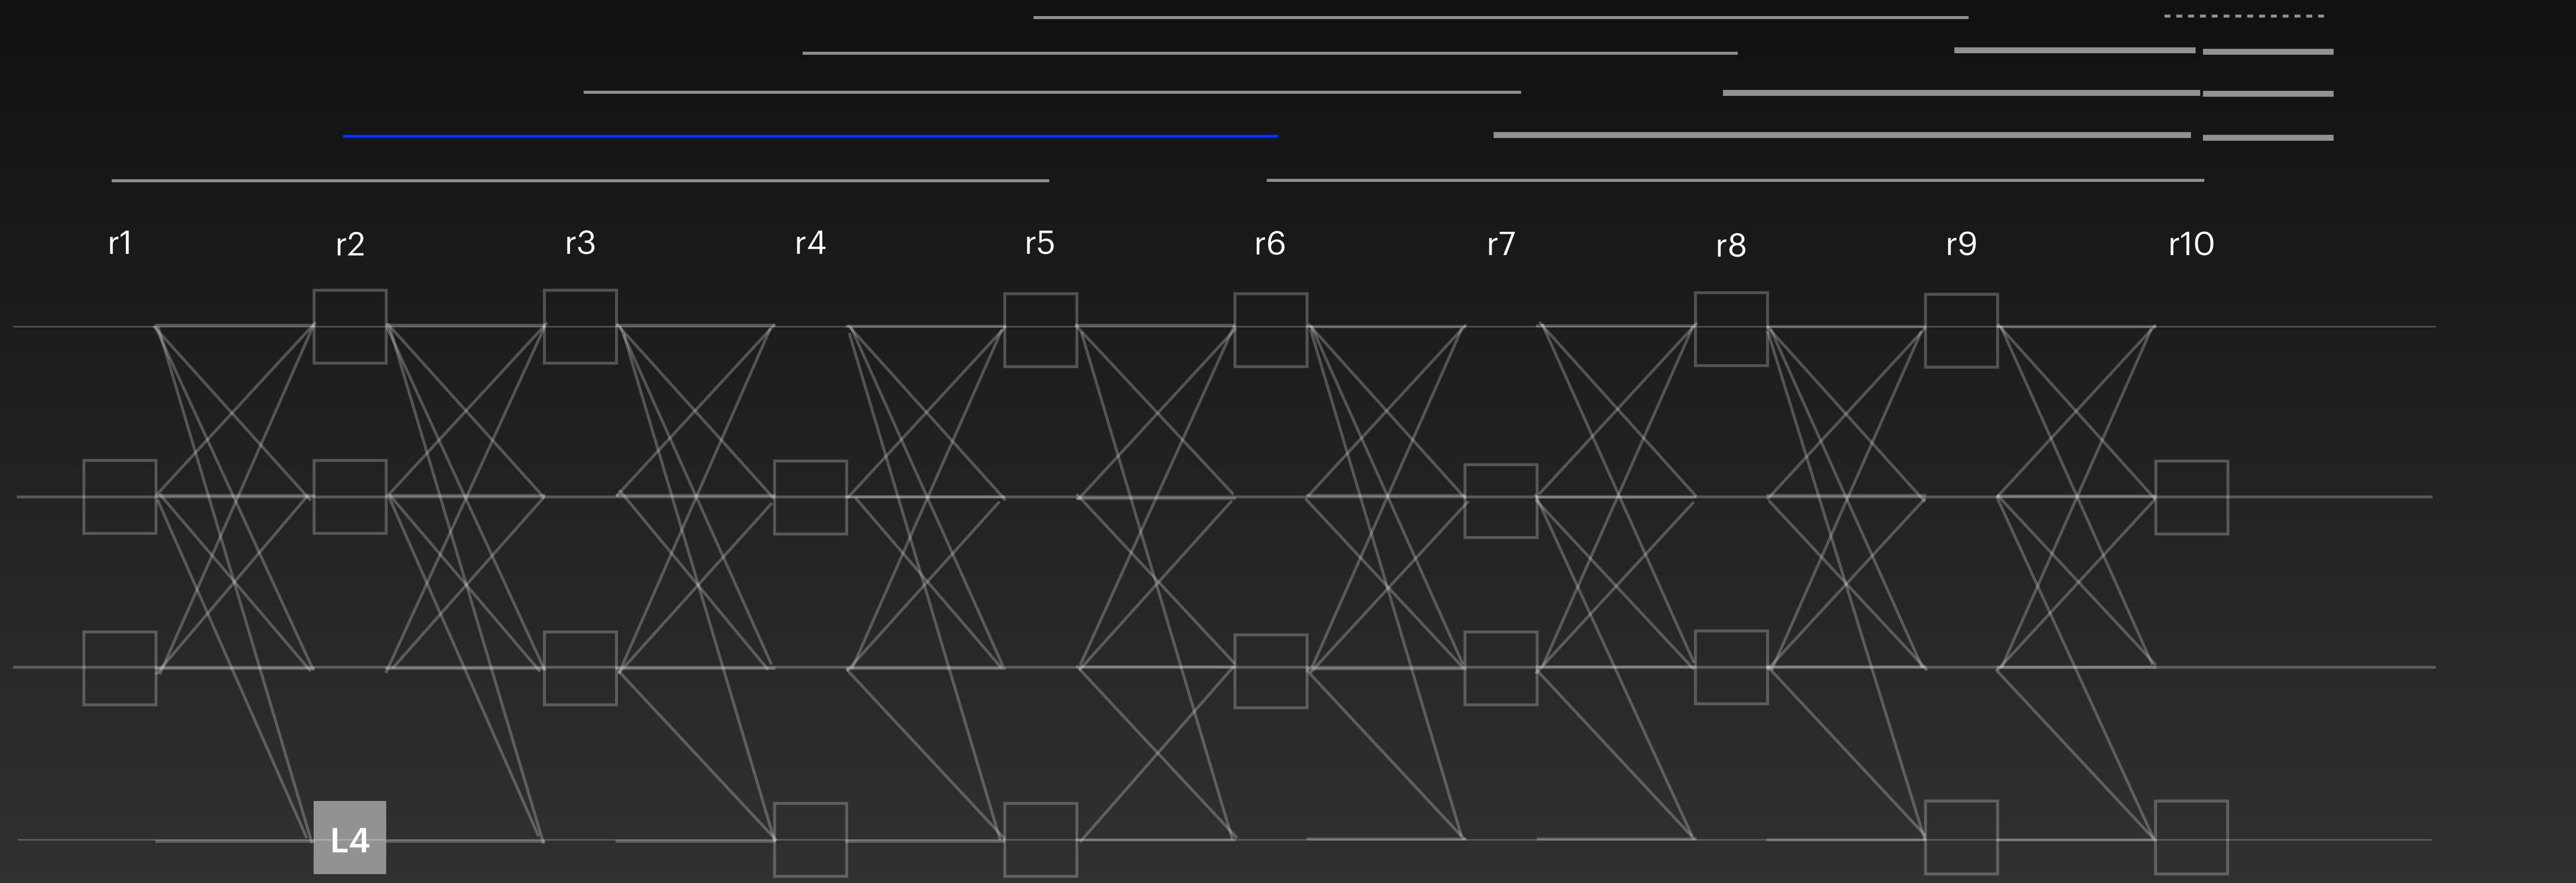
L12

.....

Apply Direct Rule



Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

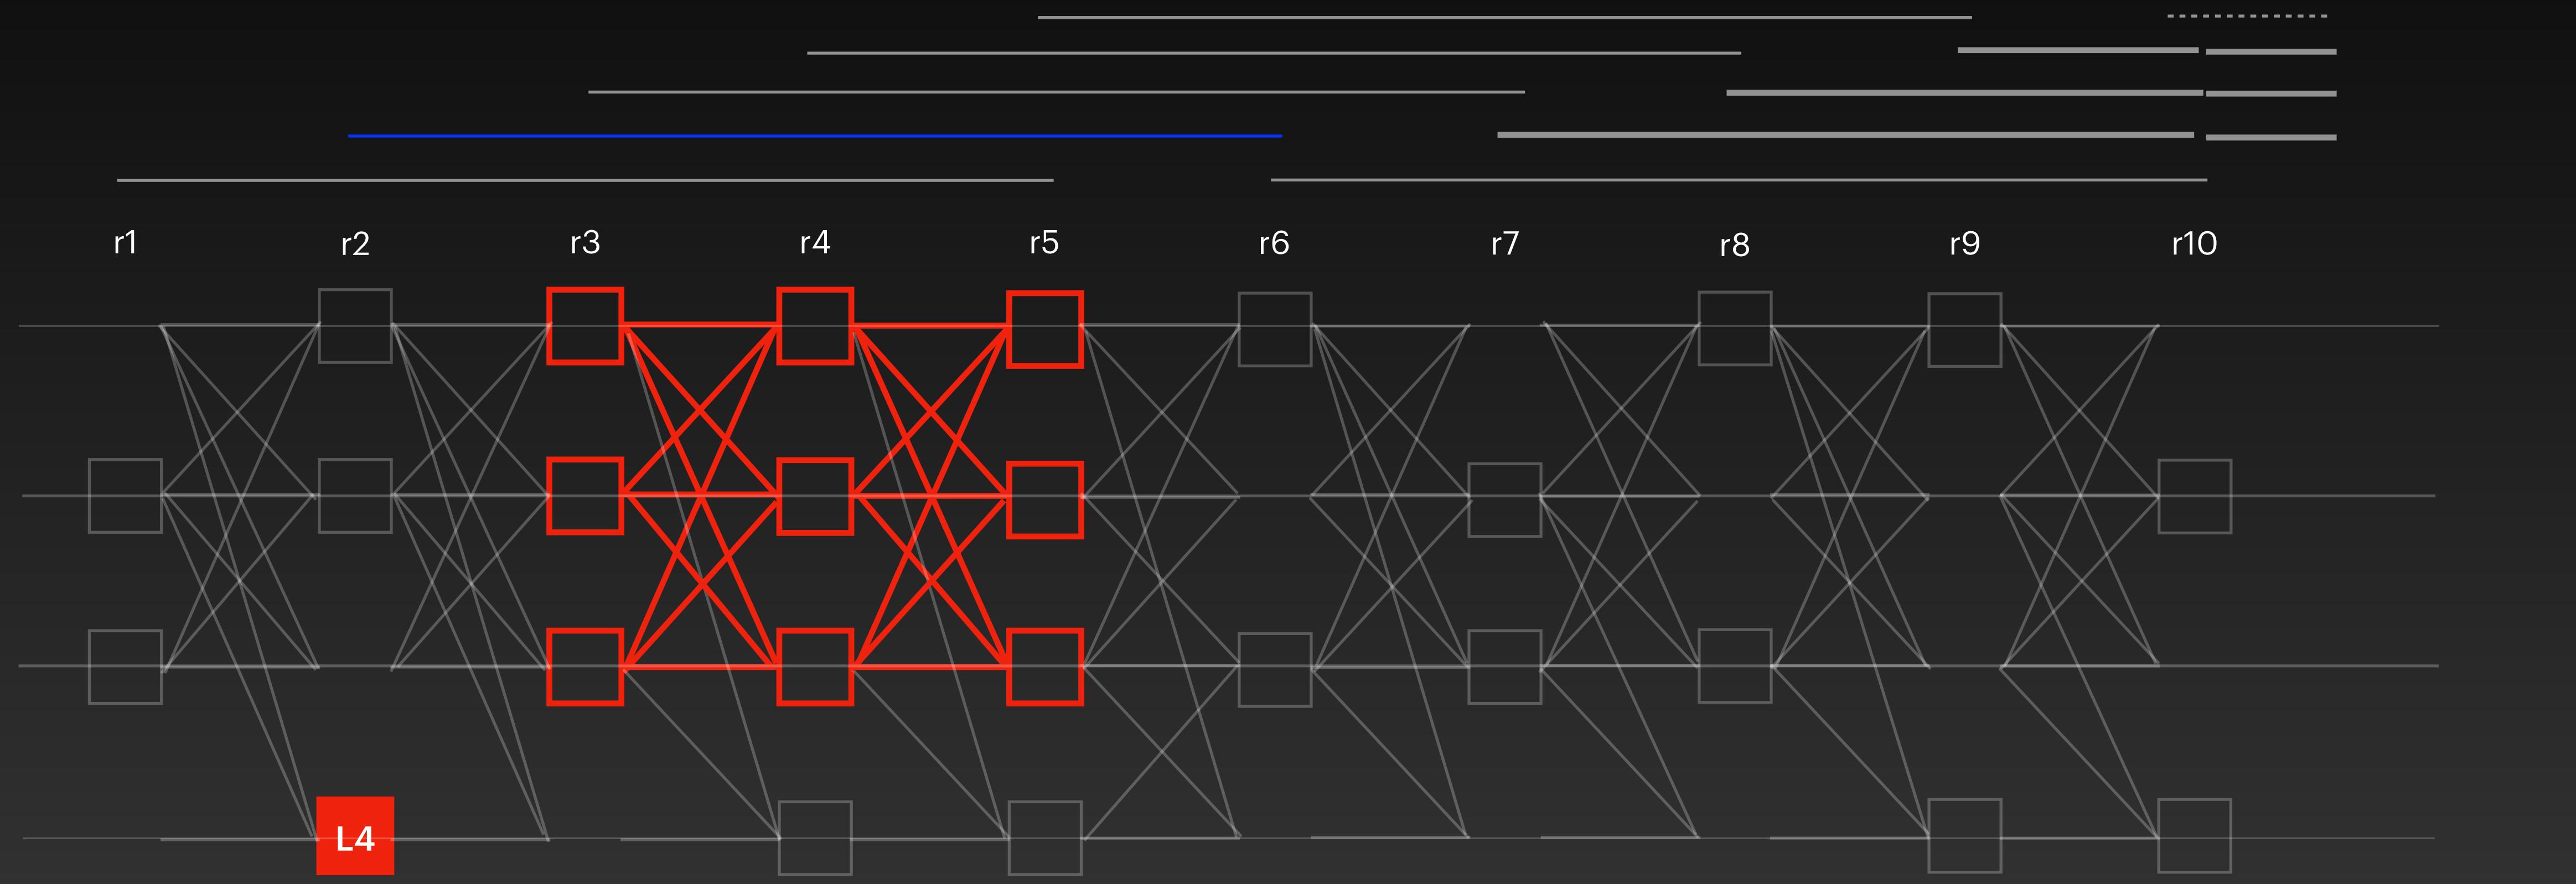
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

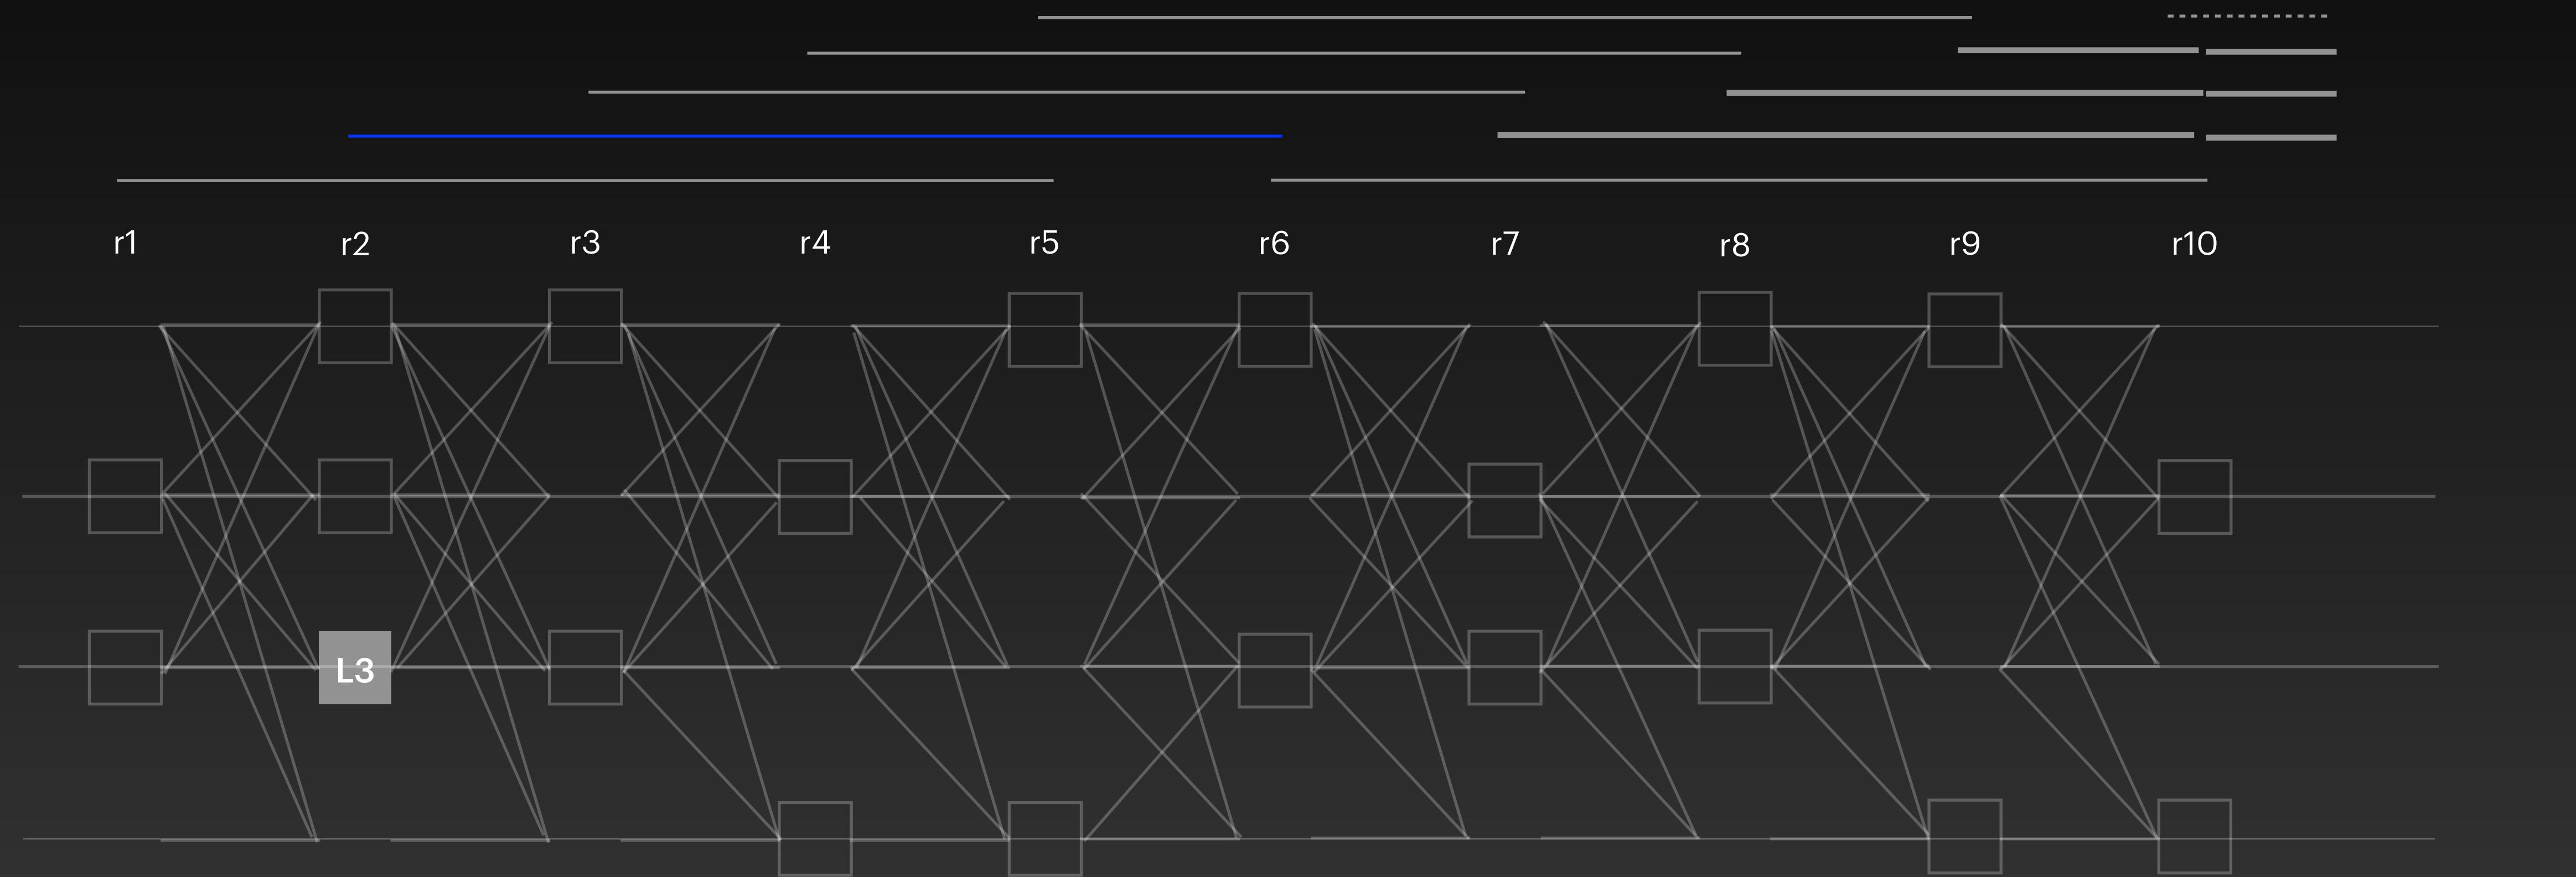
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

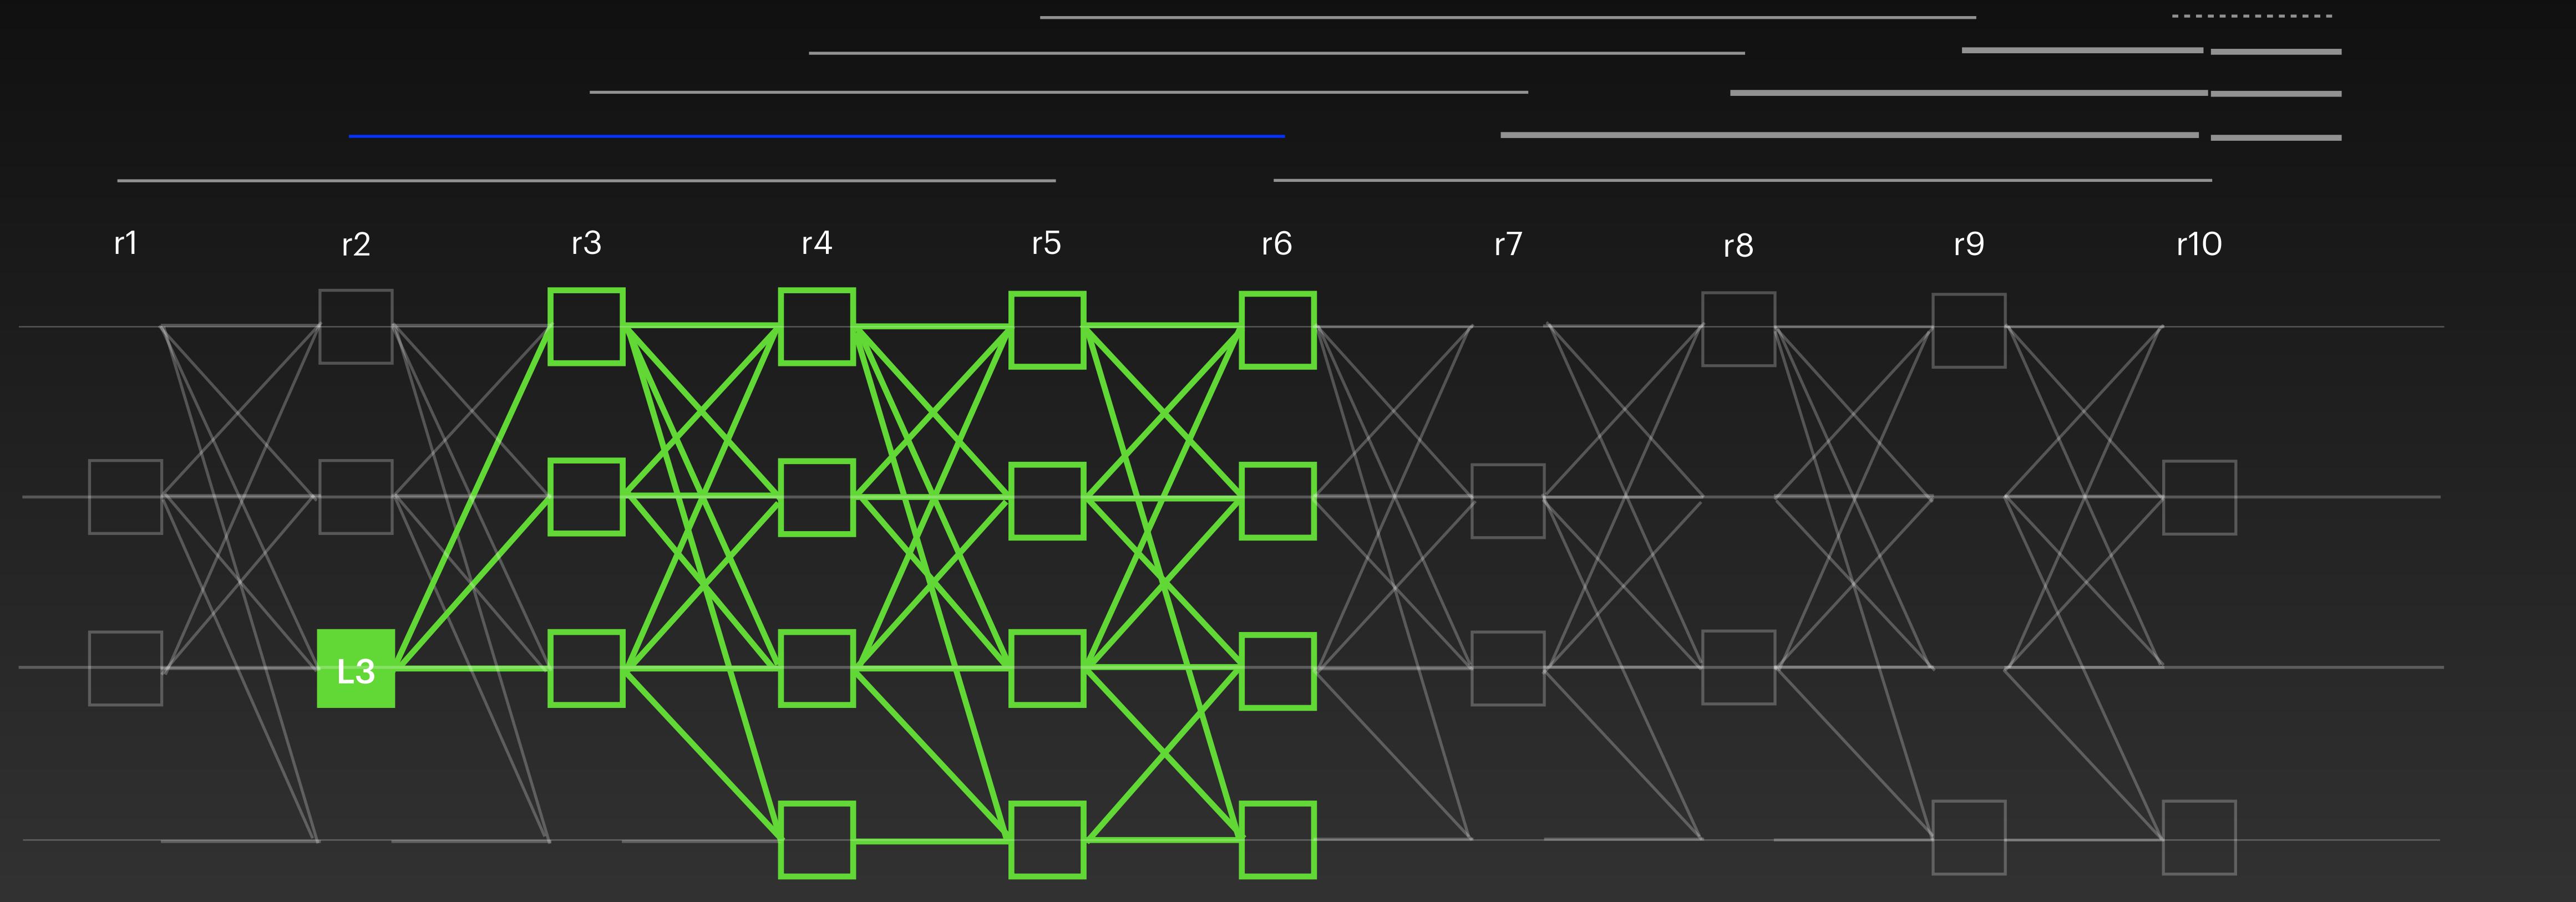
L10

L11

L12

.....

Apply Direct Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

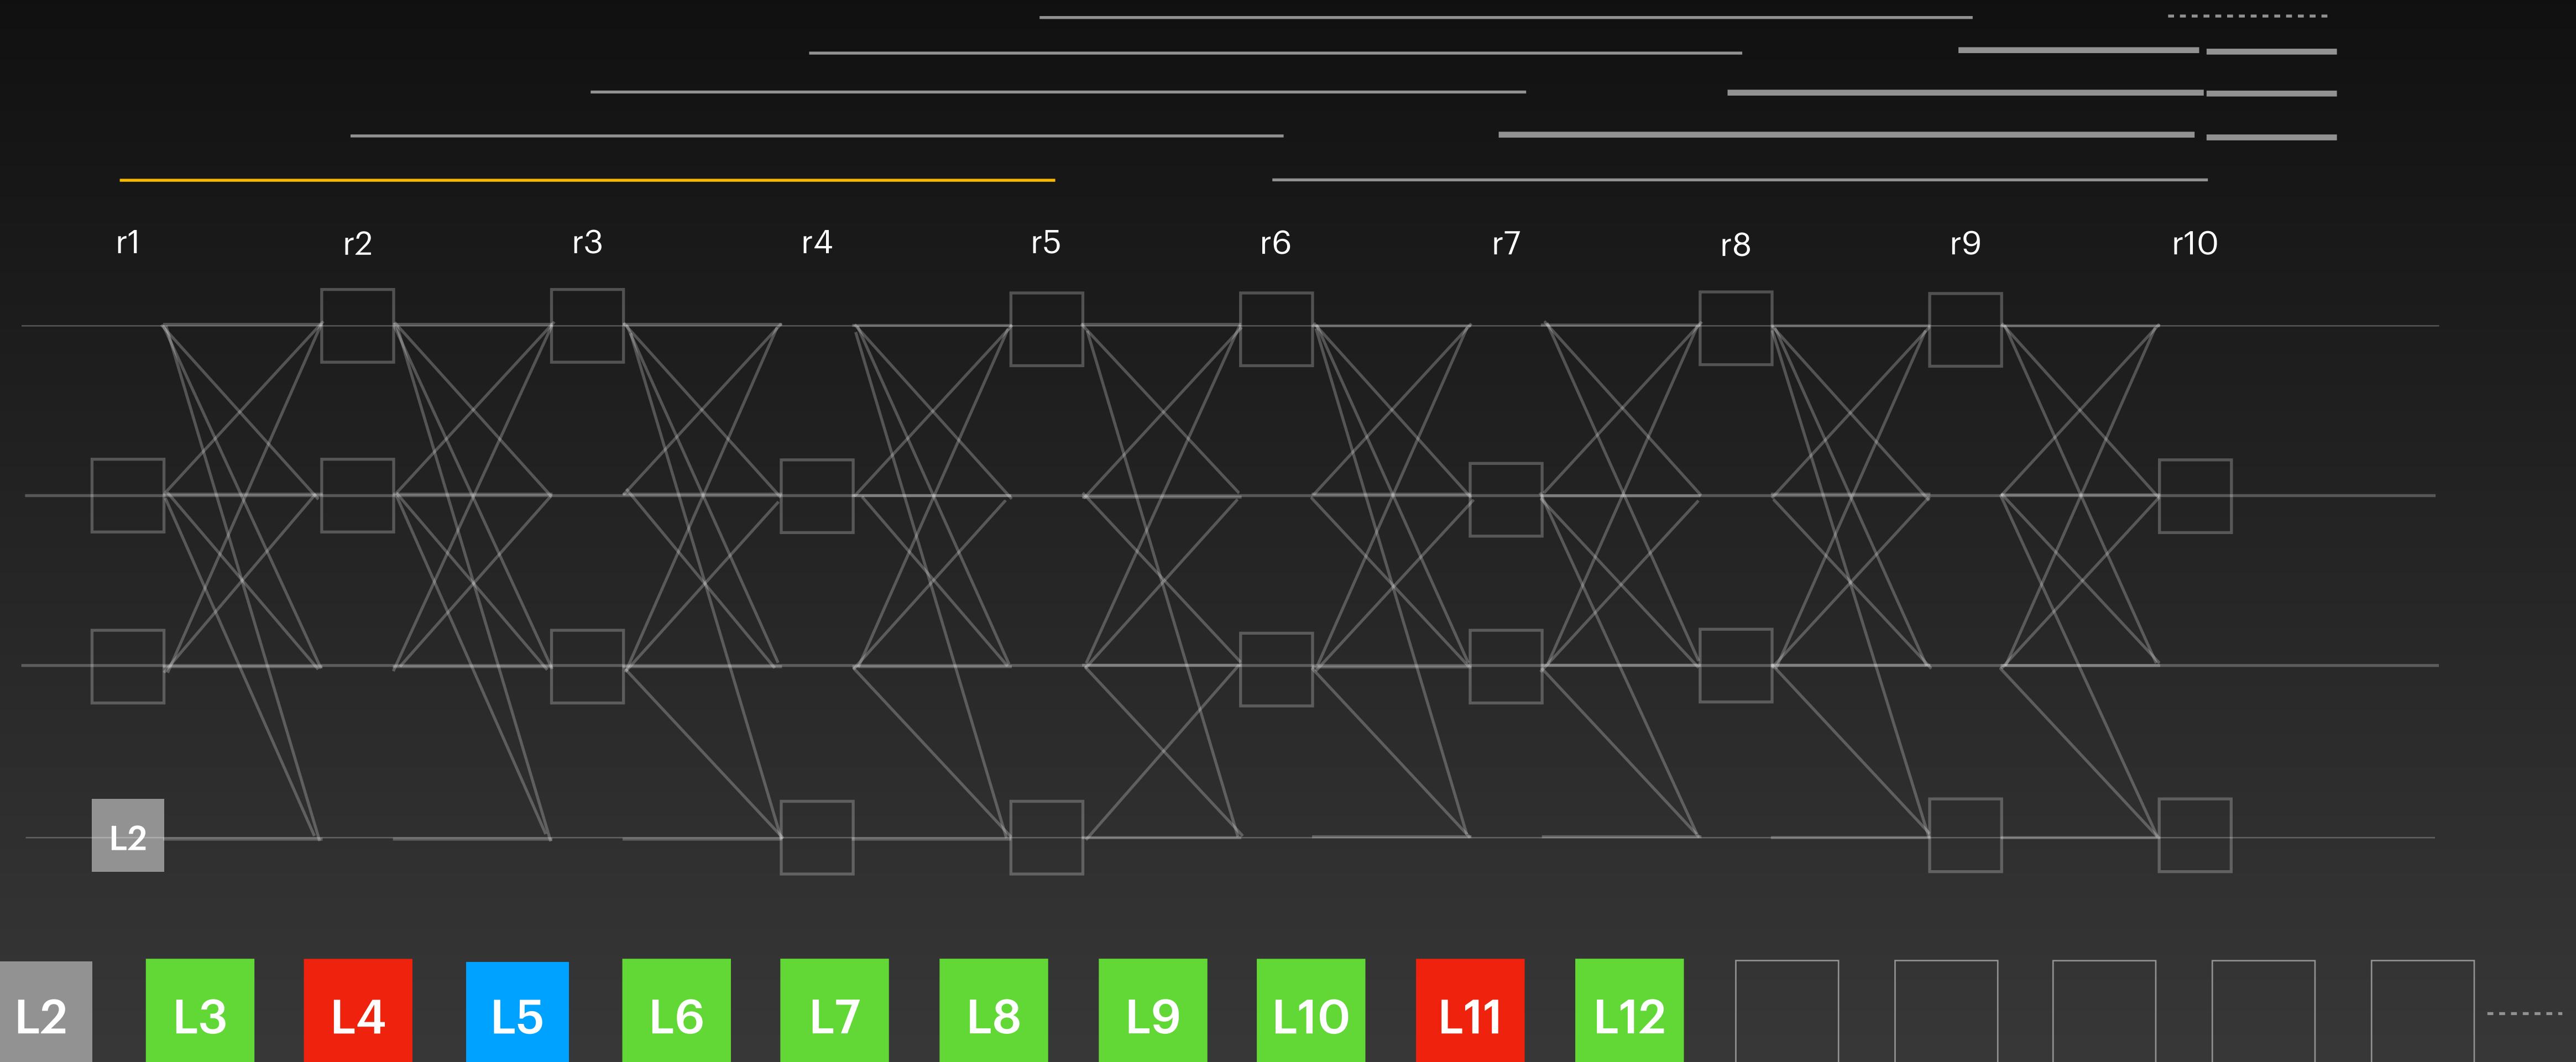
L10

L11

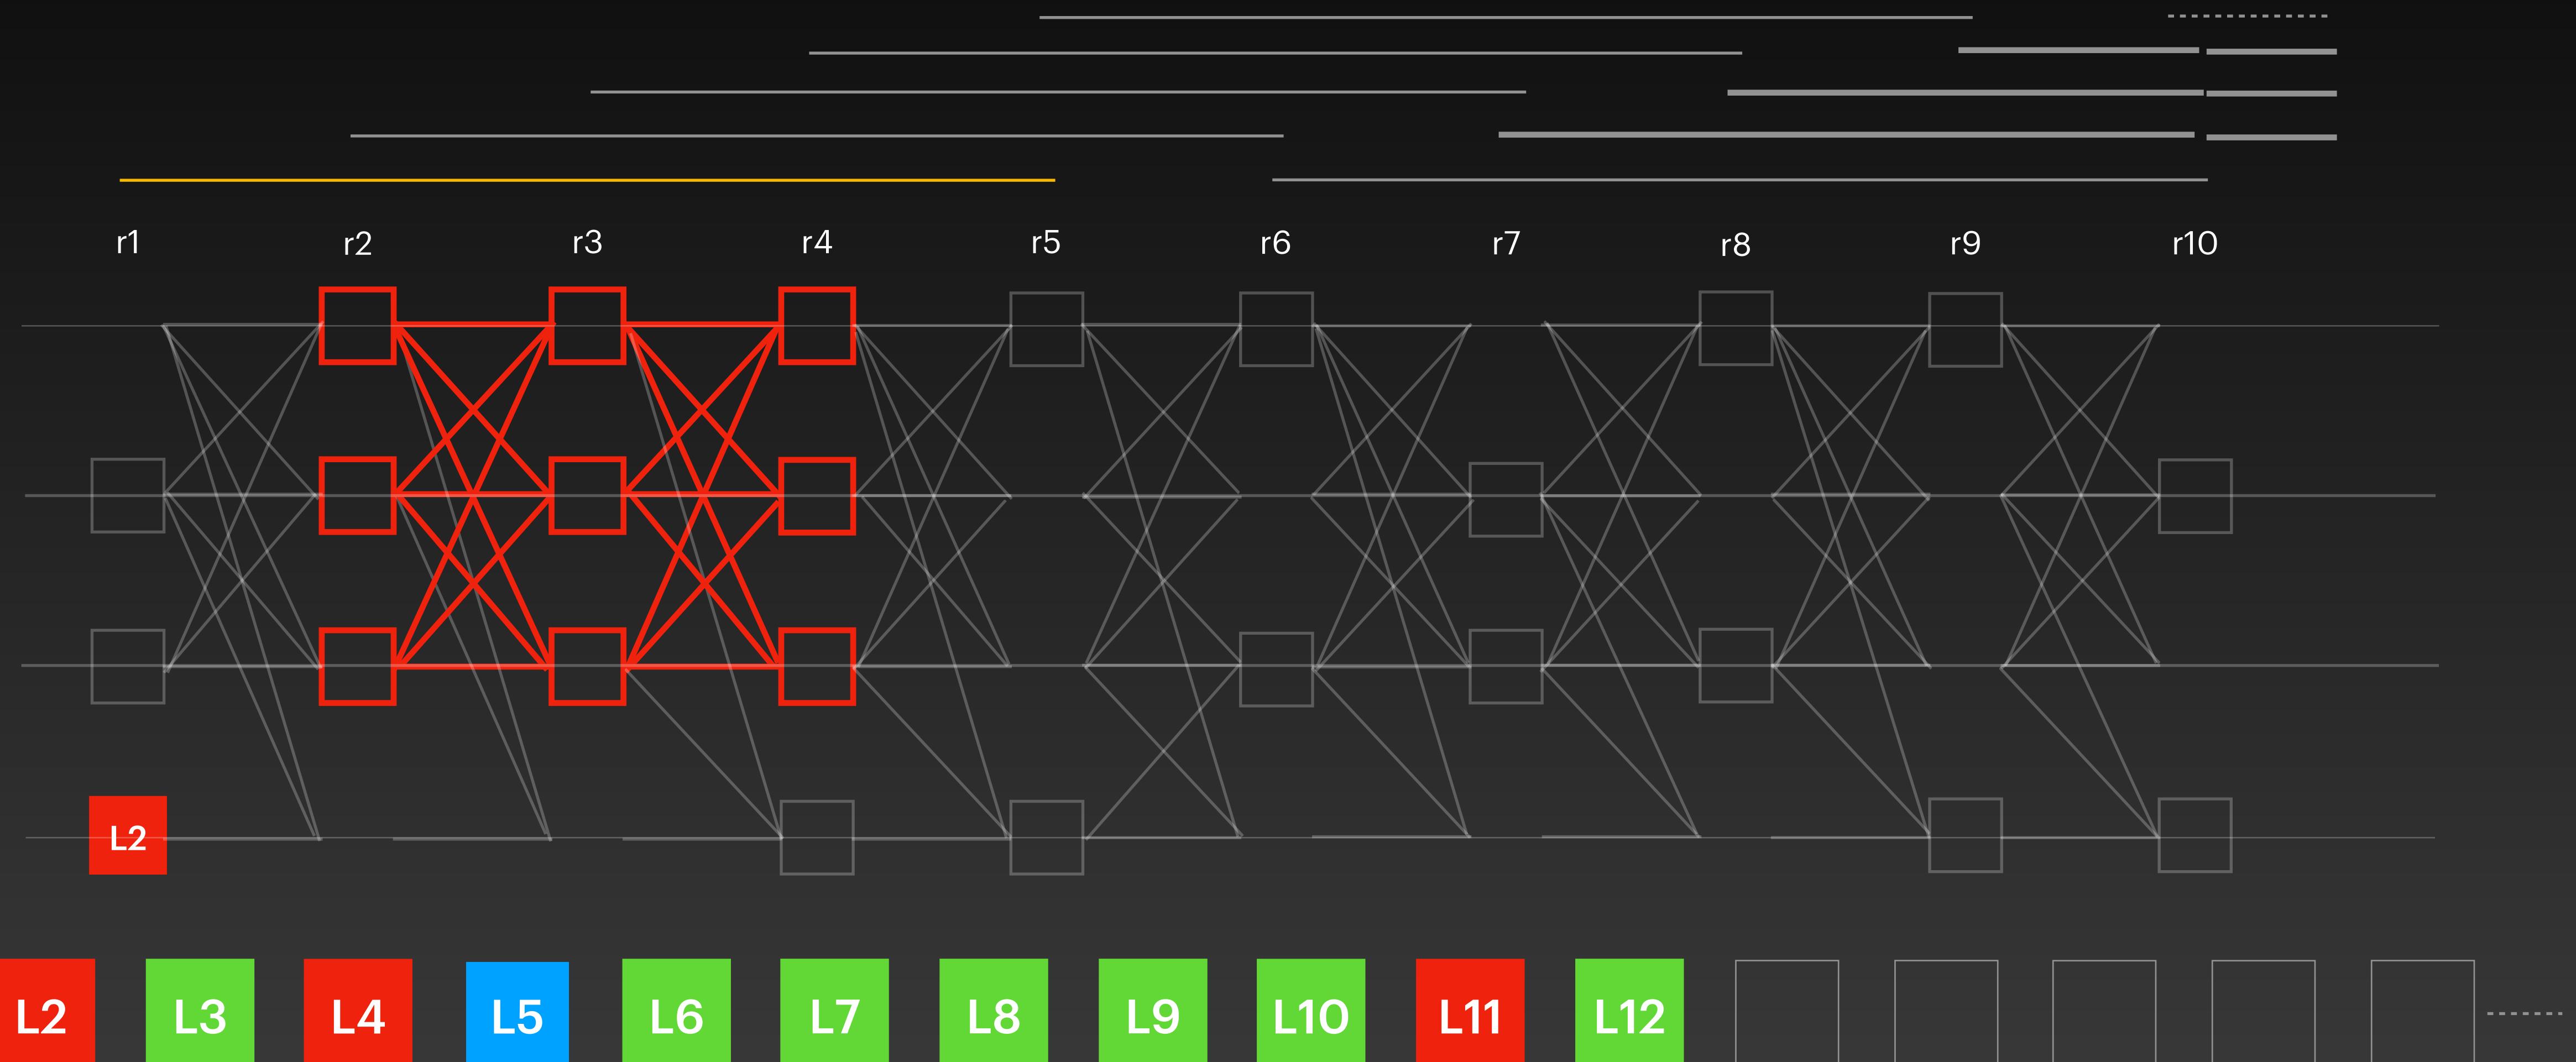
L12

...

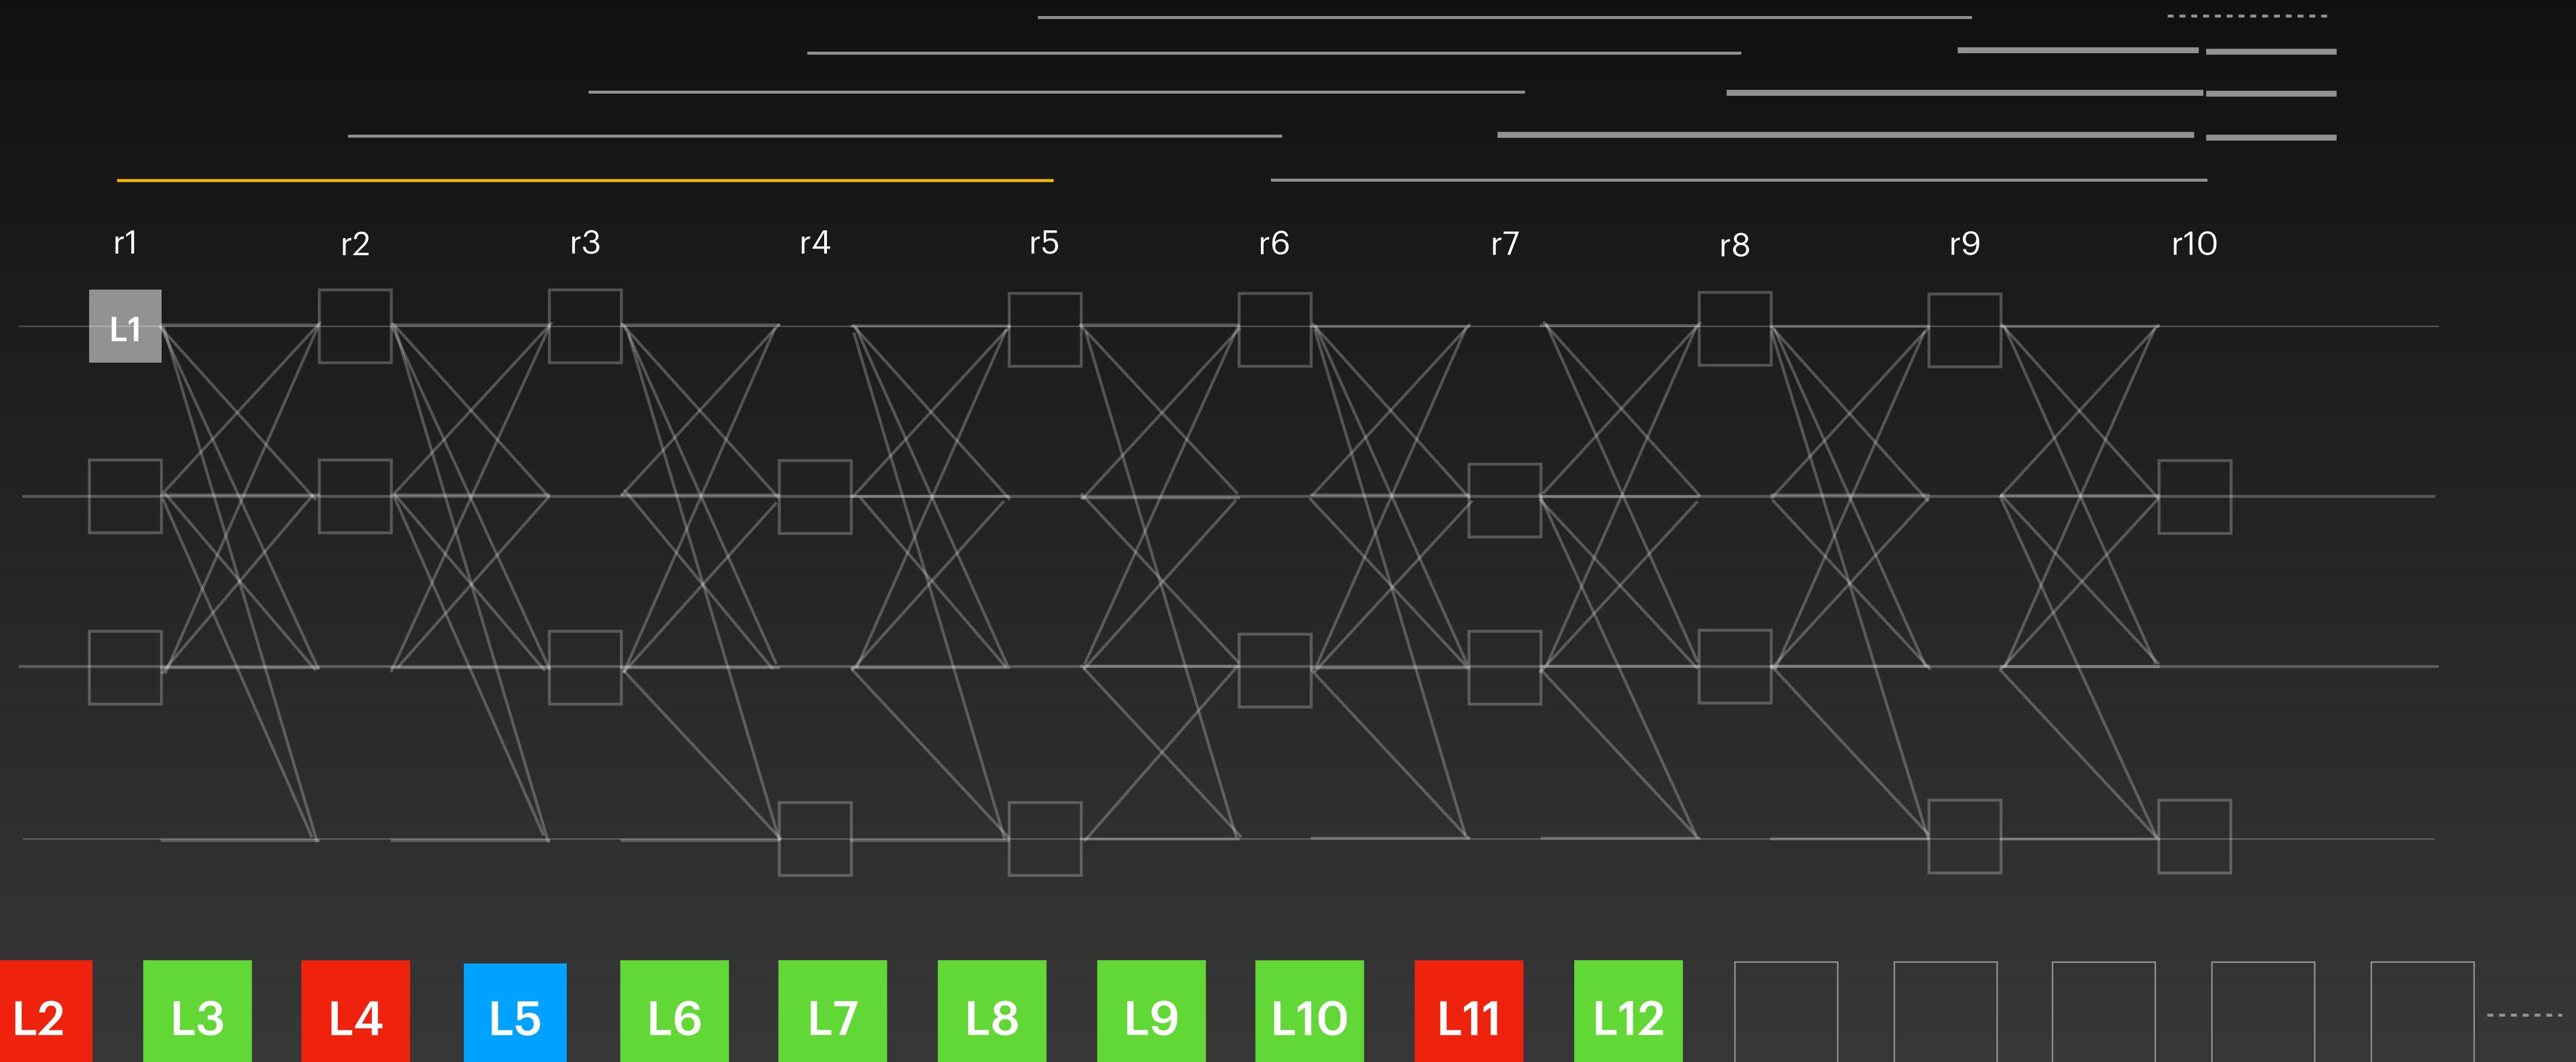
Apply Direct Rule



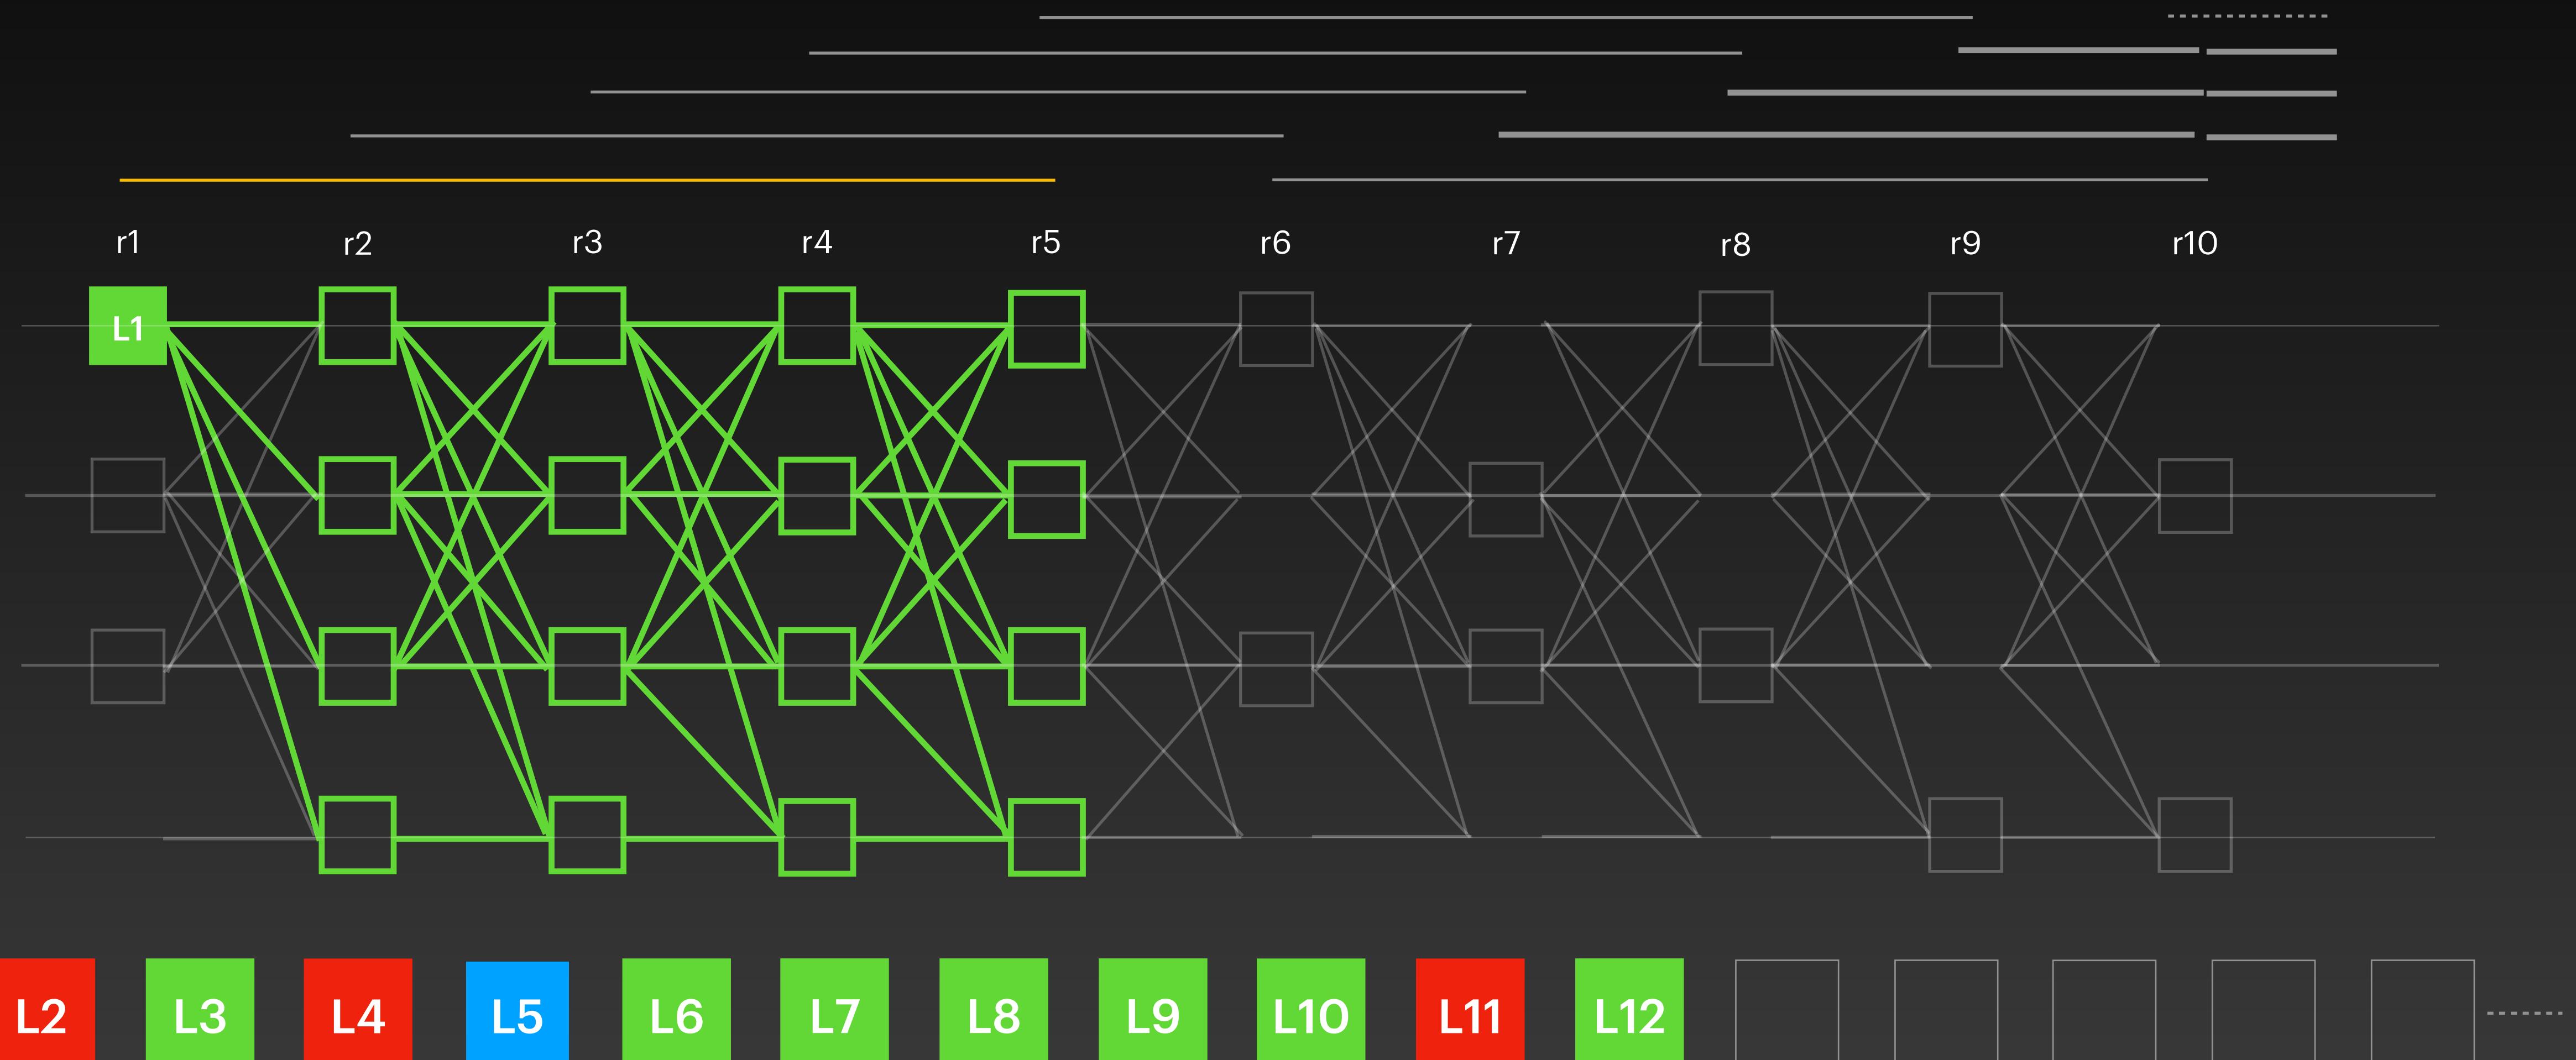
Apply Direct Rule



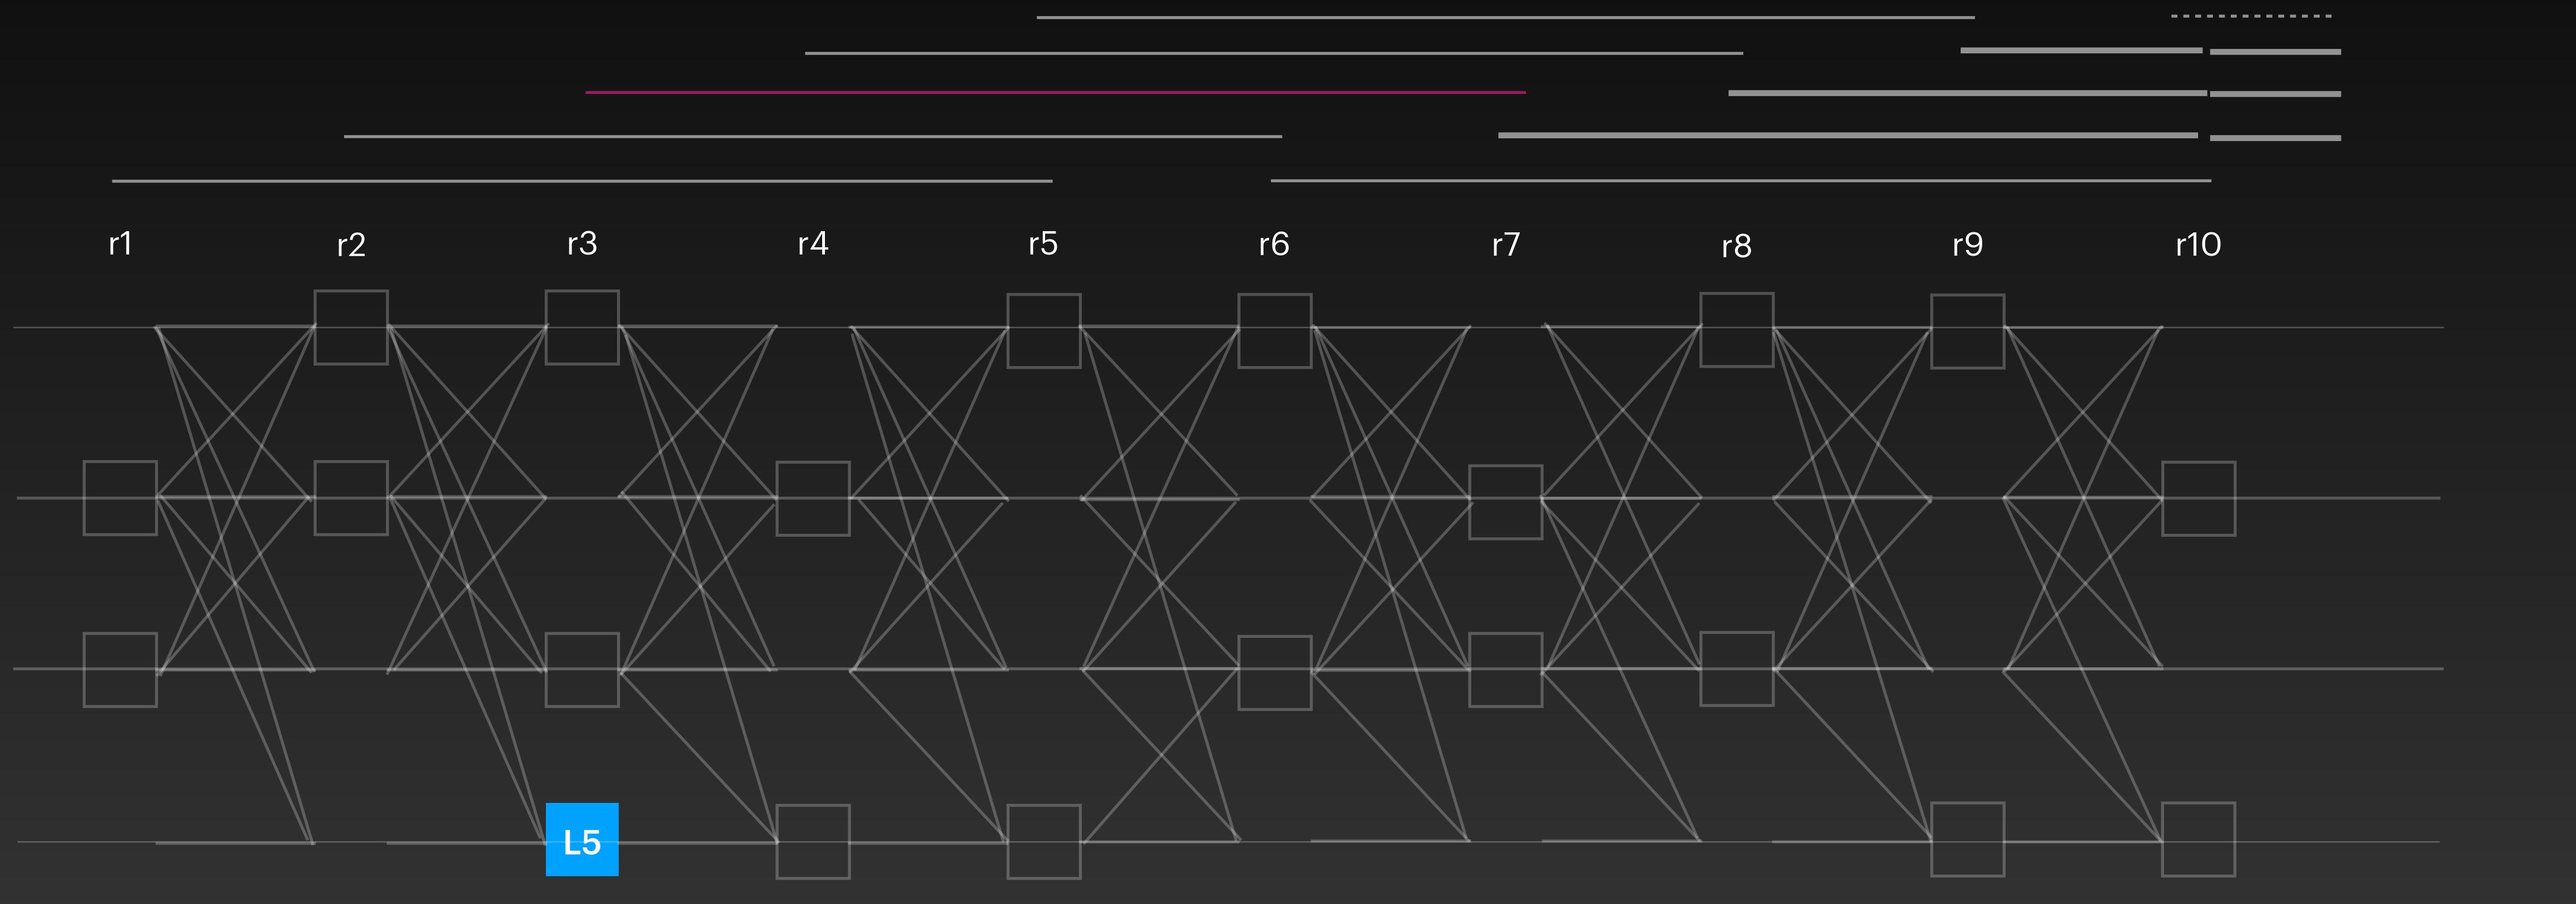
Apply Direct Rule



Apply Direct Rule



Apply Indirect Rule



L1

L2

L3

L4

L5

L6

L7

L8

L9

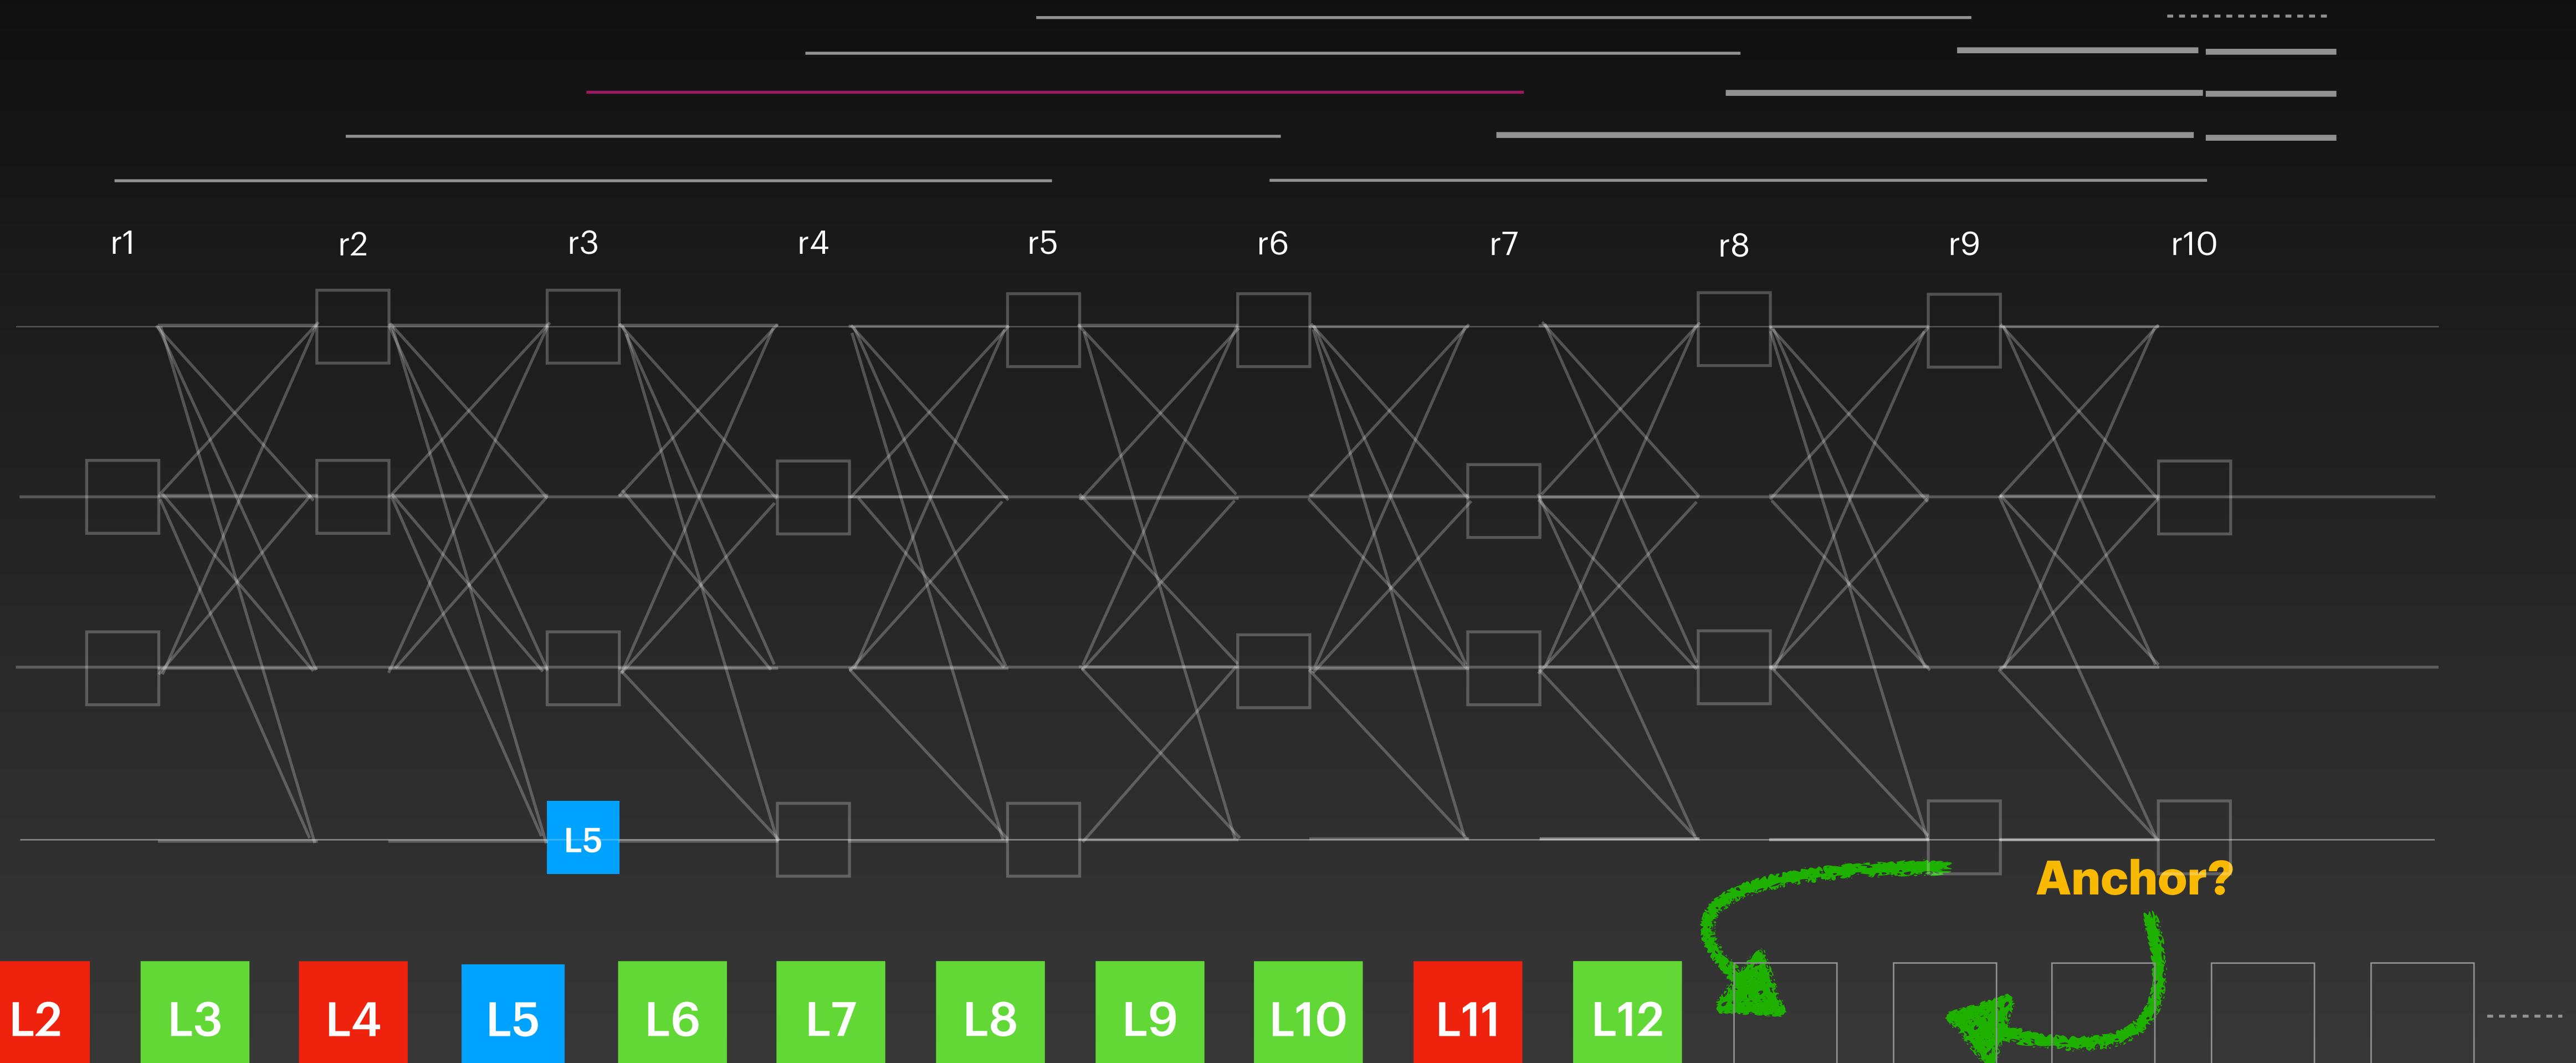
L10

L11

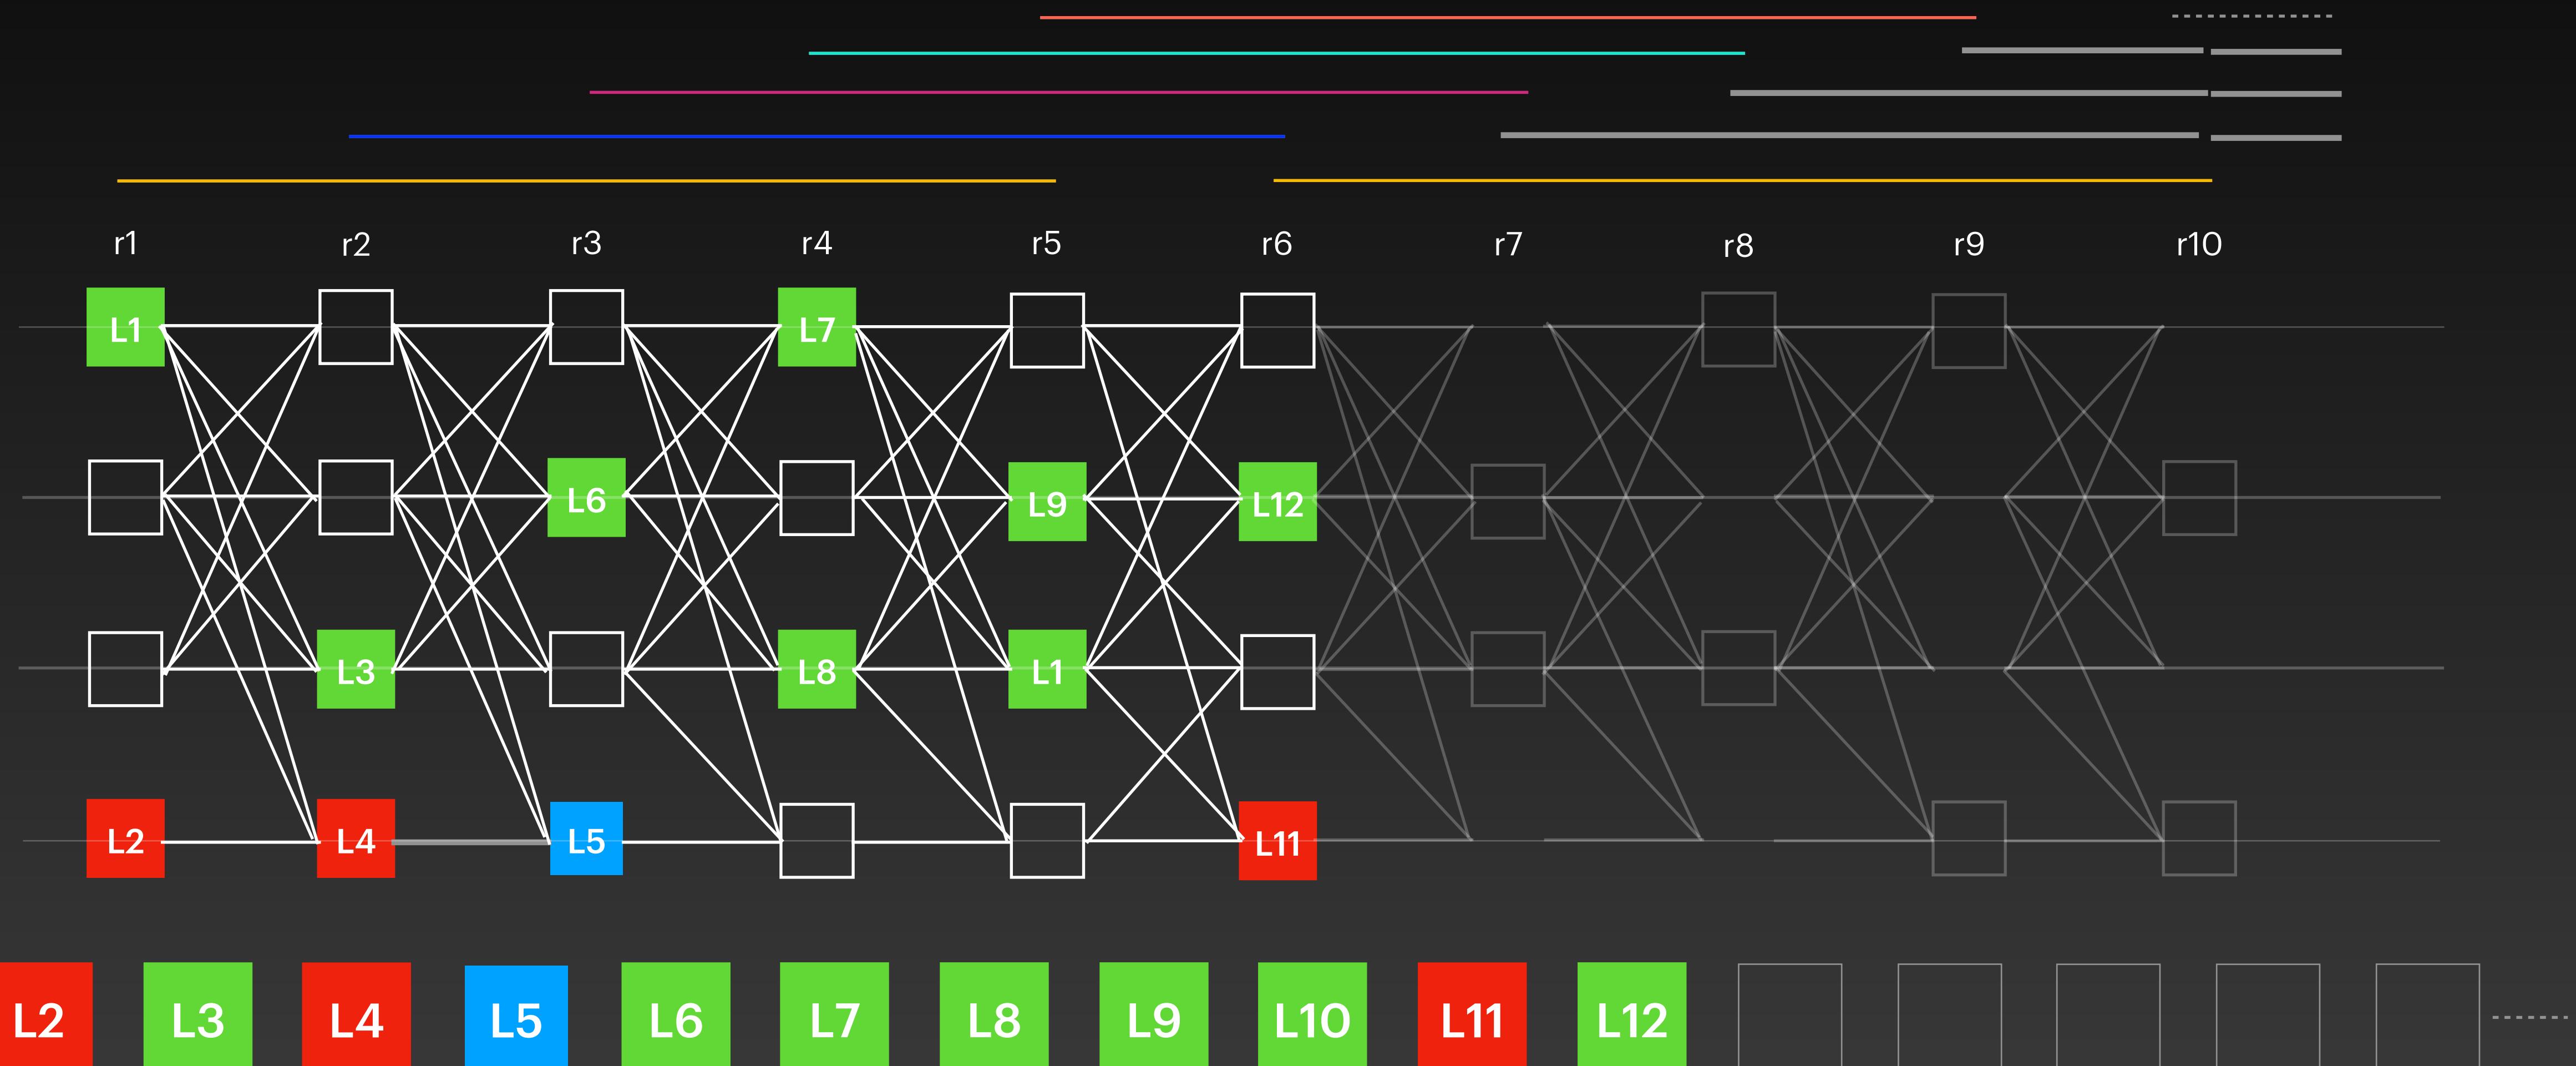
L12

.....

Need to Wait for L13 and L14

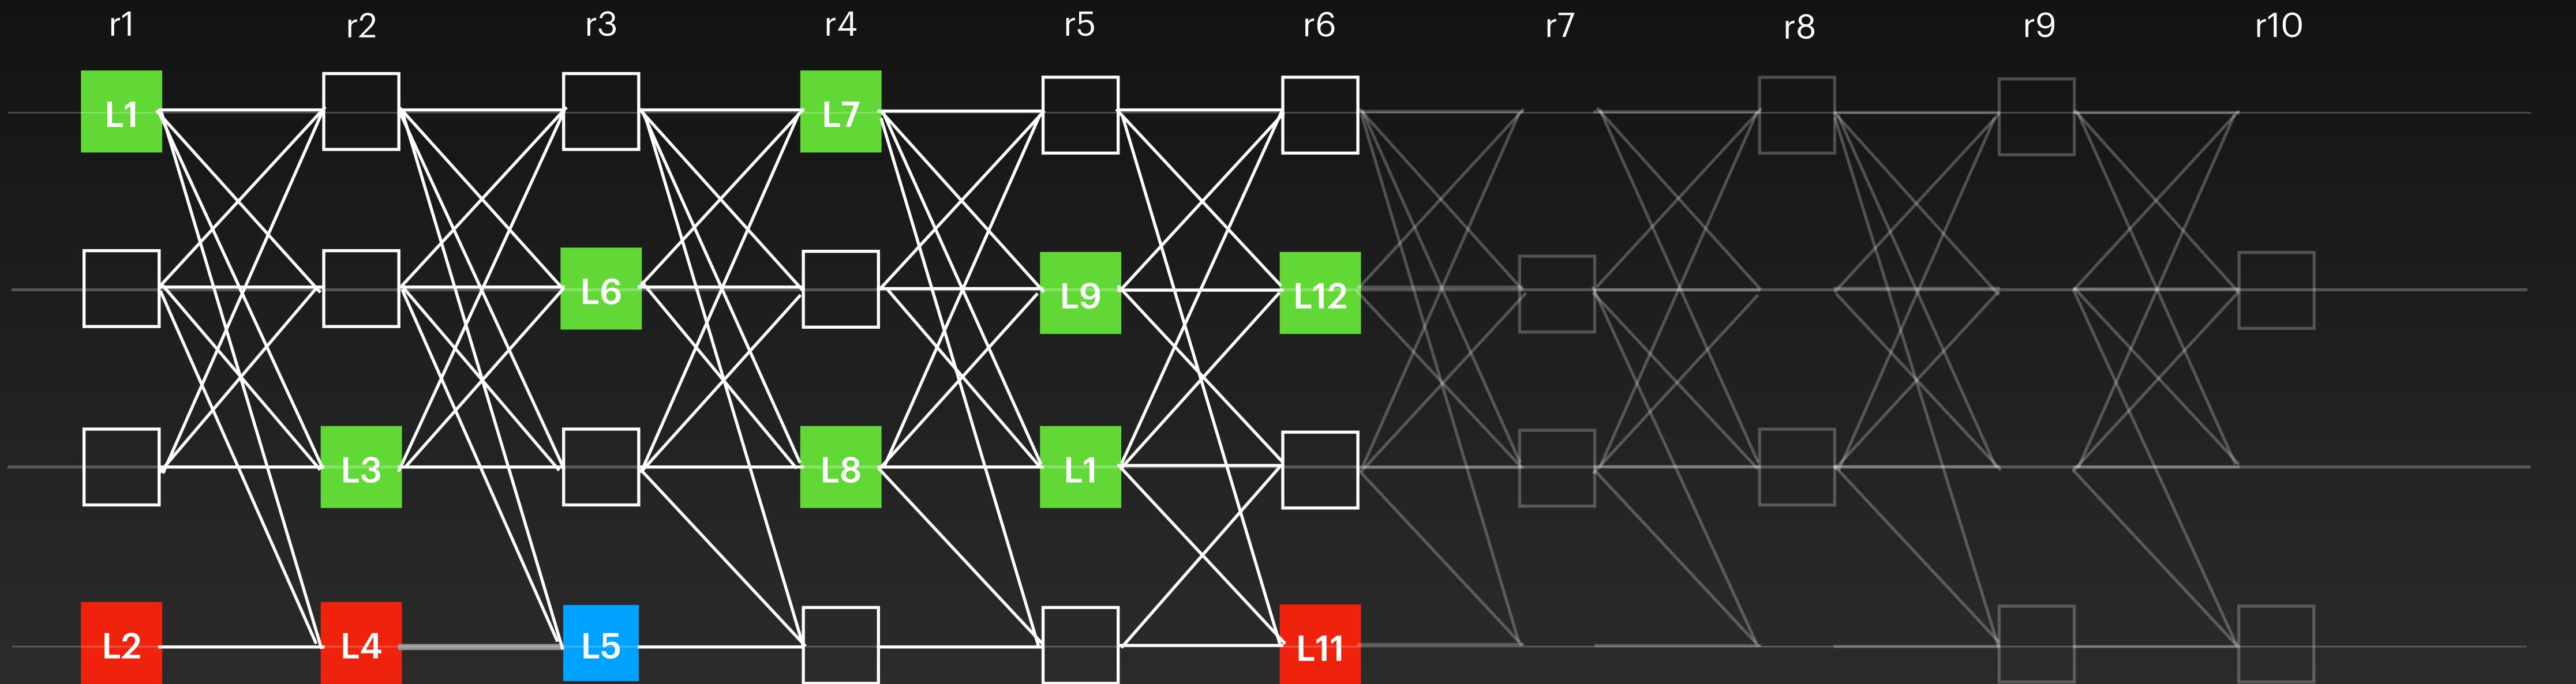


Current Status



Commit Sequence

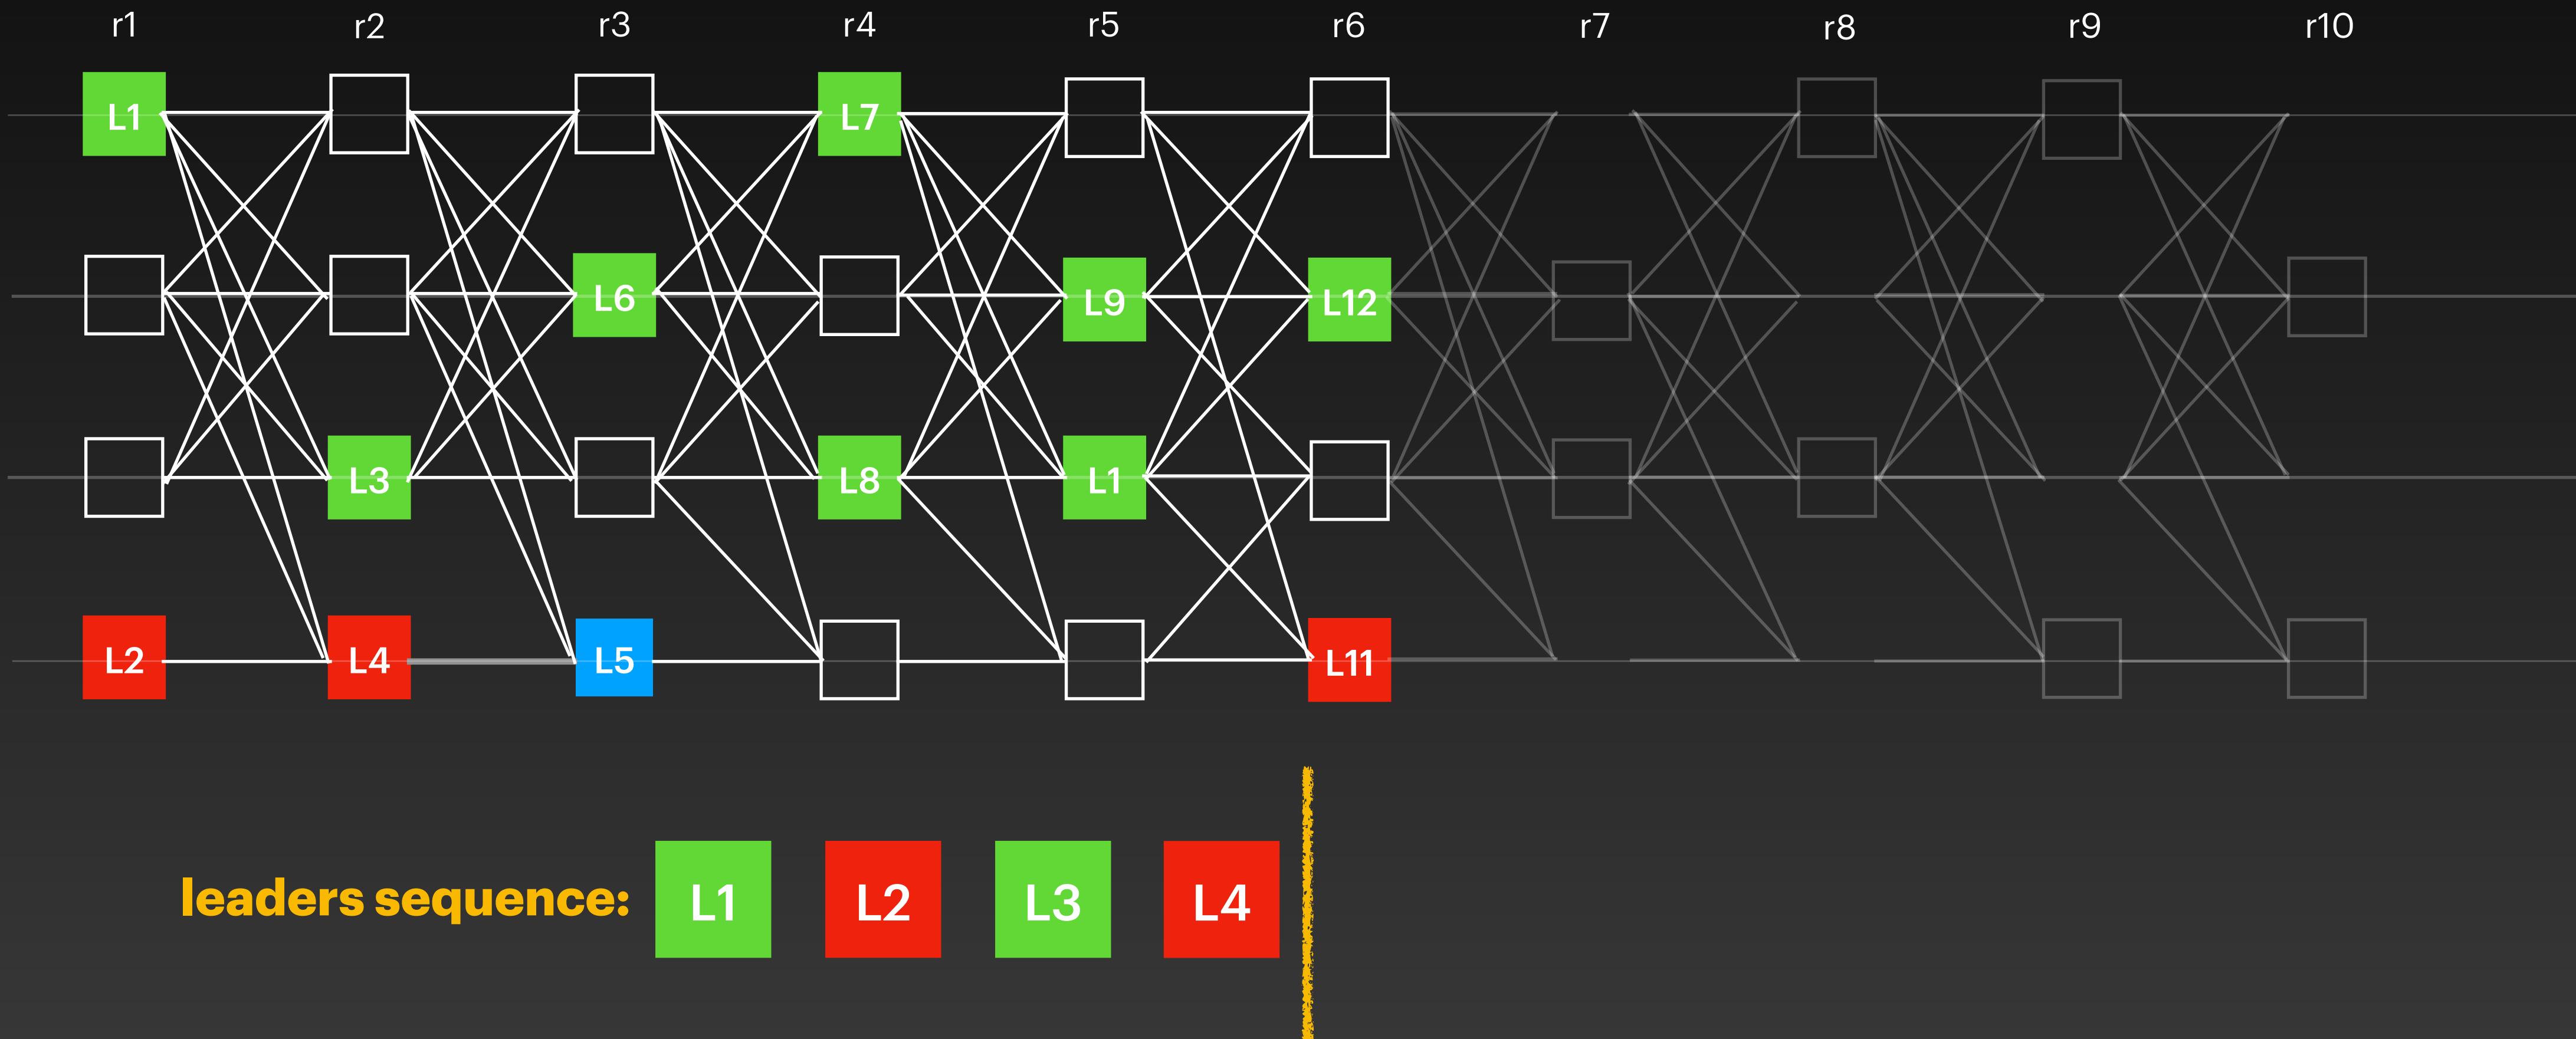
Take all leaders in order



leaders sequence: L1 L2 L3 L4 L5 L6 L7 L8 L9 L10 L11 L12

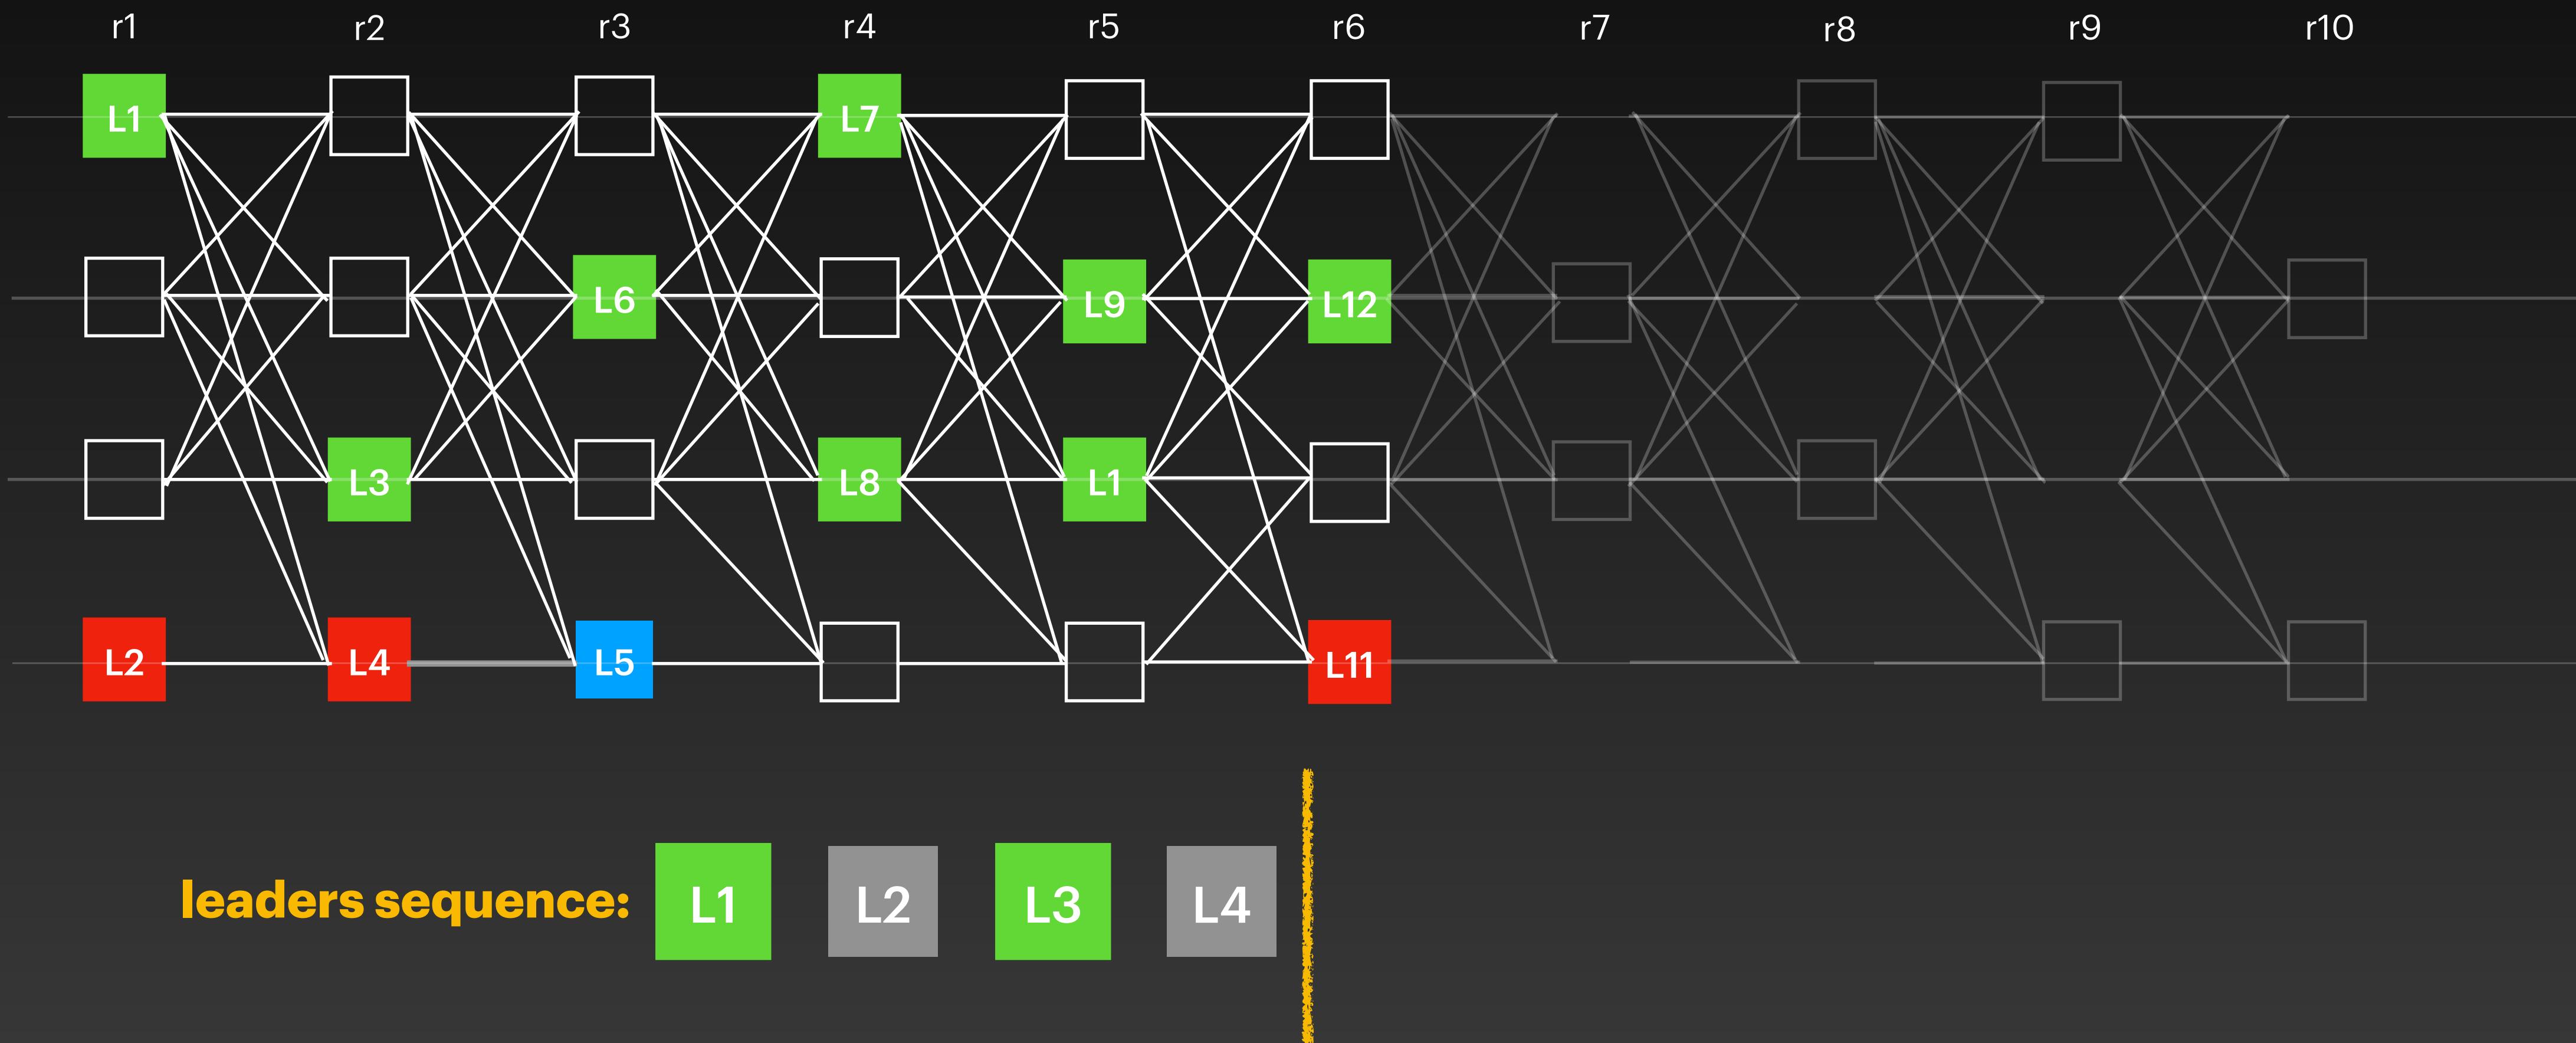
Commit Sequence

Stop at the first Undecided leader



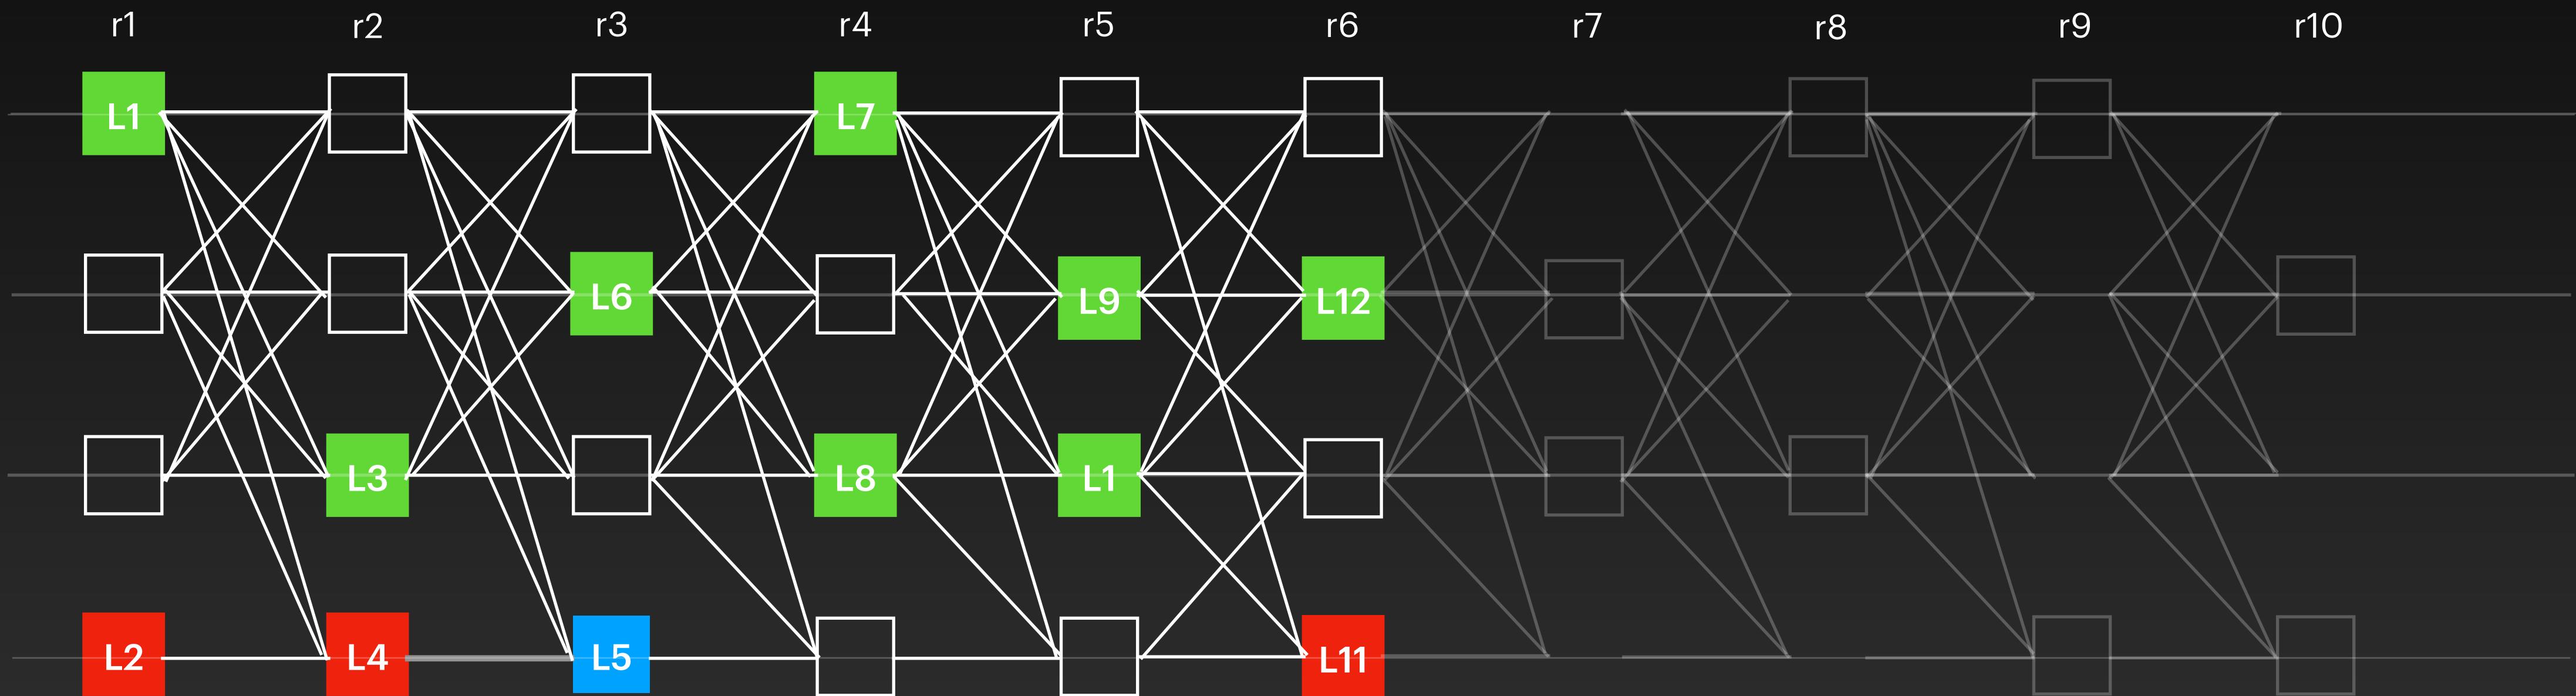
Commit Sequence

Remove skipped leaders



Commit Sequence

Final leader sequence

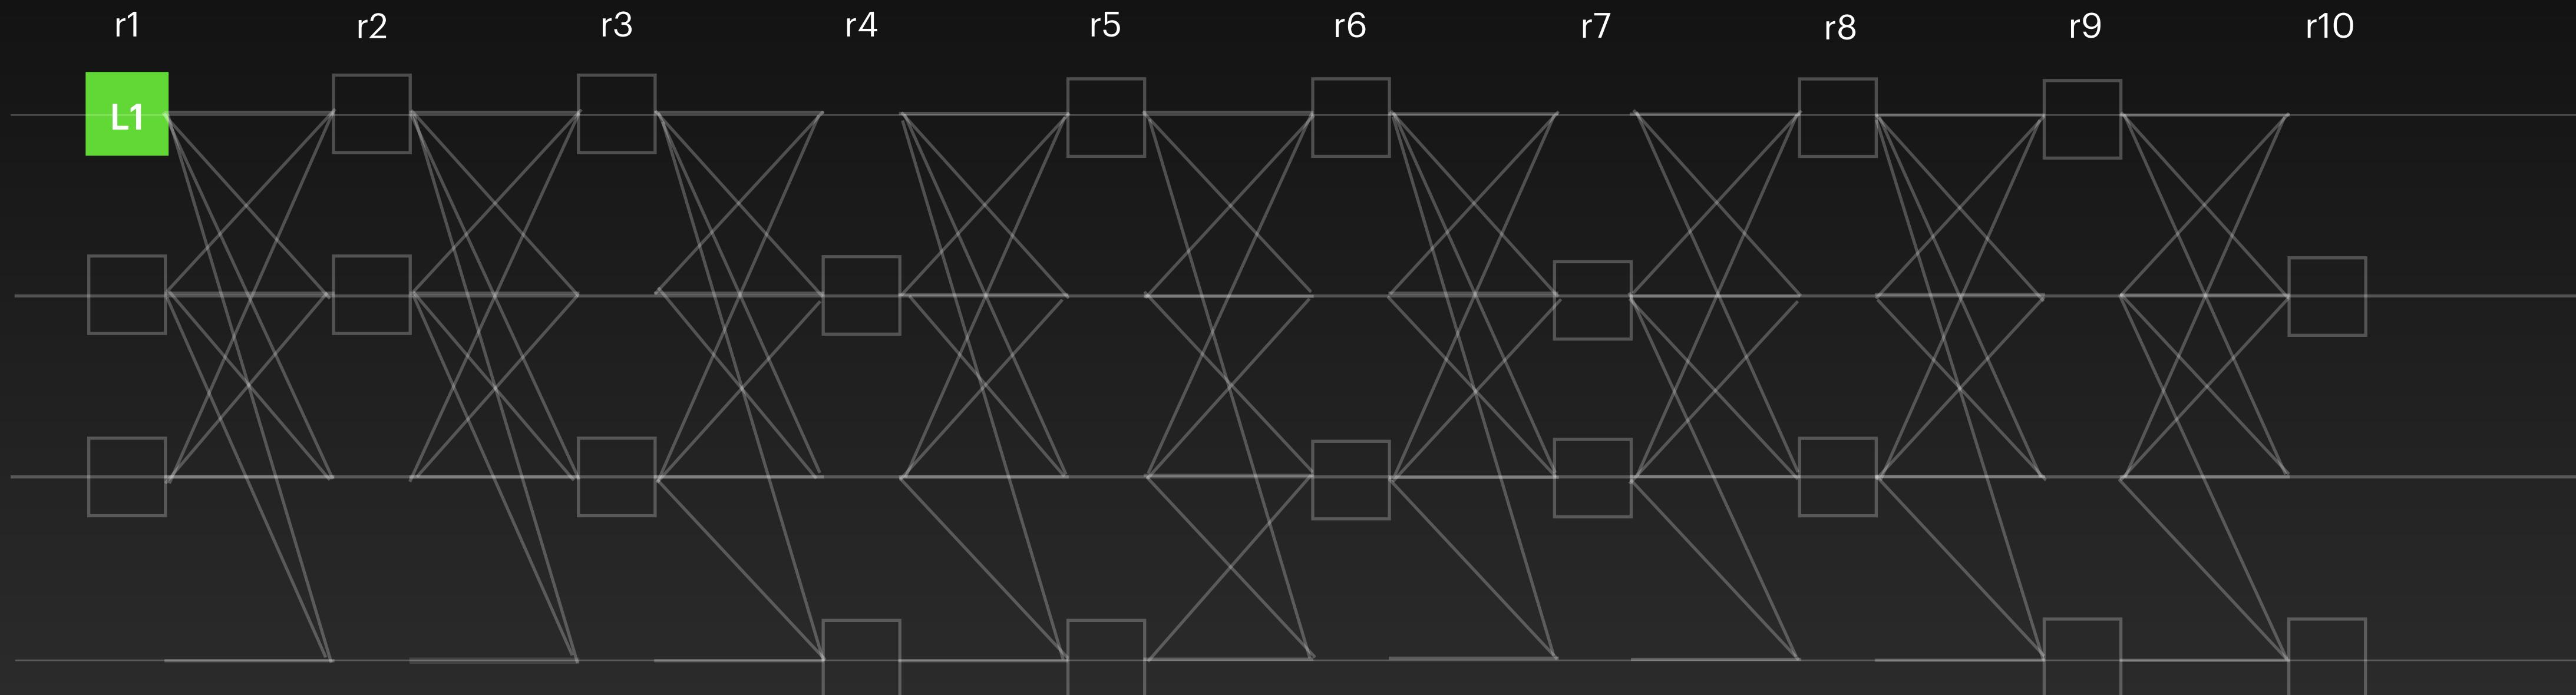


leaders sequence:



Commit Sequence

Commit sub-dag



leaders sequence:

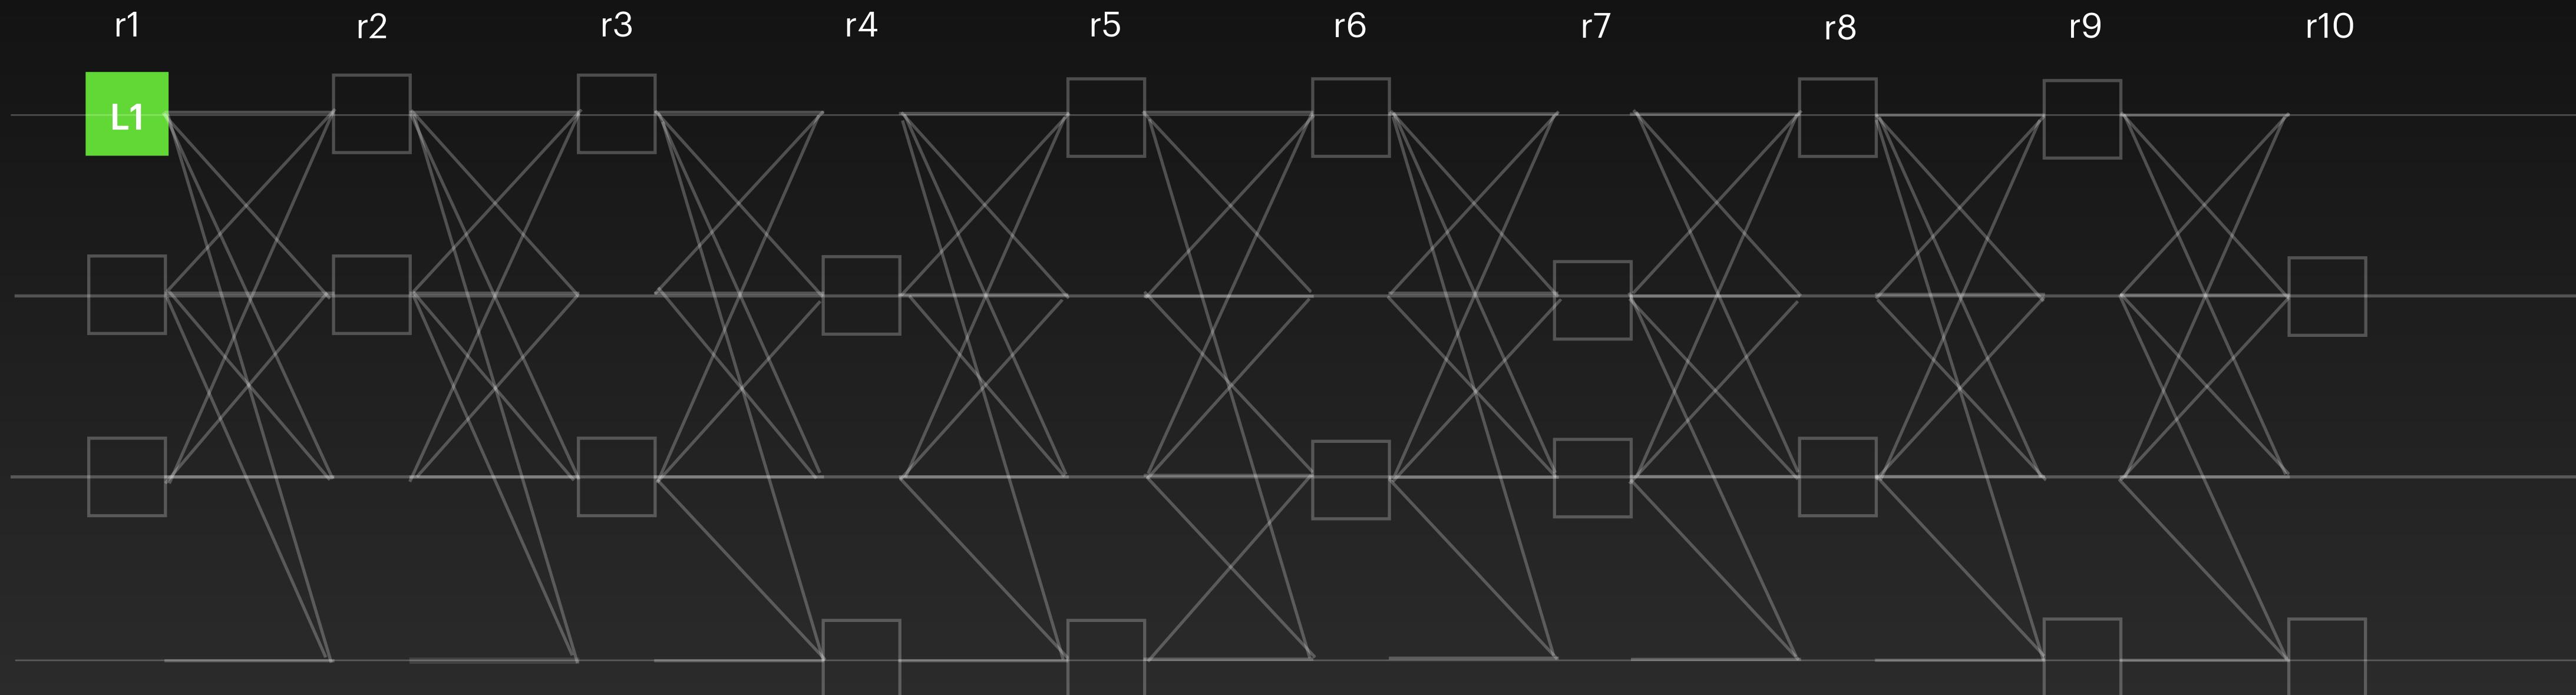
L1

L3

output sequence:

Commit Sequence

Commit sub-dag



leaders sequence:

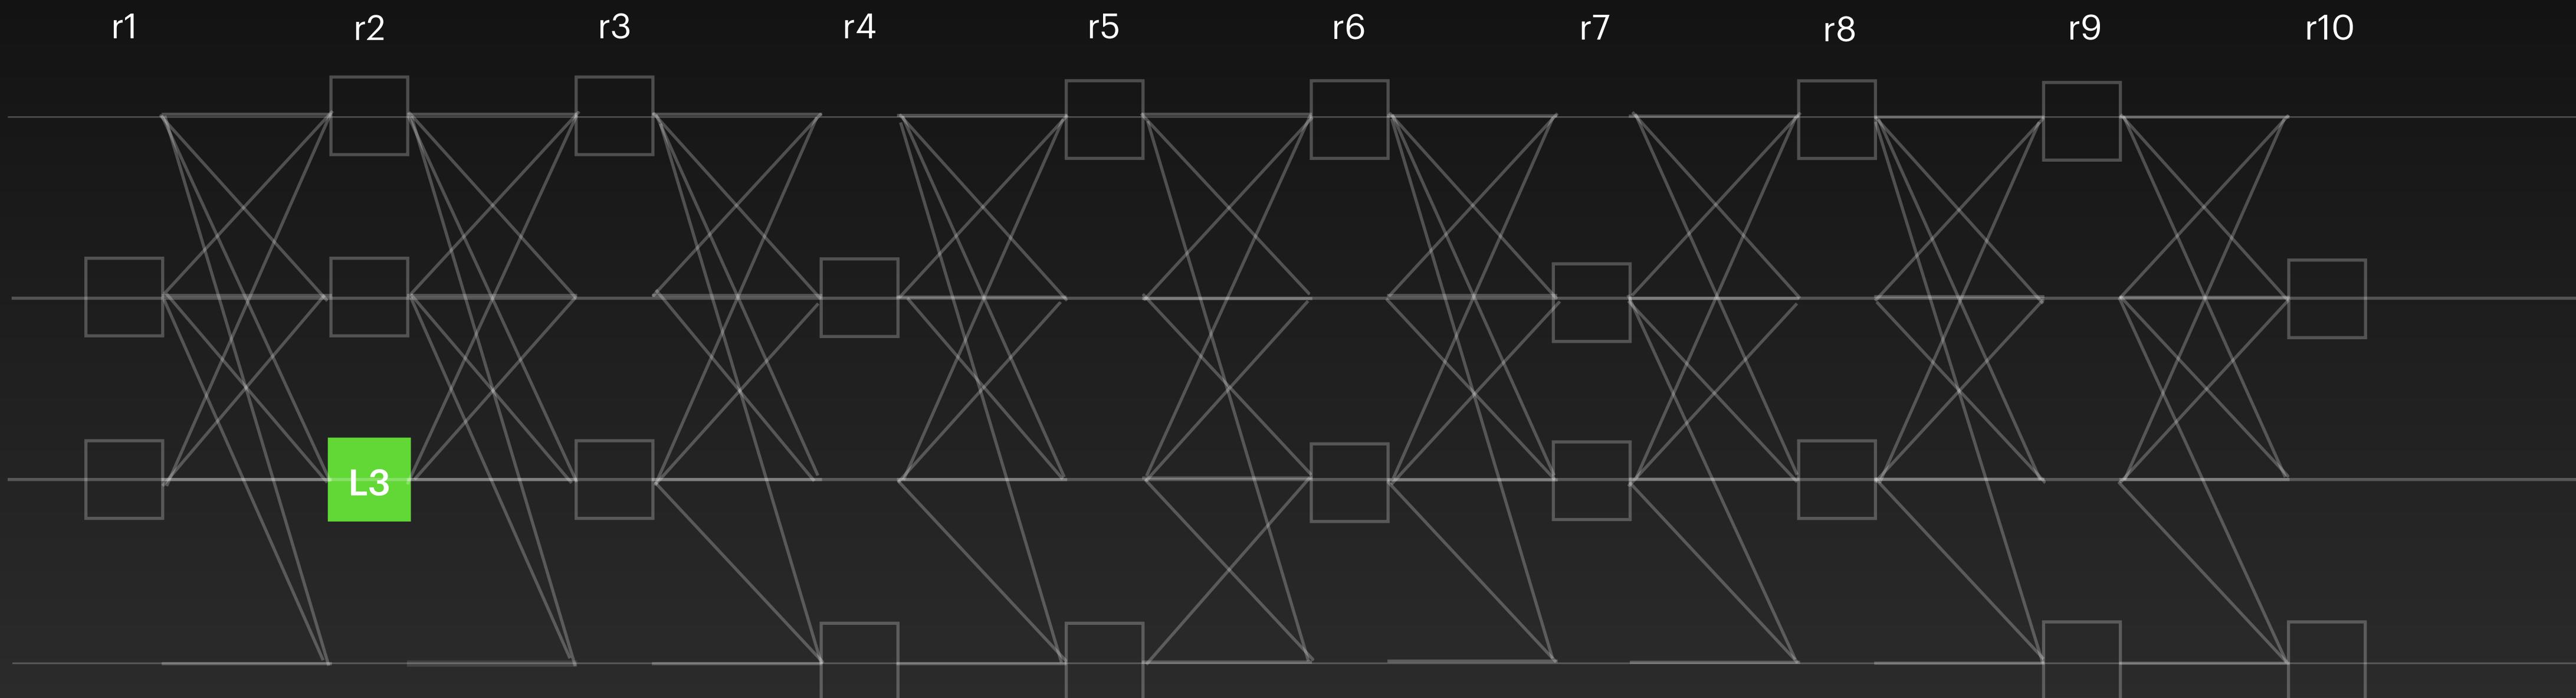
L3

output sequence:

L1

Commit Sequence

Commit sub-dag



leaders sequence:

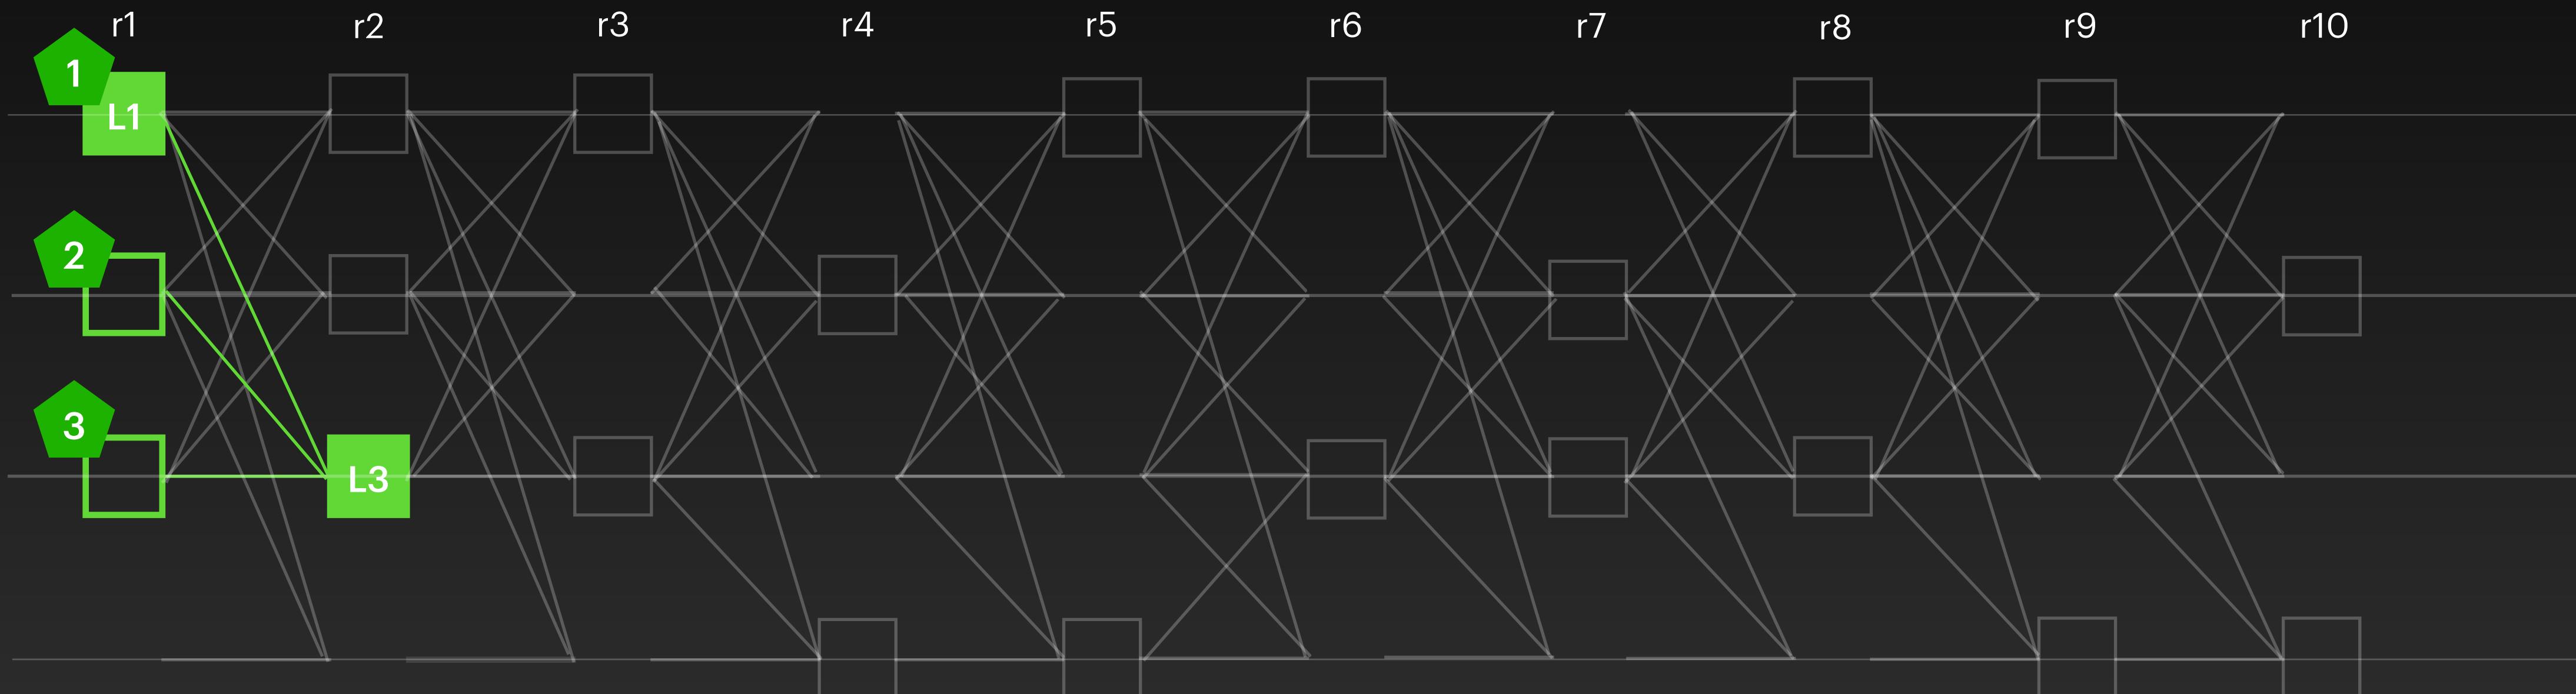
L3

output sequence:

L1

Commit Sequence

Commit sub-dag



leaders sequence:

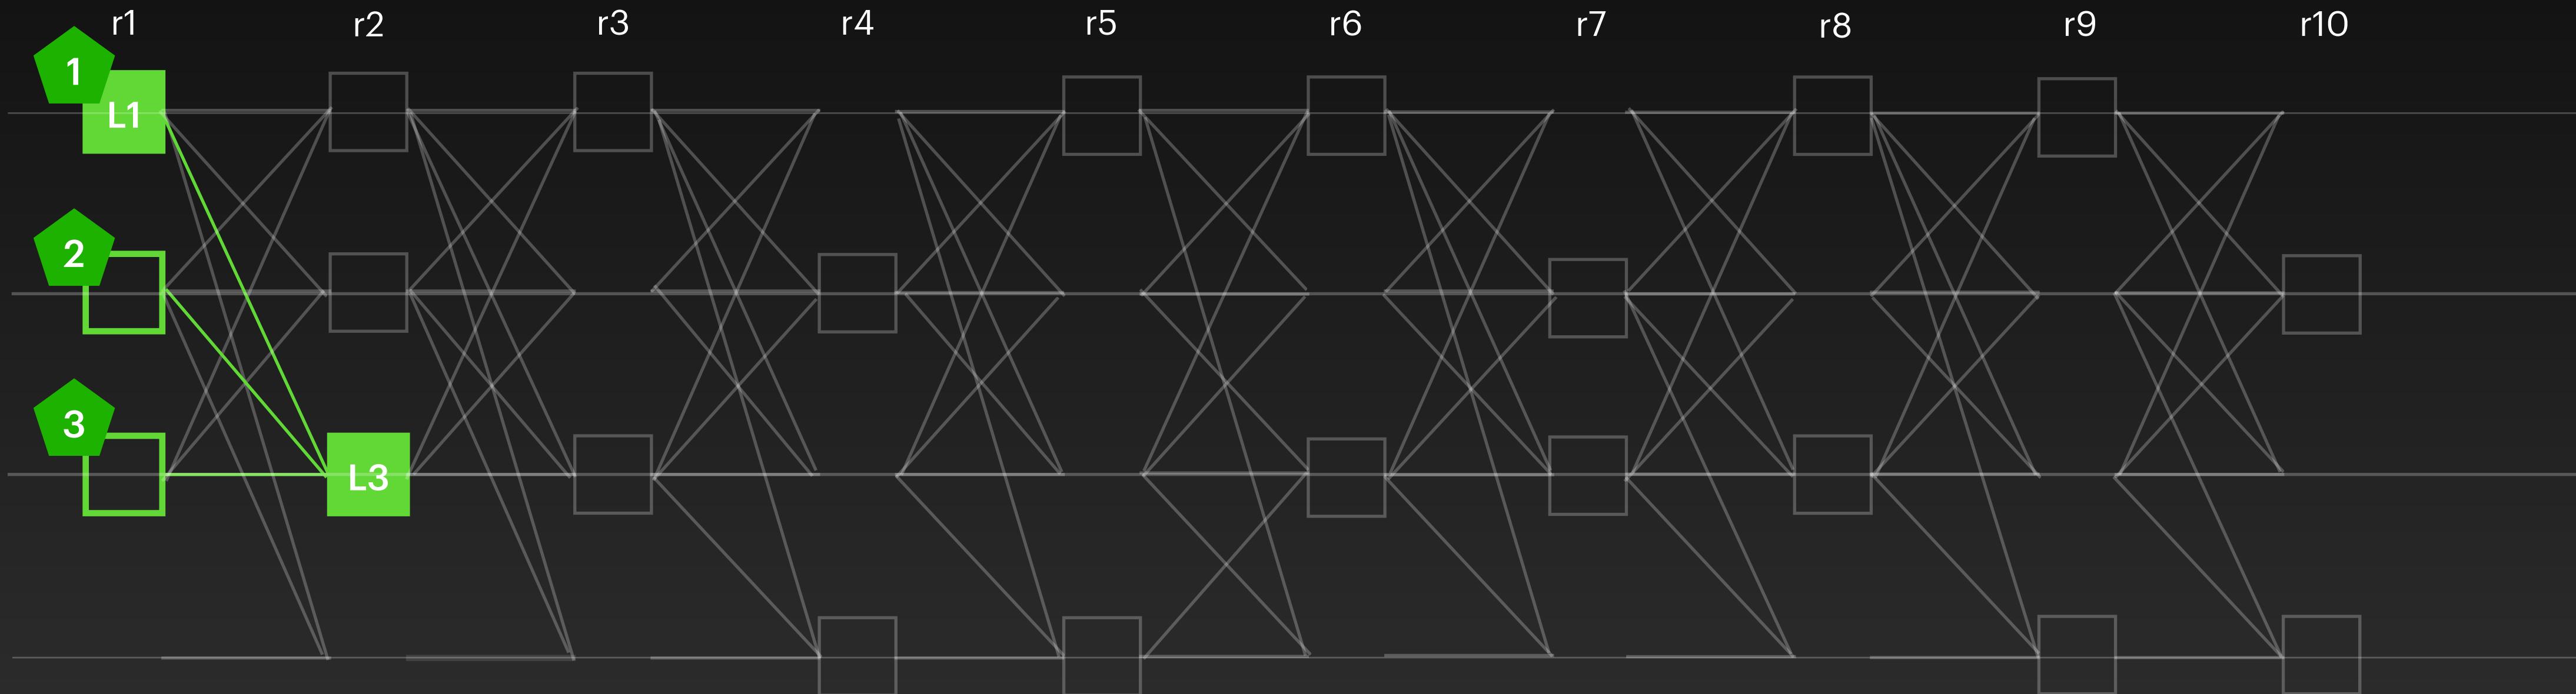
L3

output sequence:

L1

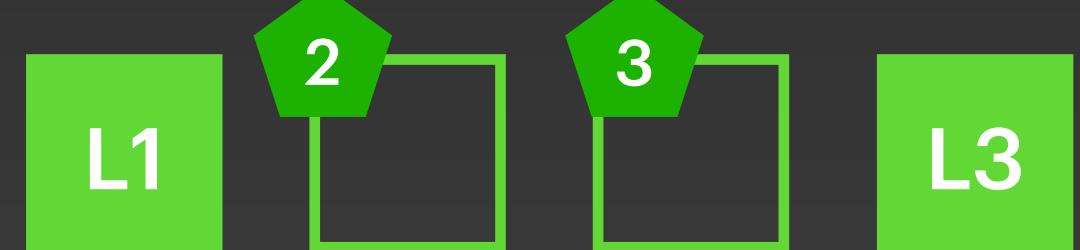
Commit Sequence

Commit sub-dag



leaders sequence:

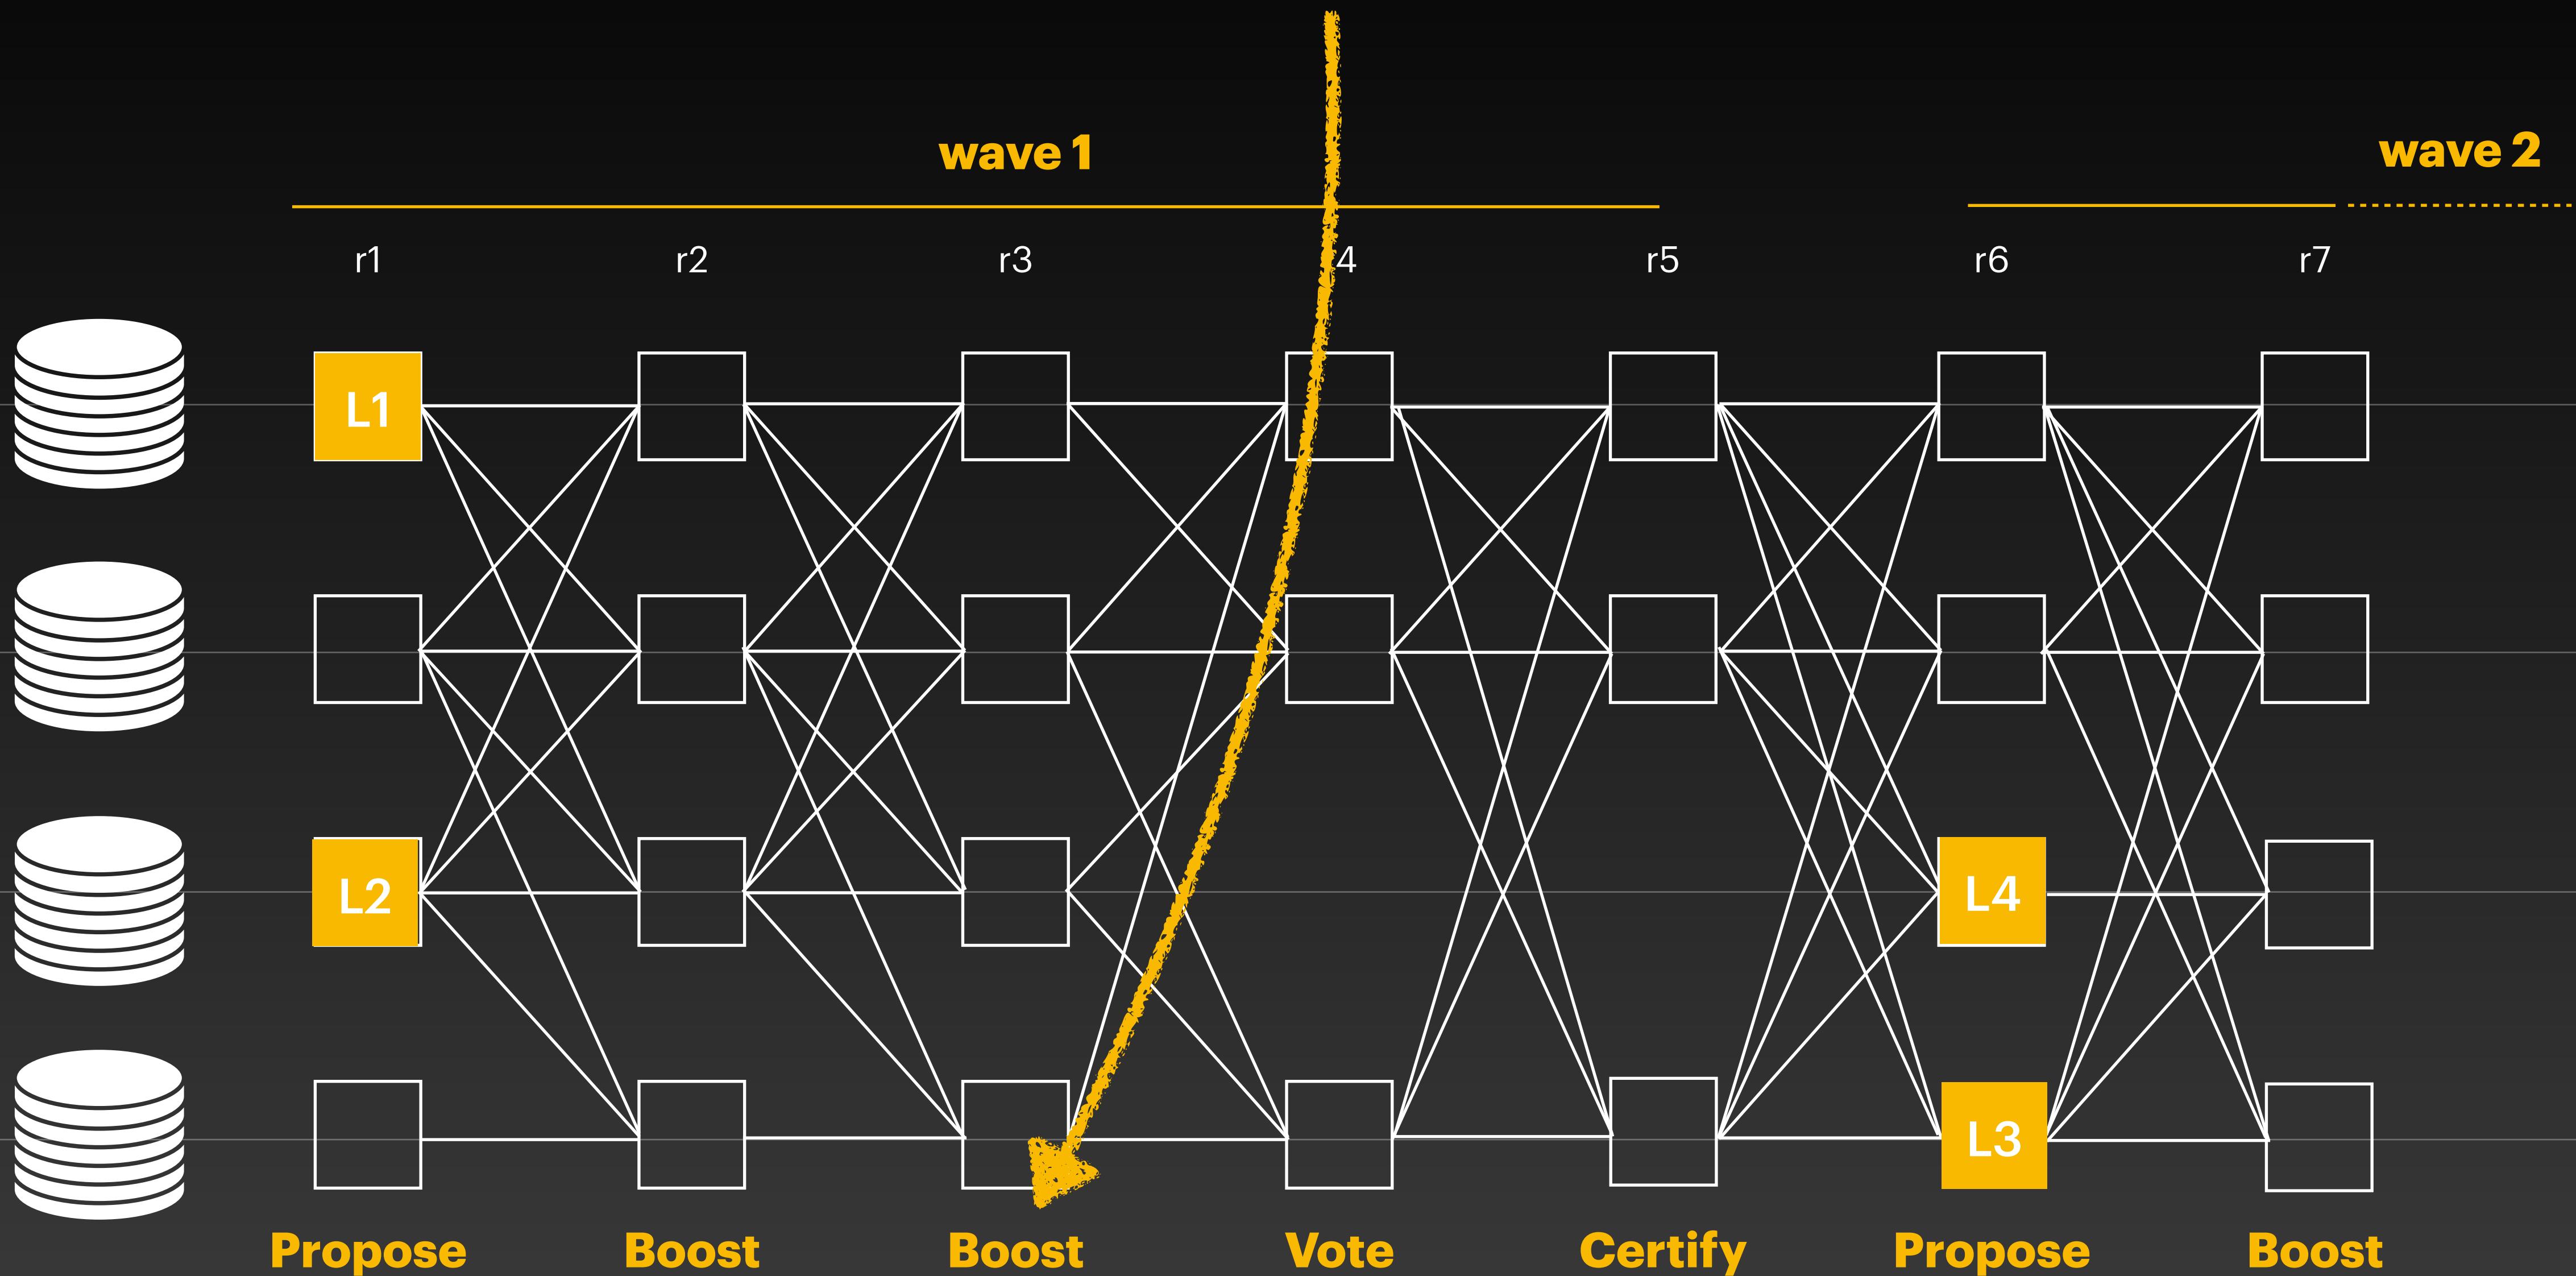
output sequence:



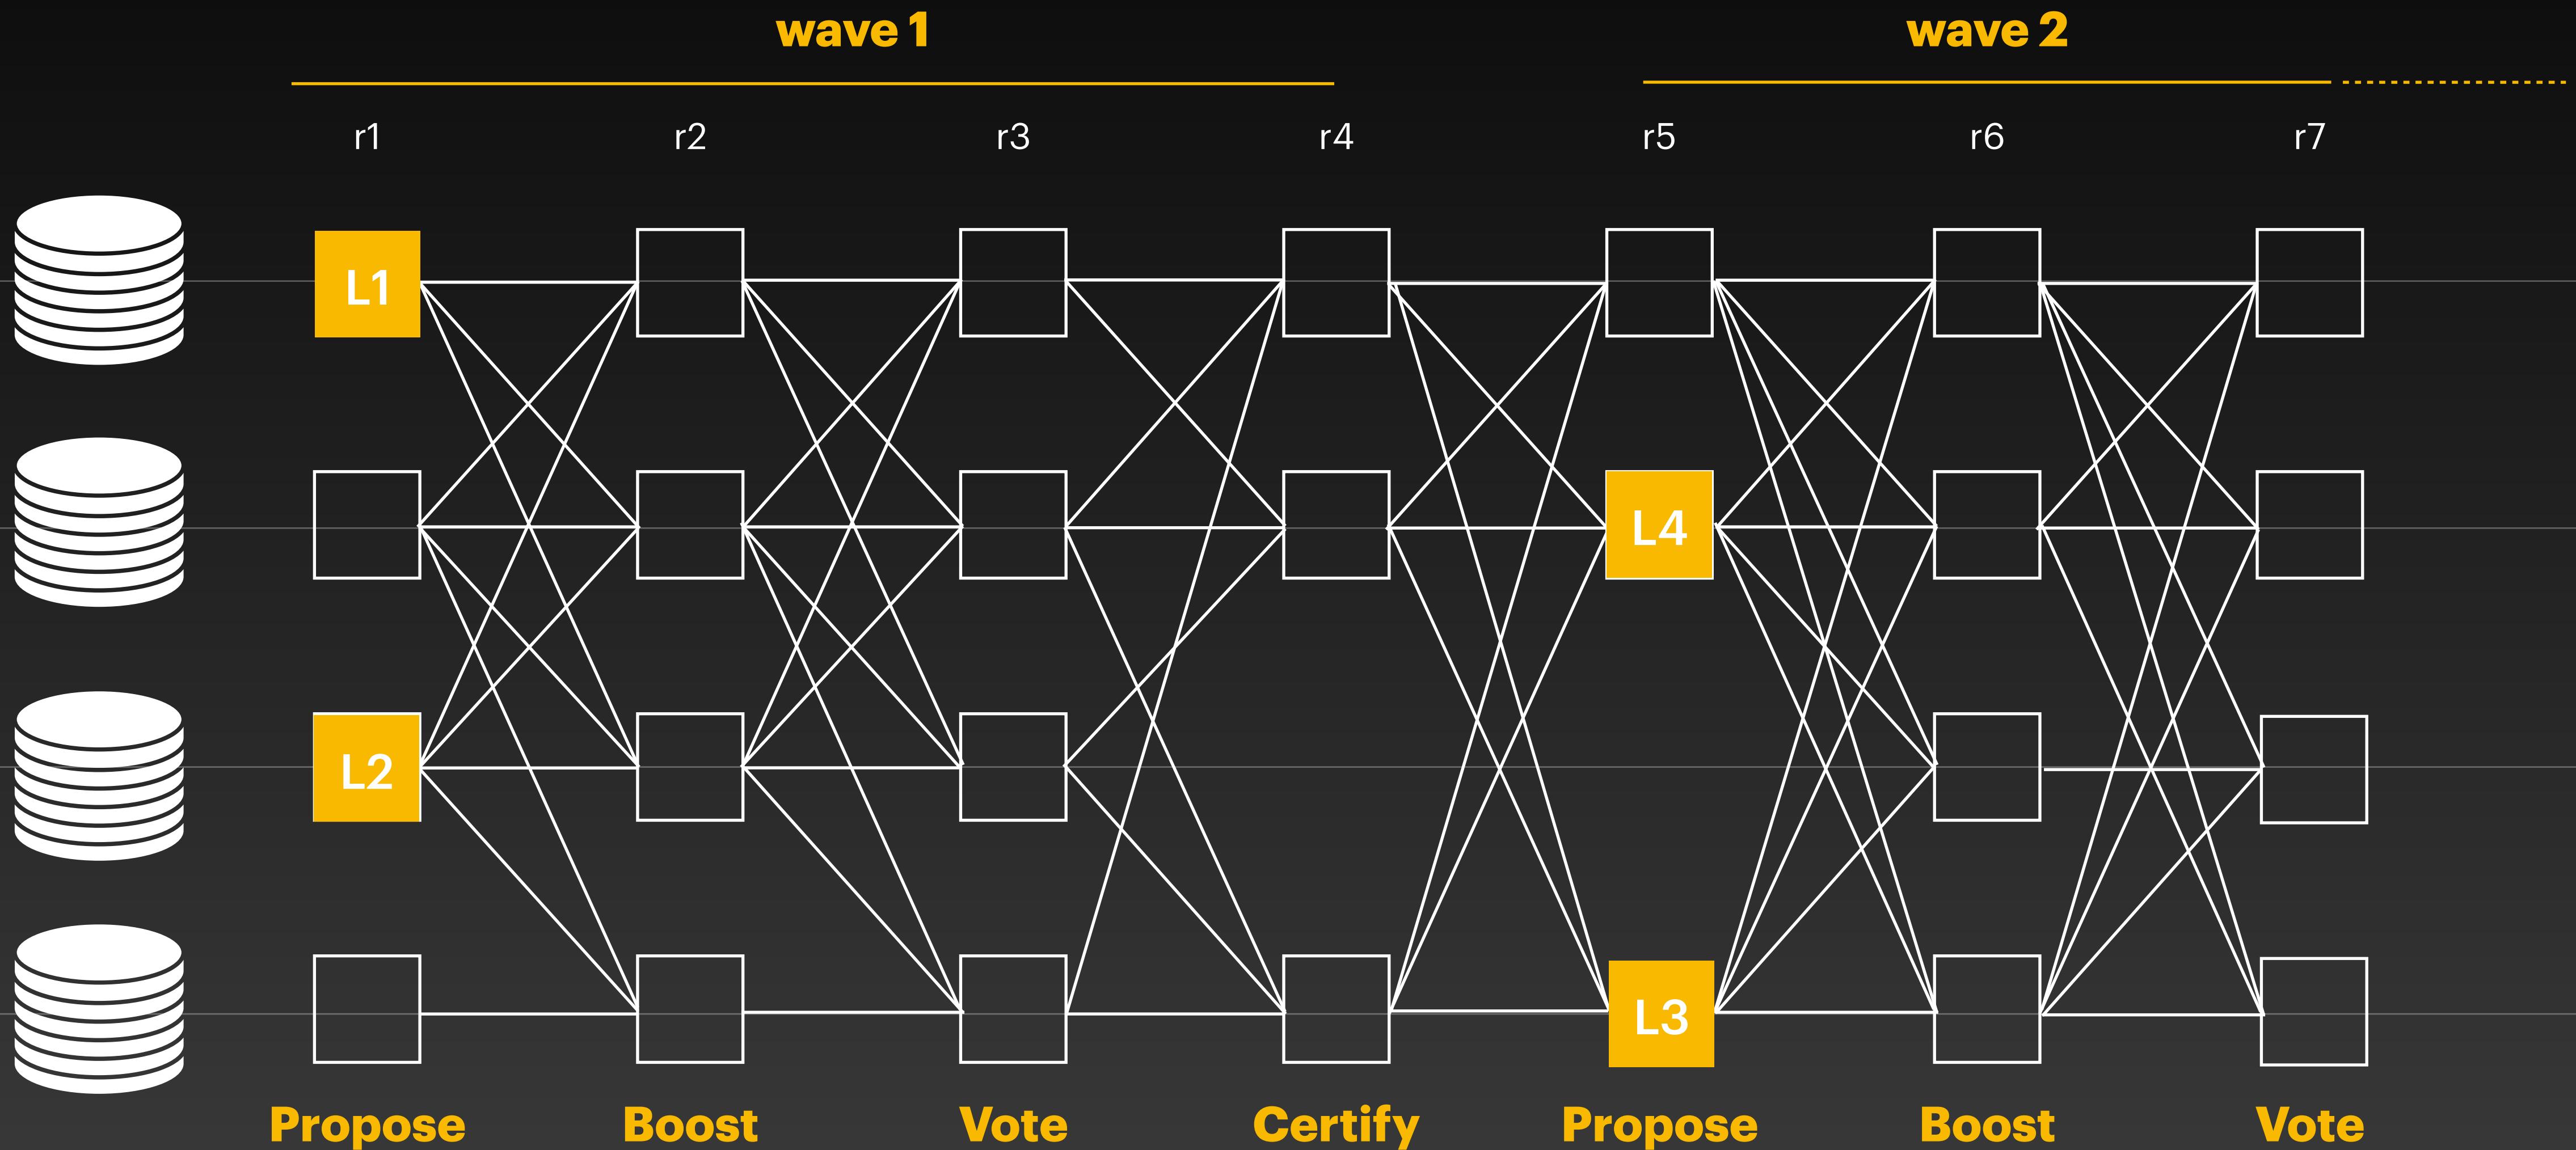
Parameterizing Mahi-Mahi

Number of Leaders + 4 or 5 Rounds per Wave

We can reduce the number of rounds per wave



Mahi-Mahi 4: 4 rounds per wave



Mahi-Mahi parameters

Select Number of Rounds and Number of Leaders per Round

5 rounds

- More round costs involved
- Higher probability of commits
- Maintains all security properties

4 rounds

- Less round costs involved
- Lower probability of commits
- Maintains all security properties

- Need to fix the size of leaders per round
- Both maintain liveness

Benchmarks

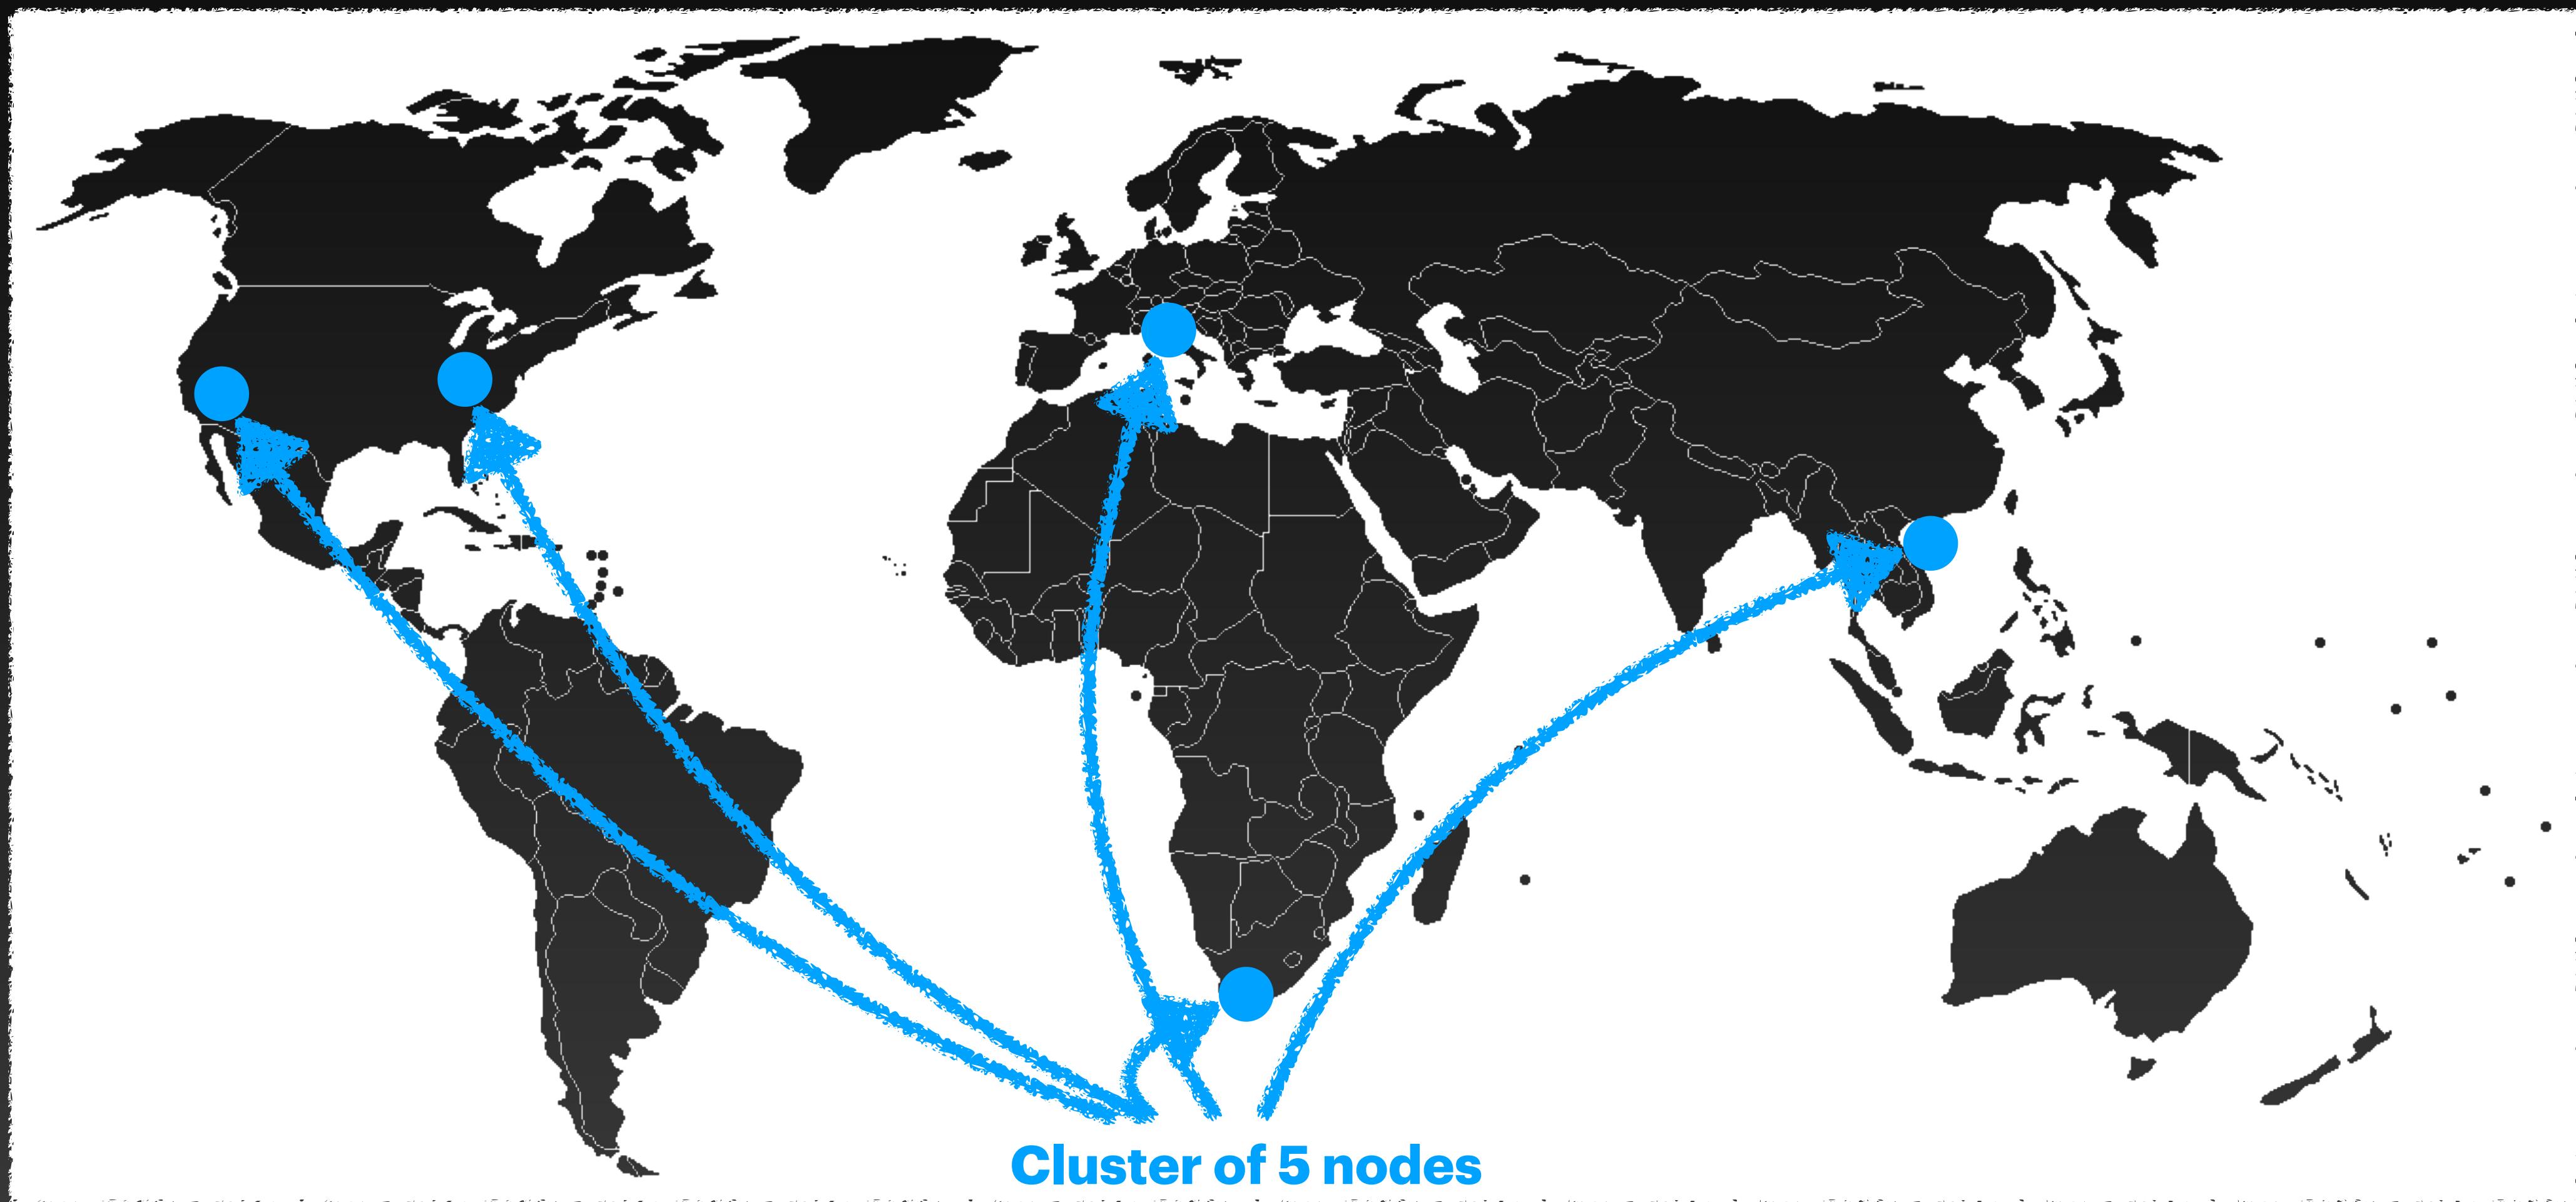
Implementation

- Written in Rust
- Networking: Tokio (TCP)
- Storage: custom WAL
- Cryptography: ed25519-consensus and blake2

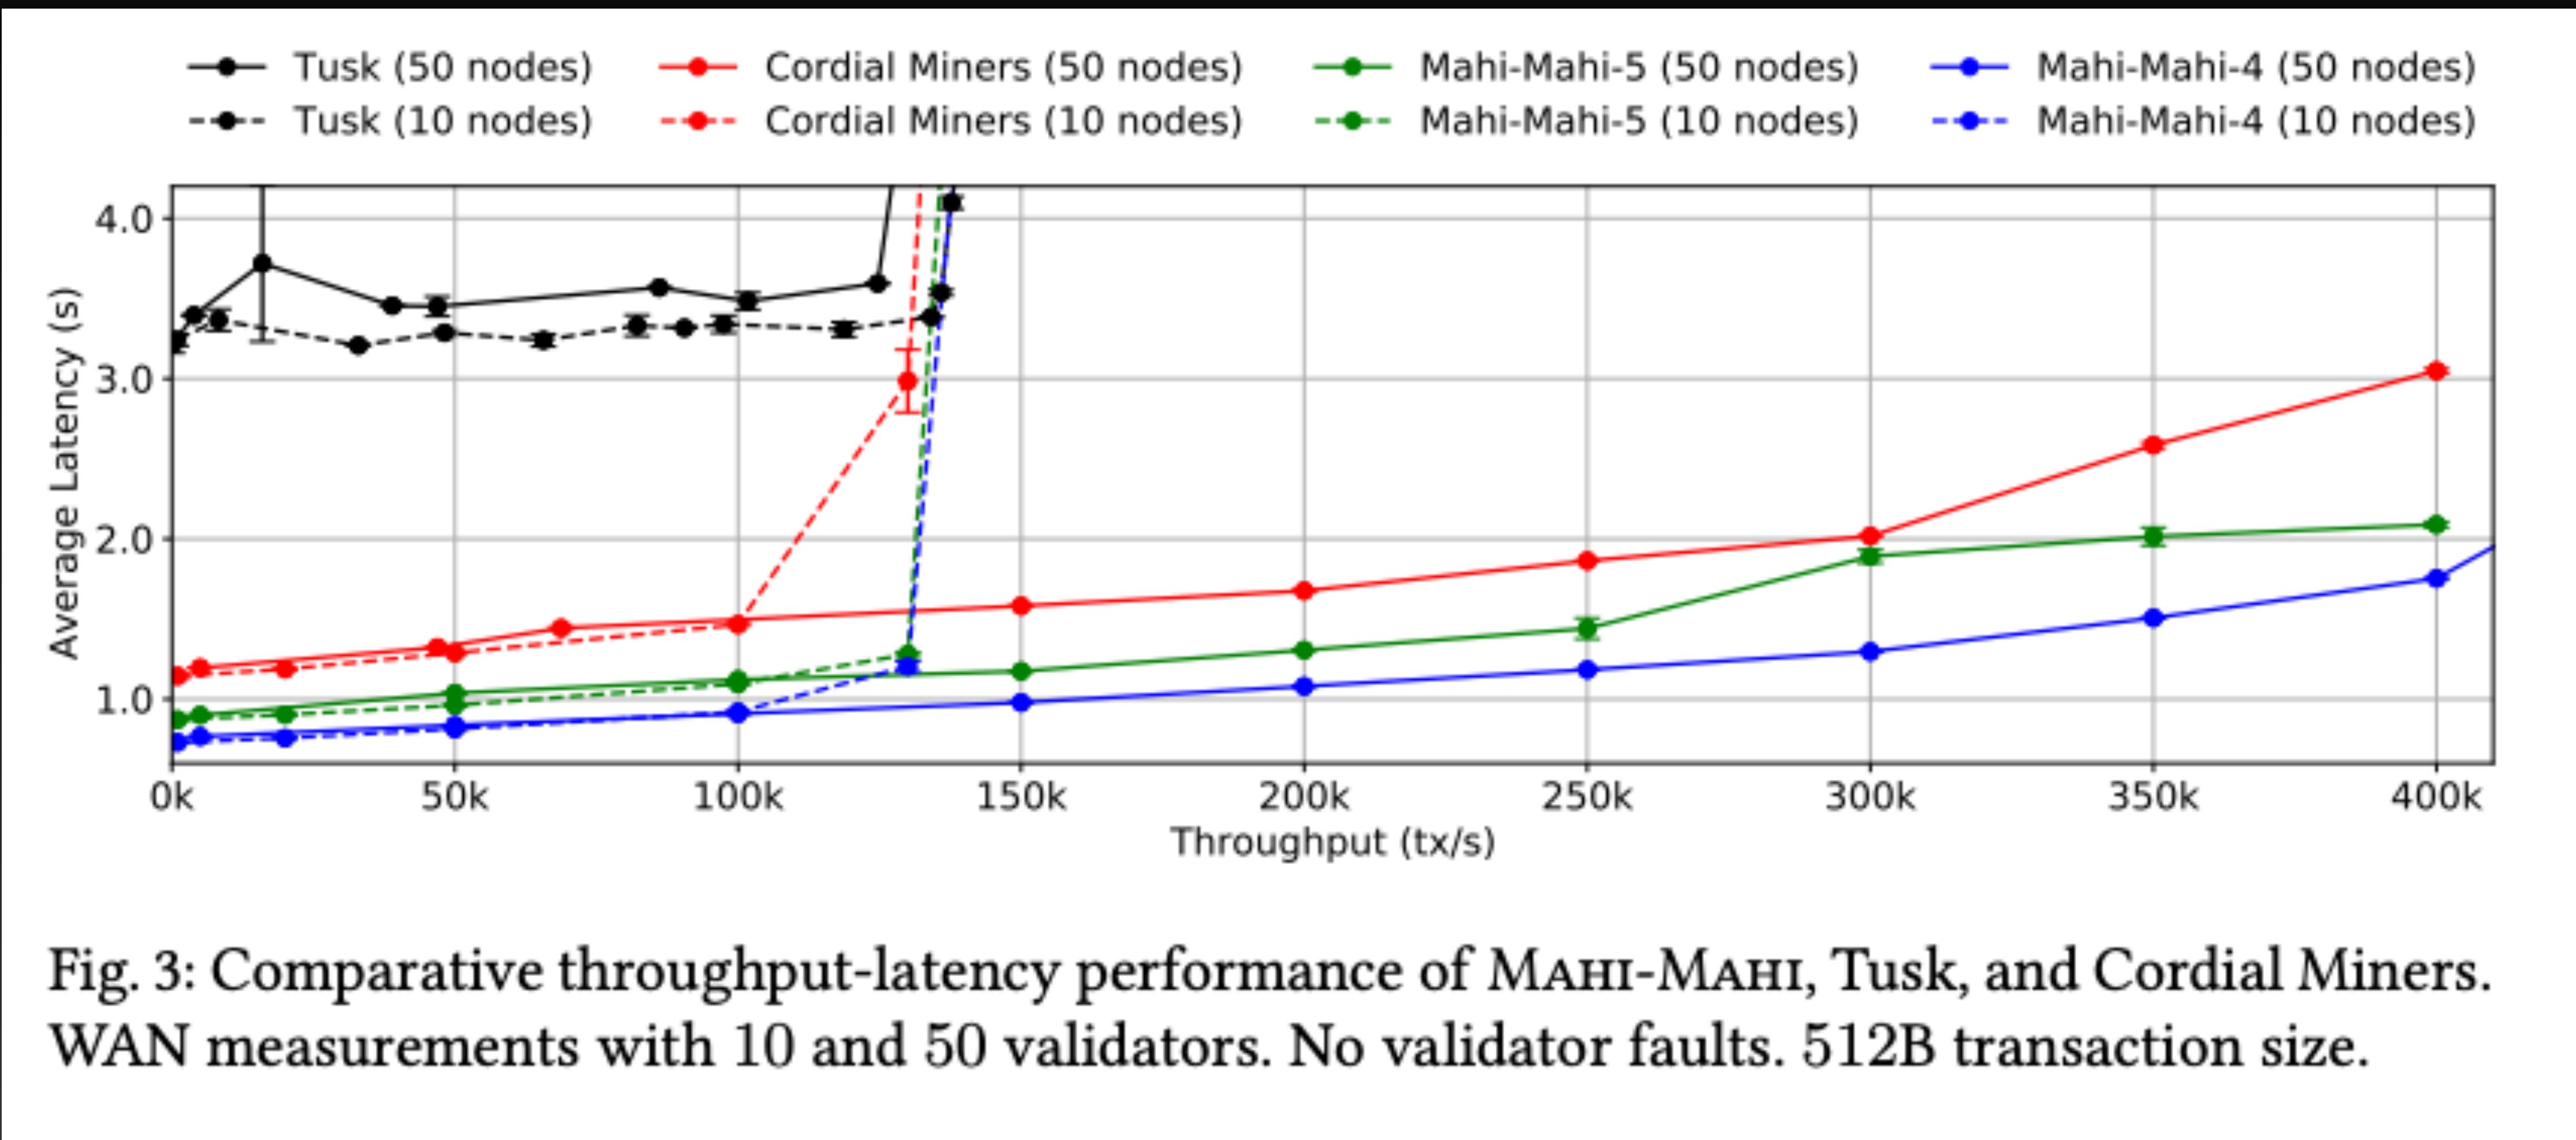
<https://github.com/PasinduTennage/mahi-mahi-consensus>

Evaluation

Experimental setup on AWS



No Fault Performance



Faulty Nodes Performance

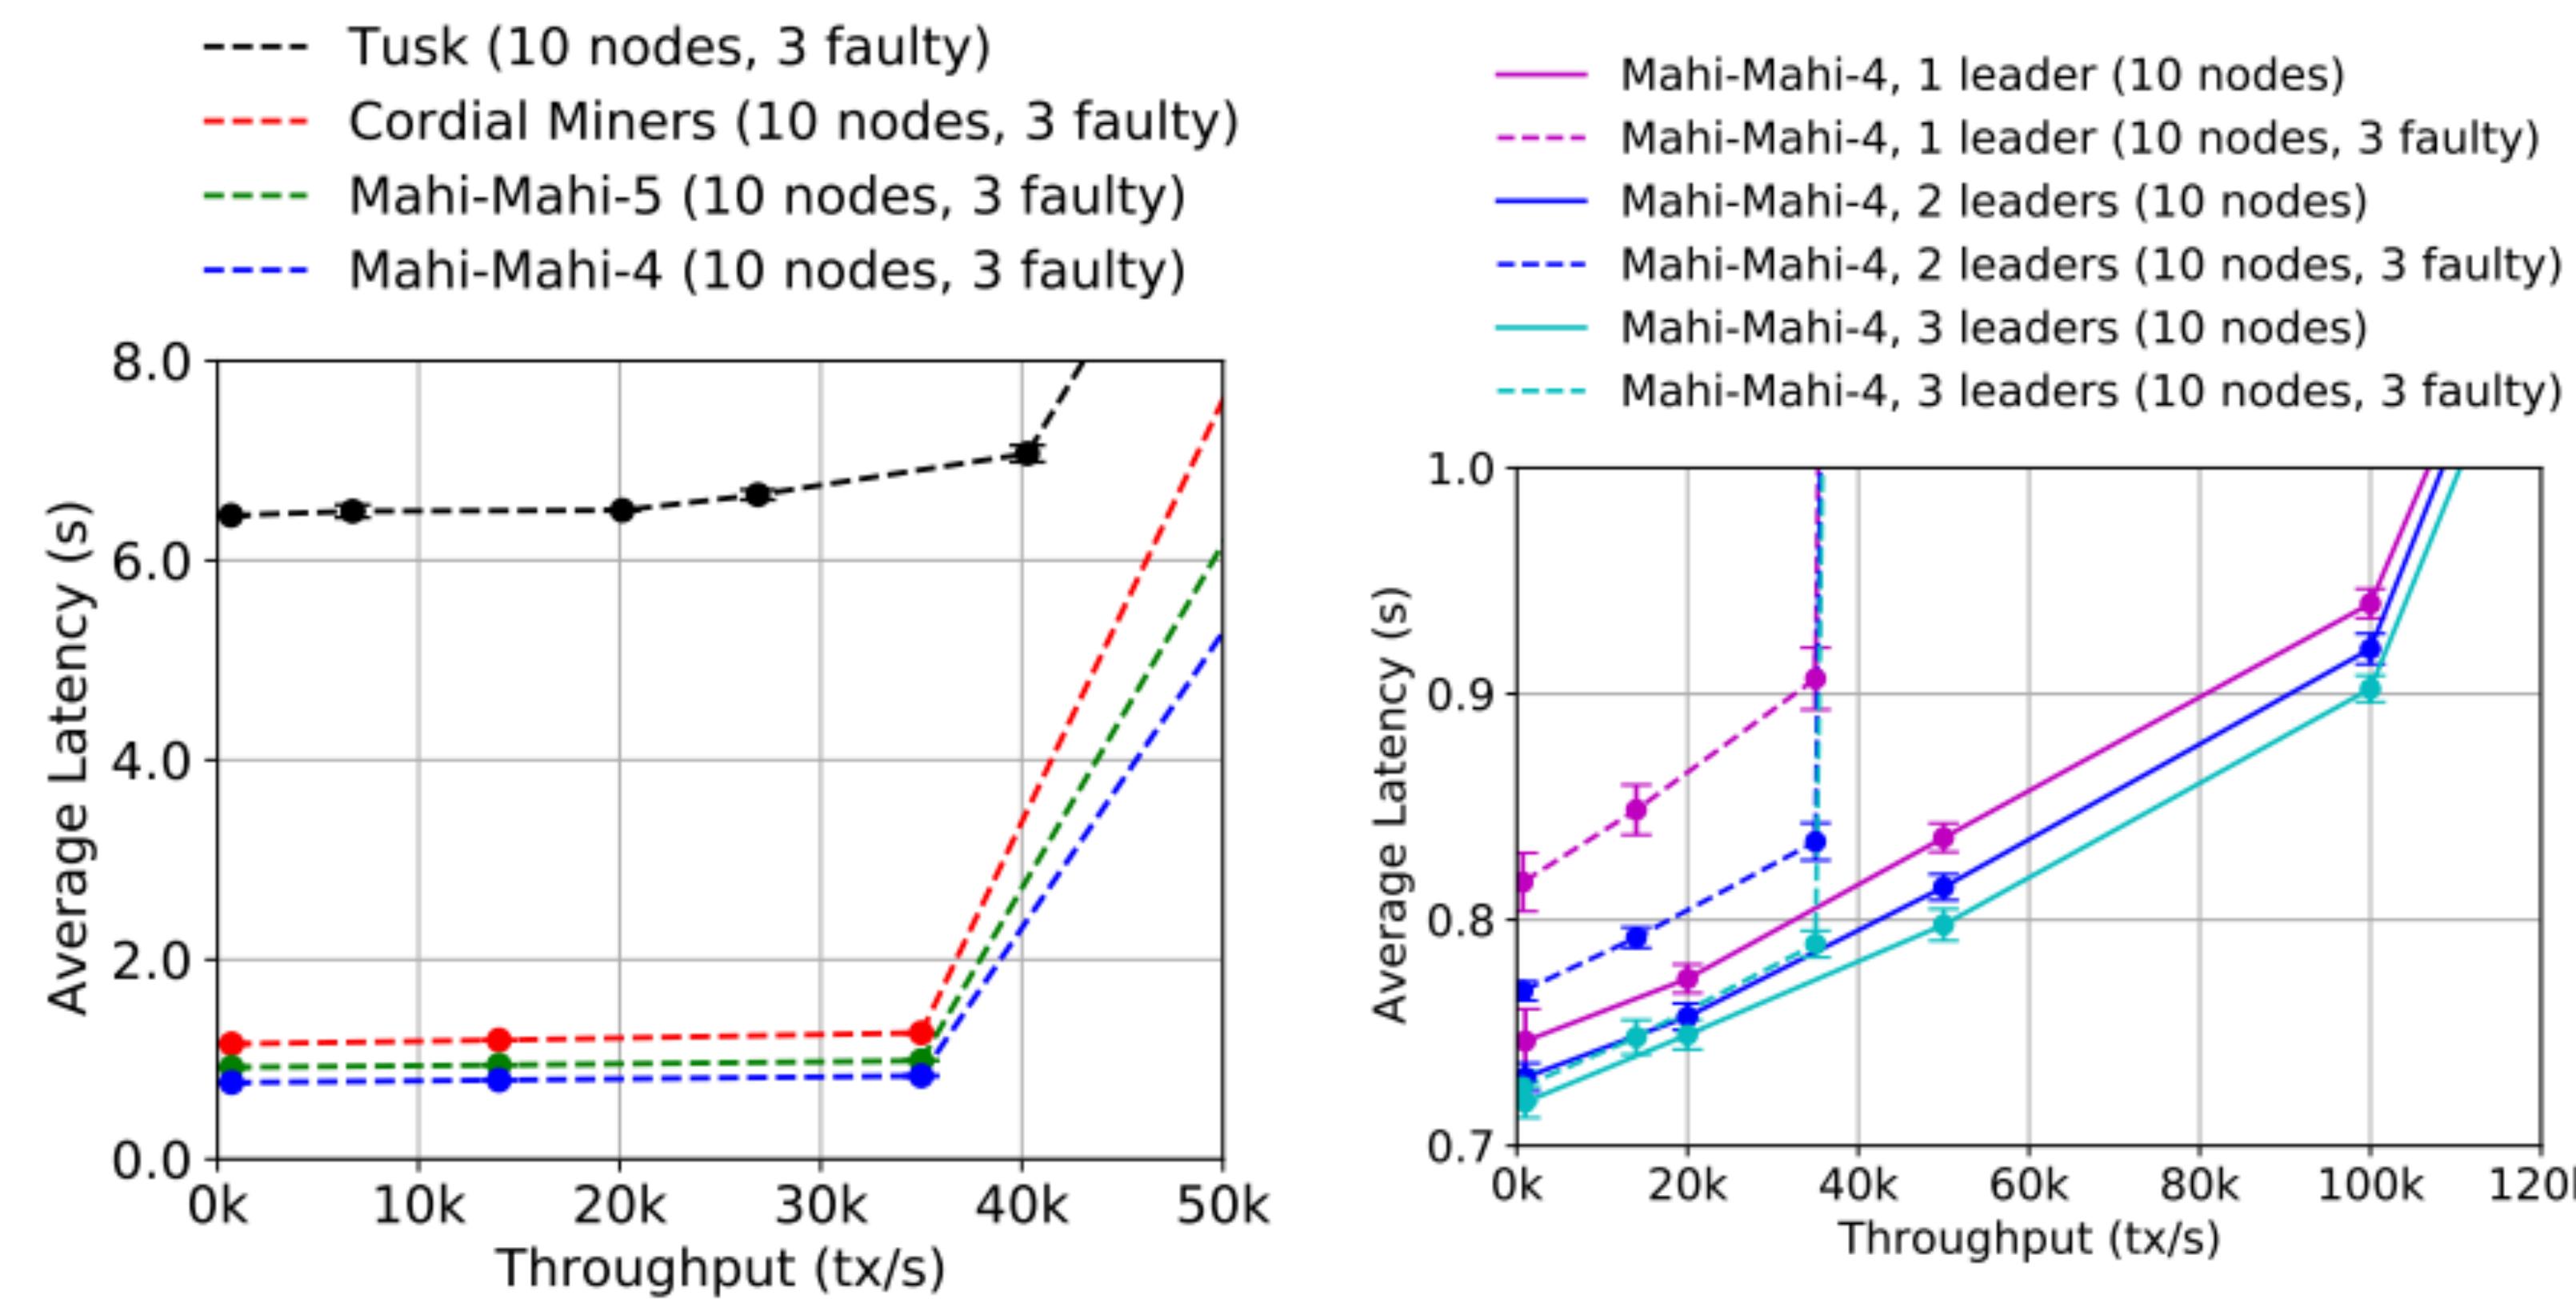
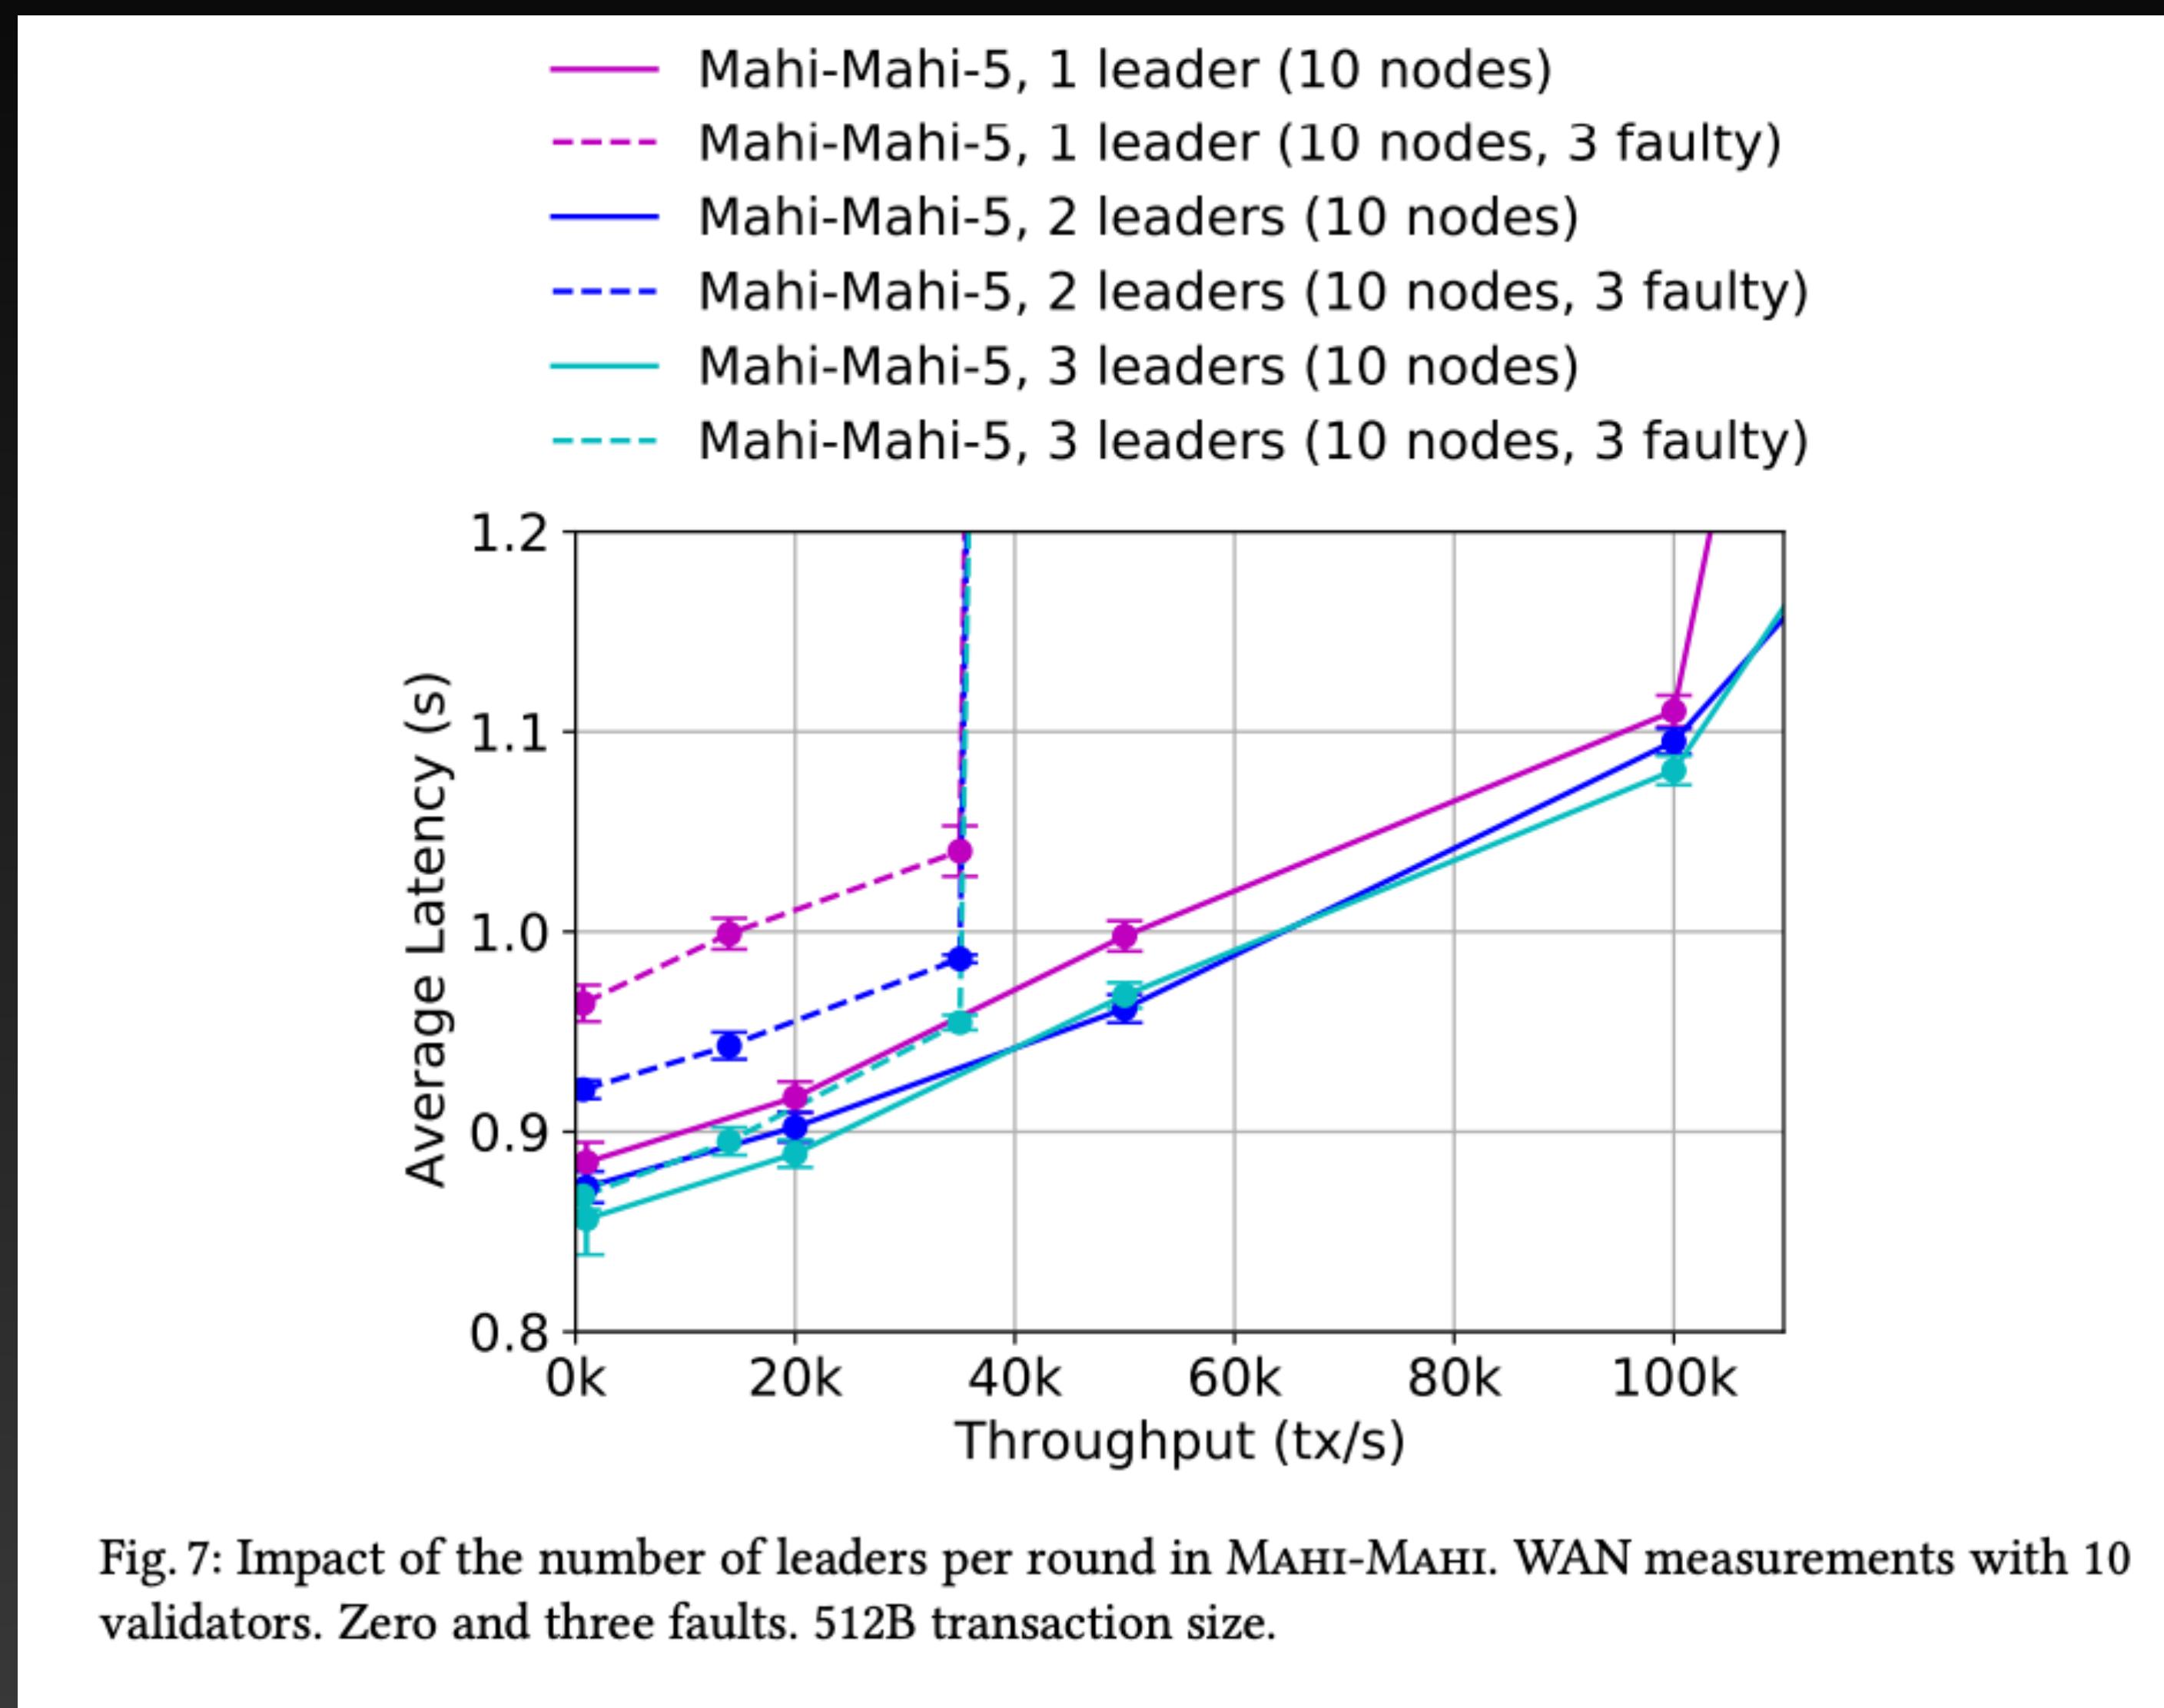


Fig. 4: Comparative throughput-latency of MAHI-MAHI, Tusk, and Cordial Miners. WAN measurements with 10 validators. Three faults. 512B transaction size.

Fig. 5: Impact of the number of leaders per round in MAHI-MAHI. WAN measurements with 10 validators. Zero and three faults. 512B transaction size.

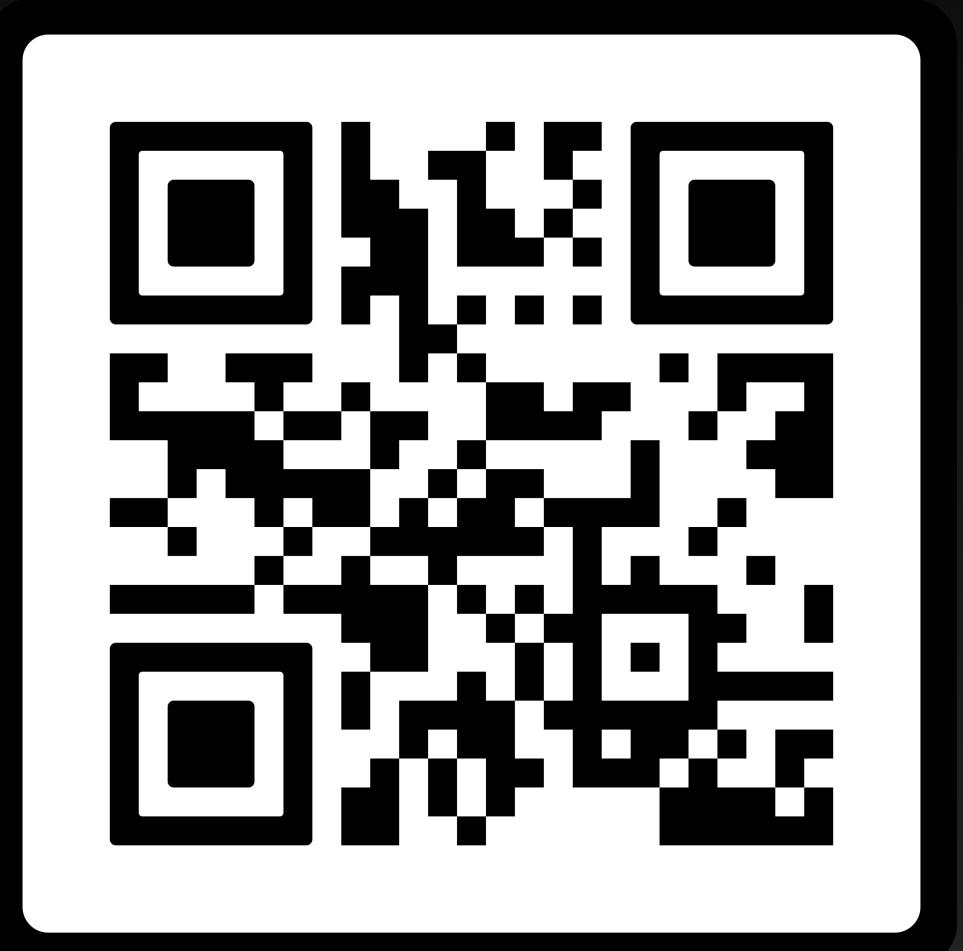
Various Mahi-Mahi Performance



Summary

Mahi-Mahi

- Low-latency asynchronous protocol
 - 100k tps with sub second latency
-
- **Paper:** <https://www.arxiv.org/abs/2410.08670>
 - **Code:** <https://github.com/PasinduTennage/mahi-mahi-consensus>



SCAN ME

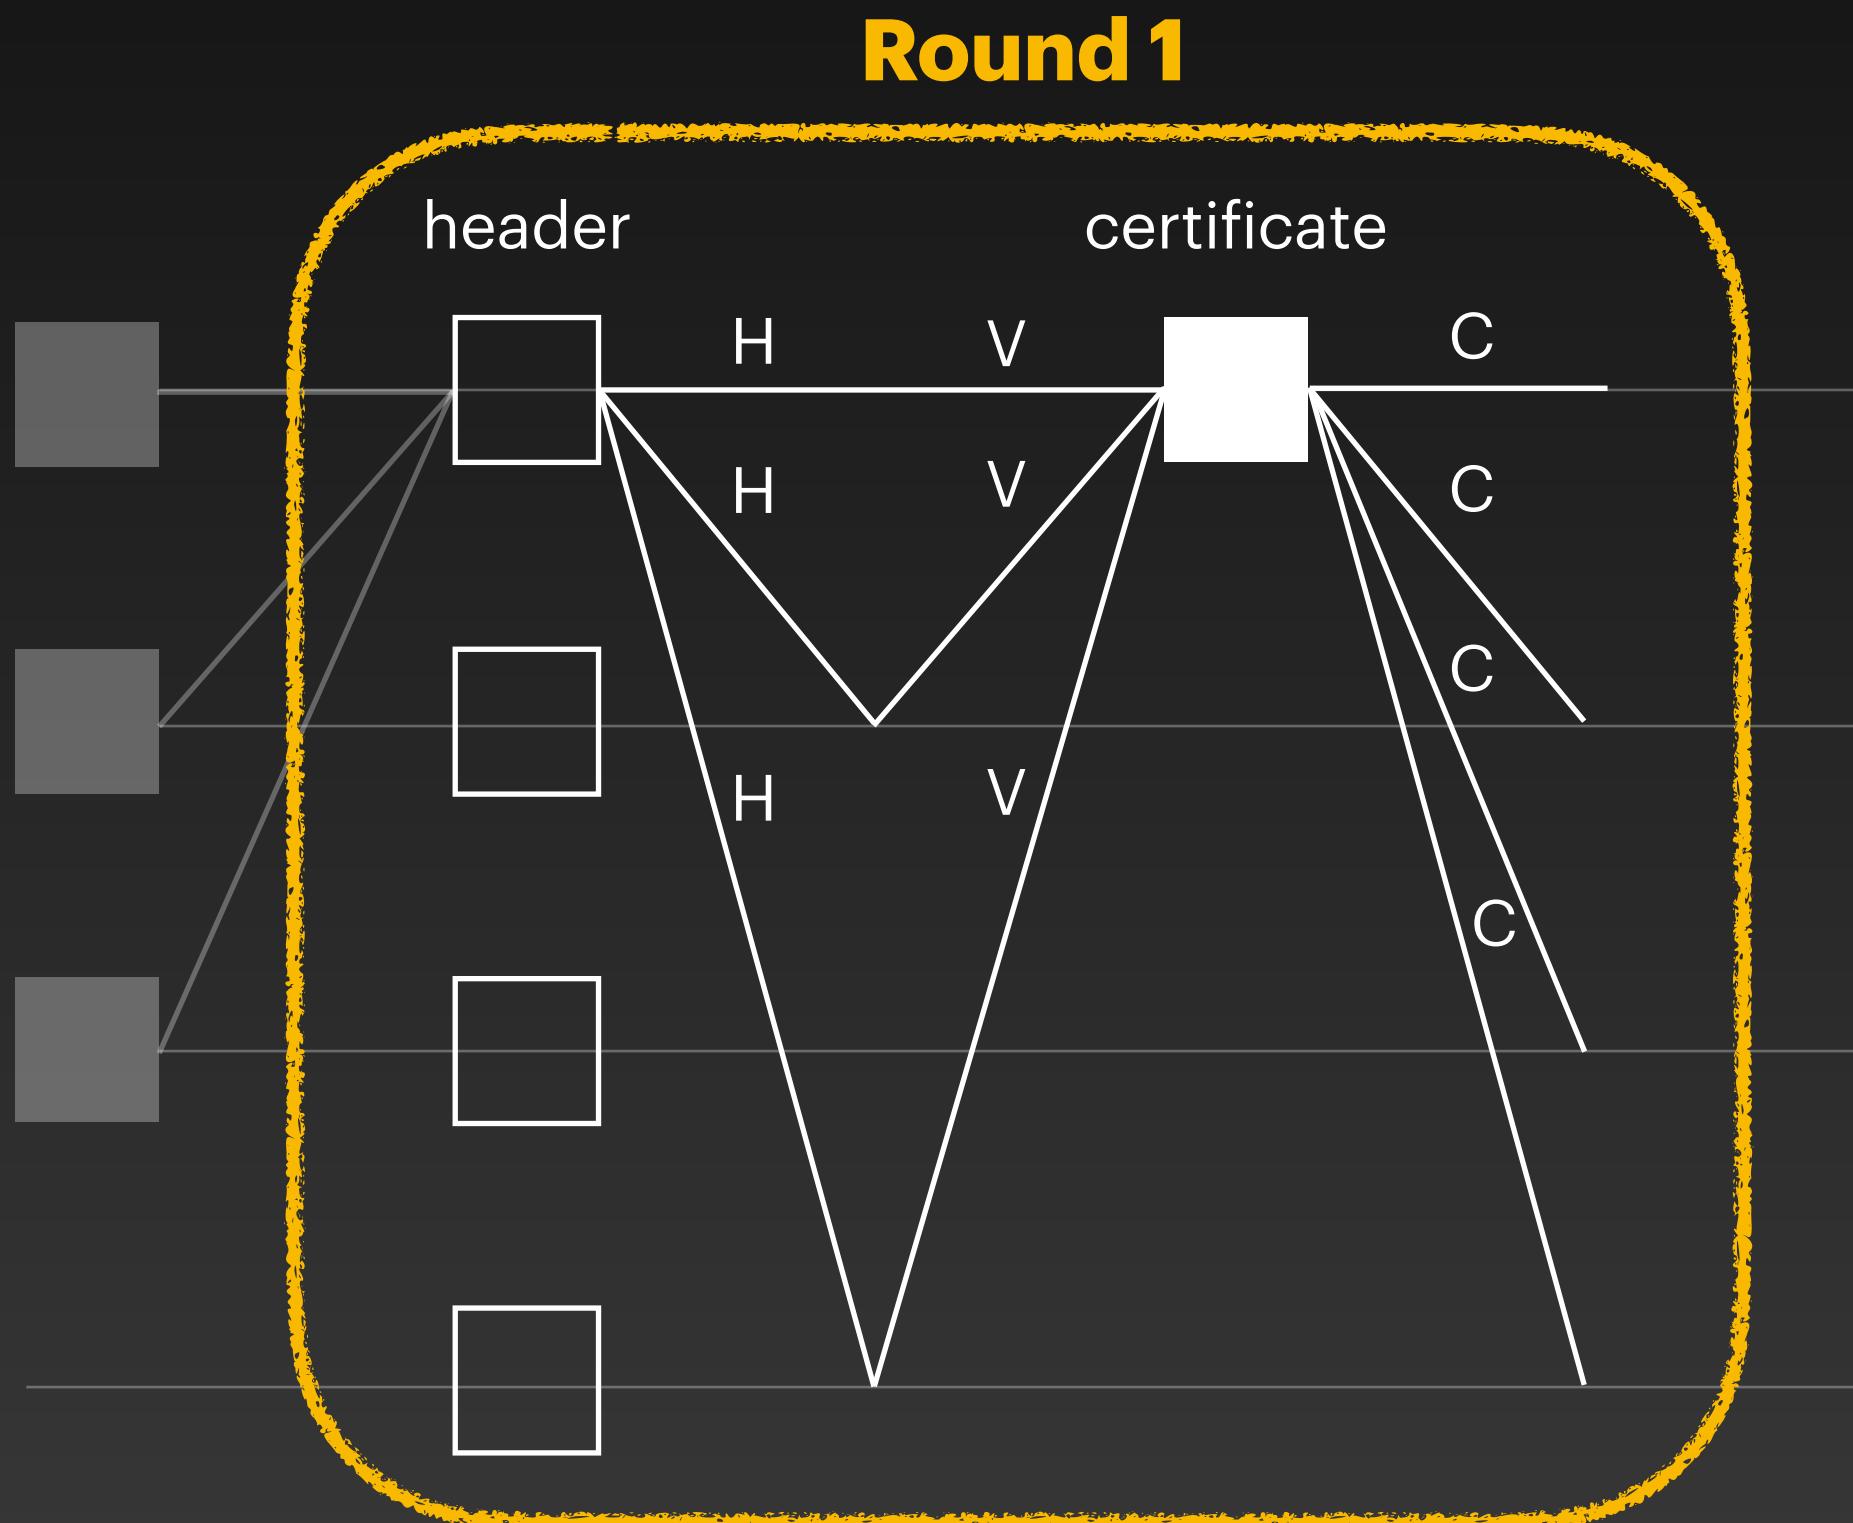
Narwhal vs Mahi-Mahi

Key differences & Insight

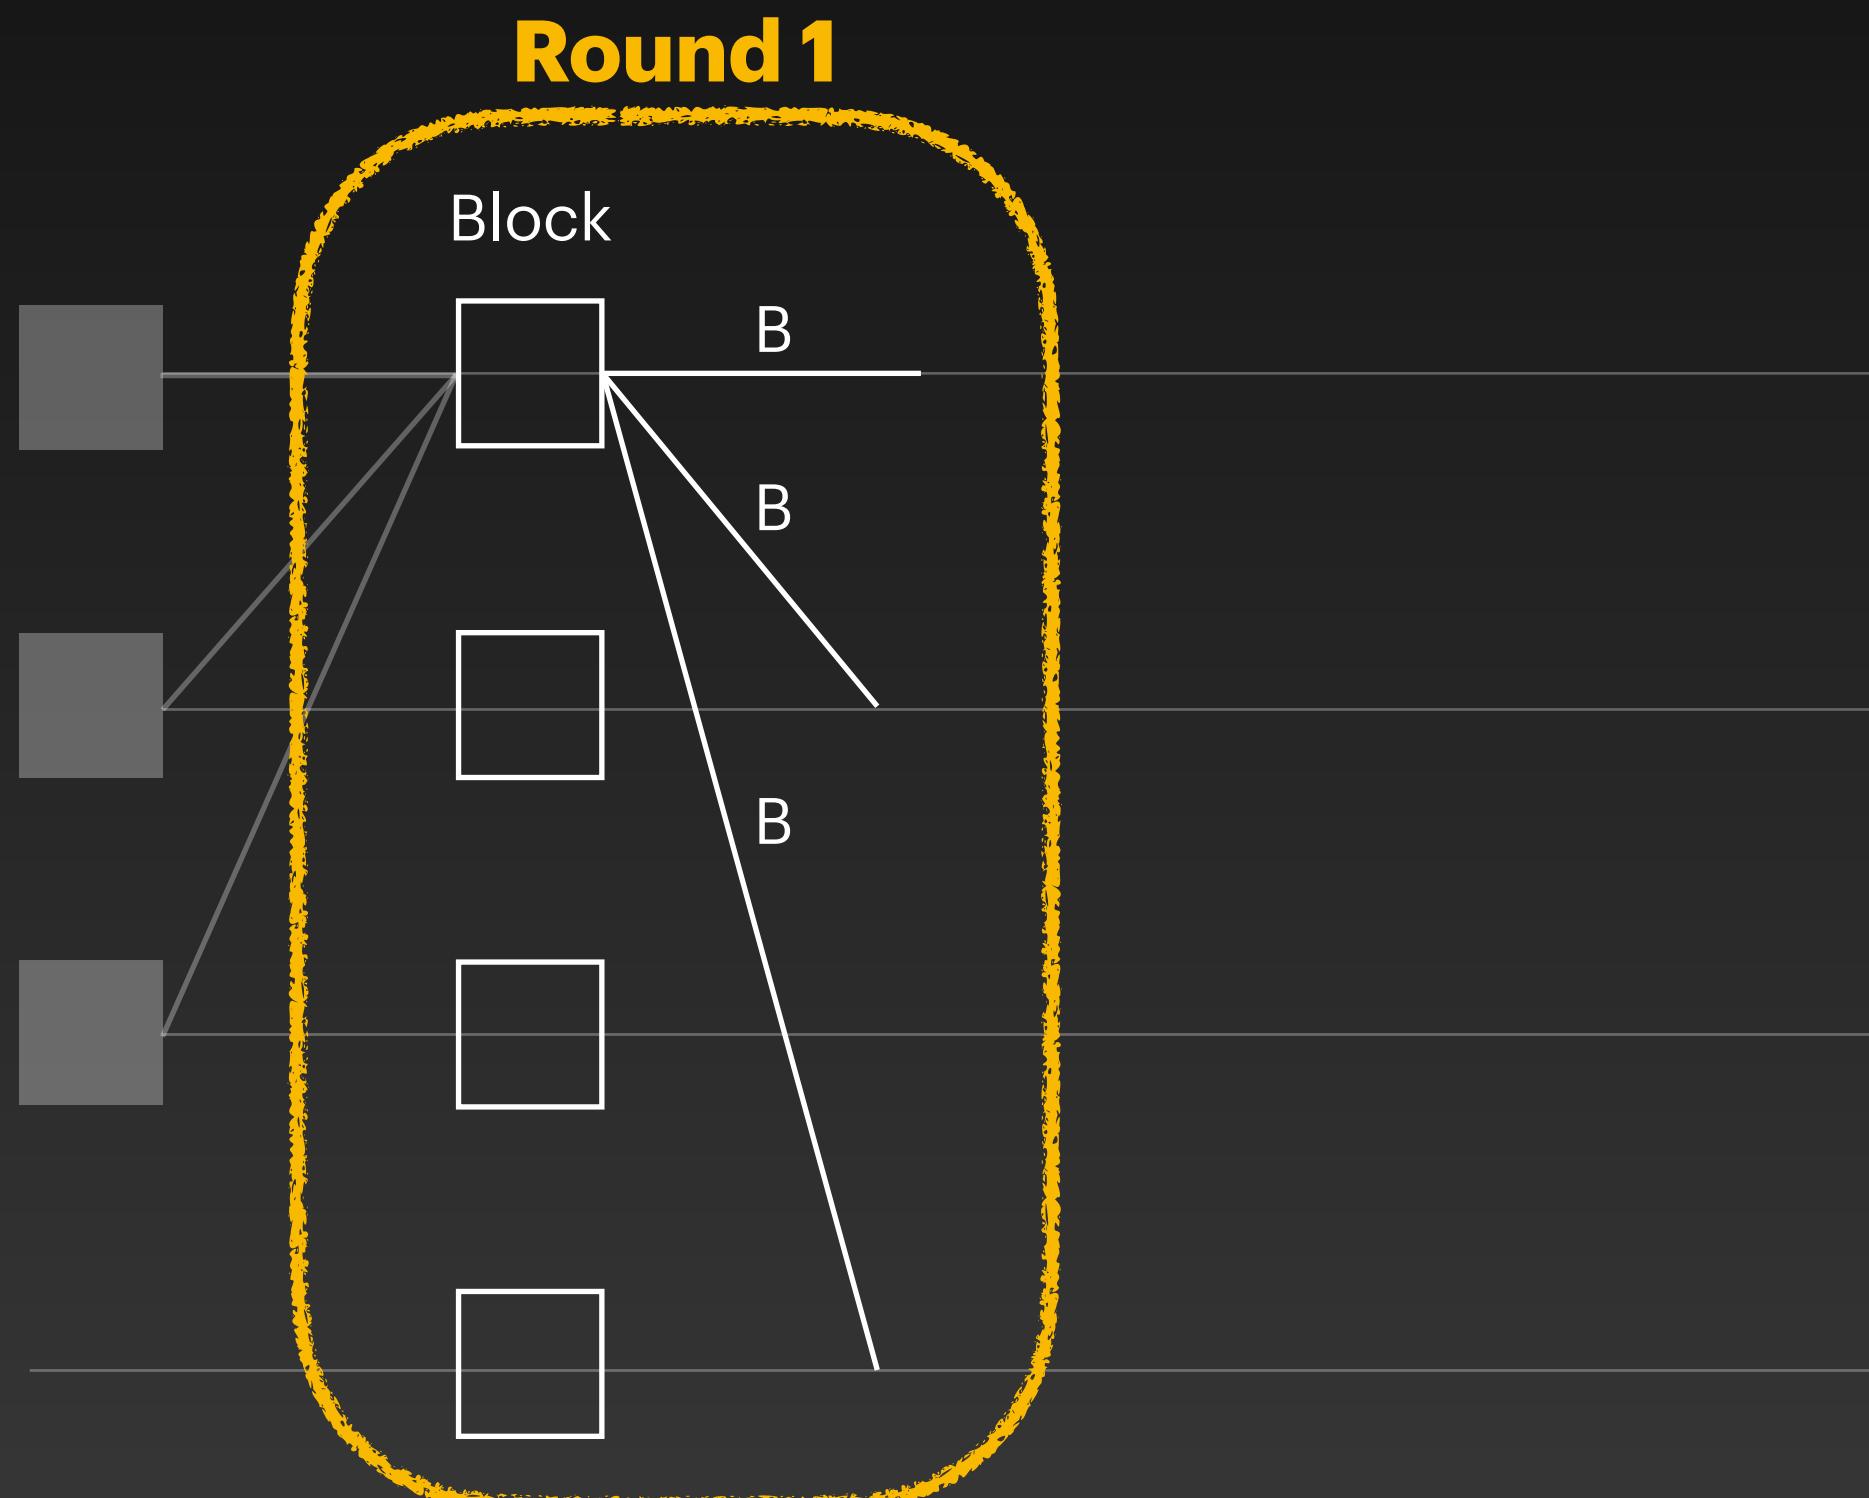
Generally kept as extra slides, anticipate CM being asked?

Narwhal vs Mysticeti

Narwhal

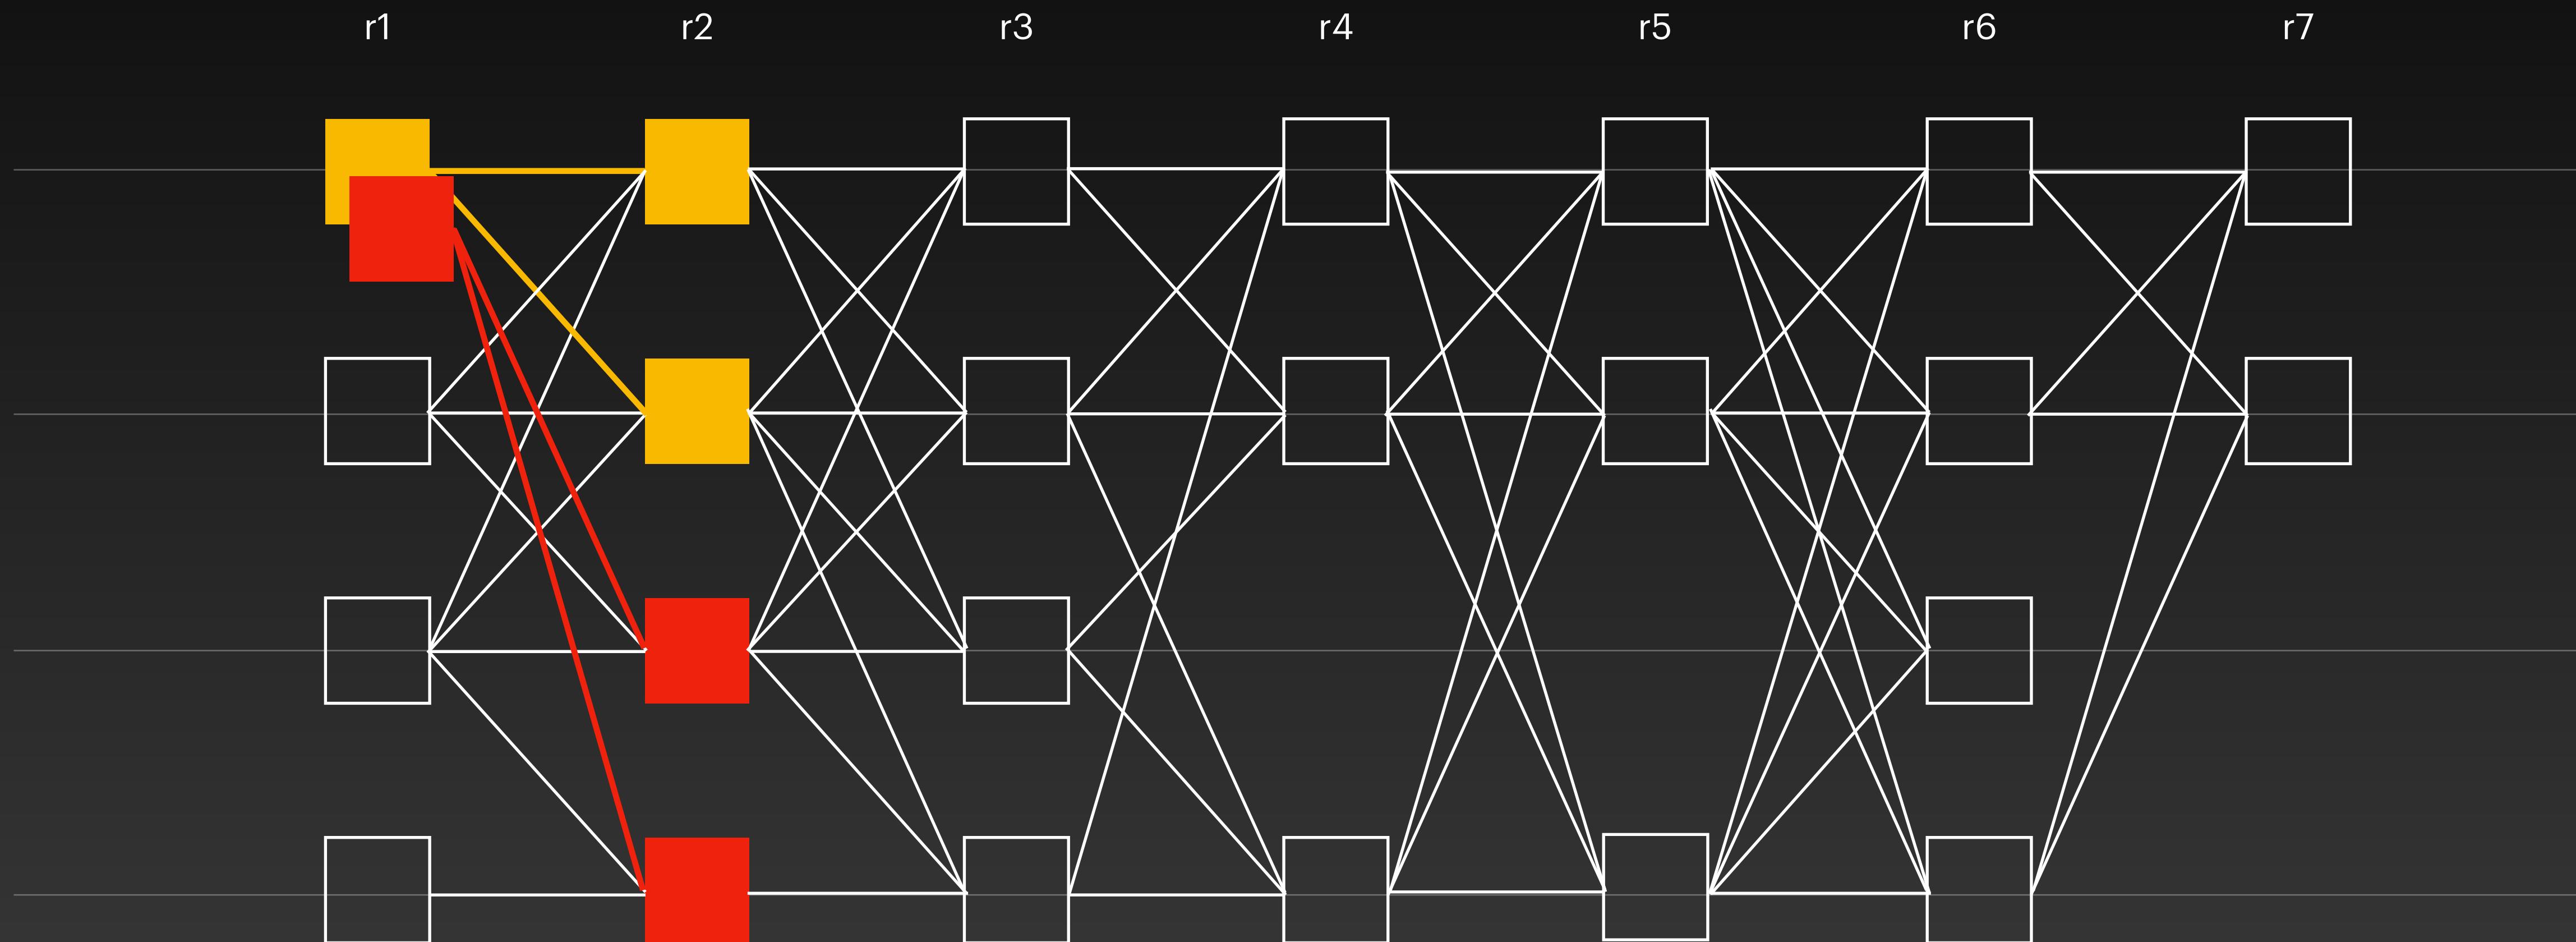


Mahi-Mahi



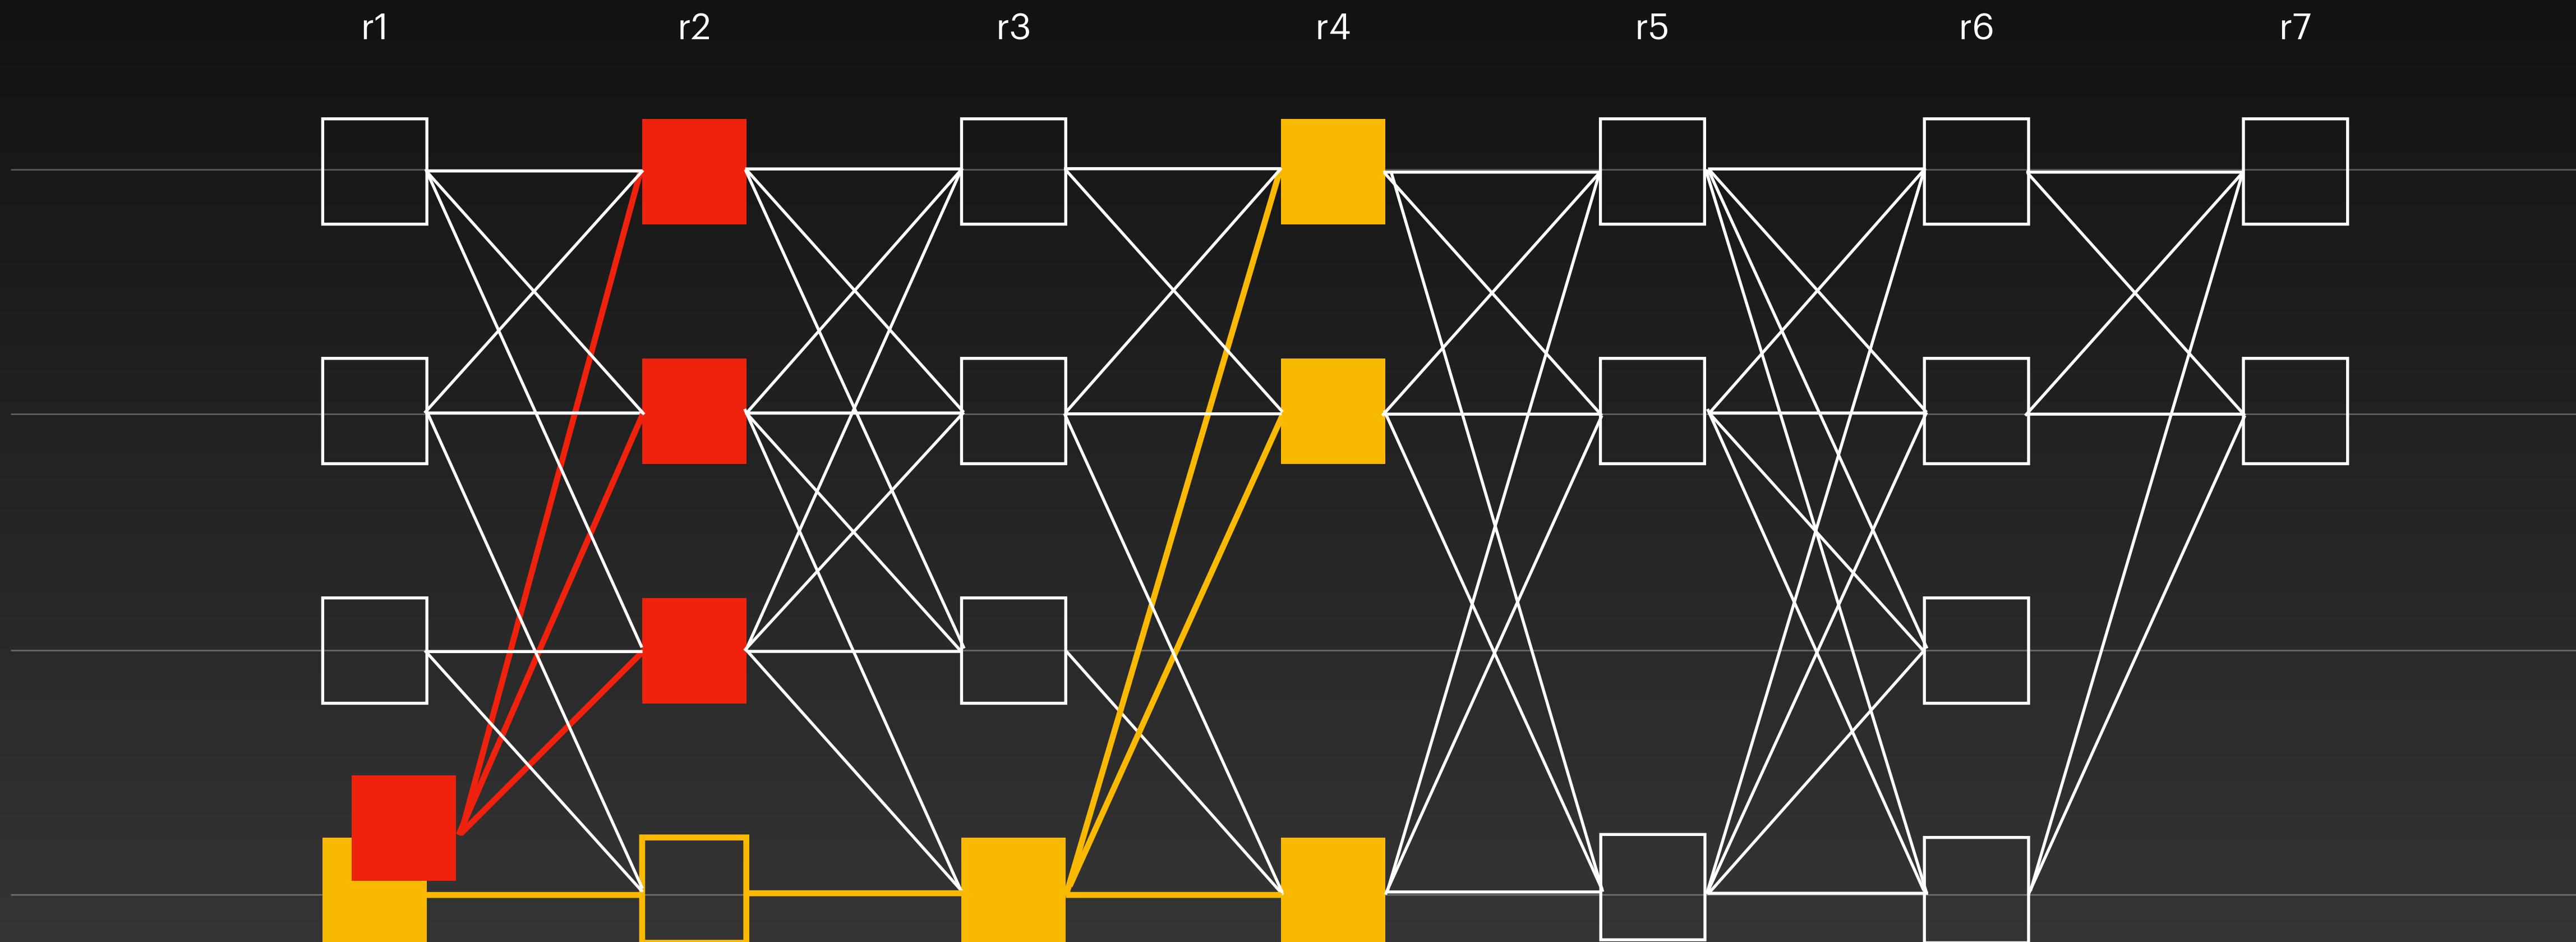
Main Challenge

Possible equivocations



Main Challenge

Possible equivocations (even with $2f+1$ support)



Decision Rules

Upon interpreting the DAG...

Bullshark

- A leader is **Commit** or not
- Either directly or indirectly
(recursion)

Mahi-Mahi

- A leader is **Commit**, **Skip**, or
Undecided
- Either directly or indirectly
(recursion)