

# Elastic Scaling Web3

Lefteris Kokoris Kogias — Myster Labs

# Elastic Scaling

- **Elastic scalability** is the ability of a system to dynamically adjust its resource usage based on workload demands.

# Example

## Video Streaming Lectures

Normal Days



# Example

## Video Streaming Lectures

Exam Period




- Video is lower quality
- Students disconnect

# Example

## Video Streaming Lectures

Morning Before the Exam




- No Service
- System overload

# Elastic Scaling

## Video Streaming Lectures

Exam Period




- Add Resources as Demand Increases
- Cost per User remains constant

# Elastic Scaling

## Video Streaming Lectures

Morning Before the Exam




- Smooth User Experience
- Better load balance even in the presence of faults

# Key Components For Elastic Scaling



**Autoscaling:** Automatically adjust the number of compute resources based on workload demands



**Load Balancing:** Distributes incoming traffic across multiple instances to ensure optimal resource utilization and performance.

# Benefits of Elastic Scaling



**Cost Efficiency:** Pay only for the resources used, minimizing idle capacity



**Performance:** Maintain consistent performance levels during peak and off-peak periods.

# The state of Web3


- Minimum Validator Requirements are high
  - Handle load spikes
  - High cost
  - Downward spiral when the load is low → **Increase fees or bankrupt**
    1. Invest in a powerful machine to be ready to handle spikes
    2. Load is low, but the cost of buying and running the machine is constant
    3. Need to charge more per transaction to break even
    4. The marginal utility of transaction drops as fees increase
    5. Load drops further

# The state of Web3

- When load is higher than the provisioned machine can handle
  - Fees and cost are no longer linked
  - It is an auction → **Pay the premium or leave**
    - Stable in the short term, but leads exit the ecosystem in the long term
  - Huge queuing delays → **Horrible UX**
    - Also leads to exit the ecosystem


# Sharding Blockchains — Design

"Omniledger: A secure, scale-out, decentralized ledger via sharding." IEEE S&P, 2018.




# Sharding Blockchains — Design

"Omniledger: A secure, scale-out, decentralized ledger via sharding." IEEE S&P, 2018.



# Sharding Blockchains — Properties

"Omniledger: A secure, scale-out, decentralized ledger via sharding." IEEE S&P, 2018.



- + Low Cost per Node
- + Scales-Out
- Fragmenting the state-space — Expensive Atomic Commit
- Susceptible to adaptive adversaries
- Security drop

# Sharding Blockchains – Challenges

“Divide and scale: Formalization of distributed ledger sharding protocols” SIROCCO, 2023





Table 1: Summarizing sharding protocol properties under our model

Protocol	Persistence	Consistency	Liveness	Scalability	Permissionless	S.-adaptive
Elastico	✓	✗	✓	✗	✓	✓
Monoxide	✓	✓	✓	✗	✓	✓
OmniLedger	✓	✓	✗	✓	✓	✓
RapidChain	✓	✓	✓	✓	✓	~
Chainspace	✓	✓	✓	✓	✗	✗


# First Step to the Solution

## Layering



# First Step to the Solution

## Layering




# Narwhal

Dag-based mempool

“Narwhal and tusk: a dag-based mempool and efficient bft consensus.” *EuroSys 2022*


# Narwhal

## The workers and the primary




# Narwhal

## The workers and the primary




# Narwhal

## The workers and the primary




# Narwhal

## The workers and the primary




# Narwhal

## The workers and the primary




# Narwhal

## The primary machine




# Narwhal

## The primary machine




# Narwhal

## The primary machine




# Narwhal

## The primary machine




# Narwhal

## The primary machine




# Narwhal

## The primary machine



# Second Step to the Solution



All You Need is DAG—PODC 21'

Narwhal and Tusk: A DAG-based Mempool and Efficient BFT Consensus—Eurosyst 22'

Bullshark: Dag bft protocols made practical — CCS 22'

Hammerhead: Leader reputation for dynamic scheduling — ICDCS 24'

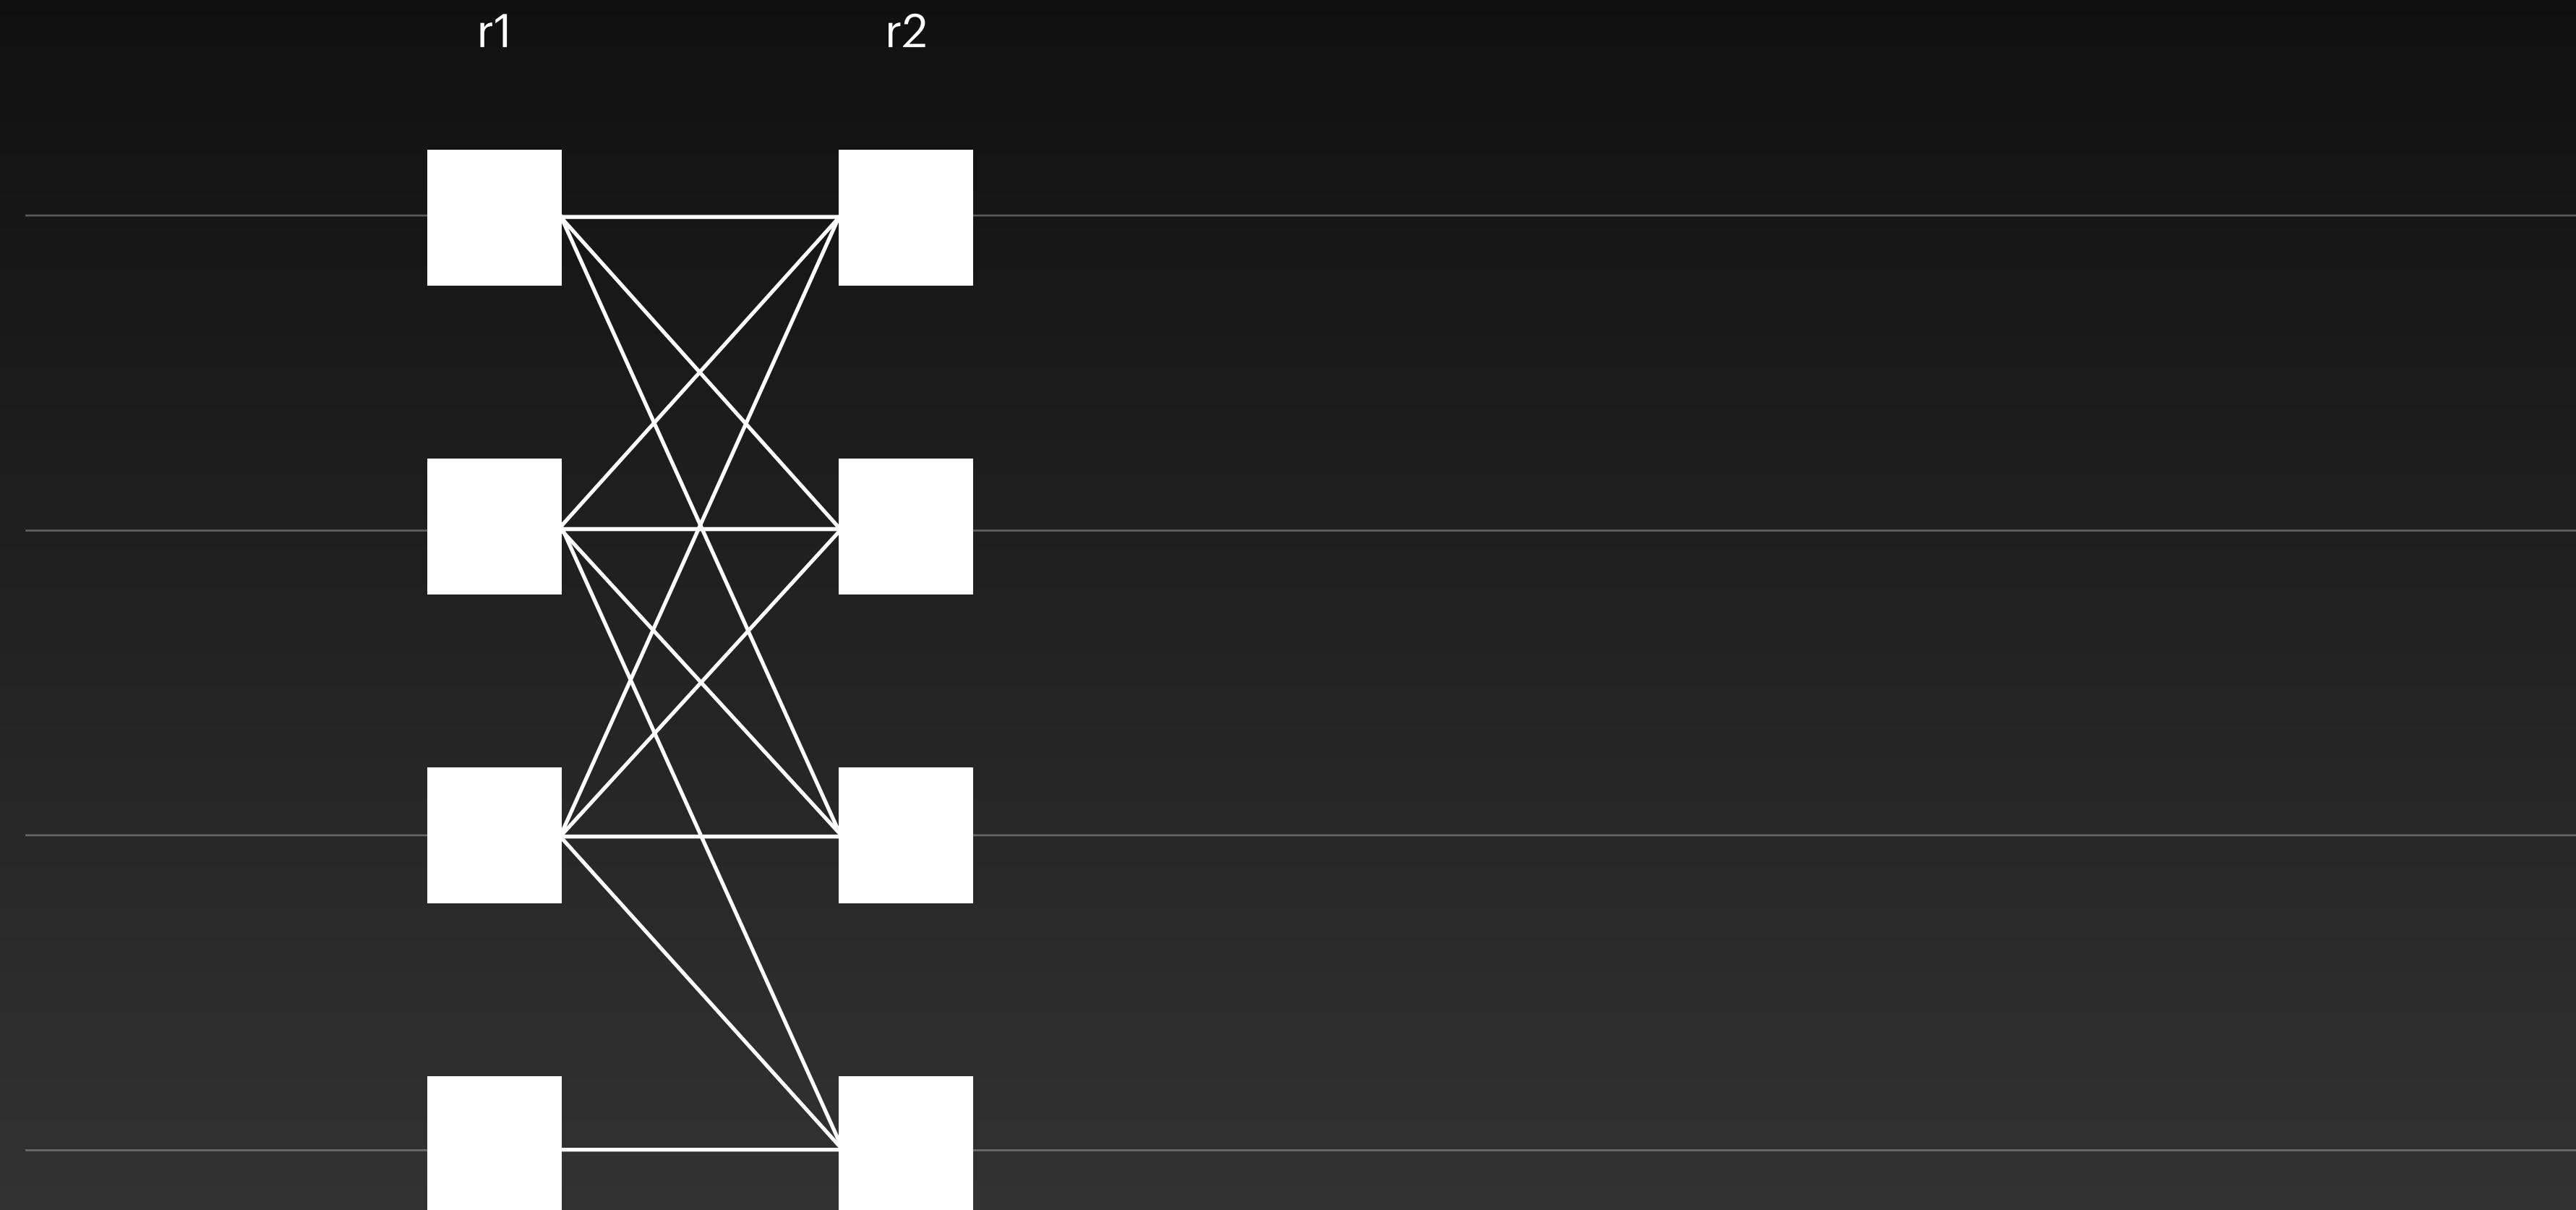
# Bullshark

Zero-message partially-synchronous consensus

\* without asynchronous fallback

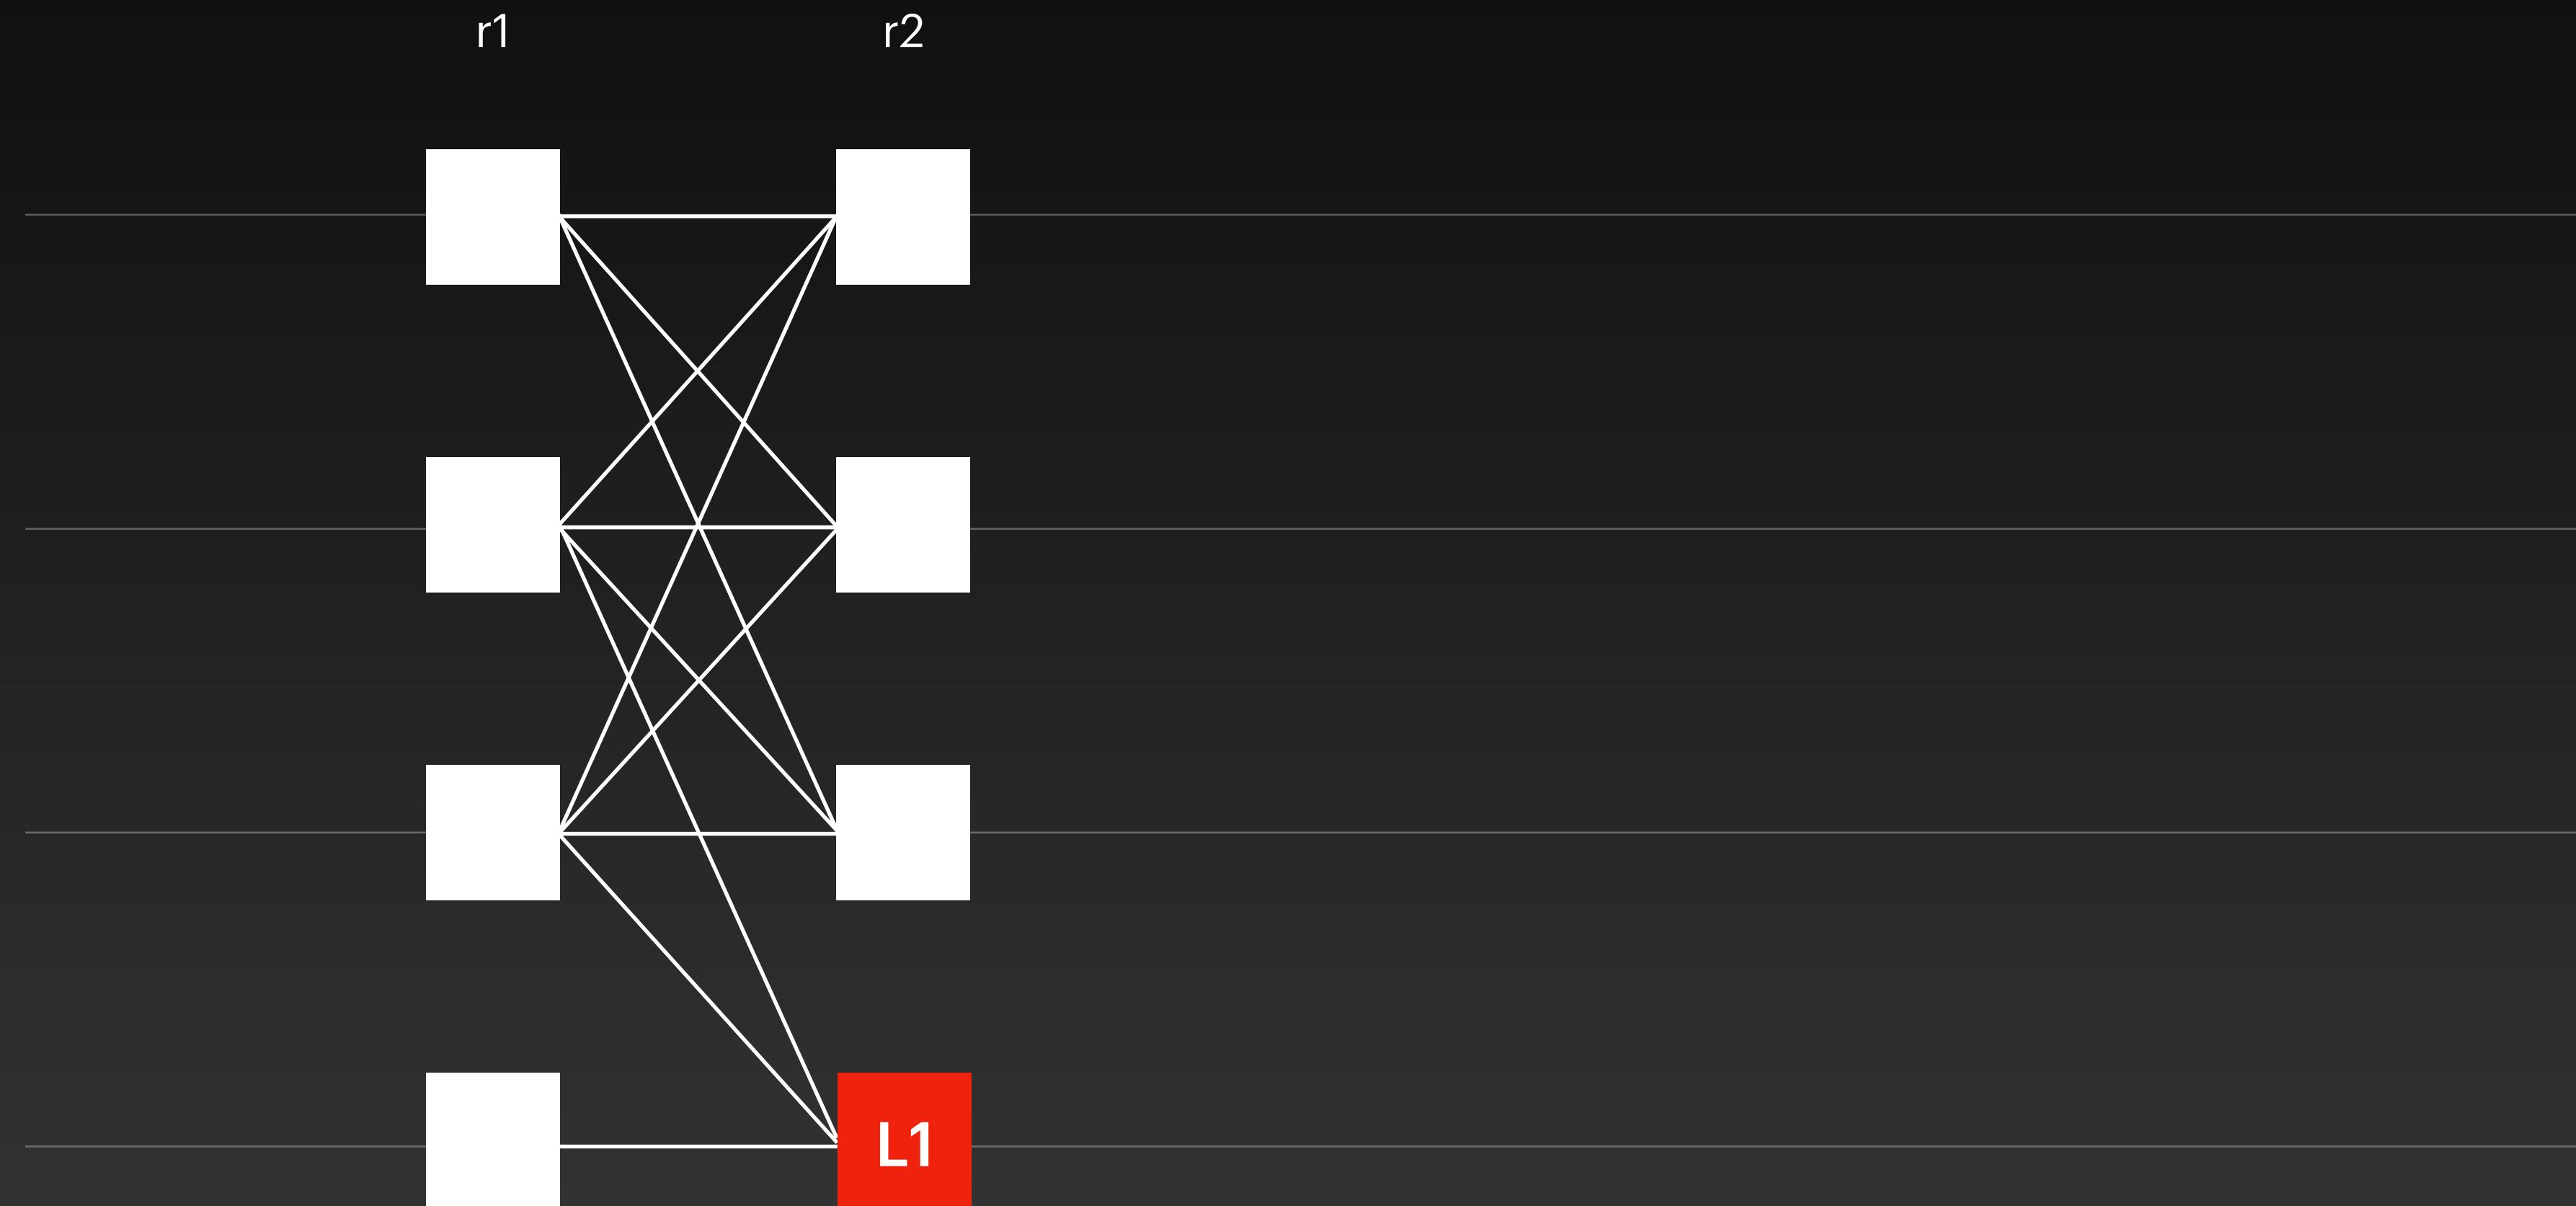
# Bullshark

## Just interpret the DAG



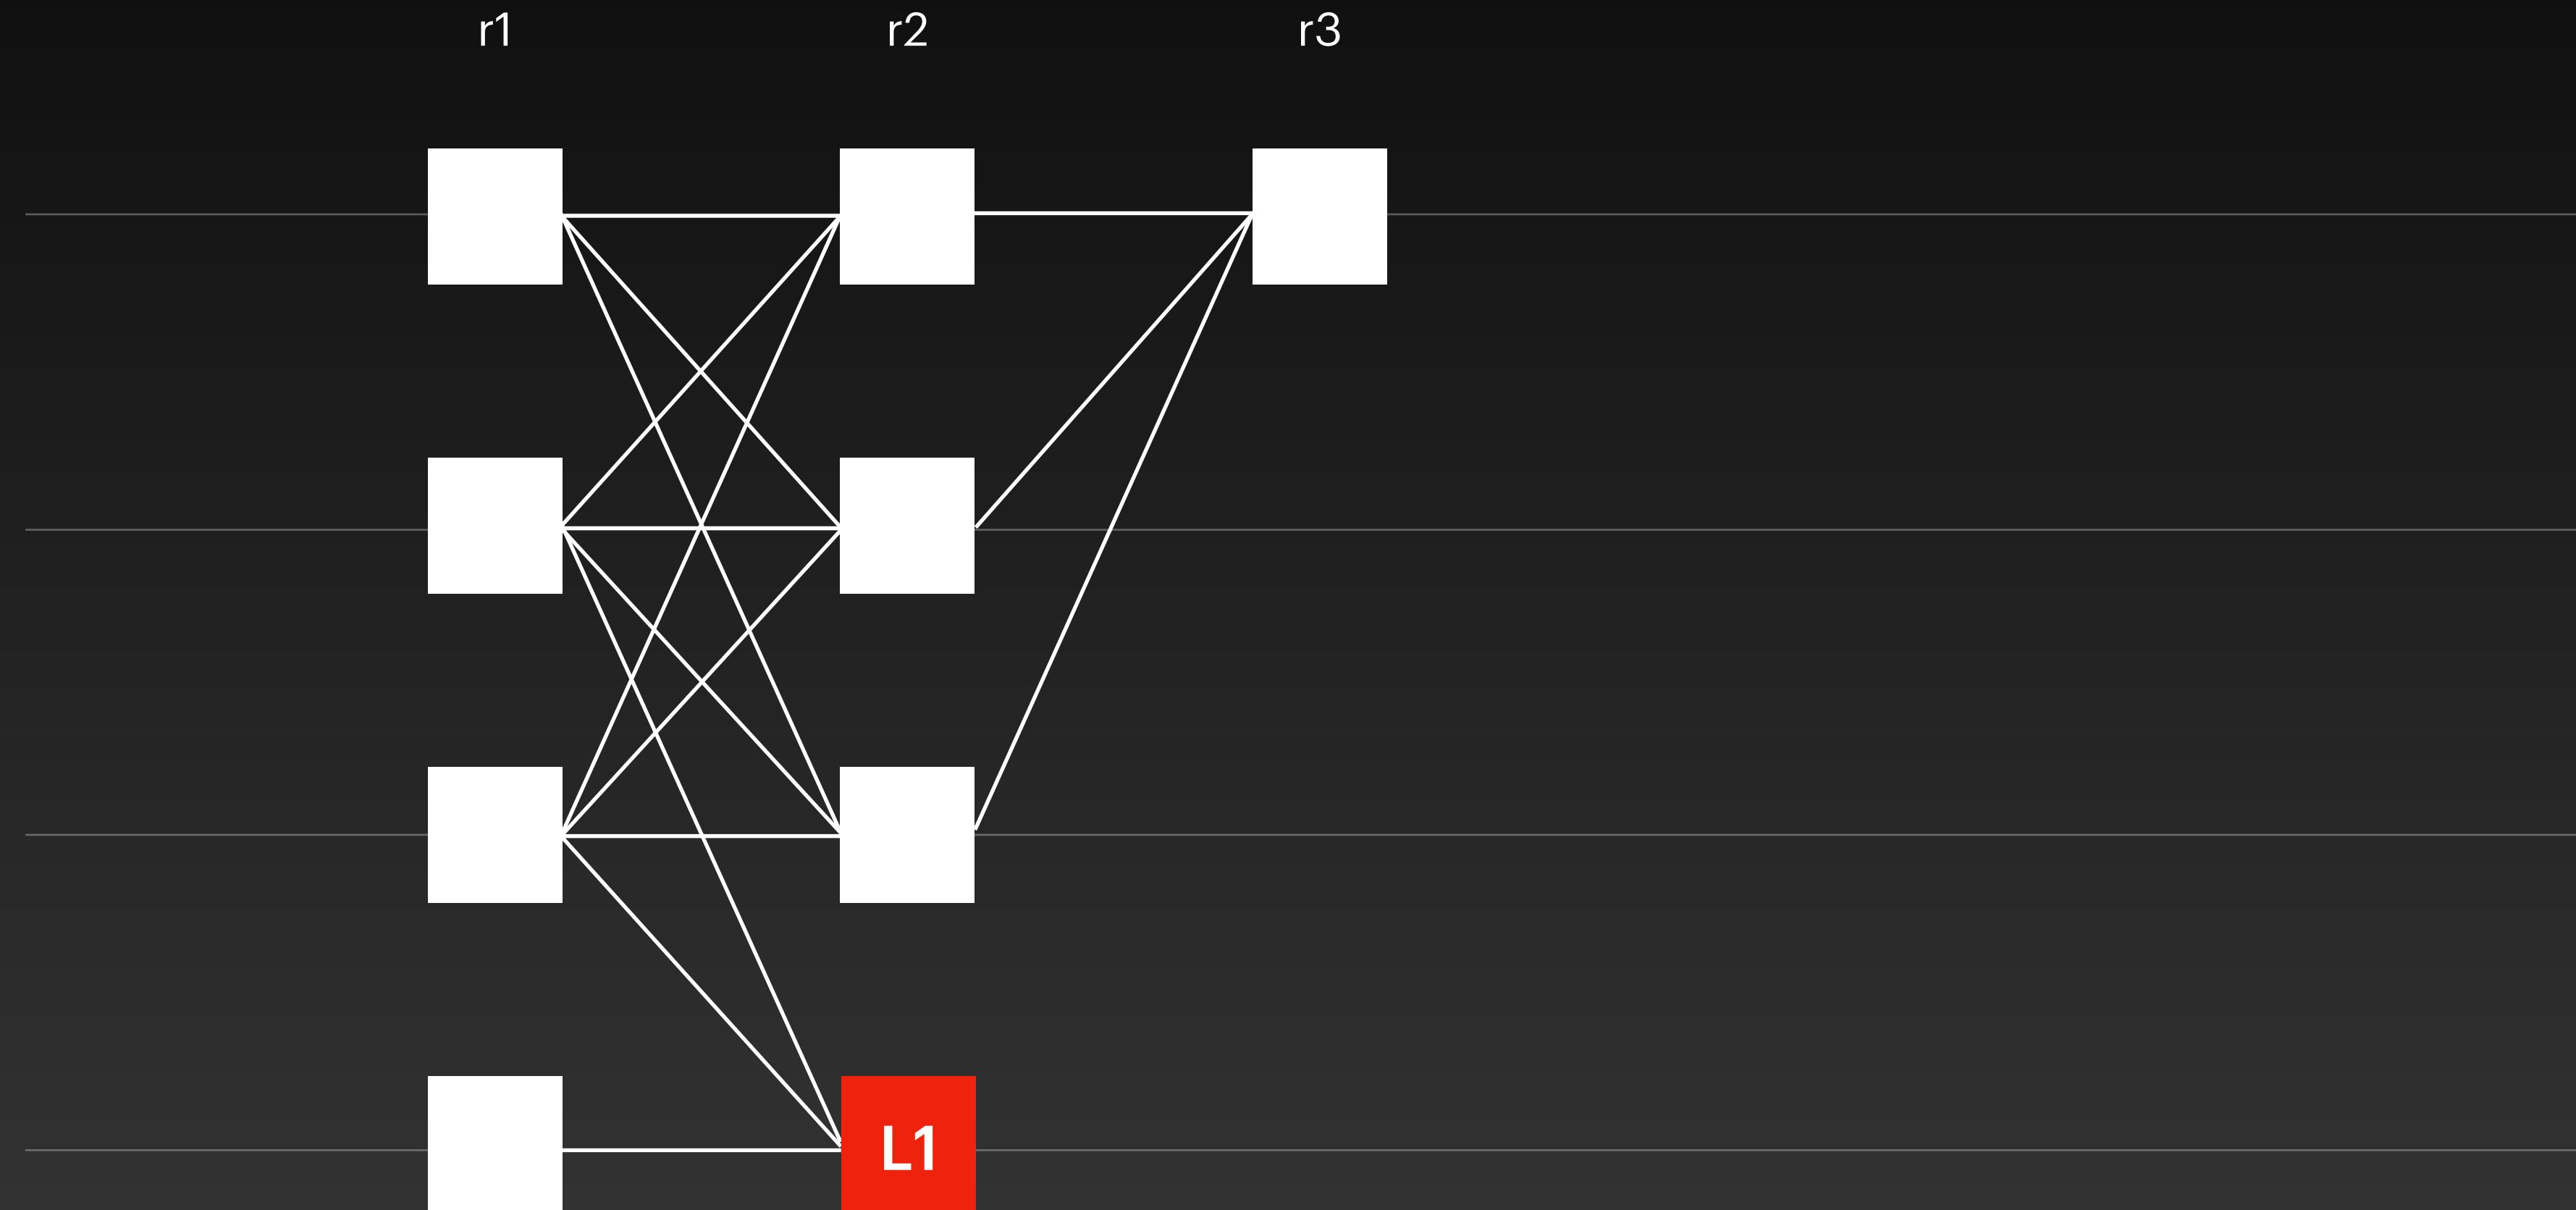
# Bullshark

## Deterministic leader every 2 rounds



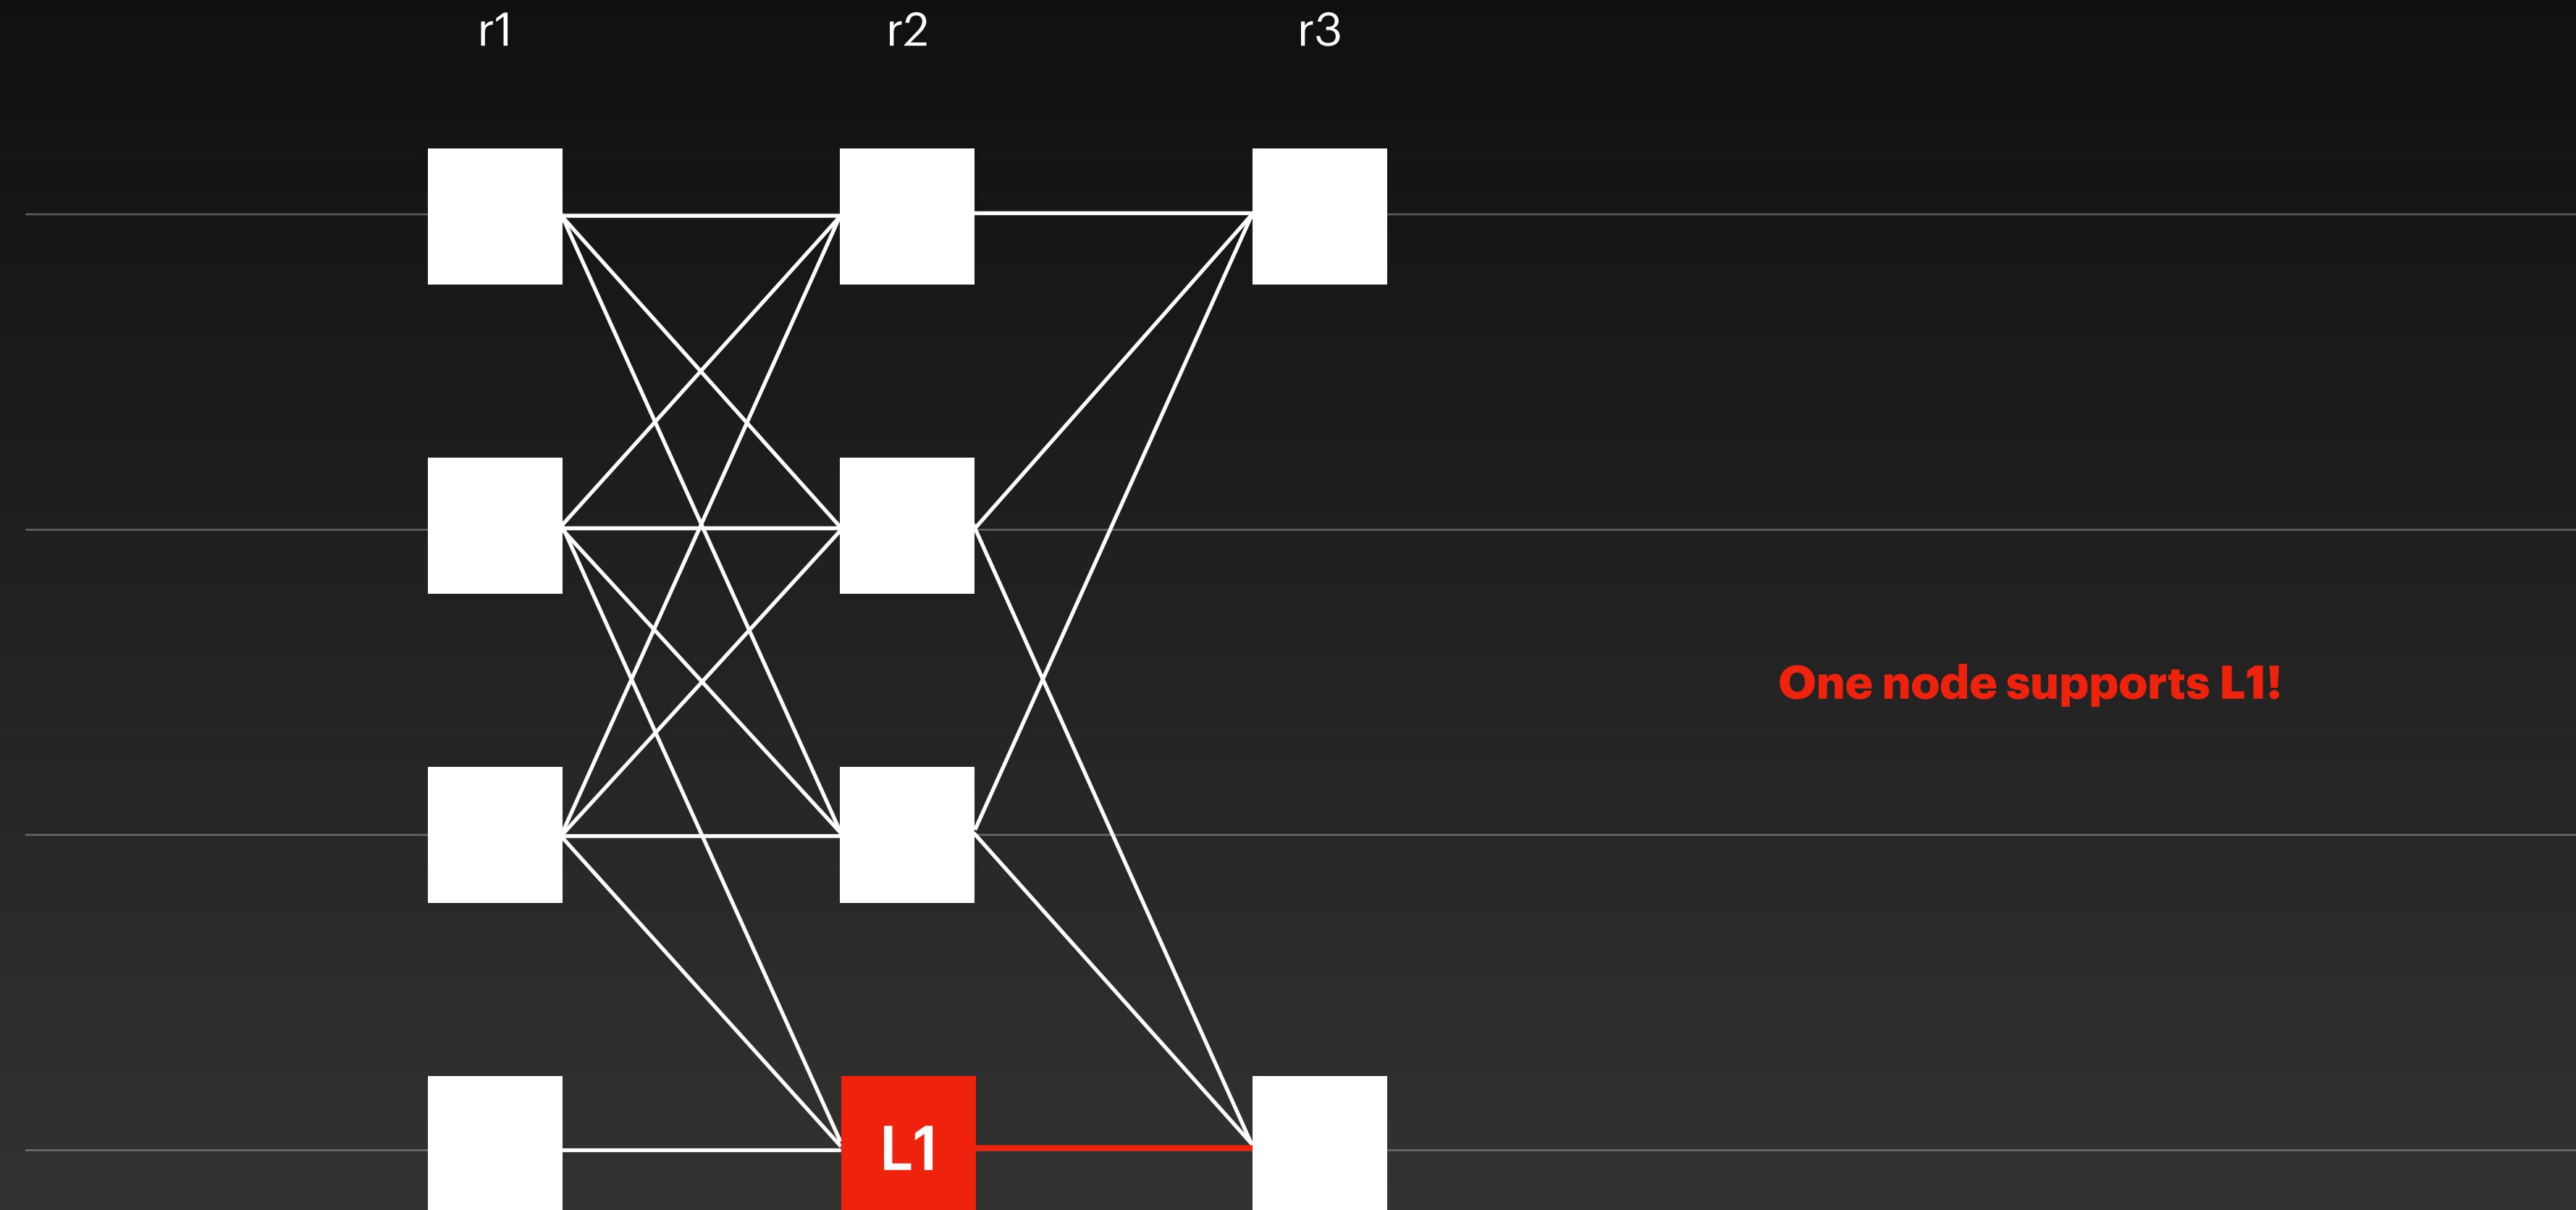
# Bullshark

## The leader needs $f+1$ links from round $r$



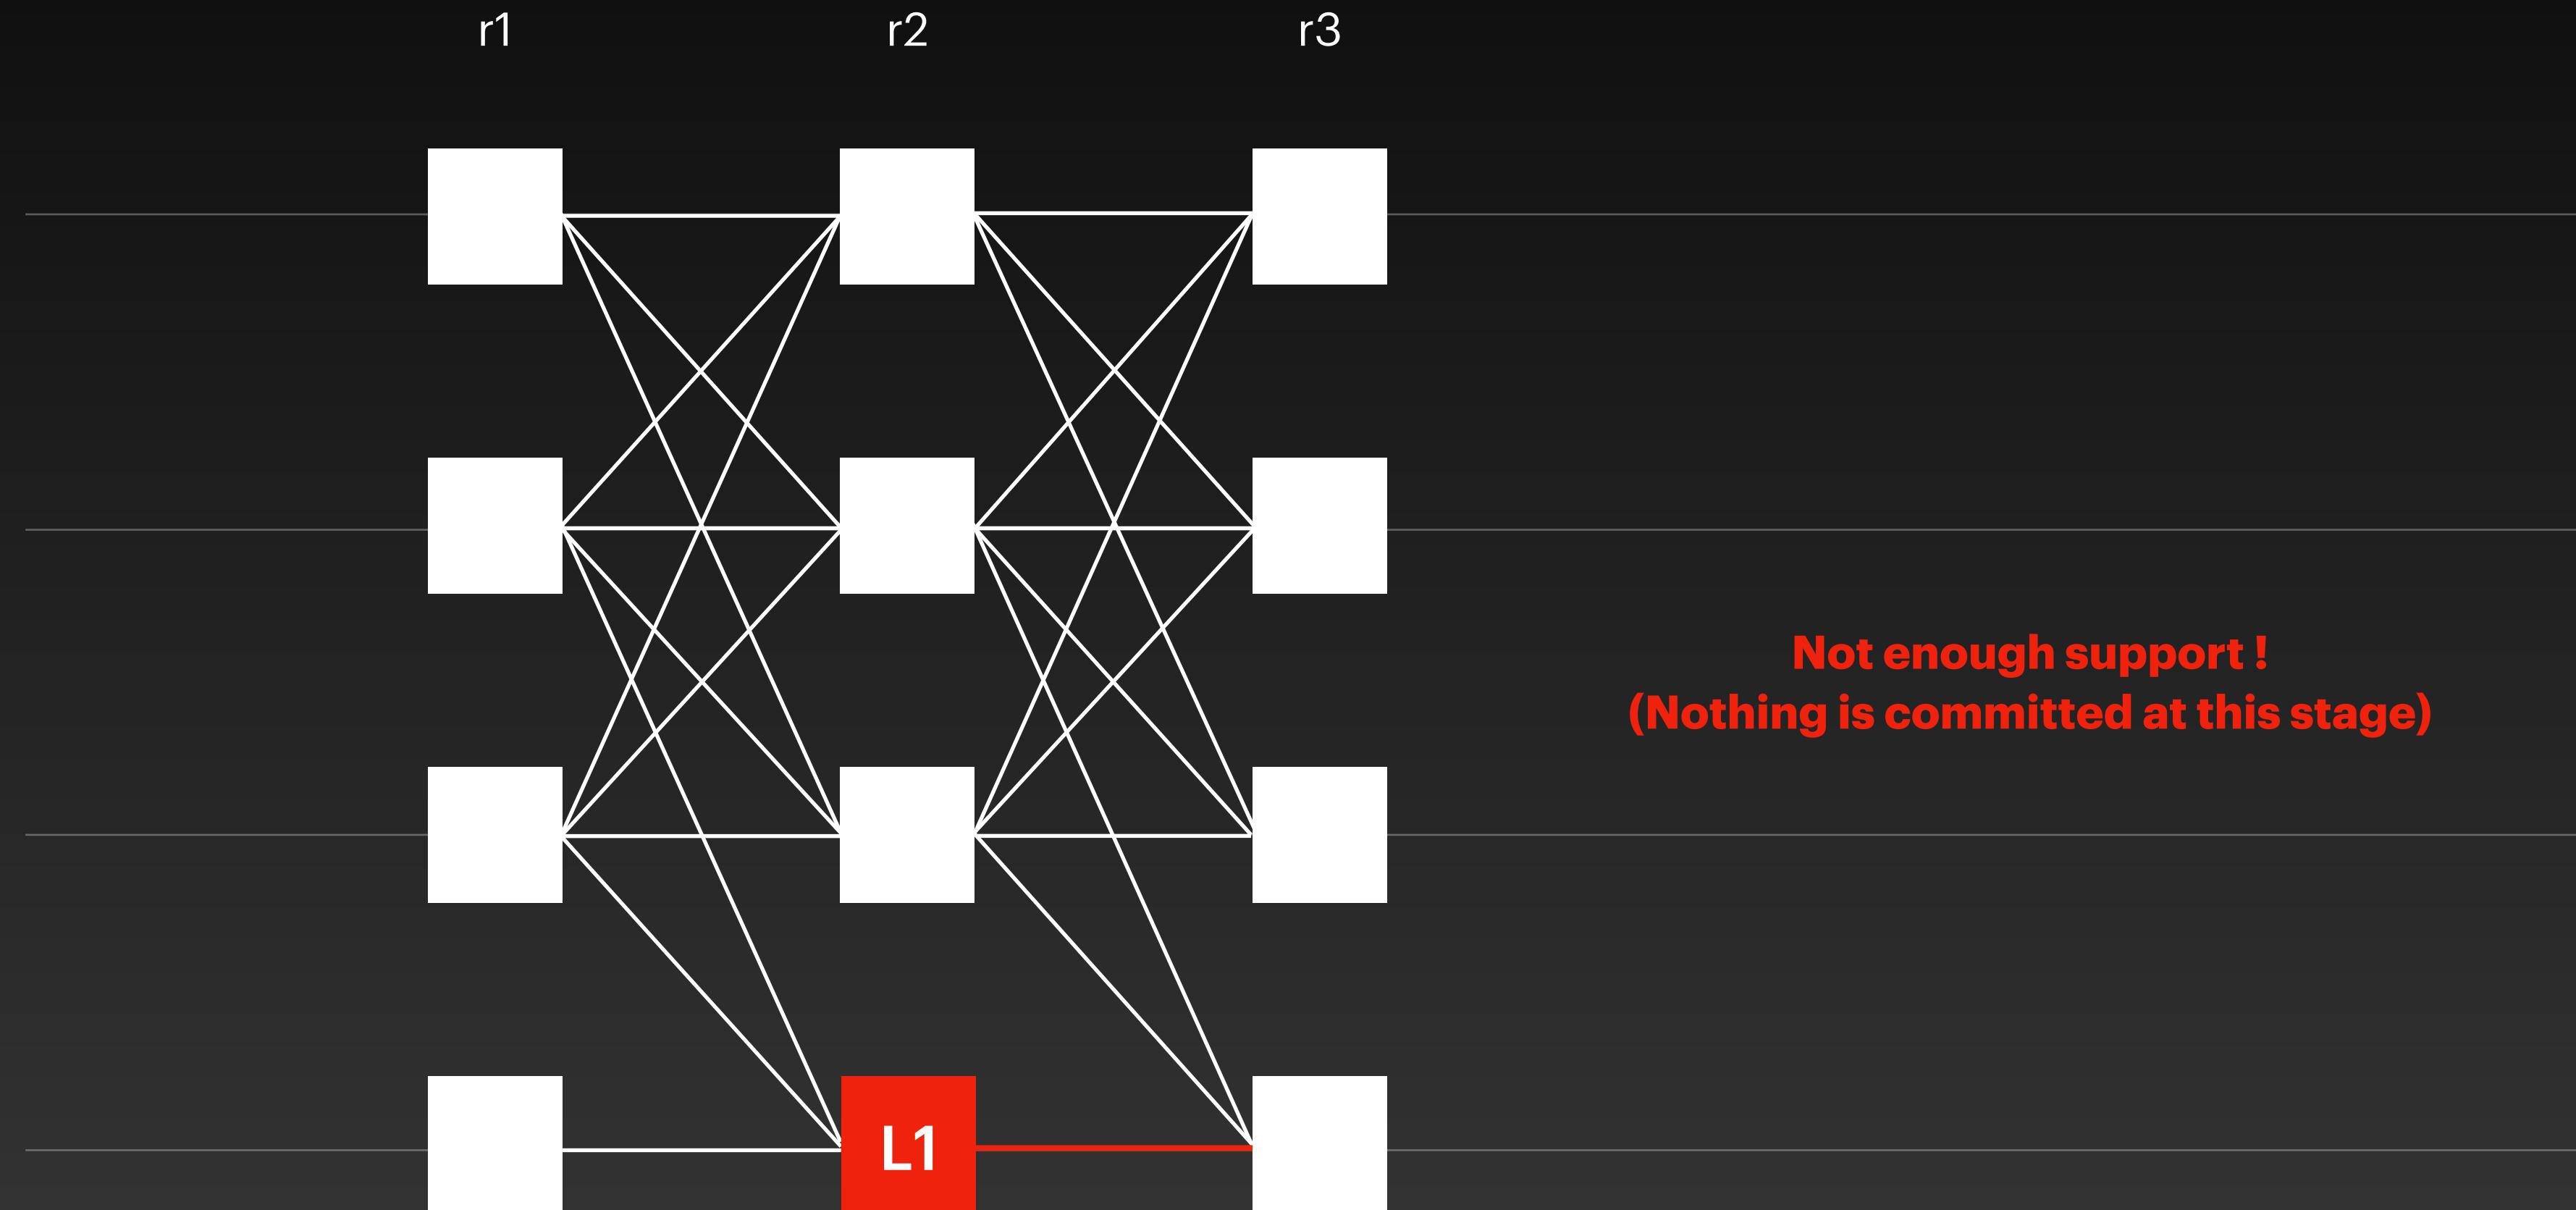
# Bullshark

## The leader needs $f+1$ links from round $r$



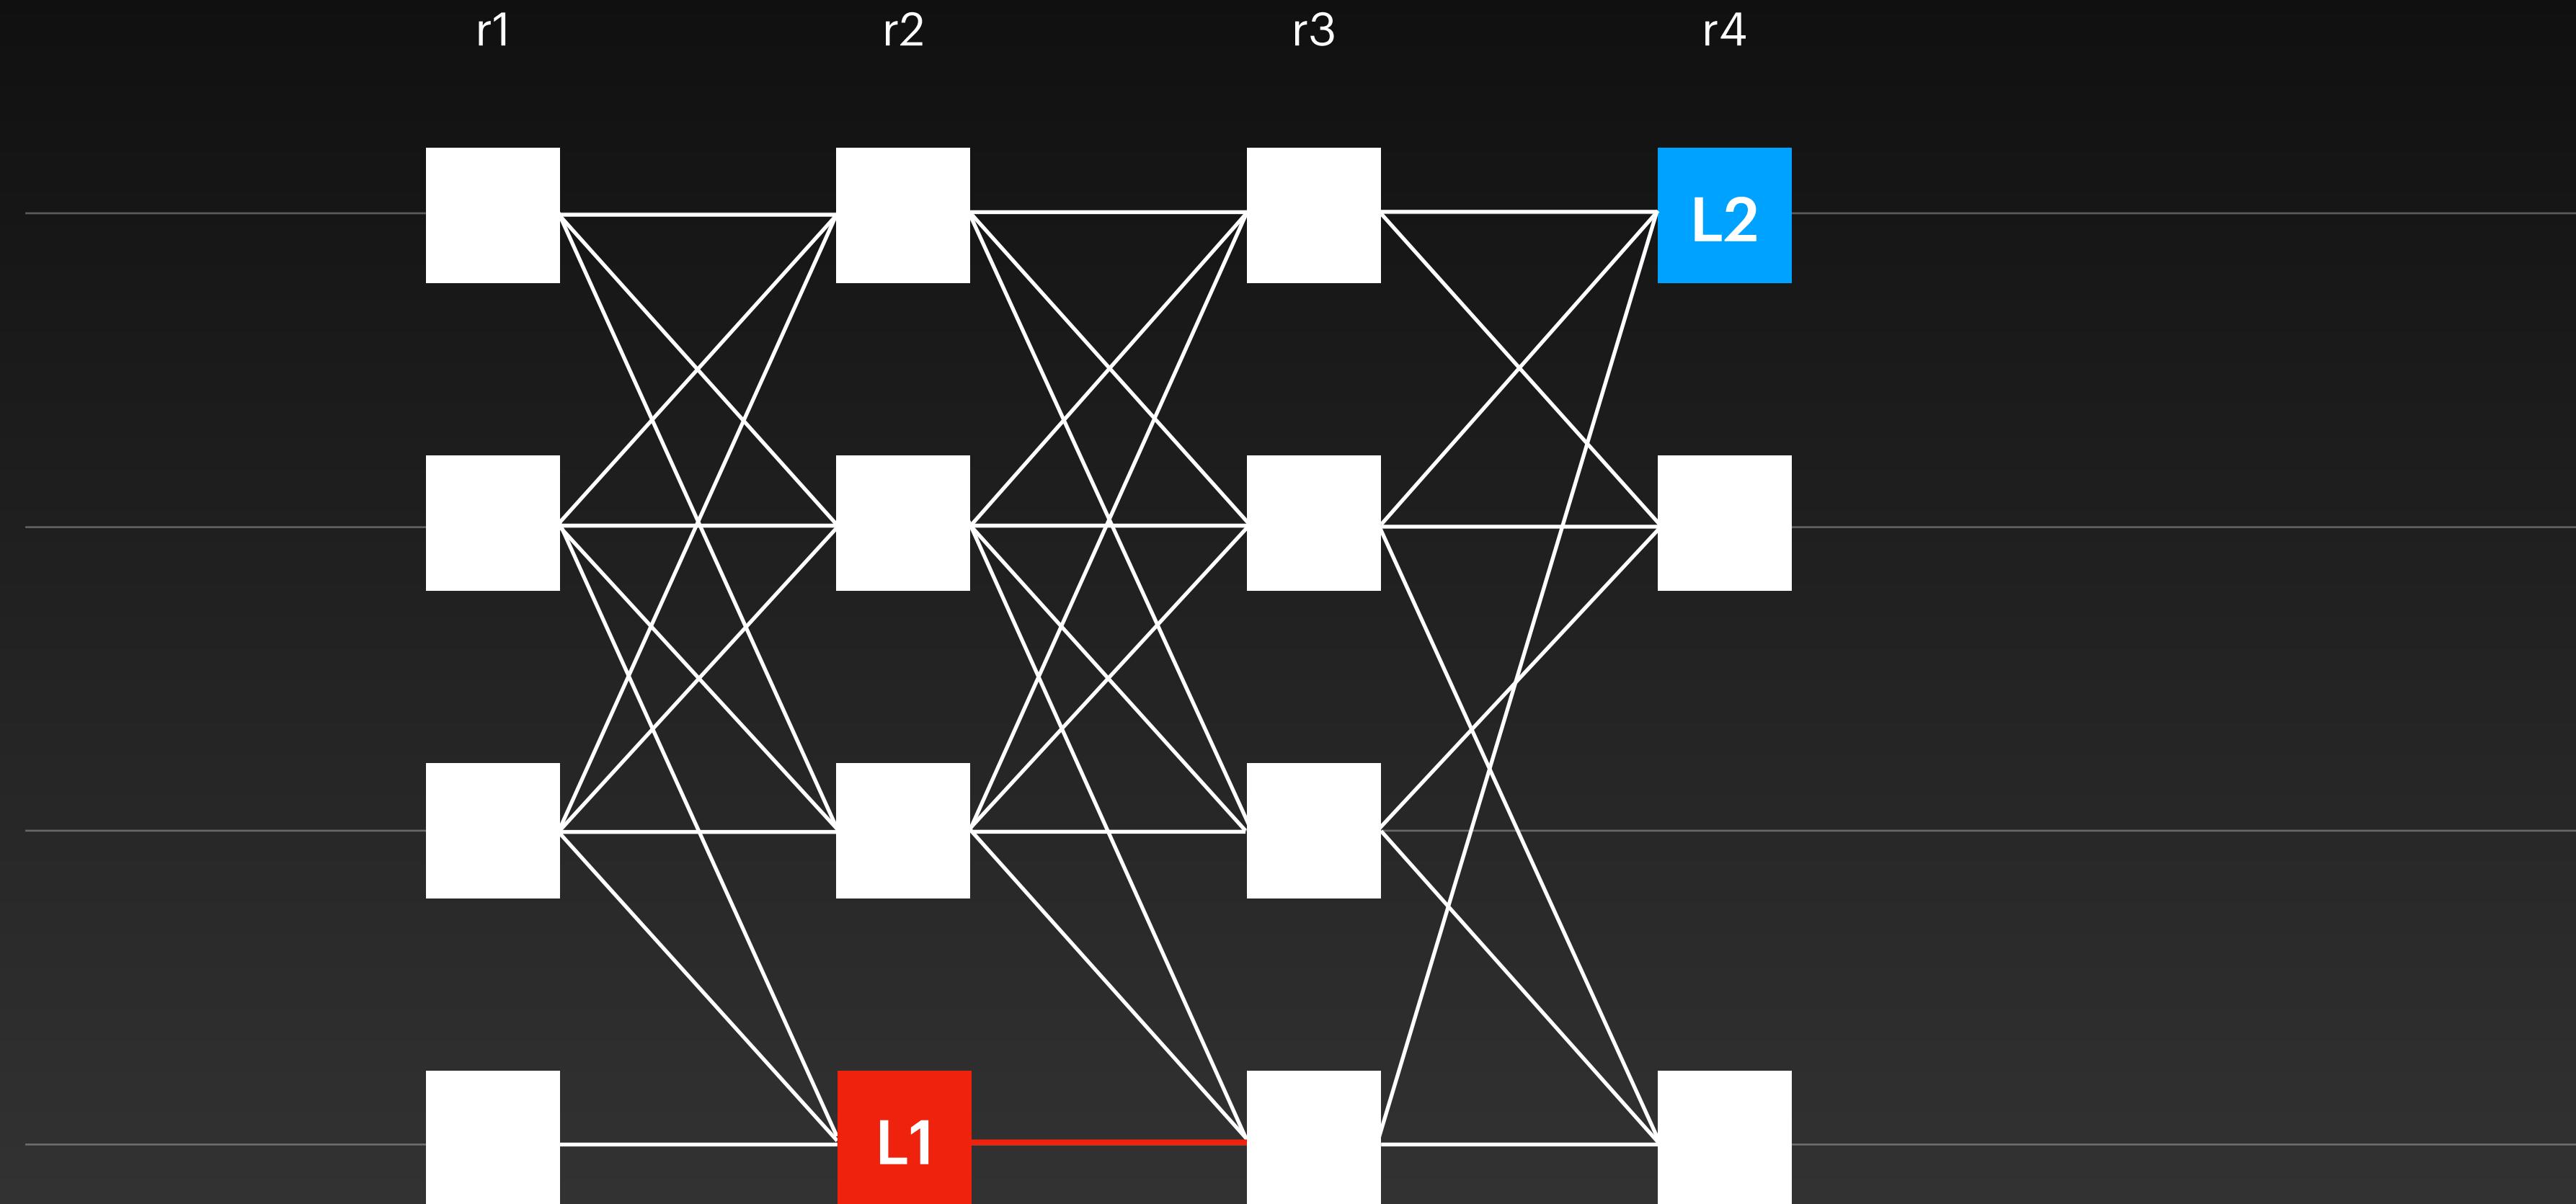
# Bullshark

## The leader needs $f+1$ links from round $r$



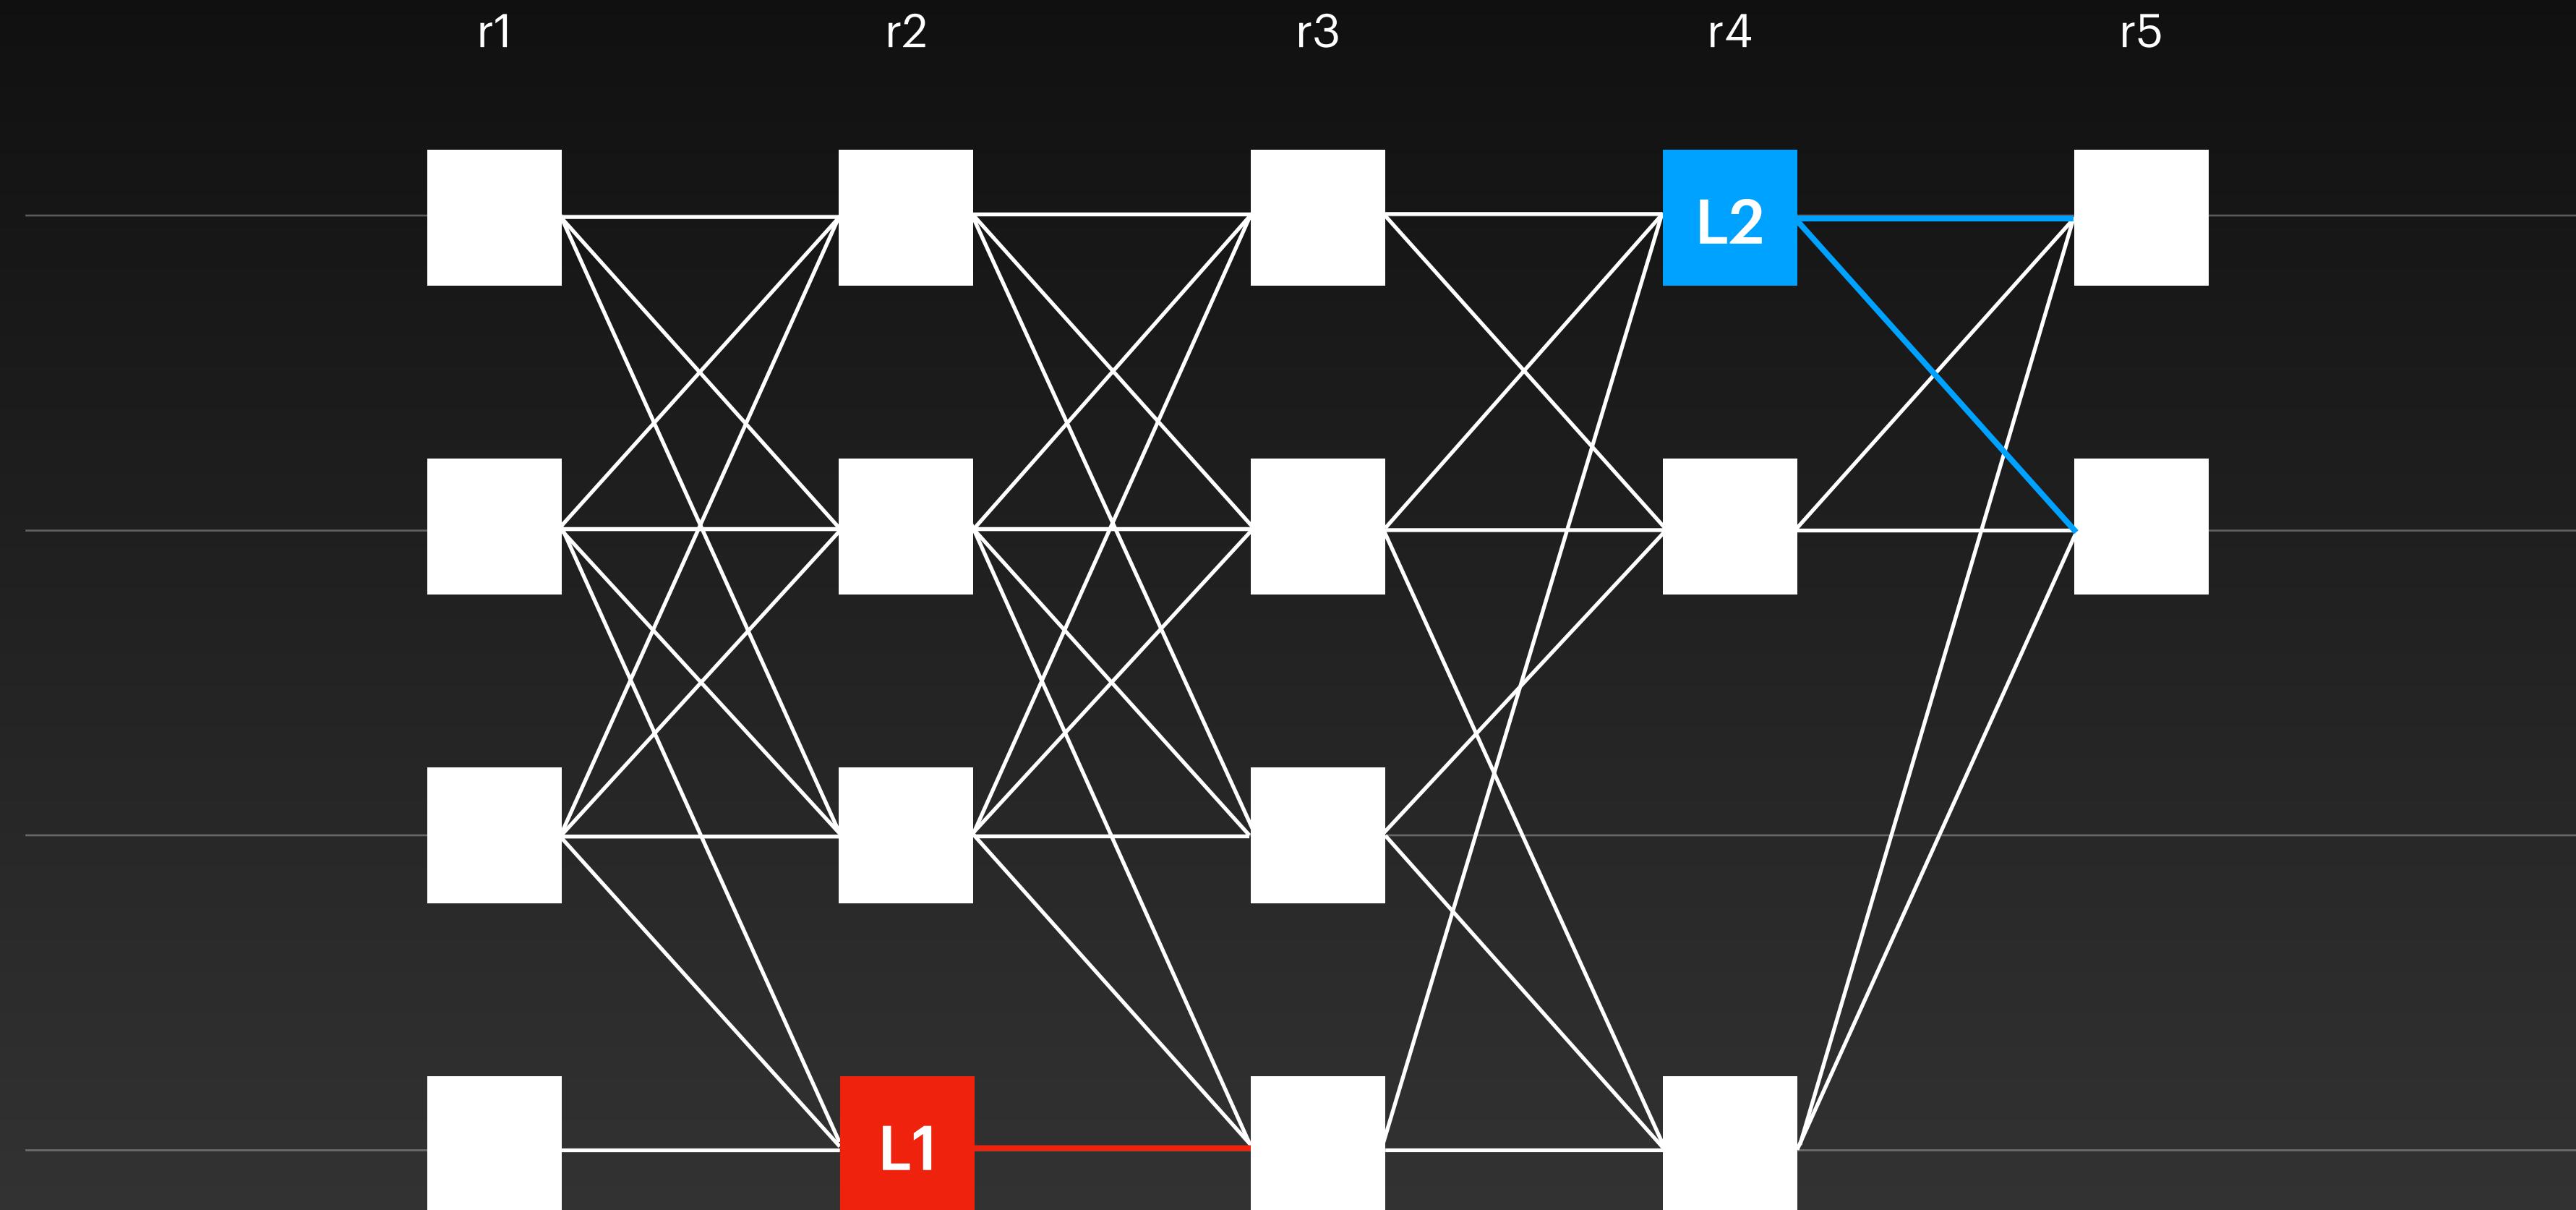
# Bullshark

## Elect the leader of r4




# Bullshark

## Leader L2 has enough support



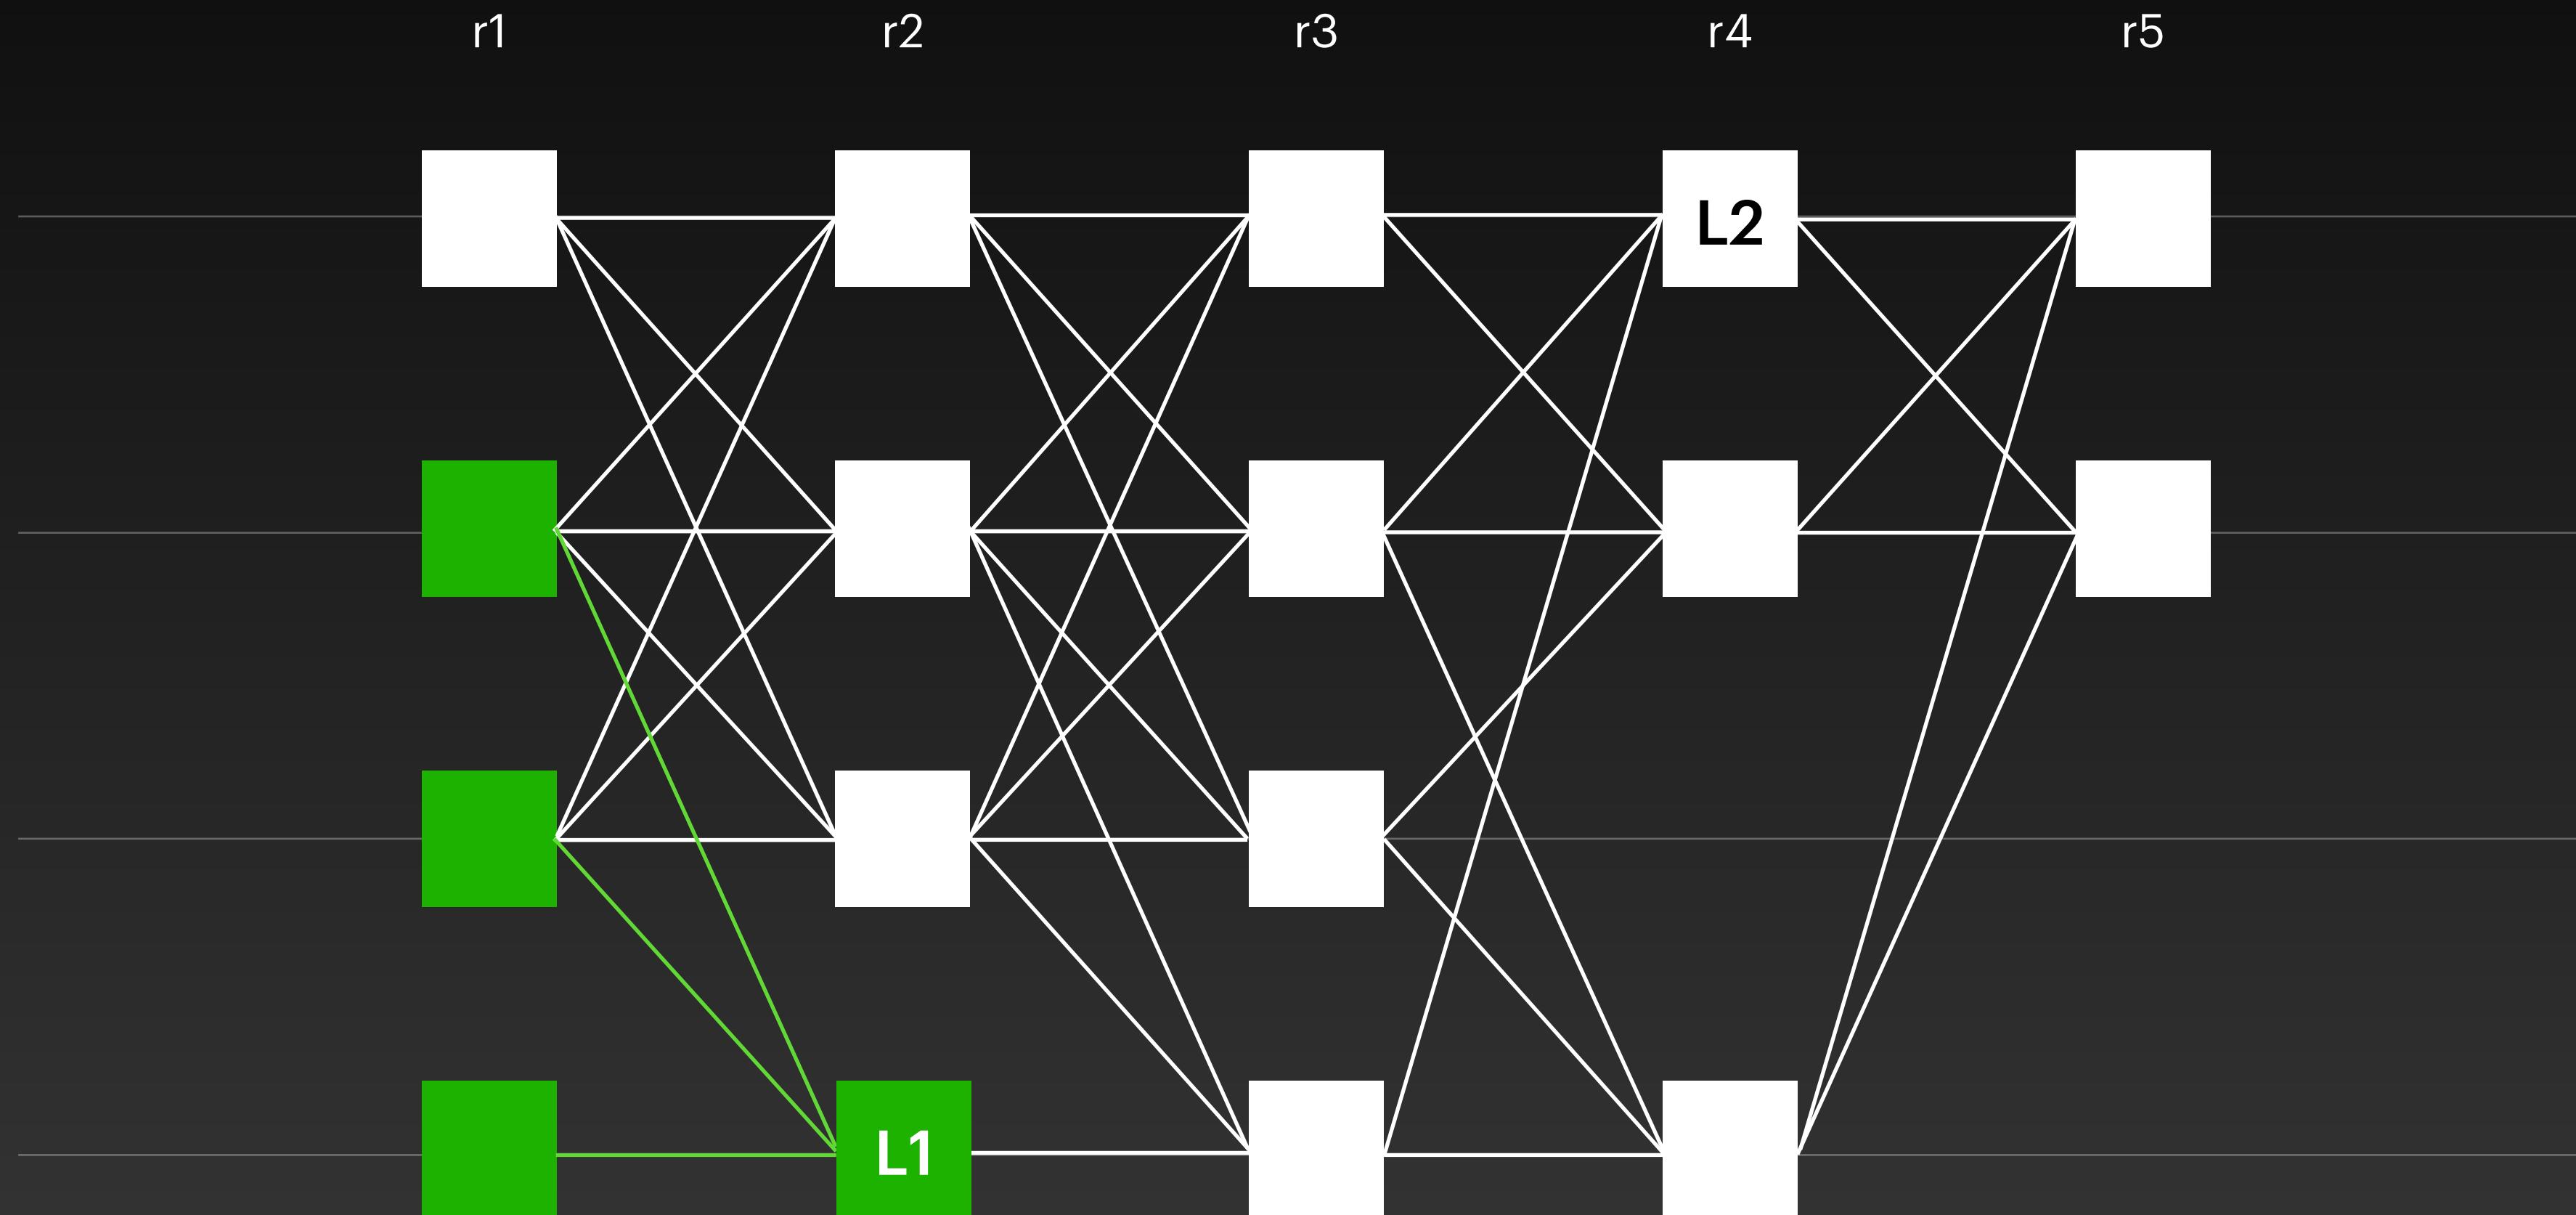
# Bullshark

## Leader L2 has links to leader L1



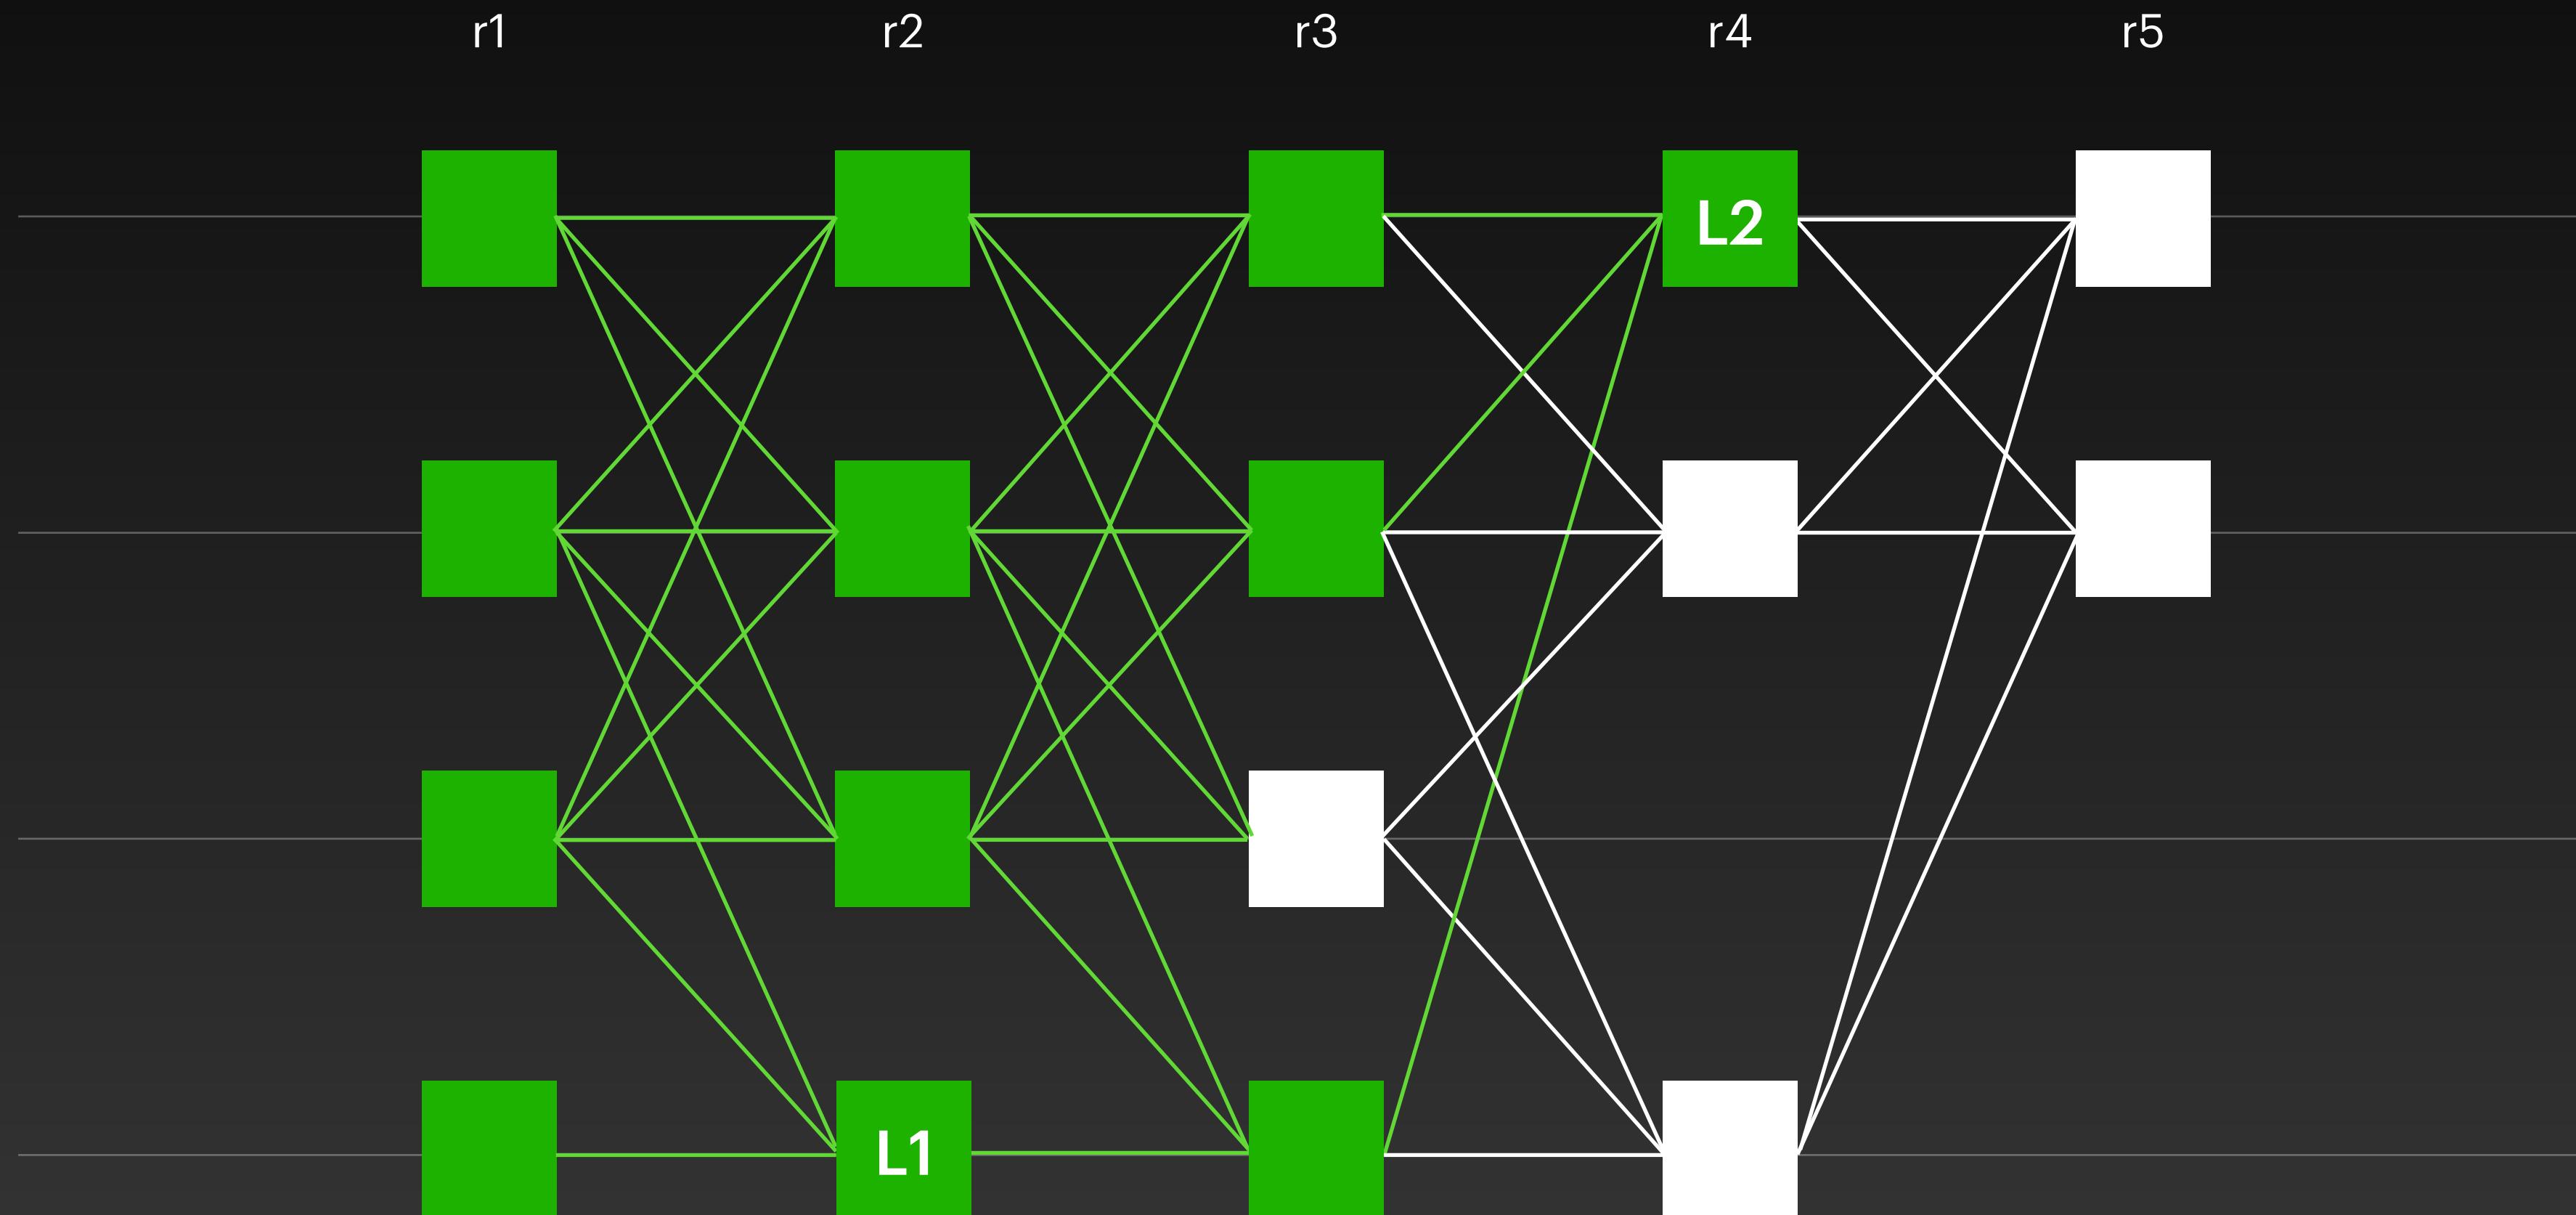
# Bullshark

## Commit all the sub-DAG of the leader




# Bullshark

## Commit all the sub-DAG of the leader




# Evaluation

## Experimental setup on AWS




# Evaluation

## Throughput latency graph



# Evaluation

## Performance under faults




# Summary

## Bullshark

- Zero-message overhead, no view-change, no common-coin
- Disseminate data with Narwhal, exploits periods of synchrony

# Are we done? Latency?




# The Mysticeti DAG

Uncertified DAG

# The Mysticeti DAG


## Block Creation



- Round number
- Author
- Payload (transactions)
- Signature

# The Mysticeti DAG


## Rule 1: Link to $2f+1$ parents



- Total nodes:  **$3f+1 = 4$**
- Quorum:  **$2f+1 = 3$**


# The Mysticeti DAG

## Rule 2: Every node waits and links to leaders




# The Mysticeti DAG

## Rule 3: All node run in parallel




# DAG Structure



# Interpreting DAG Patterns

Reminder



# Direct Decision Rule


On each leader starting from highest round:

- **Skip** if  $2f+1$  blames
- **Commit** if  $2f+1$  certificates
- **Undecided** otherwise

# Direct Decision Rule

On each leader starting from highest round:


- **Skip** if  $2f+1$  blames
- **Commit** if  $2f+1$  certificates
- **Undecided** otherwise



# Direct Decision Rule

On each leader starting from highest round:

- **Skip** if  $2f+1$  blames
- **Commit** if  $2f+1$  certificates
- **Undecided** otherwise



# Why? Crash Faults

In a year of running Sui:

- How many Byzantine faults? 0

# Why? Crash Faults

In a year of running Sui:

- How many Byzantine faults? 0
- How many Crash faults?

# Why? Crash Faults


In a year of running Sui:

- How many Byzantine faults? 0
- How many Crash faults? 😭


# Resources

- **Paper:** <https://arxiv.org/pdf/2310.14821>
- **Presentation:** <https://www.youtube.com/watch?v=JhhCxyZylx8>

# Evaluation




# Last Step to the Solution




# Sharding Over DAGs—Design

"Executing and proving over dirty ledgers." FC, 2023.



# Sharding Over DAGs – Properties


"Executing and proving over dirty ledgers." FC, 2023.




- + 51% security threshold per Shard
- + Scales-Out
- + Low Cost per Execution Node
- Fragmenting the state-space – Expensive Atomic Commit
- Susceptible to adaptive adversaries

# Pilotfish


## Distributed Transaction Execution for Lazy Blockchains



# Sequence Worker (SW)



# Execution Workers (EWs)



# Why is this Safe in Concurrency

- The ordering of dependencies is predefined from the consensus output
- Every EW knows the version of the objects they are supposed to read/write and back pressure the SW when it is not available yet

# Pilotfish → Elastic Scaling for Web3

- + Cost scales with load, but so does profit
- + Scales-Out
- + Flat state-space
- + Consistent Threat Model

# Evaluation




Figure 9: Pilotfish latency vs. throughput with simple transfers.





Figure 10: Pilotfish scalability with simple transfers.

# Evaluation



- Pilotfish over Bullshark provides the first ***end-to-end Elastic Distributed Ledger***
- Pilotfish does not employ batching → *Latencies of 20-50ms post-consensus*
- Pilotfish is co-designed with the blockchain → *Light worker recovery*



# Side-Stepping Consensus

## Consensus is not required

Coins, balances, and transfers

NFTs creation and transfers

Game logic allowing users to combine assets

Inventory management for games / metaverse

Auditable 3rd party services not trusted for safety

...

# New Architecture

## The Sui System

Consensus only when  
you need to

# New Architecture

## Architecture

### Owned Objects


- Objects that can be mutated by a single entity
- e.g., My bank account
- **Do not need consensus**

### Shared Objects

- Objects that can be mutated my multiple entities
- e.g., A global counter
- **Need consensus**


# The Sui System

## Architecture



# The Sui System

## Consensus-less Path



### **Send T1:**

Disseminate the transaction

### **Echo T1:**

Nodes check and sign  $T_1$

### **Cert T1:**


User gather  $>2/3$  signatures into a certificate and disseminate it

### **Effect T1:**

User gather  $>2/3$  effect signatures for finality

# The Sui System

## Consensus-less Path



+ Low Latency

+ Trivial to Scale Out

- Reconfiguration

- Equivocation results in loss of liveness

# Side-Stepping Consensus

## Safe reconfiguration

# Side-Stepping Consensus

## Equivocation Tolerance